

Introduction & 8 Important Problems in Modern Operating Systems

Yubin Xia

Institute of Parallel and Distributed Systems
(IPADS)

Shanghai Jiao Tong University

http://ipads.se.sjtu.edu.cn/Yubin_xia

Outline

Why Choose This Course?

Course information

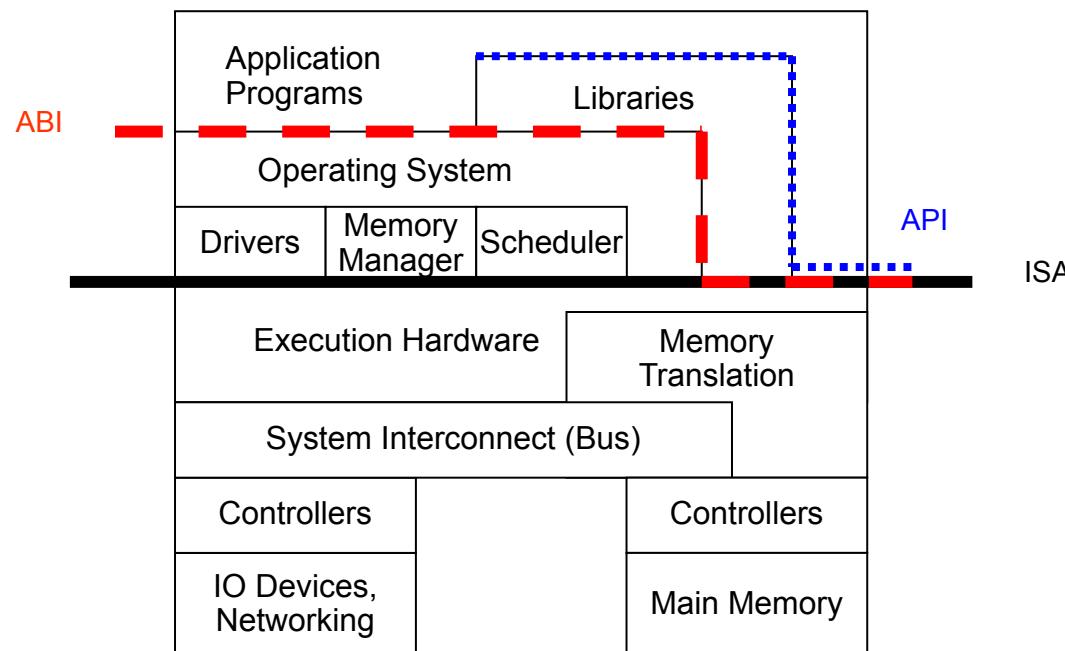
8 important problems of OS

Recall

- What is OS from a programmer's perspective?
 - A set of API, aka, system calls
 - E.g., int 0x80
- What are differences from other API?
 - The system calls are usually hardware related
 - E.g., fork, FILE ops, socket, select/epoll
- Who invokes the system calls?
 - Application itself, libraries (e.g., glibc)
 - Application through the libraries

Software Stack

- Architecture
 - Functionality and Appearance of a computer system but not implementation details
 - Level of Abstraction = Implementation layer
 - ISA, ABI, API

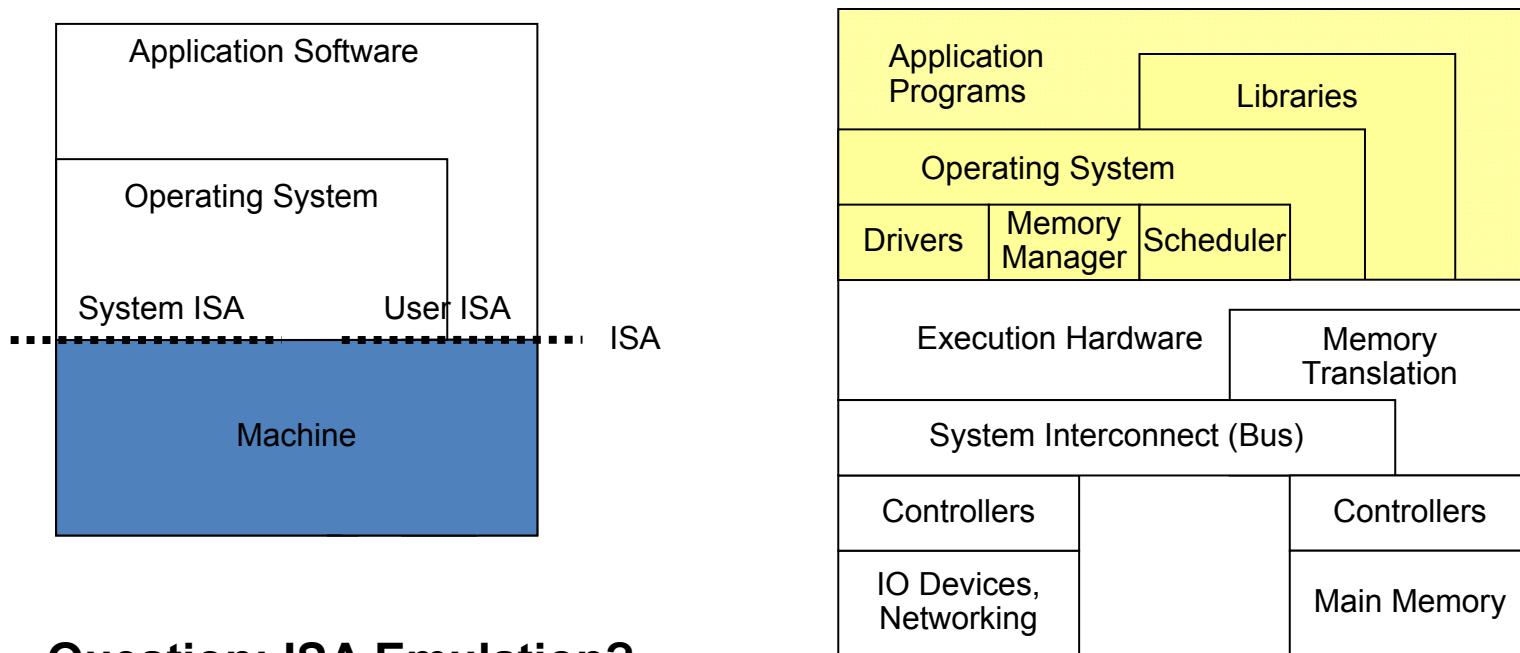


Architecture, Implementation Layers

- Implementation Layer : ISA
 - Instruction Set Architecture
 - Divides hardware and software
 - Concept of ISA originates from IBM 360
 - IBM System/360 Model (20, 40, 30, 50, 60, 62, 70, 92, 44, 57, 65, 67, 75, 91, 25, 85, 95, 195, 22) : 1964~1971
 - Various prices, processing power, processing unit, devices
 - But guarantee a *software compatibility*
 - User ISA and System ISA

Implement the ISA

- Machine from the perspective of a system
 - ISA provides interface between system and machine



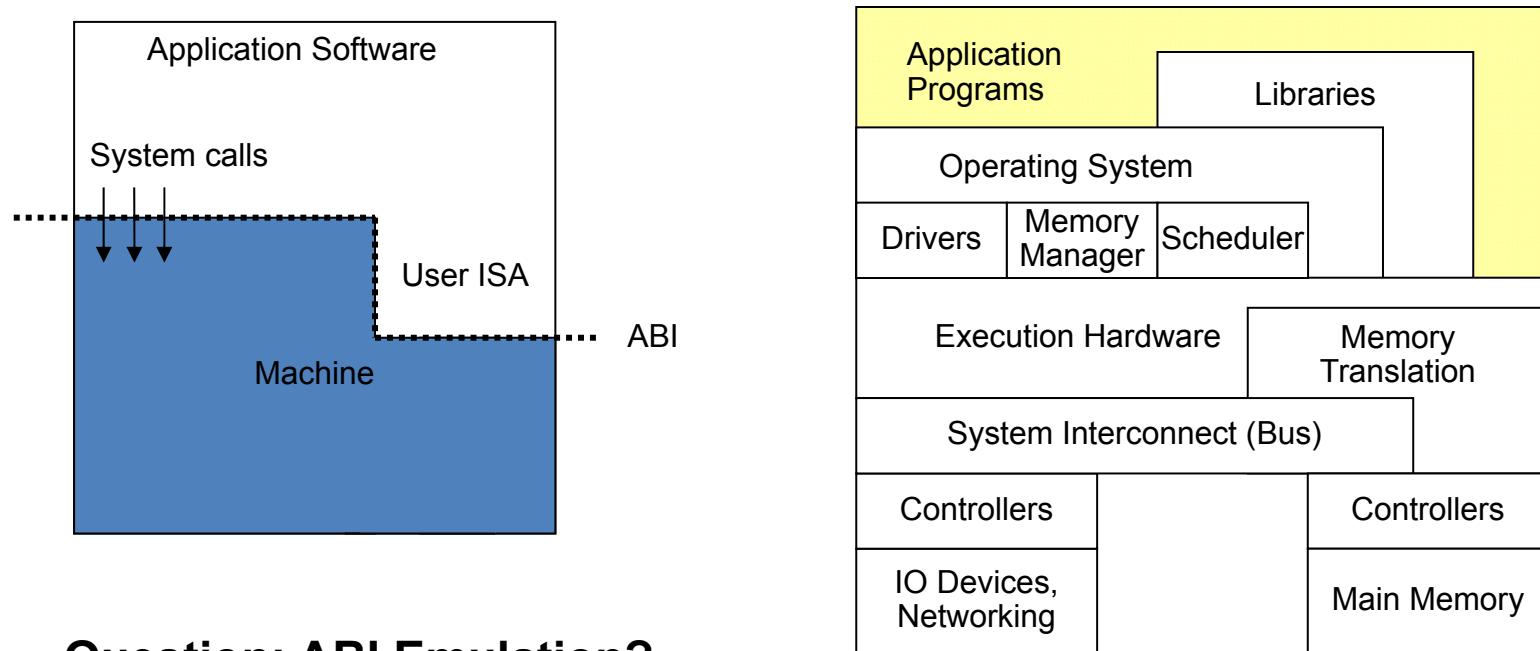
Question: ISA Emulation?

Architecture, Implementation Layers

- Implementation Layer : ABI
 - Application Binary Interface
 - Provides a program with access to the hardware resource and services available in a system
 - Consists of User ISA and System Call Interfaces

Implement the ABI

- Machine from the perspective of a process
 - ABI provides interface between process and machine



Question: ABI Emulation?

Architecture, Implementation Layers

- Implementation Layer : API
 - Application Programming Interface
 - Key element is Standard Library (or Libraries)
 - Typically defined at the source code level of High Level Language
 - **c lib** in Unix environment : supports the UNIX/C programming language

Why Study OS?

Operating systems are a maturing field

Hard to get people to switch operating systems

Then why?

Many things are OS issues

High-performance

Resource consumption

Battery life, radio spectrum

Security needs a solid foundation

New “smart” devices need new OSes

Q: What applications are highly OS-related, which are not?

Why Study OS?

OS is the core of the broad “systems area”

Lots of systems companies

Microsoft, Google, IBM, EMC, CISCO, VMware

What's Google's core?

Google Cluster, GFS, MapReduce, BigTable

Good OS knowledge is a necessity to get you good offers/promotions

at such companies and top universities

Real Machine VS. Abstraction Machine

Many courses emphasize abstraction

- Abstract data types

- Asymptotic analysis

These abstractions have limits

- Especially in the presence of bugs

- Need to understand details of underlying implementations

Useful outcomes from taking from this course

- Become more effective programmers**

- Able to find and eliminate bugs efficiently

- Able to understand and tune for program performance

- Understand design and implementation of real complex systems

- Try to build one on your own (like a turing-award level OS)

Is “simple” really simple?

A simple program?

```
#include <stdio.h>
int main() {
    printf("hello world\n");
    return 0;
}
```

How does it get executed?

How does it get translated
to machine code?

How does the program get
executed in detail?

Link, load, execute, finish?

Course Perspective

Our Course is Programmer-Centric

Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer

Enable you to

- Write programs that are more reliable and efficient

- Incorporate features that require hooks into OS

- E.g., concurrency, signal handlers

- Build an operating system!

Cover material in this course that you won't see elsewhere

Not just a course for dedicated **hackers**

Outline

Why we study OS?

Course information

8 problems of OS

Faculty Information

Instructor

XIA Yubin 夏虞斌 CHEN Rong 陈榕

Office: Room 3-402, Software Building

Office hour: 1:00pm~2:00pm, Friday

Contact:

xiayubin@sjtu.edu.cn, rongchen@sjtu.edu.cn

TA

DONG Mingkai MinGKaiDong@gmail.com

YU Qianqian kiki_yu@foxmail.com

ZHENG Yang tomsun.0.7@gmail.com

Reference Books

xv6: a simple, Unix-like teaching operating system.

By Russ Cox, Frans Kaashoek, Robert Morris.
Online

Understanding Linux Kernel.

Labs

Implement JOS: adopted from MIT 6.828 course

[Lab1](#): Booting a PC

[Lab2](#): Memory Management

[Lab3](#): User Environments

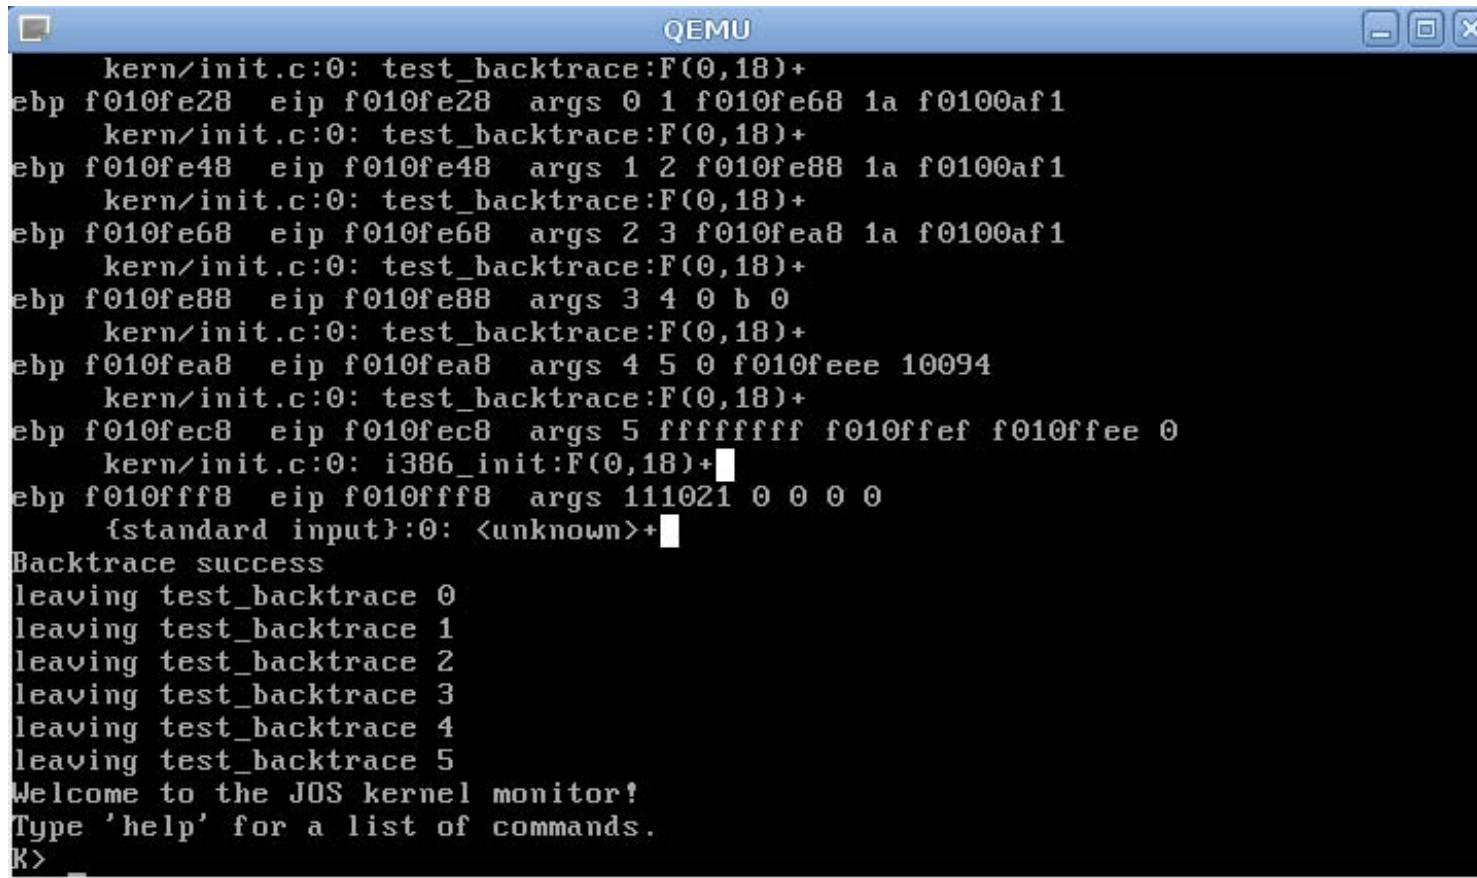
[Lab4](#): Preemptive Multitasking

*[Lab 5: XXX](#)

*[Lab 6: XXX](#)

Lab 1 Boot a PC

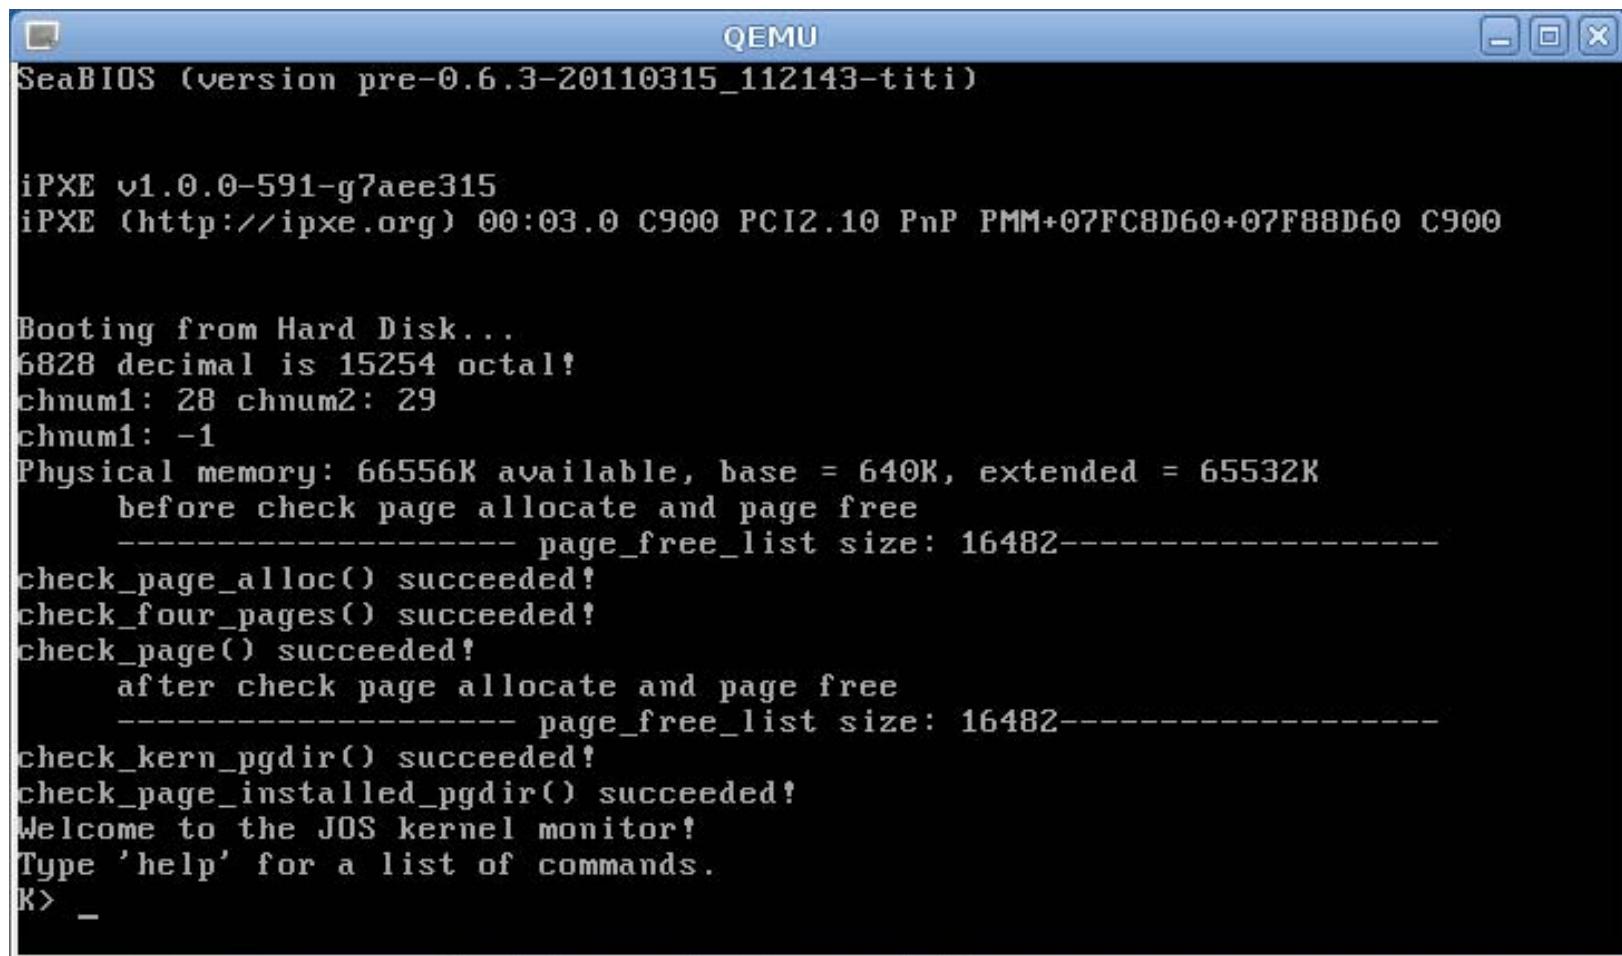
Familiar with PC boot process



```
QEMU
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fe28 eip f010fe28 args 0 1 f010fe68 1a f0100af1  
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fe48 eip f010fe48 args 1 2 f010fe88 1a f0100af1  
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fe68 eip f010fe68 args 2 3 f010fea8 1a f0100af1  
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fe88 eip f010fe88 args 3 4 0 b 0  
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fea8 eip f010fea8 args 4 5 0 f010fee 10094  
kern/init.c:0: test_backtrace:F(0,18)+  
ebp f010fec8 eip f010fec8 args 5 ffffffff f010ffef f010ffee 0  
kern/init.c:0: i386_init:F(0,18)+  
ebp f010fff8 eip f010fff8 args 111021 0 0 0 0  
{standard input}:0: <unknown>+  
Backtrace success  
leaving test_backtrace 0  
leaving test_backtrace 1  
leaving test_backtrace 2  
leaving test_backtrace 3  
leaving test_backtrace 4  
leaving test_backtrace 5  
Welcome to the JOS kernel monitor!  
Type 'help' for a list of commands.  
K>
```

Lab2 Memory Management

Setup virtual memory map for JOS



The screenshot shows a terminal window titled "QEMU" running SeaBIOS. The logs include iPXE boot information, disk booting, and memory allocation tests. The JOS kernel then initializes and welcomes the user.

```
QEMU
SeaBIOS (version pre-0.6.3-20110315_112143-titi)

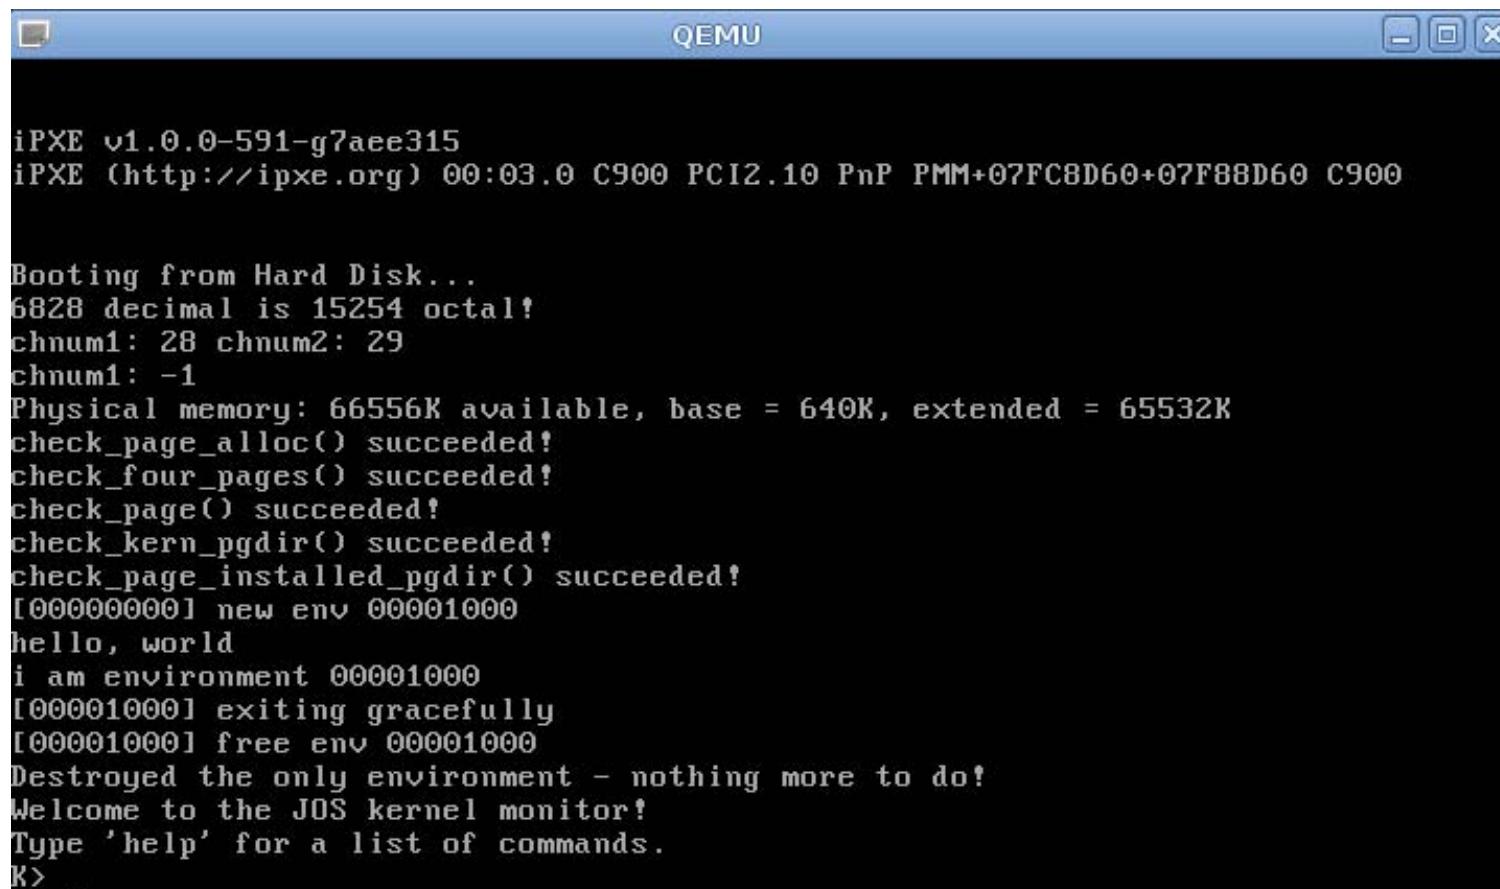
iPXE v1.0.0-591-g7aee315
iPXE (http://ipxe.org) 00:03.0 C900 PCI2.10 PnP PMM+07FC8D60+07F88D60 C900

Booting from Hard Disk...
6828 decimal is 15254 octal!
chnum1: 28 chnum2: 29
chnum1: -1
Physical memory: 66556K available, base = 640K, extended = 65532K
    before check page allocate and page free
    ----- page_free_list size: 16482-----
check_page_alloc() succeeded!
check_four_pages() succeeded!
check_page() succeeded!
    after check page allocate and page free
    ----- page_free_list size: 16482-----
check_kern_pgdir() succeeded!
check_page_installed_pgdir() succeeded!
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> _
```

Lab3: User environment

Environment in JOS = Process in Linux

In this lab, only an environment on a single core



The screenshot shows a terminal window titled "QEMU" running on a Windows operating system. The window contains the following text:

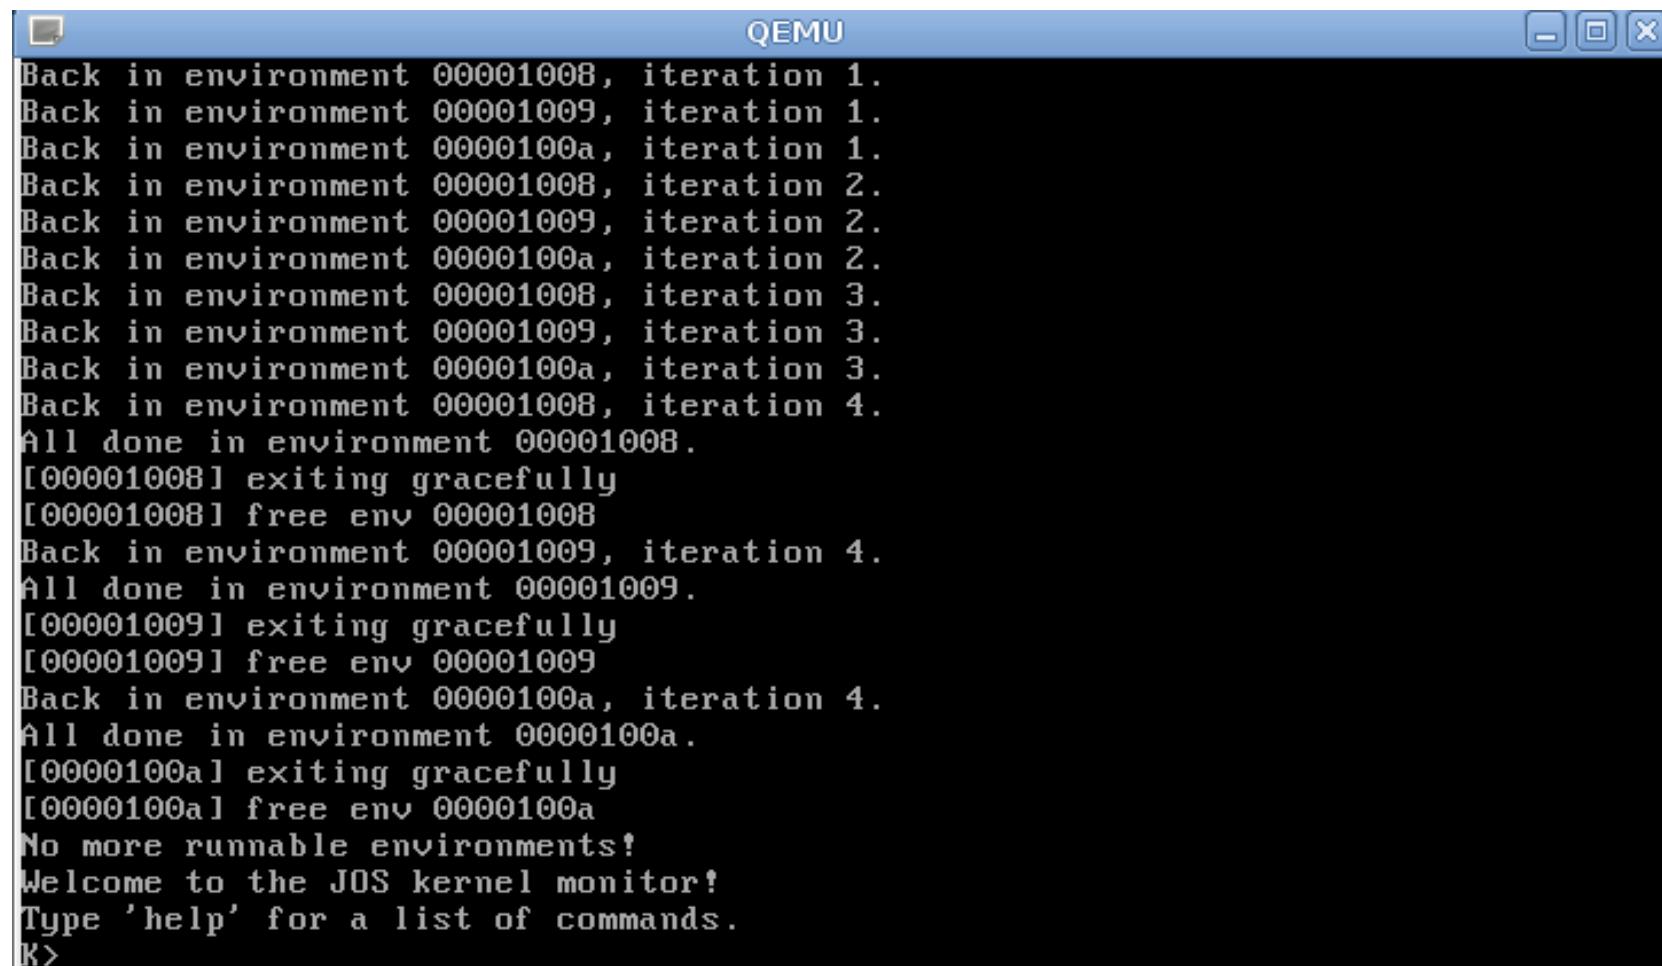
```
iPXE v1.0.0-591-g7aee315
iPXE (http://ipxe.org) 00:03.0 C900 PCI2.10 PnP PMM+07FC8D60+07F88D60 C900

Booting from Hard Disk...
6828 decimal is 15254 octal!
chnum1: 28 chnum2: 29
chnum1: -1
Physical memory: 66556K available, base = 640K, extended = 65532K
check_page_alloc() succeeded!
check_four_pages() succeeded!
check_page() succeeded!
check_kern_pgdir() succeeded!
check_page_installed_pgdir() succeeded!
[00000000] new env 00001000
hello, world
i am environment 00001000
[00001000] exiting gracefully
[00001000] free env 00001000
Destroyed the only environment - nothing more to do!
Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K> _
```

Lab4: Preemptive Multitasking

Multicore

Preemptive multitasking



The screenshot shows a terminal window titled "QEMU". The window contains the following text output:

```
Back in environment 00001008, iteration 1.  
Back in environment 00001009, iteration 1.  
Back in environment 0000100a, iteration 1.  
Back in environment 00001008, iteration 2.  
Back in environment 00001009, iteration 2.  
Back in environment 0000100a, iteration 2.  
Back in environment 00001008, iteration 3.  
Back in environment 00001009, iteration 3.  
Back in environment 0000100a, iteration 3.  
Back in environment 00001008, iteration 4.  
All done in environment 00001008.  
[00001008] exiting gracefully  
[00001008] free env 00001008  
Back in environment 00001009, iteration 4.  
All done in environment 00001009.  
[00001009] exiting gracefully  
[00001009] free env 00001009  
Back in environment 0000100a, iteration 4.  
All done in environment 0000100a.  
[0000100a] exiting gracefully  
[0000100a] free env 0000100a  
No more runnable environments!  
Welcome to the JOS kernel monitor!  
Type 'help' for a list of commands.  
K> _
```

Build your OS

Supporting code will be provided

Still a hard but interesting job

Hints: Start to work early and work hard

Lab Rules

Individual project

Don't copy any code from others

Can read but don't copy other OSes (Linux, Open/FreeBSD, etc.)

Cite any code that inspired your code

Please don't post any “攻略” or code in BBS/forum

Grading

General (Tentative)

Lab 1~4: 40%

Final Exam: 40%

Labs will be reflected in exam

Homework: 10%

Help you to understand OS/Labs

Course performance (5%)

Q&A in classroom, Presentation



Bonus (5%): Lab5 & Lab6

Must finish both in order to get at least 80% scores

8 IMPORTANT PROBLEMS IN OPERATING SYSTEMS

8 Important Problems (Not a Rank)

Scale Up

Security & Privacy

Energy Efficiency

Mobility

Write Correct (Parallel) Code

Scale Out

Non-Volatile Storage

Virtualization

#1: Scale up (or Performance Scalability)

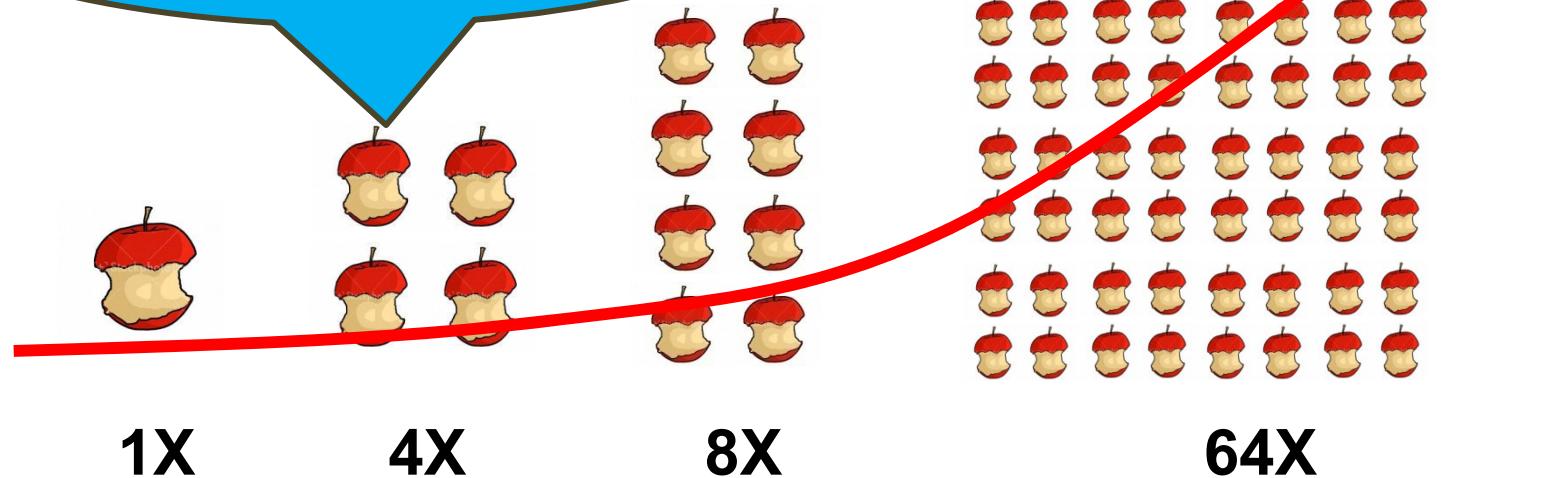
Multicore Evolution

Multicore is commercially prevalent recently

Eight cores and Twelve cores on a chip are common

Hundreds of cores on a single chip will appear in near feature

**Challenge#1: make
software performance
follow Moore's Law**



Why Operating System Matters?

Terms of operating system here

Broadly defined, including hypervisor, operating systems, runtime environments

Operating systems manages and hide resources from applications

Hide tedious and complex low-level details

Provide execution environments to apps

Operating systems determines app perf. in many cases

Many apps spend a large fraction of time in system software

Operating System Meets Multicore

(Operating) systems 20 years ago

- Enjoys the free lunch provided by hardware (Moore's Law)

- Recall the Andy-Bill's Law

Multicore: ending the free lunch!

- Software must evolve with hardware changes

OS: a vital role to bridge the gap between apps and hardware

- Need to evolve itself with multicore

- Need to evolve for apps with multicore

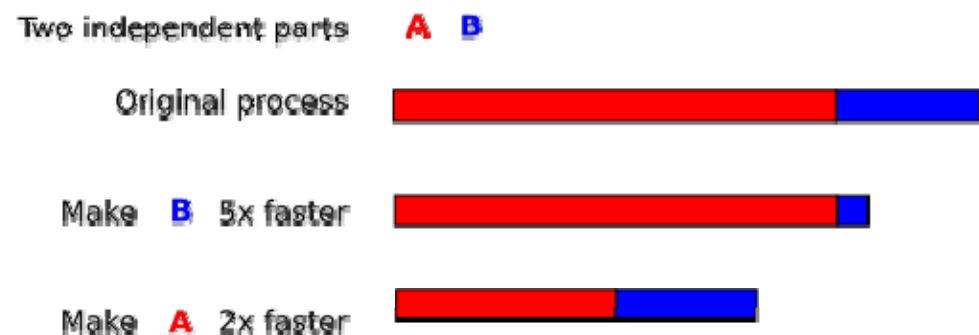
What is scalability?

Application does N times as much work on N cores
as it could on 1 core

Scalability may be limited by Amdahl's Law:

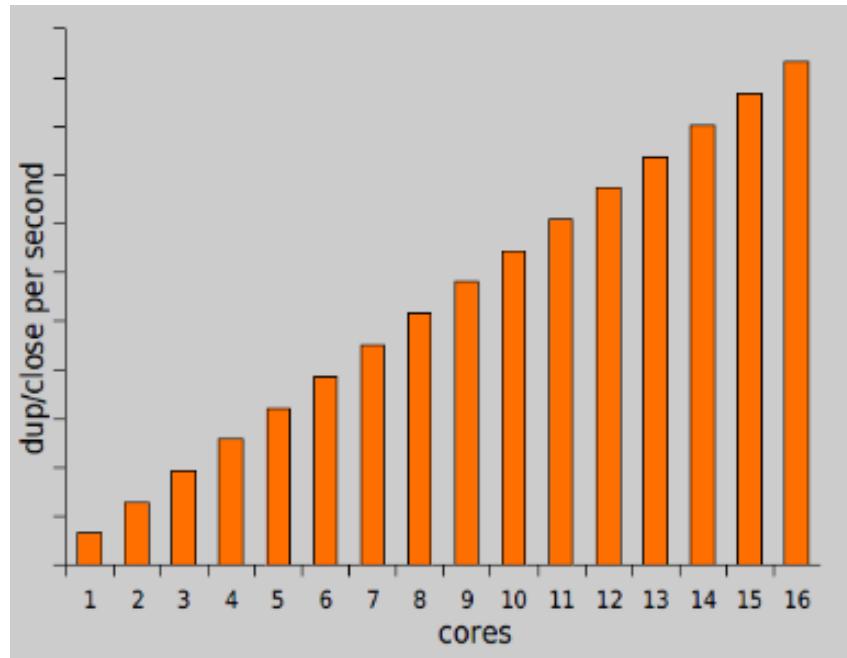
Locks, shared data structures, ... Shared hardware
(DRAM, NIC, ...)

$$\frac{1}{(1 - P) + \frac{P}{S}}$$

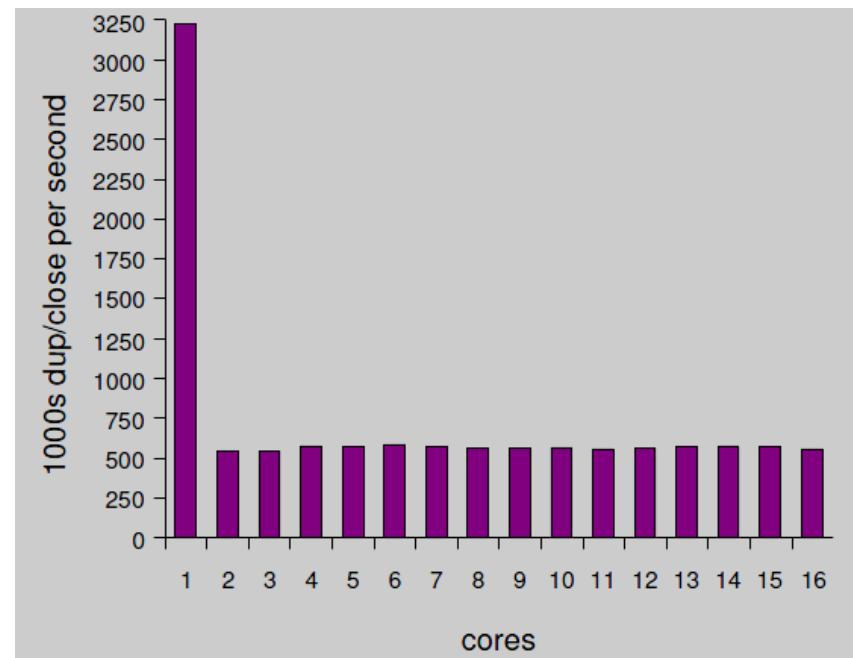


Motivating example: file descriptors

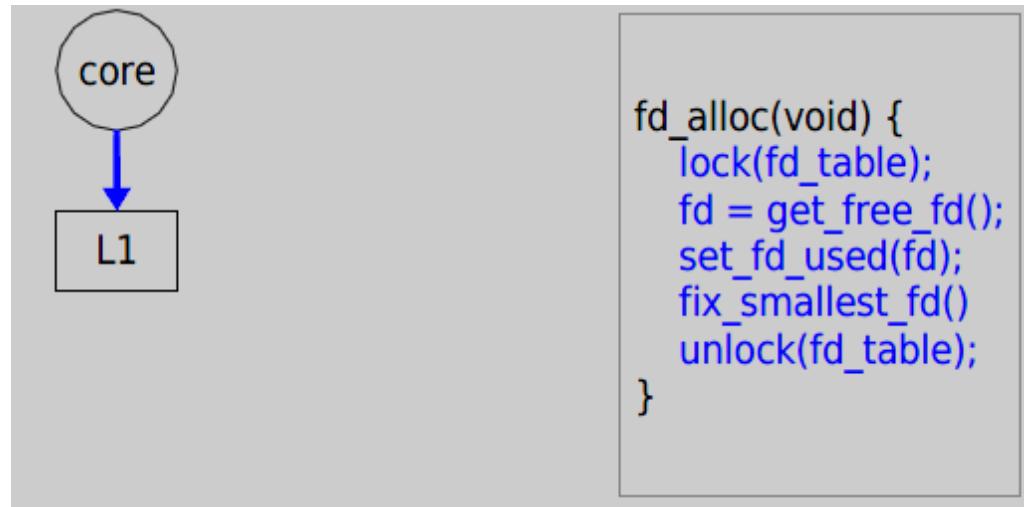
Ideal FD performance graph



Actual FD performance

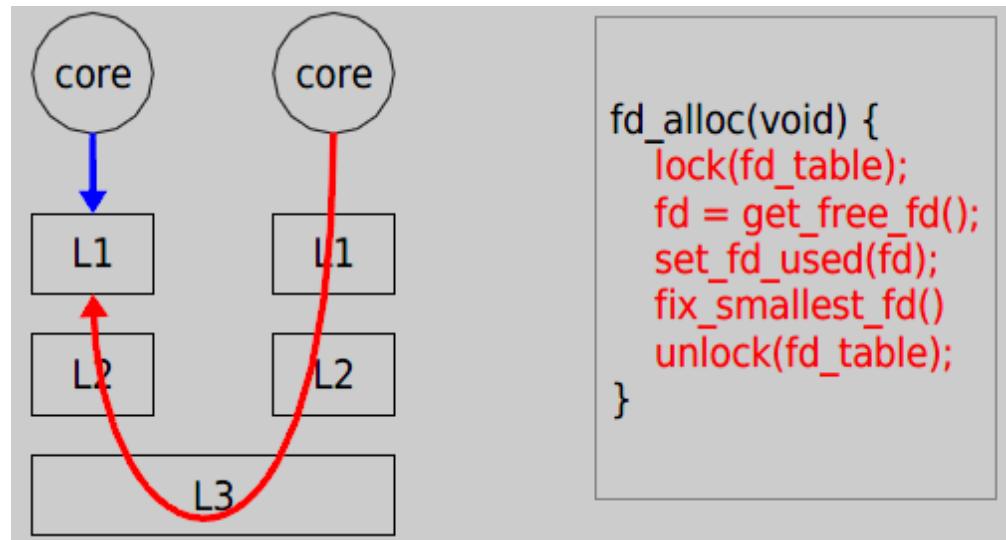


Why throughput drops?



Load fd_table data from L1 in 3 cycles.

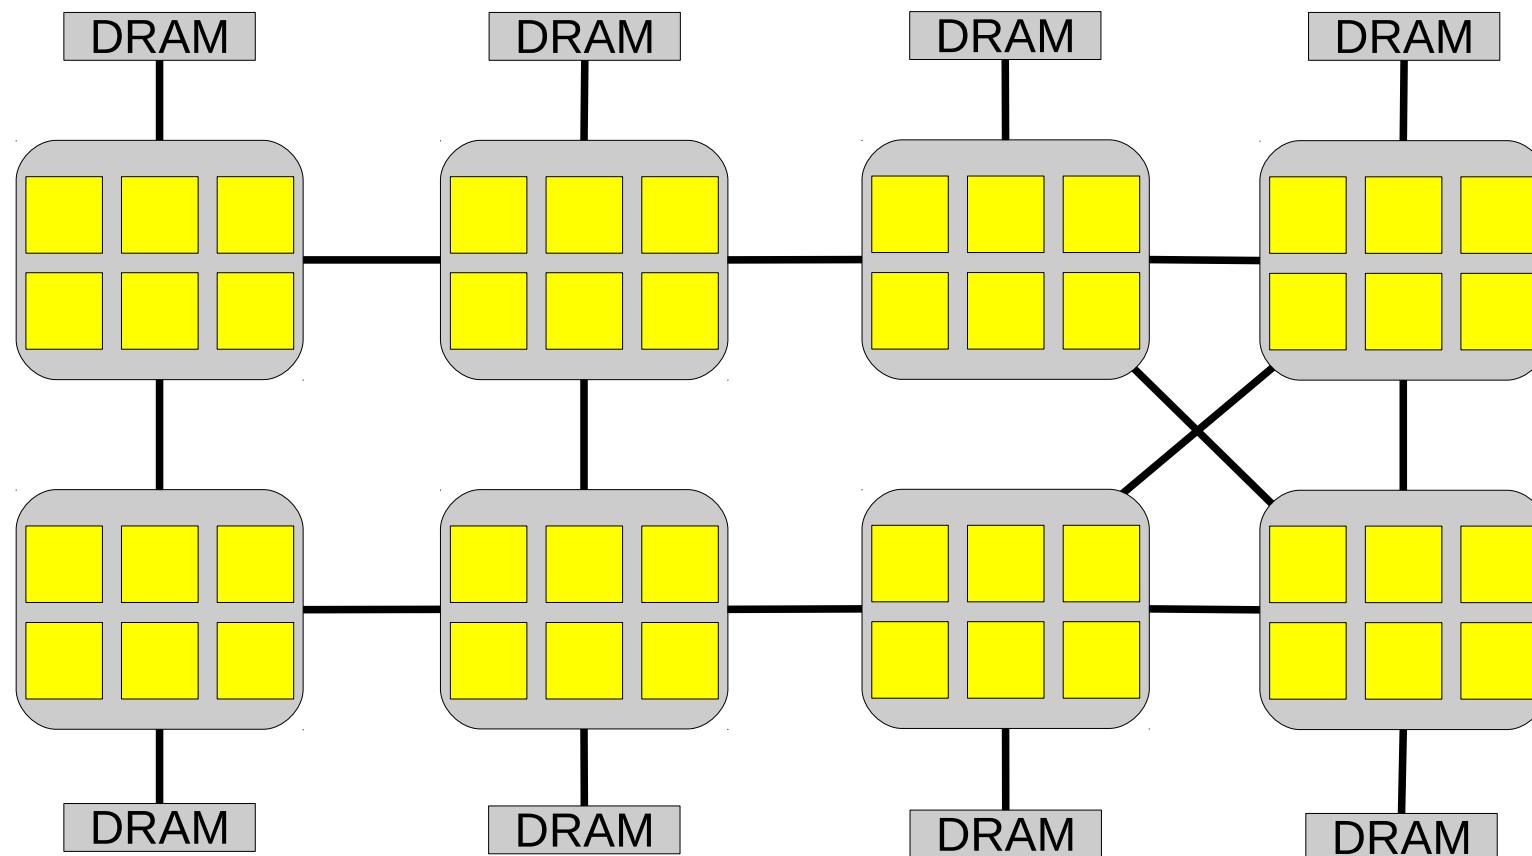
Why throughput drops?



Now it takes 121 cycles!

Off-the-shelf 48-core server

6 core x 8 chip AMD



Scale Up OS

Better abstraction to user applications

Eliminating non-scalable synchronization

Minimize sharing of common data structures

You will fill this..

Issue#2: Security & Privacy

Security



From Wikipedia:

Security is the degree of resistance to, or protection from, harm

Why Operating System matters?

Operating system (including hypervisors) lies in the lowest levels in a computer

How could you protect your application & data without properly protect your OS

Example: The do_brk() Bug

```
Asmlinkage unsigned long sys_brk(unsigned long brk) {
    ...
    /* OK, looks good – let it rip. */
    -if (do_mmap(NULL, oldbrk, newbrk-oldbrk,
    -PROT_READ|PROT_WRITE|PROT_EXEC,
    -MAP_FIXED|MAP_PRIVATE, 0) != oldbrk)
    +if (do_brk(oldbrk, newbrk-oldbrk) != oldbrk)
        goto out;
```

- do_mmap() checked things properly, do_brk() removed almost all checking
- Purpose of do_brk() was to speed up anonymous mmaps from sys_brk()

The do_brk() Fix

```
if (!len)
    return addr;

+ if ((addr + len)) > TASK_SIZE || (addr + len) < addr)
+ return -EINVAL;
+
```

- The TASK_SIZE is typically set to 0xc0000000, the start of kernel memory

Issue#3: Power Efficiency

Power Efficiency

Greener World, Better Life!



Again, why it has to do with OS?

OS control every computer and many devices!

Thus determines the power used by each computer

Google Datacenter: Finland



Power Breakdown in Server

Component	Peak Power	Count	Total
CPU [16]	40 W	2	80 W
Memory [18]	9 W	4	36 W
Disk [24]	12 W	1	12 W
PCI slots [22]	25 W	2	50 W
Motherboard	25W	1	25 W
Fan	10 W	1	10 W
System Total			213 W

Table 1: Component peak power breakdown for a typical server

12 Variable Speed Fans



138 Air Dampers



312 Light Relays



50 Electrical Sub-meters



151 Temperature Sensors



> 6,000 Sense and Control Points



Why Operating System Matters?

Dynamic voltage and frequency scaling

Adjust the frequency of CPU according to workload

Power off devices to save power

Like disk, LCD, or even memory

Problem: energy consumption doesn't scale with workloads

Static power, leakage power

New!: Near Threshold Voltage

Issue#4: Mobility

Smartphone OS

Symbian

Windows Mobile

RIM Blackberry OS

Apple iOS

Google Android

Palm WebOS

Windows Phone 7

Android Software Stack

Linux kernel

Libraries

Android run time

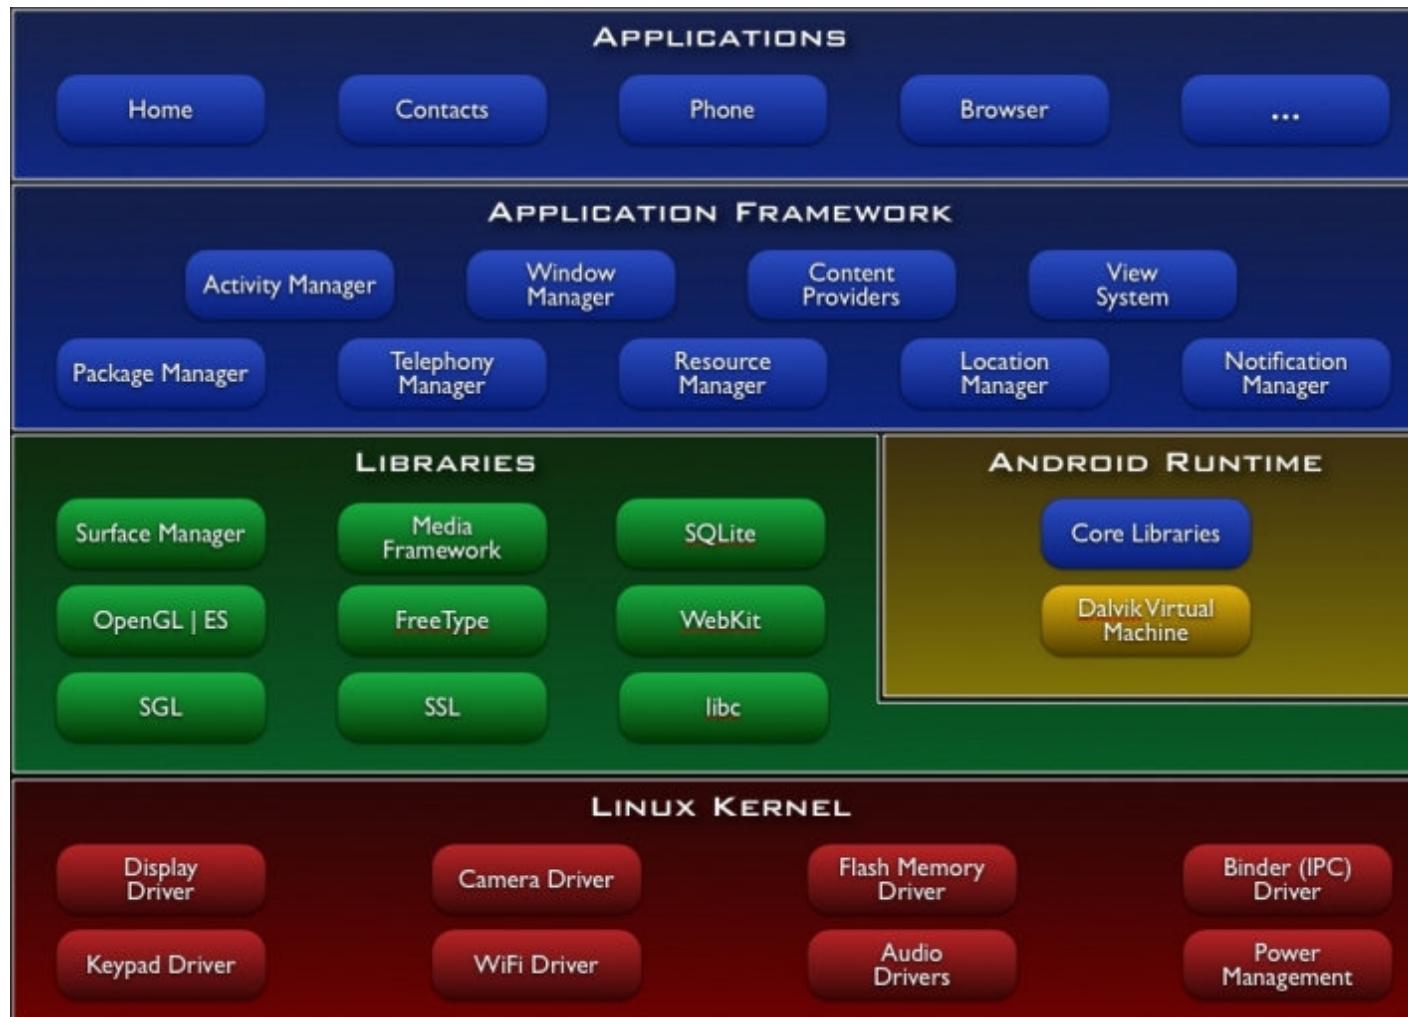
- core libraries

- Dalvik virtual machine

application layer

application protocol

Android Architecture



What's Unique for Mobile OS?

Energy Efficiency

Richer User Experience

But more limited resources

Security

More user data put into mobile OSes

Issue#5: Write Correct Parallel Code

Writing concurrent program is hard

Concurrent programs are prone to concurrency bugs

Concurrency bugs have notorious characteristics

- Non-deterministic

- Hard to test and diagnose

But in multi-core era, we need to write concurrent program to harness the power of multi-core

Quote from Dijkstra

I am convinced more than ever that this type of work is very difficult, and that every effort to do it with other than the best people is doomed to either failure or moderate success at enormous cost.

Edsger Dijkstra
The Structure of the “THE” –Multiprogramming System
1968

Easiest One

Thread 1

Thread 2

```
if (thd->proc_info) {  
    ...  
    thd->proc_info = NULL;  
  
    fputs(thd->proc_info, ...)  
  
    ...  
}
```

MySQL ha_innodb.hpp

Approaches to addressing the problem

Concurrency bug detection

- Race detection

- Atomicity violation detection

- Deadlock bug detection

Concurrent program testing

- Exhaustive testing

- Different coverage criteria proposed

Concurrent programming language/model design

- Transactional memory

What OS can do?

Reduce sources of non-determinism

Tolerate application bugs

Faithfully capture concurrency bugs for reproduce

Provide better programming interfaces to ease
programming

Operating transactions

Issue#6: Scale Out (use distributed systems)

The Datacenter as a Computer

*An Introduction to the Design of
Warehouse-Scale Machines*

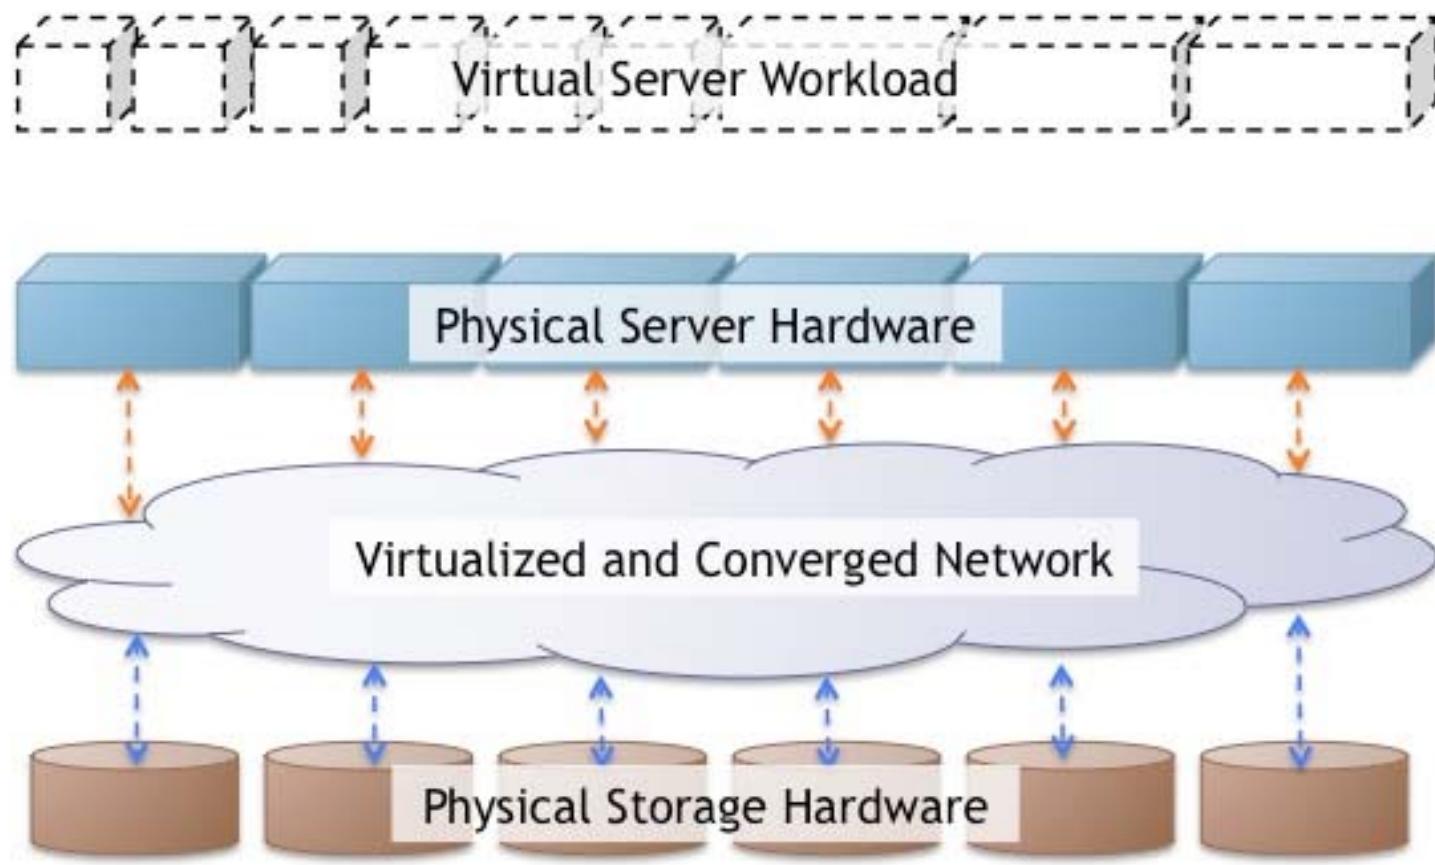
Luiz André Barroso and Urs Hözle
Google Inc.

SYNTHESIS LECTURES ON COMPUTER ARCHITECTURE # 6

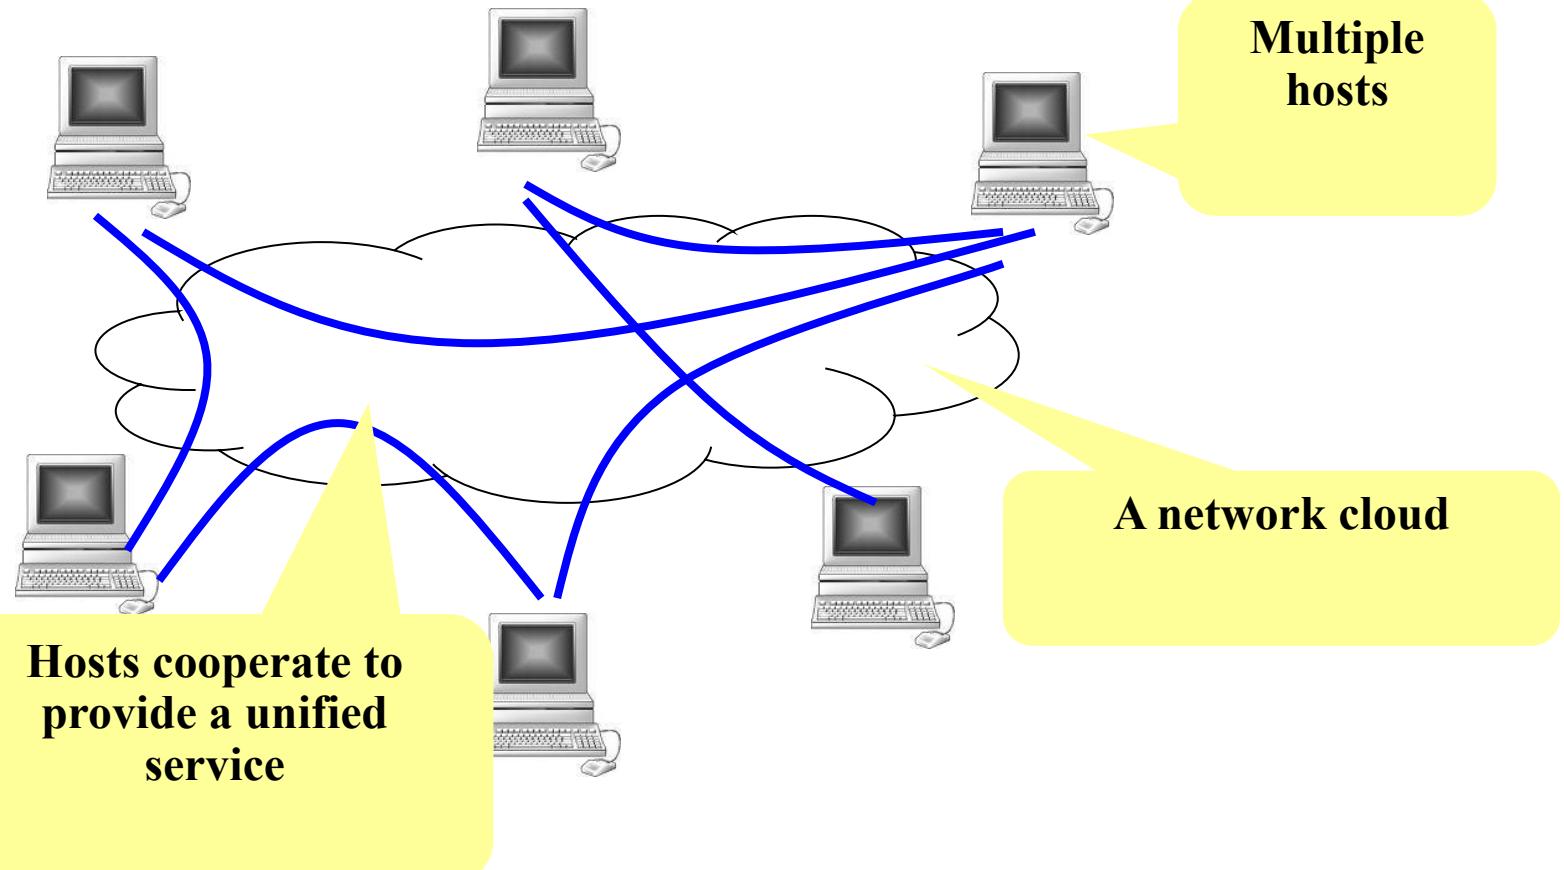


MORGAN & CLAYPOOL PUBLISHERS

The Virtual Datacenter Dream



What are distributed systems?



- Examples?

Why distributed systems? for ease-of-use

Handle geographic separation

Provide users (or applications) with location transparency:

- Web: access information with a few “clicks”

- Network file system: access files on remote servers as if they are on a local disk, share files among multiple computers

Why distributed systems? for availability

Build a reliable system out of unreliable parts

Hardware can fail: power outage, disk failures, memory corruption, network switch failures...

Software can fail: bugs, mis-configuration, upgrade ...

To achieve 99.999% availability, replicate data/computation on many hosts with automatic failover

Why distributed systems? for scalable capacity

Aggregate resources of many computers

CPU: Dryad, MapReduce, Grid computing

Bandwidth: Akamai CDN, BitTorrent

Disk: Frangipani, Google file system

Why distributed systems? for modular functionality

Only need to build a service to accomplish a single task well.

Authentication server

Backup server

Challenges

System design

What is the right **interface** or abstraction?

How to partition functions for scalability?

Consistency

How to share data consistently among multiple readers/writers?

Fault Tolerance

How to keep system available despite node or network failures?

Challenges (continued)

Different deployment scenarios

- Clusters

- Wide area distribution

- Sensor networks

Security

- How to authenticate clients or servers?

- How to defend against or audit misbehaving servers?

Implementation

- How to maximize concurrency?

- What's the bottleneck?

- How to reduce load on the bottleneck resource?

A word of warning

A distributed system is a system in which I can't do my work because some computer that I've never even heard of has failed."

-- Leslie Lamport

Issue#7: Non-Volatile Storage

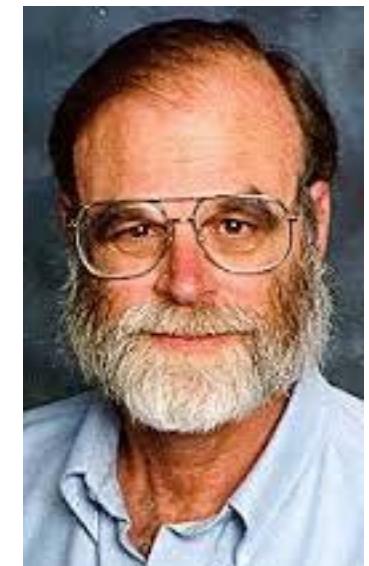
Tape is Dead Disk is Tape Flash is Disk **RAM Locality is King**

Jim Gray

Microsoft

December 2006

In memory of Jim Gray



Tape Is Dead Disk is Tape

1TB disks are available

10+ TB disks are predicted in 5 years

Unit disk cost: ~\$400 → ~\$80

But: ~ 5..15 **hours to read (sequential)**

~15..150 **days to read (random)**

Need to treat **most of disk as
Cold-storage archive**

FLASH Storage?

1995 16 Mb NAND flash chips

2005 16 Gb NAND flash

Doubled each year since 1995

Market driven by Phones, Cameras, iPod,...

Low entry-cost,

~\$30/chip → ~\$3/chip

2012 1 Tb NAND flash

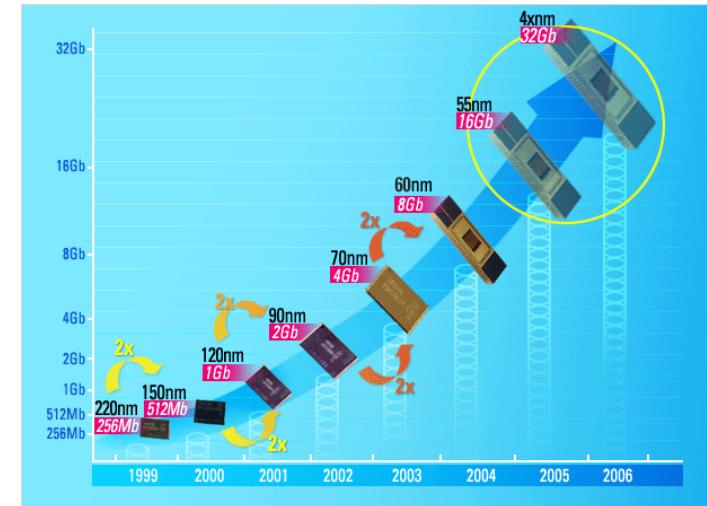
== 128 GB chip

== 1TB or 2TB “disk”

for ~\$400

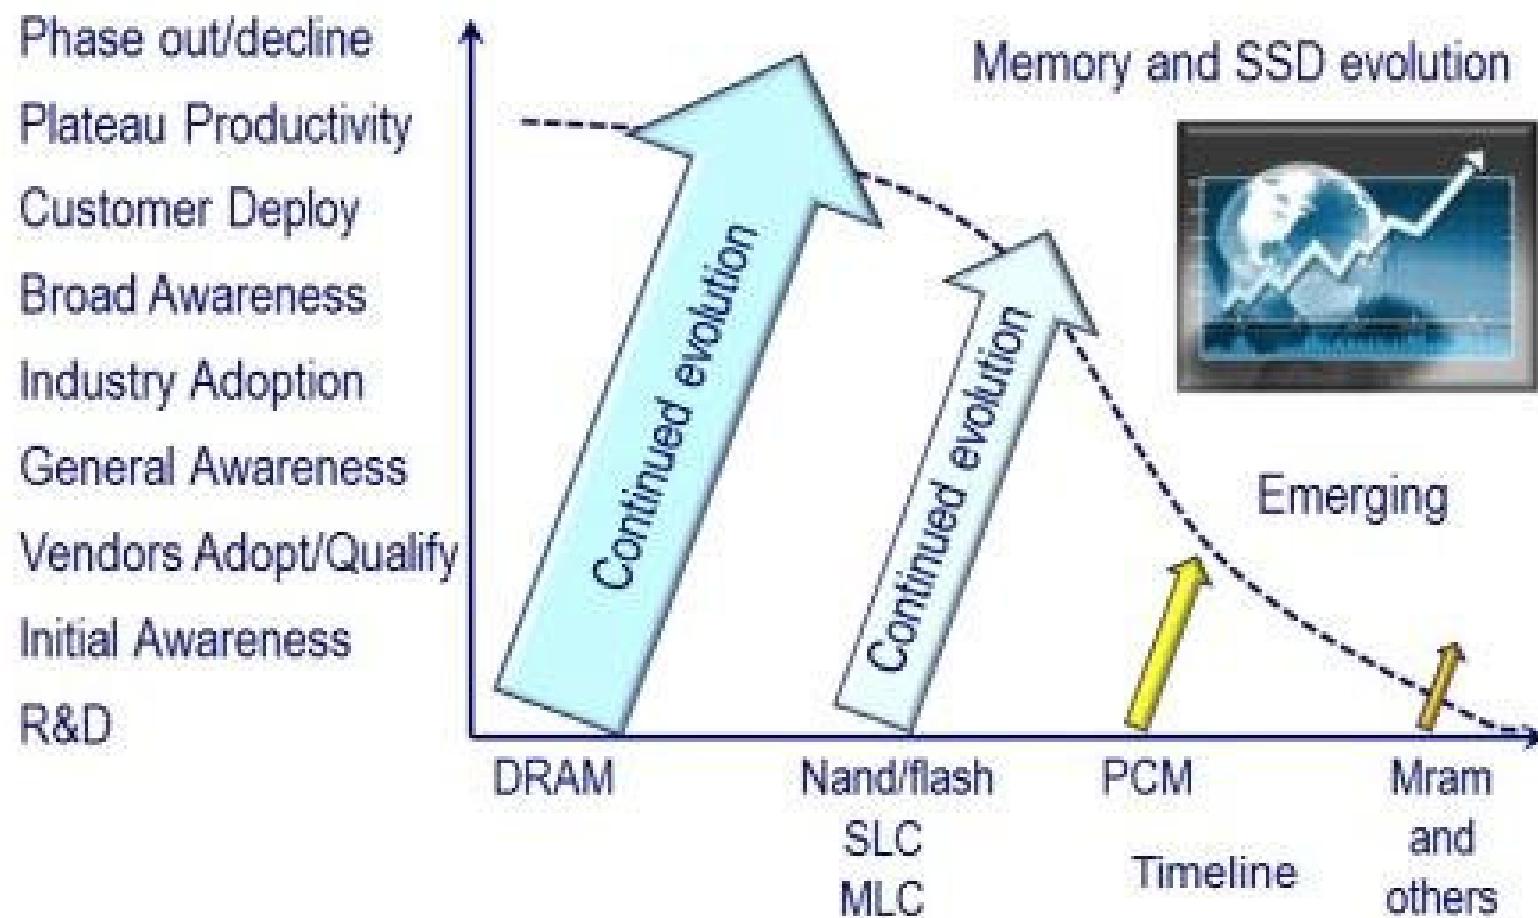
or 128GB disk for \$40

or 32GB disk for \$5



Samsung prediction

Storage Hiereary



What can OS do?

Needs to adapt to the new storage model

Non-volatility changes the way operating system manage resources

- Like crash recovery

- Recall fsck

Better I/O performance

Issue#8: Virtualization

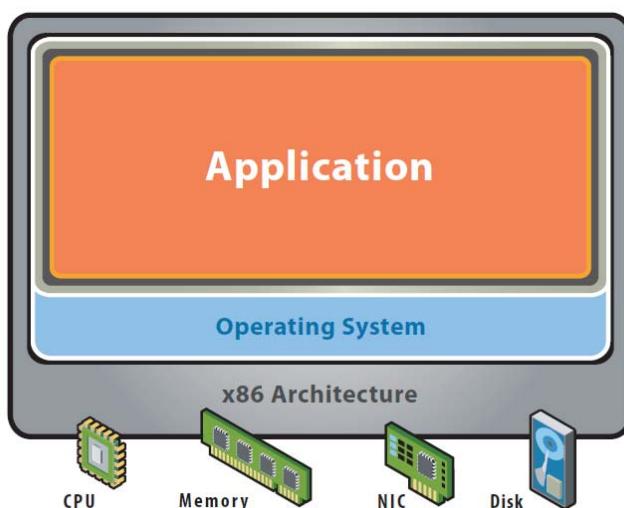
*"Any problem in computer science can
be solved with another level of
indirection."*

– *David Wheeler in Butler Lampson's
1992 ACM Turing Award speech.*

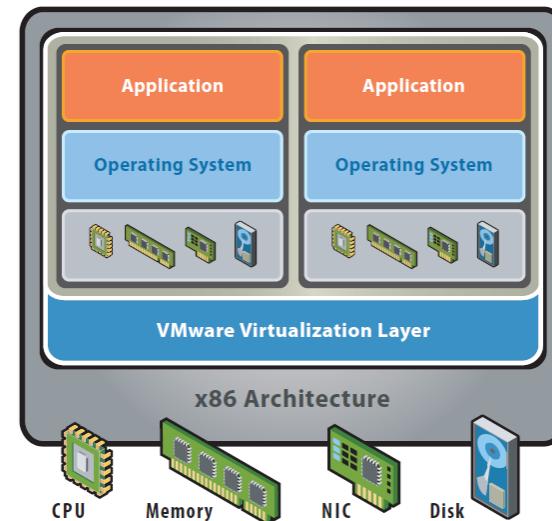
Concept of Virtualization

“Virtualization” refers to the creation of a virtual machine

A virtual machine is a software implementation of a machine (i.e. a computer) that executes programs like a physical machine



Before Virtualization



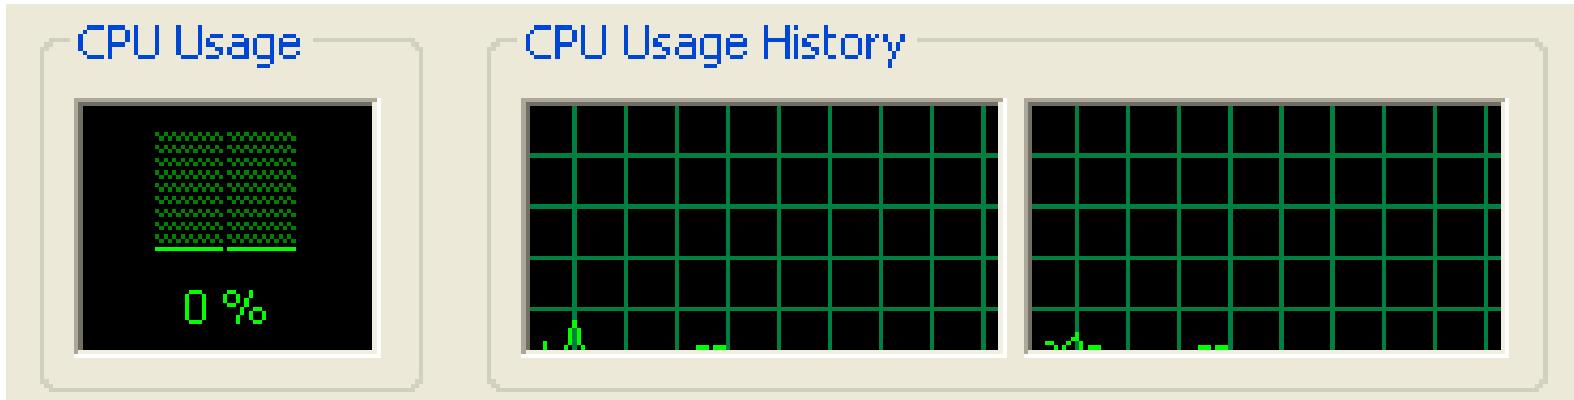
After Virtualization

Why Virtualization ?

- Decoupling the operating systems from physical hardware supports
 - Maximum resource utilization
 - Multiple OS environments can co-exist on the same computer, in strong isolation from each other
 - Server consolidation
 - Fault tolerance
 - Fault and migrate
 - Portability
 - Can use any OS
 - Manageability
 - Easy to maintain

Why Virtualize?

Too many servers for too little work



High costs and infrastructure needs

Maintenance

Networking

Floor space

Cooling

Power

Disaster Recovery



Virtualization is Essentially Another Layer of OS

How does the process and OS use hardware resources?

	process	OS
CPU	Non-privileged registers and instructions	+Privileged registers and instructions
memory	Virtual memory	+Physical memory
exceptions	Signals, errors	+Traps, interrupts
I/O	File system	Programmed I/O, DMA, interrupts

Next Class

Introduction to PC programming

Booting of an OS

Homework

- Writing your answer (in Chinese) to following questions
- What is do_brk()? How to exploit?
 - <http://www.csee.umbc.edu/courses/undergraduate/421/Spring12/02/slides/ULKV.pdf>
 - Give your own explanation

PC Programming & Booting

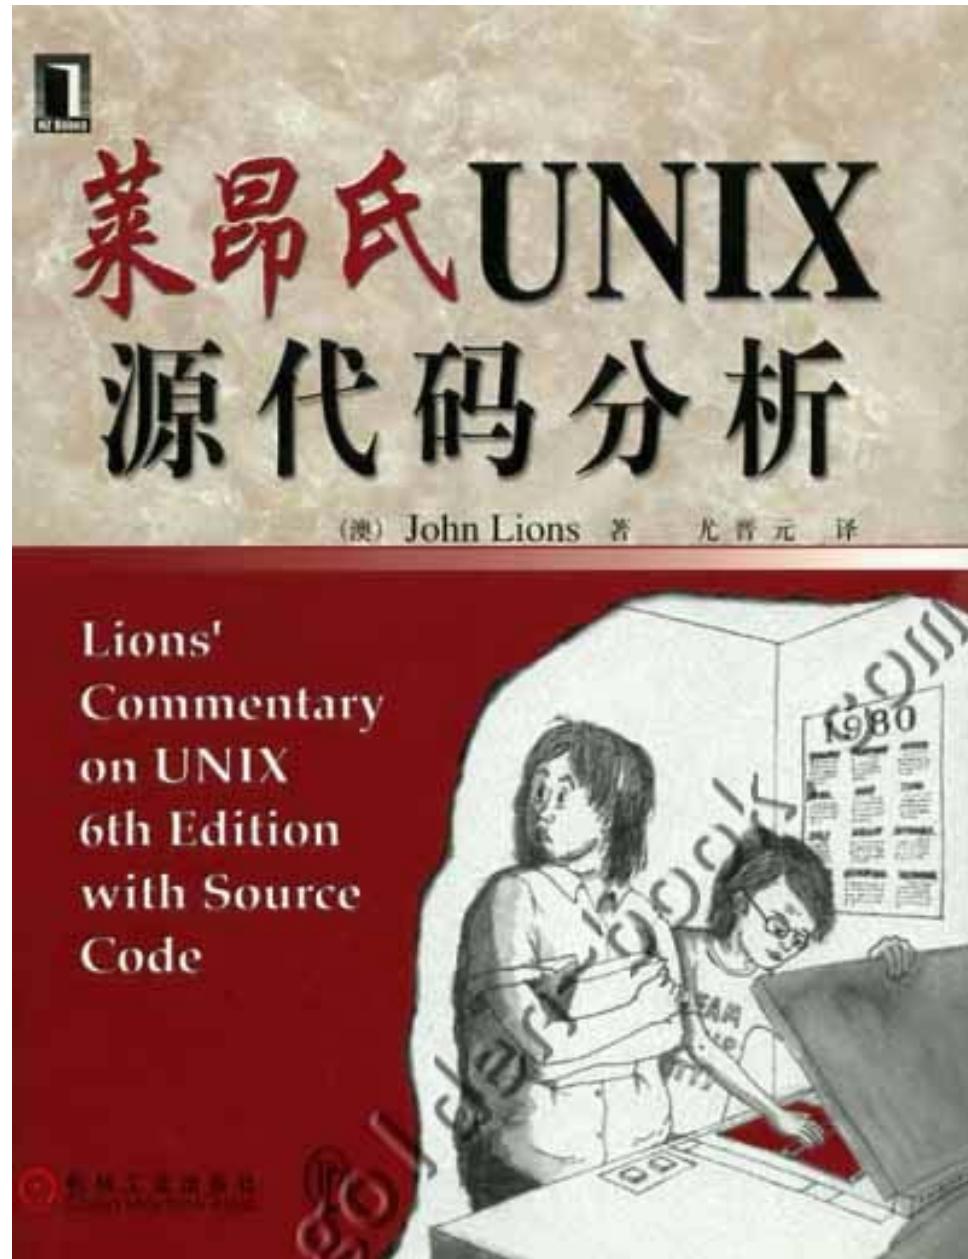
Yubin Xia & Rong Chen

Our Tools

- QA Site:
 - <http://ipads.se.sjtu.edu.cn/courses/qa>
- Reference
 - <http://pdos.csail.mit.edu/6.828/2012/xv6.html>
 - Both the code and book
 - <http://wiki.osdev.org>
 - [IA-32 Intel Architecture Software Developer's Manuals](#)
 - [Volume 3A: System Programming Guide, Part 1](#)
 - [Volume 3B: System Programming Guide, Part 2](#)

Once upon a time ...

- Lion's commentary
- Based on UNIX v6
- Which is not on x86
 - But PDP-11
- Xv6
 - Unix v6
 - For x86!
 - Runnable!



A PC



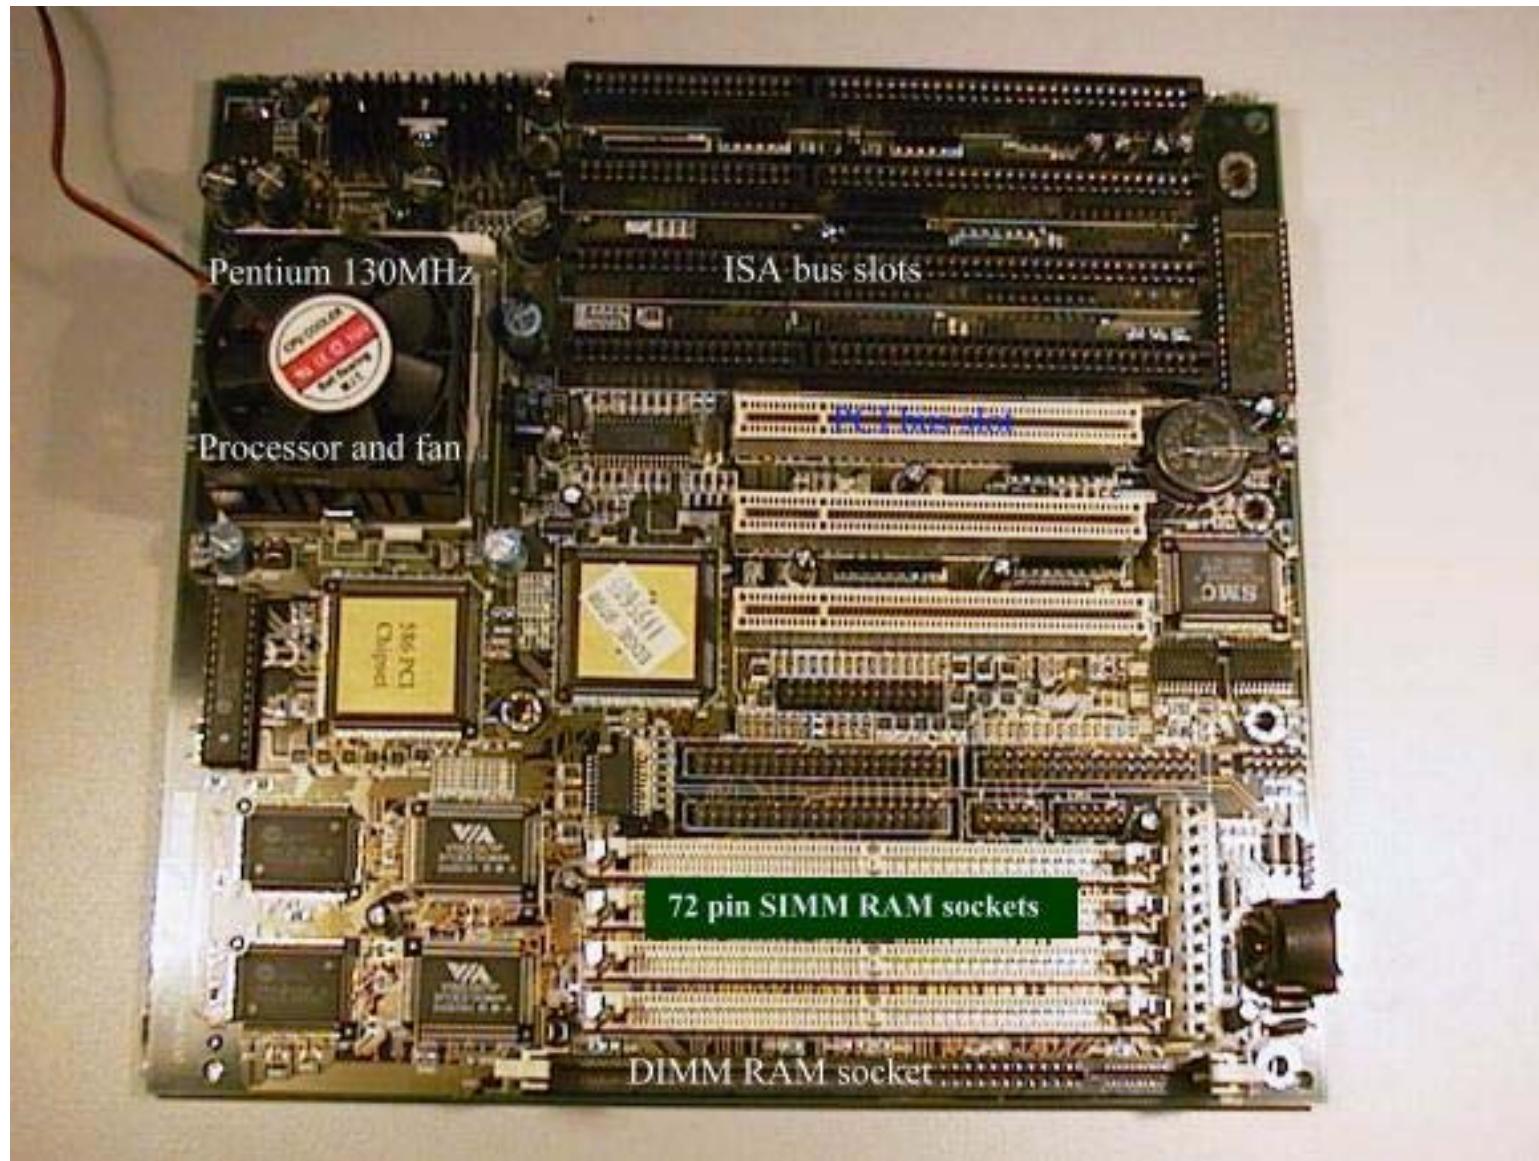
How to make it to do something useful?

Outline

- PC Architecture
- Memory
- Execution
- PC Emulation

- PC Architecture
- Memory
- Execution
- PC Emulation

PC board



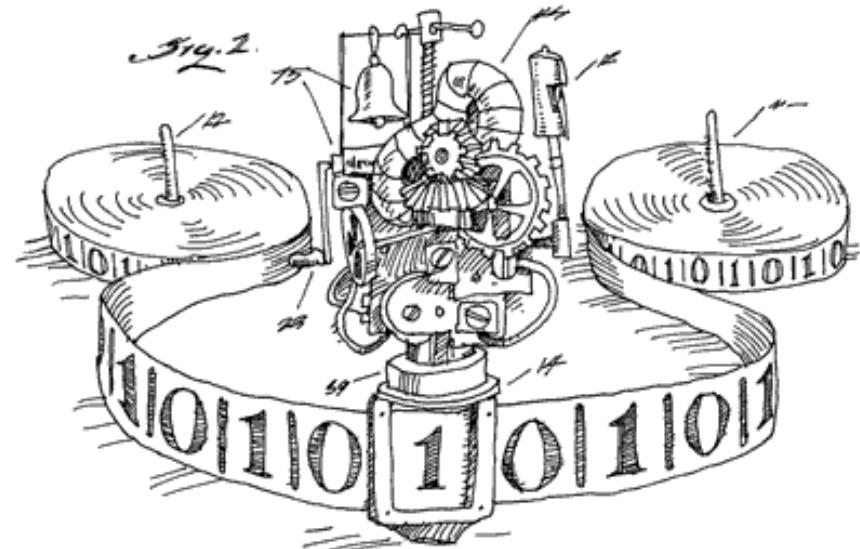
PC board - Nowadays



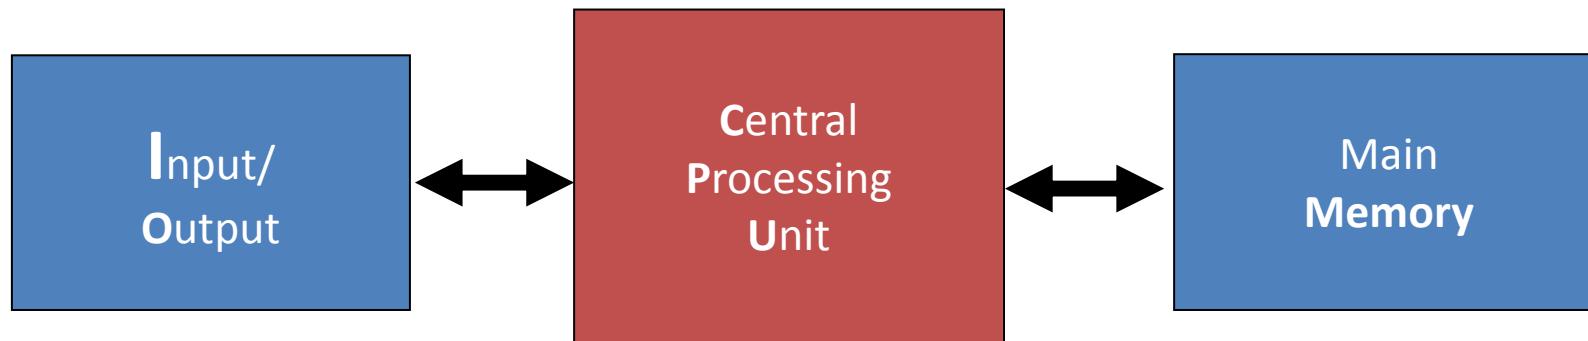
The Turing Machine

Calculate 3+2: 0000001110110000

		0000001110110000
1	0	0000001110110000
2	0	0000001110110000
3	0	0000001110110000
4	0	0000001110110000
5	0	0000001110110000
6	0	0000001110110000
7	0	0000001110110000
8	1	0000001110110000
9	1	0000001110110000
10	1	0000001110110000
11	10	0000001111110000
12	10	00000011111110000
13	10	00000011111110000
14	11	00000011111110000
15	0	000000111111100000

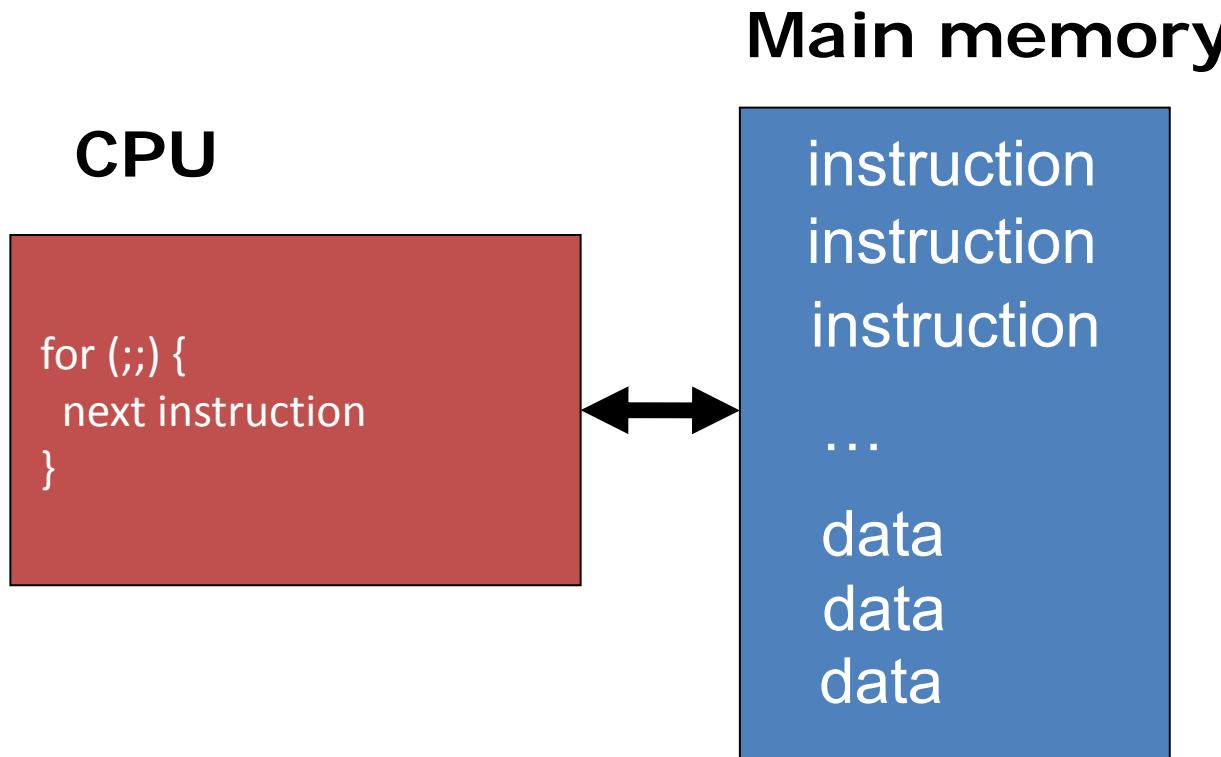


The von Neumann Model



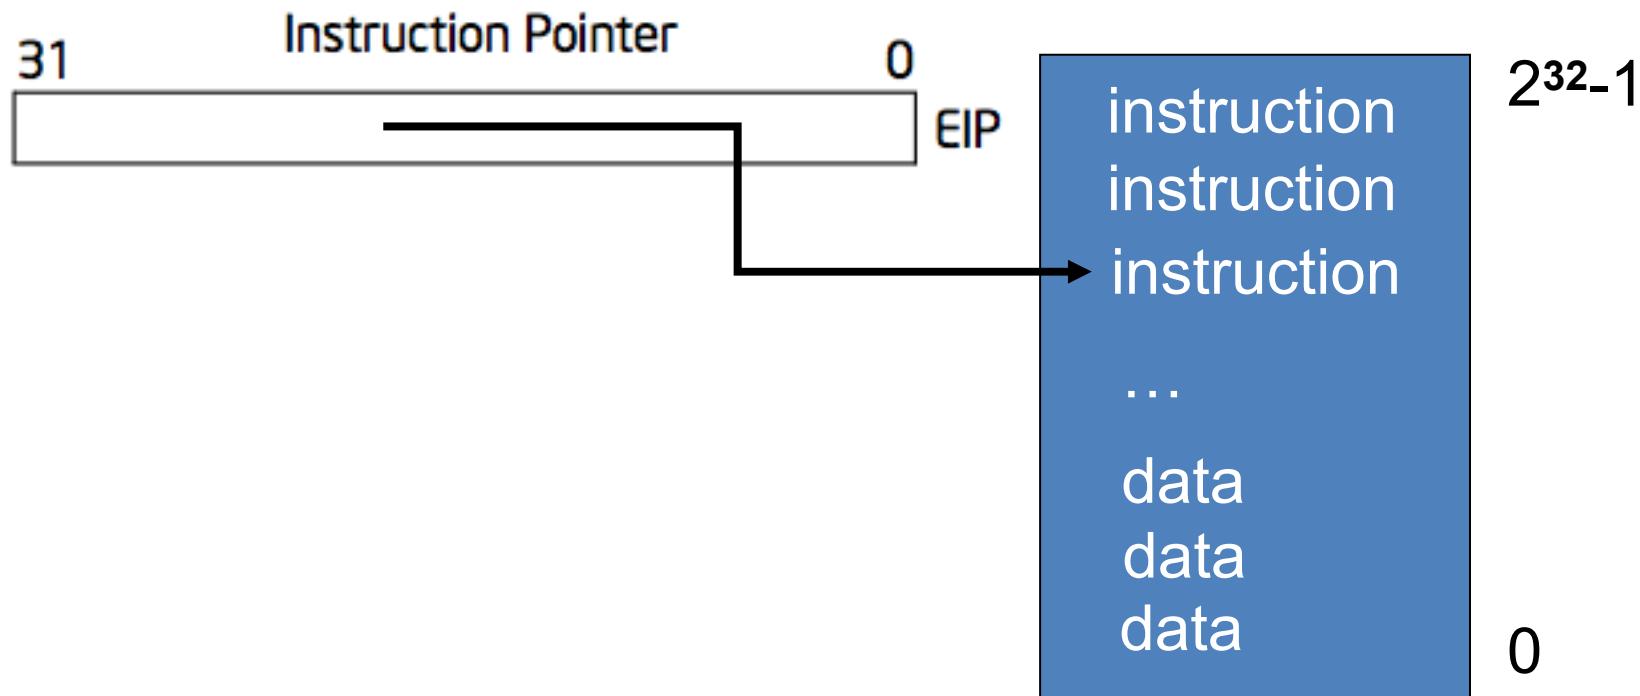
- **I/O:** communicating data to and from devices
- **CPU:** digital logic for performing computation
- **Memory:** N words of B bits

The Stored Program Computer



- **CPU** interpreter of instructions
- **Memory** holds instructions and data

x86 implementation



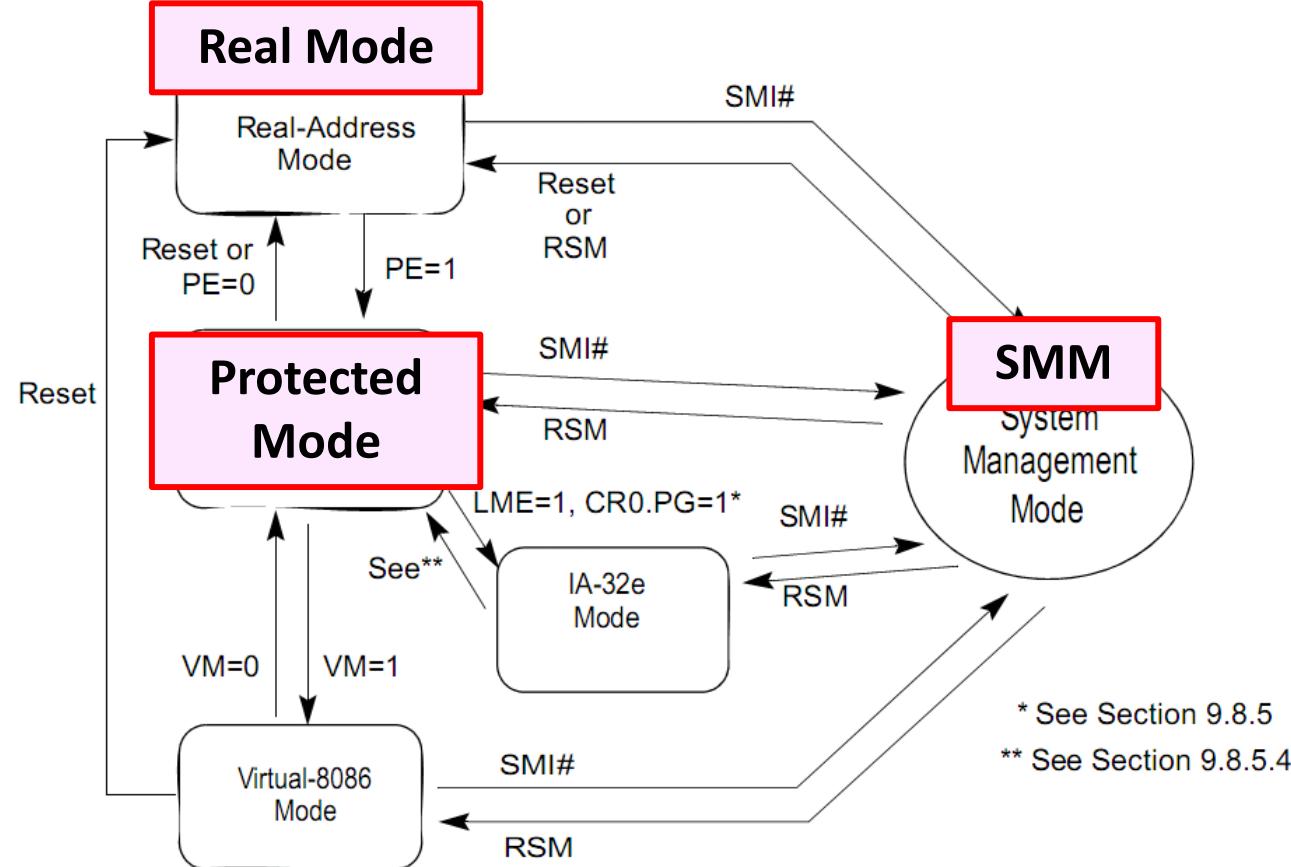
- **EIP** is incremented after each instruction
- Instructions are different length
- EIP modified by CALL, RET, JMP, and cond. JMP

System Architecture Overview

Volume 3A: System Programming Guide, Part 1

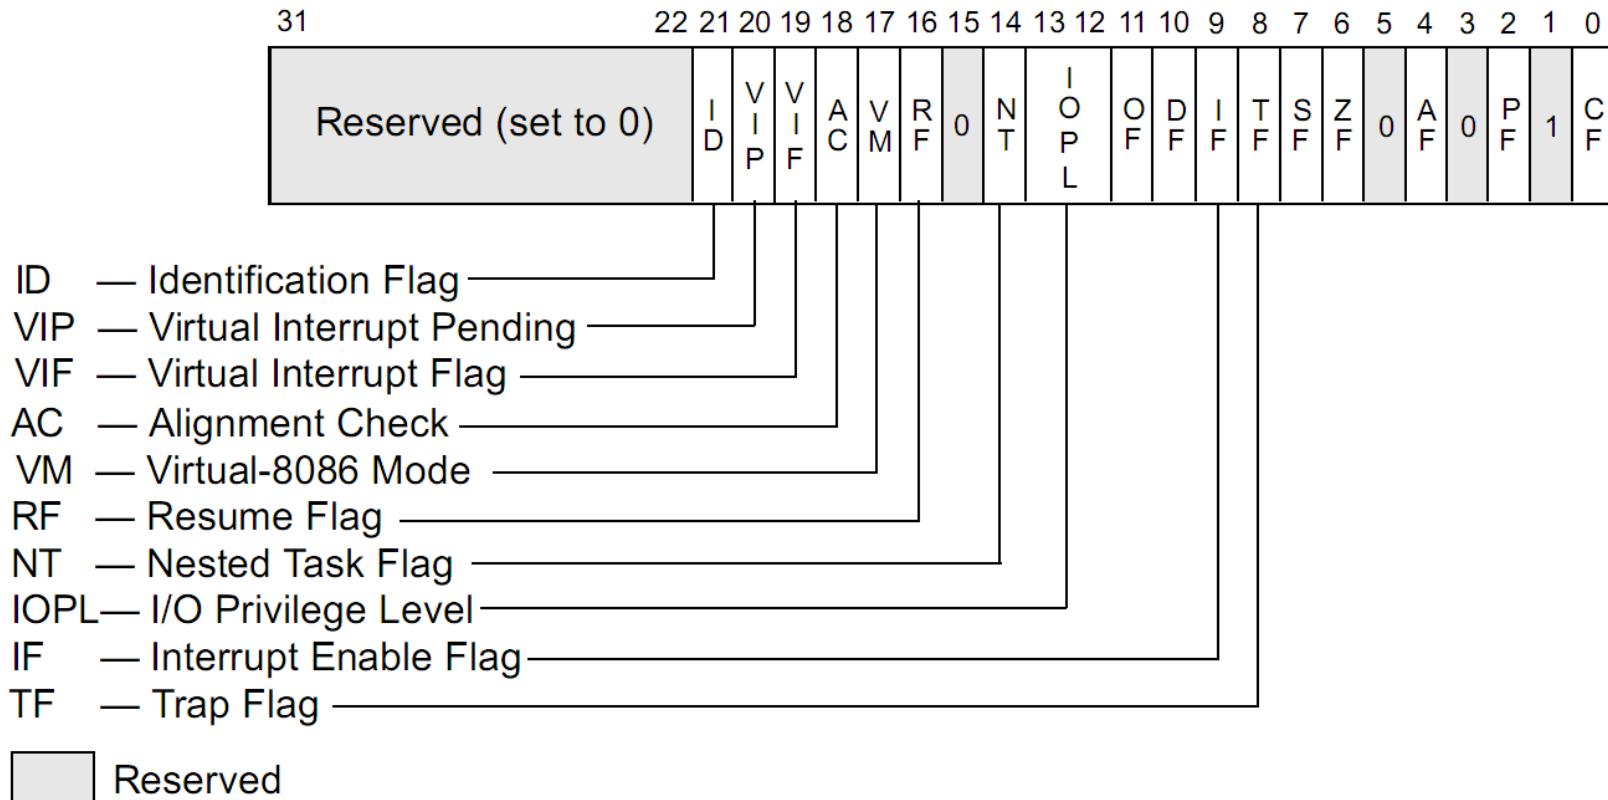
- 2.1~2.5

- Modes



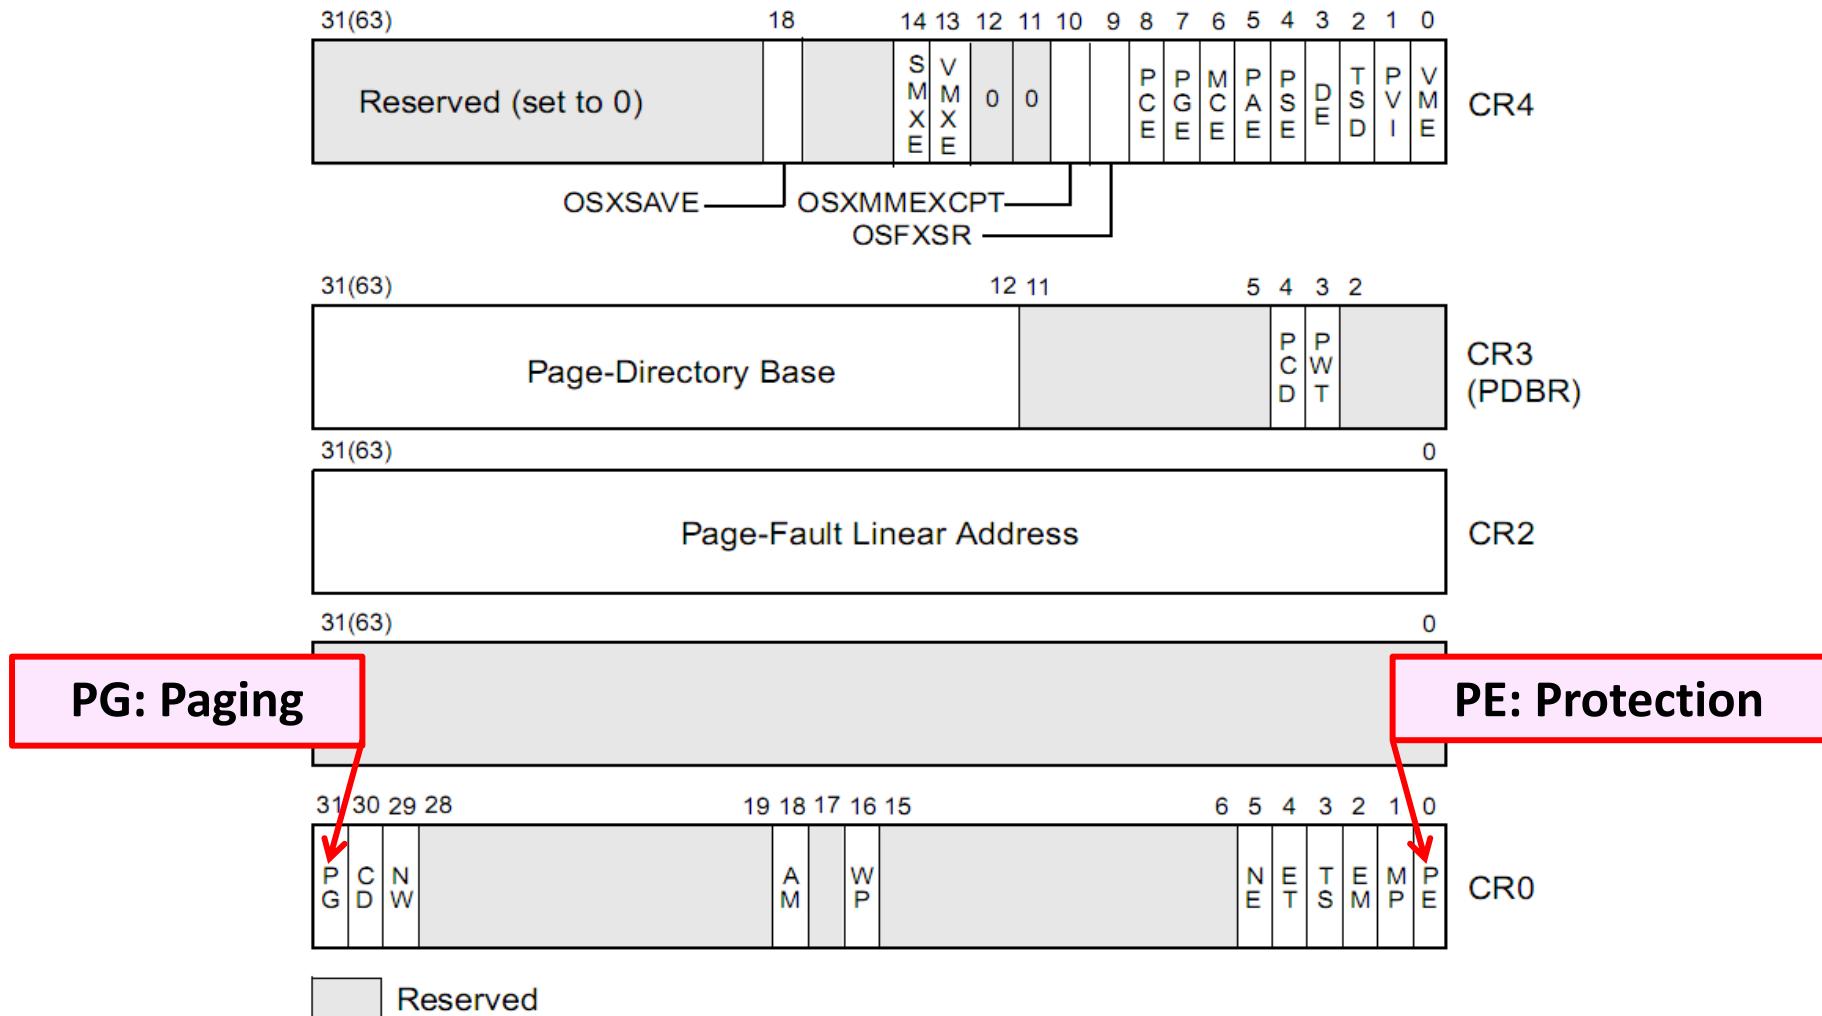
System Architecture Overview

- System Flags in the EFLAGS



System Architecture Overview

- Control Registers



System Architecture Overview

- Memory-Management Registers
 - GDTR (Global Descriptor Table Register)
 - Base Address, Limit ...
 - IDTR (Interrupt Descriptor Table Register)
 - Handler Address, Ring Level ...
 - TR (Task Register)
 - TSS

- PC Architecture
- **Memory**
- Execution
- PC Emulation

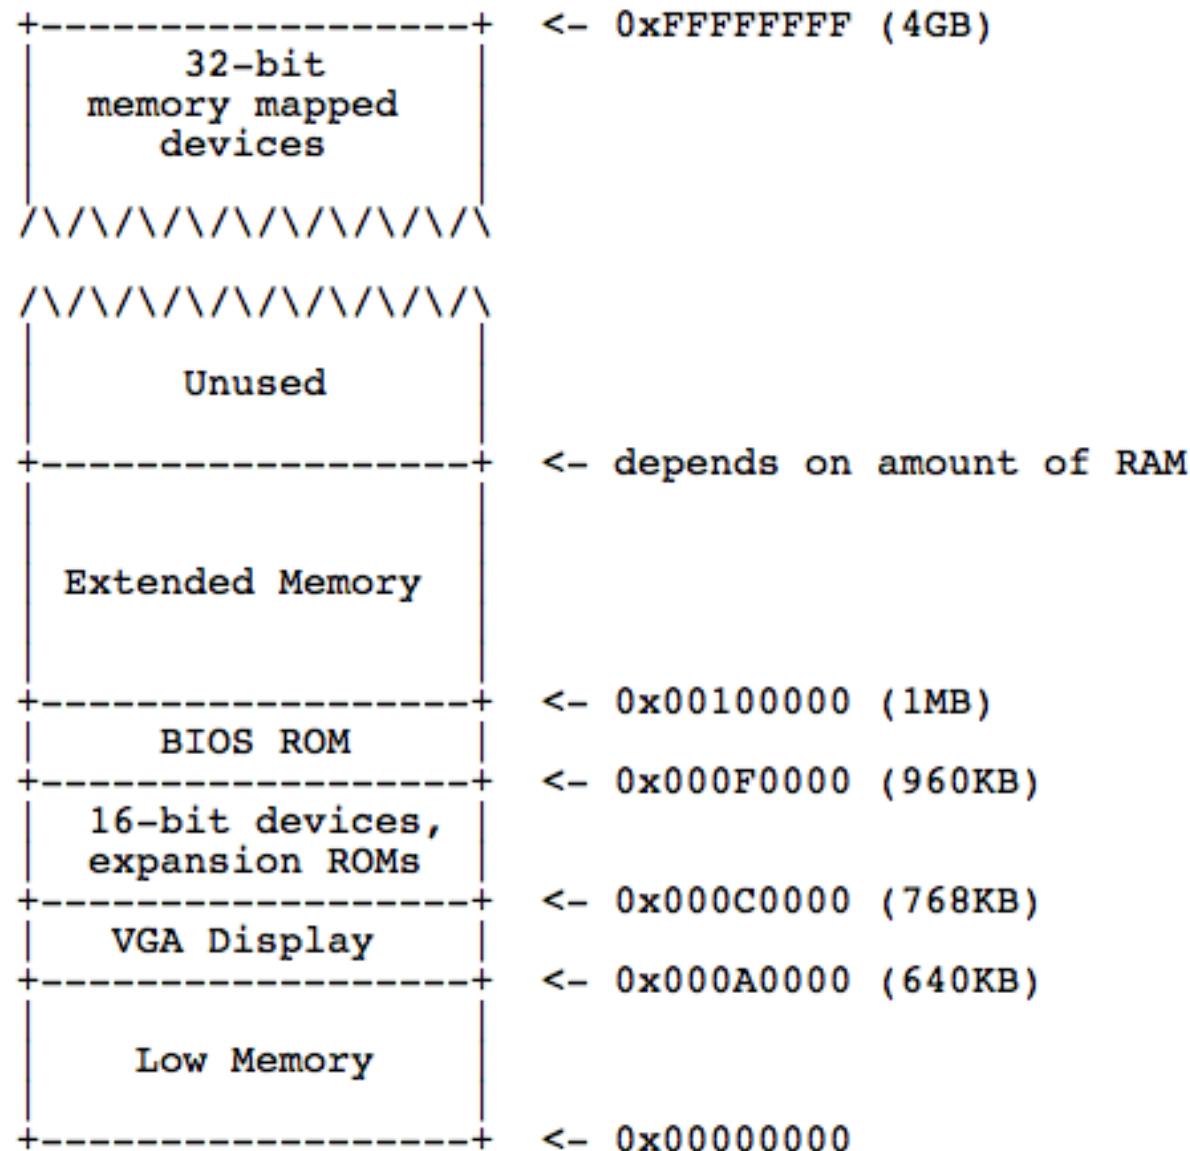
Memory Model

- 8086: 16-bits microprocessor
 - Real Mode: $\text{physical addr} = 16 * \text{segment} + \text{offset}$
 - Space: 64KB
- external address to 20-bits
 - Space: 1MB
 - The extra 4 bits come *segment registers*
 - CS: code segment, for EIP
 - SS: stack segment, for SP and BP
 - DS: data segment for load/store via other registers
 - ES: another data segment, destination for string ops
 - e.g. $CS=FOOO\ IP=FFFF0 \Rightarrow ADDR: FFFF0$

Memory Model

- 80386: 32-bit data and bus addresses
 - Protected Mode
- Now: the transition to 64-bit addresses
- Backwards compatibility:
 - Boots in 16-bit real mode, and switches to 32-bit protected mode
 - See: “*boot/boot.S*”

Memory Layout



I/O space

```
#define DATA_PORT      0x378
#define STATUS_PORT    0x379
#define BUSY 0x80
#define CONTROL_PORT  0x37A
#define STROBE 0x01

void
lpt_putc(int c)
{
    /* wait for printer to consume previous byte */
    while((inb(STATUS_PORT) & BUSY) == 0)
        ;

    /* put the byte on the parallel lines */
    outb(DATA_PORT, c);

    /* tell the printer to look at the data */
    outb(CONTROL_PORT, STROBE);
    outb(CONTROL_PORT, 0);
}
```

```
static __inline uint8_t inb(int port)
{
    uint8_t data;
    __asm __volatile("inb %w1,%0"
                    : "=a" (data) : "d" (port));
    return data;
}
```

```
/* wait for printer to consume previous byte */
while((inb(STATUS_PORT) & BUSY) == 0)
```

```
/* put the byte on the parallel lines */
```

```
outb(DATA_PORT, c);
```

```
/* tell the printer to look at the data */
```

```
outb(CONTROL_PORT, STROBE);
```

```
outb(CONTROL_PORT, 0);
```

```
static __inline void
    outb(int port, uint8_t data)
{
    __asm __volatile("outb %0,%w1"
                    : : "a" (data), "d" (port));
}
```

Memory-mapped I/O

- Use normal addresses
 - No need for special instructions
 - No 1024 limit
 - System controller routes to device
- Works like “magic” Memory
 - Addressed and accessed like memory
 - But does not behave like memory
 - Reads and writes have “**side effects**”
 - Read result can change due to external events
 - Recall: the “volatile” keyword

- PC Architecture
- Memory
- **Execution**
- PC Emulation

XV6 BOOTING CODE

```
# Start the first CPU: switch to 32-bit protected mode, jump into C.  
# The BIOS loads this code from the first sector of the hard disk into  
# memory at physical address 0x7c00 and starts executing in real mode  
# with %cs=0 %ip=7c00.  
  
.code16                                # Assemble for 16-bit mode  
.globl start  
start:  
    cli                                # BIOS enabled interrupts; disable  
  
    # Zero data segment registers DS, ES, and SS.  
    xorw    %ax,%ax                      # Set %ax to zero  
    movw    %ax,%ds                      # -> Data Segment  
    movw    %ax,%es                      # -> Extra Segment  
    movw    %ax,%ss                      # -> Stack Segment  
  
    # Physical address line A20 is tied to zero so that the first PCs  
    # with 2 MB would run software that assumed 1 MB. Undo that.  
seta20.1:  
    inb    $0x64,%al                      # Wait for not busy  
    testb   $0x2,%al  
    jnz     seta20.1  
  
    movb    $0xd1,%al                      # 0xd1 -> port 0x64  
    outb    %al,$0x64  
  
seta20.2:  
    inb    $0x64,%al                      # Wait for not busy  
    testb   $0x2,%al  
    jnz     seta20.2  
  
    movb    $0xdf,%al                      # 0xdf -> port 0x60  
    outb    %al,$0x60
```

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdtdesc
movl    %cr0, %eax
orl    $CR0_PE, %eax
movl    %eax, %cr0

//PAGEBREAK!
# Complete transition to 32-bit protected mode by using long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp   $(SEG_KCODE<<3), $start32

.code32 # Tell assembler to generate 32-bit code now.
start32:
# Set up the protected-mode data segment registers
movw   $(SEG_KDATA<<3), %ax      # Our data segment selector
movw   %ax, %ds                   # -> DS: Data Segment
movw   %ax, %es                   # -> ES: Extra Segment
movw   %ax, %ss                   # -> SS: Stack Segment
movw   $0, %ax                   # Zero segments not ready for use
movw   %ax, %fs                   # -> FS
movw   %ax, %gs                   # -> GS

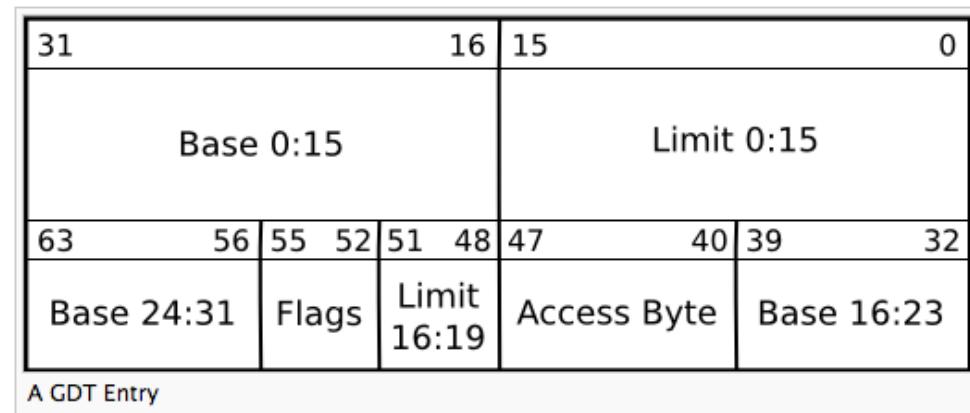
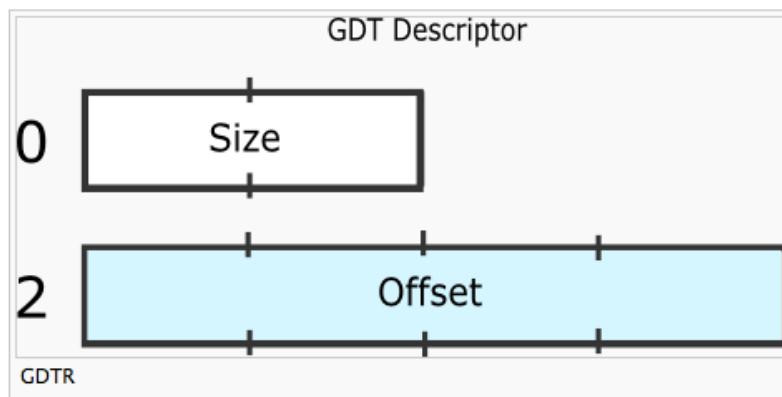
# Set up the stack pointer and call into C.
movl   $start, %esp
call   bootmain
```

```

# Bootstrap GDT
.p2align 2                                # force 4 byte alignment
gdt:
SEG_NULLASM                                # null seg
SEG_ASM(STA_X|STA_R, 0x0, 0xffffffff)      # code seg
SEG_ASM(STA_W, 0x0, 0xffffffff)              # data seg

gdtdesc:
.word  (gdtdesc - gdt - 1)                  # sizeof(gdt) - 1
.long  gdt                                  # address gdt

```



```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}
```

```
void
readsect(void *dst, uint offset)
{
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    outb(0x1F3, offset);
    outb(0x1F4, offset >> 8);
    outb(0x1F5, offset >> 16);
    outb(0x1F6, (offset >> 24) | 0xE0);
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}

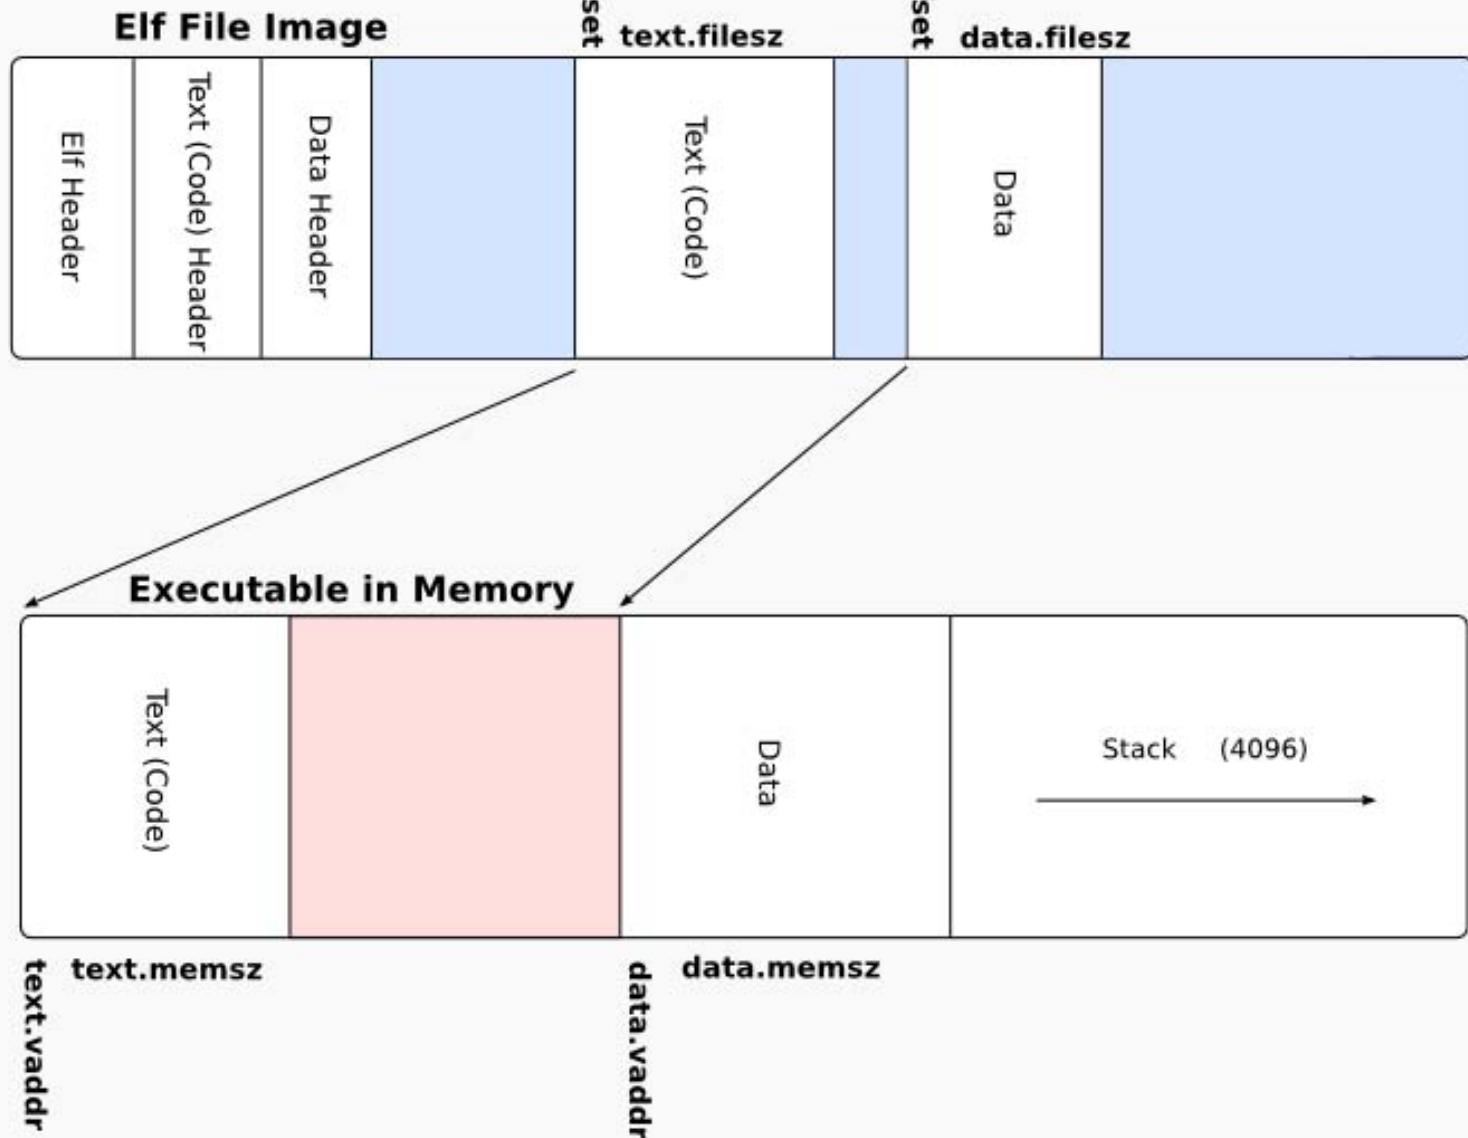
// Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
// Might copy more than asked.
void
readseg(uchar* pa, uint count, uint offset)
{
    uchar* epa;

    epa = pa + count;

    // Round down to sector boundary.
    pa -= offset % SECTSIZE;

    // Translate from bytes to sectors; kernel starts at sector 1.
    offset = (offset / SECTSIZE) + 1;

    // If this is too slow, we could read lots of sectors at a time.
    // We'd write more to memory than asked, but it doesn't matter --
    // we load in increasing order.
    for(; pa < epa; pa += SECTSIZE, offset++)
        readsect(pa, offset);
}
```



Executable image and elf binary can being mapped onto each other

```
# By convention, the _start symbol specifies the ELF entry point.
# Since we haven't set up virtual memory yet, our entry point is
# the physical address of 'entry'.
.globl _start
_start = V2P_W0(entry)

# Entering xv6 on boot processor, with paging off.
.globl entry
entry:
    # Turn on page size extension for 4Mbyte pages
    movl    %cr4, %eax
    orl    $(CR4_PSE), %eax
    movl    %eax, %cr4
    # Set page directory
    movl    $(V2P_W0(entrypgdir)), %eax
    movl    %eax, %cr3
    # Turn on paging.
    movl    %cr0, %eax
    orl    $(CR0_PG|CR0_WP), %eax
    movl    %eax, %cr0

    # Set up the stack pointer.
    movl $(stack + KSTACKSIZE), %esp

    # Jump to main(), and switch to executing at
    # high addresses. The indirect call is needed because
    # the assembler produces a PC-relative instruction
    # for a direct jump.
    mov $main, %eax
    jmp *%eax

.comm stack, KSTACKSIZE
```

```
static void startothers(void);
static void mpmain(void) __attribute__((noreturn));
extern pde_t *kpgdir;
extern char end[]; // first address after kernel loaded from ELF file

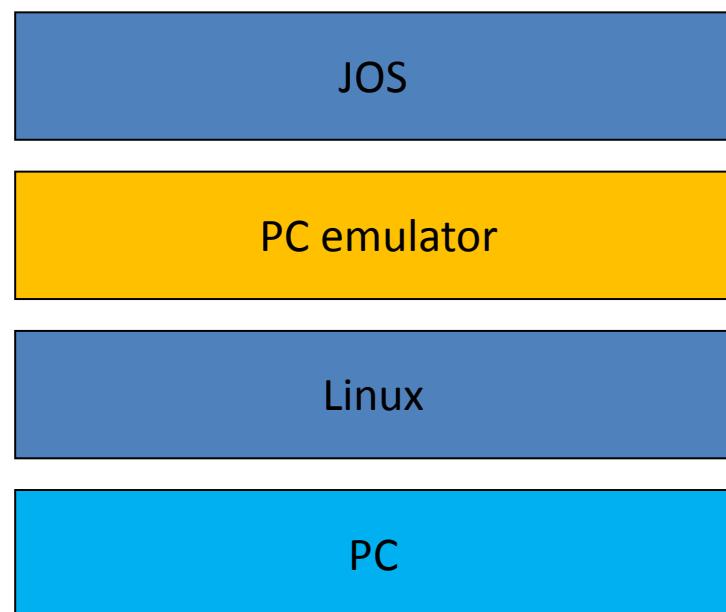
// Bootstrap processor starts running C code here.
// Allocate a real stack and switch to it, first
// doing some setup required for memory allocator to work.
int
main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // collect info about this machine
    lapicinit(mpbcpu());
    seginit(); // set up segments
    cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
    picinit(); // interrupt controller
    ioapicinit(); // another interrupt controller
    consoleinit(); // I/O devices & their interrupts
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    binit(); // buffer cache
    fileinit(); // file table
    iinit(); // inode cache
    ideinit(); // disk
    if(!ismp)
        timerinit(); // uniprocessor timer
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    // Finish setting up this processor in mpmain.
    mpmain();
}
```

- PC Architecture
- Memory
- Execution
- PC Emulation

Development using PC emulator

- Bochs PC emulator
 - **does** what a real PC **does**
 - Only implemented in software!

Runs like a normal program
on “host” operating system



Emulation of CPU

```
for (;;) {
    read_instruction();
    switch (decode_instruction_opcode()) {
        case OPCODE_ADD: OPCODE_ADD
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] + regs[src];
            break;
        case OPCODE_SUB:
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] - regs[src];
            break;
        ...
    }
    eip += instruction_length;
}
```

Emulation of Memory

```
int32_t regs[8];
#define REG_EAX 1;
#define REG_EBX 2;
#define REG_ECX 3;
...
int32_t eip;
int16_t segregs[4];
...

char mem[256*1024*1024];
```

Emulation of x86 Memory

```
uint8_t read_byte(uint32_t phys_addr) {
    if (phys_addr < LOW_MEMORY)
        return low_mem[phys_addr];
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        return rom_bios[phys_addr - 960*KB];
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        return ext_mem[phys_addr-1*MB];
    } else ...
}

void write_byte(uint32_t phys_addr, uint8_t val) {
    if (phys_addr < LOW_MEMORY)
        low_mem[phys_addr] = val;
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        ; /* ignore attempted write to ROM! */
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        ext_mem[phys_addr-1*MB] = val;
    } else ...
}
```

Low Memory

Extended Memory

Emulation of Devices

- Hard disk: using a file of the host
- VGA display: draw in a host window
- Keyboard: host's keyboard API
- Clock chip: host's clock
- Etc.

Why Emulator

- OS Test and Debug
- Increase Utilization
- Just as Why IBM's Virtualization
 - IBM's M44/44X, in 1960s

Thanks

Next time:

Intro to PC booting & OS structure

Homework

- What is the difference between user-level ISA and system-level ISA?
- What is memory addressing mode? How many modes are there (please describe them)? Why not just use one?
- Use your own word (and figures, if you want) to describe the process from power-on to BIOS end (just before kernel starts)

Xv6 Boot & OS Structure

Yubin Xia & Rong Chen

Some slides are based on Frans Kaashoek's slides
kaashoek@mit.edu

Review

- PC Architecture
 - The von Neumann Model
- X86
 - Modes: real mode / protected mode / SMM / ...
 - Registers: EFLAGS, CR[0-4], GDTR, IDTR, ...
 - Memory model: 16-bit / 20-bit / 32-bit / 36-bit / 64-bit
 - Memory-mapped I/O
 - Low memory: 0-640 KB, extended memory: above 1MB
 - Booting
 - BIOS: 0xFFFF0 (what is seg, what is offset?)
 - First instruction: 0x7c00 (what instruction?)

XV6 BOOTING CODE

Makefile

```
82 xv6.img: bootblock kernel fs.img
83     dd if=/dev/zero of=xv6.img count=10000
84     dd if=bootblock of=xv6.img conv=notrunc
85     dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
92 bootblock: bootasm.S bootmain.c
93     $(CC) $(CFLAGS) -fno-pic -O -nostdinc -I. -c bootmain.c
94     $(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c bootasm.S
95     $(LD) $(LDFLAGS) -N -e start -Ttext 0x7C00 -o bootblock.o bootasm.o bootmain.o
96     $(OBJDUMP) -S bootblock.o > bootblock.asm
97     $(OBJCOPY) -S -O binary -j .text bootblock.o bootblock
98     ./sign.pl bootblock
```

Sign.pl

```
13
14 $buf .= "\0" x (510-$n);
15 $buf .= "\x55\xAA";
16
```

```
# Start the first CPU: switch to 32-bit protected mode, jump into C.  
# The BIOS loads this code from the first sector of the hard disk into  
# memory at physical address 0x7c00 and starts executing in real mode  
# with %cs=0 %ip=7c00.  
  
.code16                                # Assemble for 16-bit mode  
.globl start  
start:  
    cli                                # BIOS enabled interrupts; disable  
  
    # Zero data segment registers DS, ES, and SS.  
    xorw    %ax,%ax                      # Set %ax to zero  
    movw    %ax,%ds                      # -> Data Segment  
    movw    %ax,%es                      # -> Extra Segment  
    movw    %ax,%ss                      # -> Stack Segment  
  
    # Physical address line A20 is tied to zero so that the first PCs  
    # with 2 MB would run software that assumed 1 MB. Undo that.  
seta20.1:  
    inb    $0x64,%al                      # Wait for not busy  
    testb   $0x2,%al  
    jnz     seta20.1  
  
    movb    $0xd1,%al                      # 0xd1 -> port 0x64  
    outb    %al,$0x64  
  
seta20.2:  
    inb    $0x64,%al                      # Wait for not busy  
    testb   $0x2,%al  
    jnz     seta20.2  
  
    movb    $0xdf,%al                      # 0xdf -> port 0x60  
    outb    %al,$0x60
```

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdtdesc
movl    %cr0, %eax
orl    $CR0_PE, %eax
movl    %eax, %cr0

//PAGEBREAK!
# Complete transition to 32-bit protected mode by using long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp   $(SEG_KCODE<<3), $start32

.code32 # Tell assembler to generate 32-bit code now.
start32:
# Set up the protected-mode data segment registers
movw   $(SEG_KDATA<<3), %ax      # Our data segment selector
movw   %ax, %ds                   # -> DS: Data Segment
movw   %ax, %es                   # -> ES: Extra Segment
movw   %ax, %ss                   # -> SS: Stack Segment
movw   $0, %ax                   # Zero segments not ready for use
movw   %ax, %fs                   # -> FS
movw   %ax, %gs                   # -> GS

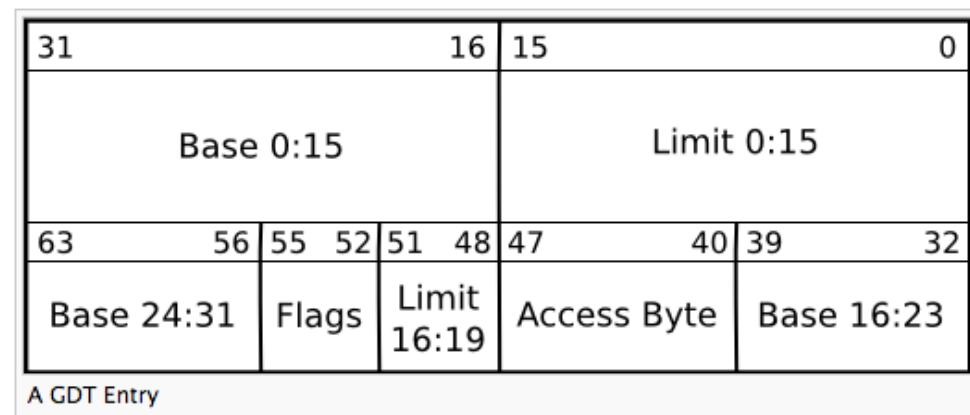
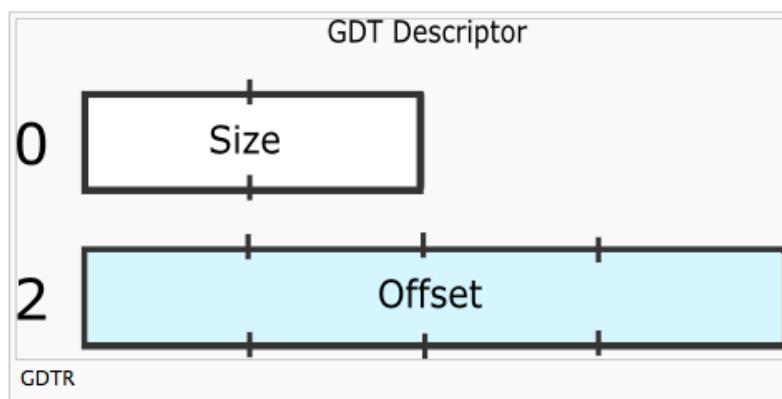
# Set up the stack pointer and call into C.
movl   $start, %esp
call   bootmain
```

```

# Bootstrap GDT
.p2align 2                                # force 4 byte alignment
gdt:
SEG_NULLASM                                # null seg
SEG_ASM(STA_X|STA_R, 0x0, 0xffffffff)      # code seg
SEG_ASM(STA_W, 0x0, 0xffffffff)              # data seg

gdtdesc:
.word  (gdtdesc - gdt - 1)                  # sizeof(gdt) - 1
.long  gdt                                  # address gdt

```



```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*)(void))(elf->entry);
    entry();
}
```

```
void
readsect(void *dst, uint offset)
{
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    outb(0x1F3, offset);
    outb(0x1F4, offset >> 8);
    outb(0x1F5, offset >> 16);
    outb(0x1F6, (offset >> 24) | 0xE0);
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}

// Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
// Might copy more than asked.
void
readseg(uchar* pa, uint count, uint offset)
{
    uchar* epa;

    epa = pa + count;

    // Round down to sector boundary.
    pa -= offset % SECTSIZE;

    // Translate from bytes to sectors; kernel starts at sector 1.
    offset = (offset / SECTSIZE) + 1;

    // If this is too slow, we could read lots of sectors at a time.
    // We'd write more to memory than asked, but it doesn't matter --
    // we load in increasing order.
    for(; pa < epa; pa += SECTSIZE, offset++)
        readsect(pa, offset);
}
```

ELF Header :

```
#define EI_NIDENT 16
typedef struct {
    unsigned char e_ident [EI_NIDENT];
    Elf32_Half e_type;
    Elf32_Half e_machine ;
    Elf32_Word e_version ;
    Elf32_Addr e_entry ;
    Elf32_Off e_phoff;
    Elf32_Off e_shoff;
    Elf32_Word e_flags ;
    Elf32_Half e_ehsize ;
    Elf32_Half e_phentsize ;
    Elf32_Half e_phnum ;
    Elf32_Half e_shentsize ;
    Elf32_Half e_shnum ;
    Elf32_Half e_shstrndx ;
}Elf32_Ehdr ;
```

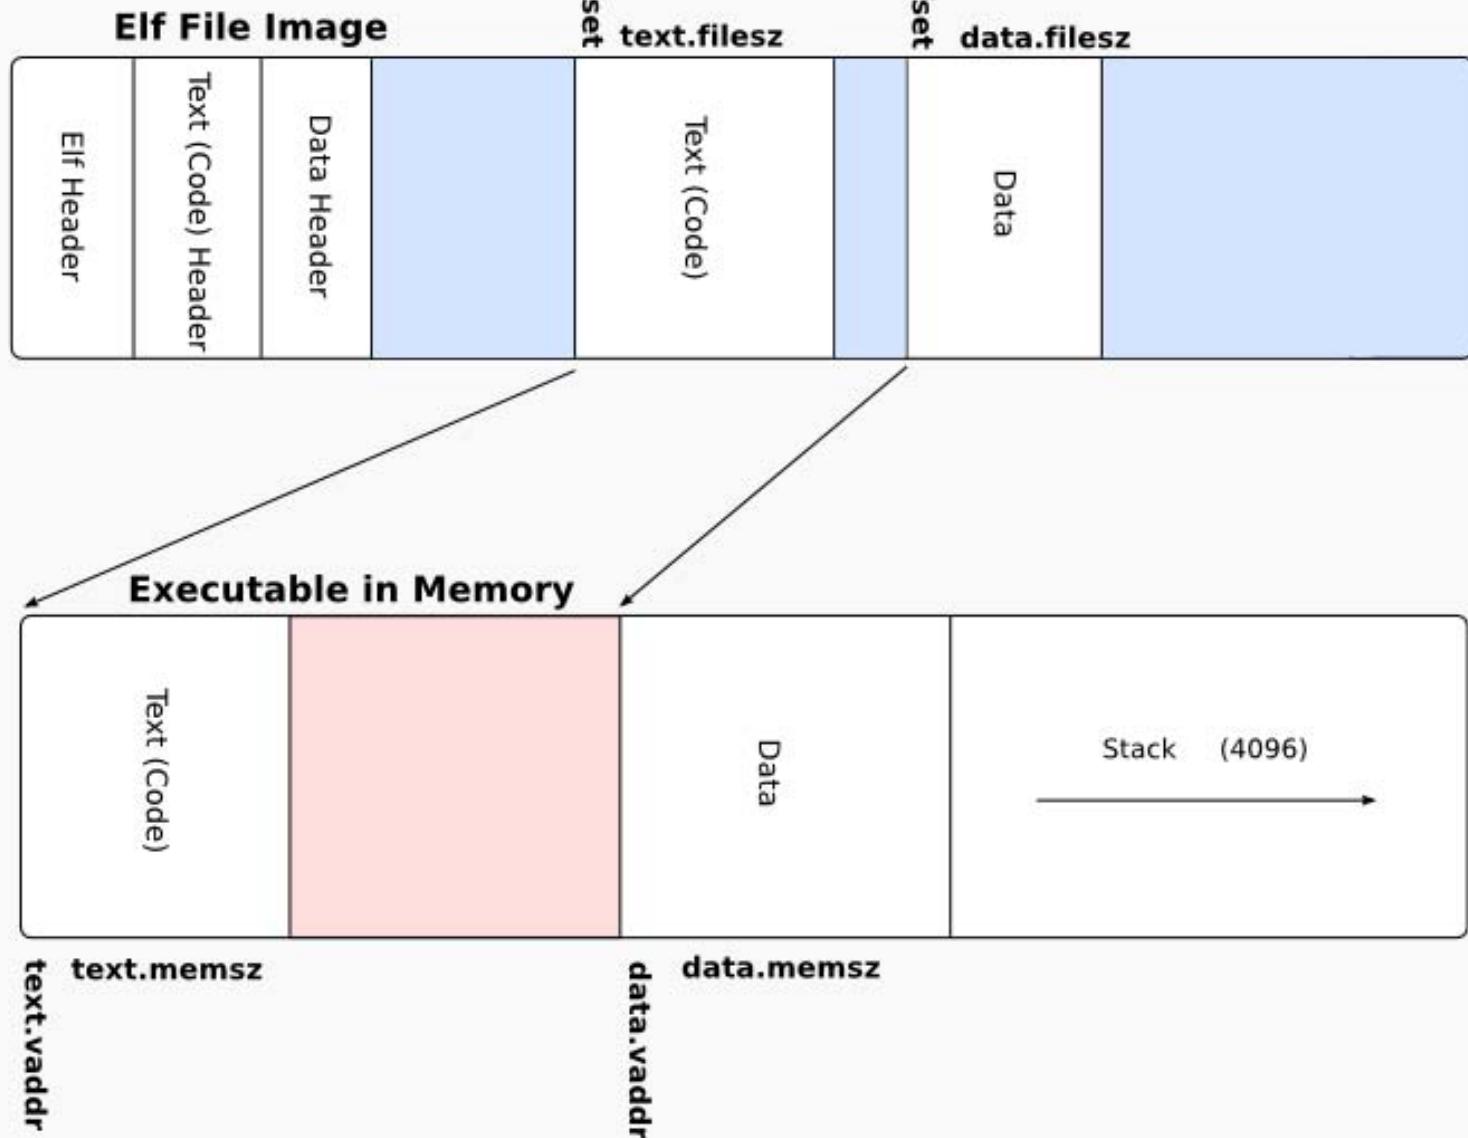
ELF header resides at the beginning and holds a road map describing the files organization

e_machine : specifies the required architecture :**0** = No machine ... **3** = Intel
80386 **4** = Motorola 68000 ...**8** = MIPS RS3000

e_entry : the virtual address to which the system first transfers control, when starting the process. If the file has no associated entry point, this member holds zero.

e_shoff : holds the **section header table's file offset** in bytes. If the file has no section header table, this member holds zero.

Name	Size	Alignment	Purpose
EIf_32_Addr	4	4	Unsigned program address
EIf32_Half	2	2	Unsigned medium integer
EIf32_Off	4	4	Unsigned file offset
EIf32_Sword	4	4	Signed large integer
EIf32_Word	4	4	Unsigned large integer
Unsigned char	1	1	Unsigned small integer



Executable image and elf binary can being mapped onto each other

```
# By convention, the _start symbol specifies the ELF entry point.  
# Since we haven't set up virtual memory yet, our entry point is  
# the physical address of 'entry'.  
.globl _start  
_start = V2P_W0(entry)  
  
# Entering xv6 on boot processor, with paging off.  
.globl entry  
entry:  

```

entry.S

```
static void startothers(void);
static void mpmain(void) __attribute__((noreturn));
extern pde_t *kpgdir;
extern char end[]; // first address after kernel loaded from ELF file

// Bootstrap processor starts running C code here.
// Allocate a real stack and switch to it, first
// doing some setup required for memory allocator to work.
int
main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // collect info about this machine
    lapicinit(mpbcpu());
    seginit(); // set up segments
    cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
    picinit(); // interrupt controller
    ioapicinit(); // another interrupt controller
    consoleinit(); // I/O devices & their interrupts
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    binit(); // buffer cache
    fileinit(); // file table
    iinit(); // inode cache
    ideinit(); // disk
    if(!ismp)
        timerinit(); // uniprocessor timer
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    // Finish setting up this processor in mpmain.
    mpmain();
}
```

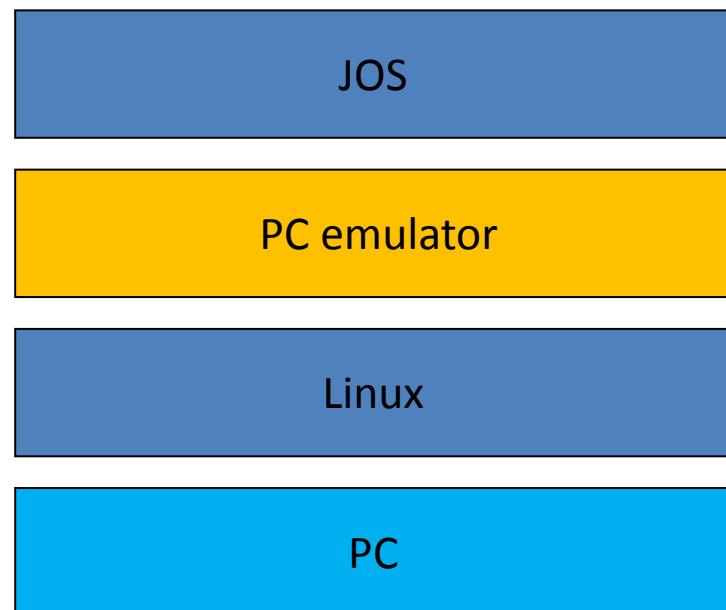
PC PROGRAMMING

- PC Architecture
- Memory
- Execution
- PC Emulation

Development using PC emulator

- Bochs PC emulator
 - **does** what a real PC **does**
 - Only implemented in software!

Runs like a normal program
on “host” operating system



Emulation of CPU

```
for (;;) {
    read_instruction();
    switch (decode_instruction_opcode()) {
        case OPCODE_ADD: OPCODE_ADD
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] + regs[src];
            break;
        case OPCODE_SUB:
            int src = decode_src_reg();
            int dst = decode_dst_reg();
            regs[dst] = regs[dst] - regs[src];
            break;
        ...
    }
    eip += instruction_length;
}
```

Emulation of Memory

```
int32_t regs[8];
#define REG_EAX 1;
#define REG_EBX 2;
#define REG_ECX 3;
...
int32_t eip;
int16_t segregs[4];
...

char mem[256*1024*1024];
```

Emulation of x86 Memory

```
uint8_t read_byte(uint32_t phys_addr) {
    if (phys_addr < LOW_MEMORY)
        return low_mem[phys_addr];
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        return rom_bios[phys_addr - 960*KB];
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        return ext_mem[phys_addr-1*MB];
    } else ...
}

void write_byte(uint32_t phys_addr, uint8_t val) {
    if (phys_addr < LOW_MEMORY)
        low_mem[phys_addr] = val;
    else if (phys_addr >= 960*KB && phys_addr < 1*MB)
        ; /* ignore attempted write to ROM! */
    else if (phys_addr >= 1*MB && phys_addr < 1*MB+EXT_MEMORY) {
        ext_mem[phys_addr-1*MB] = val;
    } else ...
}
```

Low Memory

Extended Memory

Emulation of Devices

- Hard disk: using a file of the host
- VGA display: draw in a host window
- Keyboard: host's keyboard API
- Clock chip: host's clock
- Etc.

Why Emulator

- OS Test and Debug
- Increase Utilization
- Just as Why IBM's Virtualization
 - IBM's M44/44X, in 1960s

OPERATING SYSTEM ARCHITECTURE

Operating System Design and Implementation

Design and Implementation of OS not “solvable”, but some approaches have proven successful

- Internal structure of different Operating Systems can vary widely

- Start by defining goals and specifications

- Affected by choice of hardware, type of system

Remember “Worse is better design”

Operating System Design and Implementation

User goals and System goals

User goals

operating system should be convenient to use, easy to learn, reliable, safe, and fast

System goals

operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient

Operating System Design and Implementation (Cont.)

Important principle to separate

Policy: What will be done?

Mechanism: How to do it?

Mechanisms determine how to do something

Policies decide what will be done

The separation of policy from mechanism is a very important principle

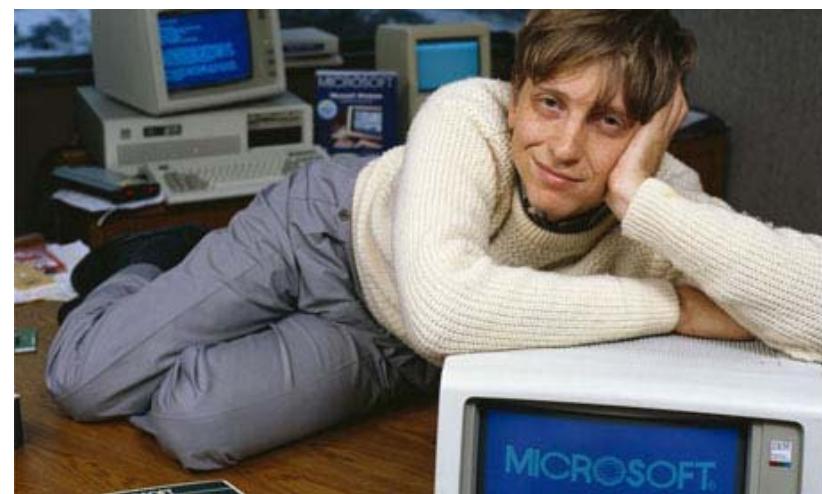
It allows maximum flexibility if policy decisions are to be changed later

Simple Structure

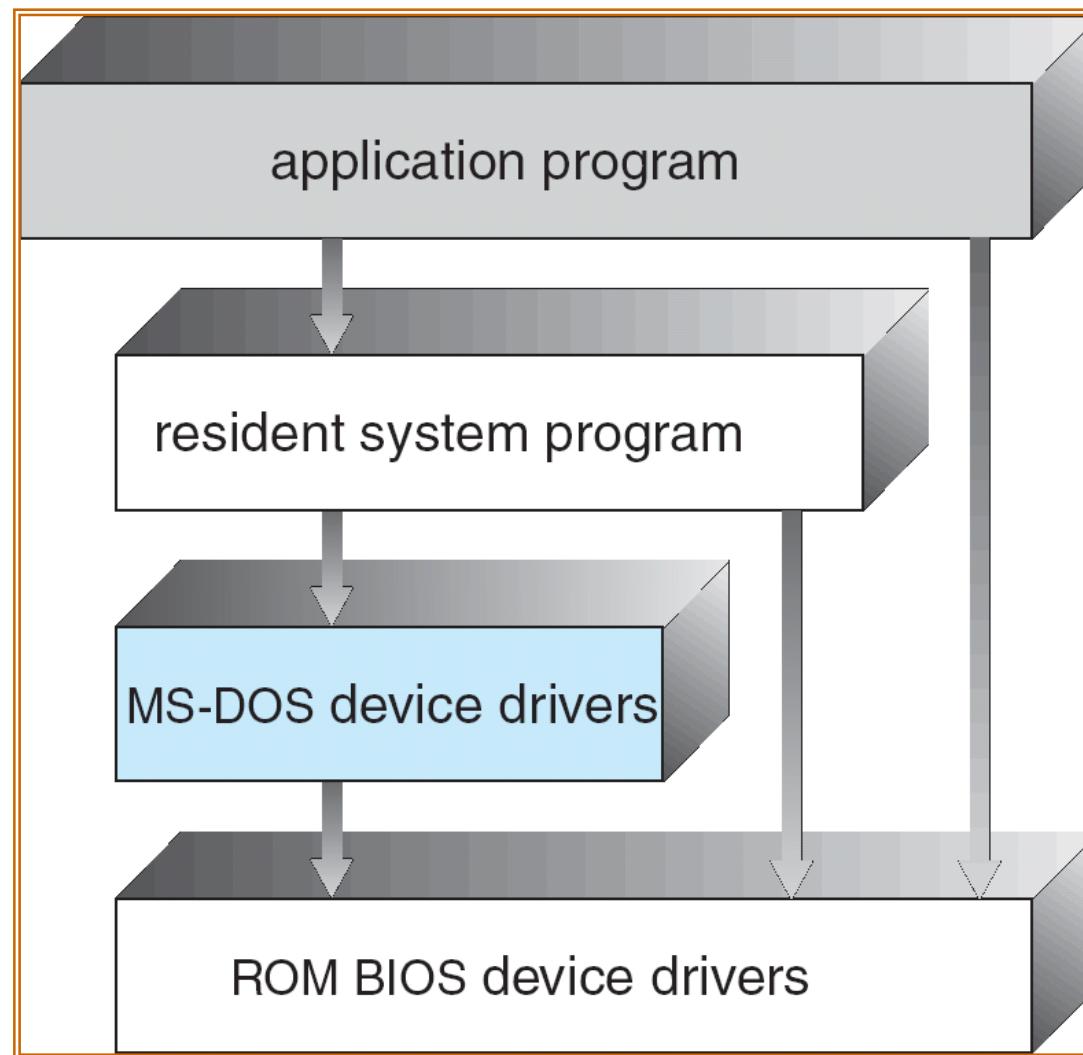
MS-DOS – written to provide the most functionality in the least space (1981~1994)

Not divided into modules

Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated



MS-DOS Layer Structure



Layered Approach

The operating system is divided into a number of layers (levels)

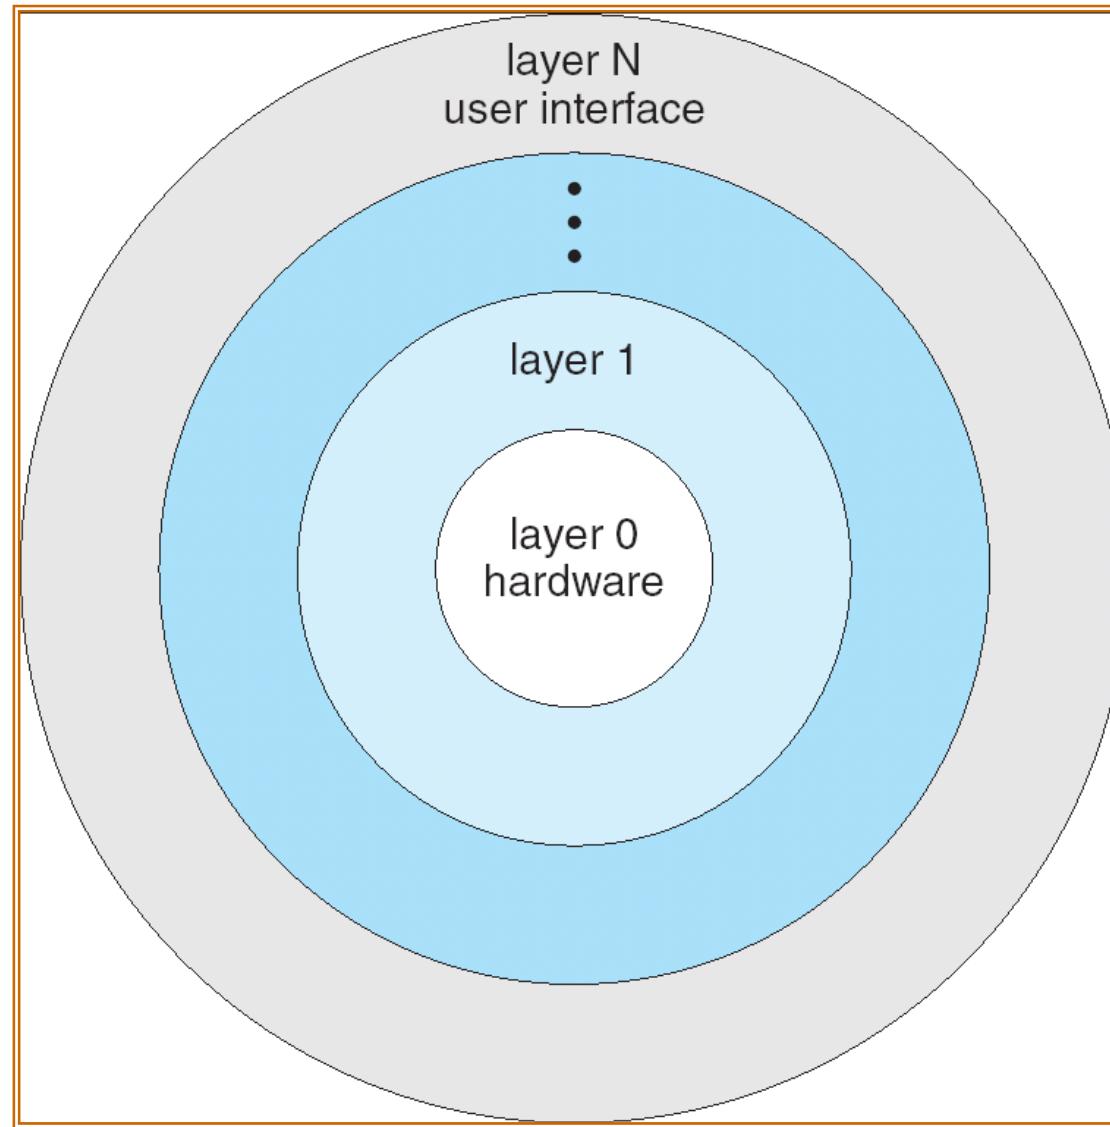
- Each built on top of lower layers.

- The bottom layer (layer 0), is the hardware;

- The highest (layer N) is the user interface.

With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers

Layered Operating System



UNIX

UNIX – limited by hardware functionality

the original UNIX operating system had limited structuring.

The UNIX OS consists of two separable parts

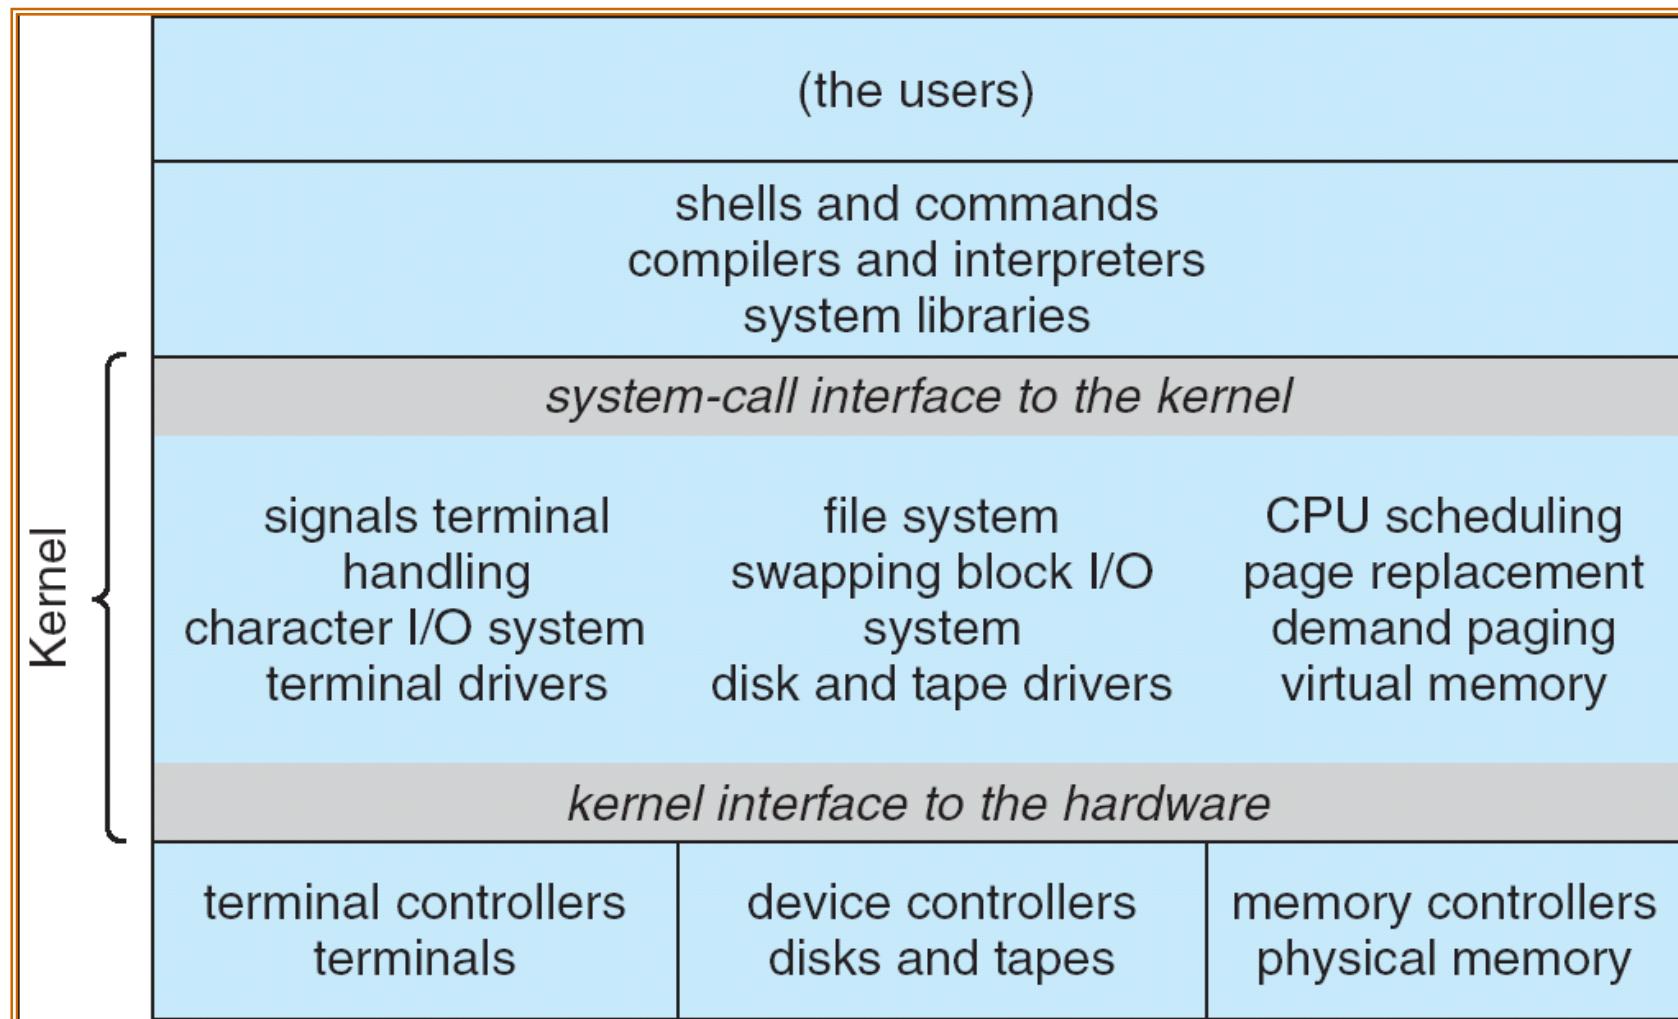
Systems programs

The kernel

Consists of everything below the system-call interface and above the physical hardware

Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level

UNIX System Structure



Microkernel System Structure

Moves as much from the kernel into “user” space

Communication takes place between user modules using message passing

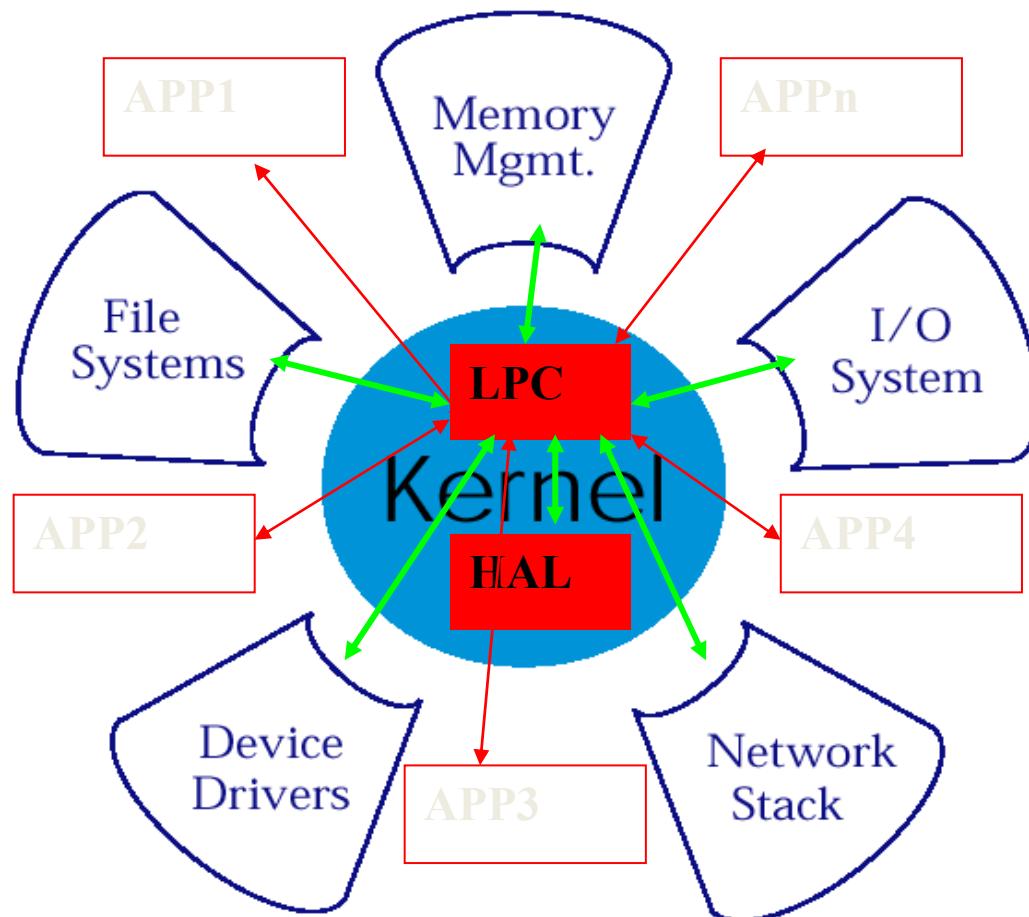
Benefits:

- Easier to extend a microkernel
- Easier to port the operating system to new architectures
- More reliable (less code is running in kernel mode)
- More secure

Detriments:

- Performance overhead of user space to kernel space communication

Microkernel Structure & Instances



Exokernel Structure

Overview

let the kernel allocate the physical resources of the machine to multiple application programs, and

let each program decide what to do with these resources.

The program can link to an operating system library (libOS) that implements OS abstractions.

Exokernel: An operating system architecture for application-level resource management, SOSP'95

Exokernel & Corey

Exokernel [SOSP'95]

Kernel classification

Motivation

Design and Implementation

Evaluation

Corey [OSDI'08]

Multi-core era and scalability problem

Sharing and Address range

Summery

Kernels

Kernel classification (wikipedia)

Kernel (architecture)	Monolithic	Microkernel	Exokernel	hybrid
--------------------------	------------	-------------	-----------	--------

Monolithic kernel

A **monolithic kernel** is an operating system architecture where the entire operating system is working in kernel space and is alone in supervisor mode.

Linux, BSD...

Kernels

Microkernel

is the near-minimum amount of software that can provide the mechanisms needed to implement an operating system (OS).

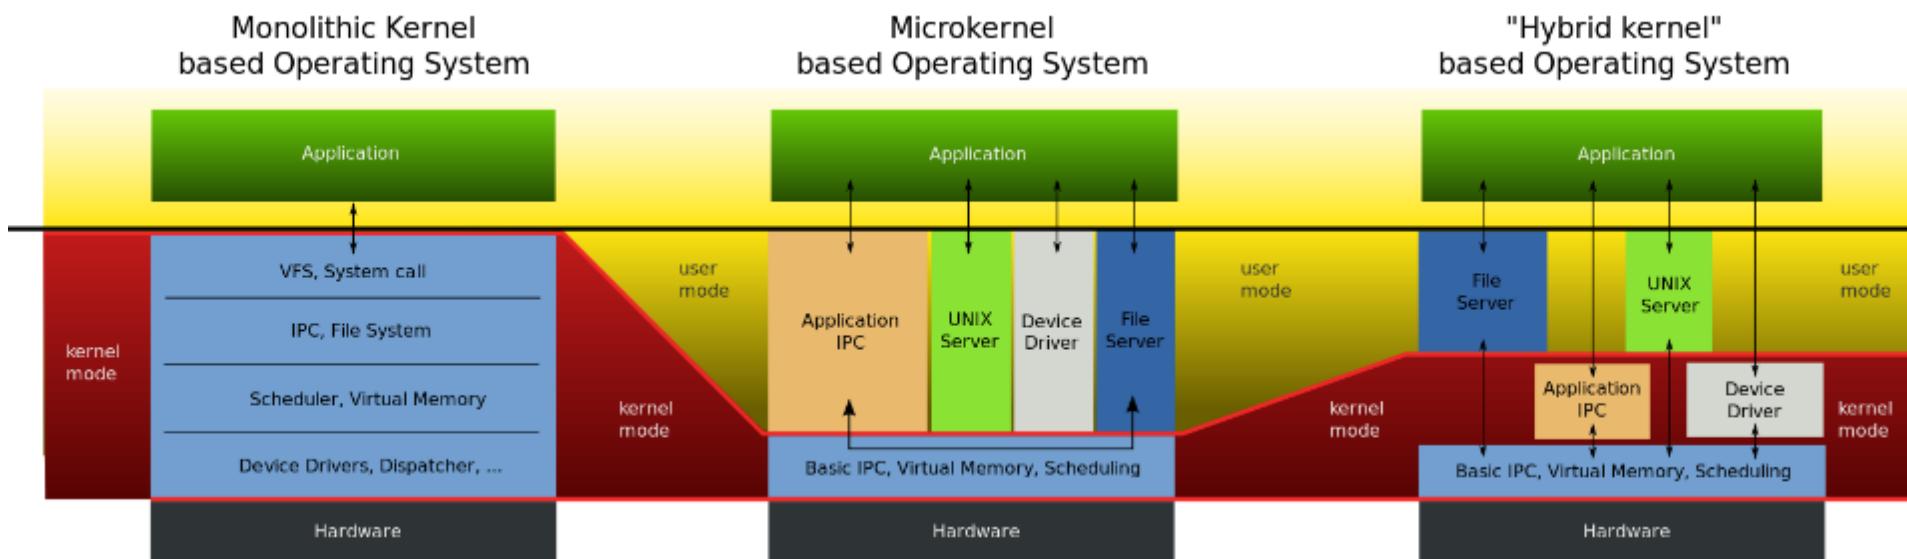
low-level address space management, thread management, and inter-process communication (IPC) and so on ...

Mach, L4 kernel

Hybrid kernel

is a kernel architecture based on combining aspects of microkernel and monolithic kernel architectures used in computer operating systems.

Windows NT kernel



Questions

What's the shortcoming of those kernels

Hint: that's the reason to create exokernel

Motivation

Exokernel's Motivation

In traditional operating systems, only privileged servers and the kernel can manage system resources

Un-trusted applications are required to interact with the hardware via some abstraction model

File systems for disk storage, *virtual address spaces* for memory, etc.

Application demands vary widely!!

An interface designed to accommodate every application must anticipate **all possible needs**

Exokernel's Idea

Give un-trusted applications as **much control** over physical resources as possible

To force as **few abstraction** as possible on developers, enabling them to make as **many decisions** as possible about hardware abstractions.

Let *the kernel* allocate the basic physical resources of the machine

Let *each program* decide what to do with these resources

Exokernel separate *protection* from *management*

They protect resources but delegate management to application

Library Operating System

Most programs to be linked with *libraries* instead of communicating with the exokernel directly

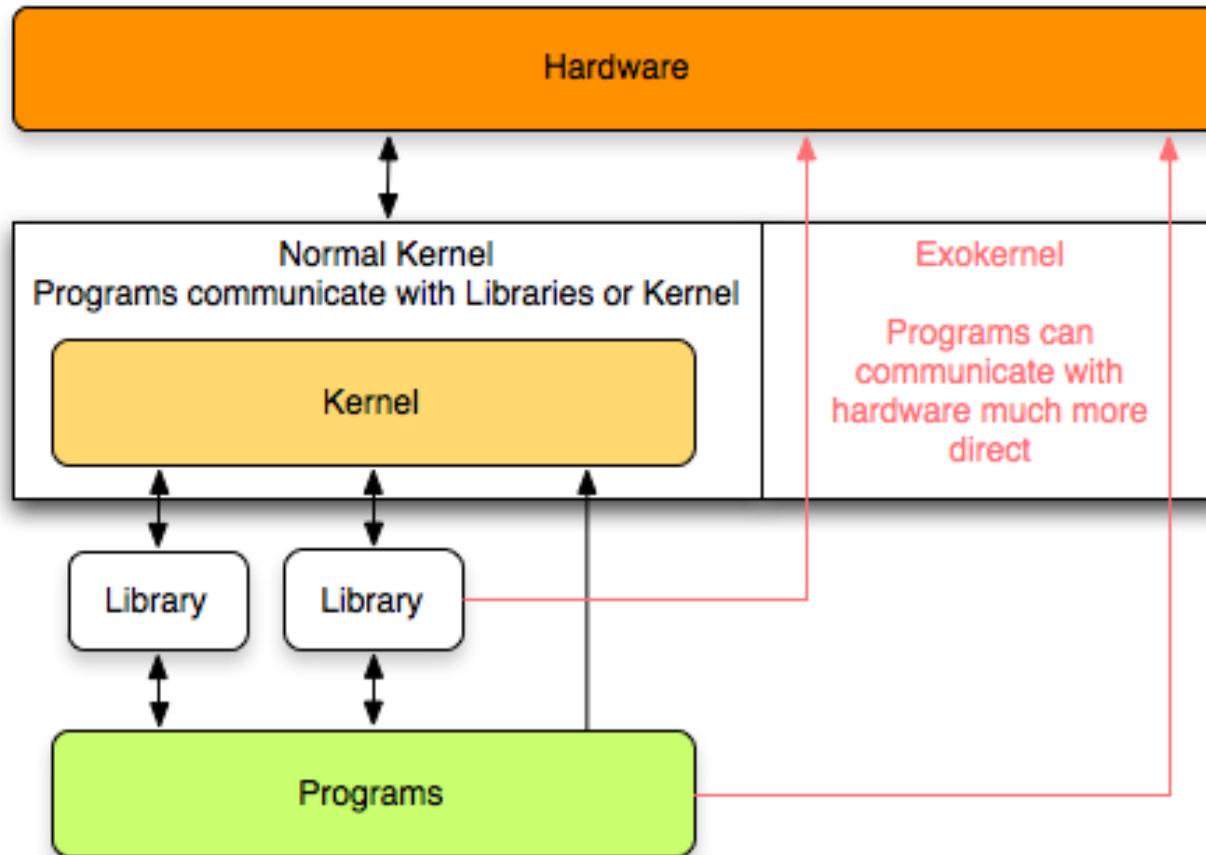
The libraries hide low-level resources

An applications can choose the library which best suits its needs, or even build its own

The kernel only ensures that the requested resource is free, and **the application** is allowed to access it.

Allow programmer to implement custom abstractions, omit unnecessary ones, most commonly to improve performance

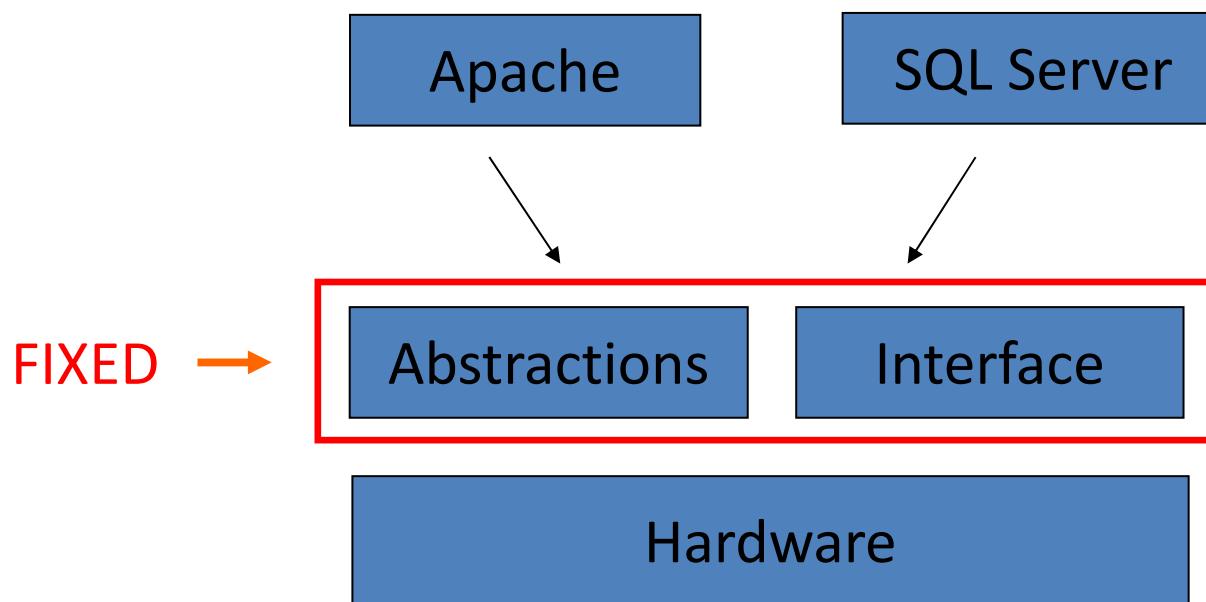
Exokernel Example 1



Exokernel give more direct access to the hardware, thus removing most abstractions

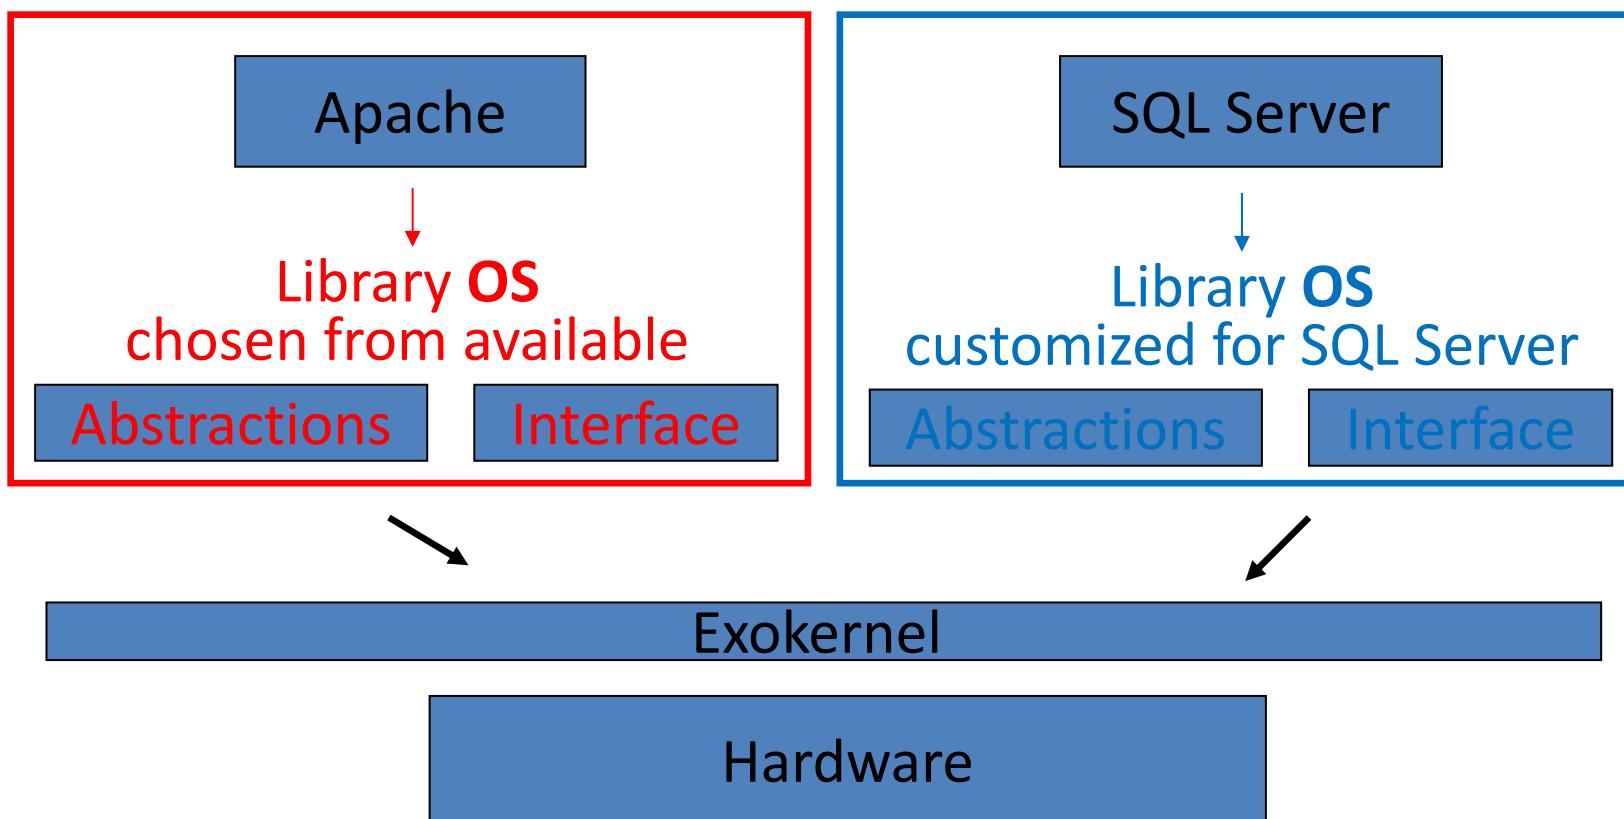
Exokernel Example 2

Traditional OS



Exokernel Example 2

Exokernel + Library OS



Exokernel Design Challenge

Kernel's new role

- Tracking ownership of resources

- Ensuring resource protection

- Revoking resource access

Three techniques

- Secure binding**

- Visible revocation**

- Abort protocol**

Secure binding

It is a protection mechanism that decouples authorization from actual use of a resource

Can improve performance

The protection checks involved in enforcing a secure binding are expressed in terms of simple operations that the kernel can implement quickly

A secure binding performs authorization only at bind time, which allows management to be decoupled from protection

Three techniques

Hardware mechanism, software caching, and downloading application code

Visible resource revocation

A way to reclaim resources and break
their(application & resources) secure binding

An exokernel uses visible revocation for most
resources

Traditionally, OS have performed revocation invisibly,
de-allocating resources without application involvement

Dialogue between an exokernel and a library OS
Library OS should organize resource lists

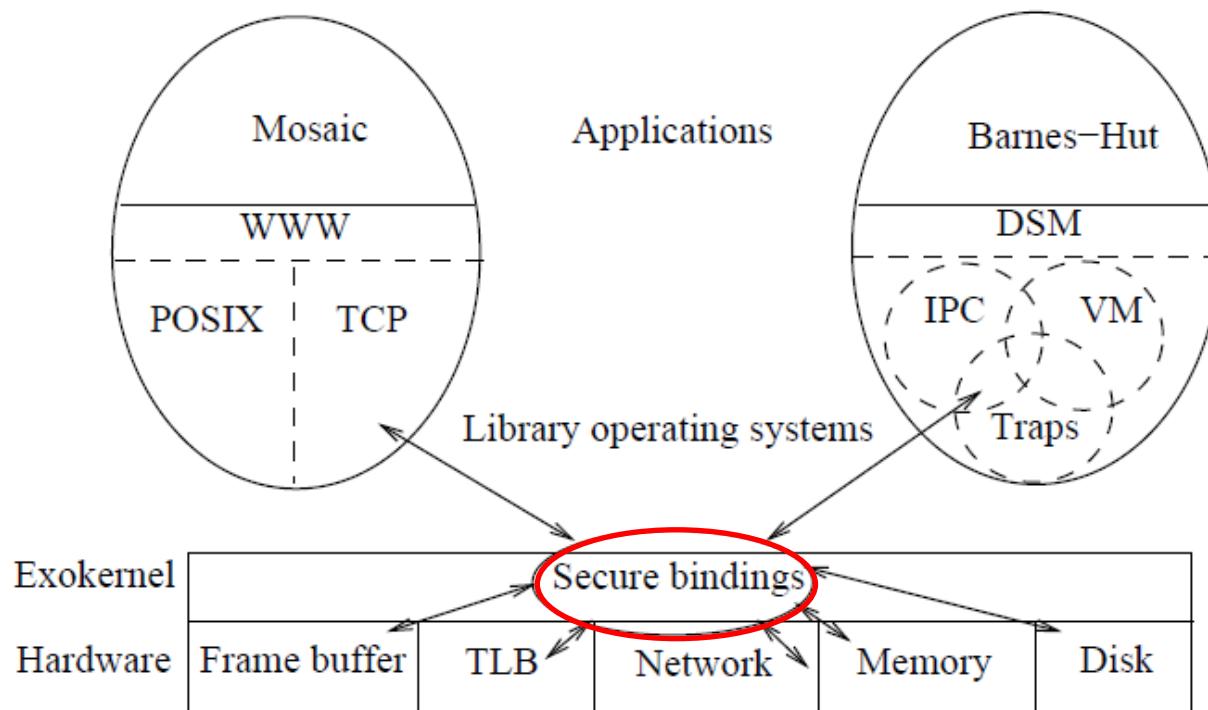
The abort protocol

An exokernel must also be able to take resources from library operating systems that fail to respond satisfactorily to revocation requests

If a library OS fails to respond quickly, the secure bindings need to be broken “**by force**”

An exokernel simply breaks all secure bindings to the resource and informs the library operating system

Exokernel Example 3



An example exokernel-based system consisting of a thin exokernel veneer that exports resources to library operating systems through secure bindings.

Implementation

Prototype (Xok / ExOS)

Xok

Exokernel

Runs on x86 (Aegis runs on DEC)

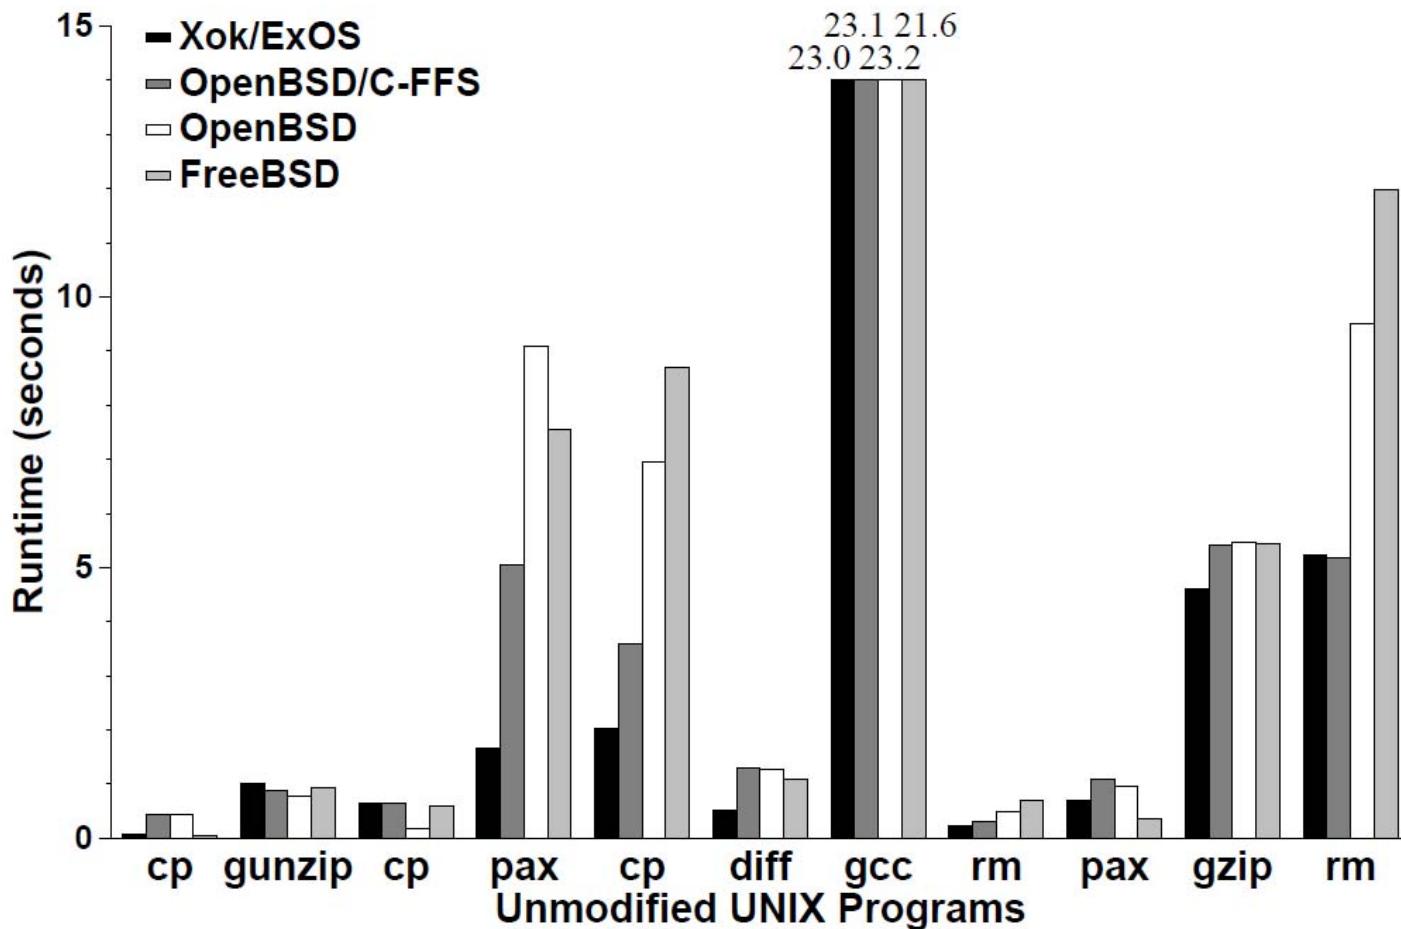
Safely multiplexes the physical resources

ExOS

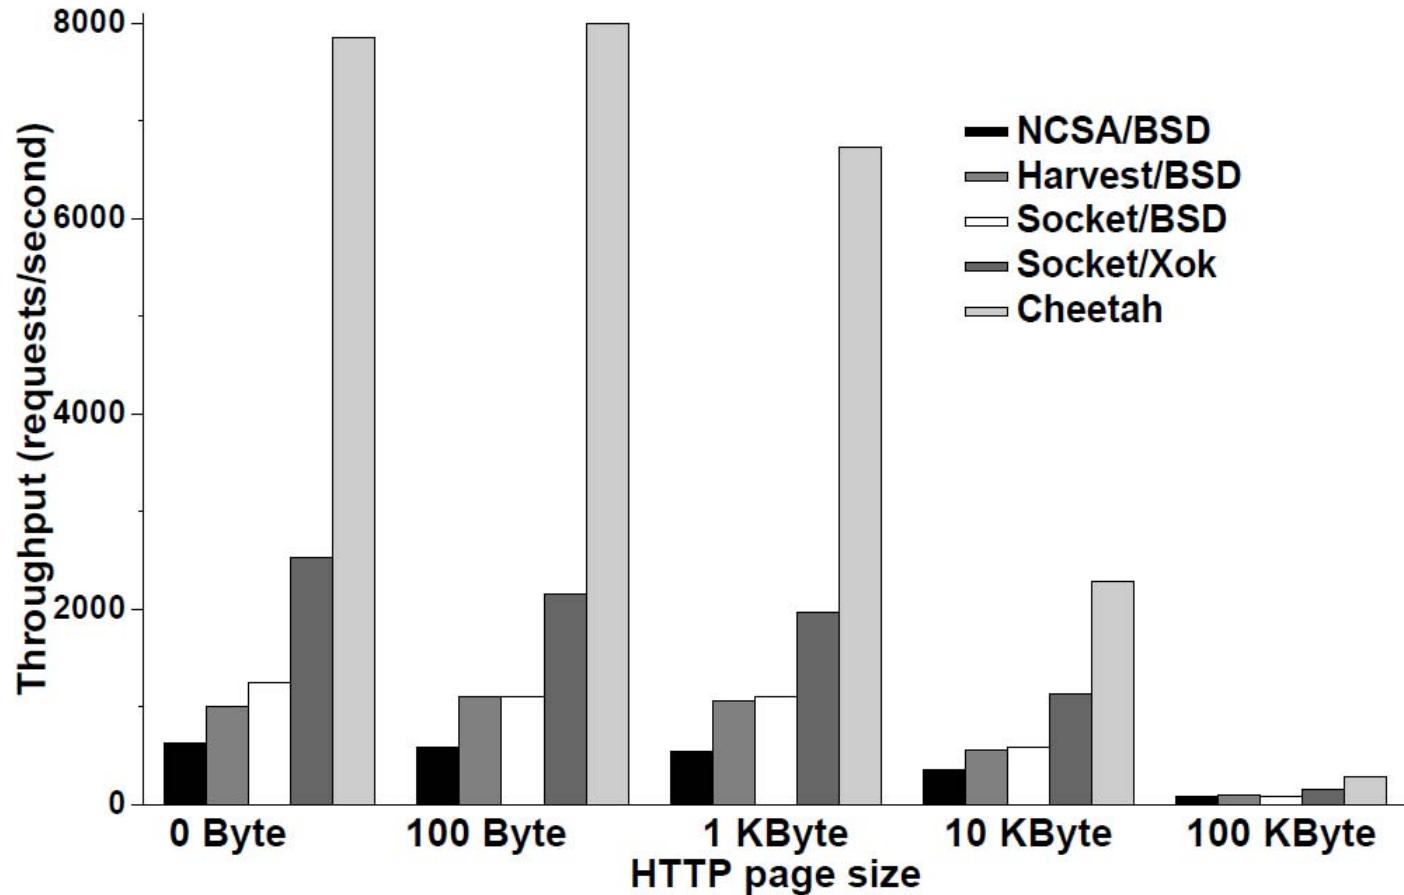
Library OS

Manages fundamental OS abstractions at application level
completely within the address space of the application
that is using it

Evaluation



Evaluation



Application-level control can significantly improve the performance of applications

Homework

Application can gain a better performance of memory and file system in exokernel compared with microkernel or monolithic kernel, why? And what's the price of performance profit?

Exokernel grants the applications more control over hardware resource. How does exokernel protect application against each other? And how to understand this goal of paper: to separate protection from management?

Open question - do you think why has not exokernel been as popular as monolithic kernel(Linux) and Hybrid kernel(Windows NT)?

References

- Exokernel: an operating system architecture for application-specific resource management, SOSP'95
- Application Performance and Flexibility on Exokernel Systems, SOSP'97
- Corey: An Operating System for Many Cores, OSDI'05

Summary

OS evolve constantly

Exokernel: application manages resources

OS Structure & XV6 VM

Recap

- OS functionalities
- OS structure

Review

- Q: what does an OS do?
- Q: what are the components of an OS?

Kernels

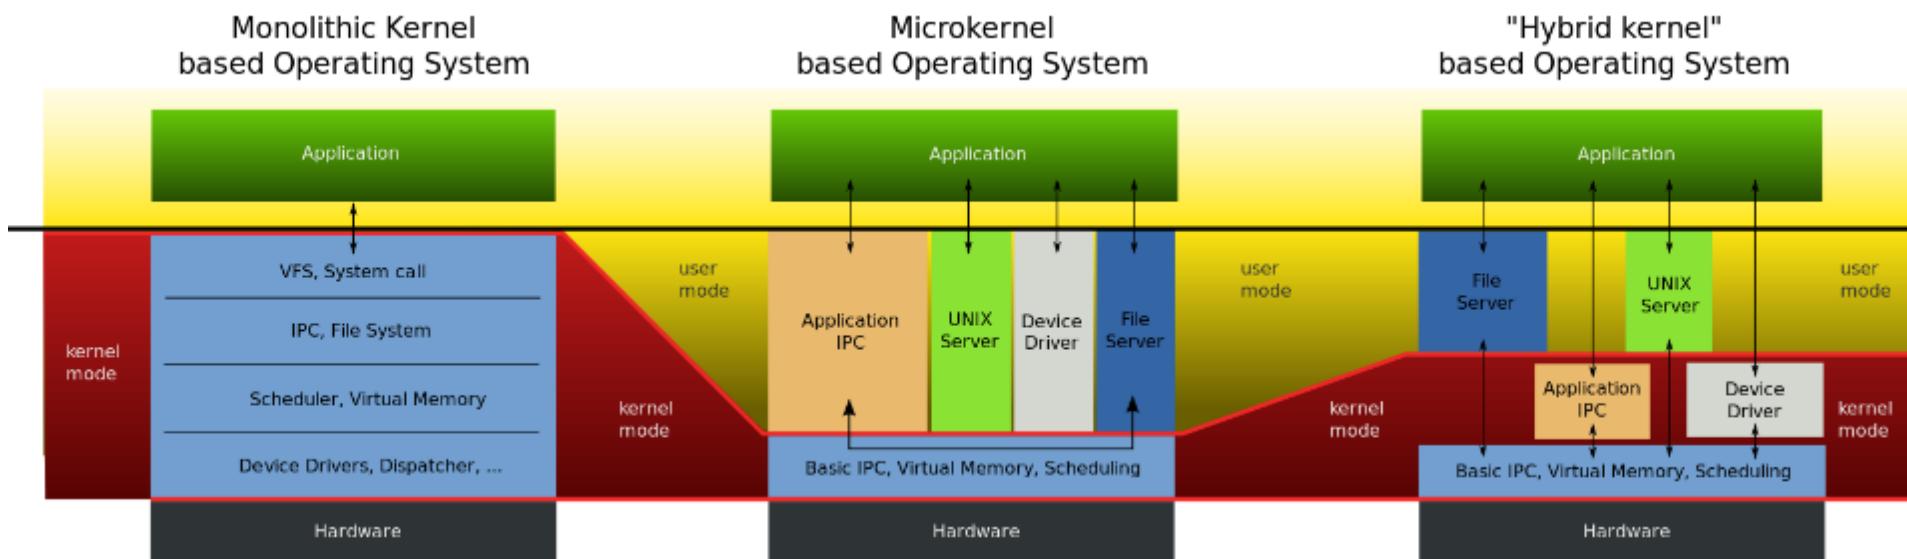
Kernel classification (wikipedia)

Kernel (architecture)	Monolithic	Microkernel	Exokernel	hybrid
--------------------------	------------	-------------	-----------	--------

Monolithic kernel

A **monolithic kernel** is an operating system architecture where the entire operating system is working in kernel space and is alone in supervisor mode.

Linux, BSD...



Kernels

Microkernel

is the near-minimum amount of software that can provide the mechanisms needed to implement an operating system (OS).

low-level address space management, thread management, and inter-process communication (IPC) and so on ...

Mach, L4 kernel

Hybrid kernel

is a kernel architecture based on combining aspects of microkernel and monolithic kernel architectures used in computer operating systems.

Windows NT kernel

Microkernel System Structure

Moves as much from the kernel into “user” space

Communication takes place between user modules using message passing

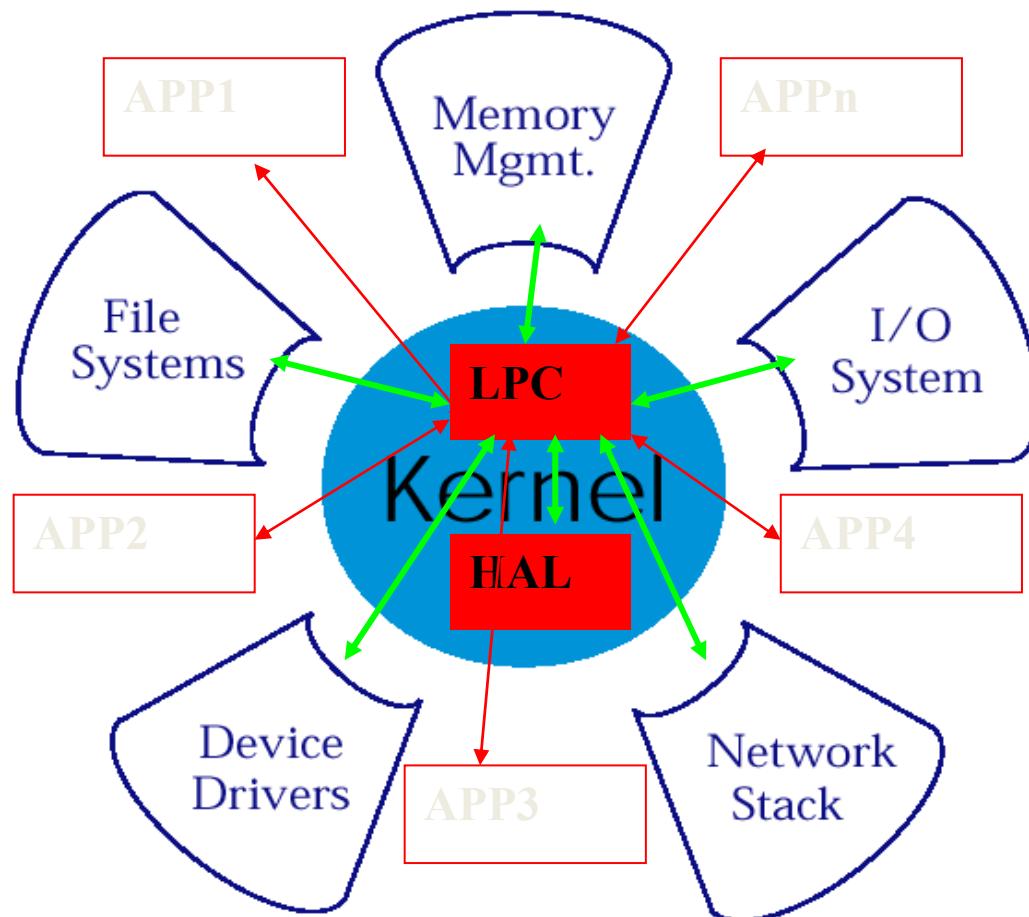
Benefits:

- Easier to extend a microkernel
- Easier to port the operating system to new architectures
- More reliable (less code is running in kernel mode)
- More secure

Detriments:

- Performance overhead of user space to kernel space communication

Microkernel Structure & Instances



Exokernel Structure

Overview

let the kernel allocate the physical resources of the machine to multiple application programs, and

let each program decide what to do with these resources.

The program can link to an operating system library (libOS) that implements OS abstractions.

Exokernel: An operating system architecture for application-level resource management, SOSP'95

Exokernel & Corey

Exokernel [SOSP'95]

Kernel classification

Motivation

Design and Implementation

Evaluation

Corey [OSDI'08]

Multi-core era and scalability problem

Sharing and Address range

Summery

Questions

What's the shortcoming of those kernels

Hint: that's the reason to create exokernel

Motivation

Exokernel's Motivation

In traditional operating systems, only privileged servers and the kernel can manage system resources

Un-trusted applications are required to interact with the hardware via some abstraction model

File systems for disk storage, *virtual address spaces* for memory, etc.

Application demands vary widely!!

An interface designed to accommodate every application must anticipate **all possible needs**

Exokernel's Idea

Give un-trusted applications as **much control** over physical resources as possible

To force as **few abstraction** as possible on developers, enabling them to make as **many decisions** as possible about hardware abstractions.

Let *the kernel* allocate the basic physical resources of the machine

Let *each program* decide what to do with these resources

Exokernel separate *protection* from *management*

They protect resources but delegate management to application

Library Operating System

Most programs to be linked with *libraries* instead of communicating with the exokernel directly

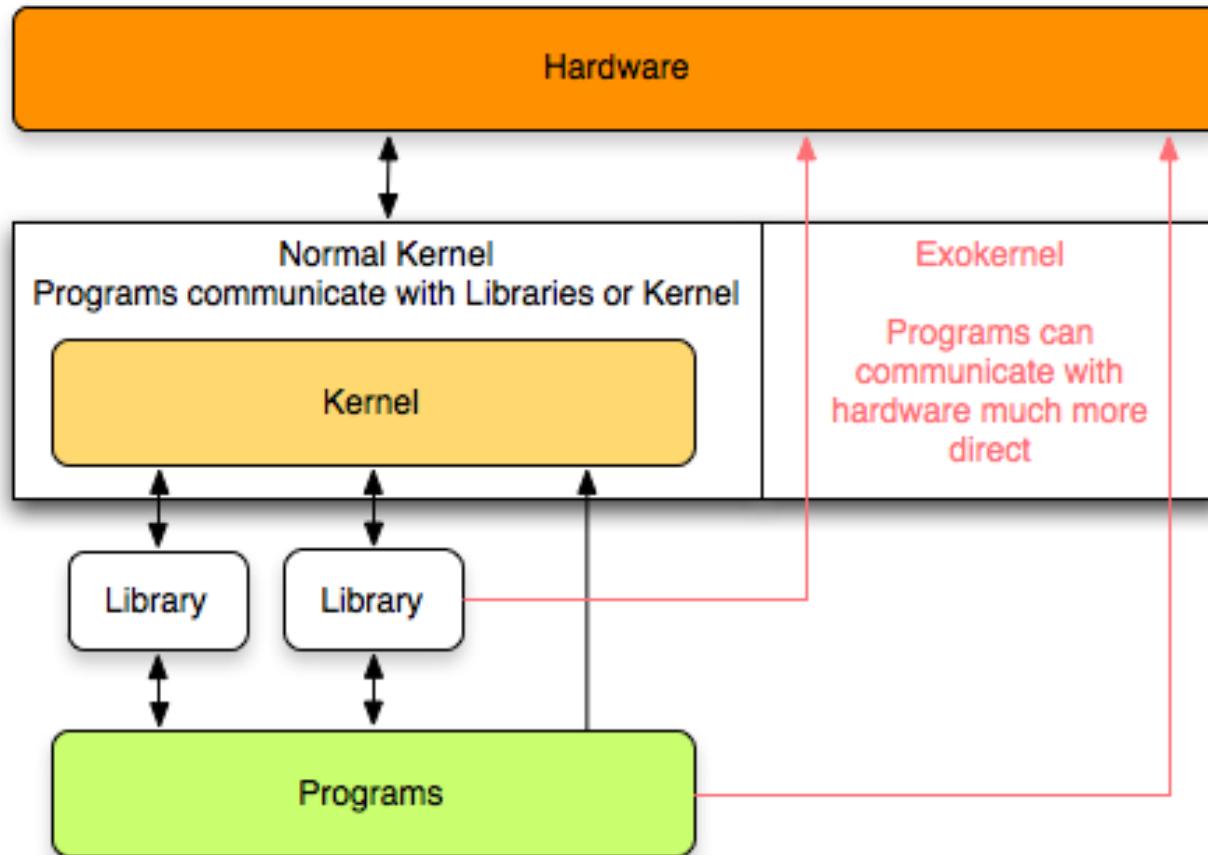
The libraries hide low-level resources

An applications can choose the library which best suits its needs, or even build its own

The kernel only ensures that the requested resource is free, and **the application** is allowed to access it.

Allow programmer to implement custom abstractions, omit unnecessary ones, most commonly to improve performance

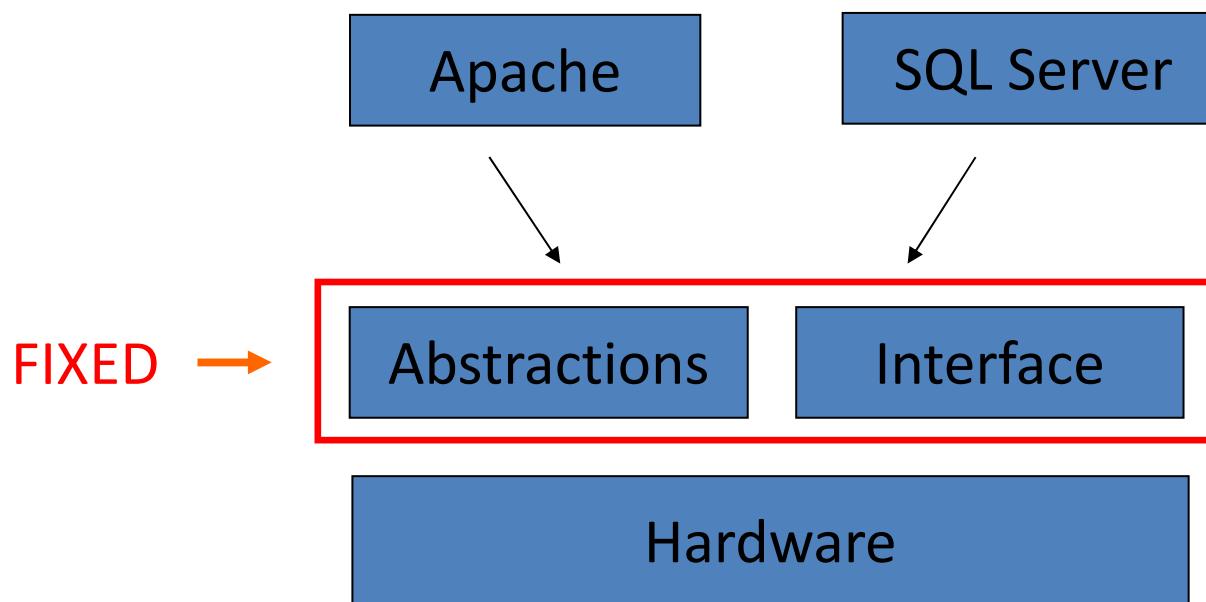
Exokernel Example 1



Exokernel give more direct access to the hardware, thus removing most abstractions

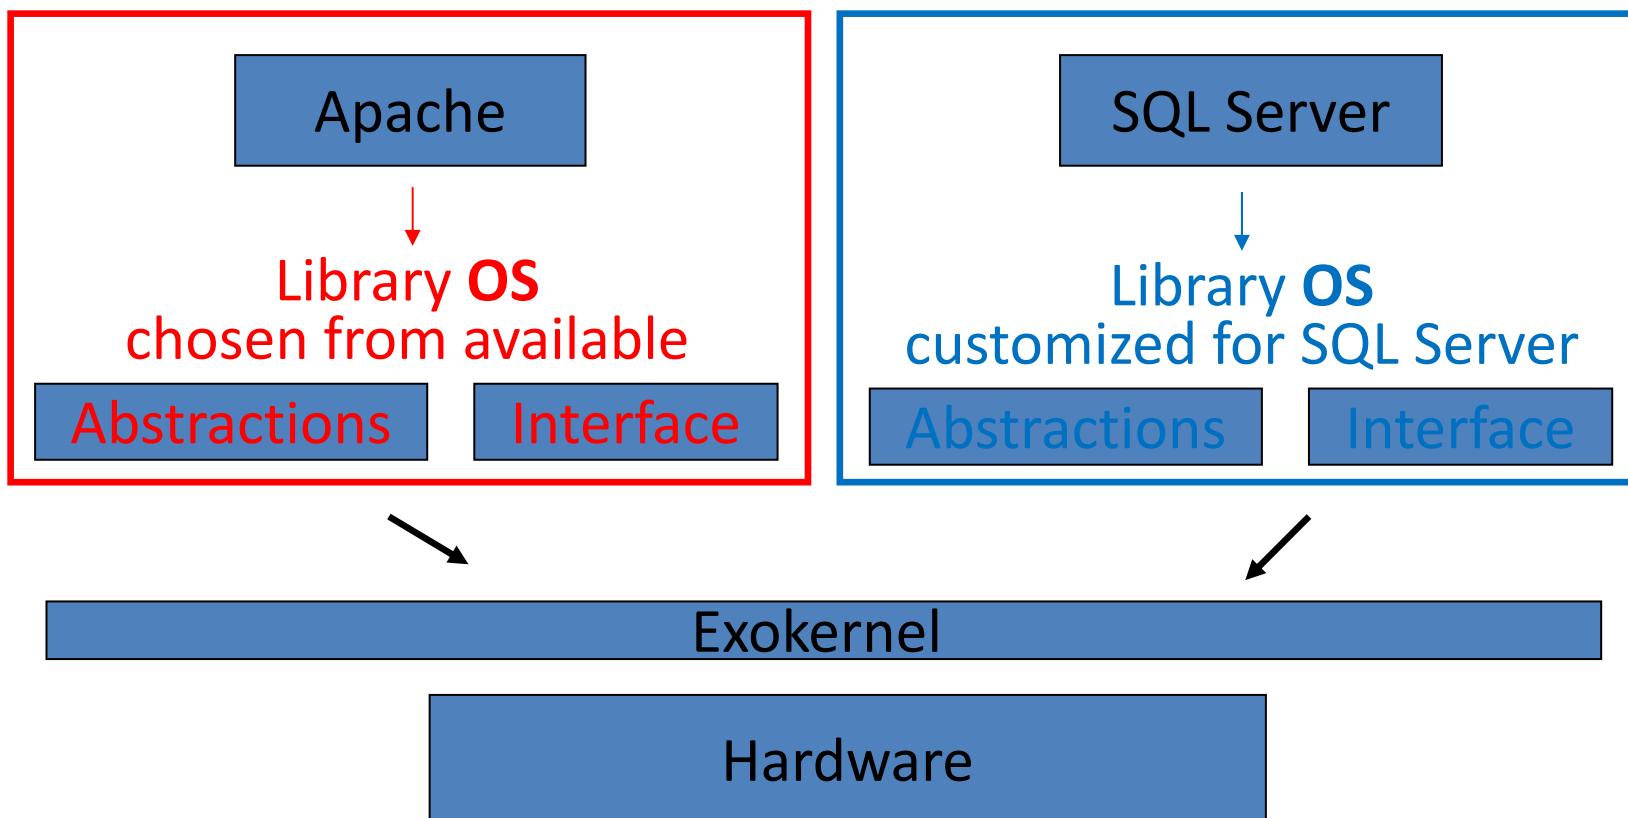
Exokernel Example 2

Traditional OS

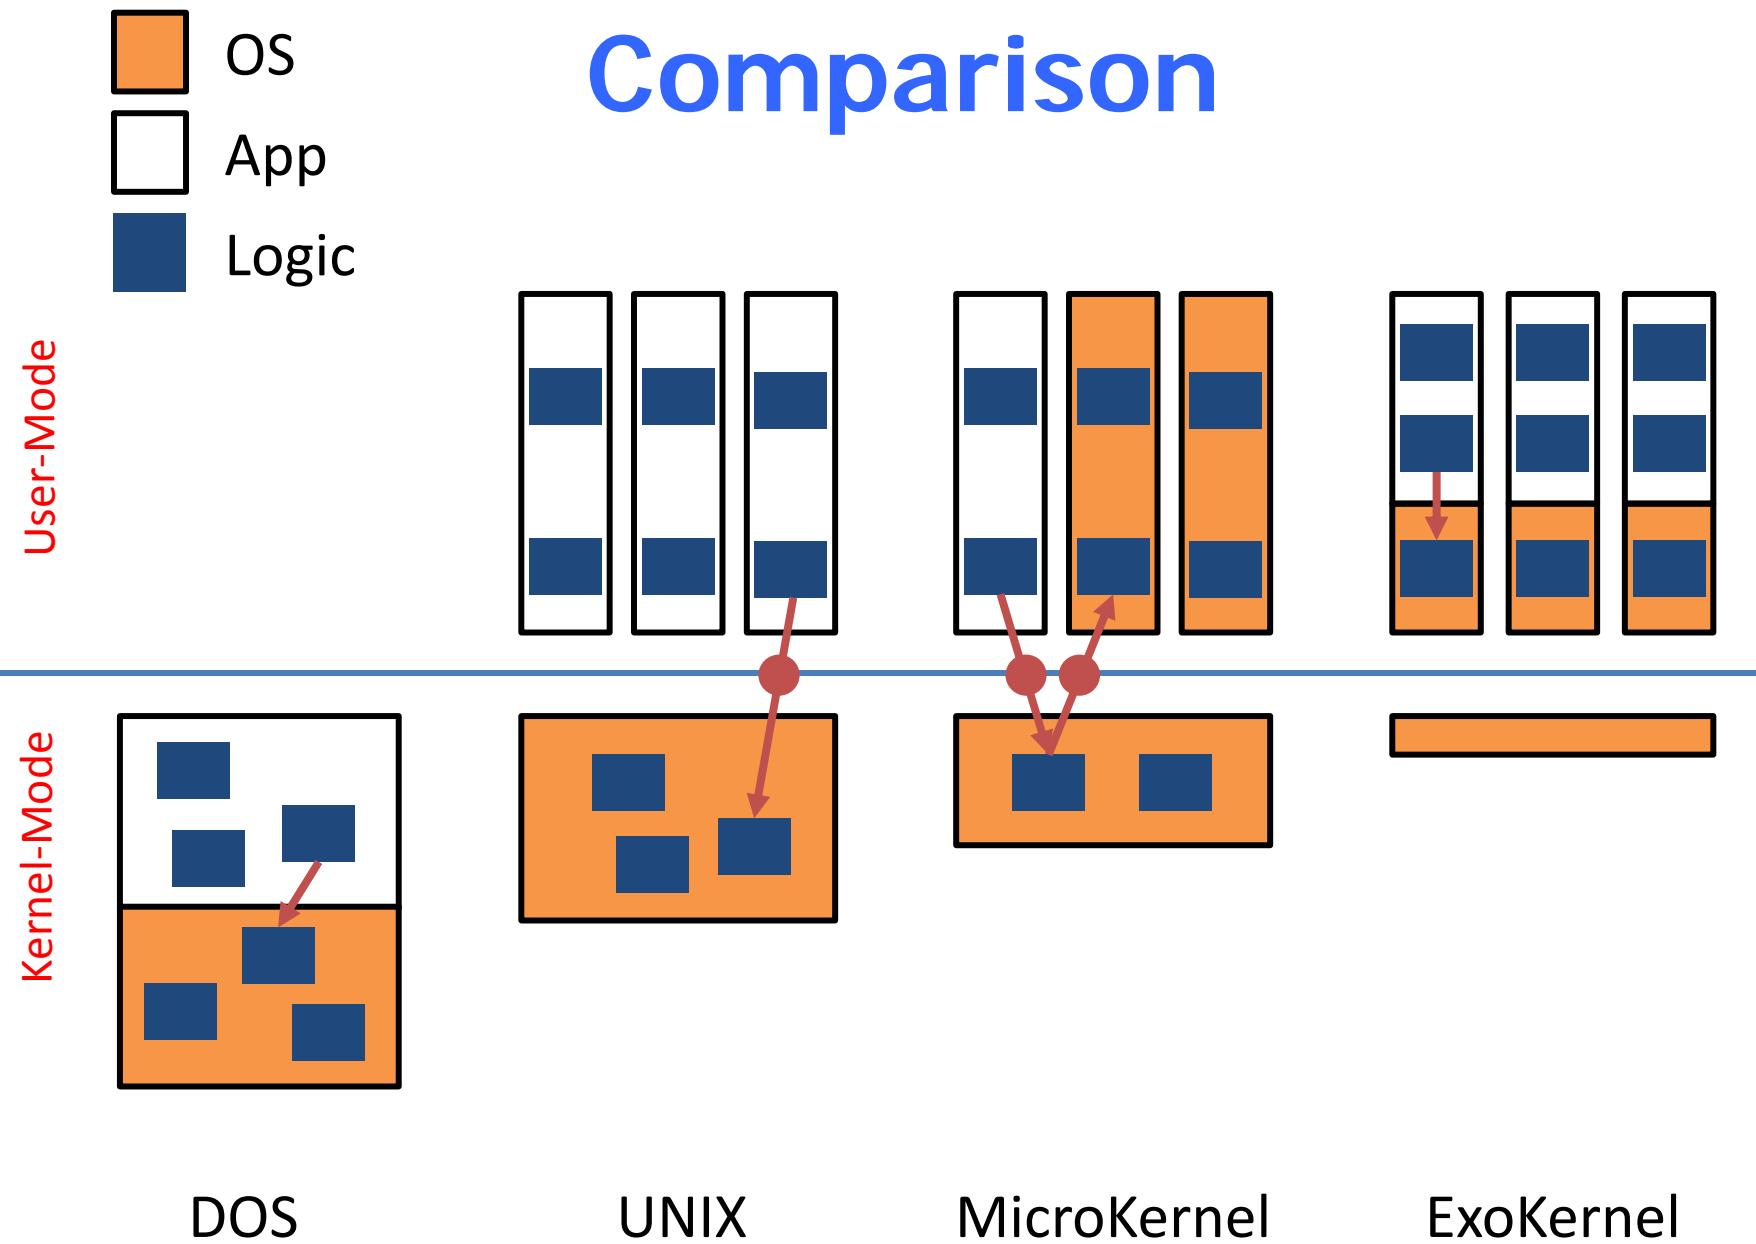


Exokernel Example 2

Exokernel + Library OS



Comparison



Exokernel Design Challenge

Kernel's new role

- Tracking ownership of resources

- Ensuring resource protection

- Revoking resource access

Three techniques

- Secure binding**

- Visible revocation**

- Abort protocol**

Secure binding

It is a protection mechanism that decouples authorization from actual use of a resource

Can improve performance

The protection checks involved in enforcing a secure binding are expressed in terms of simple operations that the kernel can implement quickly

A secure binding performs authorization only at bind time, which allows management to be decoupled from protection

Three techniques

Hardware mechanism, software caching, and downloading application code

Visible resource revocation

A way to reclaim resources and break
their(application & resources) secure binding

An exokernel uses visible revocation for most
resources

Traditionally, OS have performed revocation invisibly,
de-allocating resources without application involvement

Dialogue between an exokernel and a library OS
Library OS should organize resource lists

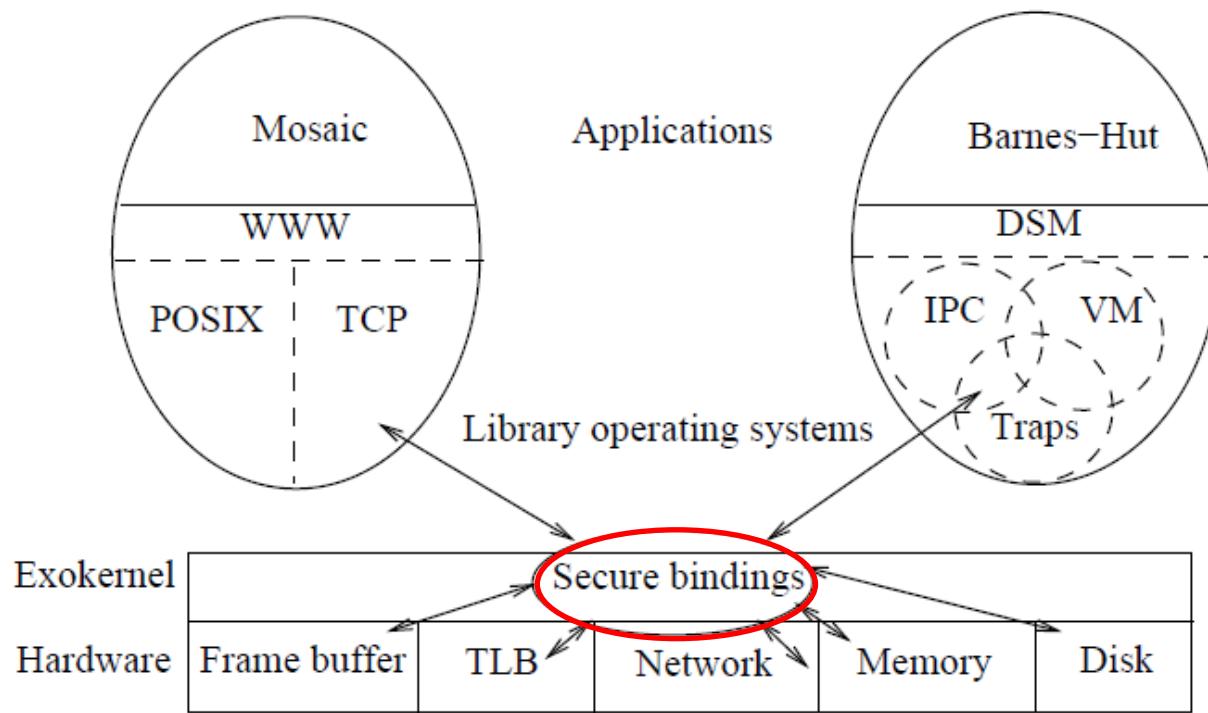
The abort protocol

An exokernel must also be able to take resources from library operating systems that fail to respond satisfactorily to revocation requests

If a library OS fails to respond quickly, the secure bindings need to be broken “**by force**”

An exokernel simply breaks all secure bindings to the resource and informs the library operating system

Exokernel Example 3



An example exokernel-based system consisting of a thin exokernel veneer that exports resources to library operating systems through secure bindings.

Implementation

Prototype (Xok / ExOS)

Xok

Exokernel

Runs on x86 (Aegis runs on DEC)

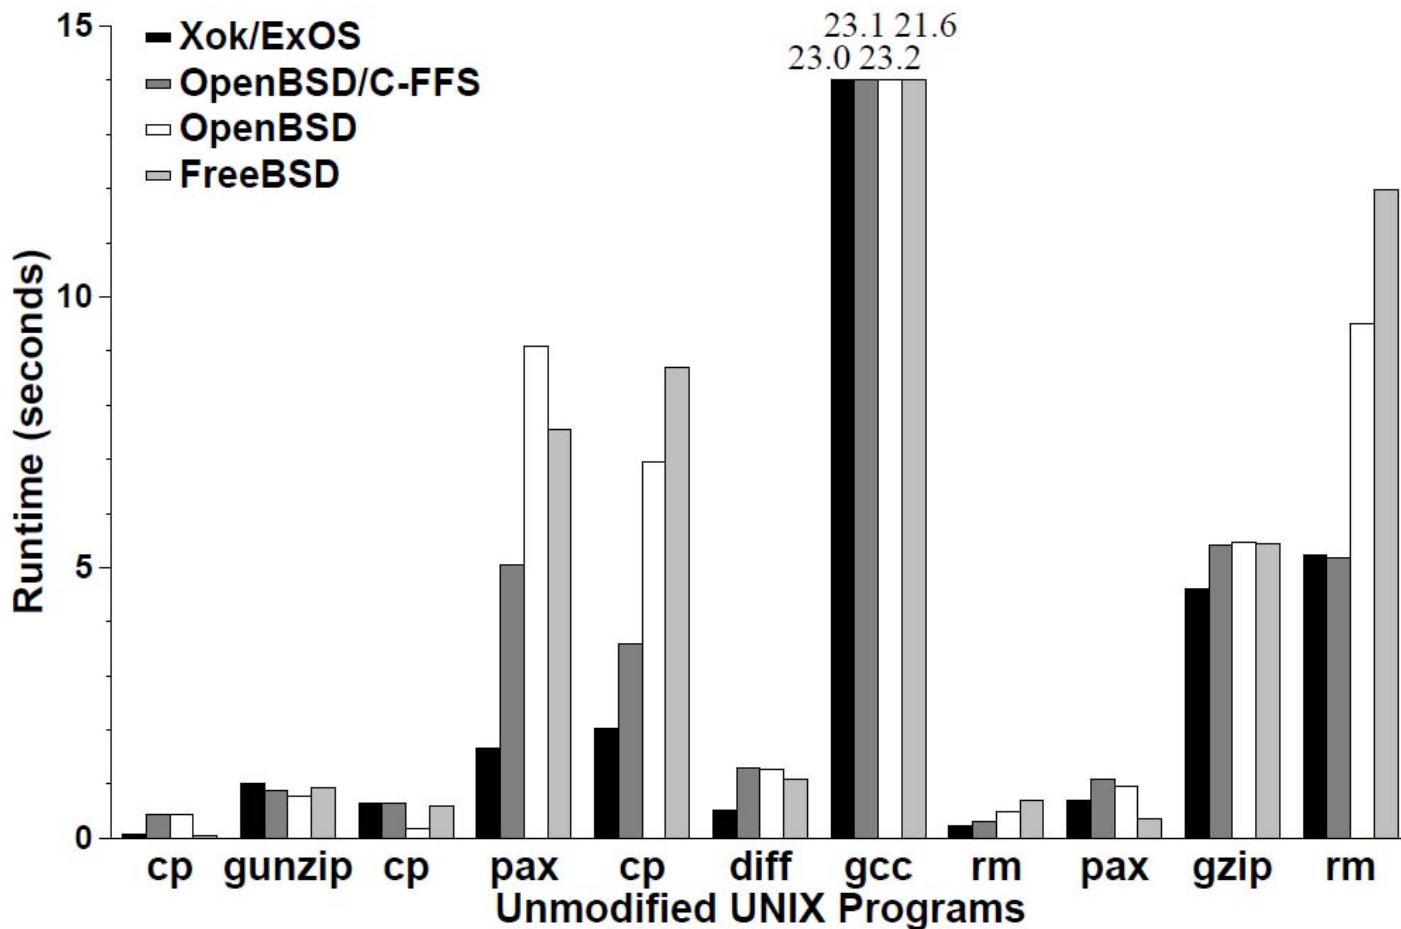
Safely multiplexes the physical resources

ExOS

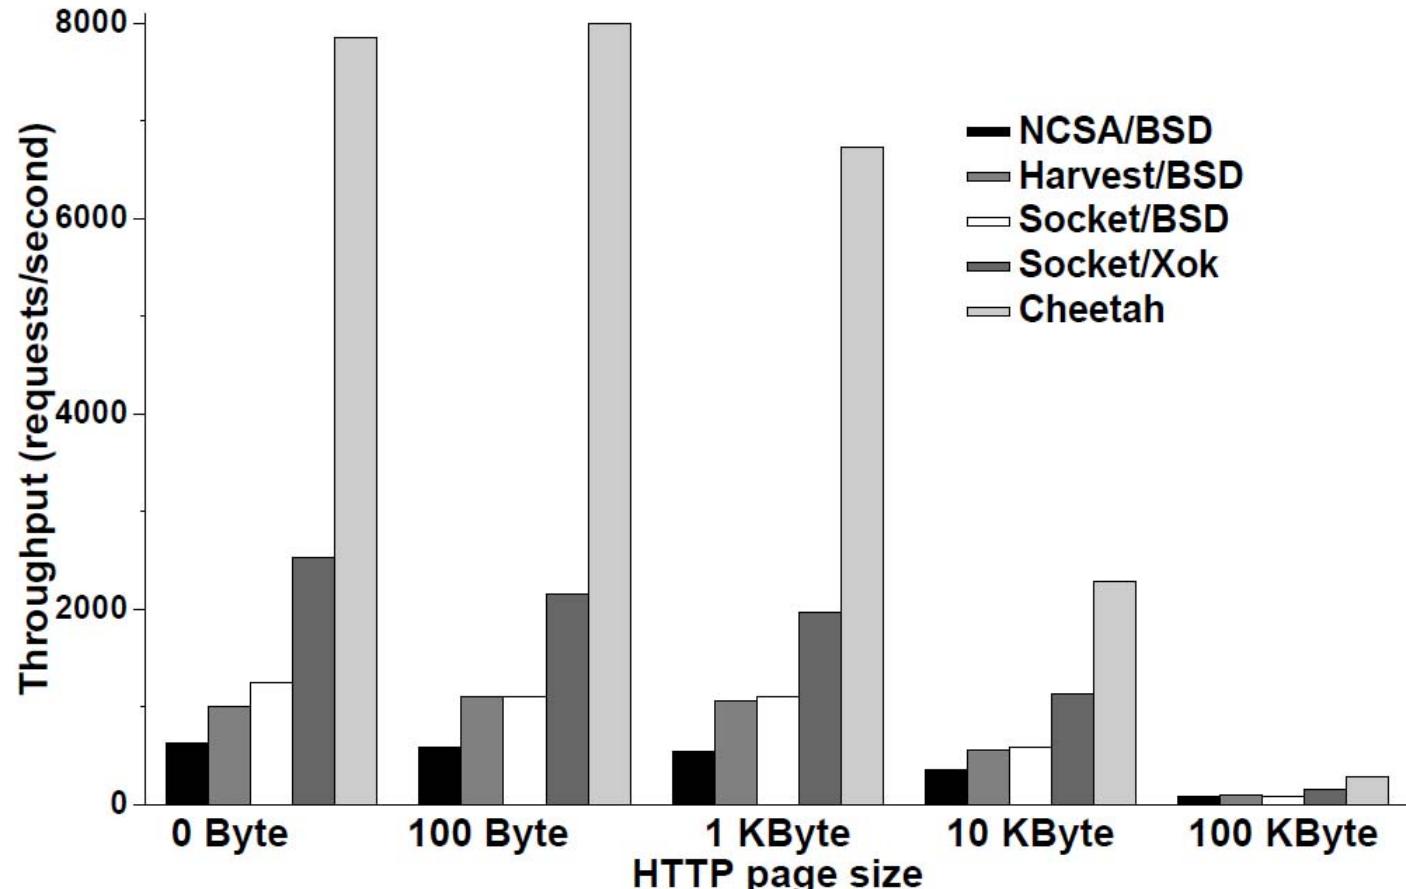
Library OS

Manages fundamental OS abstractions at application level
completely within the address space of the application
that is using it

Evaluation



Evaluation



Application-level control can significantly improve the performance of applications

Questions

Application can gain a better performance of memory and file system in exokernel compared with microkernel or monolithic kernel, why? And what's the price of performance profit?

Exokernel grants the applications more control over hardware resource. How does exokernel protect application against each other? And how to understand this goal of paper: to separate protection from management?

Open question - do you think why has not exokernel been as popular as monolithic kernel(Linux) and Hybrid kernel(Windows NT)?

References

Exokernel: an operating system architecture for application-specific resource management,
SOSP'95

Application Performance and Flexibility on
Exokernel Systems, SOSP'97

Corey: An Operating System for Many Cores,
OSDI'08

Homework

- Please read two papers with the same name
 - Are Virtual Machine Monitors Microkernels Done Right?
 - <http://research.cs.wisc.edu/areas/os/ReadingGroup/OS/papers/hotos-ukernel.pdf>
 - http://ssrg.nicta.com.au/publications/papers/Heiser_UL_06.pdf
- Please figure out what they were talking about, summarize their arguments, and give your own understanding

XV6 & VM

Rong Chen & Yubin Xia

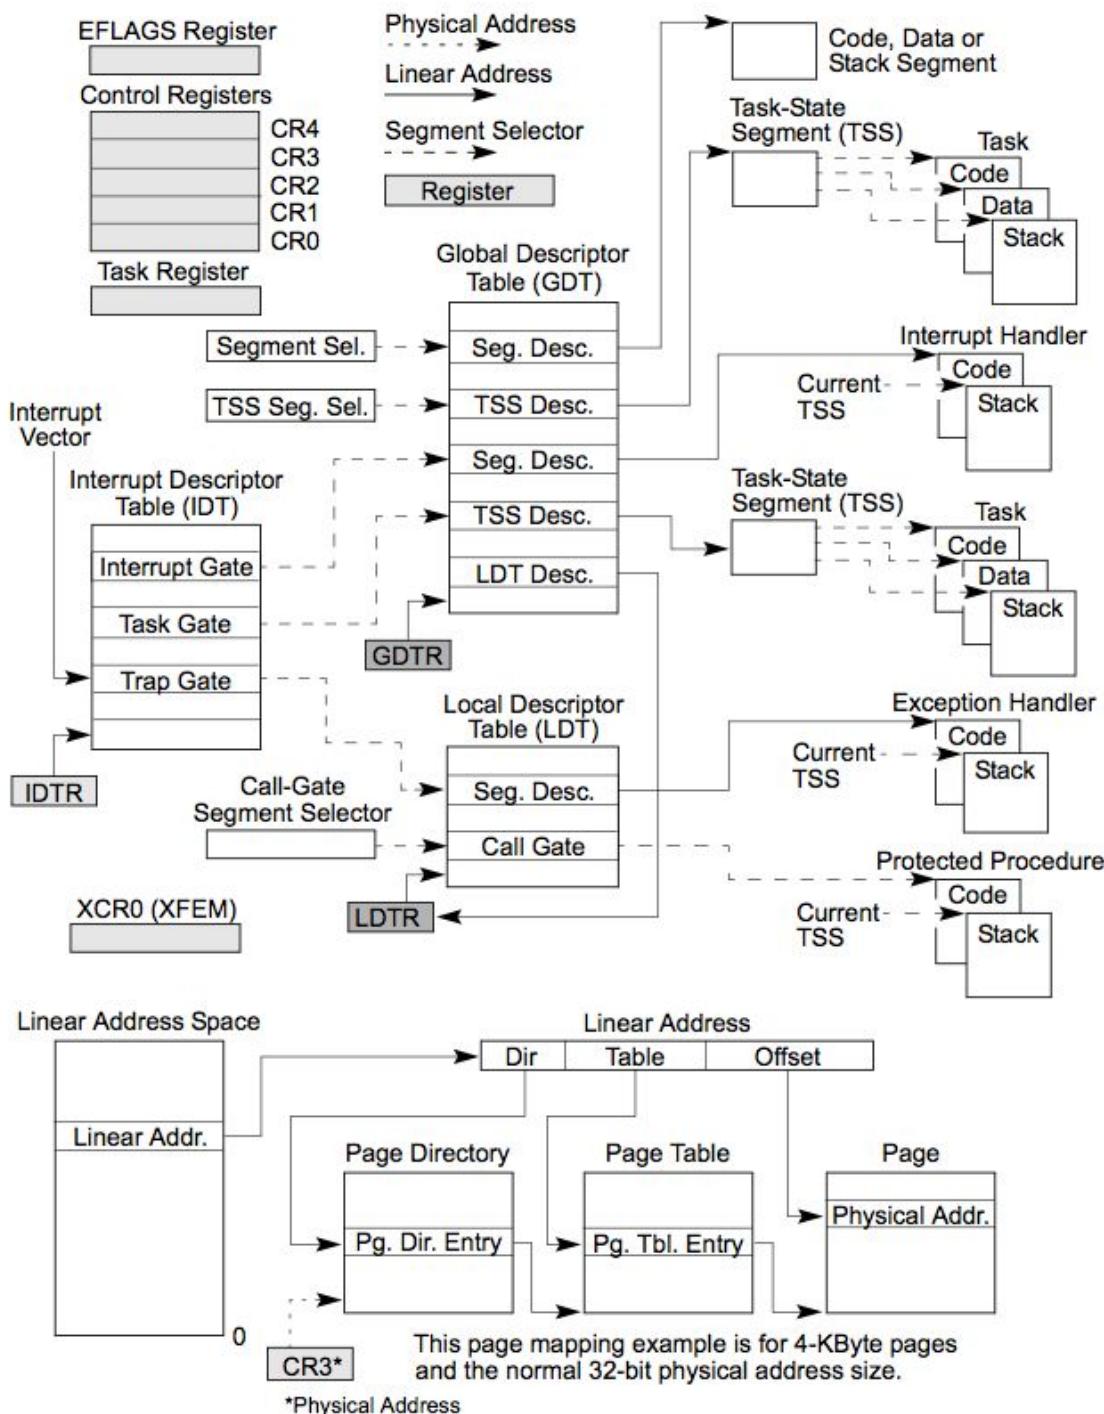
Some slides are based on Frans Kaashoek's slides
kaashoek@mit.edu

Outline

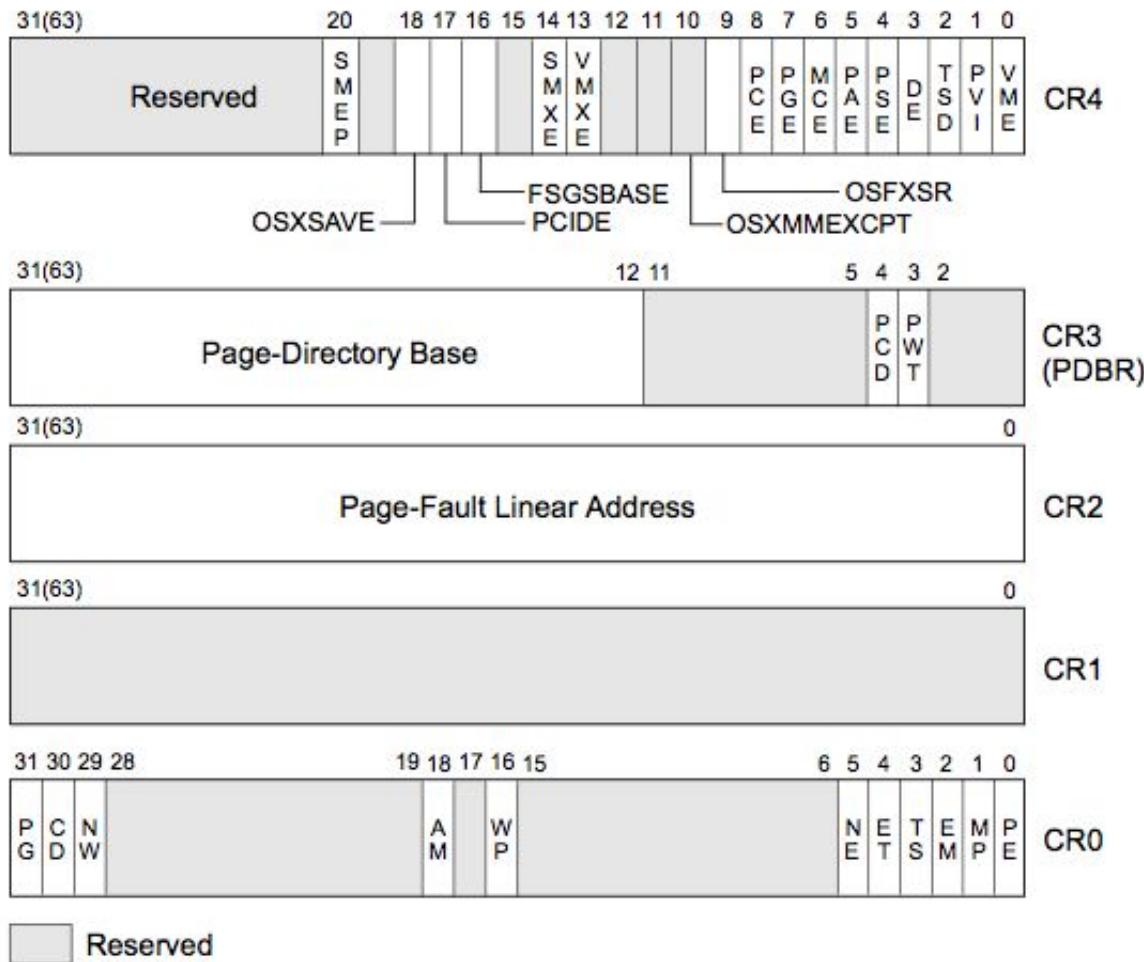
- Segment, Paging
- LA, VA, PA
- VM in xv6
 - Initial page table
 - Kernel page table
 - User page table
- Page Fault

VIRTUAL MEMORY

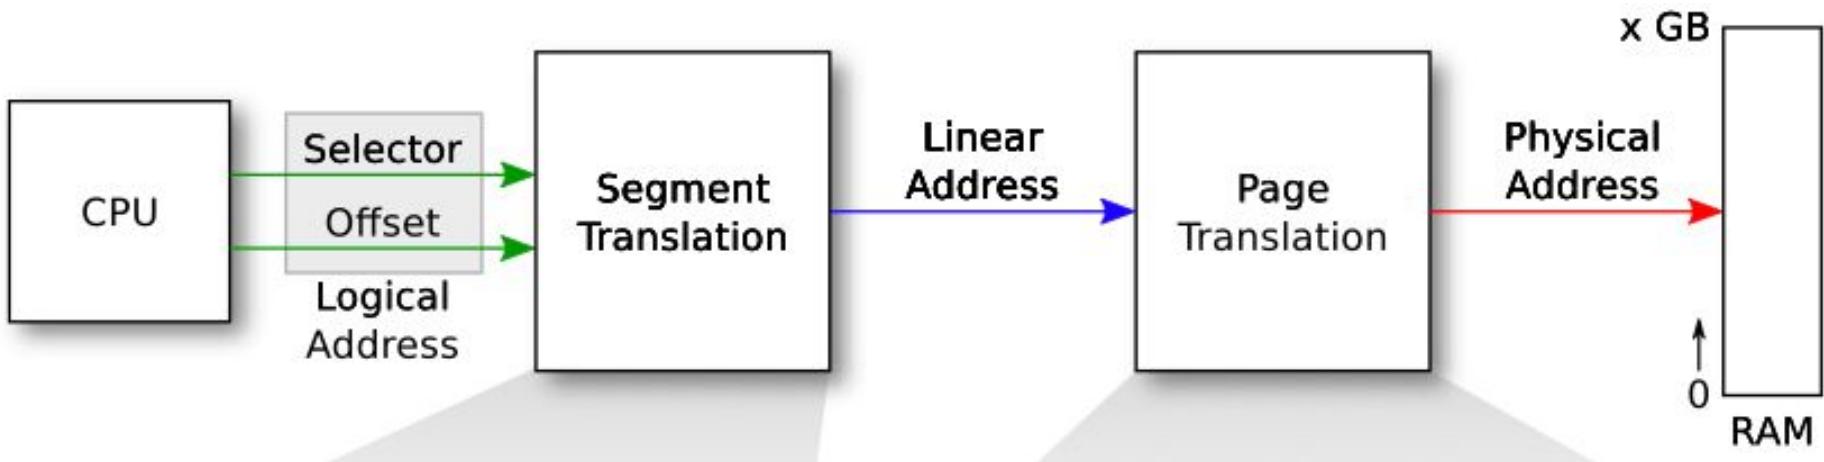
IA32 Arch Overview



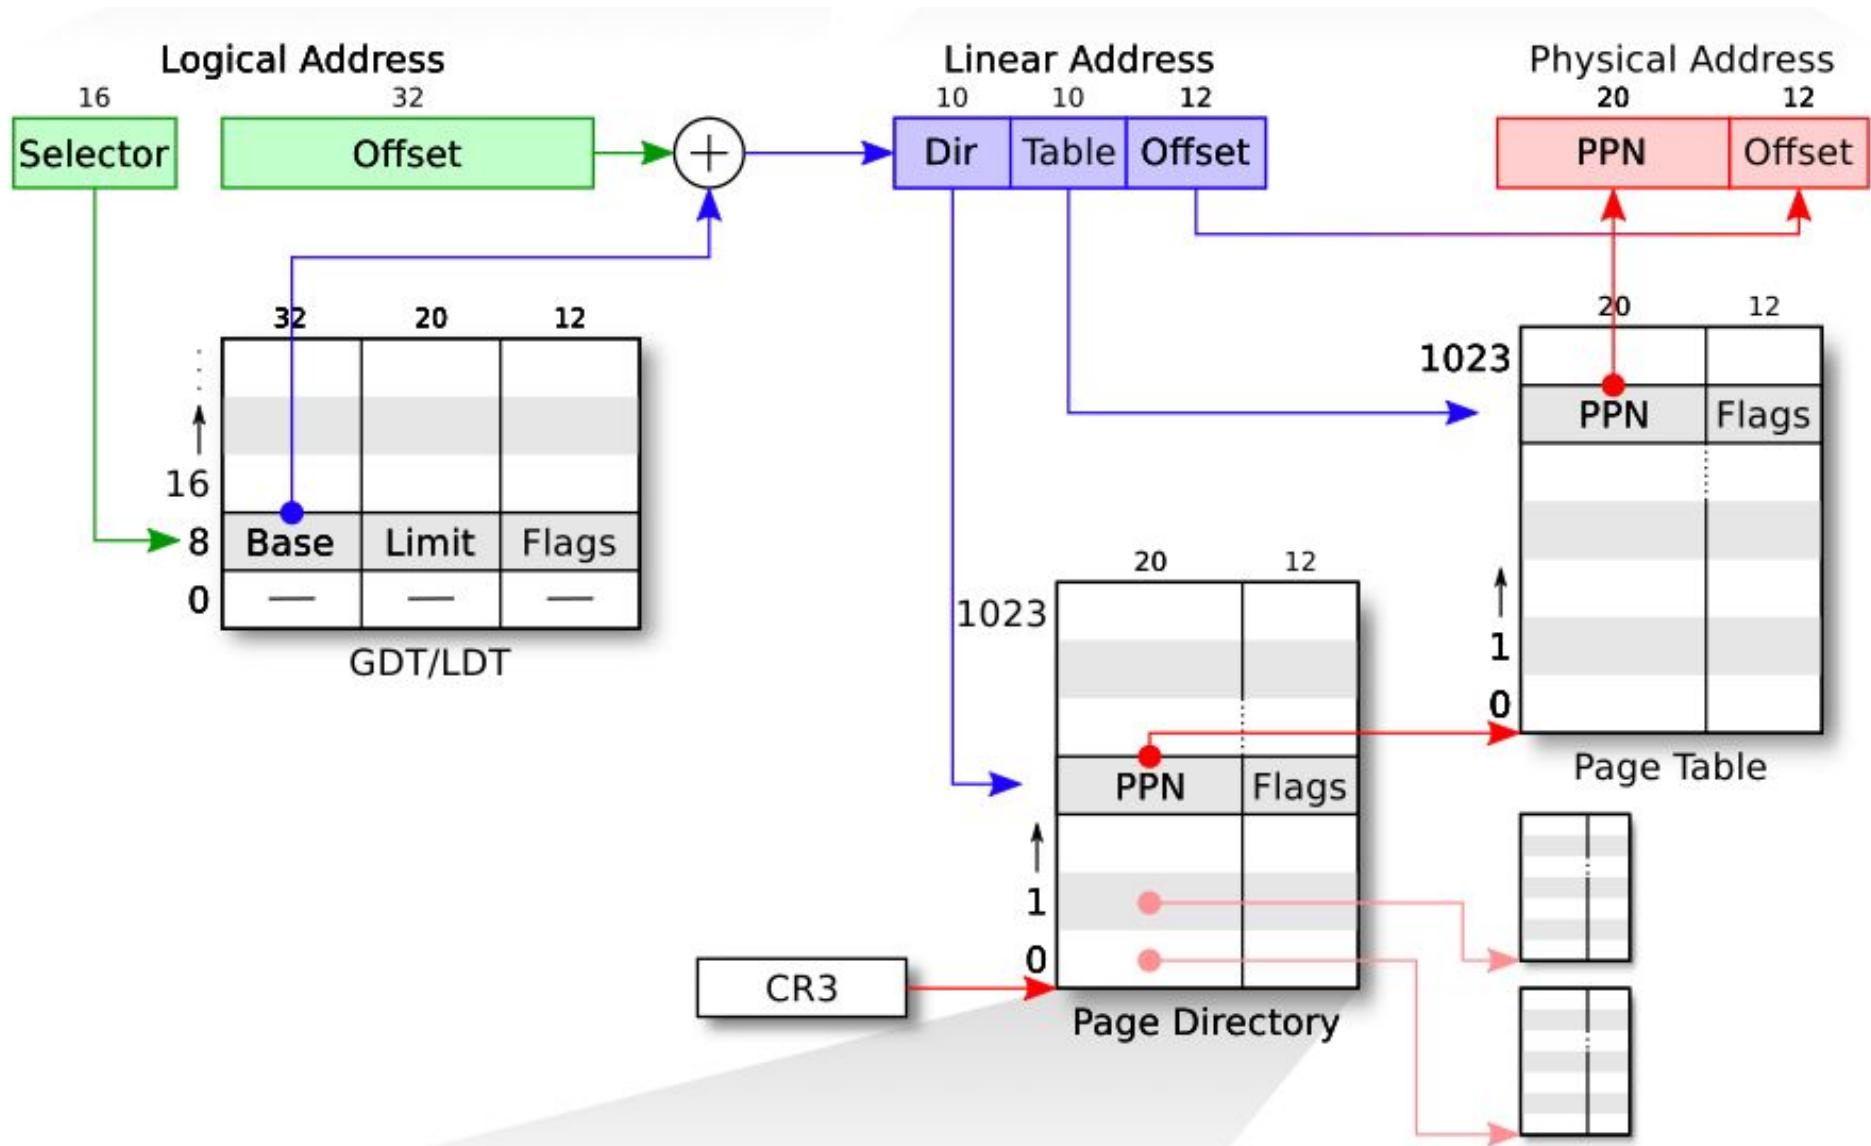
Control Registers



Protected-mode Address Translation



Protected Mode Address Translation



what's going to happen when the kernel enables segment?

- virtual address translation comes into effect?
 - presumably the enabling instruction has address segment-base + 0xffffc
 - The address may not contain any valid instruction!
 - Solution
 - Ensure that the addresses are the same before and after

Enable Segment

```
78 # Bootstrap GDT
79 .p2align 2                                # force 4 byte alignment
80 gdt:
81     SEG_NULLASM                         # null seg
82     SEG_ASM(STA_XISTA_R, 0x0, 0xffffffff) # code seg
83     SEG_ASM(STA_W, 0x0, 0xffffffff)      # data seg
84
85 gdtdesc:
86     .word    (gdtdesc - gdt - 1)          # sizeof(gdt) - 1
87     .long    gdt                          # address gdt
```

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt    gdtdesc
movl    %cr0, %eax
orl    $CR0_PE, %eax
movl    %eax, %cr0

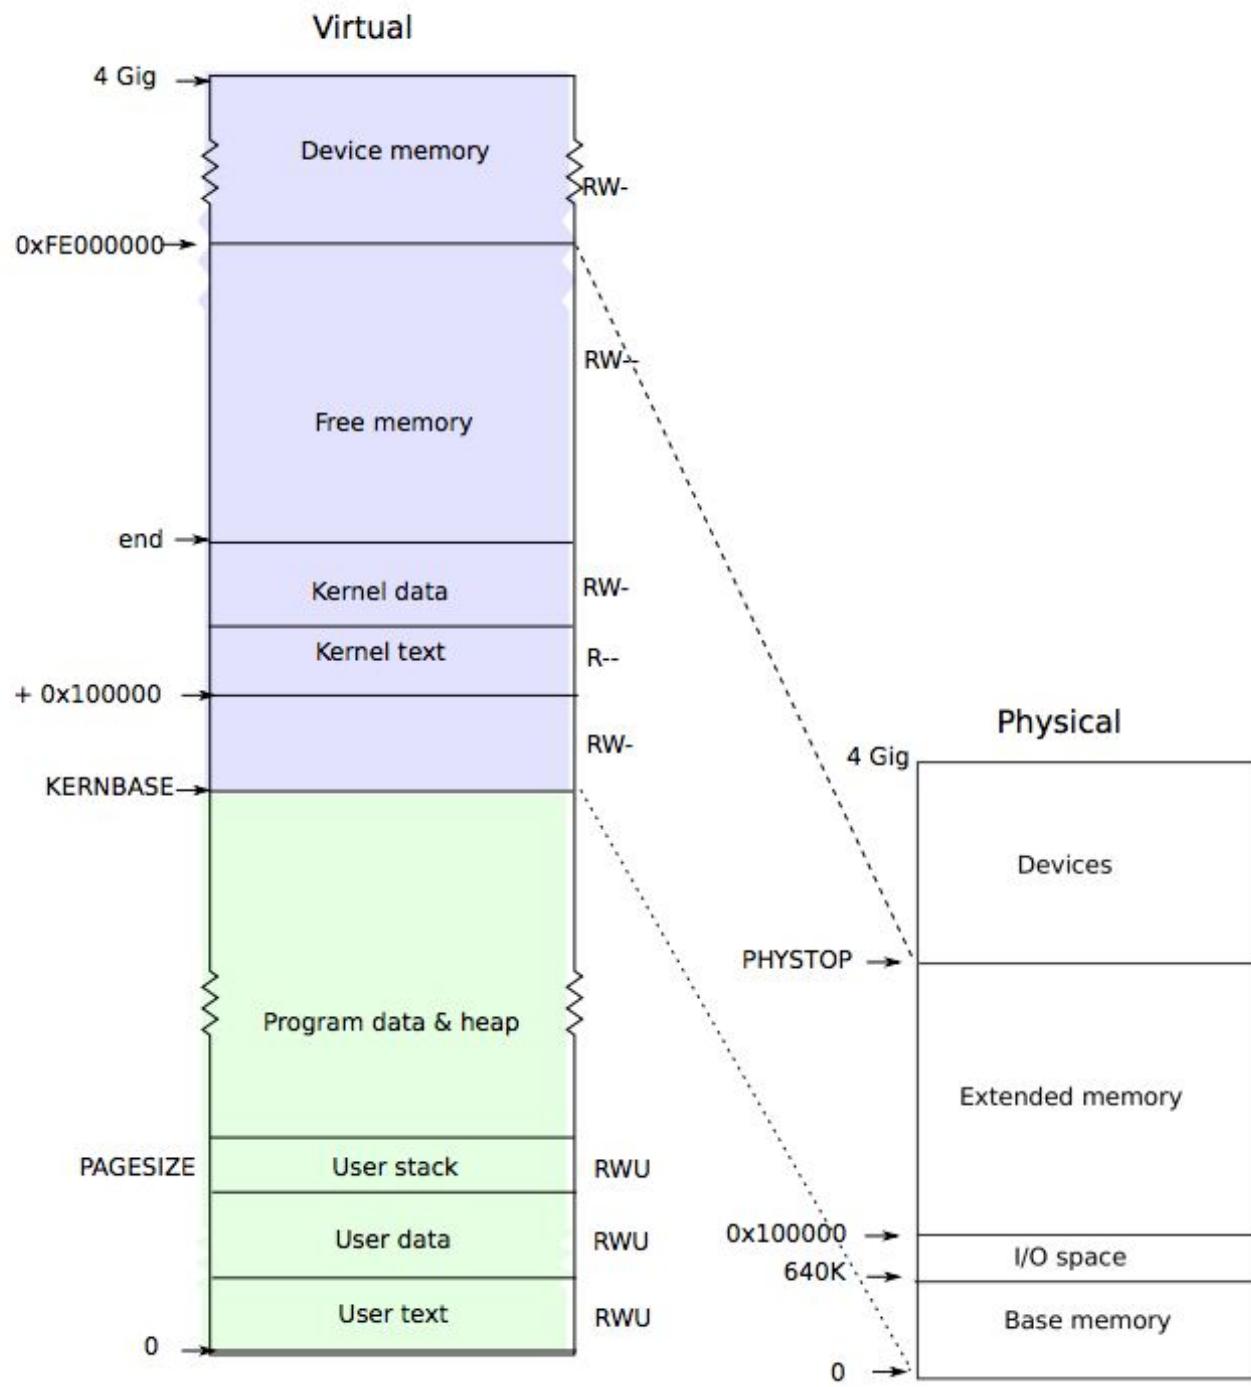
//PAGEBREAK!
# Complete transition to 32-bit protected mode by using long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp   $(SEG_KCODE<<3), start32

.code32 # Tell assembler to generate 32-bit code now.
start32:
# Set up the protected-mode data segment registers
movw   $(SEG_KDATA<<3), %ax      # Our data segment selector
movw   %ax, %ds                   # -> DS: Data Segment
movw   %ax, %es                   # -> ES: Extra Segment
movw   %ax, %ss                   # -> SS: Stack Segment
movw   $0, %ax                   # Zero segments not ready for use
movw   %ax, %fs                   # -> FS
movw   %ax, %gs                   # -> GS

# Set up the stack pointer and call into C.
movl   $start, %esp
call   bootmain
```

what's going to happen when the kernel enables paging?

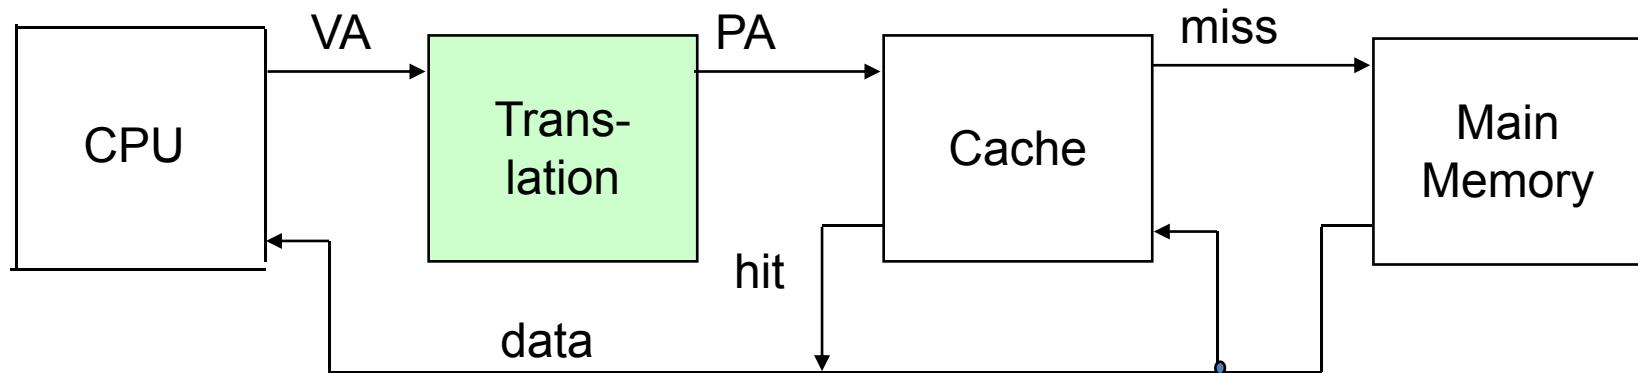
- virtual address translation comes into effect?
 - presumably the enabling instruction has address 0x100xxx
 - 0x100xxx maps to user memory, not kernel instructions!
 - solution
 - temporary page table with kernel mapped at both 0x80100000 and 0x100000
 - that's what the first few kernel instructions set up



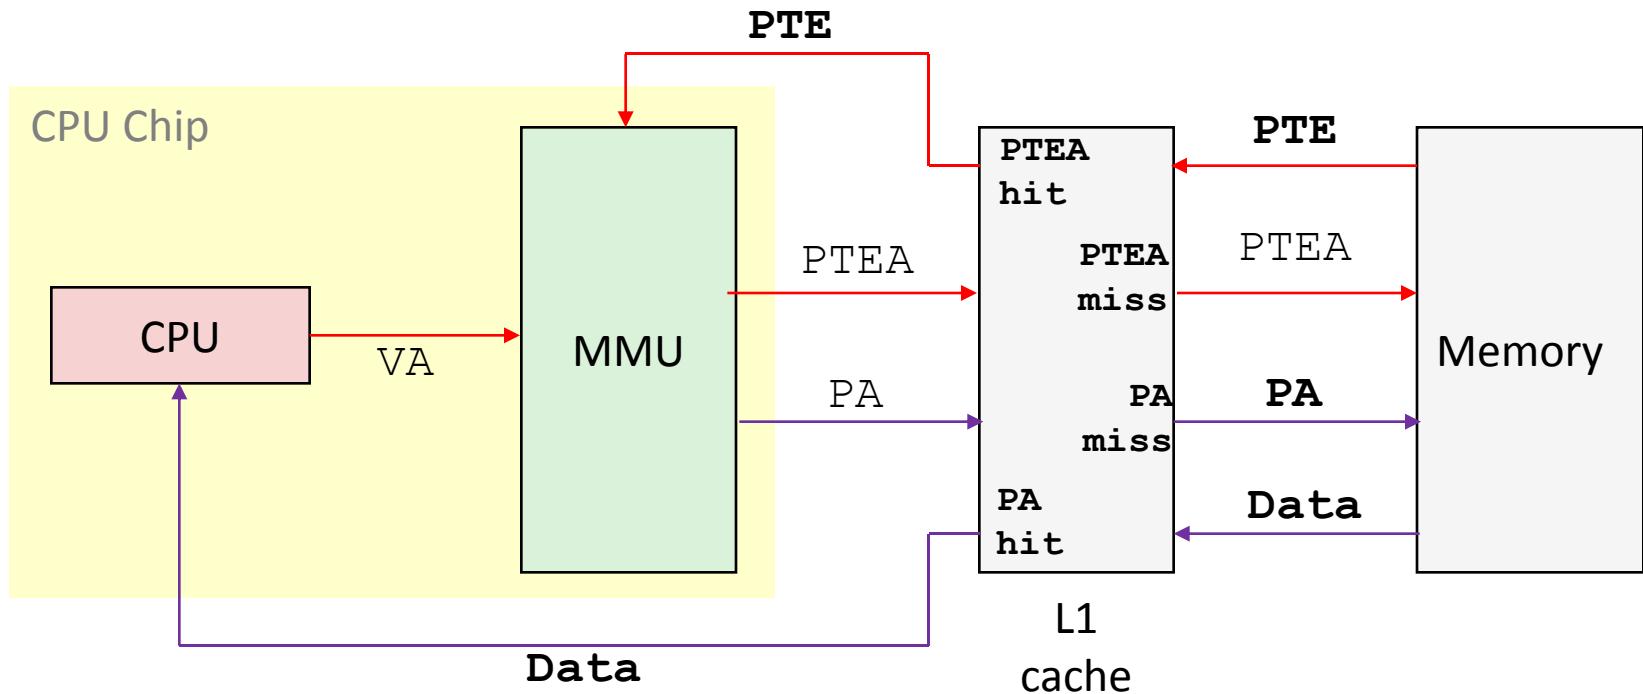
entrypgdir

```
102 // Boot page table used in entry.S and entryother.S.  
103 // Page directories (and page tables), must start on a page boundary,  
104 // hence the "__aligned__" attribute.  
105 // Use PTE_PS in page directory entry to enable 4Mbyte pages.  
106 __attribute__((__aligned__(PGSIZE)))  
107 pde_t entrypgdir[NPDENTRIES] = {  
108     // Map VA's [0, 4MB) to PA's [0, 4MB)  
109     [0] = (0) | PTE_P | PTE_W | PTE_PS,  
110     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)  
111     [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,  
112 };
```

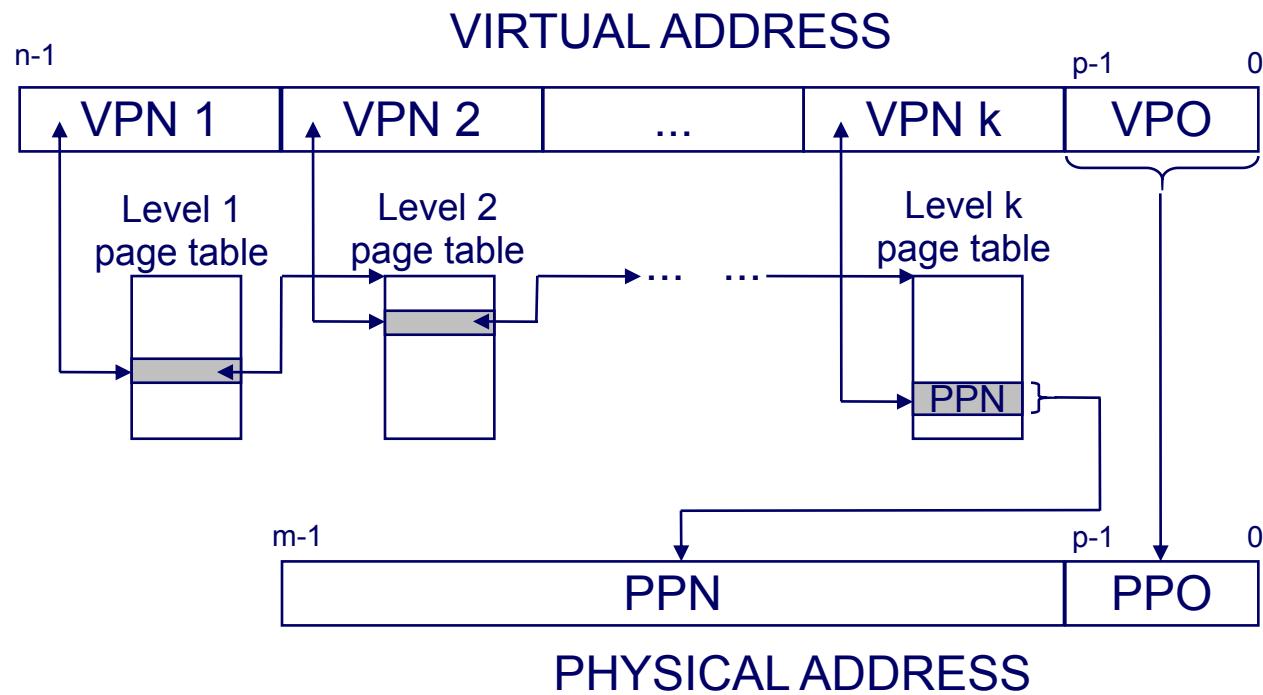
Recap: Integrating Caches and VM



Recap: Integrating Caches and VM



Translating with a k-level Page Table



Super Page: 4M Page

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0						
Address of page directory ¹												Ignored				P	P	C	W	T	Ignored								CR3								
Bits 31:22 of address of 2MB page frame				Reserved (must be 0)				Bits 39:32 of address ²				P	A	Ignored	G	1	D	A	P	P	U	R	/	1		PDE: 4MB page											
Address of page table												Ignored				Q	I	g	A	P	P	U	R	/	1		PDE: page table										
Ignored																Q											Q	PDE: not present									
Address of 4KB page frame												Ignored				G	P	A	D	A	P	P	U	R	/	1		PTE: 4KB page									
Ignored																											Q	PTE: not present									

Address translation: PAE-4k

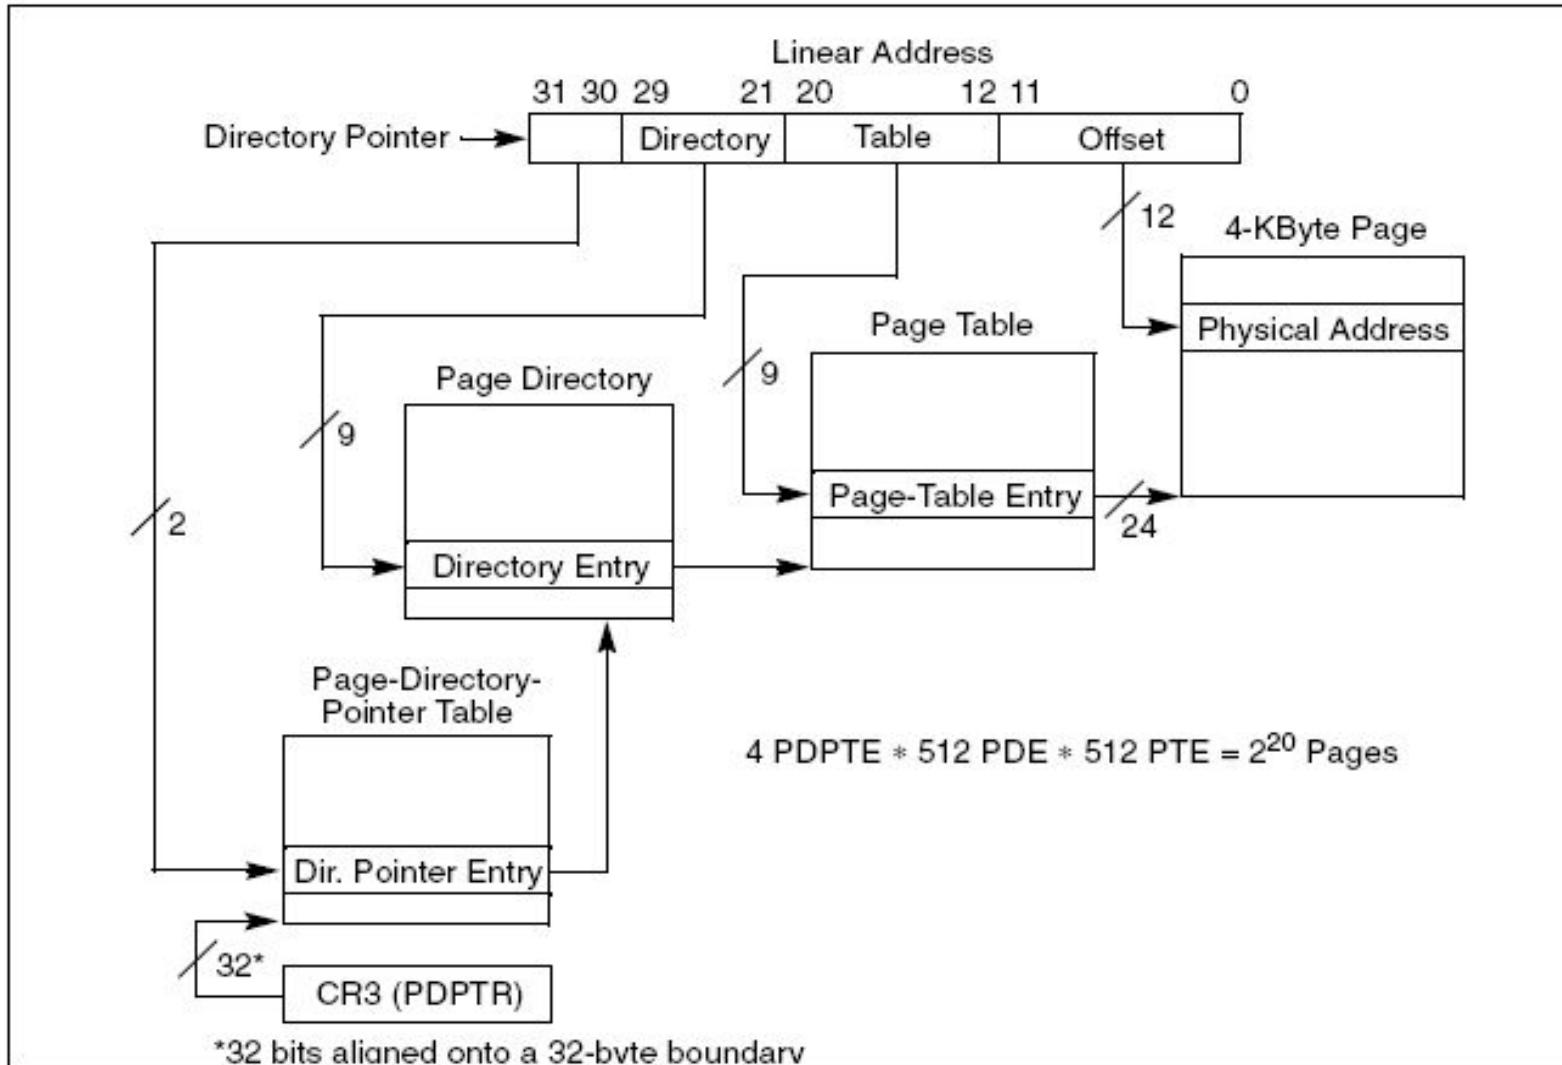


Figure 3-18. Linear Address Translation With PAE Enabled (4-KByte Pages)

Address translation: PAE-2M

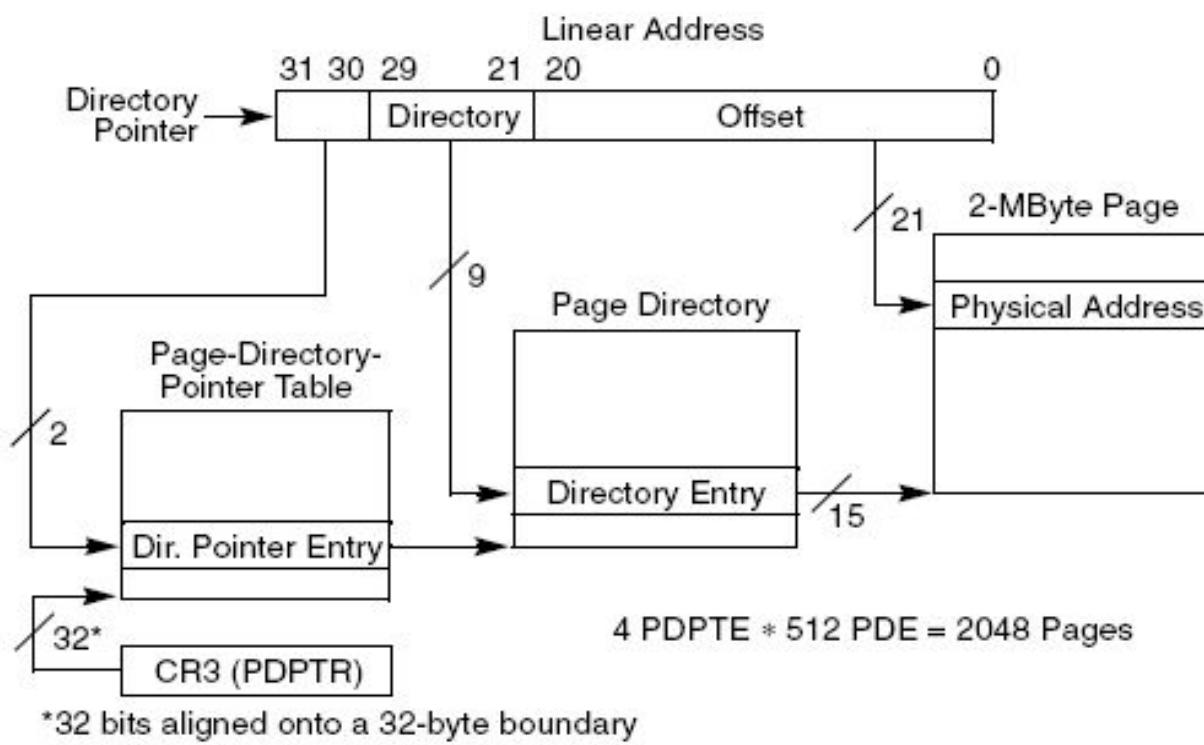


Figure 3-19. Linear Address Translation With PAE Enabled (2-MByte Pages)

Quiz#1

In **IA32** architecture, when the **CR4_PSE** bit in %cr4 register is set, then superpage mode is enabled, and the operating system can use either 4KB sized page or 4MB sized page. There are tradeoffs between small page and large page.

Consider the following scenarios:

- Dump the content of a large file of several hundred megabytes mapped to the RAM by mmap() function to another location:

```
memcpy(addr_of_new_location, addr_of_mapped_file, sizeof_file);
```

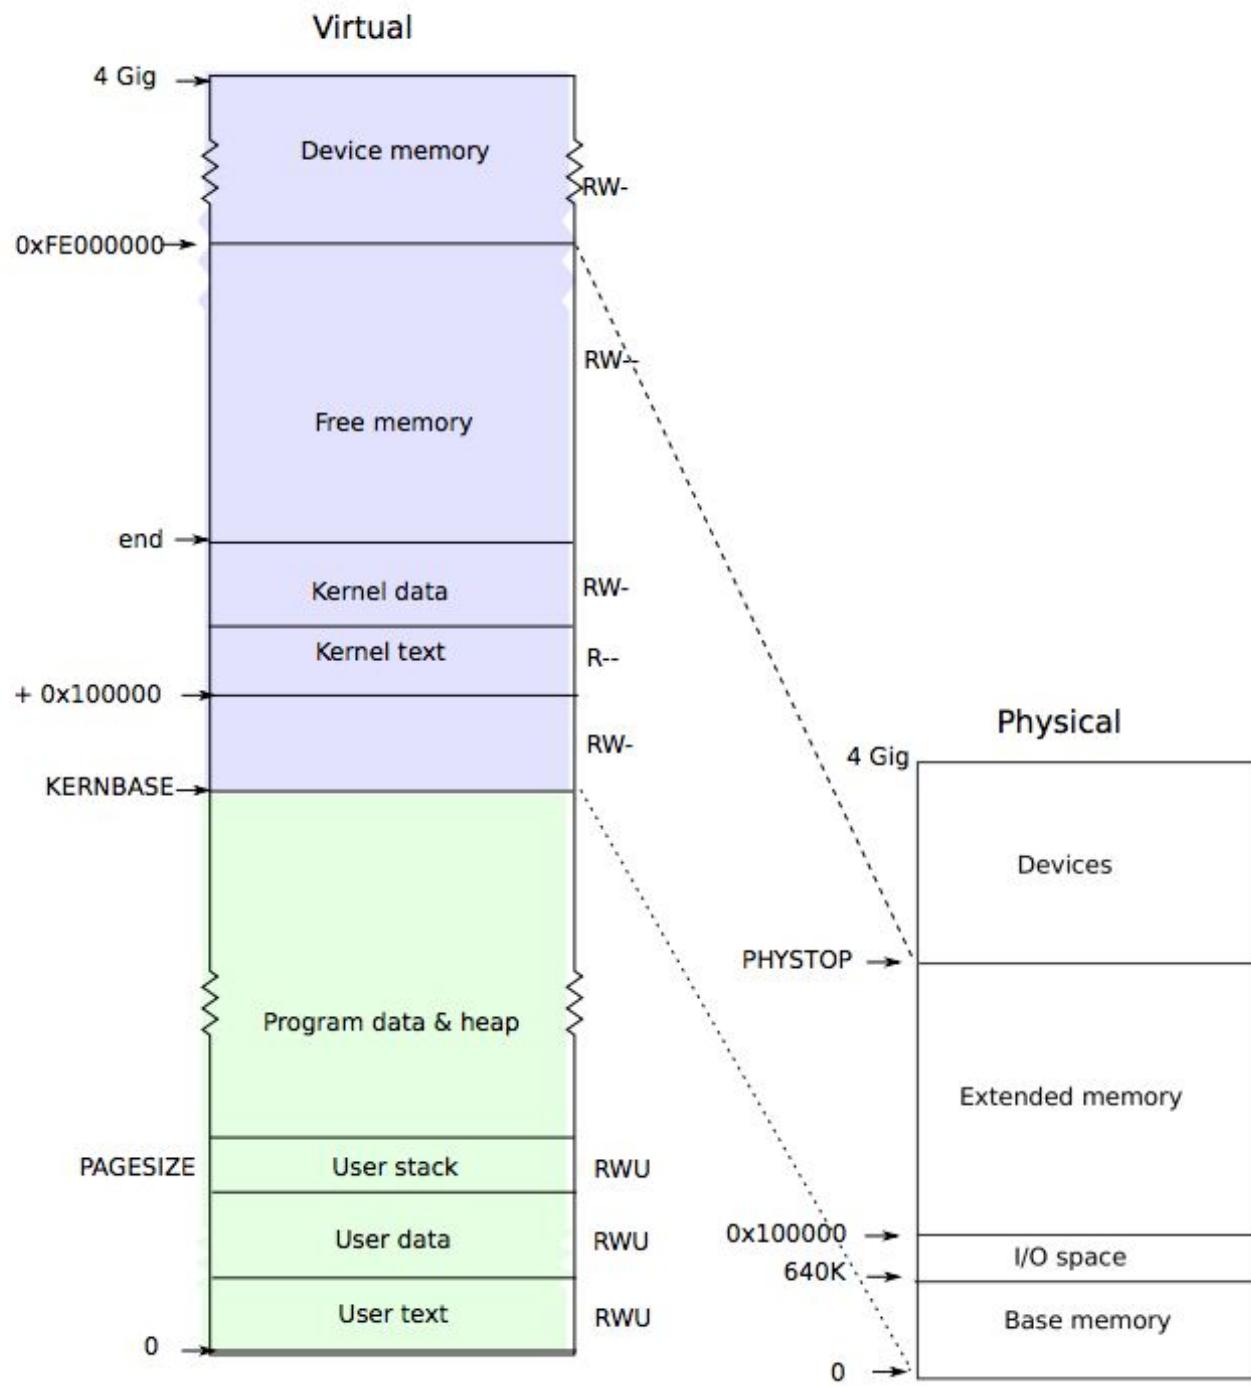
- Do the following calculation, where A, B are int pointers:

```
for(i=0; i<0x40000; ++i){  
    A[i] += B[i] & 0xFFFF;  
}
```

Which sized (4KB or 4MB) page performs better in each scenario?
Give your reasons.

```
17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc();      // kernel page table
22     mpinit();        // collect info about this machine
23     lapicinit();
24     seginit();       // set up segments
25     cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
26     picinit();       // interrupt controller
27     ioapicinit();   // another interrupt controller
28     consoleinit();  // I/O devices & their interrupts
29     uartinit();     // serial port
30     pinit();         // process table
31     tvinit();        // trap vectors
32     binit();         // buffer cache
33     fileinit();     // file table
34     iinit();         // inode cache
35     ideinit();      // disk
36     if(!ismp)
37         timerinit(); // uniprocessor timer
38     startothers();  // start other processors
39     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
40     userinit();     // first user process
41     // Finish setting up this processor in mpmain.
42     mpmain();
43 }
```

```
10 extern char data[]; // defined by kernel.ld
11 pde_t *kpgdir; // for use in scheduler()
12 struct segdesc gdt[NSEGS];
13
14 // Set up CPU's kernel segment descriptors.
15 // Run once on entry on each CPU.
16 void
17 seginit(void)
18 {
19     struct cpu *c;
20
21     // Map "logical" addresses to virtual addresses using identity map.
22     // Cannot share a CODE descriptor for both kernel and user
23     // because it would have to have DPL_USR, but the CPU forbids
24     // an interrupt from CPL=0 to CPL=3.
25     c = &cpus[cpunum()];
26     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
27     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
28     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
29     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
30
31     // Map cpu, and curproc
32     c->gdt[SEG_KCPU] = SEG(STA_W, &c->cpu, 8, 0);
33
34     lgdt(c->gdt, sizeof(c->gdt));
35     loadgs(SEG_KCPU << 3);
36
37     // Initialize cpu-local storage.
38     cpu = c;
39     proc = 0;
40 }
```



big picture of xv6's virtual addressing scheme

0x00000000:0x80000000

 user addresses below KERNBASE

0x80000000:0x80100000

 map low 1MB devices (for kernel)

0x80100000:?

 kernel instructions/data

? :0x8E000000

 224 MB of DRAM mapped here

0xFE000000:0x00000000

 more memory-mapped devices

4 KB pages

entrypgdir was simple

array of 1024 PDEs, each mapping 4 MB but 4 MB is too large a page size!

very wasteful if you have small processes xv6 programs are a few dozen kilobytes

4 MB pages require allocating full 4 MB of phys mem
solution: x86 MMU supports 4 KB pages

kvmalloc

```
17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc(); // kernel page table
22     mpinit(); // collect info about this machine

146 // Allocate one page table for the machine for the kernel address
147 // space for scheduler processes.
148 void
149 kvmalloc(void)
150 {
151     kpgdir = setupkvm();
152     switchkvm();
153 }

155 // Switch h/w page table register to the kernel-only page table,
156 // for when no process is running.
157 void
158 switchkvm(void)
159 {
160     lcr3(v2p(kpgdir)); // switch to the kernel page table
161 }
```

Setup_kvm

Usage

creates a page table

install mappings the kernel will need

will be called once for each new process

Process

must allocate PD, allocate some PTs, put mappings in PTEs

alloc allocates a physical page for the PD

returns that page's virtual address above 0x8000000

so we'll have to call V2P before installing in cr3

memset so that default is no translation (no P bit)

a call to mappages for each entry in kmap[]

```
127 // Set up kernel part of a page table.  
128 pde_t*  
129 setupkvm(void)  
130 {  
131     pde_t *pgdir;  
132     struct kmap *k;  
133  
134     if((pgdir = (pde_t*)kalloc()) == 0)  
135         return 0;  
136     memset(pgdir, 0, PGSIZE);  
137     if (p2v(PHYSTOP) > (void*)DEVSPACE)  
138         panic("PHYSTOP too high");  
139     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)  
140         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,  
141                     (uint)k->phys_start, k->perm) < 0)  
142             return 0;  
143     return pgdir;  
144 }
```

what is kmap[]?

an entry for each region of kernel virtual address space

same as address space setup from earlier in lecture

0x80000000 -> 0x00000000 (memory-mapped devices)

0x80100000 -> 0x00100000 (kernel instructions and data)

data -> data-0x80000000 (phys mem after where kernel was loaded)

DEVSPACE -> DEVSPACE (more memory mapped devices)

note no mappings below va 0x80000000 -- future user memory there

```
113 // This table defines the kernel's mappings, which are present in
114 // every process's page table.
115 static struct kmap {
116     void *virt;
117     uint phys_start;
118     uint phys_end;
119     int perm;
120 } kmap[] = {
121 { (void*)KERNBASE, 0,           EXTMEM,    PTE_W}, // I/O space
122 { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0}, // kern text+rodata
123 { (void*)data,      V2P(data),    PHYSTOP,   PTE_W}, // kern data+memory
124 { (void*)DEVSPACE, DEVSPACE,    0,          PTE_W}, // more devices
125 };
```

```
67 // Create PTEs for virtual addresses starting at va that refer to
68 // physical addresses starting at pa. va and size might not
69 // be page-aligned.
70 static int
71 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
72 {
73     char *a, *last;
74     pte_t *pte;
75
76     a = (char*)PGROUNDDOWN((uint)va);
77     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
78     for(;;){
79         if((pte = walkpgdir(pgdir, a, 1)) == 0)
80             return -1;
81         if(*pte & PTE_P)
82             panic("remap");
83         *pte = pa | perm | PTE_P;
84         if(a == last)
85             break;
86         a += PGSIZE;
87         pa += PGSIZE;
88     }
89     return 0;
90 }
```

mappages

arguments are PD, va, size, pa, perm

adds mappings from a range of va's to corresponding pa's

rounds b/c other uses pass in non-page-aligned addresses

for each page-aligned address in the range

call walkpgdir to find address of PTE need the PTE's address (not just content) b/c we want to modify

```
42 // Return the address of the PTE in page table pgdir
43 // that corresponds to virtual address va.  If alloc!=0,
44 // create any required page table pages.
45 static pte_t *
46 walkpgdir(pde_t *pgdir, const void *va, int alloc)
47 {
48     pde_t *pde;
49     pte_t *pgtab;
50
51     pde = &pgdir[PDX(va)];
52     if(*pde & PTE_P){
53         pgtab = (pte_t*)p2v(PTE_ADDR(*pde));
54     } else {
55         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
56             return 0;
57         // Make sure all those PTE_P bits are zero.
58         memset(pgtab, 0, PGSIZE);
59         // The permissions here are overly generous, but they can
60         // be further restricted by the permissions in the page table
61         // entries, if necessary.
62         *pde = v2p(pgtab) | PTE_P | PTE_W | PTE_U;
63     }
64     return &pgtab[PTX(va)];
65 }
```

```
13 static inline uint v2p(void *a) { return ((uint)(a)) - KERNBASE; }
14 static inline void *p2v(uint a) { return (void*)((a) + KERNBASE); }
```

walkpgdir

- mimics how the paging h/w finds the PTE for an address
- a little complex b/c xv6 allocates page-table pages lazily
 - might not be present, might have to allocate
- PDX extracts top ten bits
- &pgdir[PDX(va)] is the address of the relevant PDE
- now *pde is the PDE
- if PTE_P
 - the relevant page-table page already exists
 - PTE_ADDR extracts the PPN from the PDE
 - p2v() adds 0x80000000, since PTE holds physical address
- if not PTE_P
 - alloc a page-table page
 - fill in PDE with PPN -- thus v2p
- now the PTE we want is in the page-table page
 - at offset PTX(va) , which is 2nd 10 bits of va

```
16 // doing some setup required for memory allocator to work.
17 int
18 main(void)
19 {
20     kinit1(end, P2V(4*1024*1024)); // phys page allocator
21     kvmalloc();      // kernel page table
22     mpinit();        // collect info about this machine
23     lapicinit();
24     seginit();        // set up segments
25     cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
26     picinit();        // interrupt controller
27     ioapicinit();    // another interrupt controller
28     consoleinit();   // I/O devices & their interrupts
29     uartinit();       // serial port
30     pinit();          // process table
31     tvinit();         // trap vectors
32     binit();          // buffer cache
33     fileinit();       // file table
34     iinit();          // inode cache
35     ideinit();        // disk
36     if(!ismp)
37         timerinit();   // uniprocessor timer
38     startothers();    // start other processors
39     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
40     userinit();       // first user process
41     // Finish setting up this processor in mpmain.
42     mpmain();
43 }
```

```
78 void
79 userinit(void)
80 {
81     struct proc *p;
82     extern char _binary_initcode_start[], _binary_initcode_size[];
83
84     p = allocproc();
85     initproc = p;
86     if((p->pgdir = setupkvm()) == 0)
87         panic("userinit: out of memory?");
88     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
89     p->sz = PGSIZE;
90     memset(p->tf, 0, sizeof(*p->tf));
91     p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
92     p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
93     p->tf->es = p->tf->ds;
94     p->tf->ss = p->tf->ds;
95     p->tf->eflags = FL_IF;
96     p->tf->esp = PGSIZE;
97     p->tf->eip = 0; // beginning of initcode.S
98
99     safestrcpy(p->name, "initcode", sizeof(p->name));
100    p->cwd = namei("/");
101
102    p->state = RUNNABLE;
103 }
```

User VM mapping

Create a process

Inituvm : initialized user VM

Loaduvm: load program segment

Context switch

switchuvm

Grow/shrink/free VM

allocuvm

deallocuvm

Freeuvm

Fork

copyuvm

```
179 // Load the initcode into address 0 of pgdir.  
180 // sz must be less than a page.  
181 void  
182 inituvm(pde_t *pgdir, char *init, uint sz)  
183 {  
184     char *mem;  
185  
186     if(sz >= PGSIZE)  
187         panic("inituvm: more than a page");  
188     mem = kalloc();  
189     memset(mem, 0, PGSIZE);  
190     mappages(pgdir, 0, PGSIZE, v2p(mem), PTE_W|PTE_U);  
191     memmove(mem, init, sz);  
192 }
```

```
79 // Allocate one 4096-byte page of physical memory.
80 // Returns a pointer that the kernel can use.
81 // Returns 0 if the memory cannot be allocated.
82 char*
83 kalloc(void)
84 {
85     struct run *r;
86
87     if(kmem.use_lock)
88         acquire(&kmem.lock);
89     r = kmem.freelist;
90     if(r)
91         kmem.freelist = r->next;
92     if(kmem.use_lock)
93         release(&kmem.lock);
94     return (char*)r;
95 }
```

```
// Switch TSS and h/w page table to correspond to process p.
void
switchuvm(struct proc *p)
{
    pushcli();
    cpu->gdt[SEG_TSS] = SEG16(STS_T32A, &cpu->ts, sizeof(cpu->ts)-1, 0);
    cpu->gdt[SEG_TSS].s = 0;
    cpu->ts.ss0 = SEG_KDATA << 3;
    cpu->ts.esp0 = (uint)proc->kstack + KSTACKSIZE;
    ltr(SEG_TSS << 3);
    if(p->pgdir == 0)
        panic("switchuvm: no pgdir");
    lcr3(v2p(p->pgdir)); // switch to new address space
    popcli();
}
```

```
// Load a program segment into pgdir.  addr must be page-aligned
// and the pages from addr to addr+sz must already be mapped.
int
loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz)
{
    uint i, pa, n;
    pte_t *pte;

    if((uint)addr % PGSIZE != 0)
        panic("loaduvm: addr must be page aligned");
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
            panic("loaduvm: address should exist");
        pa = PTE_ADDR(*pte);
        if(sz - i < PGSIZE)
            n = sz - i;
        else
            n = PGSIZE;
        if(readi(ip, p2v(pa), offset+i, n) != n)
            return -1;
    }
    return 0;
}
```

```
// Allocate page tables and physical memory to grow process from oldsz to
// newsz, which need not be page aligned. Returns new size or 0 on error.
int
allocuvvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE){
        mem = kalloc();
        if(mem == 0){
            cprintf("allocuvvm out of memory\n");
            deallocuvvm(pgdir, newsz, oldsz);
            return 0;
        }
        memset(mem, 0, PGSIZE);
        mappages(pgdir, (char*)a, PGSIZE, v2p(mem), PTE_W|PTE_U);
    }
    return newsz;
}
```

```
// Deallocate user pages to bring the process size from oldsz to
// newsz.  oldsz and newsz need not be page-aligned, nor does newsz
// need to be less than oldsz.  oldsz can be larger than the actual
// process size.  Returns the new process size.
int
deallocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    pte_t *pte;
    uint a, pa;

    if(newsz >= oldsz)
        return oldsz;

    a = PGROUNDUP(newsz);
    for(; a < oldsz; a += PGSIZE){
        pte = walkpgdir(pgdir, (char*)a, 0);
        if(!pte)
            a += (NPTENTRIES - 1) * PGSIZE;
        else if((*pte & PTE_P) != 0){
            pa = PTE_ADDR(*pte);
            if(pa == 0)
                panic("kfree");
            char *v = p2v(pa);
            kfree(v);
            *pte = 0;
        }
    }
    return newsz;
}
```

```
// Free a page table and all the physical memory pages
// in the user part.
void
freevm(pde_t *pgdir)
{
    uint i;

    if(pgdir == 0)
        panic("freevm: no pgdir");
    deallocuvm(pgdir, KERNBASE, 0);
    for(i = 0; i < NPENTRIES; i++){
        if(pgdir[i] & PTE_P){
            char * v = p2v(PTE_ADDR(pgdir[i]));
            kfree(v);
        }
    }
    kfree((char*)pgdir);
}
```

```
// Given a parent process's page table, create a copy
// of it for a child.
pde_t*
copyuvm(pde_t *pgdir, uint sz)
{
    pde_t *d;
    pte_t *pte;
    uint pa, i;
    char *mem;

    if((d = setupkvm()) == 0)
        return 0;
    for(i = 0; i < sz; i += PGSIZE){
        if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
            panic("copyuvm: pte should exist");
        if(!(pte & PTE_P))
            panic("copyuvm: page not present");
        pa = PTE_ADDR(*pte);
        if((mem = kalloc()) == 0)
            goto bad;
        memmove(mem, (char*)p2v(pa), PGSIZE);
        if(mappages(d, (void*)i, PGSIZE, v2p(mem), PTE_W|PTE_U) < 0)
            goto bad;
    }
    return d;
}

bad:
    freevm(d);
    return 0;
}
```

```
// Map user virtual address to kernel address.
char*
uva2ka(pde_t *pgdir, char *uva)
{
    pte_t *pte;

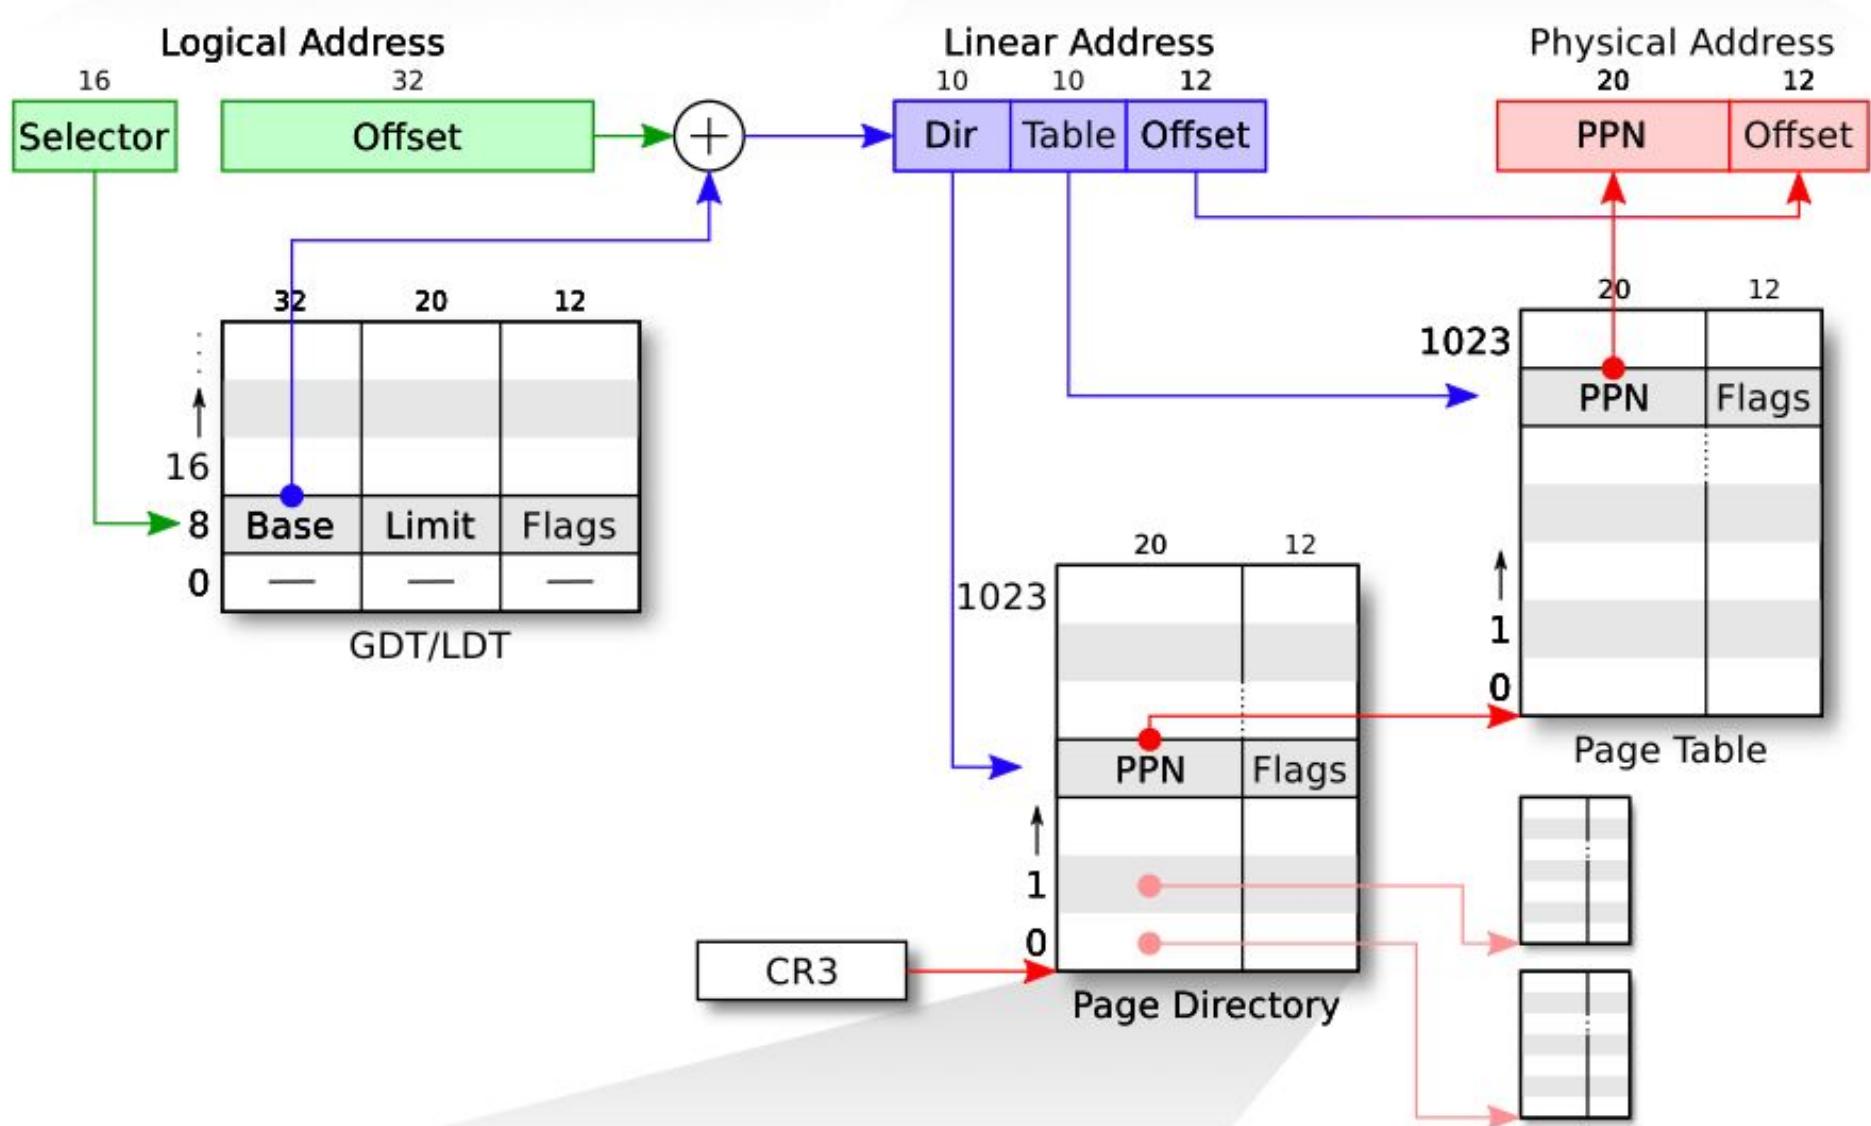
    pte = walkpgdir(pgdir, uva, 0);
    if((*pte & PTE_P) == 0)
        return 0;
    if((*pte & PTE_U) == 0)
        return 0;
    return (char*)p2v(PTE_ADDR(*pte));
}
```

Processes

Rong Chen, Yubin Xia
IPADS, SJTU

ACKs: Some slides are adapted from the textbook's original slides
and frans's os course notes

Review: LA->LA->PA



Review

- Booting: enabling segment and paging
 - Segment: set GDT & use long jmp (ljmp)
 - Paging: temporary page table (entrypgdir) with kernel mapped at both 0x80100000 and 0x100000
- Different page size: 4KB vs. 4MB(2MB)

Question in Last Class

```
39 # Switch from real to protected mode.  Use a bootstrap GDT that makes
40 # virtual addresses map directly to physical addresses so that the
41 # effective memory map doesn't change during the transition.
42 lgdt    gtdesc
43 movl    %cr0, %eax
44 orl    $CR0_PE, %eax
45 movl    %eax, %cr0
46
47 //PAGEBREAK!
48 # Complete transition to 32-bit protected mode by using long jmp
49 # to reload %cs and %eip.  The segment descriptors are set up with no
50 # translation, so that the mapping is still the identity mapping.
51 ljmp    $(SEG_KCODE<<3), $start32
52
53 .code32 # Tell assembler to generate 32-bit code now.
54 start32:
55 # Set up the protected-mode data segment registers
56 movw    $(SEG_KDATA<<3), %ax    # Our data segment selector
57 movw    %ax, %ds                # -> DS: Data Segment
58 movw    %ax, %es                # -> ES: Extra Segment
59 movw    %ax, %ss                # -> SS: Stack Segment
60 movw    $0, %ax                # Zero segments not ready for use
61 movw    %ax, %fs                # -> FS
62 movw    %ax, %gs                # -> GS
```

Explanation

- Intel's manual 9.9.1
 - The JMP or CALL instruction immediately after the MOV CR0 instruction changes the flow of execution and serializes the processor.
- Xv6-book-rev7
 - Enabling protected mode does not immediately change how the processor translates logical to physical addresses; it is only when one loads a new value into a segment register that the processor reads the GDT and changes its internal segmentation settings. One cannot directly modify %cs, so instead the code executes an ljmp (far jump) instruction (8453), which allows a code segment selector to be specified. The jump continues execution at the next line (8456) but in doing so sets %cs to refer to the code descriptor entry in gdt. That descriptor describes a 32-bit code segment, so the processor switches into 32-bit mode.

Super Page: 4M Page

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0						
Address of page directory ¹												Ignored				P	P	C	W	T	Ignored								CR3								
Bits 31:22 of address of 2MB page frame				Reserved (must be 0)				Bits 39:32 of address ²				P	A	Ignored	G	1	D	A	P	P	U	R	/	1		PDE: 4MB page											
Address of page table												Ignored				Q	I	g	A	P	P	U	R	/	1		PDE: page table										
Ignored																Q											Q	PDE: not present									
Address of 4KB page frame												Ignored				G	P	A	D	A	P	P	U	R	/	1		PTE: 4KB page									
Ignored																											Q	PTE: not present									

PROCESS IN XV6

Processes Outline

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems

Process: Big Picture

- More programs than processors
 - How to share the limited number of processors among the programs?
- Observation: most programs don't need the processor continuously
 - because they frequently have to wait for input (from user, disk, network, etc.)
- Idea: when one program must wait, it releases the processor, and gives it to another program

Threads

- Mechanism: thread of computation, an active computation.
 - A thread is an abstraction that contains
 - the minimal state that is necessary to **stop** an active and **resume** it at some point later.
 - What that state is depends on the processor
 - On x86, it is the processor registers.

Threads in xv6

```
34 // Saved registers for kernel context switches.  
35 // Don't need to save all the segment registers (%cs, etc),  
36 // because they are constant across kernel contexts.  
37 // Don't need to save %eax, %ecx, %edx, because the  
38 // x86 convention is that the caller has saved them.  
39 // Contexts are stored at the bottom of the stack they  
40 // describe; the stack pointer is the address of the context.  
41 // The layout of the context matches the layout of the stack in swtch.S  
42 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,  
43 // but it is on the stack and allocproc() manipulates it.  
44 struct context {  
45     uint edi;  
46     uint esi;  
47     uint ebx;  
48     uint ebp;  
49     uint eip;  
50 };
```

Address spaces and threads

- Address spaces and threads are in principle **independent** concepts
 - One can switch from one thread to another thread in the same address space
 - or one can switch from one thread to another thread in another address space.

Process Concepts

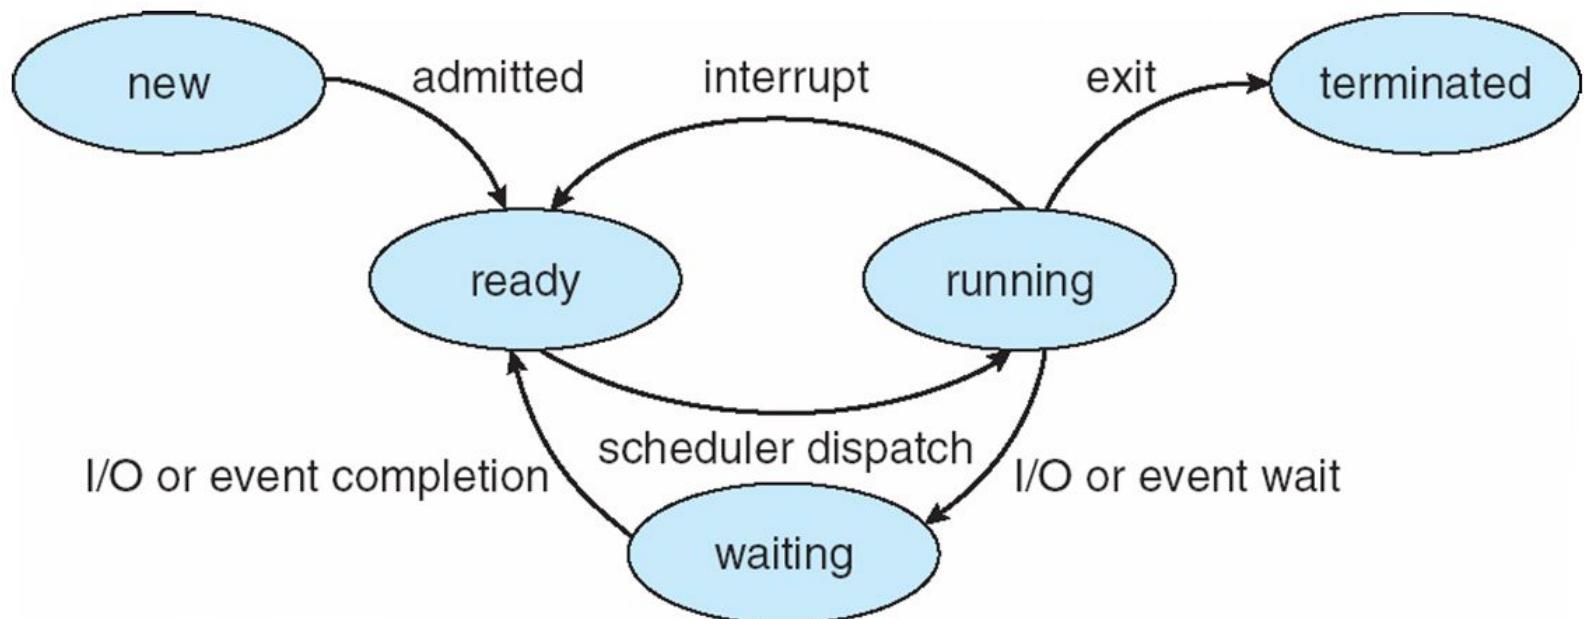
- Process: a program in execution
 - one address space + one or more threads of computation
- A process includes:
 - program counter, stack, data section
- In xv6 all *user* programs contain one thread of computation and one address space

Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution
- In xv6

```
enum proc_state { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
```

Diagram of Process State

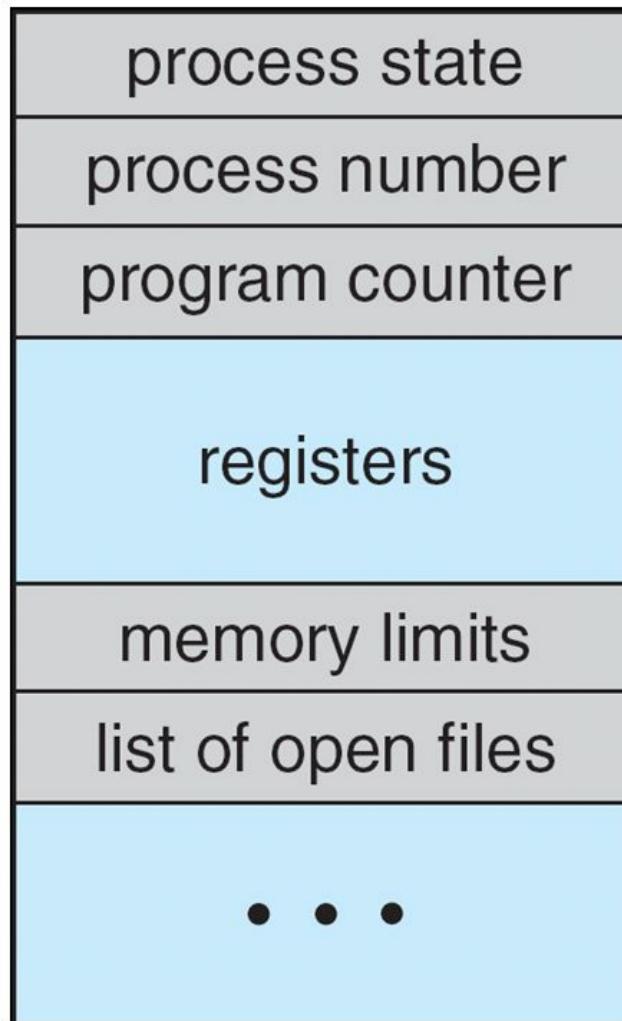


Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

Process Control Block (PCB)

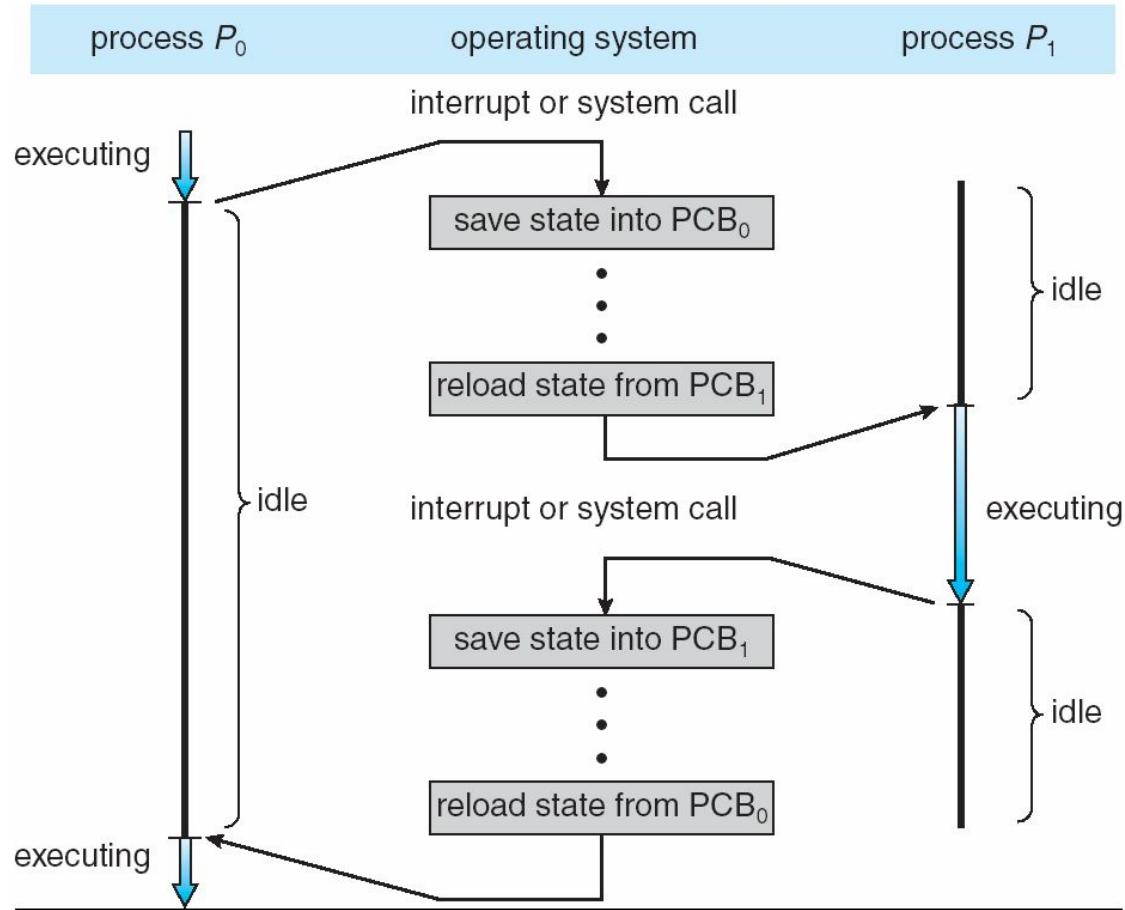


Process in xv6 (proc.c)

Context Switch

- When CPU switches to another process, the system must do it via a **context switch**
 - save the state of the old process
 - load the saved state for the new process **Context** of a process represented in the PCB
- Context-switch time is overhead
 - the system does no useful work while switching
 - Time dependent on hardware support

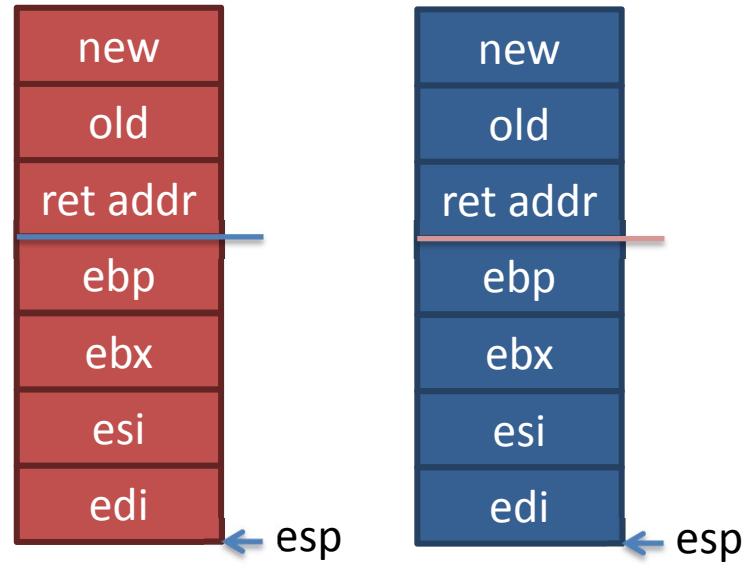
CPU Switch From Process to Process



Context Switch in xv6 (swtch.S)

```
5 # Save current register context in old
6 # and then load register context from new.
7
8 .globl swtch
9 swtch:
10    movl 4(%esp), %eax
11    movl 8(%esp), %edx
12
13    # Save old callee-save registers
14    pushl %ebp
15    pushl %ebx
16    pushl %esi
17    pushl %edi
18
19    # Switch stacks
20    movl %esp, (%eax)
21    movl %edx, %esp
22
23    # Load new callee-save registers
24    popl %edi
25    popl %esi
26    popl %ebx
27    popl %ebp
28    ret
```

```
# Context switch
#
#     void swtch(struct context **old, struct context *new);
#
```



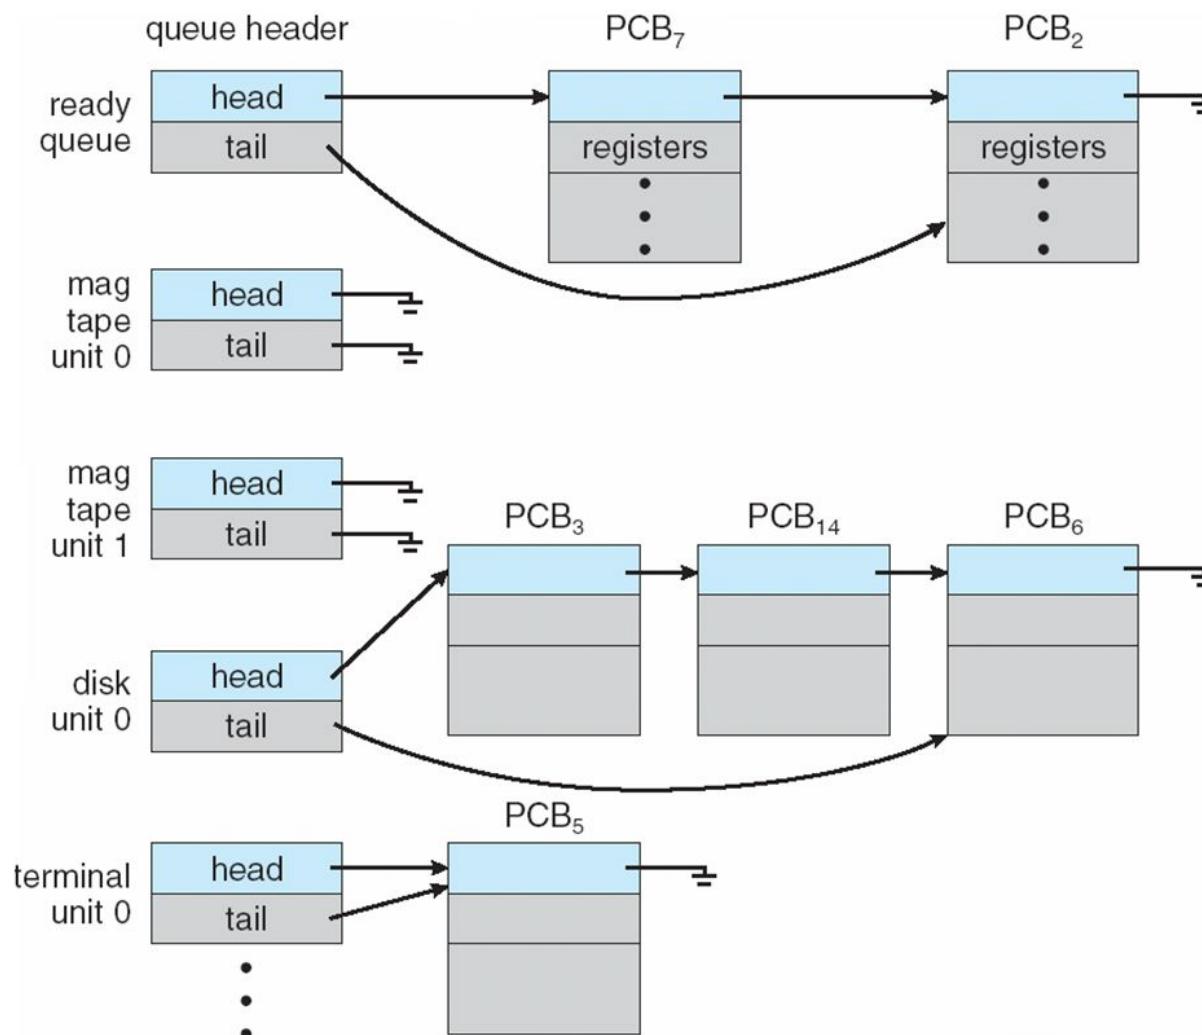
```
258 scheduler(void)
259 {
260     struct proc *p;
261
262     for(;;){
263         // Enable interrupts on this processor.
264         sti();
265
266         // Loop over process table looking for process to run.
267         acquire(&ptable.lock);
268         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
269             if(p->state != RUNNABLE)
270                 continue;
271
272             // Switch to chosen process. It is the process's job
273             // to release ptable.lock and then reacquire it
274             // before jumping back to us.
275             proc = p;
276             switchuvm(p);
277             p->state = RUNNING;
278             swtch(&cpu->scheduler, proc->context); ←
279             switchkvm(); ←
280
281             // Process is done running for now.
282             // It should have changed its p->state before coming back.
283             proc = 0;
284         }
285         release(&ptable.lock);
286     }
287 }
288 }
```

```
290 // Enter scheduler. Must hold only ptable.lock
291 // and have changed proc->state.
292 void
293 sched(void)
294 {
295     int intena;
296
297     if(!holding(&ptable.lock))
298         panic("sched ptable.lock");
299     if(cpu->ncli != 1)
300         panic("sched locks");
301     if(proc->state == RUNNING)
302         panic("sched running");
303     if(readeflags()&FL_IF)
304         panic("sched interruptible");
305     intena = cpu->intena;
306     swtch(&proc->context, cpu->scheduler); ←
307     cpu->intena = intena;
308 }
```

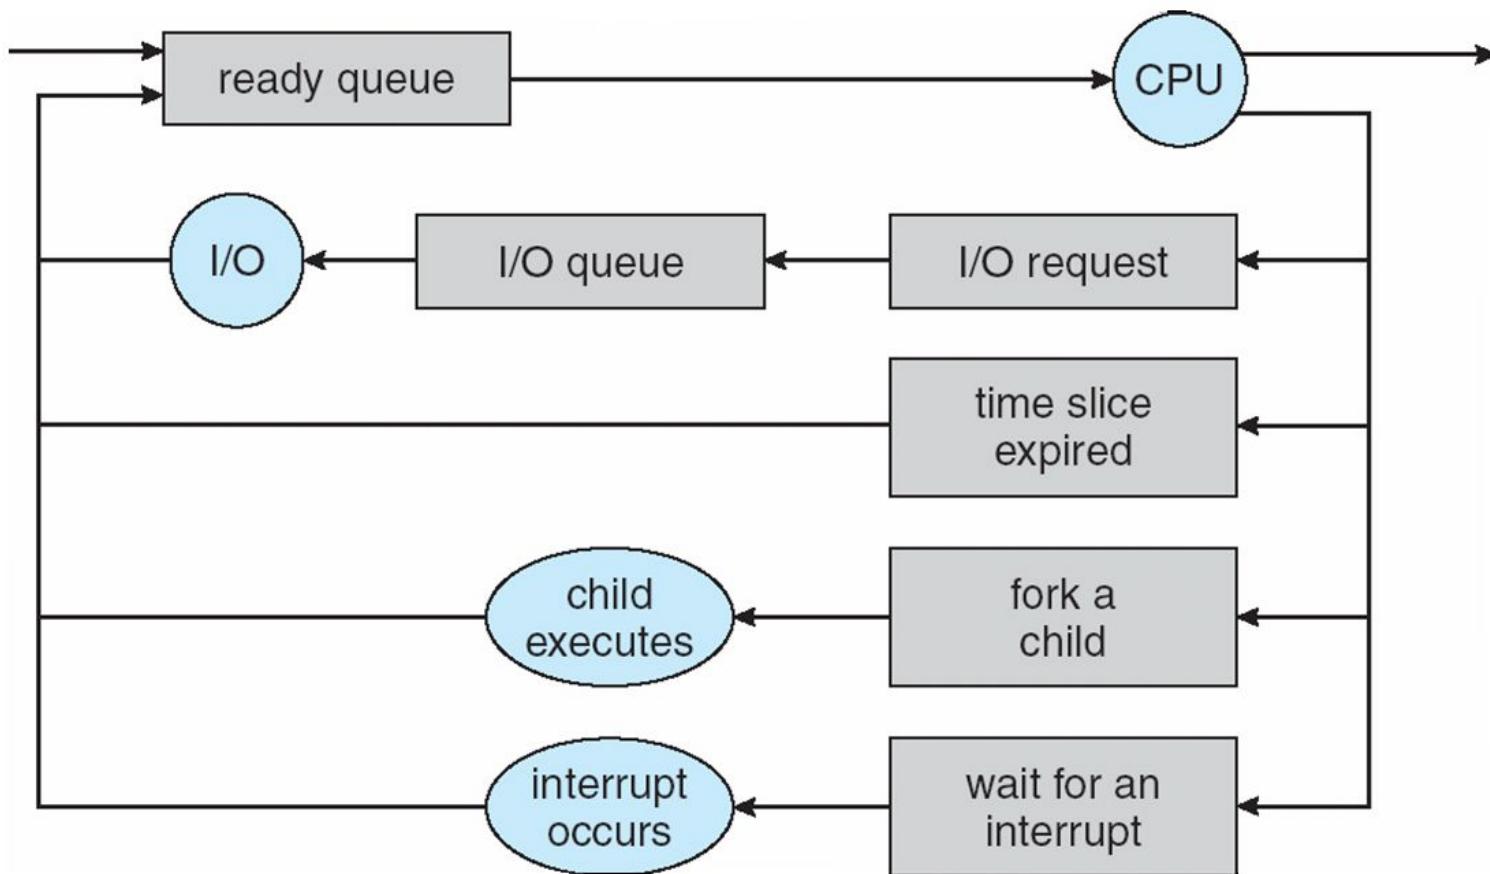
Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues

Ready Queue & Various I/O Device Queues



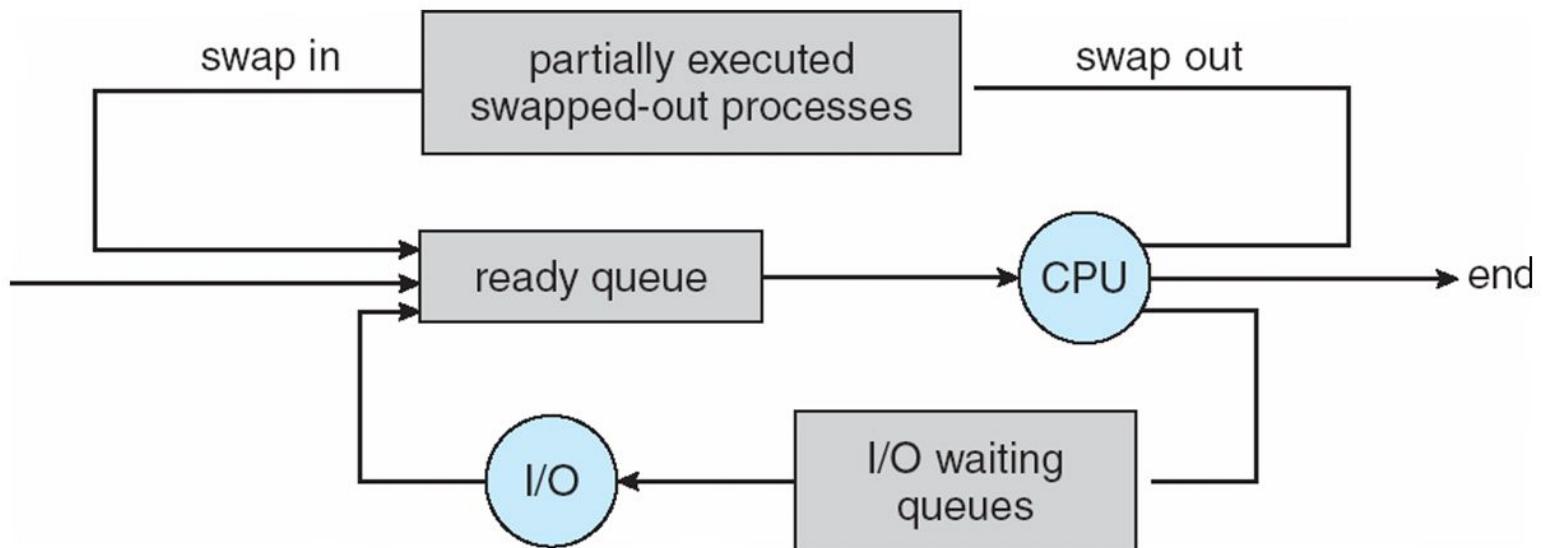
Representation of Process Scheduling



Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU

Addition of Medium Term Scheduling



Schedulers (Cont)

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*

Scheduler (cont.)

- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

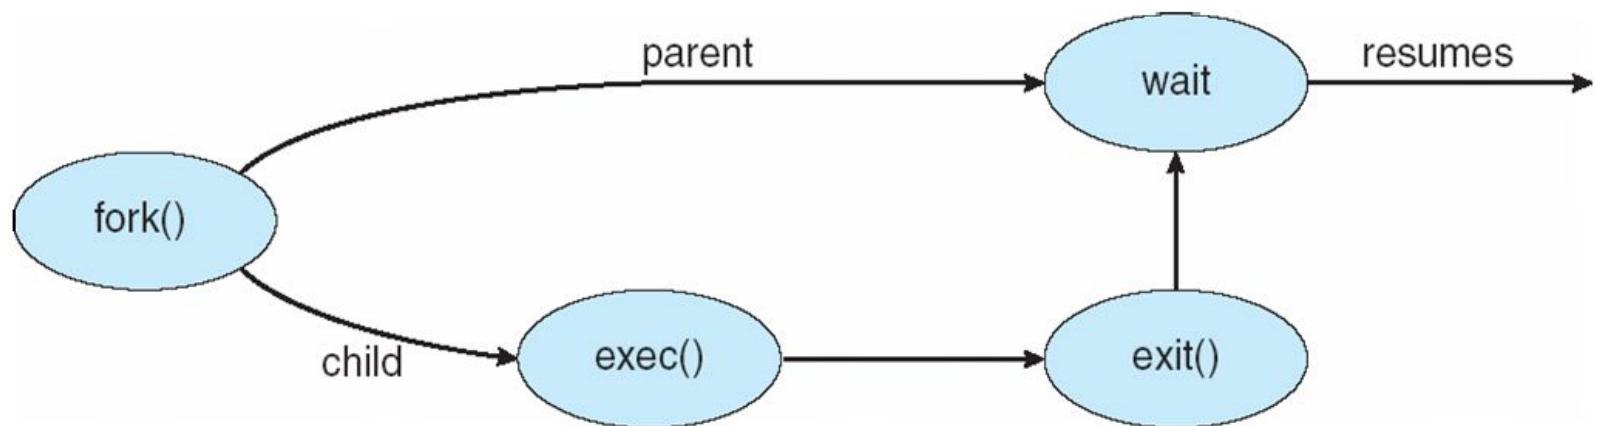
Process Creation

- **Parent** process create **children** processes
 - forming a tree of processes
 - Generally, process identified and managed via **a process identifier (pid)**
- Resource sharing policies
 - share all resources
 - share subset of parent's resources
 - share no resources

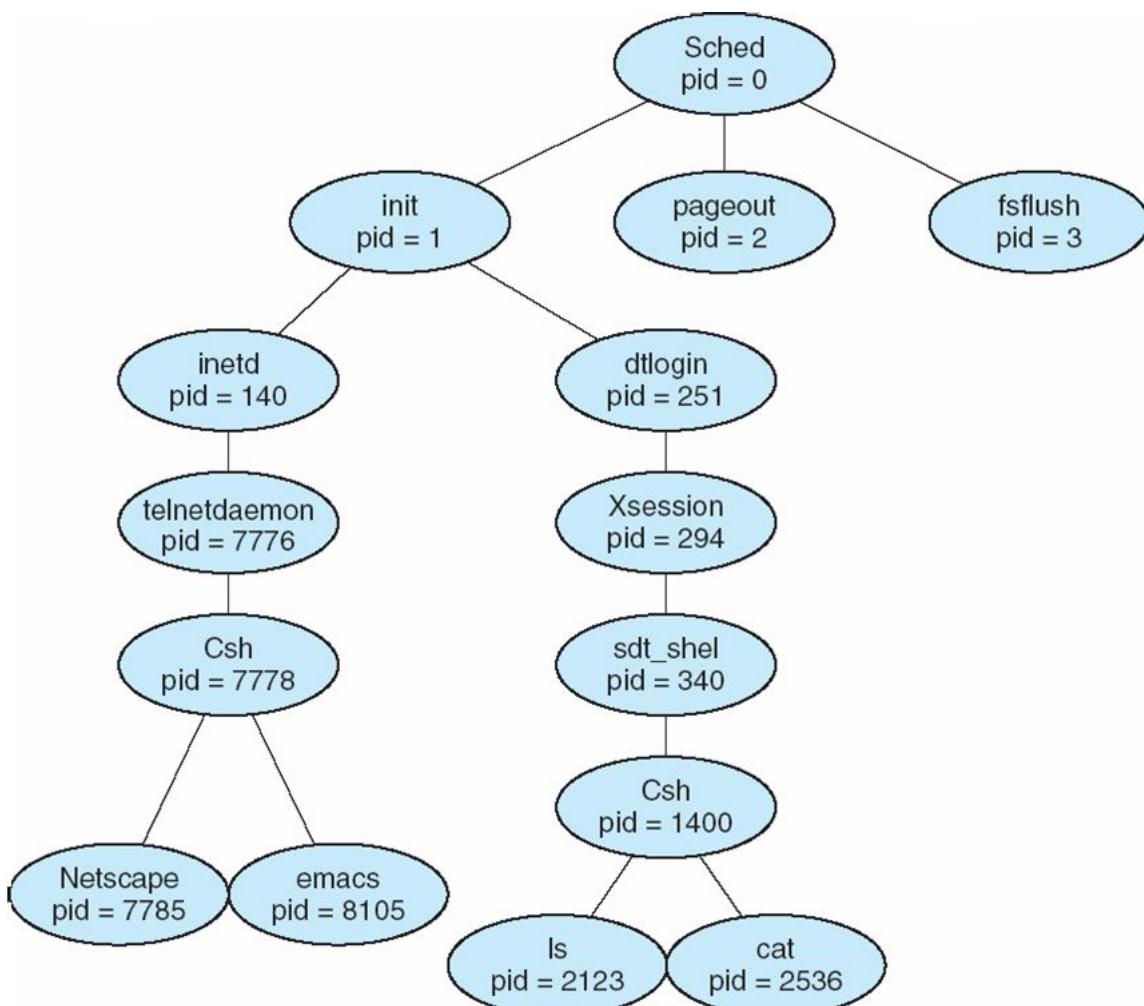
Process Creation (Cont)

- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate
- Address space
 - Child duplicate of parent's
 - Child has a program loaded into it
- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program

Process Creation



A tree of processes on a typical Solaris



Process Creation in XV6

```
125 // Create a new process copying p as the parent.
126 // Sets up stack to return as if from system call.
127 // Caller must set state of returned proc to RUNNABLE.
128 int
129 fork(void)
130 {
131     int i, pid;
132     struct proc *np;
133
134     // Allocate process.
135     if((np = allocproc()) == 0)
136         return -1;
137
138     // Copy process state from p.
139     if((np->pgdir = copyuvm(proc->pgdir, proc->sz)) == 0){
140         kfree(np->kstack);
141         np->kstack = 0;
142         np->state = UNUSED;
143         return -1;
144     }
145     np->sz = proc->sz;
146     np->parent = proc;
147     *np->tf = *proc->tf;
148
149     // Clear %eax so that fork returns 0 in the child.
150     np->tf->eax = 0;
151
152     for(i = 0; i < NOFILE; i++)
153         if(proc->ofile[i])
154             np->ofile[i] = filedup(proc->ofile[i]);
155     np->cwd = idup(proc->cwd);
156
157     pid = np->pid;
158     np->state = RUNNABLE;
159     safestrcpy(np->name, proc->name, sizeof(proc->name));
160     return pid;
161 }
```

Read following code

- allocproc()

Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required

Process Termination (Cont.)

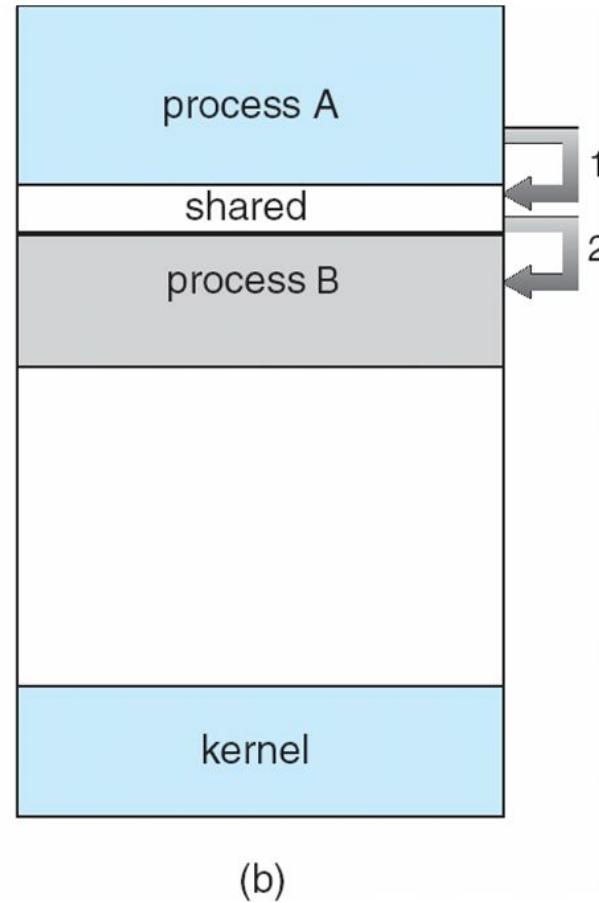
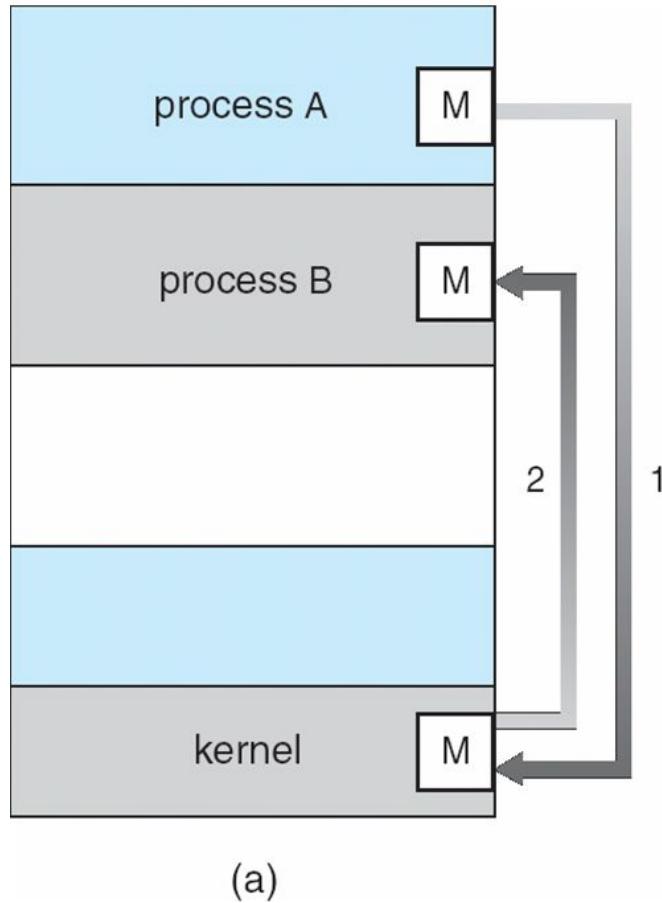
- If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**

```
167 exit(void)
168 {
169     struct proc *p;
170     int fd;
171
172     if(proc == initproc)
173         panic("init exiting");
174
175     // Close all open files.
176     for(fd = 0; fd < NOFILE; fd++){
177         if(proc->ofile[fd]){
178             fileclose(proc->ofile[fd]);
179             proc->ofile[fd] = 0;
180         }
181     }
182
183     iput(proc->cwd);
184     proc->cwd = 0;
185
186     acquire(&ptable.lock);
187
188     // Parent might be sleeping in wait().
189     wakeup1(proc->parent);
190
191     // Pass abandoned children to init.
192     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
193         if(p->parent == proc){
194             p->parent = initproc;
195             if(p->state == ZOMBIE)
196                 wakeup1(initproc);
197         }
198     }
199
200     // Jump into the scheduler, never to return.
201     proc->state = ZOMBIE;
202     sched();
203     panic("zombie exit");
204 }
```

Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Communications Models



IPC

- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - Shared memory
 - Message passing

Producer-Consumer Problem

- *Producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
}
```

```
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- Solution is correct, but can only use BUFFER_SIZE-1 elements

Bounded-Buffer – Producer

```
while (true) {
    /* Produce an item */
    while (((in = (in + 1) % BUFFER SIZE count)
            == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER SIZE;
}
```

Bounded Buffer – Consumer

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```

Example of IPC in xv6: Pipe (pipe.c)

```
12 struct pipe {  
13     struct spinlock lock;  
14     char data[PIPESIZE];  
15     uint nread;      // number of bytes read  
16     uint nwrite;    // number of bytes written  
17     int readopen;   // read fd is still open  
18     int writeopen;  // write fd is still open  
19 };
```

```
77 int
78 pipewrite(struct pipe *p, char *addr, int n)
79 {
80     int i;
81
82     acquire(&p->lock);
83     for(i = 0; i < n; i++){
84         while(p->nwrite == p->nread + PIPESIZE){ //DOC: pipewrite-full
85             if(p->readopen == 0 || proc->killed){
86                 release(&p->lock);
87                 return -1;
88             }
89             wakeup(&p->nread);
90             sleep(&p->nwrite, &p->lock); //DOC: pipewrite-sleep
91         }
92         p->data[p->nwrite++ % PIPESIZE] = addr[i];
93     }
94     wakeup(&p->nread); //DOC: pipewrite-wakeup1
95     release(&p->lock);
96     return n;
97 }
```

```
99 int
100 piperead(struct pipe *p, char *addr, int n)
101 {
102     int i;
103
104     acquire(&p->lock);
105     while(p->nread == p->nwrite && p->writeopen){ //DOC: pipe-empty
106         if(proc->killed){
107             release(&p->lock);
108             return -1;
109         }
110         sleep(&p->nread, &p->lock); //DOC: piperead-sleep
111     }
112     for(i = 0; i < n; i++){ //DOC: piperead-copy
113         if(p->nread == p->nwrite)
114             break;
115         addr[i] = p->data[p->nread++ % PIPESIZE];
116     }
117     wakeup(&p->nwrite); //DOC: piperead-wakeup
118     release(&p->lock);
119     return i;
120 }
```

Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system
 - processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(message) – message size fixed or variable
 - **receive**(message)

Message Passing (Cont.)

- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Direct Communication

- Processes must name each other explicitly:
 - **send** (P , message) – send a message to process P
 - **receive**(Q , message) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox

Indirect Communication

- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
send(A, message) – send a message to mailbox A
receive(A, message) – receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver.
Sender is notified who the receiver was.

Synchronization & Asynchronous

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking** send has the sender send the message and continue
 - **Non-blocking** receive has the receiver receive a valid message or null

Buffering

- Queue of messages attached to the link;
implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Examples of IPC Systems - POSIX

- **POSIX Shared Memory**

- Process first creates shared memory segment

```
segment id = shmget(IPC_PRIVATE, size, S  
IRUSR | S_IWUSR);
```

- Process wanting access to that shared memory must attach to it

```
shared memory = (char *) shmat(id, NULL, 0);
```

- Now the process could write to the shared memory

```
sprintf(shared memory, "Writing to shared  
memory");
```

- When done a process can detach the shared memory from its address space

```
shmdt(shared memory);
```

Examples of IPC Systems – WinXP

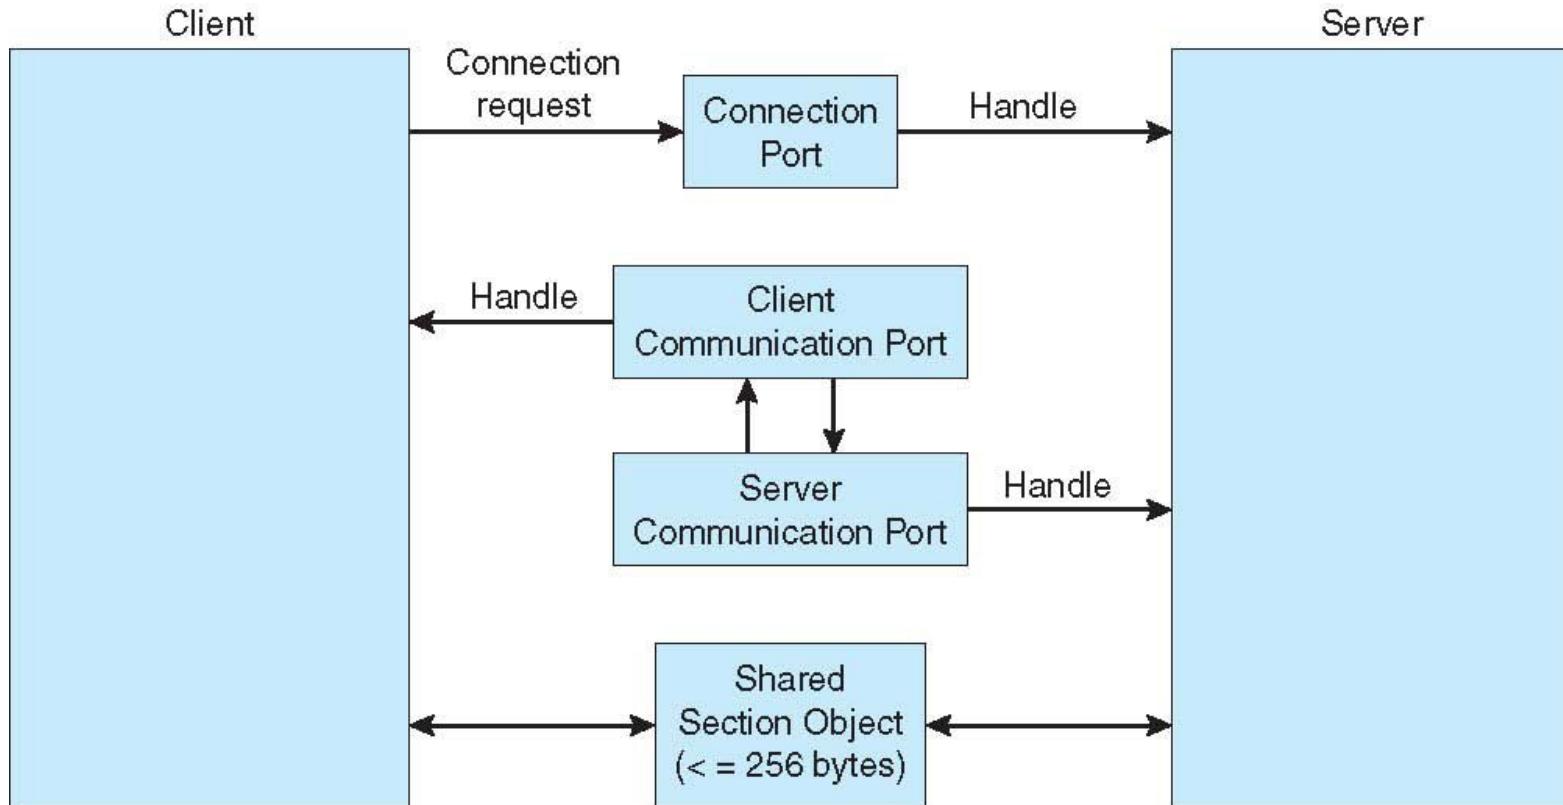
- Message-passing centric via **local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels

Examples of IEC Systems with (cont.)

- Communication works as follows:
 - The client opens a handle to the subsystem's connection port object
 - The client sends a connection request
 - The server creates two private communication ports and returns the handle to one of them to the client
 - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies

Document - Windows Communication Foundation

XP



Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

INTERRUPT, TRAP, SYS CALL

Concepts

- System call
 - a user program can ask for an operating system service
- Exception
 - refers to an illegal program action
- Interrupt
 - refers to a signal generated by a hardware device

Requirements

- Must save the processor's registers for future transparent resume
- Must be set up for execution in the kernel
- Must choose a place for the kernel to start executing
- Must be able to retrieve information about the event, e.g., system call arguments
- Must all be done securely
- Must maintain isolation of user processes and the kernel

IDT

- IDT: interrupt descriptor table
 - Has 256 entries
 - Each giving the %cs and %eip to be used when handling the corresponding interrupt
- System call
 - A program invokes *int n*
 - *n* specifies the index into the IDT

System Call

1. Fetch the n 'th descriptor from the IDT, where n is the argument of int.
2. Check that CPL in %cs is \leq DPL, where DPL is the privilege level in the descriptor.
3. Save %esp and %ss in a CPU-internal registers, but only if the target segment selector's PL $<$ CPL.
4. Load %ss and %esp from a task segment descriptor.
5. Push %ss, %esp, %eflags, %cs, %eip,
6. Clear some bits of %eflags
7. Set %cs and %eip to the values in the descriptor

Kernel Stack

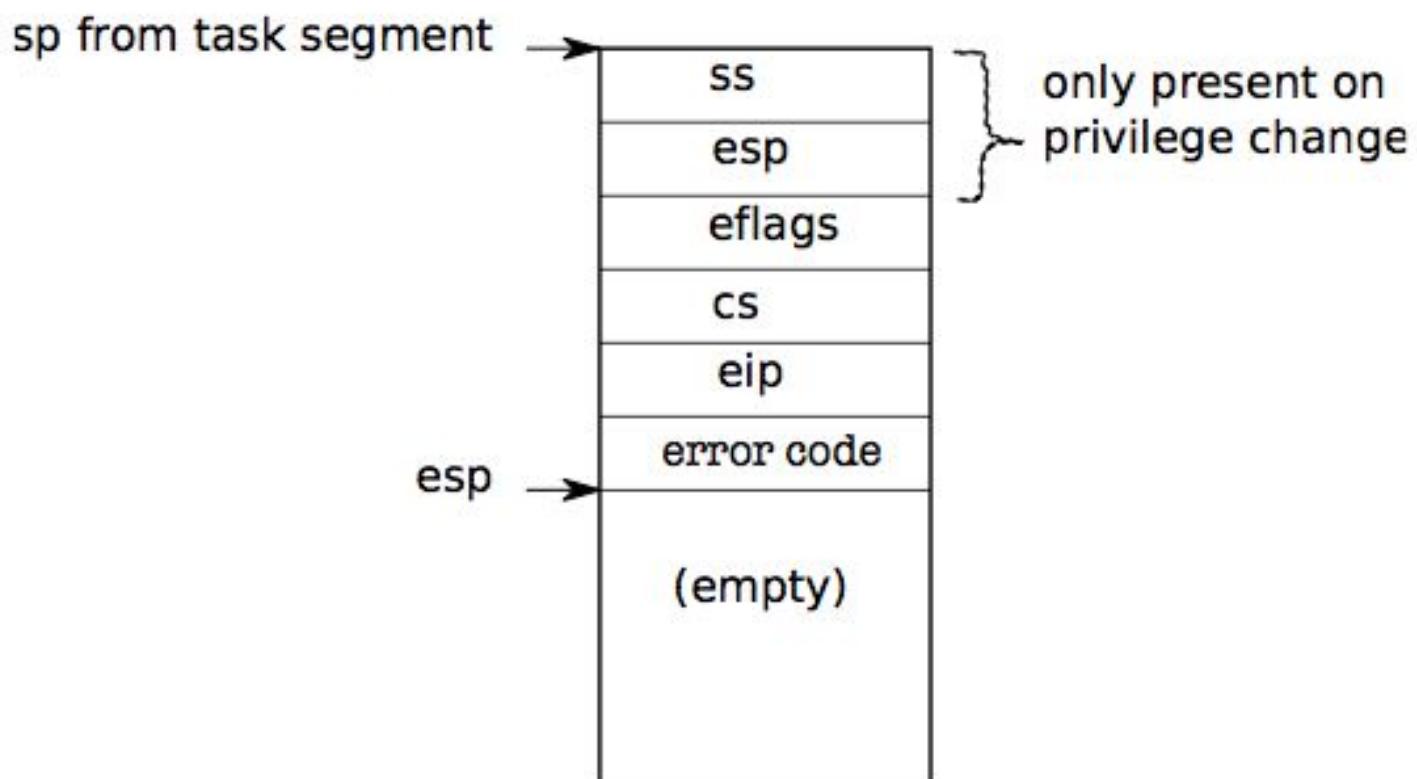


Figure 3-1. Kernel stack after an int instruction.

tvinit

```
17 void
18 tvinit(void)
19 {
20     int i;
21
22     for(i = 0; i < 256; i++)
23         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
24     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
25
26     initlock(&tickslock, "time");
27 }
```

IPC

Yubin Xia, Rong Chen
IPADS, SJTU

ACKs: Some slides are adapted from the textbook's original slides
and Frans's OS course notes

Review: Threads in xv6

```
34 // Saved registers for kernel context switches.  
35 // Don't need to save all the segment registers (%cs, etc),  
36 // because they are constant across kernel contexts.  
37 // Don't need to save %eax, %ecx, %edx, because the  
38 // x86 convention is that the caller has saved them.  
39 // Contexts are stored at the bottom of the stack they  
40 // describe; the stack pointer is the address of the context.  
41 // The layout of the context matches the layout of the stack in swtch.S  
42 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,  
43 // but it is on the stack and allocproc() manipulates it.  
44 struct context {  
45     uint edi;  
46     uint esi;  
47     uint ebx;  
48     uint ebp;  
49     uint eip;  
50 };
```

Review: Process in xv6 (proc.c)

Review: Context Switch in xv6

```
5 # Save current register context in old
6 # and then load register context from new.
7
8 .globl swtch
9 swtch:
10    movl 4(%esp), %eax
11    movl 8(%esp), %edx
12
13    # Save old callee-save registers
14    pushl %ebp
15    pushl %ebx
16    pushl %esi
17    pushl %edi
18
19    # Switch stacks
20    movl %esp, (%eax)
21    movl %edx, %esp
22
23    # Load new callee-save registers
24    popl %edi
25    popl %esi
26    popl %ebx
27    popl %ebp
28    ret
```

```
# Context switch
#
# void swtch(struct context **old, struct context *new);
#
```

The diagram illustrates the stack state during a context switch between two contexts, 'new' and 'old'. It shows two vertical stacks represented as columns of colored boxes. The left column represents the 'new' context, and the right column represents the 'old' context. Arrows point from the labels 'new' and 'old' to their respective columns.

new	old
ret addr	ret addr
ebp	ebp
ebx	ebx
esi	esi
edi	edi

The stack grows downwards. In the 'new' context, the stack pointer (esp) points to the bottom of the stack, which contains the values for edi, esi, ebx, ebp, and the return address. In the 'old' context, the stack pointer (esp) also points to the bottom of the stack, which contains the same values. A blue arrow labeled 'esp' points to the bottom of the 'new' stack, and another blue arrow labeled 'esp' points to the bottom of the 'old' stack.

Review: Context Switch in xv6

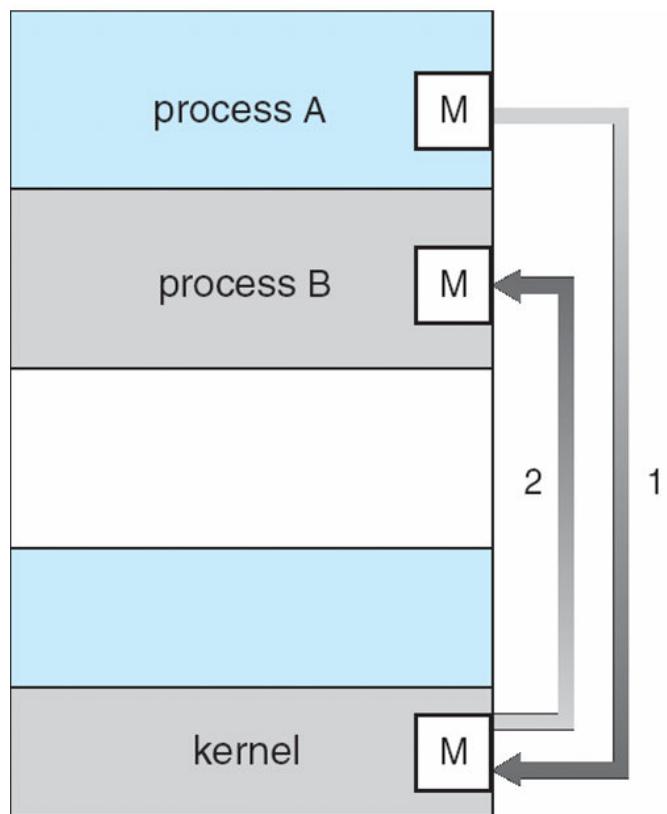
```
258 scheduler(void)
259 {
260     struct proc *p;
261
262     for(;;){
263         // Enable interrupts on this processor.
264         sti();
265
266         // Loop over process table looking for process to run.
267         acquire(&ptable.lock);
268         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
269             if(p->state != RUNNABLE)
270                 continue;
271
272             // Switch to chosen process. It is the process's job
273             // to release ptable.lock and then reacquire it
274             // before jumping back to us.
275             proc = p;
276             switchuvm(p);
277             p->state = RUNNING;
278             swtch(&cpu->scheduler, proc->context); ←
279             switchkvm(); →
280
281             // Process is done running for now.
282             // It should have changed its p->state before coming back.
283             proc = 0;
284         }
285         release(&ptable.lock);
286     }
287 }
```

```
290 // Enter scheduler. Must hold only ptable.lock
291 // and have changed proc->state.
292 void
293 sched(void)
294 {
295     int intena;
296
297     if(!holding(&ptable.lock))
298         panic("sched ptable.lock");
299     if(cpu->ncli != 1)
300         panic("sched locks");
301     if(proc->state == RUNNING)
302         panic("sched running");
303     if(readeflags()&FL_IF)
304         panic("sched interruptible");
305     intena = cpu->intena;
306     swtch(&proc->context, cpu->scheduler); →
307     cpu->intena = intena;
308 }
```

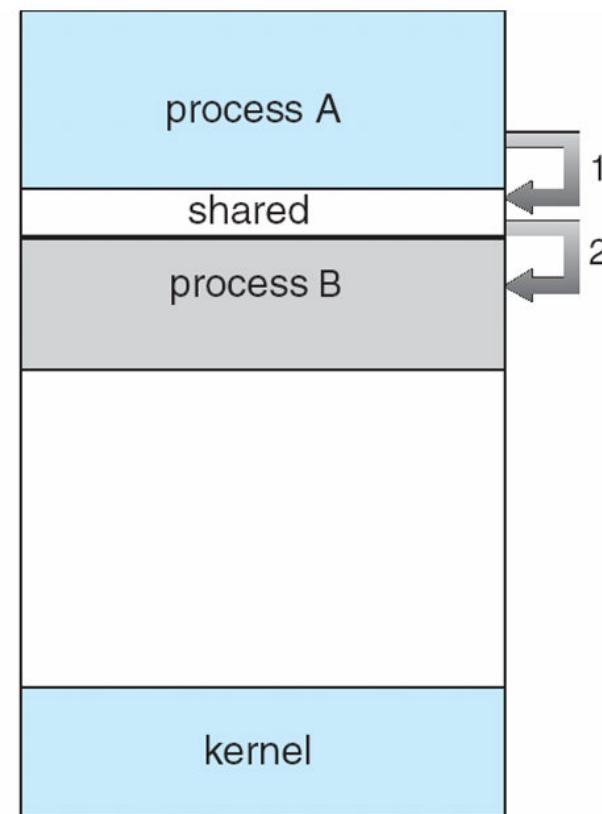
Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

Communications Models



(a)



(b)

IPC

- Cooperating processes need **inter-process communication (IPC)**
- Two models of IPC
 - Shared memory
 - Message passing
- In POSIX
 - Pipe, msg passing, signal, shared memory, semaphore, etc.

Examples of IPC Systems - POSIX

- POSIX Shared Memory

- Process first creates shared memory segment

```
segment id = shmget( IPC PRIVATE, size, S  
IRUSR | S IWUSR );
```

- Process wanting access to that shared memory must attach to it

```
shared memory = (char *) shmat(id, NULL, 0);
```

- Now the process could write to the shared memory

```
sprintf(shared memory, "Writing to shared  
memory");
```

- When done a process can detach the shared memory from its address space

```
shmdt(shared memory);
```

Shared Memory Control

- The **shmctl** system call permits the user to perform a number of generalized control operations on an existing shared memory segment
 - The first, shmid, is a valid shared memory segment identifier generated by a prior `shmget` system call.
 - The second argument, cmd, specifies the operation `shmctl` is to perform.
 - The third argument, buf, is a reference to a structure of the type `shmid_ds`.

Using a File as Shared Memory

- mmap system call can be used to map a file to a process's virtual memory address space.
- In many ways mmap is more flexible than its shared memory system call counterpart.
- Once a mapping has been established, standard system calls rather than specialized system calls can be used to manipulate the shared memory object.
- Unlike memory, the contents of a file are nonvolatile and will remain available even after a system has been shut down (and rebooted).

Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system
 - processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(message) – message size fixed or variable
 - **receive**(message)

Message Passing (Cont.)

- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Direct Communication

- Processes must name each other explicitly:
 - **send** (P , message) – send a msg to process P
 - **receive**(Q , message) – receive a msg from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox

Indirect Communication

- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization & Asynchronous

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking** send has the sender send the message and continue
 - **Non-blocking** receive has the receiver receive a valid message or null

Buffering

- Queue of messages attached to the link;
implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Examples of IPC Systems – WinXP

- Message-passing centric via **local procedure call (LPC) facility**
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels

Producer-Consumer Problem

- *Producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
```

```
    . . .
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

- Solution is correct, but can only use $\text{BUFFER_SIZE}-1$ elements

Bounded-Buffer – Producer

```
while (true) {
    /* Produce an item */
    while (((in = (in + 1) % BUFFER SIZE count)
            == out)
        ; /* do nothing -- no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER SIZE;
}
```

Bounded Buffer – Consumer

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```

Example of IPC in xv6: Pipe (pipe.c)

```
12 struct pipe {  
13     struct spinlock lock;  
14     char data[PIPESIZE];  
15     uint nread;      // number of bytes read  
16     uint nwrite;    // number of bytes written  
17     int readopen;   // read fd is still open  
18     int writeopen;  // write fd is still open  
19 };
```

```
77 int
78 pipewrite(struct pipe *p, char *addr, int n)
79 {
80     int i;
81
82     acquire(&p->lock);
83     for(i = 0; i < n; i++){
84         while(p->nwrite == p->nread + PIPESIZE){ //DOC: pipewrite-full
85             if(p->readopen == 0 || proc->killed){
86                 release(&p->lock);
87                 return -1;
88             }
89             wakeup(&p->nread);
90             sleep(&p->nwrite, &p->lock); //DOC: pipewrite-sleep
91         }
92         p->data[p->nwrite++ % PIPESIZE] = addr[i];
93     }
94     wakeup(&p->nread); //DOC: pipewrite-wakeup1
95     release(&p->lock);
96     return n;
97 }
```

```
99 int
100 piperead(struct pipe *p, char *addr, int n)
101 {
102     int i;
103
104     acquire(&p->lock);
105     while(p->nread == p->nwrite && p->writeopen){ //DOC: pipe-empty
106         if(proc->killed){
107             release(&p->lock);
108             return -1;
109         }
110         sleep(&p->nread, &p->lock); //DOC: piperead-sleep
111     }
112     for(i = 0; i < n; i++){ //DOC: piperead-copy
113         if(p->nread == p->nwrite)
114             break;
115         addr[i] = p->data[p->nread++ % PIPESIZE];
116     }
117     wakeup(&p->nwrite); //DOC: piperead-wakeup
118     release(&p->lock);
119     return i;
120 }
```

Systems

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

INTERRUPT, TRAP, SYS CALL

Concepts

- System call
 - a user program can ask for an operating system service
- Exception
 - refers to an illegal program action
- Interrupt
 - refers to a signal generated by a hardware device

Requirements

- Must save the processor's registers for future transparent resume
- Must be set up for execution in the kernel
- Must choose a place for the kernel to start executing
- Must be able to retrieve information about the event, e.g., system call arguments
- Must all be done securely
- Must maintain isolation of user processes and the kernel

System Call

1. Fetch the n 'th descriptor from the IDT, where n is the argument of int.
2. Check that CPL in %cs is \leq DPL, where DPL is the privilege level in the descriptor.
3. Save %esp and %ss in a CPU-internal registers, but only if the target segment selector's PL $<$ CPL.
4. Load %ss and %esp from a task segment descriptor.
5. Push %ss, %esp, %eflags, %cs, %eip,
6. Clear some bits of %eflags
7. Set %cs and %eip to the values in the descriptor

Kernel Stack

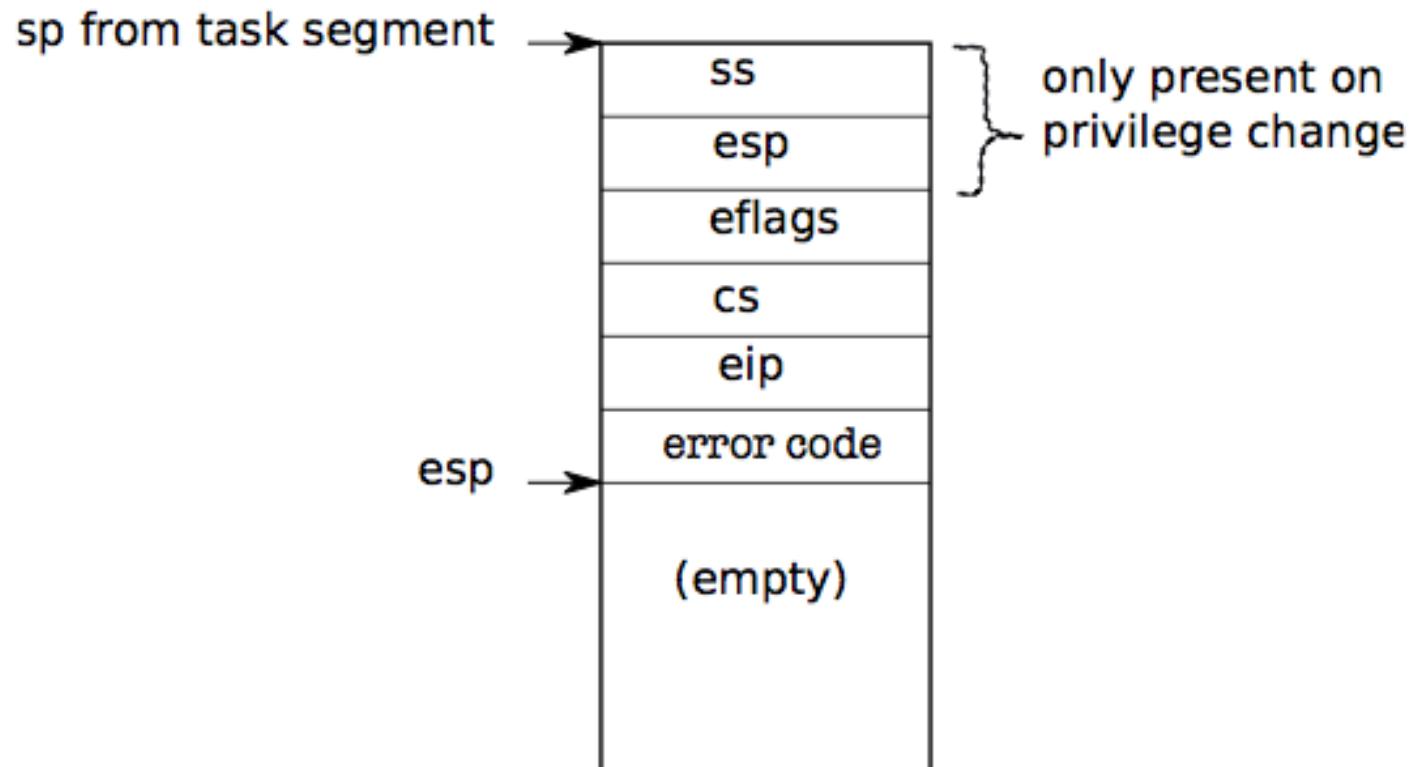


Figure 3-1. Kernel stack after an int instruction.

Interrupts

- Interrupts are similar to system calls, except devices generate them at any time
- There is hardware on the motherboard to signal the CPU when a device needs attention (e.g., the user has typed a character on the keyboard)
- We must program the device to generate an interrupt, and arrange that a CPU receives the interrupt.

Vectors, IDT

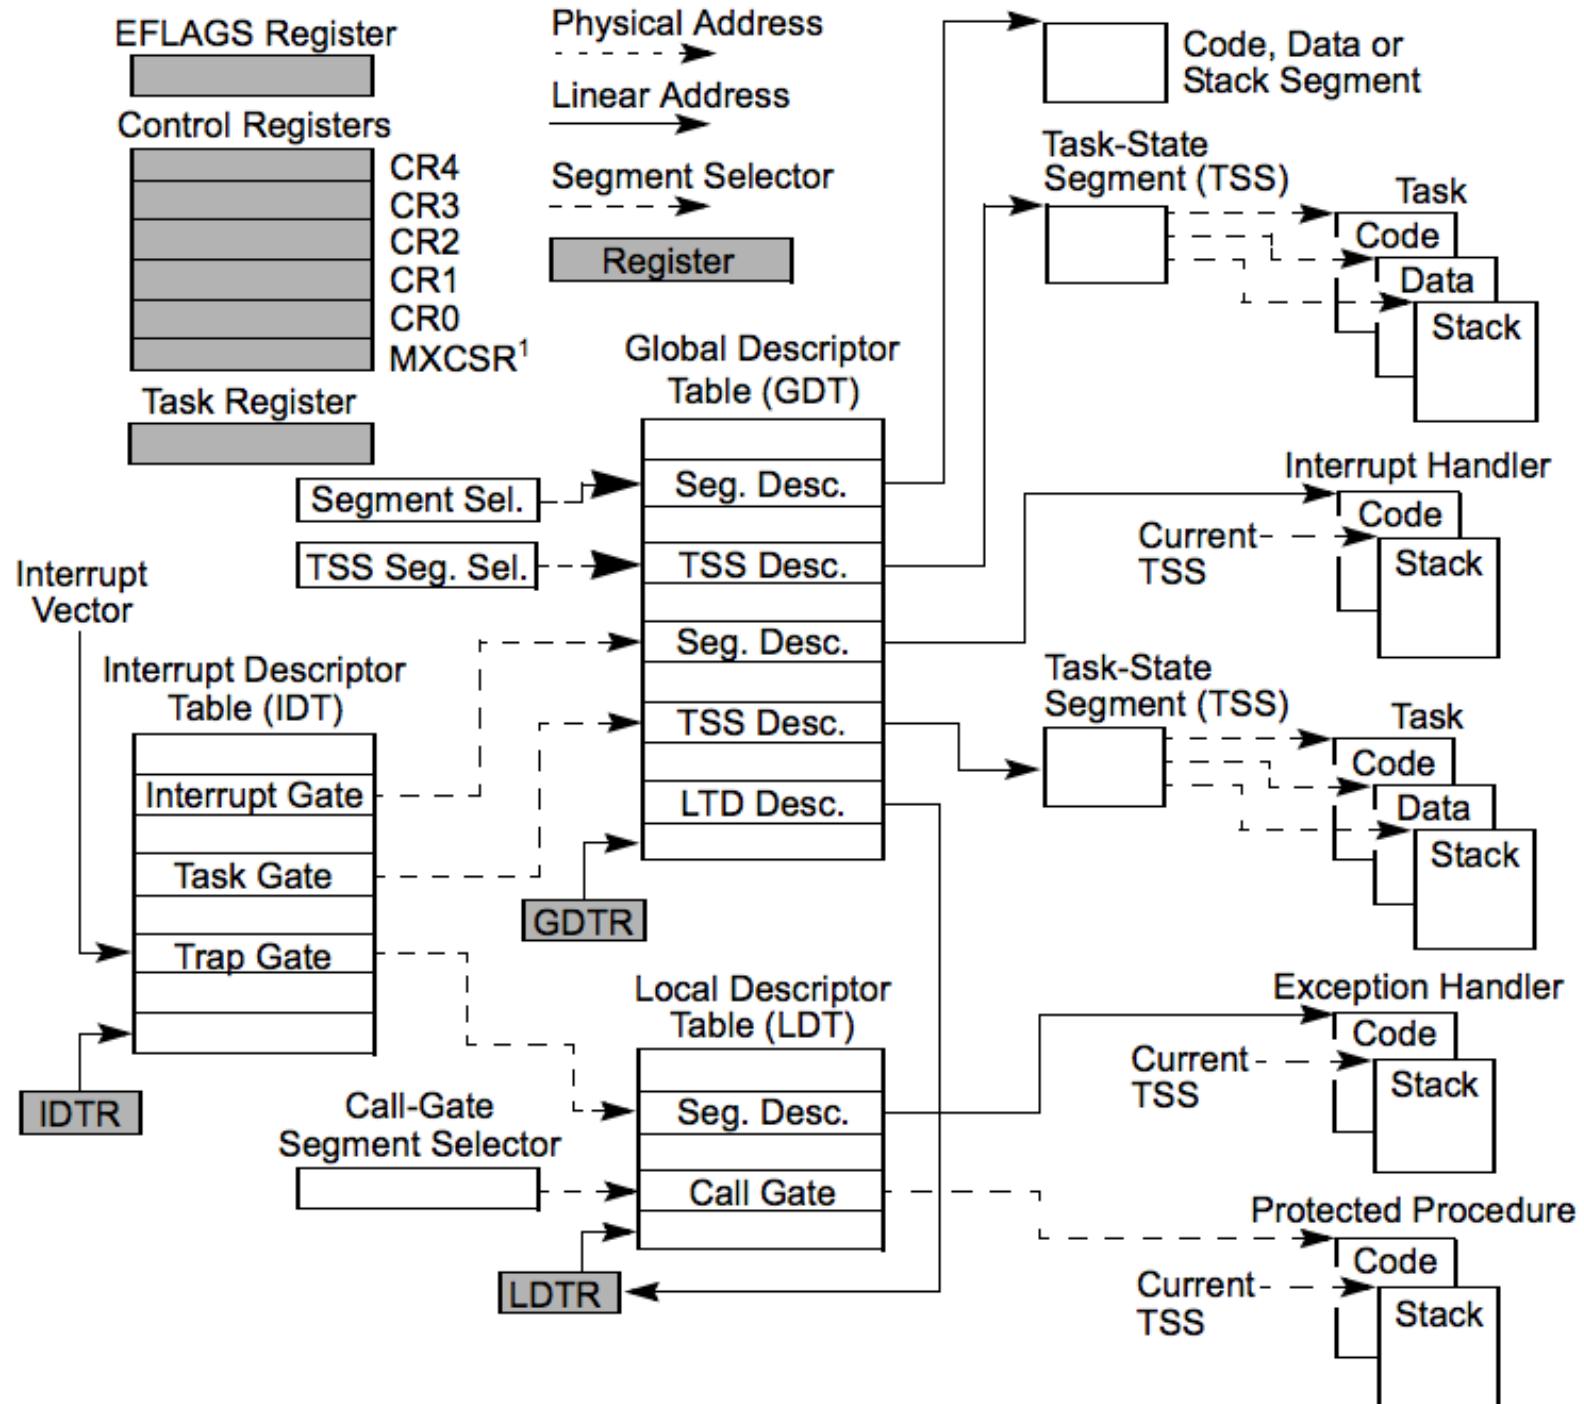
- Vector: index (0-255) into descriptor table (IDT)
 - Special register: idtr points to table (use lidt to load)
- IDT: table of “gate descriptors”
 - Segment selector + offset for handler
 - Descriptor Privilege Level (DPL)
 - Gates (slightly different ways of entering kernel)
 - **Task gate**: includes TSS to transfer to (not used by Linux)
 - **Interrupt gate**: disables further interrupts
 - **Trap gate**: further interrupts still allowed
- Vector assignments
 - Exceptions, NMI are fixed
 - Maskable interrupts can be assigned as needed

Gates

- Interrupt gate
 - E.g., vector 32 in XV6: interrupt
 - Clears IF, so no recursive interrupts
- Trap gate
 - E.g., vector 64 in XV6: system call
 - Interrupt
- Task Gate
 - Includes TSS to transfer to (not used by Linux)

```
17 void
18 tvinit(void)
19 {
20     int i;
21
22     for(i = 0; i < 256; i++)
23         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
24     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
25
26     initlock(&tickslock, "time");
27 }
```

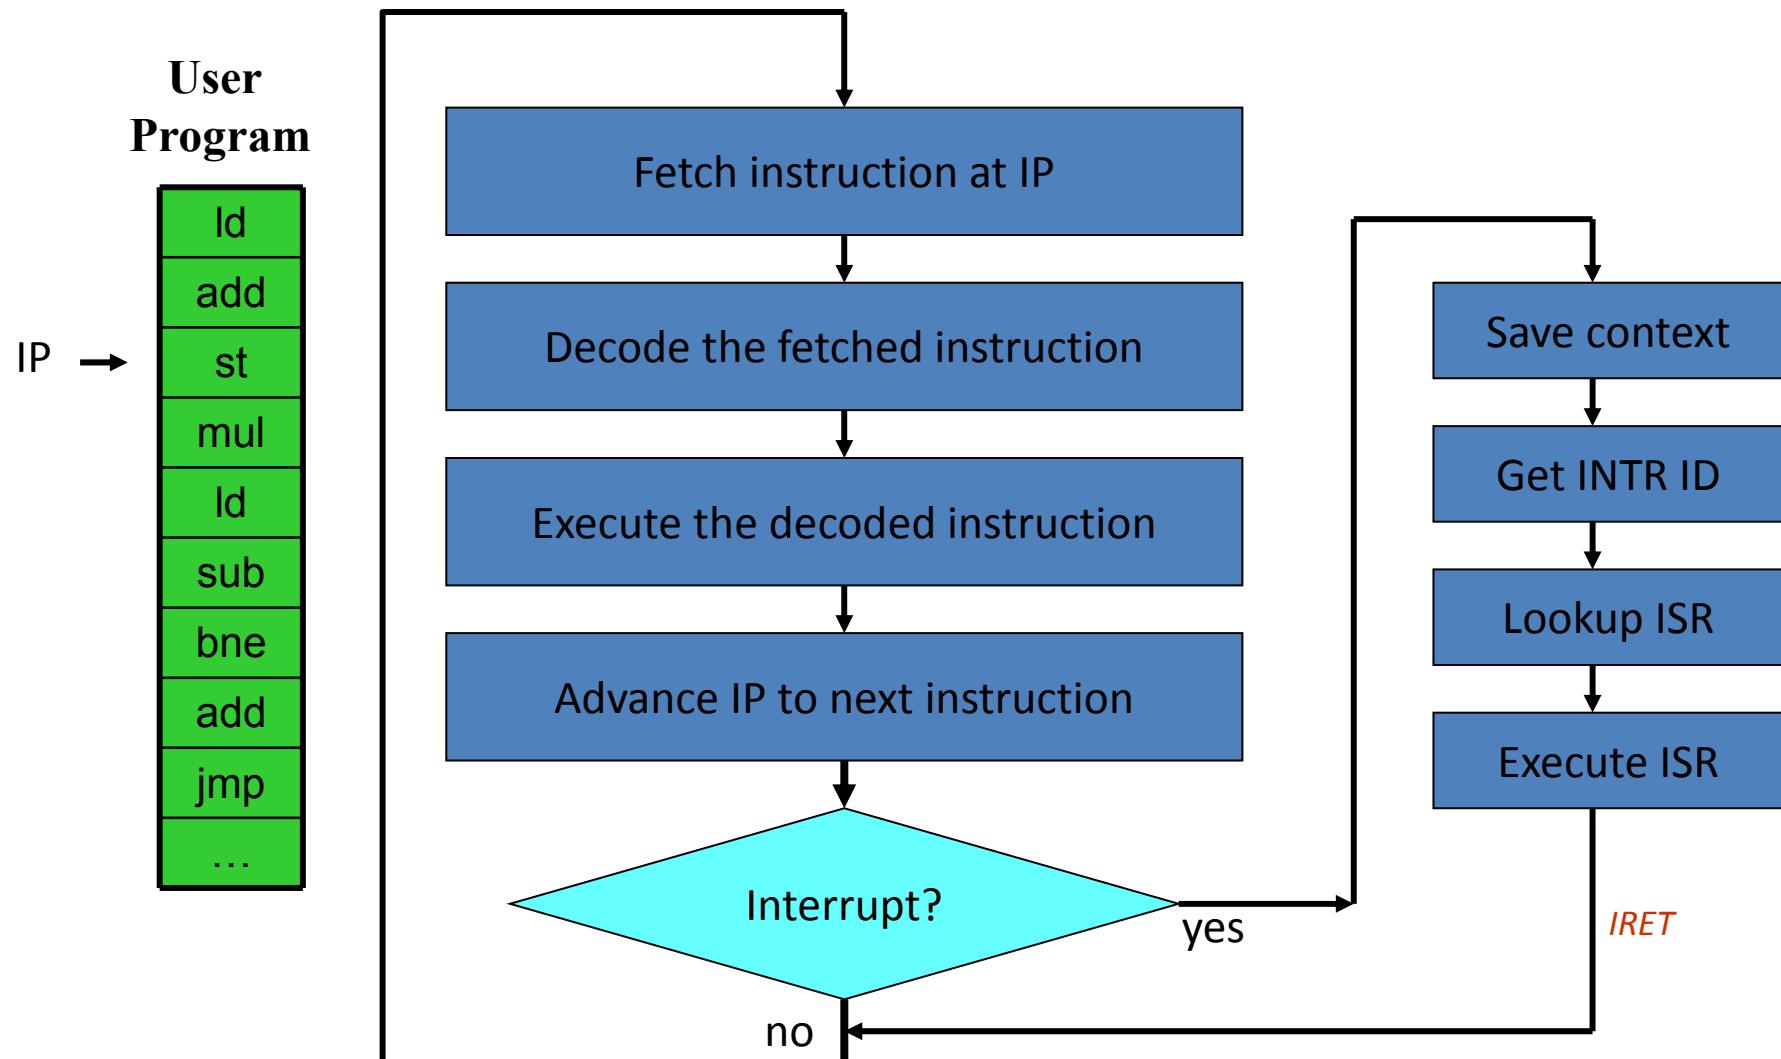
```
213 #define SETGATE(gate, istrap, sel, off, d) \
214 { \
215     (gate).off_15_0 = (uint)(off) & 0xffff; \
216     (gate).cs = (sel); \
217     (gate).args = 0; \
218     (gate).rsv1 = 0; \
219     (gate).type = (istrap) ? STS_TG32 : STS_IG32; \
220     (gate).s = 0; \
221     (gate).dpl = (d); \
222     (gate).p = 1; \
223     (gate).off_31_16 = (uint)(off) >> 16; \
224 }
```



Intel-Reserved ID-Numbers

- Of the 256 possible interrupt ID numbers, Intel reserves the first 32 for ‘exceptions’
- OS’s such as Linux are free to use the remaining 224 available interrupt ID numbers for their own purposes (e.g., for service-requests from external devices, or for other purposes such as system-calls)
- Examples:
 - 0: divide-overflow fault
 - 6: Undefined Opcode
 - 7: Coprocessor Not Available
 - 11: Segment-Not-Present fault
 - 12: Stack fault
 - 13: General Protection Exception
 - 14: Page-Fault Exception

CPU's 'fetch-execute' cycle



```
10 #define IO_TIMER1          0x040           // 8253 Timer #1
11
12 // Frequency of all three count-down timers;
13 // (TIMER_FREQ/freq) is the appropriate count
14 // to generate a frequency of freq Hz.
15
16 #define TIMER_FREQ         1193182
17 #define TIMER_DIV(x)       ((TIMER_FREQ+(x)/2)/(x))
18
19 #define TIMER_MODE         (IO_TIMER1 + 3) // timer mode port
20 #define TIMER_SEL0          0x00      // select counter 0
21 #define TIMER_RATEGEN        0x04      // mode 2, rate generator
22 #define TIMER_16BIT          0x30      // r/w counter 16 bits, LSB first
23
24 void
25 timerinit(void)
26 {
27     // Interrupt 100 times/sec.
28     outb(TIMER_MODE, TIMER_SEL0 | TIMER_RATEGEN | TIMER_16BIT);
29     outb(IO_TIMER1, TIMER_DIV(100) % 256);
30     outb(IO_TIMER1, TIMER_DIV(100) / 256);
31     picenable(IRQ_TIMER);
32 }
```

Timer Interrupt

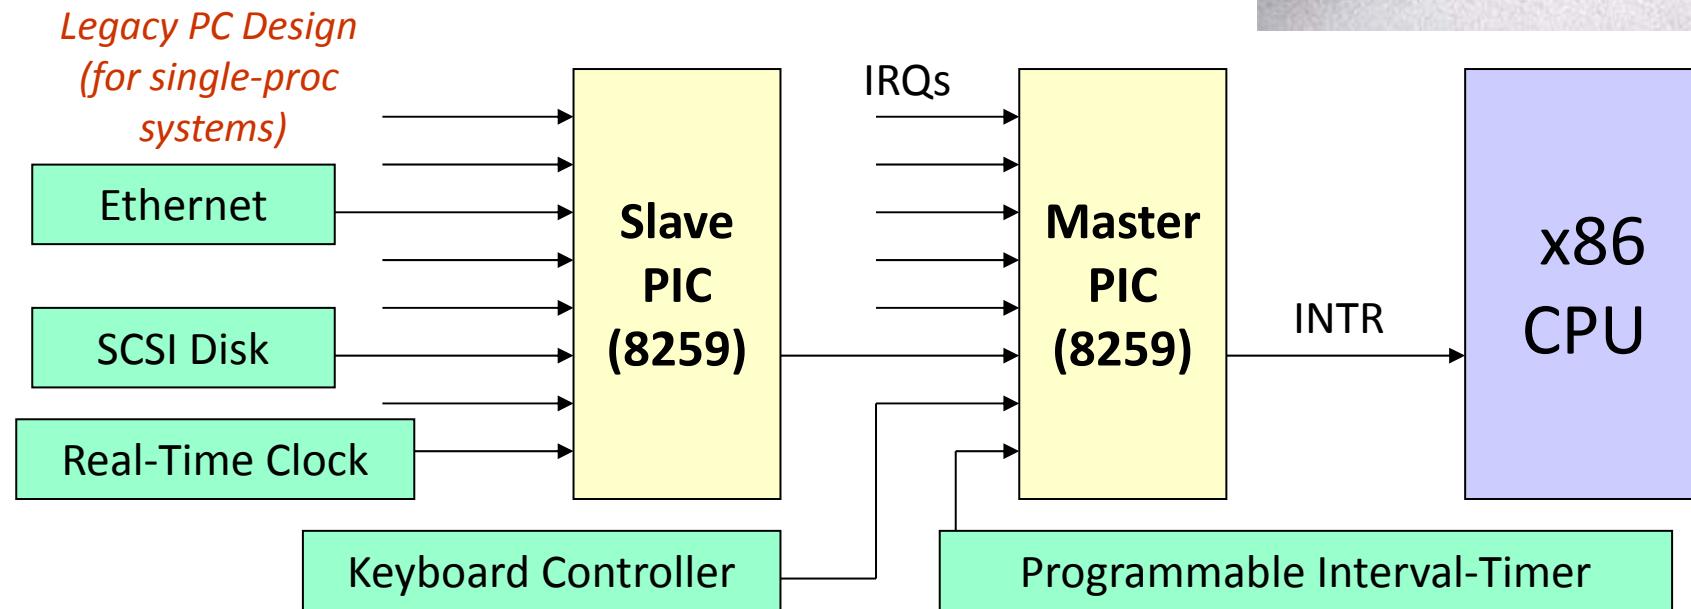
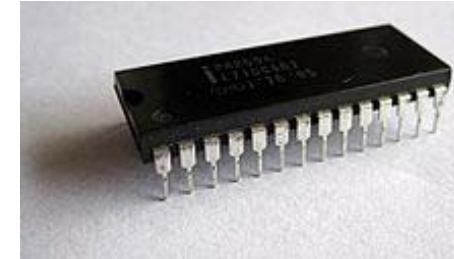
```
49 switch(tf->trapno){  
50 case T_IRQ0 + IRQ_TIMER:  
51     if(cpu->id == 0){  
52         acquire(&tickslock);  
53         ticks++;  
54         wakeup(&ticks);  
55         release(&tickslock);  
56     }  
57     lapiceoi();  
58     break;
```

```
58 int  
59 sys_sleep(void)  
60 {  
61     int n;  
62     uint ticks0;  
63  
64     if(argint(0, &n) < 0)  
65         return -1;  
66     acquire(&tickslock);  
67     ticks0 = ticks;  
68     while(ticks - ticks0 < n){  
69         if(proc->killed){  
70             release(&tickslock);  
71             return -1;  
72         }  
73         sleep(&ticks, &tickslock);  
74     }  
75     release(&tickslock);  
76     return 0;  
77 }
```

Early Boards

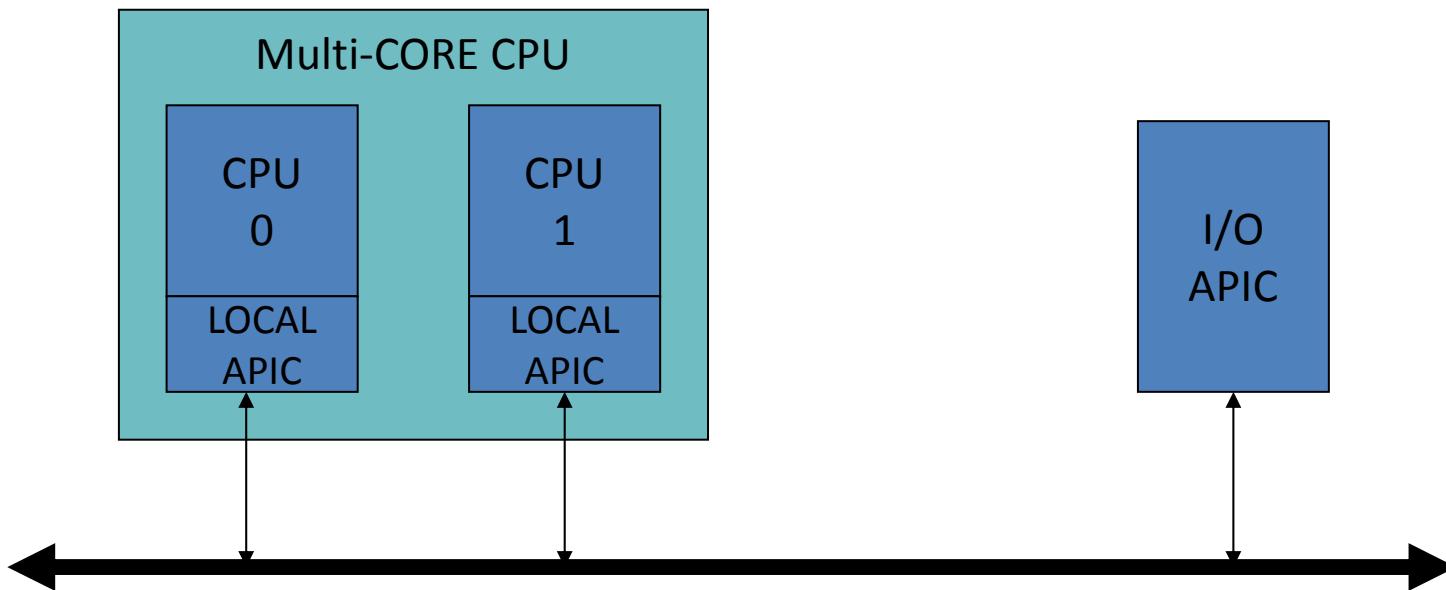
- PIC (Programmable Interrupt Controller)
 - picirq.c
 - 8259A chip
 - Each PIC handles 8 interrupts, and multiplex them on the interrupt pin of the processor
 - PIC can be cascaded
 - Master: 0 to 7
 - Slave: 8 to 15
 - Read timer.c for timer init
- XV6 works correctly on uni-processor

Interrupt Hardware



- I/O devices have (unique or shared) *Interrupt Request Lines* (IRQs)
- IRQs are mapped by special hardware to *interrupt vectors*, and passed to the CPU
- This hardware is called a *Programmable Interrupt Controller* (PIC)

Multiple Logical Processors



Advanced Programmable Interrupt Controller is needed to perform ‘routing’ of I/O requests from peripherals to CPUs

(The legacy PICs are masked when the APICs are enabled)

APIC, IO-APIC, LAPIC

- Advanced PIC (APIC) for SMP systems
 - Used in all modern systems
 - Interrupts “routed” to CPU over system bus
 - IPI: inter-processor interrupt
- Local APIC (LAPIC) versus “frontend” IO-APIC
 - Devices connect to front-end IO-APIC
 - IO-APIC communicates (over bus) with Local APIC
- Interrupt routing
 - Allows broadcast or selective routing of interrupts
 - Ability to distribute interrupt handling load
 - Routes to lowest priority process
 - Special register: Task Priority Register (TPR), i.e., locking
 - Arbitrates (round-robin) if equal priority

Multi-processor PC Boards

- Two parts in XV6
 - IO APIC, for the I/O system
 - ioapic.c
 - Local APIC, for each processor
 - lapic.c
- XV6 is designed for multi-processor
 - Each processor must be programmed to receive interrupts

Assigning IRQs to Devices

- IRQ assignment is hardware-dependent
 - Sometimes it's hardwired, sometimes it's set physically, sometimes it's programmable
 - PCI bus usually assigns IRQs at boot
- Some IRQs are fixed by the architecture
 - IRQ0: Interval timer
 - IRQ2: Cascade pin for 8259A
- Linux device drivers request IRQs when the device is opened
 - Note: especially useful for dynamically-loaded drivers, such as for USB or PCMCIA devices
 - Two devices that aren't used at the same time can share an IRQ, even if the hardware doesn't support simultaneous sharing

Assigning Vectors to IRQs

- Vector: index (0-255) into interrupt descriptor table
- Vectors usually IRQ# + 32
 - Below 32 reserved for non-maskable intr & exceptions
 - Maskable interrupts can be assigned as needed
 - Vector 128 used for syscall
 - Vectors 251-255 used for IPI

```
31 // Initialize the 8259A interrupt controllers.
32 void
33 picinit(void)
34 {
35     // mask all interrupts
36     outb(IO_PIC1+1, 0xFF);
37     outb(IO_PIC2+1, 0xFF);
38
39     // Set up master (8259A-1)
40
41     // ICW1: 0001g0hi
42     //     g: 0 = edge triggering, 1 = level triggering
43     //     h: 0 = cascaded PICs, 1 = master only
44     //     i: 0 = no ICW4, 1 = ICW4 required
45     outb(IO_PIC1, 0x11);
46
47     // ICW2: Vector offset
48     outb(IO_PIC1+1, T_IRQ0);
49
50     // ICW3: (master PIC) bit mask of IR lines connected to slaves
51     //         (slave PIC) 3-bit # of slave's connection to master
52     outb(IO_PIC1+1, 1<<IRQ_SLAVE);
```

End of Interrupt

- Issued to the PIC chips at the end of an IRQ-based interrupt routine
 - If the IRQ came from the Master PIC, it is sufficient to issue this command only to the Master PIC;
 - If the IRQ came from the Slave PIC, it is necessary to issue the command to both PIC chips

```
#define PIC_EOI          0x20      /* End-of-interrupt

void PIC_sendEOI(unsigned char irq)
{
    if(irq >= 8)
        outb(PIC2_COMMAND, PIC_EOI);

    outb(PIC1_COMMAND, PIC_EOI);
}
```

Nested Interrupts

- What if a second interrupt occurs while an interrupt routine is executing?
- Generally a good thing to permit that — is it possible?
- And why is it a good thing?

Maximizing Parallelism

- You want to keep all I/O devices as busy as possible
- In general, an I/O interrupt represents the end of an operation; another request should be issued as soon as possible
- Most devices don't interfere with each others' data structures; there's no reason to block out other devices

Handling Nested Interrupts

- As soon as possible, unmask the global interrupt
- As soon as reasonable, re-enable interrupts from that IRQ
- But that isn't always a great idea, since it could cause re-entry to the same handler
- IRQ-specific mask is not enabled during interrupt-handling

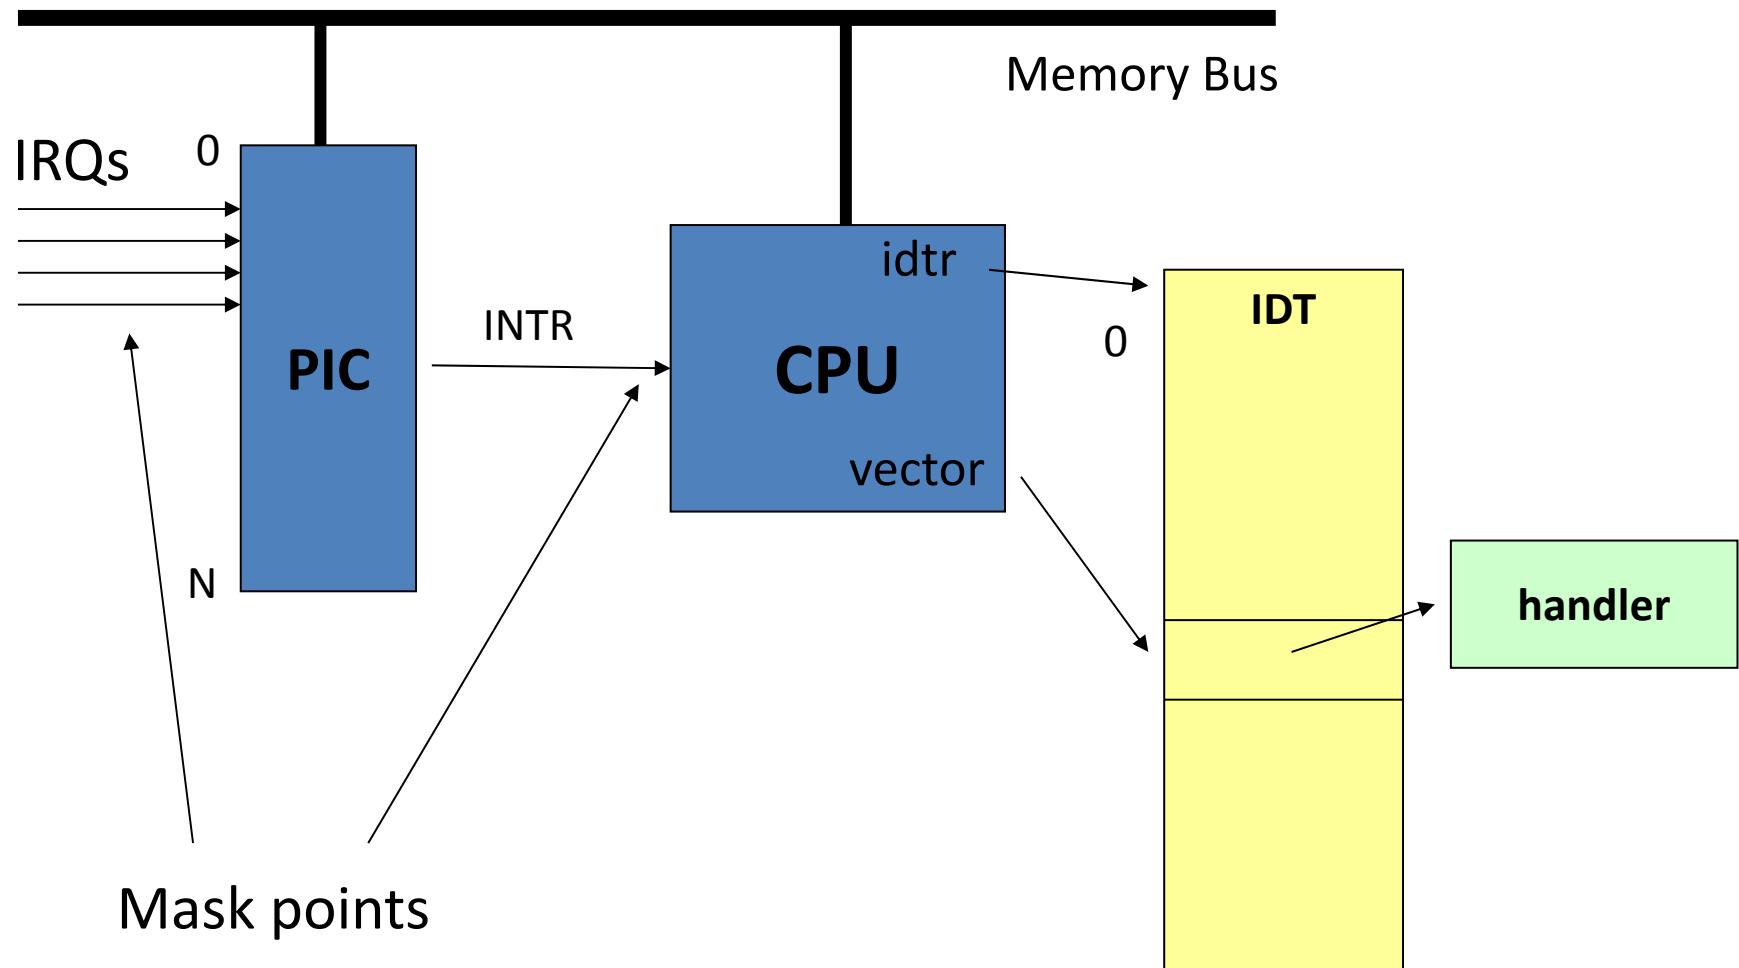
Nested Execution

- Interrupts can be interrupted
 - By different interrupts; handlers need not be reentrant
 - Small portions execute with interrupts disabled
 - Interrupts remain pending until acked by CPU
- Exceptions can be interrupted
 - By interrupts (devices needing service)
- Exceptions can nest two levels deep
 - Exceptions indicate coding error
 - Exception code (kernel code) shouldn't have bugs
 - Page fault is possible (trying to touch user data)

Interrupt Masking

- Two different types: global and per-IRQ
- Global — delays all interrupts
- Selective — individual IRQs can be masked selectively
- Selective masking is usually what's needed — interference most common from two interrupts of the same type
- NMI (Non-Maskable Interrupt)

Putting It All Together



/proc/interrupts

```
$ cat /proc/interrupts
CPU0
 0: 865119901      IO-APIC-edge  timer
 1:        4      IO-APIC-edge  keyboard
 2:        0        XT-PIC       cascade
 8:        1      IO-APIC-edge   rtc
12:       20      IO-APIC-edge PS/2 Mouse
14:  6532494      IO-APIC-edge   ide0
15:       34      IO-APIC-edge   ide1
16:       0      IO-APIC-level  usb-uhci
19:       0      IO-APIC-level  usb-uhci
23:       0      IO-APIC-level ehci-hcd
32:       40      IO-APIC-level   ioc0
33:       40      IO-APIC-level   ioc1
48: 273306628      IO-APIC-level    eth0
NMI:       0
ERR:       0
```

- Columns: IRQ, count, interrupt controller, devices

More in /proc/pci:

```
$ cat /proc/pci
PCI devices found:
Bus 0, device 0, function 0:
    Host bridge: PCI device 8086:2550 (Intel Corp.) (rev 3).
        Prefetchable 32 bit memory at 0xe8000000 [0xebffff].

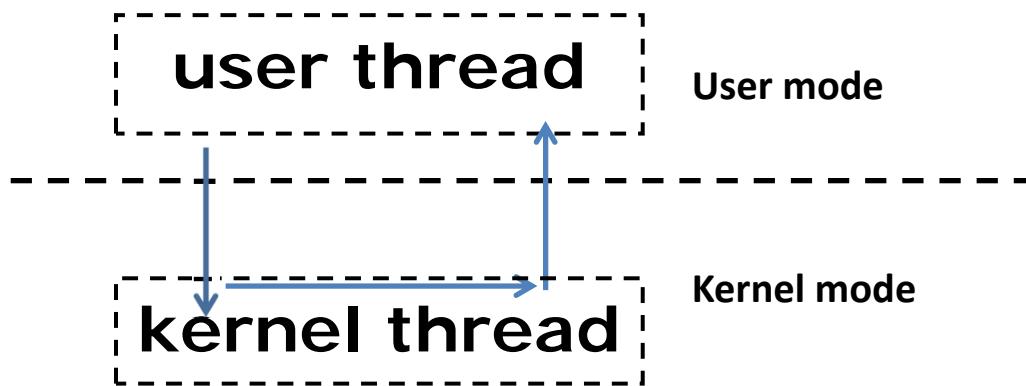
Bus 0, device 29, function 1:
    USB Controller: Intel Corp. 82801DB USB (Hub #2) (rev 2).
        IRQ 19.
        I/O at 0xd400 [0xd41f].
Bus 0, device 31, function 1:
    IDE interface: Intel Corp. 82801DB ICH4 IDE (rev 2).
        IRQ 16.
        I/O at 0xf000 [0xf00f].
        Non-prefetchable 32 bit memory at 0x80000000 [0x800003ff].
Bus 3, device 1, function 0:
    Ethernet controller: Broadcom NetXtreme BCM5703X Gigabit Eth (rev 2).
        IRQ 48.
        Master Capable. Latency=64. Min Gnt=64.
        Non-prefetchable 64 bit memory at 0xf7000000 [0xf700ffff].
```

Summary: Interrupts vs Exceptions

- Varying terminology but for Intel:
 - **Interrupt** (asynchronous, device generated)
 - Maskable: device-generated, associated with IRQs (interrupt request lines); may be temporarily disabled (still pending)
 - Nonmaskable: some critical hardware failures
 - **Exceptions** (synchronous)
 - Processor-detected
 - Faults – correctable (restartable); e.g. page fault
 - Traps – no reexecution needed; e.g. breakpoint
 - Aborts – severe error; process usually terminated (by signal)
 - Programmed exceptions (**software interrupts**)
 - int (system call), int3 (breakpoint)
 - into (overflow), bounds (address check)

Linux System Call

- One system call will cause two context switches between **User** and **Kernel**



Questions

- 1) How can this mode reduce the frequency of mode switching? (2')
- 2) The simplest way to implement this mode is use one kernel thread to process all system calls. Describe how it works on a single core machine. (Hint: how do the user threads and the kernel thread cooperate?) (6')
- 3) What is the drawback of using just one kernel thread? (2')

Questions

- 4) How to mitigate this drawback? (3') (Please provide the details)

- 5) In a multi-core machine, we can dedicate one core to only run the kernel threads and the rest cores to run user threads. Please describe the pros and cons. (3')

Flexible System Call

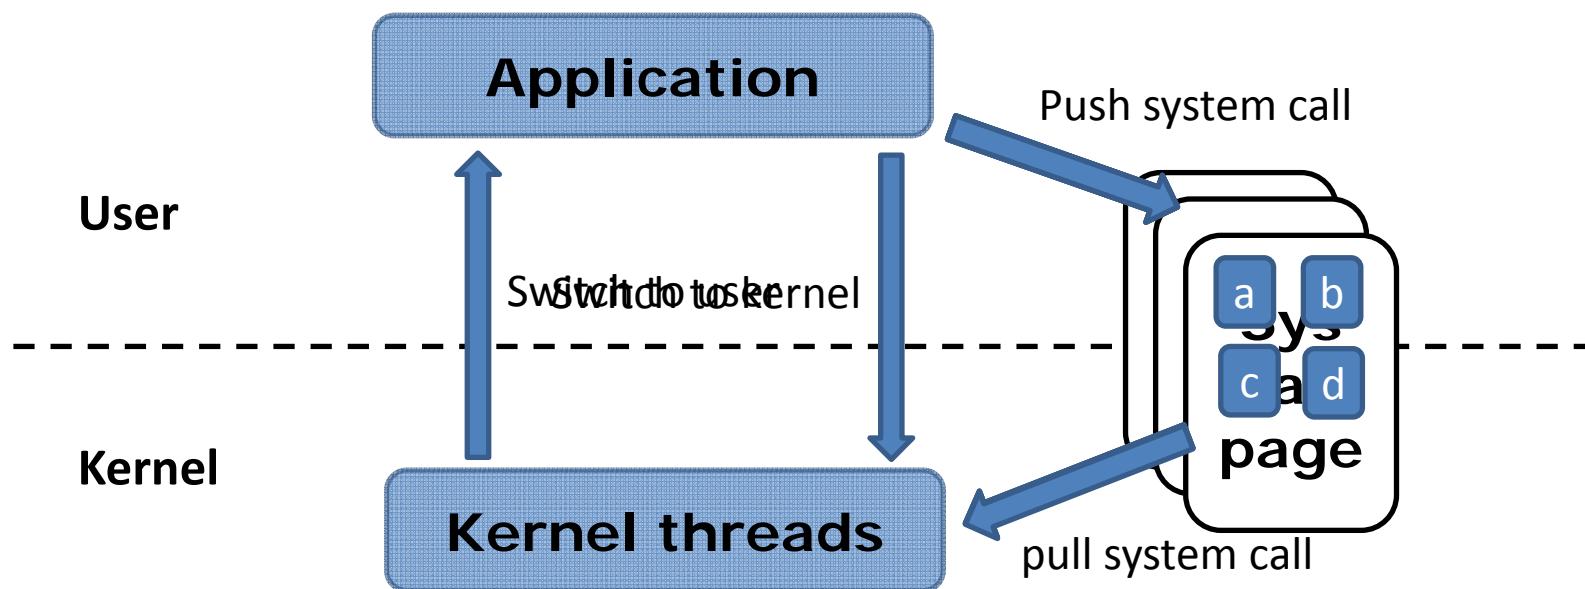
- Problem 2-2
 - New system call mechanism – **Flexible System Call**
 - no direct mapping between user threads and kernel threads
 - *system call page*
 - User threads can push the system call requests into the system call page
 - kernel threads will poll the system call requests out the system call page

Flexible System Call

- Basic Idea
 - “Flexible System Call Scheduling with Exception-Less System Calls”
 - University of Toronto
 - OSDI 2010

Flexible System Call

- Basic Idea

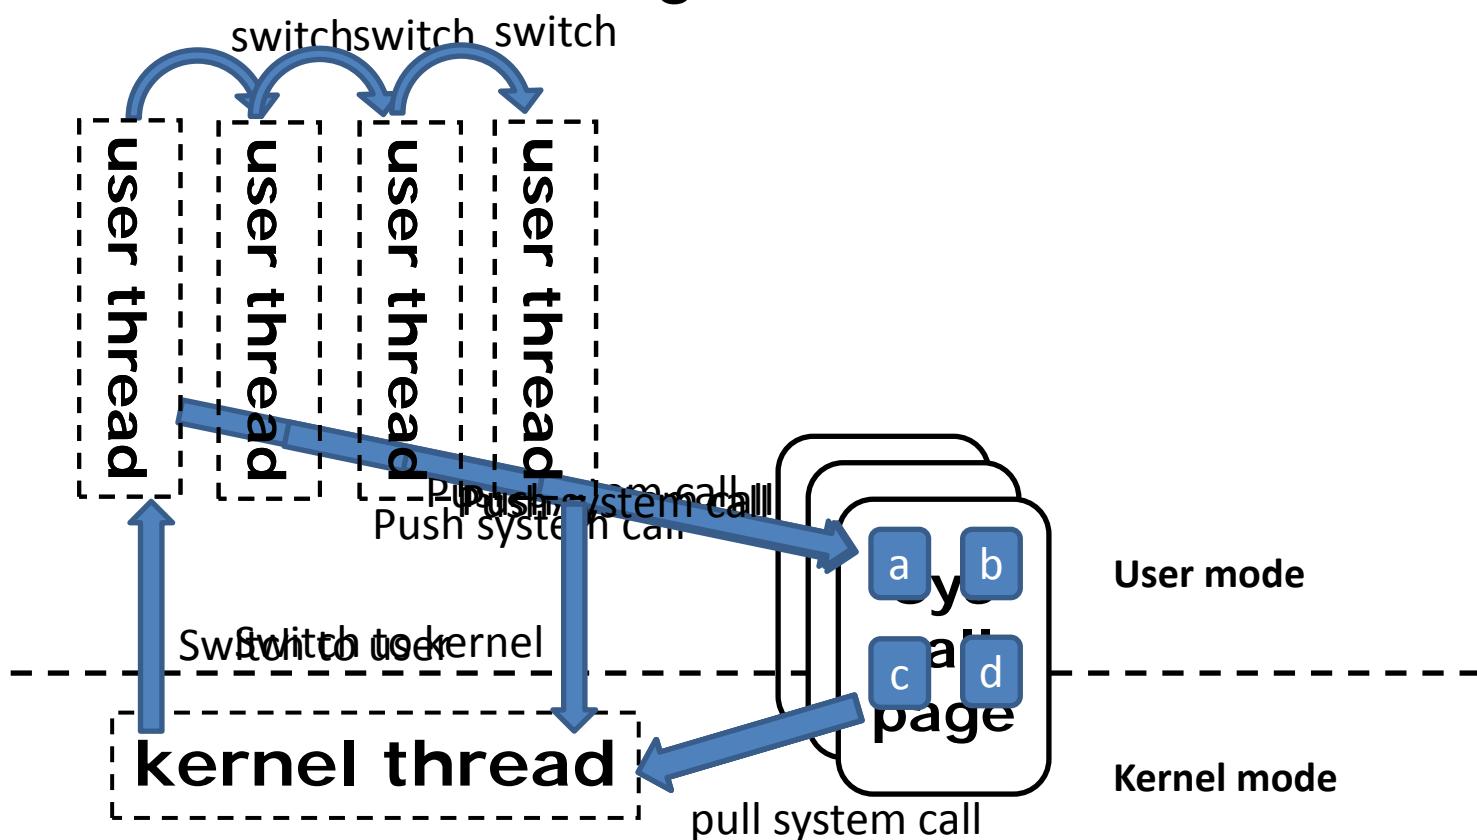


Problem 2-2.1

- How can this mode reduce the frequency of mode switching?
 - Batching system calls using system call page
 - During one switch-around
 - Kernel threads can process multiple system call requests together

Problem 2-2.2

- How it works in a single core machine?



Problem 2-2.3

- What is the drawback of using just one kernel thread?
 - Blocked system calls (e.g. read from disk file) can block the kernel thread
 - It will block the whole machine (not the process)

Problem 2-2.4

- How to mitigate this drawback?
 - We can use **multiple kernel threads**.
 - When one kernel thread is blocked schedule another to continue processing the left system call requests

Problem 2-2.5

- In a multi-core machine, we can dedicate one core to only run the kernel threads and the rest cores to run user threads. Please describe the pros and cons

Problem 2-2.5-Sol

- Pros:
 - No need to mode switching
 - Most of kernel data/code will not be shared across cores
- Cons:
 - Remote data exchanges between user/kernel threads are expensive (cross cores data access)
 - Kernel thread core maybe overloaded
 - Kernel thread core maybe idle

Interrupt, Trap, SYS CALL

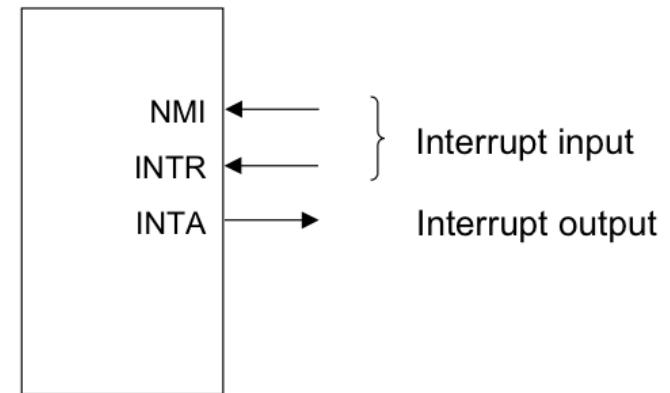
Haibo Chen, Yubin Xia
IPADS, SJTU

ACKs: Some slides are adapted from the textbook's original slides
and frans's os course notes

Interrupt Sources

1. Hardware interrupt: external

- *Nonmaskable interrupt (NMI)* input pin
- *Interrupt (INTR)* input pin



1. Software interrupt: execution of the Interrupt instruction, INT

2. Error condition: If some error condition occur by the execution of an instruction. E.g.,

- *divide-by-zero interrupt*: If you attempt to divide an operand by zero, the 8086 will automatically interrupt the currently executing program

Terms

- Vector, Interrupt vector, Trap number
- IRQ: Interrupt Request
- Soft IRQ: A mechanism to implement bottom half
- Interrupt, trap, fault, exception
- Software interrupt / system call
- IDT: Interrupt Descriptor Table
- ISP: Interrupt Service Procedure

Requirements

- Must save the processor's registers for future transparent resume
- Must be set up for execution in the kernel
- Must choose a place for the kernel to start executing
- Must be able to retrieve information about the event, e.g., system call arguments
- Must all be done securely
- Must maintain isolation of user processes and the kernel

Interrupts

- Interrupts are similar to system calls, except devices generate them at any time
- There is hardware on the motherboard to signal the CPU when a device needs attention (e.g., the user has typed a character on the keyboard)
- We must program the device to generate an interrupt, and arrange that a CPU receives the interrupt.

Types of Interrupts

- Asynchronous
 - From external source, such as I/O device
 - Not related to instruction being executed
- Synchronous (also called *exceptions*)
 - *Processor-detected* exceptions:
 - *Faults* — correctable; offending instruction is *retried*
 - *Traps* — often for debugging; instruction is *not retried*
 - *Aborts* — major error (hardware failure)
 - *Programmed* exceptions:
 - Requests for kernel intervention (software interrupt/syscalls)

System Call

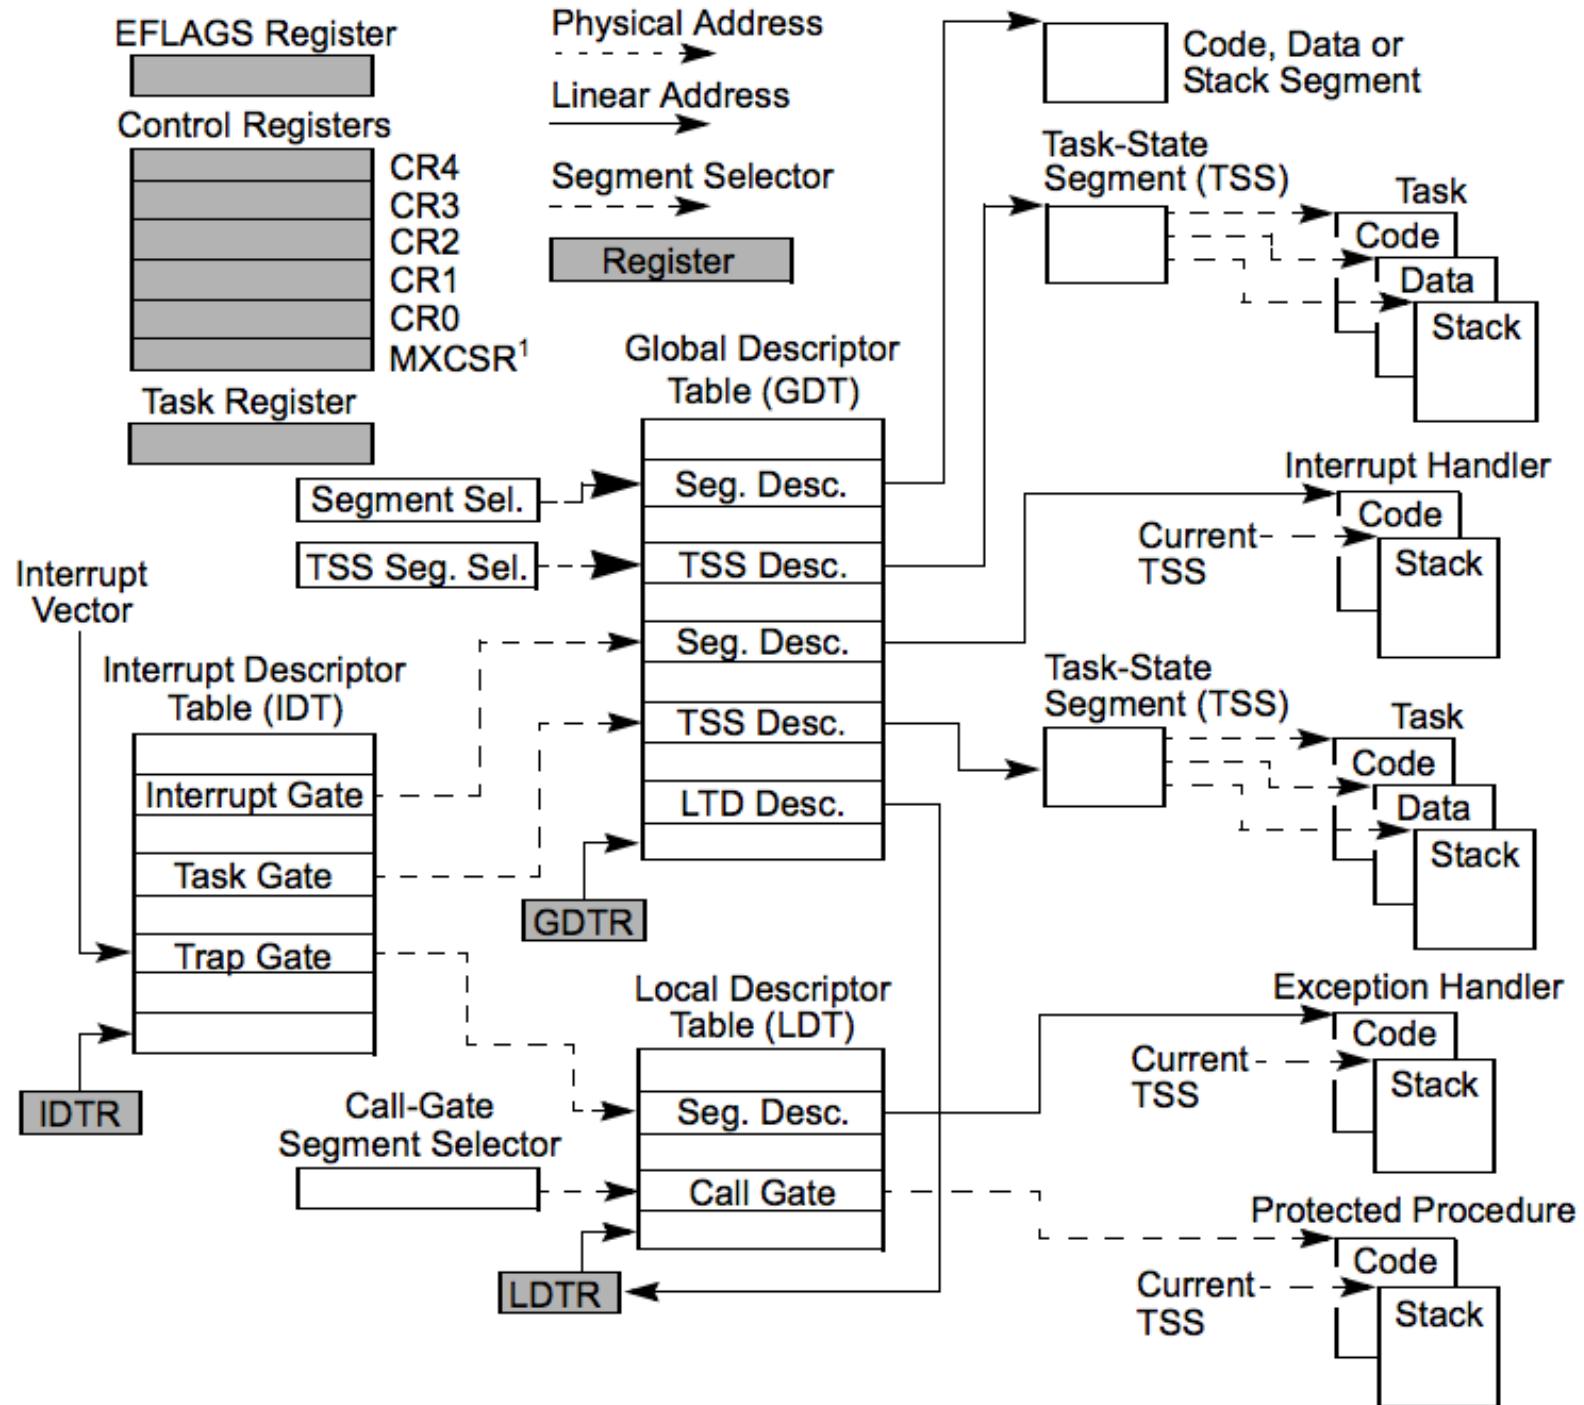
1. Fetch the n 'th descriptor from the IDT, where n is the argument of int.
2. Check that CPL in %cs is \leq DPL, where DPL is the privilege level in the descriptor.
3. Save %esp and %ss in a CPU-internal registers, but only if the target segment selector's PL $<$ CPL.
4. Load %ss and %esp from a task segment descriptor.
5. Push %ss, %esp, %eflags, %cs, %eip,
6. Clear some bits of %eflags
7. Set %cs and %eip to the values in the descriptor

Vectors, IDT

- Vector: index (0-255) into descriptor table (IDT)
 - Special register: idtr points to table (use lidt to load)
- IDT: table of “gate descriptors”
 - Segment selector + offset for handler
 - Descriptor Privilege Level (DPL)
 - Gates (slightly different ways of entering kernel)
 - **Task gate**: includes TSS to transfer to (not used by Linux)
 - **Interrupt gate**: disables further interrupts
 - **Trap gate**: further interrupts still allowed
- Vector assignments
 - Exceptions, NMI are fixed
 - Maskable interrupts can be assigned as needed

```
17 void
18 tvinit(void)
19 {
20     int i;
21
22     for(i = 0; i < 256; i++)
23         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
24     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
25
26     initlock(&tickslock, "time");
27 }
```

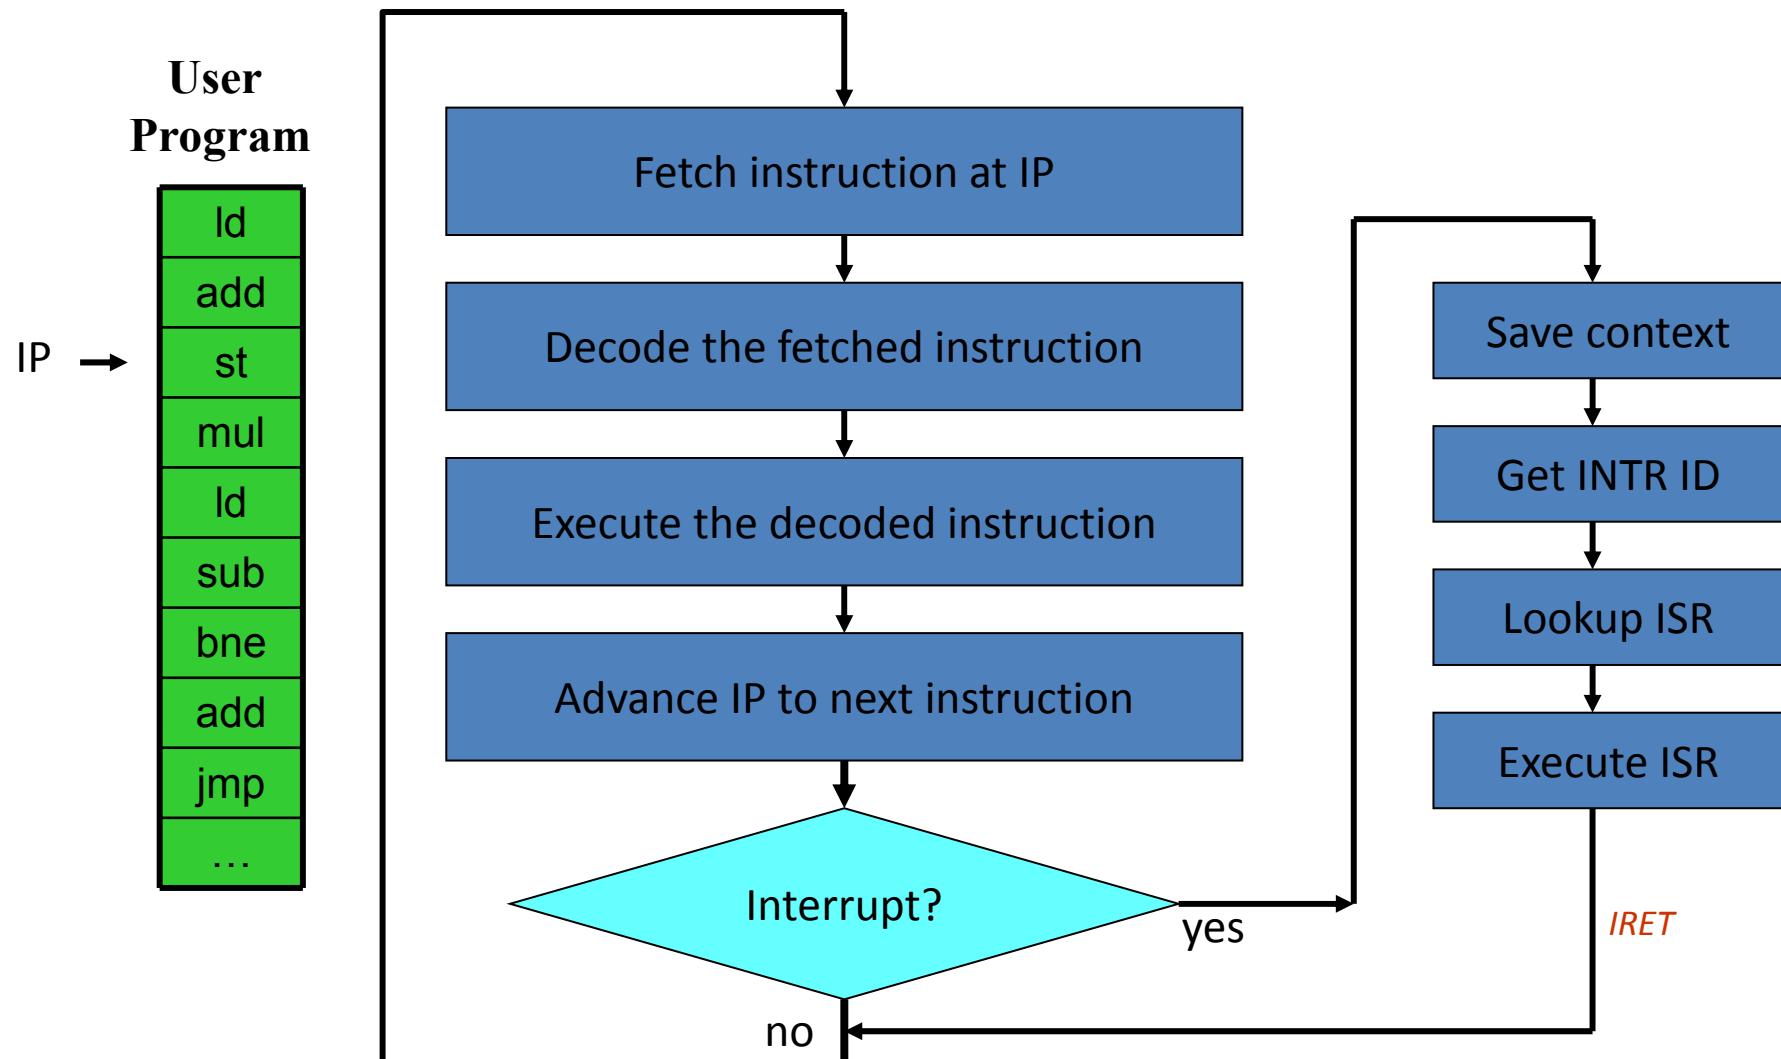
```
213 #define SETGATE(gate, istrap, sel, off, d) \
214 { \
215     (gate).off_15_0 = (uint)(off) & 0xffff; \
216     (gate).cs = (sel); \
217     (gate).args = 0; \
218     (gate).rsv1 = 0; \
219     (gate).type = (istrap) ? STS_TG32 : STS_IG32; \
220     (gate).s = 0; \
221     (gate).dpl = (d); \
222     (gate).p = 1; \
223     (gate).off_31_16 = (uint)(off) >> 16; \
224 }
```



Intel-Reserved ID-Numbers

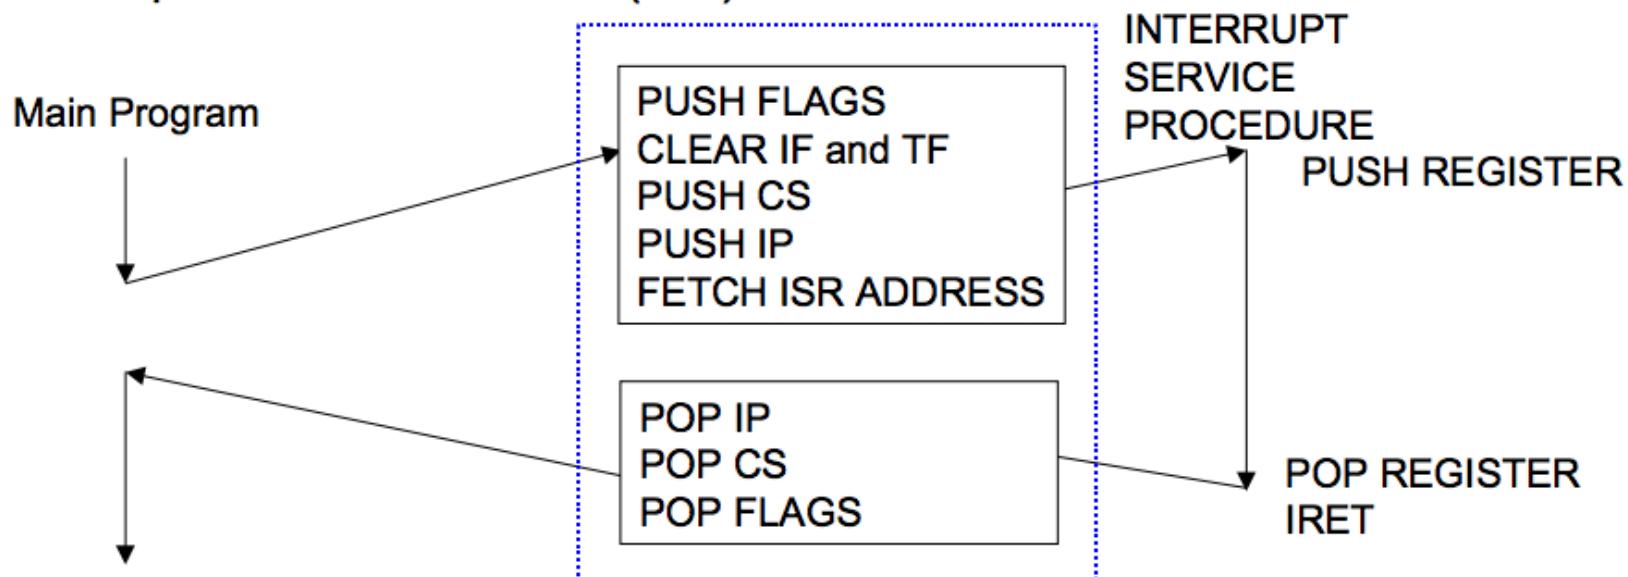
- Of the 256 possible interrupt ID numbers, Intel reserves the first 32 for ‘exceptions’
- OS’s such as Linux are free to use the remaining 224 available interrupt ID numbers for their own purposes (e.g., for service-requests from external devices, or for other purposes such as system-calls)
- Examples:
 - 0: divide-overflow fault
 - 6: Undefined Opcode
 - 7: Coprocessor Not Available
 - 11: Segment-Not-Present fault
 - 12: Stack fault
 - 13: General Protection Exception
 - 14: Page-Fault Exception

CPU's 'fetch-execute' cycle



Interrupt Response Sequence

1. PUSH Flag Register
2. Clear IF and TF
3. PUSH CS
4. PUSH IP
5. Fetch Interrupt vector contents and place into both IP and CS to start the Interrupt Service Procedure (ISP)



```
10 #define IO_TIMER1          0x040           // 8253 Timer #1
11
12 // Frequency of all three count-down timers;
13 // (TIMER_FREQ/freq) is the appropriate count
14 // to generate a frequency of freq Hz.
15
16 #define TIMER_FREQ         1193182
17 #define TIMER_DIV(x)       ((TIMER_FREQ+(x)/2)/(x))
18
19 #define TIMER_MODE         (IO_TIMER1 + 3) // timer mode port
20 #define TIMER_SEL0          0x00      // select counter 0
21 #define TIMER_RATEGEN        0x04      // mode 2, rate generator
22 #define TIMER_16BIT          0x30      // r/w counter 16 bits, LSB first
23
24 void
25 timerinit(void)
26 {
27     // Interrupt 100 times/sec.
28     outb(TIMER_MODE, TIMER_SEL0 | TIMER_RATEGEN | TIMER_16BIT);
29     outb(IO_TIMER1, TIMER_DIV(100) % 256);
30     outb(IO_TIMER1, TIMER_DIV(100) / 256);
31     picenable(IRQ_TIMER);
32 }
```

Timer Interrupt

```
49 switch(tf->trapno){  
50 case T_IRQ0 + IRQ_TIMER:  
51     if(cpu->id == 0){  
52         acquire(&tickslock);  
53         ticks++;  
54         wakeup(&ticks);  
55         release(&tickslock);  
56     }  
57     lapiceoi();  
58     break;
```

```
58 int  
59 sys_sleep(void)  
60 {  
61     int n;  
62     uint ticks0;  
63  
64     if(argint(0, &n) < 0)  
65         return -1;  
66     acquire(&tickslock);  
67     ticks0 = ticks;  
68     while(ticks - ticks0 < n){  
69         if(proc->killed){  
70             release(&tickslock);  
71             return -1;  
72         }  
73         sleep(&ticks, &tickslock);  
74     }  
75     release(&tickslock);  
76     return 0;  
77 }
```

Sleep & Wakeup

```
378 // Wake up all processes sleeping on chan.  
379 // The ptable lock must be held.  
380 static void  
381 wakeup1(void *chan)  
382 {  
383     struct proc *p;  
384  
385     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
386         if(p->state == SLEEPING && p->chan == chan)  
387             p->state = RUNNABLE;  
388 }  
389  
390 // Wake up all processes sleeping on chan.  
391 void  
392 wakeup(void *chan)  
393 {  
394     acquire(&ptable.lock);  
395     wakeup1(chan);  
396     release(&ptable.lock);  
397 }
```

```
340 // Atomically release lock and sleep on chan.  
341 // Reacquires lock when awakened.  
342 void  
343 sleep(void *chan, struct spinlock *lk)  
344 {  
345     if(proc == 0)  
346         panic("sleep");  
347  
348     if(lk == 0)  
349         panic("sleep without lk");  
350  
351     // Must acquire ptable.lock in order to  
352     // change p->state and then call sched.  
353     // Once we hold ptable.lock, we can be  
354     // guaranteed that we won't miss any wakeup  
355     // (wakeup runs with ptable.lock locked),  
356     // so it's okay to release lk.  
357     if(lk != &ptable.lock){ //DOC: sleeplock0  
358         acquire(&ptable.lock); //DOC: sleeplock1  
359         release(lk);  
360     }  
361  
362     // Go to sleep.  
363     proc->chan = chan;  
364     proc->state = SLEEPING;  
365     sched();  
366  
367     // Tidy up.  
368     proc->chan = 0;  
369  
370     // Reacquire original lock.  
371     if(lk != &ptable.lock){ //DOC: sleeplock2  
372         release(&ptable.lock);  
373         acquire(lk);  
374     }  
375 }
```

Early Boards

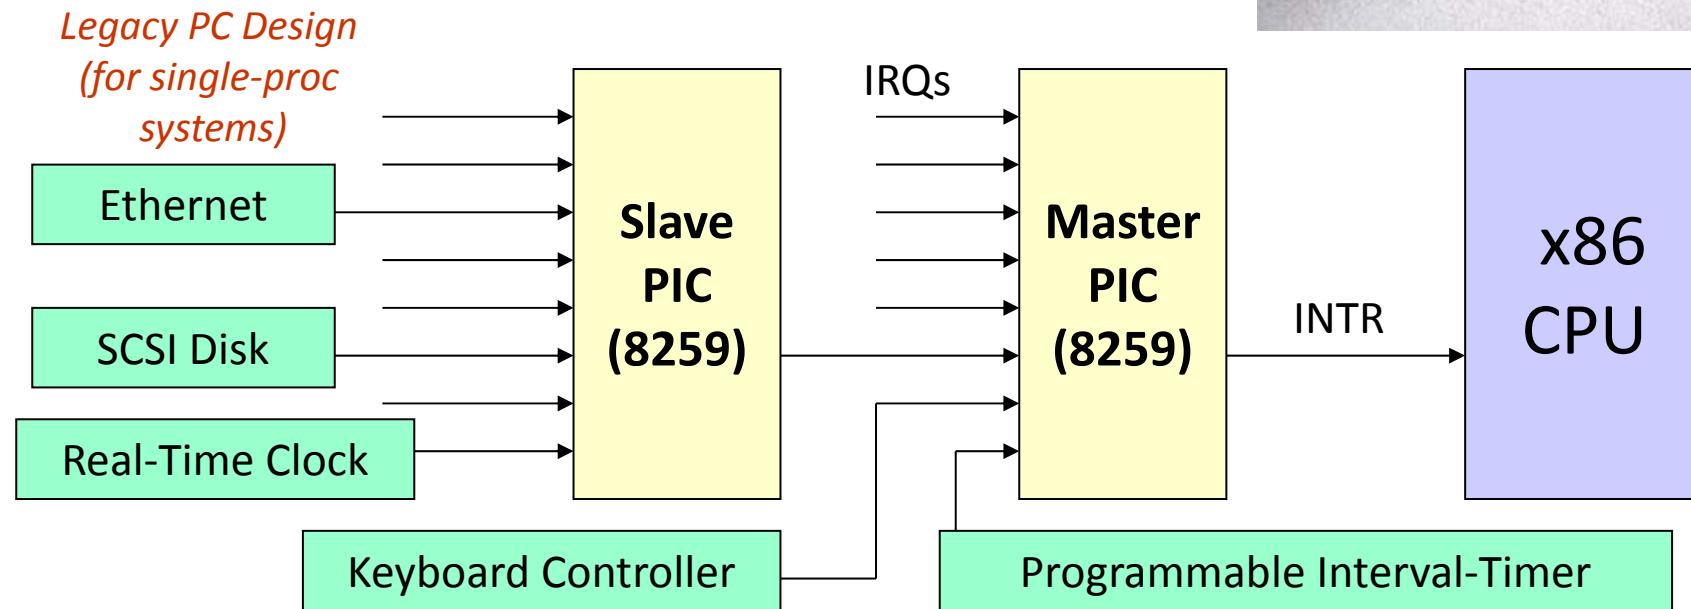
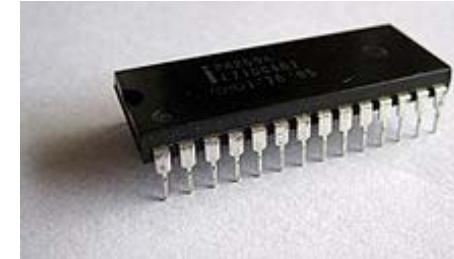
- PIC (Programmable Interrupt Controller)
 - picirq.c
 - 8259A chip
 - Each PIC handles 8 interrupts, and multiplex them on the interrupt pin of the processor
 - PIC can be cascaded
 - Master: 0 to 7
 - Slave: 8 to 15
 - Read timer.c for timer init

8259A Programmable Interrupt Controller (PIC)

- 8259A is a 28-pin integrated circuit which was designed specifically for the 8088/8086 microprocessors

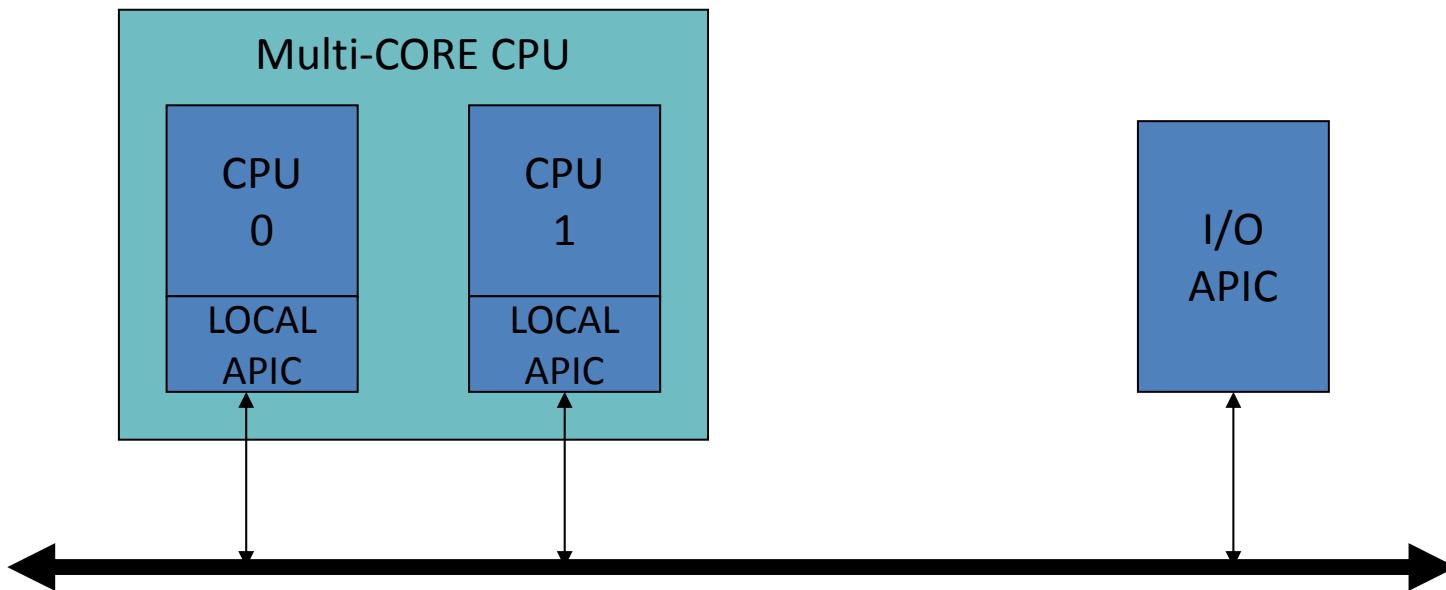
11	D0	18	IR0
10	D1	19	IR1
9	D2	20	IR2
8	D3	21	IR3
7	D4	22	IR4
6	D5	23	IR5
5	D6	24	IR6
4	D7	25	IR7
27	A0		
1	CS		
3	RD		
2	WR		
16	SP/EN	12	CAS0
17	INT	13	CAS1
26	INTA	15	CAS2

Interrupt Hardware



- I/O devices have (unique or shared) *Interrupt Request Lines* (IRQs)
- IRQs are mapped by special hardware to *interrupt vectors*, and passed to the CPU
- This hardware is called a *Programmable Interrupt Controller* (PIC)

Multiple Logical Processors



Advanced Programmable Interrupt Controller is needed to perform ‘routing’ of I/O requests from peripherals to CPUs

(The legacy PICs are masked when the APICs are enabled)

APIC, IO-APIC, LAPIC

- Advanced PIC (APIC) for SMP systems
 - Used in all modern systems
 - Interrupts “routed” to CPU over system bus
 - IPI: inter-processor interrupt
- Local APIC (LAPIC) versus “frontend” IO-APIC
 - Devices connect to front-end IO-APIC
 - IO-APIC communicates (over bus) with Local APIC
- Interrupt routing
 - Allows broadcast or selective routing of interrupts
 - Ability to distribute interrupt handling load
 - Routes to lowest priority process
 - Special register: Task Priority Register (TPR), i.e., locking
 - Arbitrates (round-robin) if equal priority

Interrupt Priority

- When different type of interrupt (ie. software, NMI, INTR or exceptions) occur at the same time, the one with the highest priority is handled first. Intel microprocessors use the following order of priority:

Interrupt Type	Priority
divide error interrupt, INT n, INTO	highest
NMI	
INTR	
TRAP flag (single step)	lowest

Multi-processor PC Boards

- Two parts in XV6
 - IO APIC, for the I/O system
 - ioapic.c
 - Local APIC, for each processor
 - lapic.c
- XV6 is designed for multi-processor
 - Each processor must be programmed to receive interrupts

Assigning IRQs to Devices

- IRQ assignment is hardware-dependent
 - Sometimes it's hardwired, sometimes it's set physically, sometimes it's programmable
 - PCI bus usually assigns IRQs at boot
- Some IRQs are fixed by the architecture
 - IRQ0: Interval timer
 - IRQ2: Cascade pin for 8259A
- Linux device drivers request IRQs when the device is opened
 - Note: especially useful for dynamically-loaded drivers, such as for USB or PCMCIA devices
 - Two devices that aren't used at the same time can share an IRQ, even if the hardware doesn't support simultaneous sharing

Assigning Vectors to IRQs

- Vector: index (0-255) into interrupt descriptor table
- Vectors usually IRQ# + 32
 - Below 32 reserved for non-maskable intr & exceptions
 - Maskable interrupts can be assigned as needed
 - Vector 128 used for syscall
 - Vectors 251-255 used for IPI

End of Interrupt

- Issued to the PIC chips at the end of an IRQ-based interrupt routine
 - If the IRQ came from the Master PIC, it is sufficient to issue this command only to the Master PIC;
 - If the IRQ came from the Slave PIC, it is necessary to issue the command to both PIC chips

```
#define PIC_EOI          0x20      /* End-of-interrupt

void PIC_sendEOI(unsigned char irq)
{
    if(irq >= 8)
        outb(PIC2_COMMAND, PIC_EOI);

    outb(PIC1_COMMAND, PIC_EOI);
}
```

Nested Interrupts

- What if a second interrupt occurs while an interrupt routine is executing?
- Generally a good thing to permit that — is it possible?
- And why is it a good thing?

Maximizing Parallelism

- You want to keep all I/O devices as busy as possible
- In general, an I/O interrupt represents the end of an operation; another request should be issued as soon as possible
- Most devices don't interfere with each others' data structures; there's no reason to block out other devices

Handling Nested Interrupts

- As soon as possible, unmask the global interrupt
- As soon as reasonable, re-enable interrupts from that IRQ
- But that isn't always a great idea, since it could cause re-entry to the same handler
- IRQ-specific mask is not enabled during interrupt-handling

Nested Execution

- Interrupts can be interrupted
 - By different interrupts; handlers need not be reentrant
 - Small portions execute with interrupts disabled
 - Interrupts remain pending until acked by CPU
- Exceptions can be interrupted
 - By interrupts (devices needing service)
- Exceptions can nest two levels deep
 - Exceptions indicate coding error
 - Exception code (kernel code) shouldn't have bugs
 - Page fault is possible (trying to touch user data)

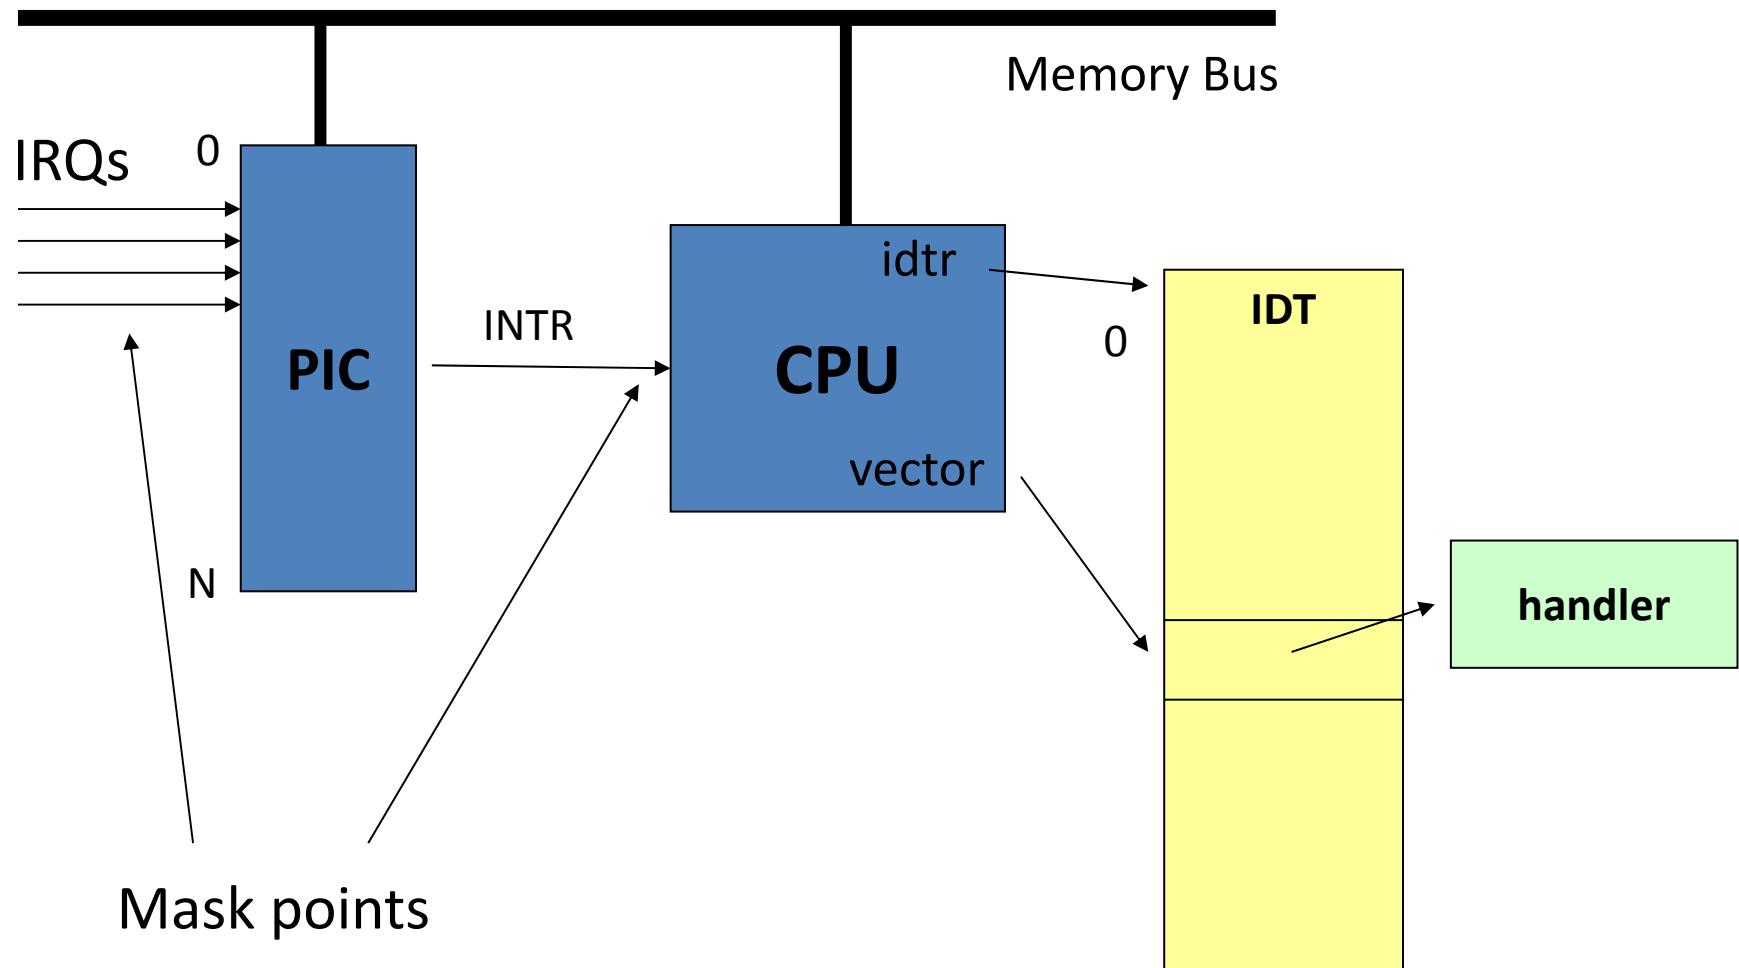
Interrupt Masking

- Two different types: global and per-IRQ
 - Global — delays all interrupts
 - Selective — individual IRQs can be masked selectively
 - Selective masking is usually what's needed — interference most common from two interrupts of the same type
- NMI (Non-Maskable Interrupt)

Triple Fault

- Things never to do in an OS #1: Swap out the page swapping code (triple-fault here we come) —Kemp
- When a fault occurs, the CPU invokes an exception handler.
- If a fault occurs while trying to invoke the exception handler, that's called a double fault, which the CPU tries to handle with yet another exception handler.
- If that invocation results in a fault too, the system reboots with a triple fault.

Putting It All Together



/proc/interrupts

```
$ cat /proc/interrupts
CPU0
 0: 865119901      IO-APIC-edge  timer
 1:        4      IO-APIC-edge  keyboard
 2:        0        XT-PIC       cascade
 8:        1      IO-APIC-edge   rtc
12:       20      IO-APIC-edge PS/2 Mouse
14:  6532494      IO-APIC-edge   ide0
15:       34      IO-APIC-edge   ide1
16:       0      IO-APIC-level  usb-uhci
19:       0      IO-APIC-level  usb-uhci
23:       0      IO-APIC-level ehci-hcd
32:       40      IO-APIC-level   ioc0
33:       40      IO-APIC-level   ioc1
48: 273306628      IO-APIC-level    eth0
NMI:       0
ERR:       0
```

- Columns: IRQ, count, interrupt controller, devices

More in /proc/pci:

```
$ cat /proc/pci
PCI devices found:
Bus 0, device 0, function 0:
    Host bridge: PCI device 8086:2550 (Intel Corp.) (rev 3).
        Prefetchable 32 bit memory at 0xe8000000 [0xebffff].

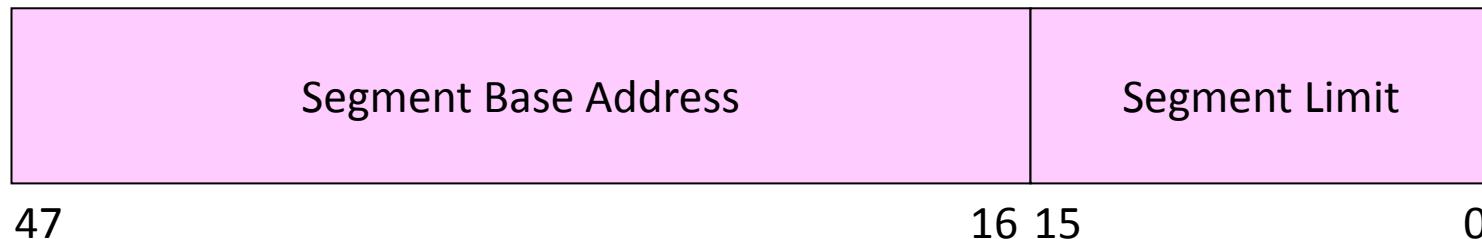
Bus 0, device 29, function 1:
    USB Controller: Intel Corp. 82801DB USB (Hub #2) (rev 2).
        IRQ 19.
        I/O at 0xd400 [0xd41f].
Bus 0, device 31, function 1:
    IDE interface: Intel Corp. 82801DB ICH4 IDE (rev 2).
        IRQ 16.
        I/O at 0xf000 [0xf00f].
        Non-prefetchable 32 bit memory at 0x80000000 [0x800003ff].
Bus 3, device 1, function 0:
    Ethernet controller: Broadcom NetXtreme BCM5703X Gigabit Eth (rev 2).
        IRQ 48.
        Master Capable. Latency=64. Min Gnt=64.
        Non-prefetchable 64 bit memory at 0xf7000000 [0xf700ffff].
```

Three crucial data-structures

- **The Global Descriptor Table (GDT)**
defines the system's memory-segments and their access-privileges, which the CPU has the duty to enforce
- **The Interrupt Descriptor Table (IDT)**
defines entry-points for the various code-routines that will handle all 'interrupts' and 'exceptions'
- **The Task-State Segment (TSS)**
holds the values for registers SS and ESP that will get loaded by the CPU upon entering kernel-mode

How does CPU find GDT/IDT?

- Two dedicated registers: GDTR and IDTR
- Both have identical 48-bit formats:



Kernel must setup these registers during system startup (set-and-forget)

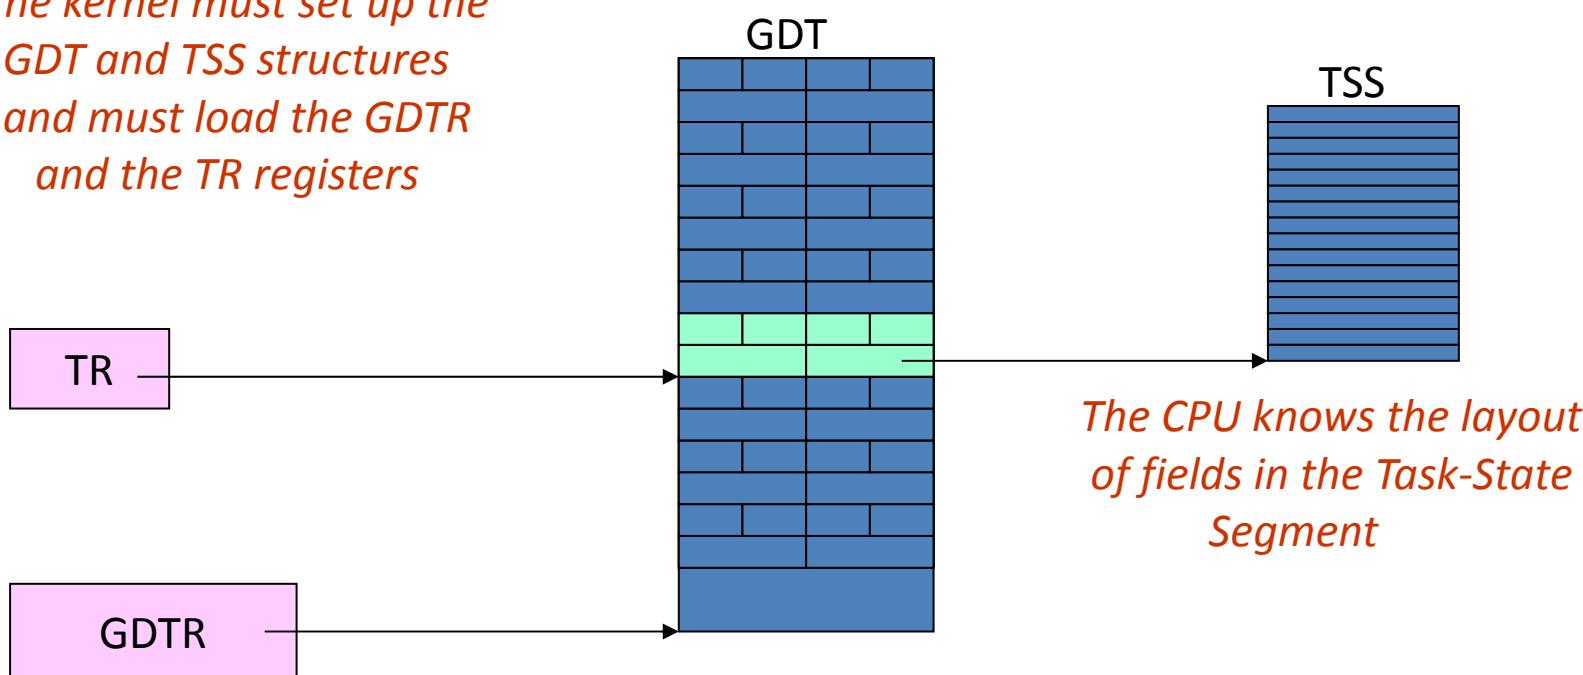
Privileged instructions: LGDT and LIDT used to set these register-values

Unprivileged instructions: SGDT and SIDT used for reading register-values

How does CPU find the TSS?

- Dedicated system segment-register TR holds a descriptor's offset into the GDT

The kernel must set up the GDT and TSS structures and must load the GDTR and the TR registers



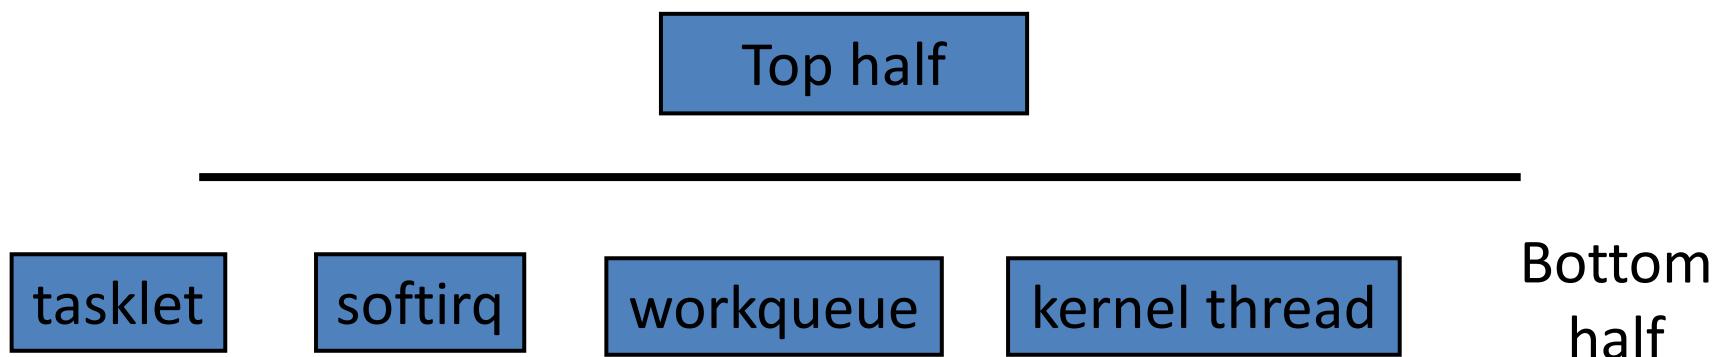
BOTTOM HALF

Summary: Interrupts vs Exceptions

- Varying terminology but for Intel:
 - **Interrupt** (asynchronous, device generated)
 - Maskable: device-generated, associated with IRQs (interrupt request lines); may be temporarily disabled (still pending)
 - Nonmaskable: some critical hardware failures
 - **Exceptions** (synchronous, from software)
 - Processor-detected
 - Faults – correctable (restartable); e.g. page fault
 - Traps – no reexecution needed; e.g. breakpoint
 - Aborts – severe error; process usually terminated (by signal)
 - Programmed exceptions (**software interrupts**)
 - int (system call), int3 (breakpoint)
 - into (overflow), bounds (address check)

Interrupt Handling Philosophy

- Do as little as possible in the interrupt handler
- Defer non-critical actions till later
- Structure: top and bottom halves
 - Top-half: do minimum work and return (ISR)
 - Bottom-half: deferred processing (softirqs, tasklets, workqueues, kernel threads)



Top Half: Do it Now!

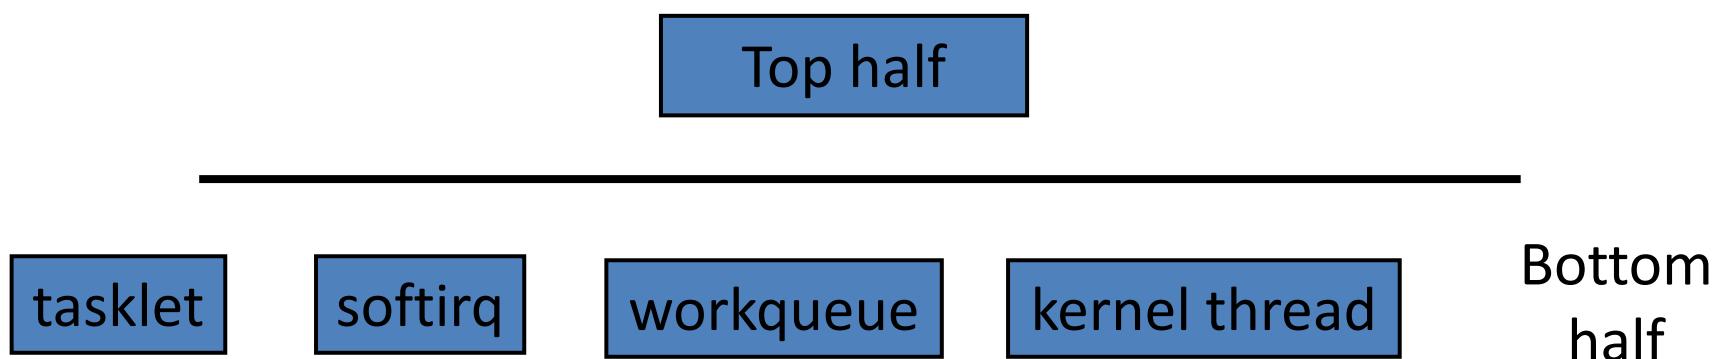
- Technically *is* the interrupt handler
- Perform minimal, common functions: save registers, unmask *other* interrupts. Eventually, undoes that: restores registers, returns to previous context.
 - Often written in assembler
- IRQ is typically masked for duration of top half
- Most important: call proper interrupt handler provided in device drivers (C program)
- Don't want to do too much here
 - IRQs are masked for part of the time
 - Don't want stack to get too big
- Typically queue the request and set a flag for deferred processing in a bottom half

Top Half: Find the Handler

- On modern hardware, multiple I/O devices can share a single IRQ and hence interrupt vector
- First differentiator is the interrupt *vector*
- Multiple *interrupt service routines* (ISR) can be associated with a vector
- Each device's ISR for that IRQ is called
- Device determines whether IRQ is for it

Bottom Half: Do it Later!

- Mechanisms to defer work to later:
 - *softirqs*
 - *tasklets* (built on top of softirqs)
 - *work queues*
 - *kernel threads*
- All can be interrupted



Warning: No Process Context

- Interrupts (as opposed to exceptions) are not associated with particular instructions
- They're also not associated with a given process (user program)
- The currently-running process, at the time of the interrupt, has no relationship whatsoever to that interrupt
- Interrupt handlers cannot sleep!

What Can't You Do?

- You cannot sleep
 - or call something that *might* sleep
- You cannot refer to `current`
- You cannot allocate memory with `GPF_KERNEL` (which can sleep), you must use `GPF_ATOMIC` (which can fail)
- You cannot call `schedule()`
- You cannot do a `down()` semaphore call
 - However, you *can* do an `up()`
- You cannot transfer data to/from user space
 - E.g., `copy_to_user()`, `copy_from_user()`

Interrupt Stack

- When an interrupt occurs, what stack is used?
 - Exceptions: The *kernel stack* of the current process, whatever it is, is used (There’s always some process running — the “idle” process, if nothing else)
 - Interrupts: hard IRQ stack (1 per processor)
 - SoftIRQs: soft IRQ stack (1 per processor)
- These stacks are configured in the IDT and TSS at boot time by the kernel

Softirqs

- **Statically** allocated: specified at kernel compile time
- Limited number:

<i>Priority</i>	<i>Type</i>
0	High-priority tasklets
1	Timer interrupts
2	Network transmission
3	Network reception
4	Block devices
5	Regular tasklets

When Do Softirqs Run?

- Run at various points by the kernel:
 - After system calls
 - After exceptions
 - After interrupts (top halves/IRQs, including the timer intr)
 - When the scheduler runs ksoftirqd
- Softirq routines can be executed simultaneously on multiple CPUs:
 - Code must be re-entrant
 - Code must do its own locking as needed
- Hardware interrupts always enabled when softirqs are running

Rescheduling Softirqs

- A softirq routine can reschedule itself
- This could starve user-level processes
- Softirq scheduler only runs a limited number of requests at a time
- The rest are executed by a kernel thread, `ksoftirqd`, which competes with user processes for CPU time

Tasklets

- Built on top of softirqs
- Can be created and destroyed dynamically
- Run on the CPU that scheduled it (cache affinity)
- Individual tasklets are locked during execution; no problem about re-entrancy, and no need for locking by the code
- Tasklets can run in parallel on multiple CPUs
 - *Same* tasklet can only run on one CPU
- Were once the preferred mechanism for most deferred activity, now changing

The Trouble with Tasklets

- Hard to get right
- One has to be careful about sleeping
- They run at higher priority than other tasks in the systems
- Can produce uncontrolled latency if coded badly
- Ongoing discussion about eliminating tasklets
- Will likely slowly fade over time

Work Queues

- Always run by kernel threads
 - Are scheduled by the scheduler
- Softirqs and tasklets run in an interrupt context; work queues have a pseudo-process context
 - i.e., have a kernel context but no user context
- Because they have a pseudo-process context, they can sleep
 - Work queues are shared by multiple devices
 - Thus, sleeping will delay other work on the queue
- However, they are kernel-only; there is no user mode associated with it
 - Don't try copying data into/out of user space

Kernel Threads

- Always operate in kernel mode
 - Again, no user context
- 2.6.30 introduced the notion of *threaded interrupt handlers*
 - Imported from the realtime tree
 - `request_threaded_irq()`
 - Now each bottom half has its own context, unlike work queues
 - Idea is to eventually replace tasklets and work queues

Comparing Approaches

	ISR	SoftIRQ	Tasklet	WorkQueue	KThread
Will disable all interrupts?	Briefly	No	No	No	No
Will disable other instances of self?	Yes	Yes	No	No	No
Higher priority than regular scheduled tasks?	Yes	Yes*	Yes*	No	No
Will be run on same processor as ISR?	N/A	Yes	Yes	Yes	Maybe
More than one run can on same CPU?	No	No	No	Yes	Yes
Same one can run on multiple CPUs?	Yes	Yes	No	Yes	Yes
Full context switch?	No	No	No	Yes	Yes
Can sleep? (Has own kernel stack)	No	No	No	Yes	Yes
Can access user space?	No	No	No	No	No

*Within limits, can be run by ksoftirqd

I/O

Yubin Xia, Rong Chen
IPADS, SJTU

ACKs: Some slides are adapted from the textbook's original slides
and frans's os course notes

Review

- Interrupt concepts
 - IDT, Trap Num, Interrupt vector, IRQ, System call number, ...
 - Vectors usually IRQ# + 32
 - Triple fault
- Interrupt process
 - Device, PIC (8259A), IDT, trap()
 - CPU checks interrupt every cycle, if interrupt enabled
 - Masked interrupt, umaskable interrupt

Intel-Reserved ID-Numbers

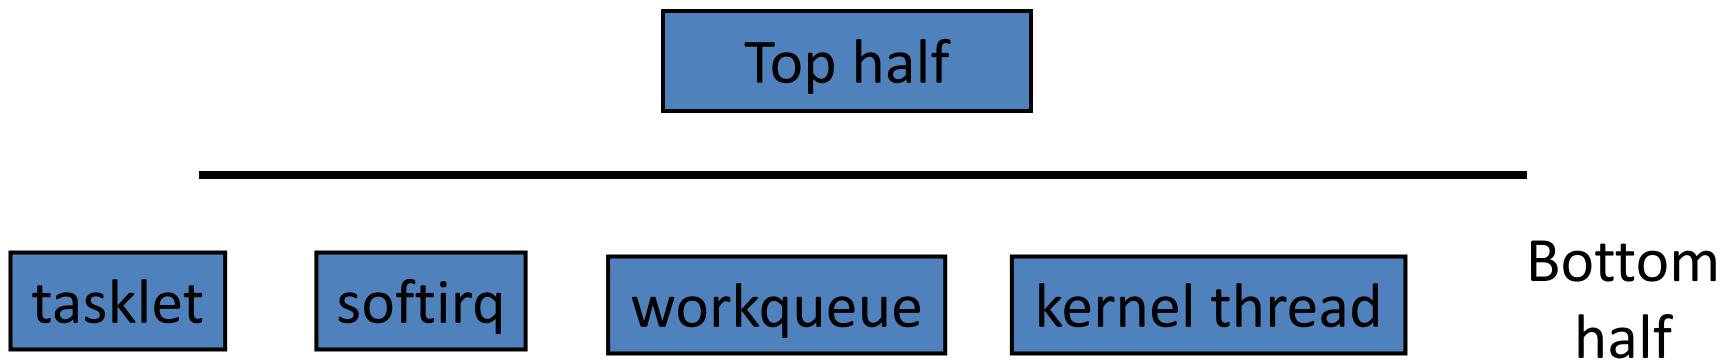
- Of the 256 possible interrupt ID numbers, Intel reserves the first 32 for ‘exceptions’
- OS’s such as Linux are free to use the remaining 224 available interrupt ID numbers for their own purposes (e.g., for service-requests from external devices, or for other purposes such as system-calls)
- Examples:
 - 0: divide-overflow fault
 - 6: Undefined Opcode
 - 7: Coprocessor Not Available
 - 11: Segment-Not-Present fault
 - 12: Stack fault
 - 13: General Protection Exception
 - 14: Page-Fault Exception

Summary: Interrupts vs Exceptions

- Varying terminology but for Intel:
 - **Interrupt** (asynchronous, device generated)
 - Maskable: device-generated, associated with IRQs (interrupt request lines); may be temporarily disabled (still pending)
 - Nonmaskable: some critical hardware failures
 - **Exceptions** (synchronous, from software)
 - Processor-detected
 - Faults – correctable (restartable); e.g. page fault
 - Traps – no reexecution needed; e.g. breakpoint
 - Aborts – severe error; process usually terminated (by signal)
 - Programmed exceptions (**software interrupts**)
 - int (system call), int3 (breakpoint)
 - into (overflow), bounds (address check)

Interrupt Handling Philosophy

- Do as little as possible in the interrupt handler
- Defer non-critical actions till later
- Structure: top and bottom halves
 - Top-half: do minimum work and return (ISR)
 - Bottom-half: deferred processing (softirqs, tasklets, workqueues, kernel threads)



Warning: No Process Context

- Interrupts (as opposed to exceptions) are not associated with particular instructions
- They're also not associated with a given process (user program)
- The currently-running process, at the time of the interrupt, as no relationship whatsoever to that interrupt
- Interrupt handlers cannot sleep!

What Can't You Do in an Interrupt Handler?

- You cannot sleep
 - or call something that *might* sleep
- You cannot refer to `current`
- You cannot allocate memory with `GPF_KERNEL` (which can sleep), you must use `GPF_ATOMIC` (which can fail)
- You cannot call `schedule()`
- You cannot do a `down()` semaphore call
 - However, you *can* do an `up()`
- You cannot transfer data to/from user space
 - E.g., `copy_to_user()`, `copy_from_user()`

Interrupt Stack

- When an interrupt occurs, what stack is used?
 - Exceptions: The *kernel stack* of the current process, whatever it is, is used (There’s always some process running — the “idle” process, if nothing else)
 - Interrupts: hard IRQ stack (1 per processor)
 - SoftIRQs: soft IRQ stack (1 per processor)
- These stacks are configured in the IDT and TSS at boot time by the kernel

Softirqs

- **Statically** allocated: specified at kernel compile time
- Limited number:

<i>Priority</i>	<i>Type</i>
0	High-priority tasklets
1	Timer interrupts
2	Network transmission
3	Network reception
4	Block devices
5	Regular tasklets

When Do Softirqs Run?

- Run at various points by the kernel:
 - After system calls
 - After exceptions
 - After interrupts (top halves/IRQs, including the timer intr)
 - When the scheduler runs ksoftirqd
- Softirq routines can be executed simultaneously on multiple CPUs:
 - Code must be re-entrant
 - Code must do its own locking as needed
- Hardware interrupts always enabled when softirqs are running

Rescheduling Softirqs

- A softirq routine can reschedule itself
- This could starve user-level processes
- Softirq scheduler only runs a limited number of requests at a time
- The rest are executed by a kernel thread, `ksoftirqd`, which competes with user processes for CPU time

Tasklets

- Built on top of softirqs
- Can be created and destroyed dynamically
- Run on the CPU that scheduled it (cache affinity)
- Individual tasklets are locked during execution; no problem about re-entrancy, and no need for locking by the code
- Tasklets can run in parallel on multiple CPUs
 - *Same* tasklet can only run on one CPU
- Were once the preferred mechanism for most deferred activity, now changing

The Trouble with Tasklets

- Hard to get right
- One has to be careful about sleeping
- They run at higher priority than other tasks in the systems
- Can produce uncontrolled latency if coded badly
- Ongoing discussion about eliminating tasklets
- Will likely slowly fade over time

Work Queues

- Always run by kernel threads
 - Are scheduled by the scheduler
- Softirqs and tasklets run in an interrupt context; work queues have a pseudo-process context
 - i.e., have a kernel context but no user context
- Because they have a pseudo-process context, they can sleep
 - Work queues are shared by multiple devices
 - Thus, sleeping will delay other work on the queue
- However, they are kernel-only; there is no user mode associated with it
 - Don't try copying data into/out of user space

Kernel Threads

- Always operate in kernel mode
 - Again, no user context
- 2.6.30 introduced the notion of *threaded interrupt handlers*
 - Imported from the realtime tree
 - `request_threaded_irq()`
 - Now each bottom half has its own context, unlike work queues
 - Idea is to eventually replace tasklets and work queues

Comparing Approaches

	ISR	SoftIRQ	Tasklet	WorkQueue	KThread
Will disable all interrupts?	Briefly	No	No	No	No
Will disable other instances of self?	Yes	Yes	No	No	No
Higher priority than regular scheduled tasks?	Yes	Yes*	Yes*	No	No
Will be run on same processor as ISR?	N/A	Yes	Yes	Yes	Maybe
More than one run can on same CPU?	No	No	No	Yes	Yes
Same one can run on multiple CPUs?	Yes	Yes	No	Yes	Yes
Full context switch?	No	No	No	Yes	Yes
Can sleep? (Has own kernel stack)	No	No	No	Yes	Yes
Can access user space?	No	No	No	No	No

*Within limits, can be run by ksoftirqd

FLEXSC

Flexible System Call

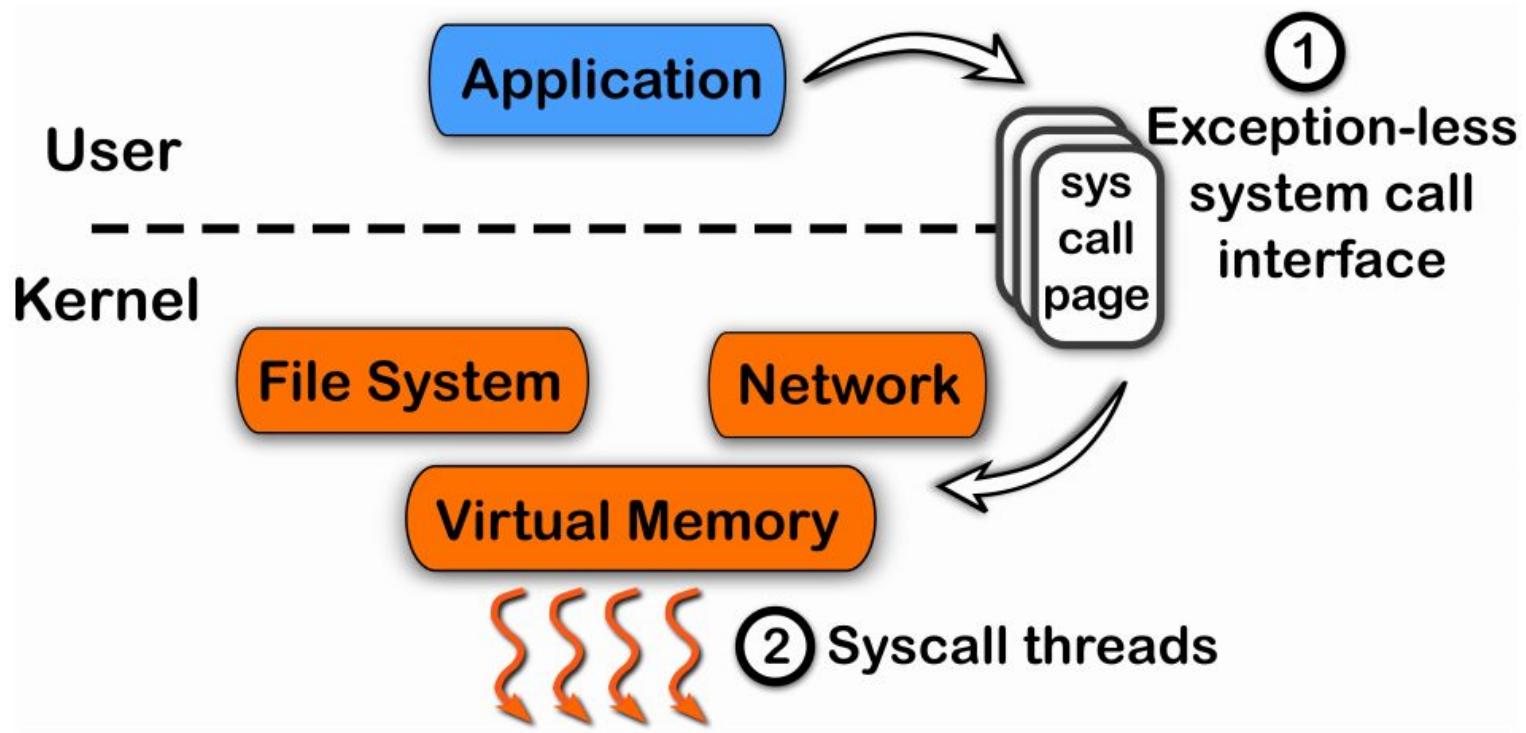
- Basic Idea
 - “Flexible System Call Scheduling with Exception-Less System Calls”
 - University of Toronto
 - OSDI 2010

Flexible System Call

- New system call mechanism – **Flexible System Call**
 - no direct mapping between user threads and kernel threads
 - *system call page*
 - User threads can push the system call requests into the system call page
 - kernel threads will poll the system call requests out the system call page

FlexSC: Another Way to System Call

- Exception-less syscall: remove synchronicity by decoupling invocation from execution



Exception-less System Call

```
write(fd, buf, 4096);
```

```
entry = free_syscall_entry();
```

```
/* write syscall */
entry->syscall = 1;
entry->num_args = 3;
entry->args[0] = fd;
entry->args[1] = buf;
entry->args[2] = 4096;
entry->status = SUBMIT;
```

```
while (entry->status != DONE)
    do_something_else();
```

```
return entry->return_code;
```

syscall number	number of args	args 0 ... 6	status	return code
			⋮	

```
write(fd, buf, 4096);
```



```
entry = free_syscall_entry();
```

```
/* write syscall */  
entry->syscall = 1;  
entry->num_args = 3;  
entry->args[0] = fd;  
entry->args[1] = buf;  
entry->args[2] = 4096;  
entry->status = SUBMIT;
```

```
while (entry->status != DONE)  
    do_something_else();
```

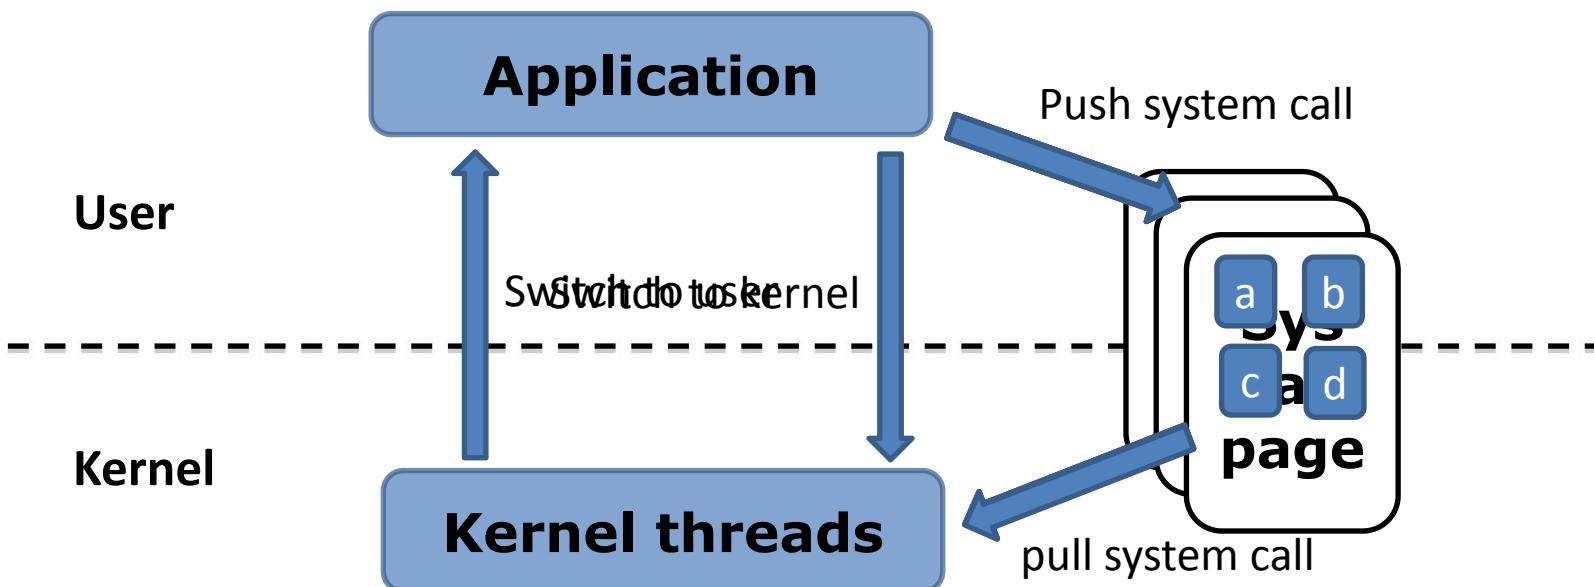
```
return entry->return_code;
```

syscall number	number of args	args 0 ... 6	status	return code
1	3	fd, buf, 4096	DONE	4096

Two orange arrows pointing upwards from the "args" and "status" columns of the last row in the table, indicating the mapping of the syscall arguments and status back to the program's memory.

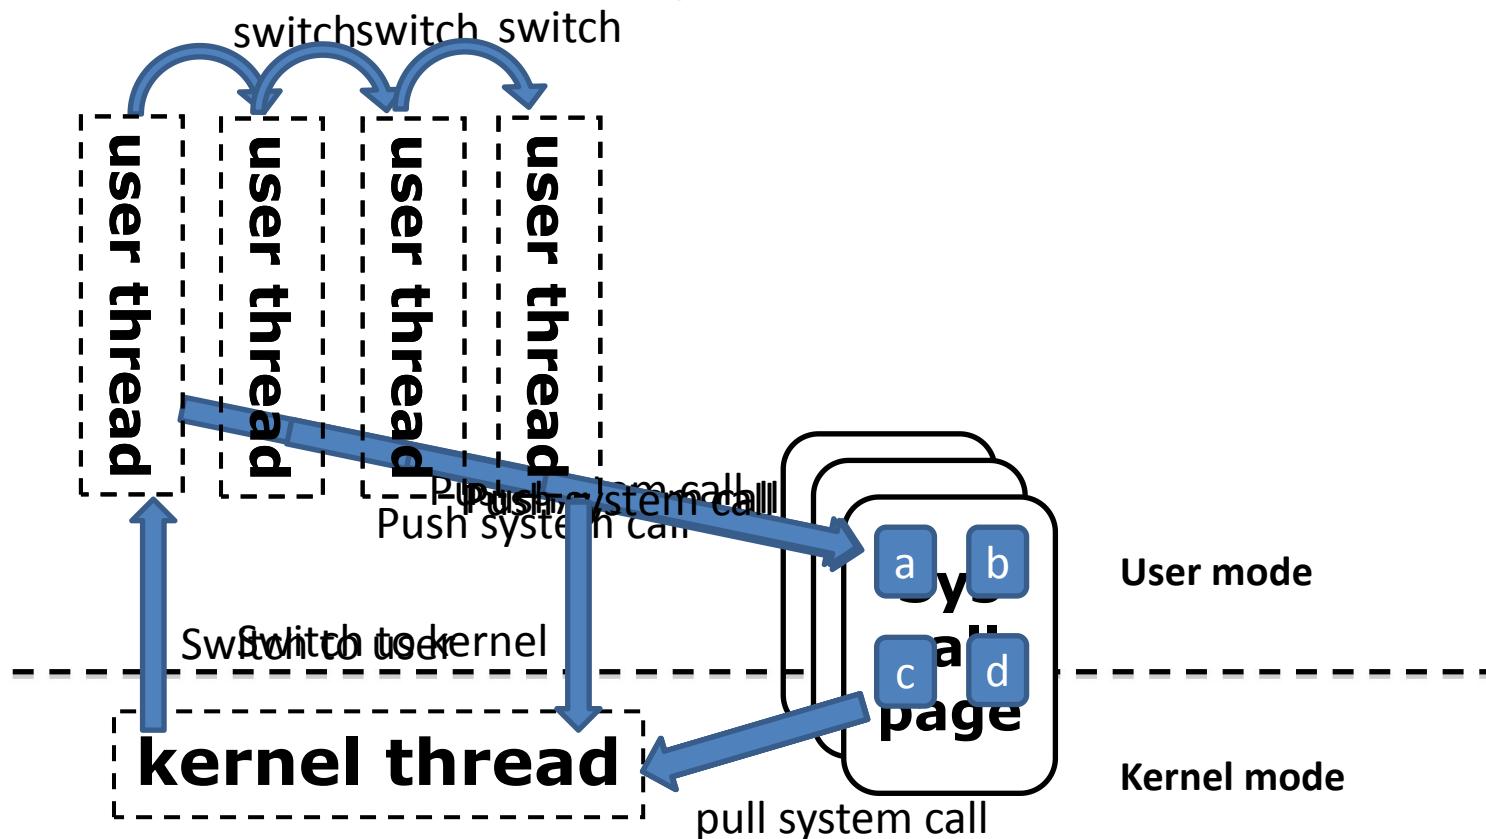
Flexible System Call

- Basic Idea



Running on a single core

- How it works in a single core machine?

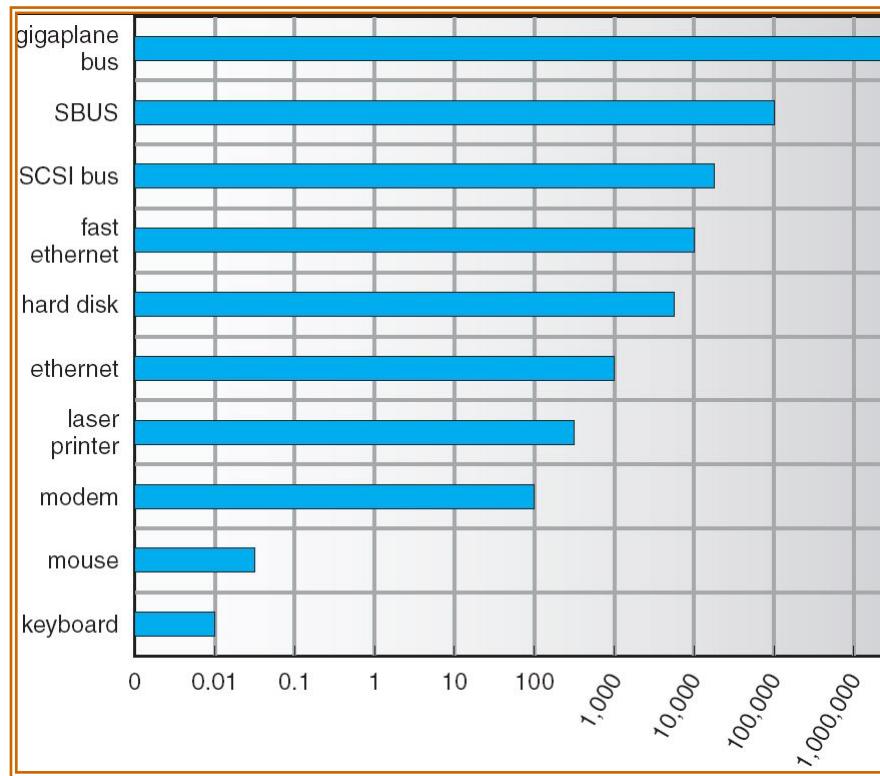


I/O

Why do we Need an I/O Subsystem

- Thousands of devices, each slightly different
 - How can we standardize the interfaces to these devices?
- Devices are **unreliable**
 - Media failures and transmission errors
 - How can we make them reliable?
- Devices are **unpredictable** and/or **slow**
 - How can we manage them if we don't know what they will do or how they will perform?

Device Rates Vary Many Orders of Magnitude



- I/O subsystem better handle this wide range
 - Better not have high overhead/byte for fast devices!
 - Better not waste time waiting for slow devices

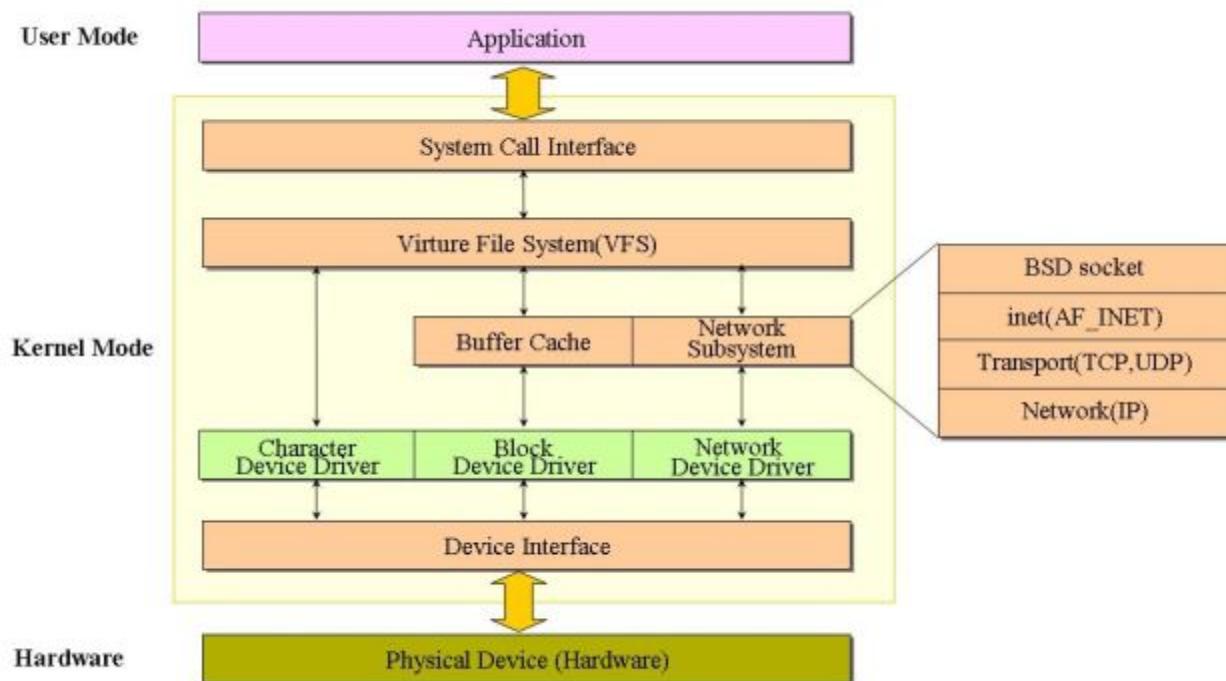
Goal of I/O Subsystem

- Provides uniform interfaces, despite wide range of different devices
 - This code works on many different devices:

```
FILE fd = fopen("/dev/something", "rw");
for (int i = 0; i < 10; i++) {
    fprintf(fd, "Count %d\n", i);
}
close(fd);
```
 - Why?
 - Because code that controls **devices** (“**device driver**”) implements standard interface.
- Provides a layer of abstraction of I/O hardware
 - Manages and interacts with hardware
 - Hides hardware and operation details

Three Types of Device Interfaces

- Three common types of device interfaces
 - Character devices
 - Block devices
 - Network devices



Character Devices

- Character Devices
 - Example: keyboard/mouse, serial port, some USB devices
 - Sequential access, single characters at a time
 - I/O commands: get(), put(), etc.
 - Often use open file interface and semantics

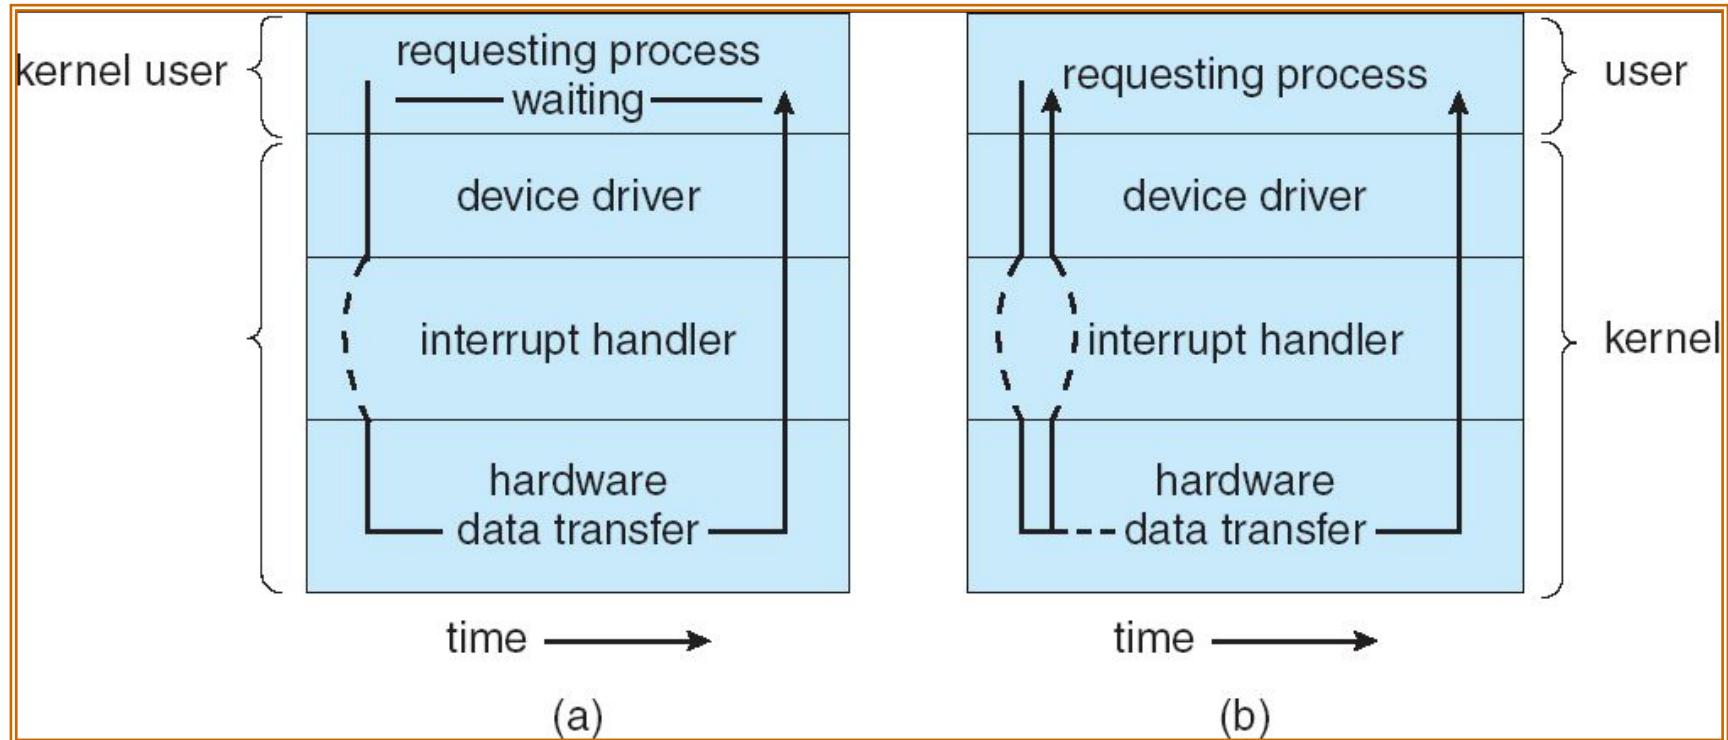
Block Devices and Network Devices

- Block Devices
 - Example: disk drive, tape drive, DVD-ROM
 - Uniform block I/O interface to access blocks of data
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- Network Devices
 - Examples: Ethernet, wireless, bluetooth
 - Different enough from block/character to have own interface
 - Provide special networking interface for supporting various network protocols
 - For example, send/receive network packets

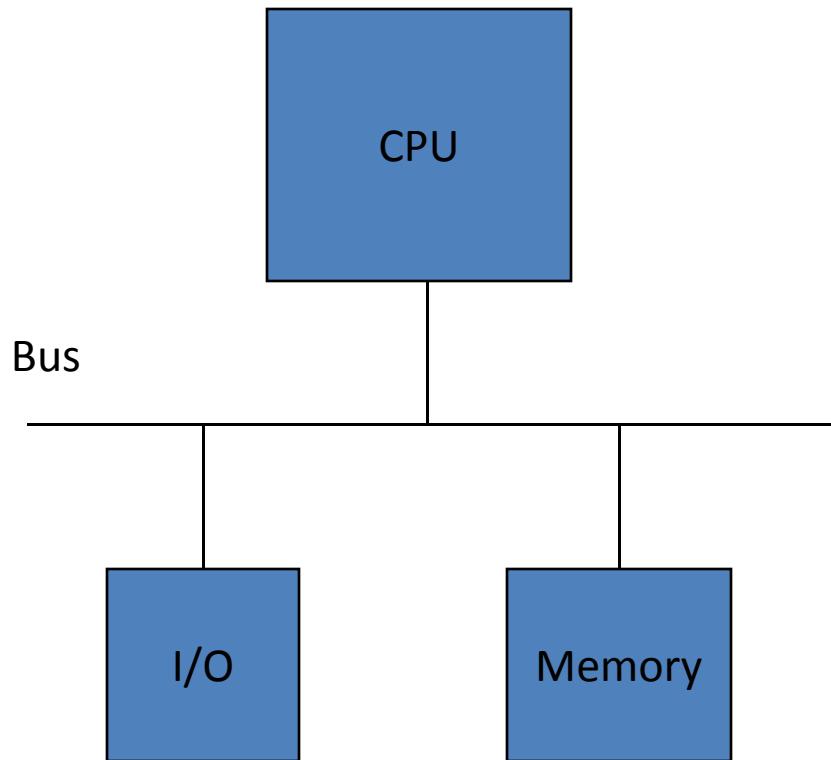
How Does User Deal with Timing?

- **Blocking I/O:** “Wait”
 - When request data (e.g. `read()` system call), put process in waiting state until data is ready
 - When write data (e.g. `write()` system call), put process in waiting state until device is ready for data
- **Non-blocking I/O:** “Don’t Wait”
 - Returns immediately from `read` or `write` request with count of bytes successfully transferred
 - `Read` may return nothing, `write` may write nothing
- **Asynchronous I/O:** “Tell Me Later”
 - When request data, take pointer to user’s buffer, return immediately; later kernel fills buffer and notifies user
 - When send data, take pointer to user’s buffer, return immediately; later kernel takes data and notifies user

Synchronous and Asynchronous I/O

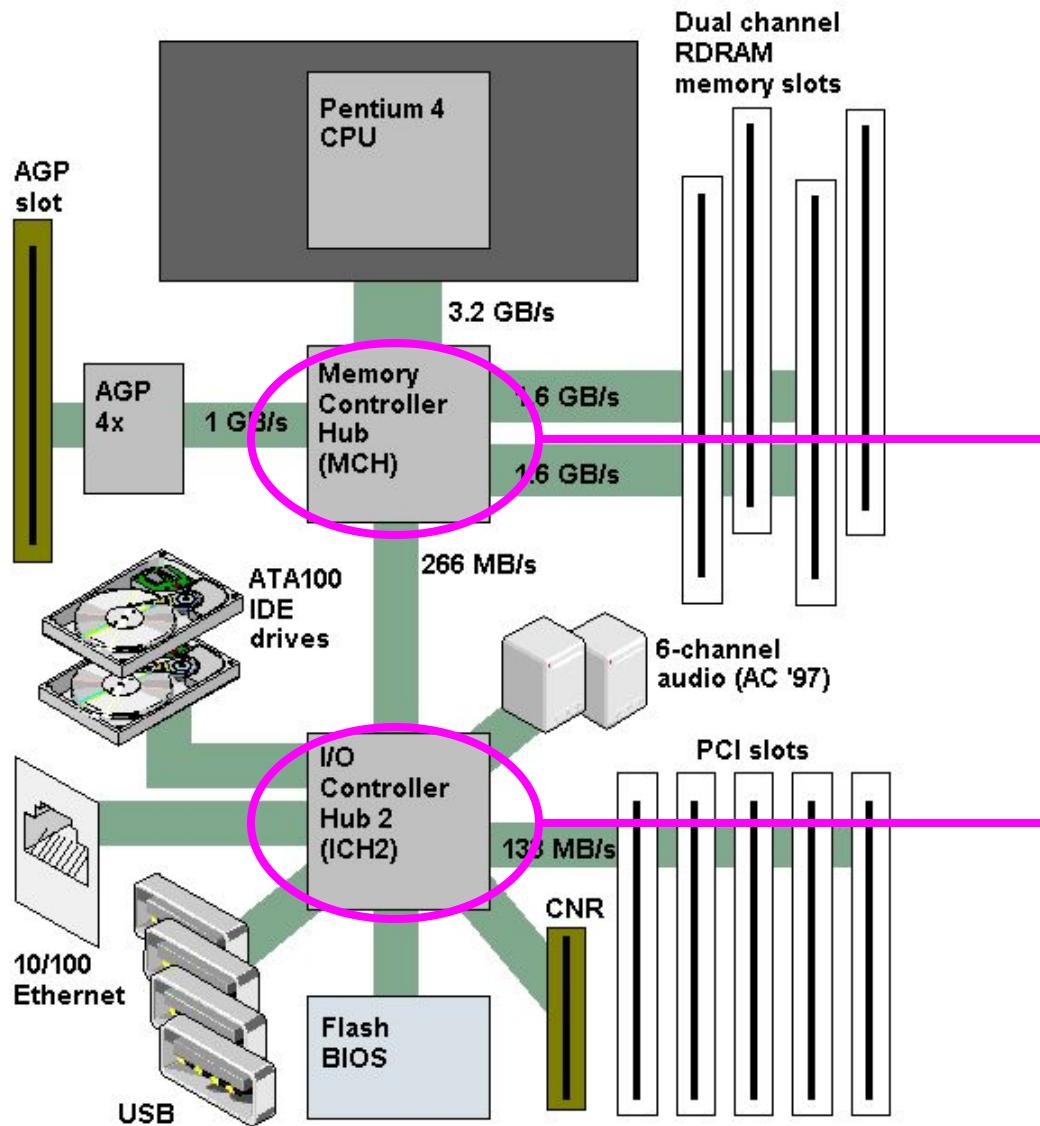


I/O Architecture: The Very Abstract View



I/O Architecture: A Modern Example

From Computer Desktop Encyclopedia
© 2001 The Computer Language Co. Inc.



“Northbridge”

- Memory
- AGP/PCI-Express
- Built-in display

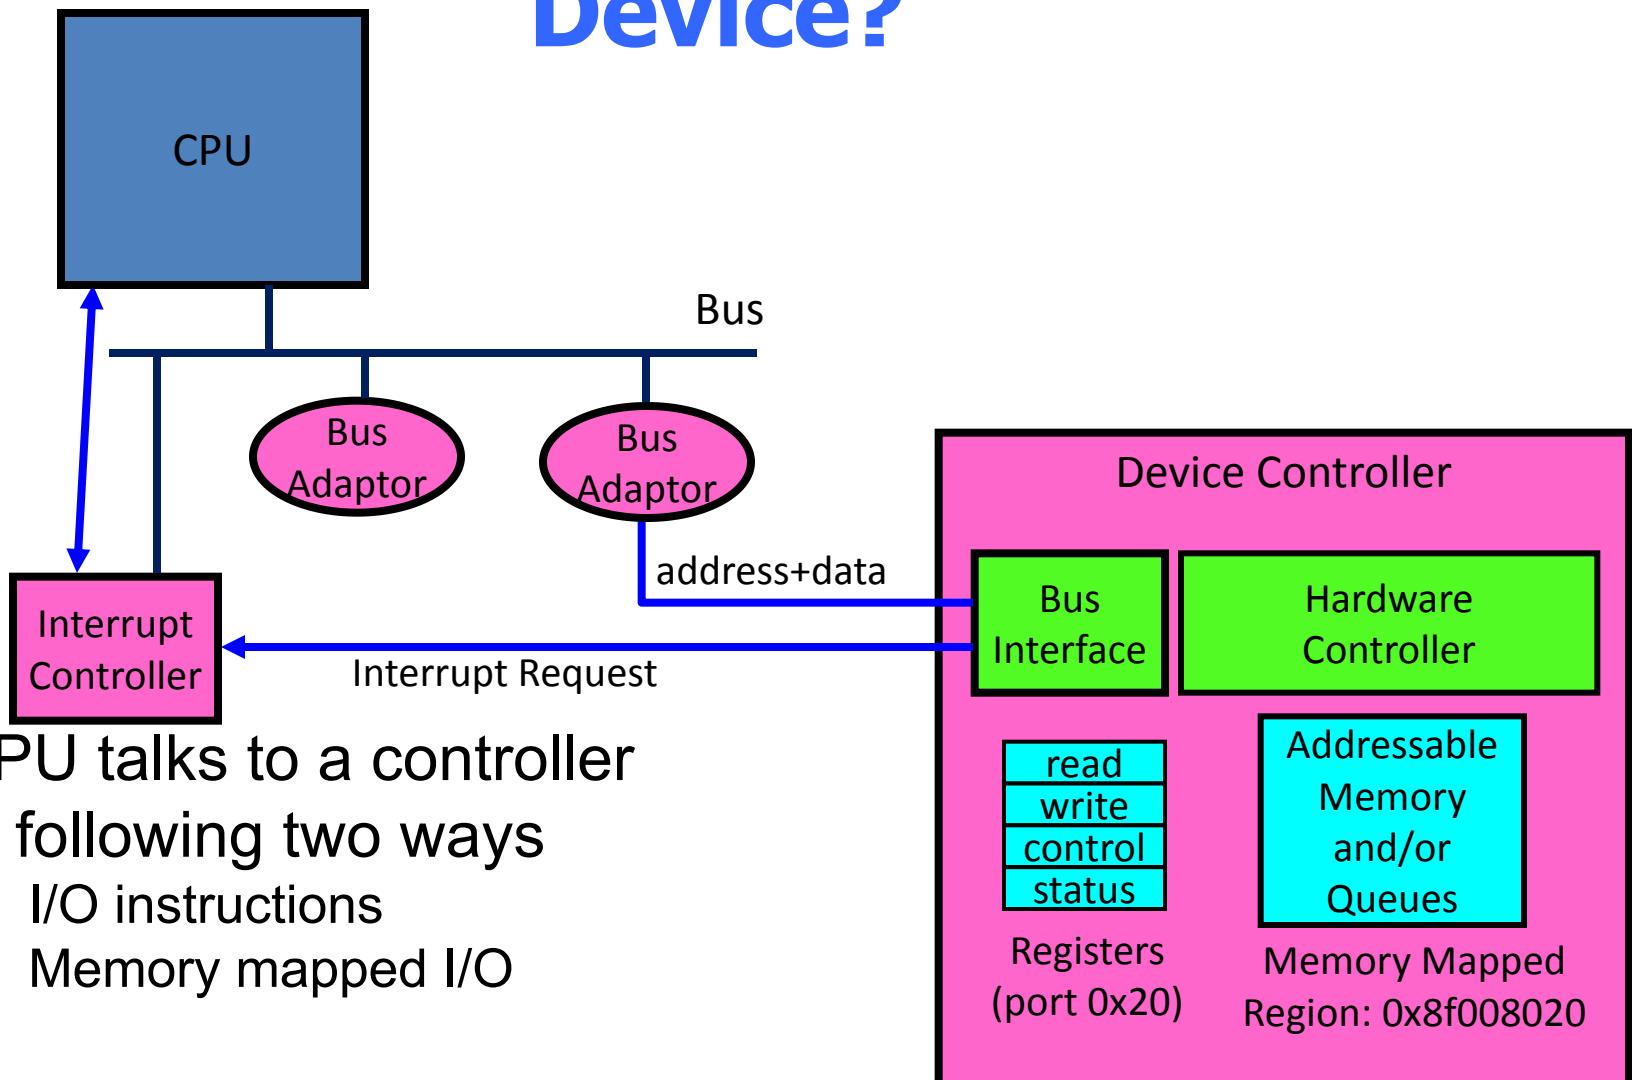
“Southbridge”

- ATA/IDE
- PCI bus
- USB/Firewire bus
- Serial/Parallel ports
- DMA controller
- Interrupt controller
- RTC, ACPI, BIOS, ...

I/O Hardware

- I/O controllers
 - Interface between CPU and I/O devices
 - Provides CPU with special instructions and registers
- I/O addresses
 - “Names” for CPU to control the I/O hardware
 - Memory locations or port numbers
- OS mechanism
 - Use I/O instruction and I/O address to control a device
 - 3 types of interactions with I/O hardware: polling, interrupt-driven, and DMA

How does CPU Actually Connect to Device?

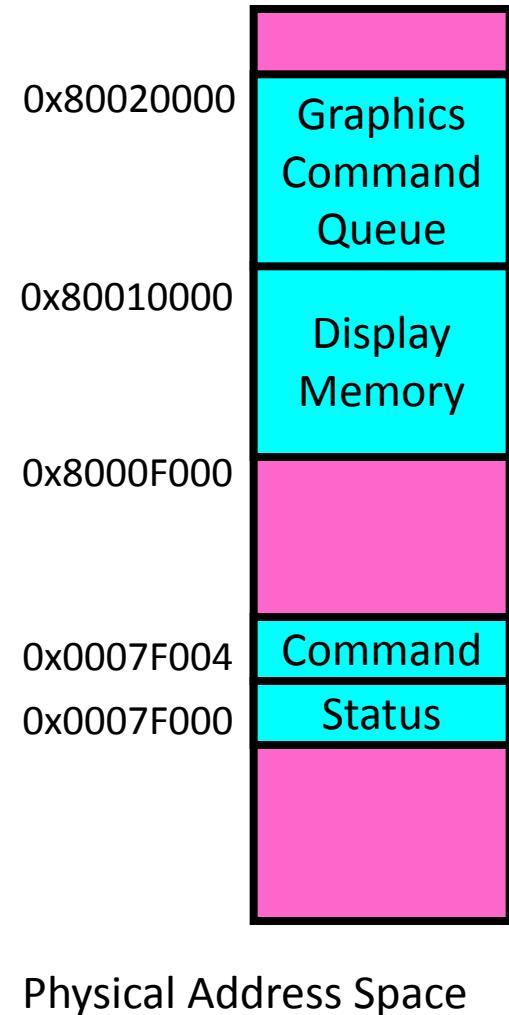


I/O Instructions and Memory-Mapped I/O

- I/O instructions
 - Access device's registers through I/O port numbers
 - Special CPU instructions dealing with I/O
 - Example from the Intel architecture: out 0x21,AL
- Memory mapped I/O (MMIO)
 - Device's registers/memory appear in CPU's physical address space
 - I/O accomplished with memory load/store instructions
 - Mapped by MMU, addresses set by hardware jumpers or programming at boot time
 - Can be protected with page tables

Example: Memory-Mapped Display Controller

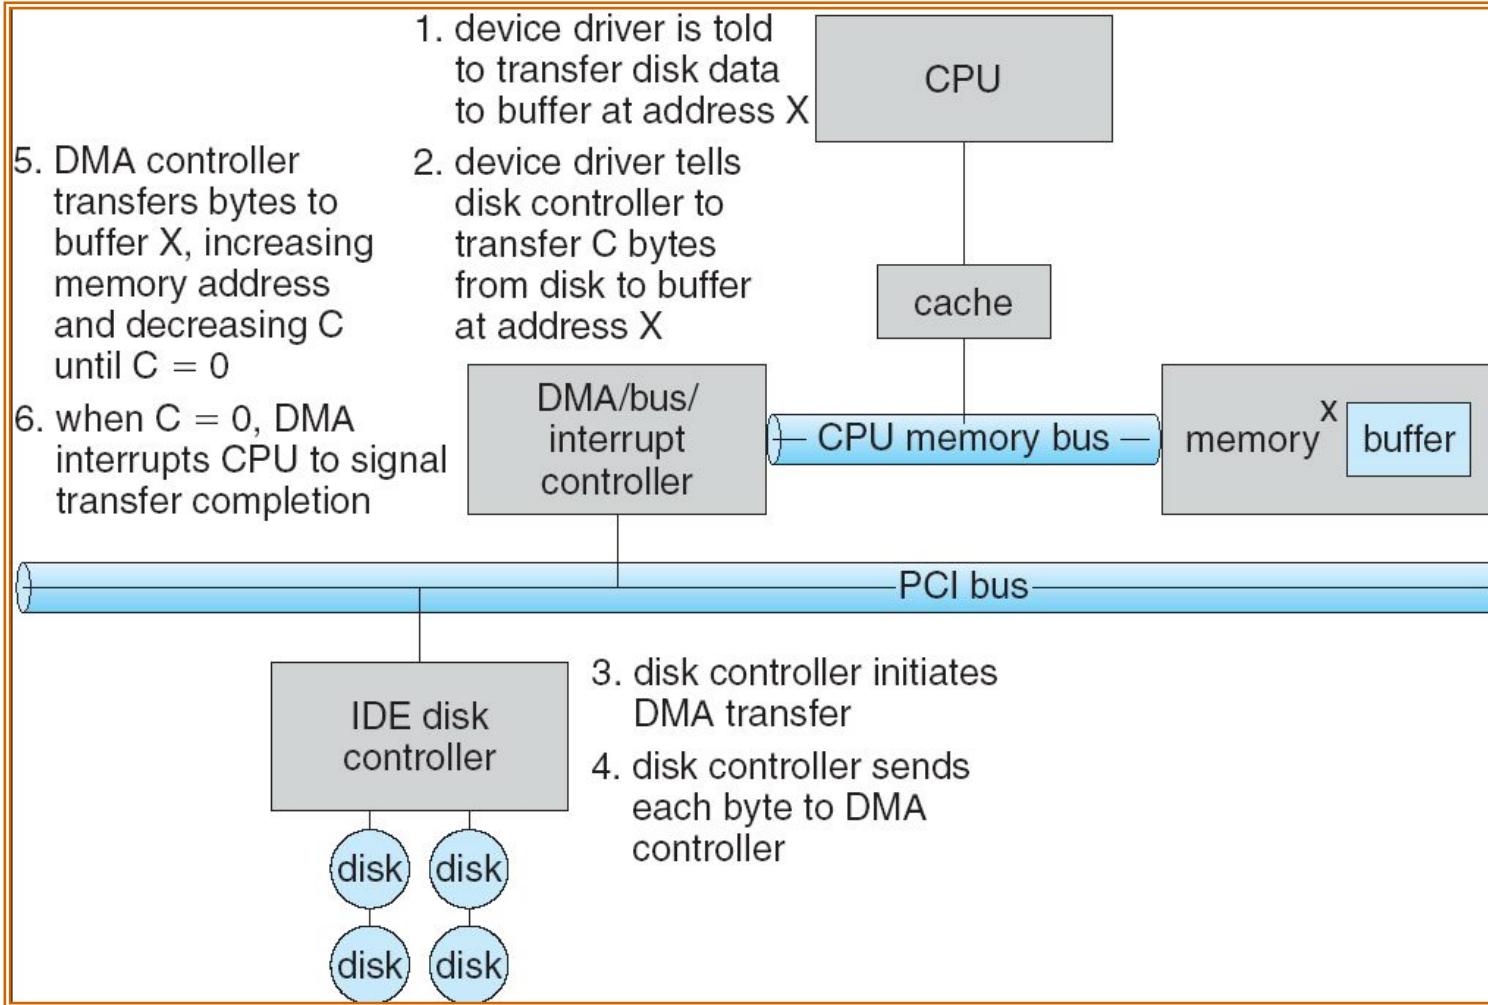
- Simply writing to display memory changes image on screen
 - Addr: 0x8000F000—0x8000FFFF
 - Also called the “frame buffer”
- One may write graphics description to command-queue area
 - E.g., enter a set of triangles that describe some scene
- Writing to the command register may cause on-board graphics hardware to do something
 - Like execute commands in the command-queue area
 - Say render the above scene



Transferring Data To/From Controller

- Programmed I/O (PIO):
 - Each byte transferred via processor in/out or load/store
 - Pro: Simple hardware, easy to program
 - Con: Consumes processor cycles proportional to data size
 - For small/simple I/O
- Direct Memory Access (DMA):
 - Give controller access to memory bus
 - Ask it to transfer data to/from memory directly
 - Pro: device transfers data without burdening CPU
 - Con: need setup
 - For high throughput I/O

Steps in a DMA Transfer



I/O Device Notifying the OS

- The OS needs to know when:
 - The I/O device has completed an operation
 - The I/O operation has encountered an error
- Two methods
 - Polling
 - Interrupt-driven
- Polling:
 - I/O device puts completion/error information in device-specific status register
 - OS periodically checks the status register
 - Pro: low overhead
 - Con: may waste many cycles on polling if infrequent or unpredictable I/O operations

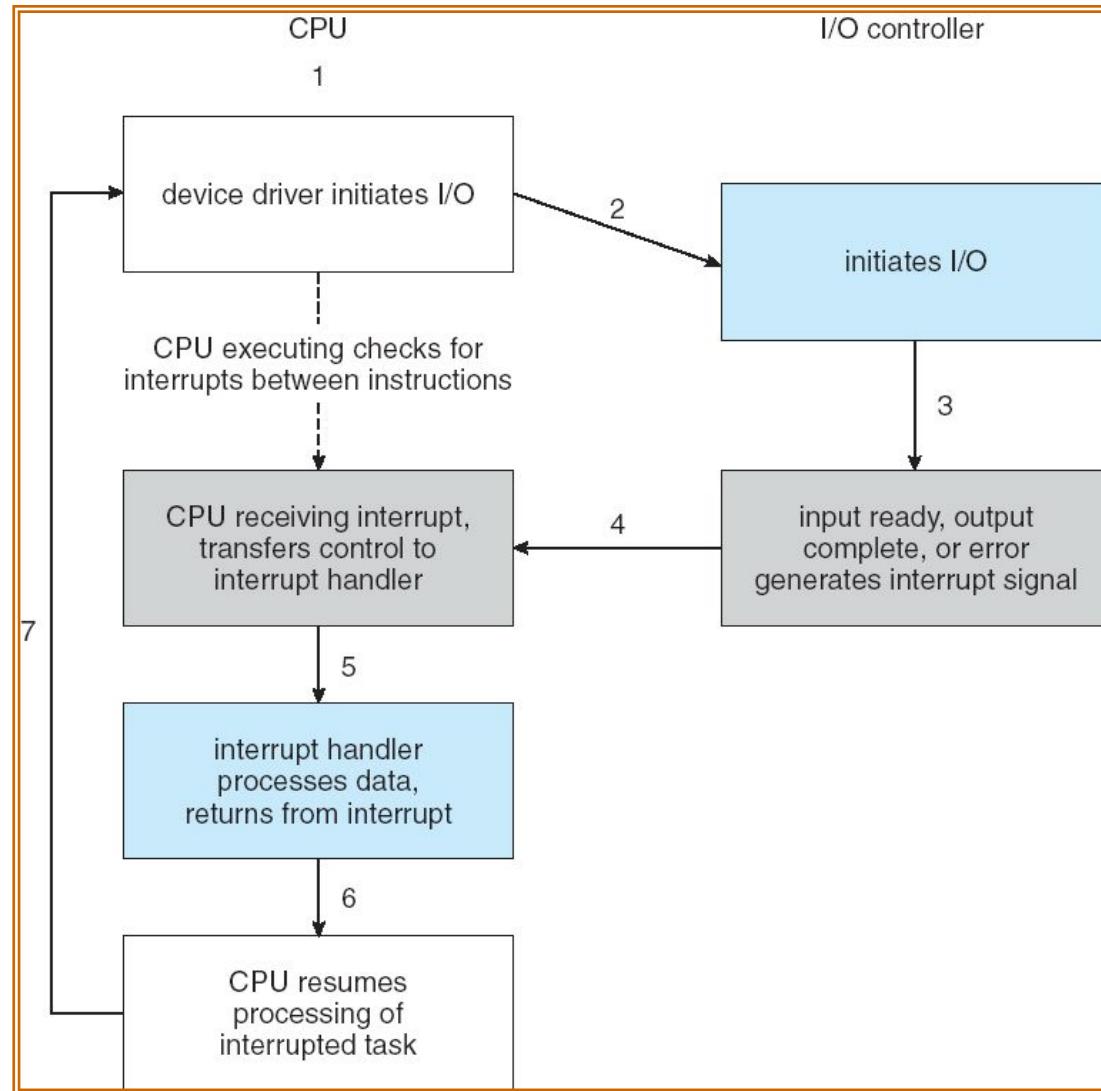
Interrupt Driven

- CPU sets up interrupt handler vector before I/O
- CPU issues I/O request and continues other tasks
- I/O device processes the I/O request
- I/O device triggers CPU interrupt-request line
- Interrupt handler receives interrupts and dispatches to correct handler

Interrupt Driven

- Pro: handles unpredictable events well
- Con: interrupts relatively high overhead
- Some devices may combine both polling and interrupt-driven
 - High-bandwidth network device example:
 - interrupt for first incoming packet
 - polling for following packets until hardware empty

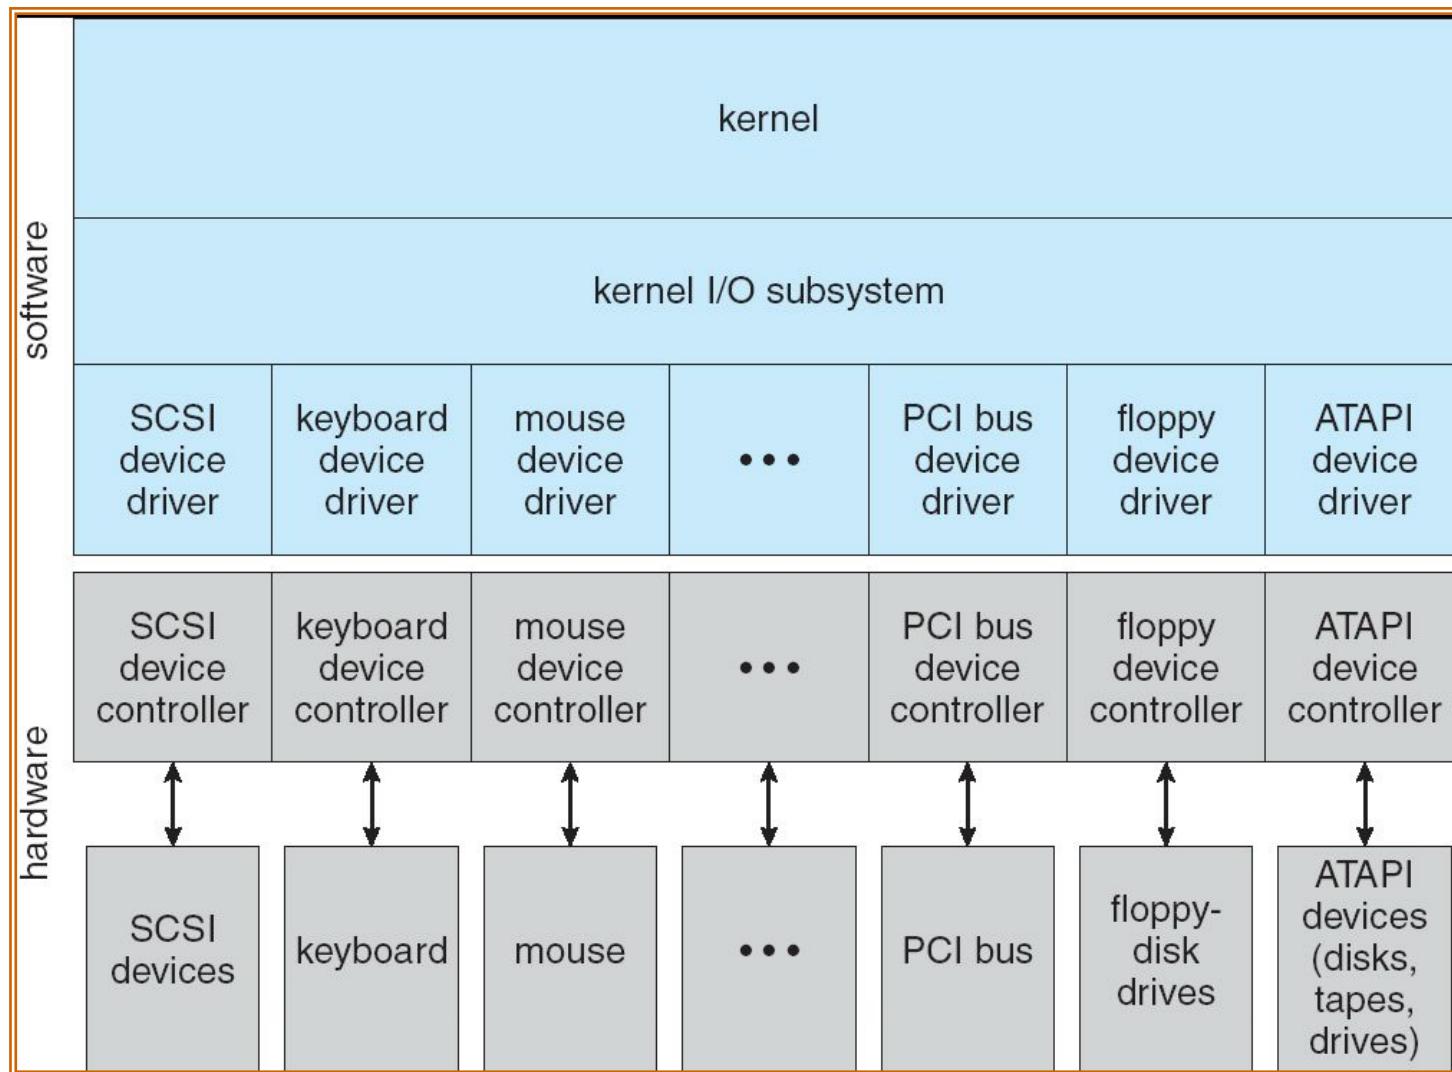
Interrupt-Driven I/O Cycle



Direct Memory Access (DMA)

- OS steps to use DMA:
 - Set up DMA channel (address, registers, direction)
 - Prepare and register interrupt service routine
 - Tell the DMA controller to start
 - When transfer completes, controller will interrupt

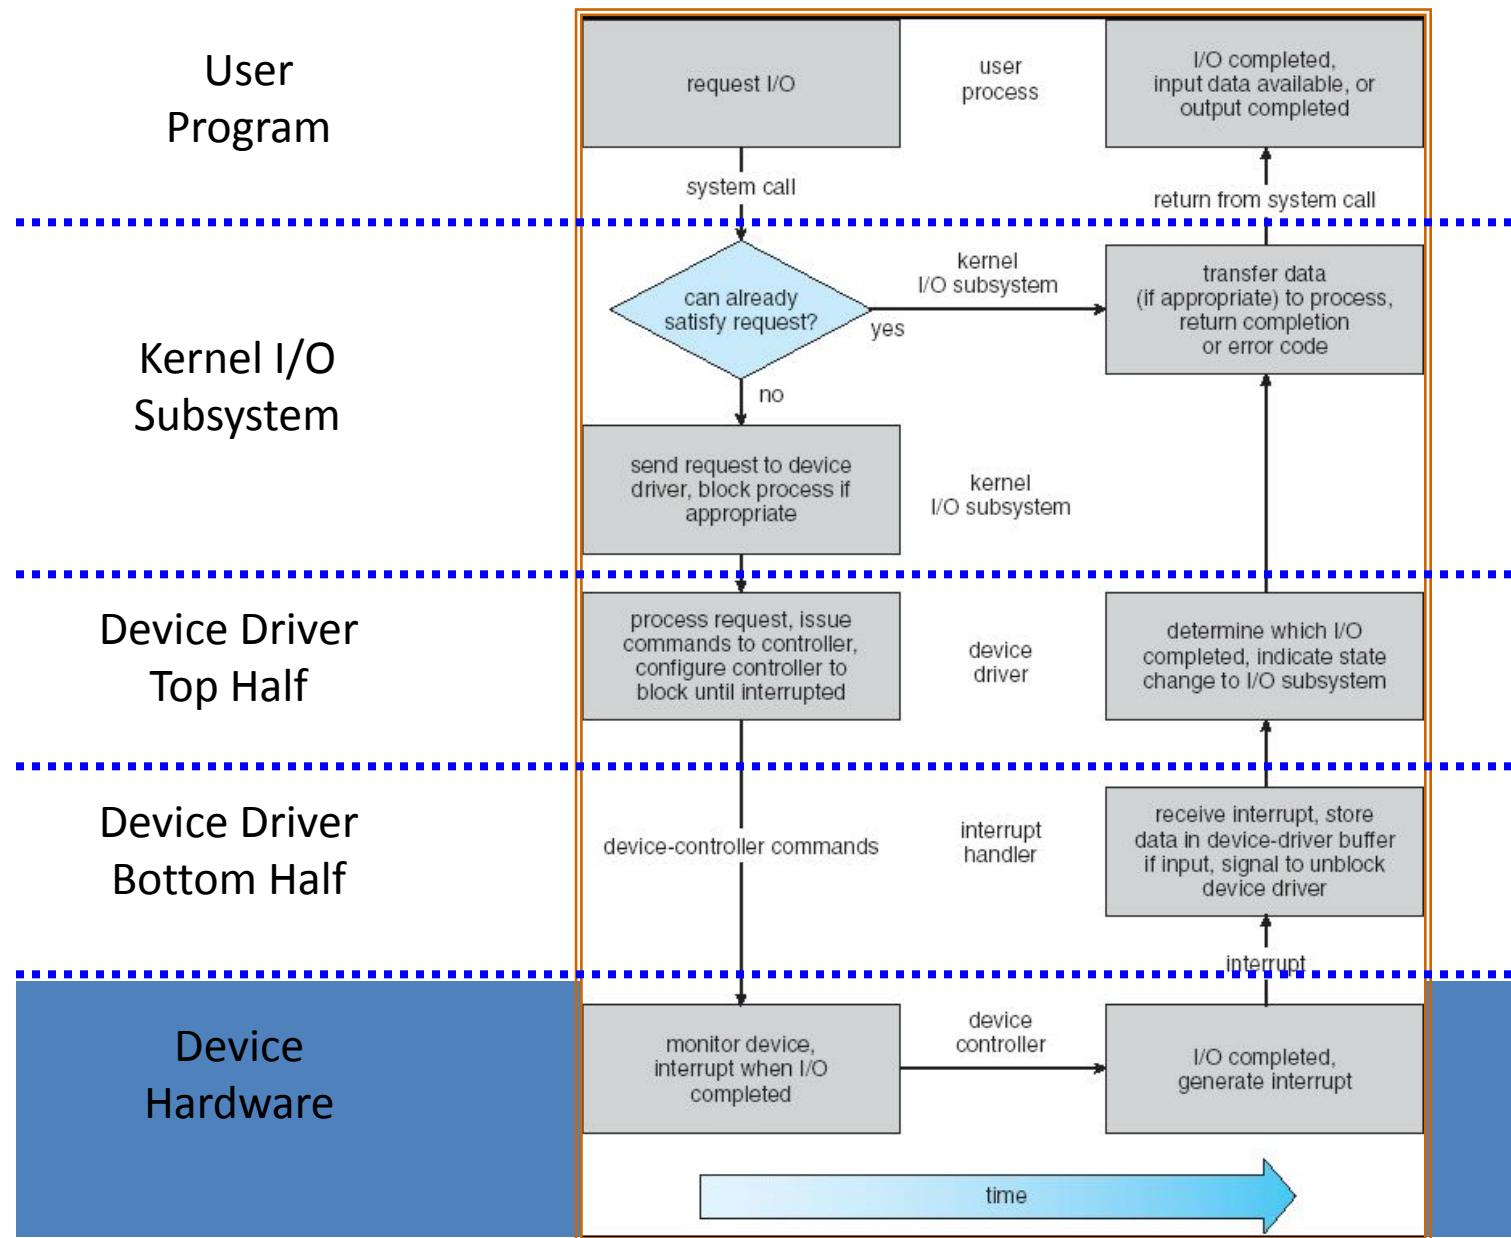
A Kernel I/O Structure



Device Drivers

- Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard internal interface
 - Same kernel I/O system can interact easily with different device drivers
 - Special device-specific configuration supported with the ioctl() system call
- Device drivers typically divided into two pieces:
 - Top half
 - Bottom half

Life Cycle of An I/O Request



Example: Char Device

```
52 // Print to the console. only understands %d, %x, %p, %s.  
53 void  
54 cprintf(char *fmt, ...)  
55 {  
56     int i, c, locking;  
57     uint *argp;  
58     char *s;  
59  
60     locking = cons.locking;  
61     if(locking)  
62         acquire(&cons.lock);  
63  
64     if (fmt == 0)  
65         panic("null fmt");  
66  
67     argp = (uint*)(void*)(&fmt + 1);  
68     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){  
69         if(c != '%'){  
70             consputc(c);  
71             continue;  
72         }  
73         c = fmt[++i] & 0xff;  
74         if(c == 0)  
75             break;  
76         switch(c){  
77             case 'd':  
78                 printint(*argp++, 10, 1);  
79                 break;  
80             case 'x':
```

```
160 void  
161 consputc(int c)  
162 {  
163     if(panicked){  
164         cli();  
165         for(;;)  
166             ;  
167     }  
168  
169     if(c == BACKSPACE){  
170         uartputc('\b'); uartputc(' '); uartputc('\b');  
171     } else  
172         uartputc(c);  
173     cgaputc(c);  
174 }
```

```
51 void  
52 uartputc(int c)  
53 {  
54     int i;  
55  
56     if(!uart)  
57         return;  
58     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)  
59         microdelay(10);  
60     outb(COM1+0, c);  
61 }
```

CGA Operation

```
129 static void
130 cgaputc(int c)
131 {
132     int pos;
133
134     // Cursor position: col + 80*row.
135     outb(CRTPORT, 14);
136     pos = inb(CRTPORT+1) << 8;
137     outb(CRTPORT, 15);
138     pos |= inb(CRTPORT+1);
139
140     if(c == '\n')
141         pos += 80 - pos%80;
142     else if(c == BACKSPACE){
143         if(pos > 0) --pos;
144     } else
145         crt[pos++] = (c&0xff) | 0x0700; // black on white
146
147     if((pos/80) >= 24){ // Scroll up.
148         memmove(crt, crt+80, sizeof(crt[0])*23*80);
149         pos -= 80;
150         memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
151     }
152
153     outb(CRTPORT, 14);
154     outb(CRTPORT+1, pos>>8);
155     outb(CRTPORT, 15);
156     outb(CRTPORT+1, pos);
157     crt[pos] = ' ' | 0x0700;
158 }
```

Recall : read

- readi()->bread()
 - Call bget() to read from buffer cache first
 - If not found, call iderw() to read from disk, and mark as valid.

```
95 // Return a B_BUSY buf with the contents of the indicated disk sector.
96 struct buf*
97 bread(uint dev, uint sector)
98 {
99     struct buf *b;
100
101     b = bget(dev, sector);
102     if(!(b->flags & B_VALID))
103         iderw(b);
104     return b;
105 }
```

Xv6 disk driver design & Implementation

- Design Overview
 - Support an IDE driver
 - Use I/O Instructions instead of Memory Mapped I/O
 - in, out
 - Asynchronous I/O model
 - Use a simple **queue** of I/O request
 - Use **interrupt** to notify the available of data

IDE Specification

- Xv6 uses a relative old IDE specification
 - Not PCI negotiation
- IDE Ports
 - 0x1F0-0x1F7
 - 0x1f0: write/read port
 - 0x1f2: number of sectors
 - 0x1f3-0x1f5: sector number
 - 0x1f6: diskno and sector number
 - 0x1f7: command registers, status bit
 - 0x3F6 interrupt control line

I/O instruction in xv6

```
3 static inline uchar
4 inb(ushort port)
5 {
6     uchar data;
7
8     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
9     return data;
10 }
```

```
21 static inline void
22 outb(ushort port, uchar data)
23 {
24     asm volatile("out %0,%1" : : "a" (data), "d" (port));
25 }
```

Polling Mechanism

```
32 // Wait for IDE disk to become ready.  
33 static int  
34 idewait(int checkerr)  
35 {  
36     int r;  
37  
38     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)  
39         ;  
40     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)  
41         return -1;  
42     return 0;  
43 }
```

IDE Initialization

```
45 void
46 ideinit(void)
47 {
48     int i;
49
50     initlock(&idelock, "ide");
51     picenable(IRQ_IDE);
52     ioapicenable(IRQ_IDE, ncpu - 1);
53     idewait(0);
54
55     // Check if disk 1 is present
56     outb(0x1f6, 0xe0 | (1<<4));
57     for(i=0; i<1000; i++){
58         if(inb(0x1f7) != 0){
59             havedisk1 = 1;
60             break;
61         }
62     }
63
64     // Switch back to disk 0.
65     outb(0x1f6, 0xe0 | (0<<4));
66 }
```

Synchronous I/O (iderw)

```
122 // Sync buf with disk.  
123 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.  
124 // Else if B_VALID is not set, read buf from disk, set B_VALID.  
125 void  
126 iderw(struct buf *b)  
127 {  
128     struct buf **pp;  
129  
130     if(!(b->flags & B_BUSY))  
131         panic("iderw: buf not busy");  
132     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)  
133         panic("iderw: nothing to do");  
134     if(b->dev != 0 && !havedisk1)  
135         panic("iderw: ide disk 1 not present");  
136  
137     acquire(&idelock); //DOC:acquire-lock  
138  
139     // Append b to idequeue.  
140     b->qnext = 0;  
141     for(pp=&idequeue; *pp; pp=&(*pp)->qnext) //DOC:insert-queue  
142         ;  
143     *pp = b;  
144  
145     // Start disk if necessary.  
146     if(idequeue == b)  
147         idestart(b);  
148  
149     // Wait for request to finish.  
150     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){  
151         sleep(b, &idelock);  
152     }  
153  
154     release(&idelock);  
155 }
```

```
340 // Atomically release lock and sleep on chan.  
341 // Reacquires lock when awakened.  
342 void  
343 sleep(void *chan, struct spinlock *lk)  
344 {  
345     if(proc == 0)  
346         panic("sleep");  
347  
348     if(lk == 0)  
349         panic("sleep without lk");  
350  
351     // Must acquire ptable.lock in order to  
352     // change p->state and then call sched.  
353     // Once we hold ptable.lock, we can be  
354     // guaranteed that we won't miss any wakeup  
355     // (wakeup runs with ptable.lock locked),  
356     // so it's okay to release lk.  
357     if(lk != &ptable.lock){ //DOC: sleeplock0  
358         acquire(&ptable.lock); //DOC: sleeplock1  
359         release(lk);  
360     }  
361  
362     // Go to sleep.  
363     proc->chan = chan;  
364     proc->state = SLEEPING;  
365     sched();  
366  
367     // Tidy up.  
368     proc->chan = 0;  
369  
370     // Reacquire original lock.  
371     if(lk != &ptable.lock){ //DOC: sleeplock2  
372         release(&ptable.lock);  
373         acquire(lk);  
374     }  
375 }
```

Sleep and Wakeup Mechanism in xv6

```
378 // Wake up all processes sleeping on chan.  
379 // The ptable lock must be held.  
380 static void  
381 wakeup1(void *chan)  
382 {  
383     struct proc *p;  
384  
385     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
386         if(p->state == SLEEPING && p->chan == chan)  
387             p->state = RUNNABLE;  
388 }  
389  
390 // Wake up all processes sleeping on chan.  
391 void  
392 wakeup(void *chan)  
393 {  
394     acquire(&ptable.lock);  
395     wakeup1(chan);  
396     release(&ptable.lock);  
397 }
```

```
68 // Start the request for b. Caller must hold idelock.
69 static void
70 idestart(struct buf *b)
71 {
72     if(b == 0)
73         panic("idestart");
74
75     idewait(0);
76     outb(0x3f6, 0); // generate interrupt
77     outb(0x1f2, 1); // number of sectors
78     outb(0x1f3, b->sector & 0xff);
79     outb(0x1f4, (b->sector >> 8) & 0xff);
80     outb(0x1f5, (b->sector >> 16) & 0xff);
81     outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((b->sector>>24)&0x0f));
82     if(b->flags & B_DIRTY){
83         outb(0x1f7, IDE_CMD_WRITE);
84         outsl(0x1f0, b->data, 512/4);
85     } else {
86         outb(0x1f7, IDE_CMD_READ);
87     }
88 }
```

```
90 // Interrupt handler.
91 void
92 ideintr(void)
93 {
94     struct buf *b;
95
96     // First queued buffer is the active request.
97     acquire(&idelock);
98     if((b = idequeue) == 0){
99         release(&idelock);
100        // cprintf("spurious IDE interrupt\n");
101        return;
102    }
103    idequeue = b->qnext;
104
105    // Read data if needed.
106    if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
107        insl(0x1f0, b->data, 512/4);
108
109    // Wake process waiting for this buf.
110    b->flags |= B_VALID;
111    b->flags &= ~B_DIRTY;
112    wakeup(b);
113
114    // Start disk on next buf in queue.
115    if(idequeue != 0)
116        idestart(idequeue);
117
118    release(&idelock);
119 }
```

```
12 static inline void
13 insl(int port, void *addr, int cnt)
14 {
15     asm volatile("cld; rep insl" :
16                 "=D" (addr), "=c" (cnt) :
17                 "d" (port), "0" (addr), "1" (cnt) :
18                 "memory", "cc");
19 }
```

IDE.c

- The driver can only handle 1 operation at a time
 - Only send the buffer at the front of the queue
 - Others are simply waiting their turn
- ide rw() maintains a queue of requests
 - Adds the buffer b to the end of the queue
 - If b is the first, then call ide start(), if not, just wait
 - Then just sleep, instead of polling
- ide intr() will handle the first request
 - Then start other request if any

I/O

Yubin Xia, Rong Chen
IPADS, SJTU

ACKs: Some slides are adapted from the textbook's original slides
and frans's os course notes

Review

- FlexSC: Exception-less syscall: remove synchronicity by decoupling invocation from execution
- Why I/O sub-system: Thousands of devices, each slightly different
- Unified interface: file
 - Open, read, write, seek, close, ...
 - Avoid too generalization
- Three types: block, char, network
- Blocking VS. non-blocking
- Sync VS. Async
- PIO, MMIO, DMA
- Top-half & bottom-half

I/O Device Notifying the OS

- The OS needs to know when:
 - The I/O device has completed an operation
 - The I/O operation has encountered an error
- Two methods
 - Polling
 - Interrupt-driven
- Polling:
 - I/O device puts completion/error information in device-specific status register
 - OS periodically checks the status register
 - Pro: low overhead
 - Con: may waste many cycles on polling if infrequent or unpredictable I/O operations

Interrupt Driven

- CPU sets up interrupt handler vector before I/O
- CPU issues I/O request and continues other tasks
- I/O device processes the I/O request
- I/O device triggers CPU interrupt-request line
- Interrupt handler receives interrupts and dispatches to correct handler

Interrupt Driven

- Pro: handles unpredictable events well
- Con: interrupts relatively high overhead
- Some devices may combine both polling and interrupt-driven
 - High-bandwidth network device example:
 - interrupt for first incoming packet
 - polling for following packets until hardware empty

Press 'a' to display 'a'

```
3 .globl alltraps
4 .globl vector0
5 vector0:
6   pushl $0
7   pushl $0
8   jmp alltraps
9 .globl vector1
10 vector1:
11  pushl $0
12  pushl $1
13  jmp alltraps
14 .globl vector2
15 vector2:
16  pushl $0
17  pushl $2
18  jmp alltraps
19 .globl vector3
20 vector3:
21  pushl $0
22  pushl $3
23  jmp alltraps
```

```
3   # vectors.S sends all traps here.
4 .globl alltraps
5 alltraps:
6   # Build trap frame.
7   pushl %ds
8   pushl %es
9   pushl %fs
10  pushl %gs
11  pushal
12
13  # Set up data and per-cpu segments.
14  movw $(SEG_KDATA<<3), %ax
15  movw %ax, %ds
16  movw %ax, %es
17  movw $(SEG_KCPU<<3), %ax
18  movw %ax, %fs
19  movw %ax, %gs
20
21  # Call trap(tf), where tf=%esp
22  pushl %esp
23  call trap
24  addl $4, %esp
25
26  # Return falls through to trapret...
27 .globl trapret
28 trapret:
29  popal
30  popl %gs
31  popl %fs
32  popl %es
33  popl %ds
34  addl $0x8, %esp  # trapno and errcode
35  iret
```

```
36 void
37 trap(struct trapframe *tf)
38 {
39     if(tf->trapno == T_SYSCALL){
40         if(proc->killed)
41             exit();
42         proc->tf = tf;
43         syscall();
44         if(proc->killed)
45             exit();
46         return;
47     }
48
49     switch(tf->trapno){
50     case T_IRQ0 + IRQ_TIMER:
51         if(cpu->id == 0){
52             acquire(&tickslock);
53             ticks++;
54             wakeup(&ticks);
55             release(&tickslock);
56         }
57         lapiceoi();
58         break;
59     case T_IRQ0 + IRQ_IDE:
60         ideintr();
61         lapiceoi();
62         break;
63     case T_IRQ0 + IRQ_IDE+1:
64         // Bochs generates spurious IDE1 interrupts.
65         break;
66     case T_IRQ0 + IRQ_KBD:
67         kbdintr();
68         lapiceoi();
69         break;
70     case T_IRQ0 + IRQ_COM1:
```

```
46 void
47 kbdintr(void)
48 {
49     consoleintr(kbdgetc());
50 }
```



Function pointer

```
187 void
188 consoleintr(int (*getc)(void))
189 {
190     int c;
191
192     acquire(&input.lock);
193     while((c = getc()) >= 0){
194         switch(c){
195             case C('P'): // Process listing.
196                 procdump();
197                 break;
198             case C('U'): // Kill line.
199                 while(input.e != input.w &&
200                       input.buf[(input.e-1) % INPUT_BUF] != '\n'){
201                     input.e--;
202                     consputc(BACKSPACE);
203                 }
204                 break;
205             case C('H'): case '\x7f': // Backspace
206                 if(input.e != input.w){
207                     input.e--;
208                     consputc(BACKSPACE);
209                 }
210                 break;
211             default:
212                 if(c != 0 && input.e-input.r < INPUT_BUF){
213                     c = (c == '\r') ? '\n' : c;
214                     input.buf[input.e++ % INPUT_BUF] = c;
215                     consputc(c);
216                     if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
217                         input.w = input.e;
218                         wakeup(&input.r);
219                     }
220                 }
221                 break;
222             }
223         }
224     release(&input.lock);
225 }
```

```

6 int
7 kbdgetc(void)
8 {
9     static uint shift;
10    static uchar *charcode[4] = {
11        normalmap, shiftmap, ctlmap, ctlmap
12    };
13    uint st, data, c;
14
15    st = inb(KBSTATP);
16    if((st & KBS_DIB) == 0)
17        return -1;
18    data = inb(KBDATAP);
19
20    if(data == 0xE0){
21        shift |= E0ESC;
22        return 0;
23    } else if(data & 0x80){
24        // Key released
25        data = (shift & E0ESC ? data : data & 0x7F);
26        shift &= ~(shiftcode[data] | E0ESC);
27        return 0;
28    } else if(shift & E0ESC){
29        // Last character was an E0 escape; or with 0x80
30        data |= 0x80;
31        shift &= ~E0ESC;
32    }
33
34    shift |= shiftcode[data];
35    shift ^= togglecode[data];
36    c = charcode[shift & (CTL | SHIFT)][data];
37    if(shift & CAPSLOCK){
38        if('a' <= c && c <= 'z')
39            c += 'A' - 'a';
40        else if('A' <= c && c <= 'Z')
41            c += 'a' - 'A';
42    }
43    return c;
44 }

```

3	#define KBSTATP	0x64	// kbd controller status port(I)
4	#define KBS_DIB	0x01	// kbd data in buffer
5	#define KBDATAP	0x60	// kbd data port(I)
6			
7	#define NO	0	
8			
9	#define SHIFT	(1<<0)	
10	#define CTL	(1<<1)	
11	#define ALT	(1<<2)	
12			
13	#define CAPSLOCK	(1<<3)	
14	#define NUMLOCK	(1<<4)	
15	#define SCROLLOCK	(1<<5)	
16			
17	#define E0ESC	(1<<6)	
18			
19	// Special keycodes		
20	#define KEY_HOME	0xE0	
21	#define KEY_END	0xE1	
22	#define KEY_UP	0xE2	
23	#define KEY_DN	0xE3	
24	#define KEY_LF	0xE4	
25	#define KEY_RT	0xE5	
26	#define KEY_PGUP	0xE6	
27	#define KEY_PGDN	0xE7	
28	#define KEY_INS	0xE8	
29	#define KEY_DEL	0xE9	

```
160 void  
161 consputc(int c)  
162 {  
163     if(panicked){  
164         cli();  
165         for(;;)  
166             ;  
167     }  
168  
169     if(c == BACKSPACE){  
170         uartputc('\b'); uartputc(' '); uartputc('\b');  
171     } else  
172         uartputc(c);  
173     cgaputc(c);  
174 }
```

```
53 void  
54 cprintf(char *fmt, ...)  
55 {  
56     int i, c, locking;  
57     uint *argp;  
58     char *s;  
59  
60     locking = cons.locking;  
61     if(locking)  
62         acquire(&cons.lock);  
63  
64     if(fmt == 0)  
65         panic("null fmt");  
66  
67     argp = (uint*)(void*)(&fmt + 1);  
68     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){  
69         if(c != '%'){  
70             consputc(c);  
71             continue;  
72         }
```

```
129 static void  
130 cgaputc(int c)  
131 {  
132     int pos;  
133  
134     // Cursor position: col + 80*row.  
135     outb(CRTPORT, 14);  
136     pos = inb(CRTPORT+1) << 8;  
137     outb(CRTPORT, 15);  
138     pos |= inb(CRTPORT+1);  
139  
140     if(c == '\n')  
141         pos += 80 - pos%80;  
142     else if(c == BACKSPACE){  
143         if(pos > 0) --pos;  
144     } else  
145         crt[pos++] = (c&0xff) | 0x0700; // black on white  
146  
147     if((pos/80) >= 24){ // Scroll up.  
148         memmove(crt, crt+80, sizeof(crt[0])*23*80);  
149         pos -= 80;  
150         memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));  
151     }  
152  
153     outb(CRTPORT, 14);  
154     outb(CRTPORT+1, pos>>8);  
155     outb(CRTPORT, 15);  
156     outb(CRTPORT+1, pos);  
157     crt[pos] = ' ' | 0x0700;  
158 }
```

Recall : read

- readi()->bread()
 - Call bget() to read from buffer cache first
 - If not found, call iderw() to read from disk, and mark as valid.

```
95 // Return a B_BUSY buf with the contents of the indicated disk sector.
96 struct buf*
97 bread(uint dev, uint sector)
98 {
99     struct buf *b;
100
101    b = bget(dev, sector);
102    if(!(b->flags & B_VALID))
103        iderw(b);
104    return b;
105 }
```

Xv6 disk driver design & Implementation

- Design Overview
 - Support an IDE driver
 - Use I/O Instructions instead of Memory Mapped I/O
 - in, out
 - Asynchronous I/O model
 - Use a simple **queue** of I/O request
 - Use **interrupt** to notify the available of data

IDE Specification

- Xv6 uses a relative old IDE specification
 - Not PCI negotiation
- IDE Ports
 - 0x1F0-0x1F7
 - 0x1f0: write/read port
 - 0x1f2: number of sectors
 - 0x1f3-0x1f5: sector number
 - 0x1f6: diskno and sector number
 - 0x1f7: command registers, status bit
 - 0x3F6 interrupt control line

I/O instruction in xv6

```
3 static inline uchar
4 inb(ushort port)
5 {
6     uchar data;
7
8     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
9     return data;
10 }
```

```
21 static inline void
22 outb(ushort port, uchar data)
23 {
24     asm volatile("out %0,%1" : : "a" (data), "d" (port));
25 }
```

Polling Mechanism

```
32 // Wait for IDE disk to become ready.  
33 static int  
34 idewait(int checkerr)  
35 {  
36     int r;  
37  
38     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)  
39         ;  
40     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)  
41         return -1;  
42     return 0;  
43 }
```

IDE Initialization

```
45 void
46 ideinit(void)
47 {
48     int i;
49
50     initlock(&idelock, "ide");
51     picenable(IRQ_IDE);
52     ioapicenable(IRQ_IDE, ncpu - 1);
53     idewait(0);
54
55     // Check if disk 1 is present
56     outb(0x1f6, 0xe0 | (1<<4));
57     for(i=0; i<1000; i++){
58         if(inb(0x1f7) != 0){
59             havedisk1 = 1;
60             break;
61         }
62     }
63
64     // Switch back to disk 0.
65     outb(0x1f6, 0xe0 | (0<<4));
66 }
```

```

340 // Atomically release lock and sleep on chan.
341 // Reacquires lock when awakened.
342 void
343 sleep(void *chan, struct spinlock *lk)
344 {
345     if(proc == 0)
346         panic("sleep");
347
348     if(lk == 0)
349         panic("sleep without lk");
350
351     // Must acquire ptable.lock in order to
352     // change p->state and then call sched.
353     // Once we hold ptable.lock, we can be
354     // guaranteed that we won't miss any wakeup
355     // (wakeup runs with ptable.lock locked),
356     // so it's okay to release lk.
357     if(lk != &ptable.lock){ //DOC: sleeplock0
358         acquire(&ptable.lock); //DOC: sleeplock1
359         release(lk);
360     }
361
362     // Go to sleep.
363     proc->chan = chan;
364     proc->state = SLEEPING;
365     sched();
366
367     // Tidy up.
368     proc->chan = 0;
369
370     // Reacquire original lock.
371     if(lk != &ptable.lock){ //DOC: sleeplock2
372         release(&ptable.lock);
373         acquire(lk);
374     }
375 }

```

Sleep and Wakeup Mechanism in xv6

```

378 // Wake up all processes sleeping on chan.
379 // The ptable lock must be held.
380 static void
381 wakeup1(void *chan)
382 {
383     struct proc *p;
384
385     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
386         if(p->state == SLEEPING && p->chan == chan)
387             p->state = RUNNABLE;
388 }
389
390 // Wake up all processes sleeping on chan.
391 void
392 wakeup(void *chan)
393 {
394     acquire(&ptable.lock);
395     wakeup1(chan);
396     release(&ptable.lock);
397 }

```

```

122 // Sync buf with disk.
123 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
124 // Else if B_VALID is not set, read buf from disk, set B_VALID.
125 void
126 iderw(struct buf *b)
127 {
128     struct buf **pp;
129
130     if(!(b->flags & B_BUSY))
131         panic("iderw: buf not busy");
132     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
133         panic("iderw: nothing to do");
134     if(b->dev != 0 && !havedisk1)
135         panic("iderw: ide disk 1 not present");
136
137     acquire(&idelock); //DOC:acquire-lock
138
139     // Append b to idequeue.
140     b->qnext = 0;
141     for(pp=&idequeue; *pp; pp=&(*pp)->qnext) //DOC:insert-queue
142         ;
143     *pp = b;
144
145     // Start disk if necessary.
146     if(idequeue == b)
147         idestart(b);
148
149     // Wait for request to finish.
150     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
151         sleep(b, &idelock);
152     }
153
154     release(&idelock);
155 }

```

Synchronous I/O (iderw)

```
68 // Start the request for b. Caller must hold idelock.
69 static void
70 idestart(struct buf *b)
71 {
72     if(b == 0)
73         panic("idestart");
74
75     idewait(0);
76     outb(0x3f6, 0); // generate interrupt
77     outb(0x1f2, 1); // number of sectors
78     outb(0x1f3, b->sector & 0xff);
79     outb(0x1f4, (b->sector >> 8) & 0xff);
80     outb(0x1f5, (b->sector >> 16) & 0xff);
81     outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((b->sector>>24)&0x0f));
82     if(b->flags & B_DIRTY){
83         outb(0x1f7, IDE_CMD_WRITE);
84         outsl(0x1f0, b->data, 512/4);
85     } else {
86         outb(0x1f7, IDE_CMD_READ);
87     }
88 }
```

```
90 // Interrupt handler.
91 void
92 ideintr(void)
93 {
94     struct buf *b;
95
96     // First queued buffer is the active request.
97     acquire(&idelock);
98     if((b = idequeue) == 0){
99         release(&idelock);
100        // cprintf("spurious IDE interrupt\n");
101        return;
102    }
103    idequeue = b->qnext;
104
105    // Read data if needed.
106    if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
107        insl(0x1f0, b->data, 512/4);
108
109    // Wake process waiting for this buf.
110    b->flags |= B_VALID;
111    b->flags &= ~B_DIRTY;
112    wakeup(b);
113
114    // Start disk on next buf in queue.
115    if(idequeue != 0)
116        idestart(idequeue);
117
118    release(&idelock);
119 }
```

```
12 static inline void
13 insl(int port, void *addr, int cnt)
14 {
15     asm volatile("cld; rep insl" :
16                 "=D" (addr), "=c" (cnt) :
17                 "d" (port), "0" (addr), "1" (cnt) :
18                 "memory", "cc");
19 }
```

IDE.c

- The driver can only handle 1 operation at a time
 - Only send the buffer at the front of the queue
 - Others are simply waiting their turn
- iderw() maintains a queue of requests
 - Adds the buffer b to the end of the queue
 - If b is the first, then call idestart(), if not, just wait
 - Then just sleep, instead of polling
- ideintr() will handle the first request
 - Then start other request if any

XV6, FAT32, NTFS, Ext4

Yubin Xia & Rong Chen

Xv6 FS Design

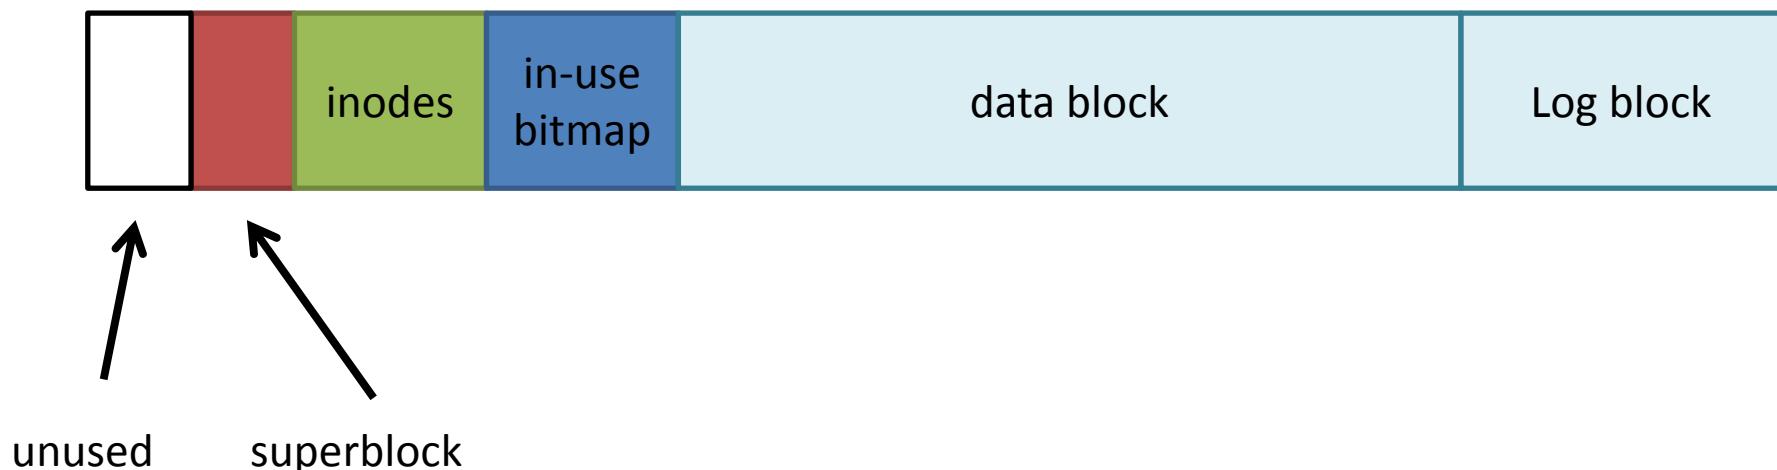
- Implements a minimal Unix file system interface
 - Superblock
 - Inode
 - Dentry
- Limitations
 - not pay attention to file system layout
 - doesn't do any disk scheduling.
 - Its cache is write-through
 - simplifies keeping disk data structures consistent
 - bad for performance.

xv6 File System

- No disk scheduling
- Write-through
 - Simplifies keeping on disk data structures consistent
 - Bad for performance
- Block size: 512 bytes (*BSIZE*)

Disk Layout

- Block 0: unused
- Block 1: superblock
- Block 2+
 - inode ($\text{ninodes} / \text{inodes_per_block}$)
 - in-use bitmap
 - data block
 - Log block



Recap: Free-Space List Bit Vector

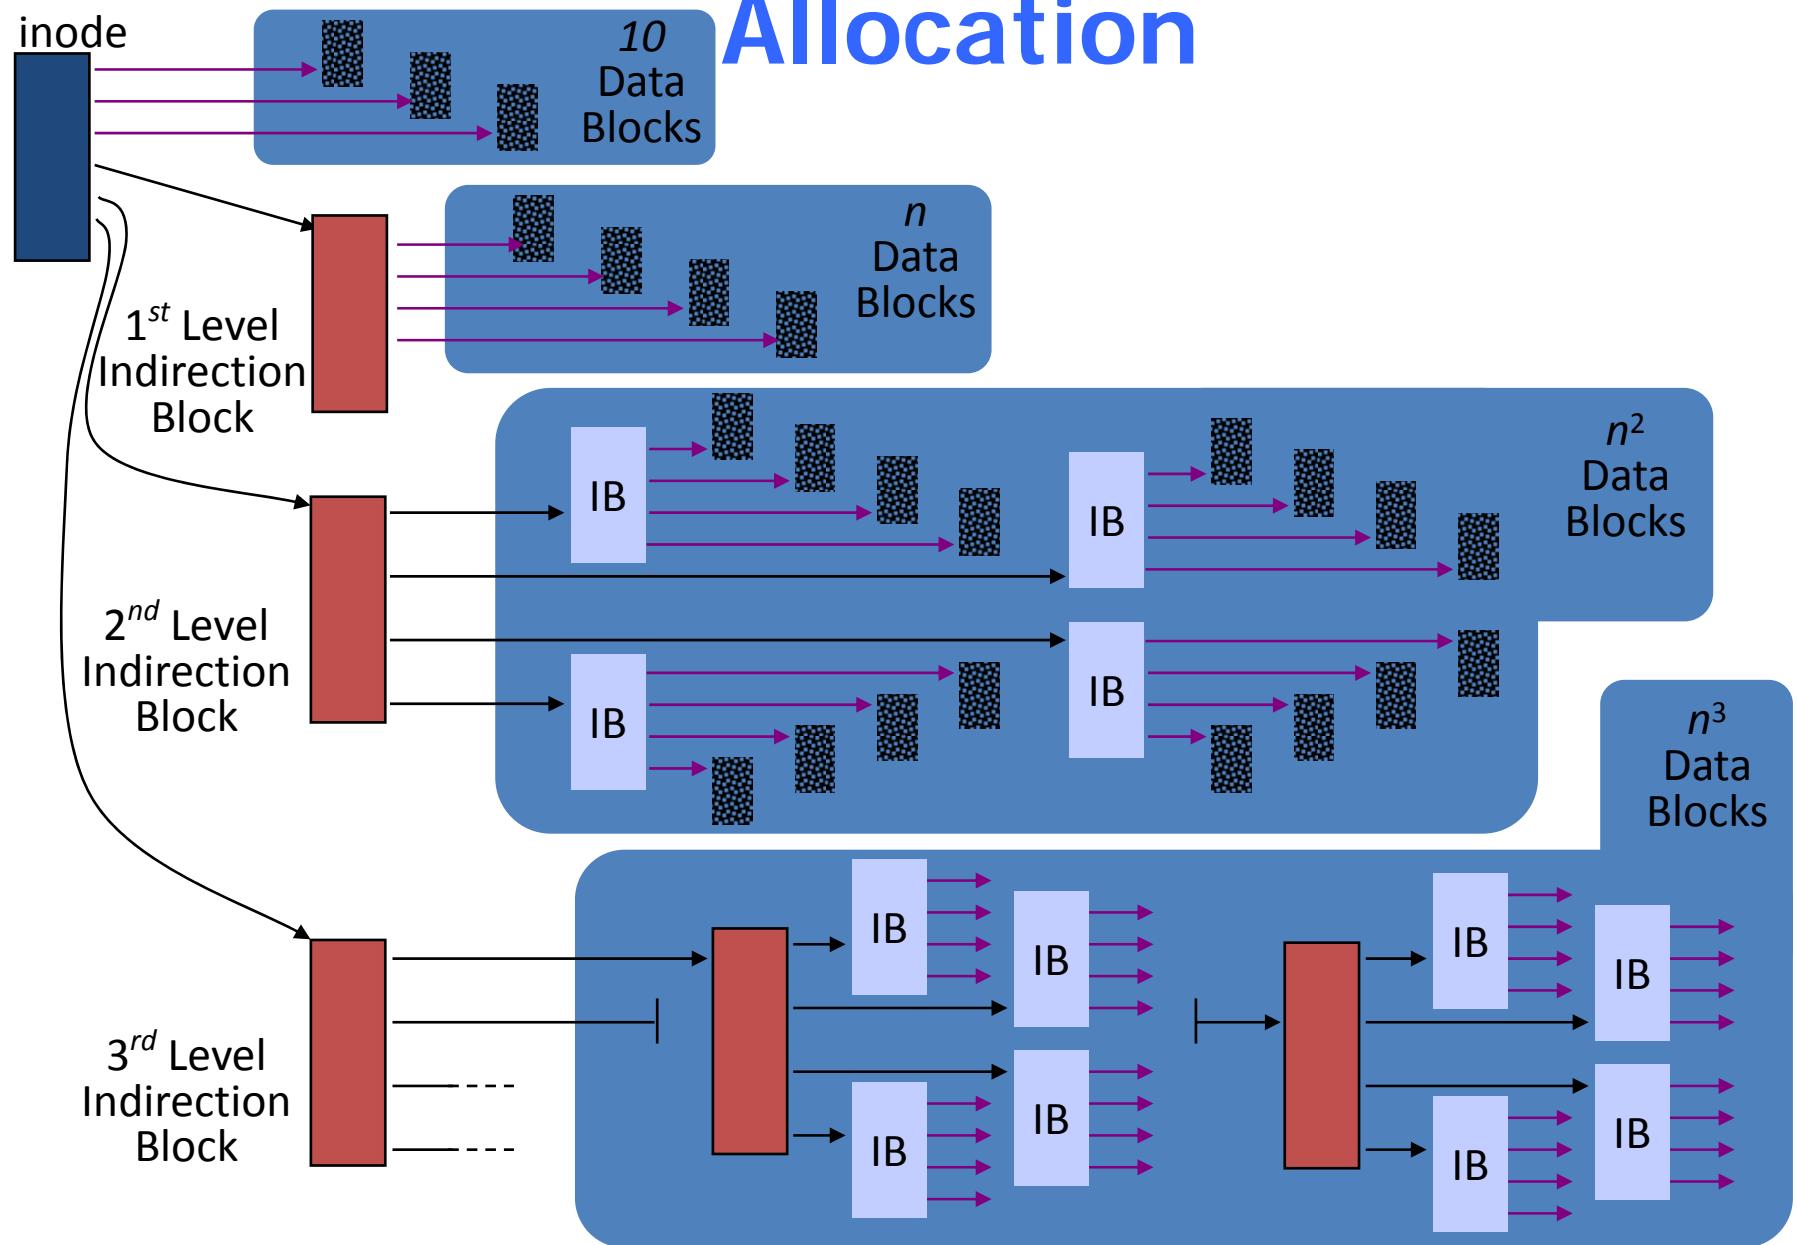
- Represent the list of free blocks as a bit vector:
 - 1111111111111100111010101110111...
 - If bit $i = 0$ then block i is free, otherwise it is allocated
- Simple to use but this can be a big vector:
 - 160GB disk \rightarrow 40M blocks \rightarrow 5MB worth of bits
 - However, if free sectors are uniformly distributed across the disk
 - then the expected number of bits that must be scanned before finding a “0” is n/r , where
 - n = total number of blocks on the disk
 - r = number of free blocks
 - If a disk is 90% full, then the average number of bits to be scanned is 10, independent of the size of the disk

Managing Disk Blocks

- Use a bitmap to maintain availability of disk blocks

```
51 // Allocate a zeroed disk block.
52 static uint
53 balloc(uint dev)
54 {
55     int b, bi, m;
56     struct buf *bp;
57     struct superblock sb;
58
59     bp = 0;
60     readsb(dev, &sb);
61     for(b = 0; b < sb.size; b += BPB){
62         bp = bread(dev, BBLOCK(b, sb.ninodes));
63         for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
64             m = 1 << (bi % 8);
65             if((bp->data[bi/8] & m) == 0){ // Is block free?
66                 bp->data[bi/8] |= m; // Mark block in use.
67                 log_write(bp);
68                 brelse(bp);
69                 bzero(dev, b + bi);
70                 return b + bi;
71             }
72         }
73         brelse(bp);
74     }
75     panic("balloc: out of blocks");
76 }
```

Recap: Multi-level Indexed Allocation



Files

- Represented by an **inode**

```
[fs.h]
#define NADDRS      (NDIRECT+1)
#define NDIRECT     12
#define INDIRECT    12
#define NINDIRECT   (BSIZE / sizeof(uint))
#define MAXFILE     (NDIRECT + NINDIRECT)

// On-disk inode structure
struct dinode {
    short type;          // File type
    short major;         // Major device number (T_DEV only)
    short minor;         // Minor device number (T_DEV only)
    short nlink;         // Number of links to inode in file system
    uint size;           // Size of file (bytes)
    uint addrs[NDIRECT+1]; // Data block addresses
};
```

Files (cont.)

- File size
 - **Small**: up to 12 block addresses ($NADDRS - 1$)
 - **Large**: Last address in the inode as a disk address for a block with 128 disk addresses (512 / 4)
 - At most $12 * 512 + 128 * 512$ bytes

Directories

- Directories are special files
 - Files with a bit of structure to them

```
// Directory is a file containing a sequence of dirent structures.  
#define DIRSIZ 14  
  
struct dirent {  
    ushort inum;  
    char name[DIRSIZ];  
};
```

Allocating an inode

```
173 // Allocate a new inode with the given type on device dev.  
174 // A free inode has a type of zero.  
175 struct inode*  
176 ialloc(uint dev, short type)  
177 {  
178     int inum;  
179     struct buf *bp;  
180     struct dinode *dip;  
181     struct superblock sb;  
182  
183     readsb(dev, &sb);  
184  
185     for(inum = 1; inum < sb.ninodes; inum++){  
186         bp = bread(dev, IBLLOCK(inum));  
187         dip = (struct dinode*)bp->data + inum%IPB;  
188         if(dip->type == 0){ // a free inode  
189             memset(dip, 0, sizeof(*dip));  
190             dip->type = type;  
191             log_write(bp); // mark it allocated on the disk  
192             brelse(bp);  
193             return iguret(dev, inum);  
194         }  
195         brelse(bp);  
196     }  
197     panic("ialloc: no inodes");  
198 }
```

Case Study: read(fd, buf, 100)

- sys_read()->fileread()->readi()
 - 1. If it's a special device file, use predefined operation
 - 2. call bread() to get file buffer

```
66 int
67 sys_read(void)
68 {
69     struct file *f;
70     int n;
71     char *p;
72
73     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
74         return -1;
75     return fileread(f, p, n);
76 }
```

```
94 // Read from file f.
95 int
96 fileread(struct file *f, char *addr, int n)
97 {
98     int r;
99
100    if(f->readable == 0)
101        return -1;
102    if(f->type == FD_PIPE)
103        return piperead(f->pipe, addr, n);
104    if(f->type == FD_INODE){
105        ilock(f->ip);
106        if((r = readi(f->ip, addr, f->off, n)) > 0)
107            f->off += r;
108        iunlock(f->ip);
109        return r;
110    }
111    panic("fileread");
112 }
```

```
433 // Read data from inode.
434 int
435 readi(struct inode *ip, char *dst, uint off, uint n)
436 {
437     uint tot, m;
438     struct buf *bp;
439
440     if(ip->type == T_DEV){
441         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
442             return -1;
443         return devsw[ip->major].read(ip, dst, n);
444     }
445
446     if(off > ip->size || off + n < off)
447         return -1;
448     if(off + n > ip->size)
449         n = ip->size - off;
450
451     for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
452         bp = bread(ip->dev, bmap(ip, off/BSIZE));
453         m = min(n - tot, BSIZE - off%BSIZE);
454         memmove(dst, bp->data + off%BSIZE, m);
455         brelse(bp);
456     }
457     return n;
458 }
```

Case Study: read (cont.)

- readi()->bread()
 - **Call bget() to read from buffer cache first**
 - If not found, call ide_rw() to read from disk, and mark as valid.

```
95 // Return a B_BUSY buf with the contents of the indicated disk sector.
96 struct buf*
97 bread(uint dev, uint sector)
98 {
99     struct buf *b;
100
101     b = bget(dev, sector);
102     if(!(b->flags & B_VALID))
103         iderw(b);
104     return b;
105 }
```

Case Study: read (cont.)

- `readi()->bread()->bget()`
 - Implements a simple buffer cache of recently-read disk blocks
 - 1. Look if the requested block is in the **cache**
 - 2. If some other process has locked the block, **wait** until it releases
 - 3. If it is not in the cache, find a cache entry to hold the block
 - 4. Mark it ours, but not valid

```
639 struct inode*
640 namei(char *path)
641 {
642     char name[DIRSIZ];
643     return namex(path, 0, name);
644 }
645
646 struct inode*
647 nameiparent(char *path, char *name)
648 {
649     return namex(path, 1, name);
650 }
```

```
601 // Look up and return the inode for a path name.  
602 // If parent != 0, return the inode for the parent and copy the final  
603 // path element into name, which must have room for DIRSIZ bytes.  
604 static struct inode*  
605 namex(char *path, int nameiparent, char *name)  
606 {  
607     struct inode *ip, *next;  
608  
609     if(*path == '/')  
610         ip = igit(ROOTDEV, ROOTINO);  
611     else  
612         ip = idup(proc->cwd);  
613  
614     while((path = skipelem(path, name)) != 0){  
615         ilock(ip);  
616         if(ip->type != T_DIR){  
617             iunlockput(ip);  
618             return 0;  
619         }  
620         if(nameiparent && *path == '\0'){  
621             // Stop one level early.  
622             iunlock(ip);  
623             return ip;  
624         }  
625         if((next = dirlookup(ip, name, 0)) == 0){  
626             iunlockput(ip);  
627             return 0;  
628         }  
629         iunlockput(ip);  
630         ip = next;  
631     }  
632     if(nameiparent){  
633         iput(ip);  
634         return 0;  
635     }  
636     return ip;  
637 }
```

```
570 // Examples:  
571 //    skipelem("a/bb/c", name) = "bb/c", setting name = "a"  
572 //    skipelem("//a//bb", name) = "bb", setting name = "a"  
573 //    skipelem("a", name) = "", setting name = "a"  
574 //    skipelem("", name) = skipelem("////", name) = 0
```

```

312 // Drop a reference to an in-memory inode.
313 // If that was the last reference, the inode cache entry can
314 // be recycled.
315 // If that was the last reference and the inode has no links
316 // to it, free the inode (and its content) on disk.
317 void
318 iput(struct inode *ip)
319 {
320     acquire(&icache.lock);
321     if(ip->ref == 1 && (ip->flags & I_VALID) && ip->nlink == 0){
322         // inode has no links: truncate and free inode.
323         if(ip->flags & I_BUSY)
324             panic("iput busy");
325         ip->flags |= I_BUSY;
326         release(&icache.lock);
327         itrunc(ip);
328         ip->type = 0;
329         iupdate(ip);
330         acquire(&icache.lock);
331         ip->flags = 0;
332         wakeup(ip);
333     }
334     ip->ref--;
335     release(&icache.lock);
336 }

```

```

219 // Find the inode with number inum on device dev
220 // and return the in-memory copy. Does not lock
221 // the inode and does not read it from disk.
222 static struct inode*
223 iget(uint dev, uint inum)
224 {
225     struct inode *ip, *empty;
226     acquire(&icache.lock);
227     // Is the inode already cached?
228     empty = 0;
229     for(ip = &icache.inode[0]; ip < &icache.inode[NINODE]; ip++){
230         if(ip->ref > 0 && ip->dev == dev && ip->inum == inum){
231             ip->ref++;
232             release(&icache.lock);
233             return ip;
234         }
235     }
236     if(empty == 0 && ip->ref == 0) // Remember empty slot.
237         empty = ip;
238     }
239     // Recycle an inode cache entry.
240     if(empty == 0)
241         panic("iget: no inodes");
242     ip = empty;
243     ip->dev = dev;
244     ip->inum = inum;
245     ip->ref = 1;
246     ip->flags = 0;
247     release(&icache.lock);
248     return ip;
249 }

```

```
503 // Look for a directory entry in a directory.
504 // If found, set *poff to byte offset of entry.
505 struct inode*
506 dirlookup(struct inode *dp, char *name, uint *poff)
507 {
508     uint off, inum;
509     struct dirent de;
510
511     if(dp->type != T_DIR)
512         panic("dirlookup not DIR");
513
514     for(off = 0; off < dp->size; off += sizeof(de)){
515         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
516             panic("dirlink read");
517         if(de.inum == 0)
518             continue;
519         if(namecmp(name, de.name) == 0){
520             // entry matches path element
521             if(poff)
522                 *poff = off;
523             inum = de.inum;
524             return igit(dp->dev, inum);
525         }
526     }
527
528     return 0;
529 }
```

EXT4 FILE SYSTEM

File-size upper limits for data block addressing

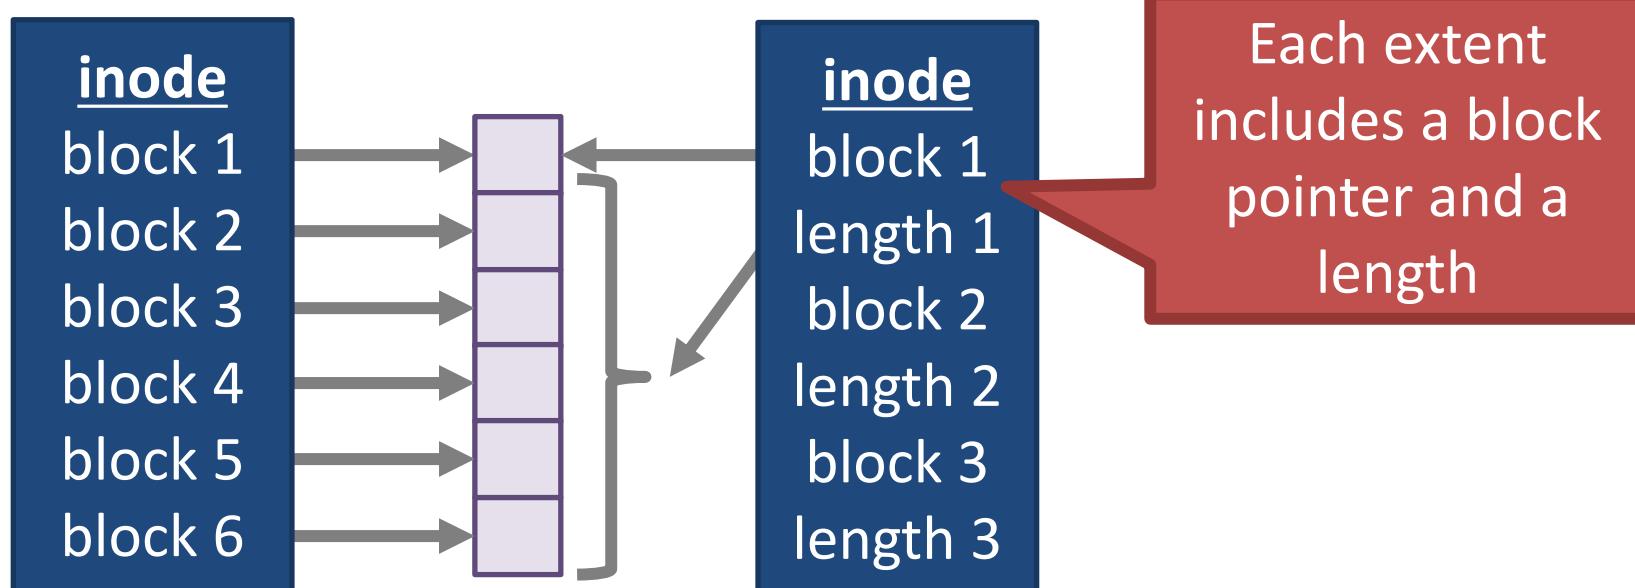
Block size	Direct	1-Indirect	2-Indirect	3-Indirect
1,024	12 KB	268 KB	64.26 MB	16.06 GB
2,048	24 KB	1.02 MB	513.02 MB	256.5 GB
4,096	48 KB	4.04 MB	4 GB	~ 4 TB

Extent

- Indirect block maps are incredibly inefficient for large files
 - One extra block read (and seek) every 1024 blocks
 - Really obvious when deleting big CD/DVD image files
- An extent is a single descriptor for a range of contiguous blocks
 - an efficient way to represent large file
 - Better CPU utilization, fewer metadata IOs

From Pointers to Extents

- Modern file systems try hard to minimize fragmentation
 - Since it results in many seeks, thus low performance
- **Extents** are better suited for contiguous files



Implementing Extents

- ext4 and NTFS use extents
- ext4 inodes include 4 extents instead of block pointers
 - Each extent can address at most 128MB of contiguous space (assuming 4KB blocks)
 - If more extents are needed, a data block is allocated
 - Similar to a block of indirect pointers

Revisiting Directories

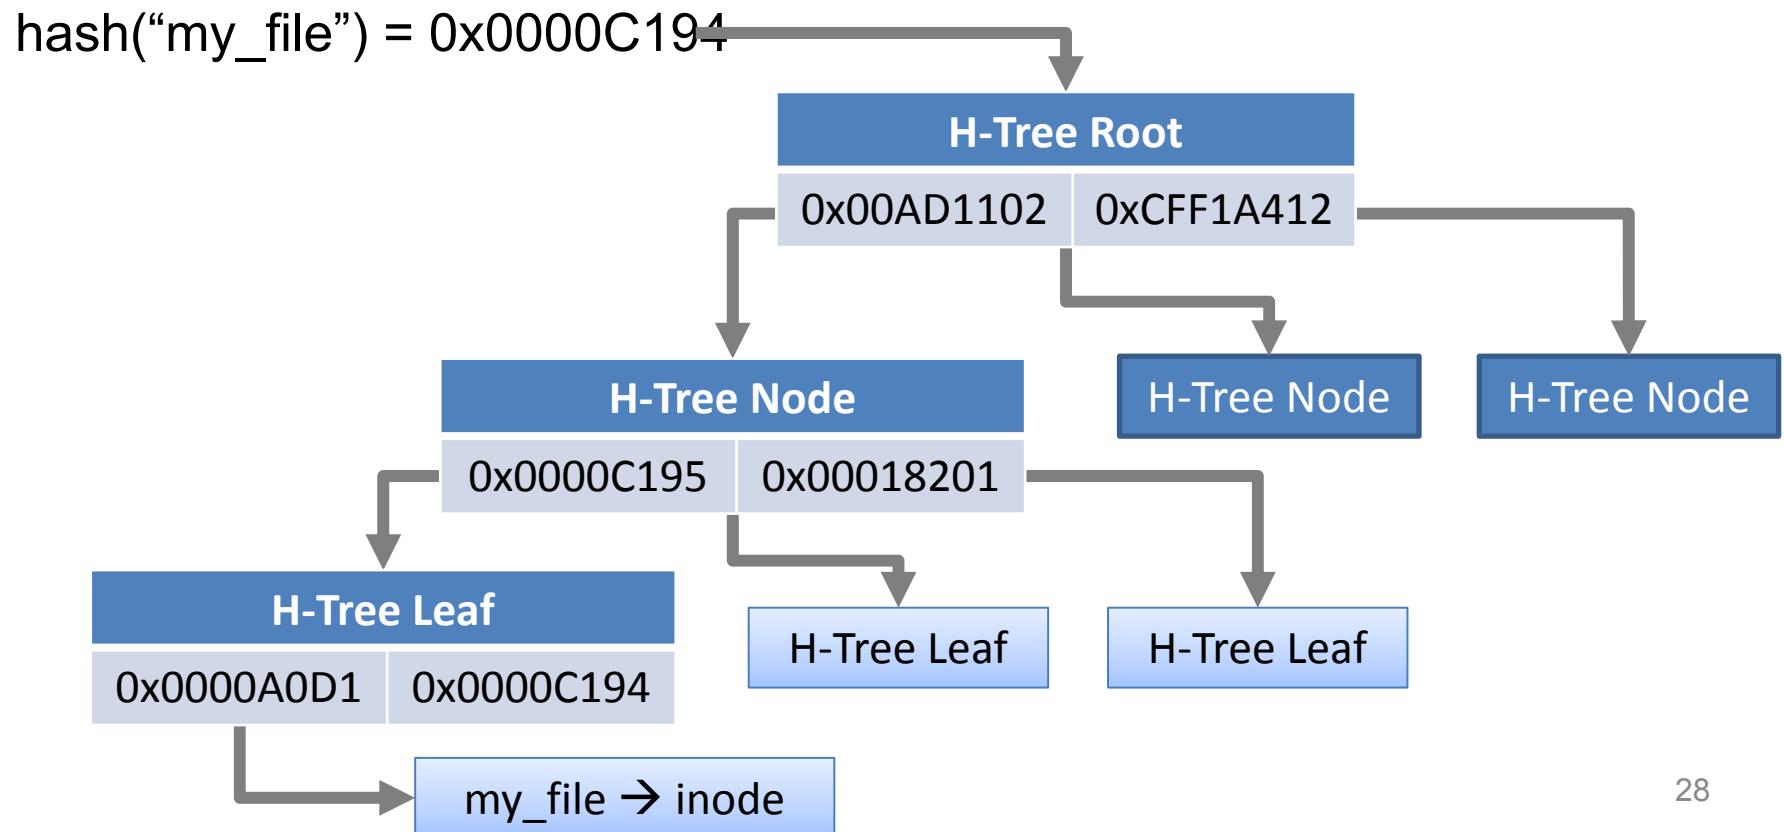
- In ext, ext2, and ext3, each directory is a file with a list of entries
 - Entries are not stored in sorted order
 - Some entries may be blank, if they have been deleted
- Problem: searching for files in large directories takes $O(n)$ time
 - Practically, you can't store $>10K$ files in a directory
 - It takes way too long to locate and open files

From Lists to B-Trees

- ext4 and NTFS encode directories as B-Trees to improve lookup time to $O(\log N)$
- A B-Tree is a type of balanced tree that is optimized for storage on disk
 - Items are stored in sorted order in blocks
- Suppose items i and j are in the root of the tree
 - The root must have 3 children, since it has 2 items
 - The three child groups contain items $a < i$, $i < a < j$, and $a > j$

Example B-Tree

- ext4 uses a B-Tree variant known as a H-Tree
 - The *H* stands for *hash* (sometime called B+Tree)
- Suppose you try to `open("my_file", "r")`



ext4: The Good and the Bad

- The good – ext4 (and NTFS) supports:
 - All of the basic file system functionality we require
 - Improved performance from ext3's block groups
 - Additional performance gains from extents and B-Tree directory files
- The bad:
 - Next-gen file systems have even nicer features
 - Copy-on-write semantics (btrfs and ZFS)

FAT FILE SYSTEM

Clusters and Sectors

- Sector: smallest storage unit on disk
 - 512 bytes
- Cluster: smallest allocated disk space to hold file
 - Data clusters are located after metadata of partition
 - Different cluster sizes depending on volume size

Volume Size	FAT32 Cluster Size
<32MB	Not Supported
32MB ~ 64MB	512 bytes
65MB ~ 128MB	1KB
129MB ~ 256MB	2KB
257MB ~ 8GB	4KB
8GB ~ 16GB	8KB
16GB ~ 32GB	16KB

Default FAT Cluster Sizes

Organization of an FAT Volume

Boot Sector	Reserved Sectors	FAT 1	FAT 2 (Duplicate)	Root Folder	Other Folders and All Files
-------------	------------------	-------	-------------------	-------------	-----------------------------

Boot Sector

- Layout of the volume
- File system structure
- Boot code

Root Folder

- Describes the files and folders in the root of the partition

FAT 1

- Original FAT

FAT 2 (Duplicate)

- Backup copy of FAT

Other Folder and All Files

- Contains the data for the files and folders within the file system.

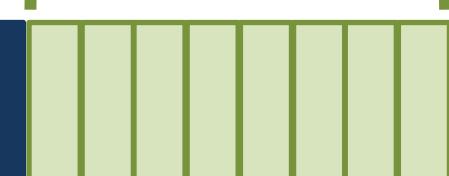
- Stores basic info about the file system
- FAT version, location of boot files
- Total number of blocks
- Index of the root directory in the FAT

- File allocation table (FAT)
- Marks which blocks are free or in-use
- **Linked-list structure** to manage large files

- Store file and directory data
- Each block is a fixed size (4KB – 64KB)
- Files may span multiple blocks

Disk

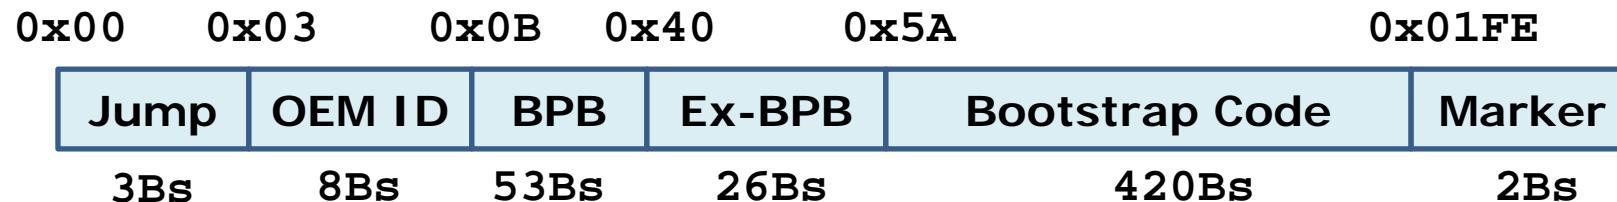
Super Block



FAT Boot Sector

- FAT Boot Sector
 - Locate at the first logical sector of each partition
 - Created when you format a volume
 - End with a 2-byte sector marker (always 0x55AA)
- Component
 - An x86-based CPU jump instruction
 - Original Equipment Manufacturer Identification (OEM ID)
 - BIOS Parameter Block (BPB)
 - Extended BPB
 - Executable boot code

Component of Boot Sector



- Jump: skip the next several non-executable bytes
- OEM ID: a string of characters
 - the name and version number of OS that formatted the volumes
- BPB & Ex-BPB: information of volume
 - Byte Per Sector, Sectors Per Cluster, Number of FAT, Root Entries, Media Descriptor, Sectors Per Fat,
- Bootstrap Code: executable code to start the OS
- Marker: end of boot sector

FAT 1 and FAT 2

- File Allocation Table
 - identifies all clusters as
 - unused, cluster in use by a file, bad cluster, last cluster in a file
 - FAT 2 is used to consistency-checking program

FAT Root Folder

- FAT Root Folder Structure

Root Folder Entry	Size	Description
Name	11B	Name in 8.3 format
Attribute Byte	1B	Information about entry (e.g. archive, system, hidden, and read-only)
Creation Time and Date	5B	Time and date file was create
Last Access Date	2B	Date file was last accessed
Last Modification Time and Date	4B	Time and data file was last modified
First Cluster	2B	Starting cluster number in the file allocation table
File Size	4B	Size of the file

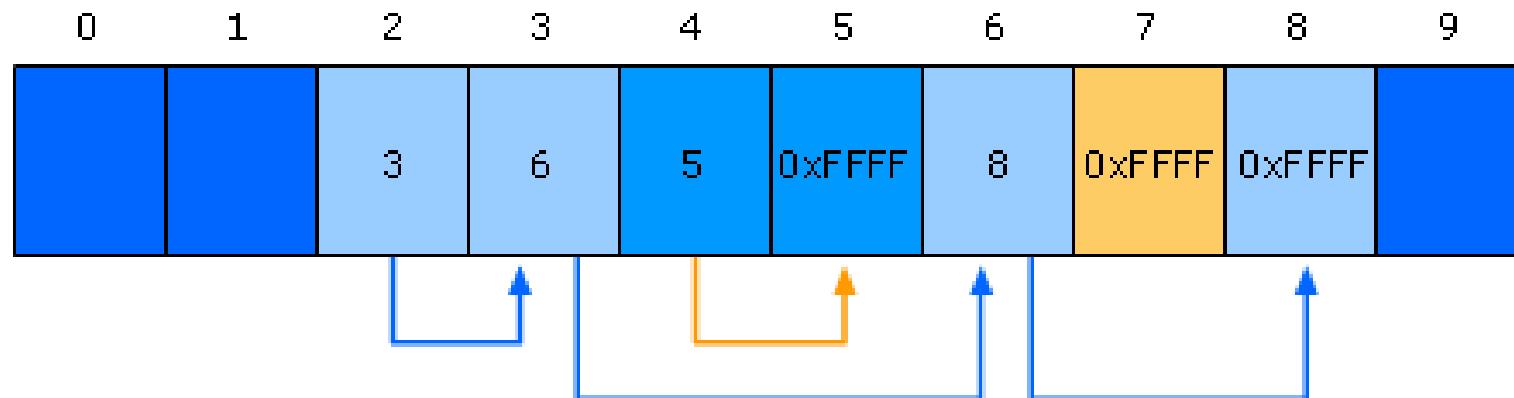
- **FAT File System**
 - FAT Architecture
 - FAT Physical Structure
 - FAT Processes and Interactions

Startup Process

- BIOS and CPU initiate the power-on self test
- BIOS find the boot device
- BIOS loads the first physical sector (**MBR**)
- Transfer CPU execution to MBR
- MBR scans the partition table
- Load a copy of the **Boot Sector** for active partition
- Transfers CPU execution to Boot Sector

File Processing on FAT Clusters

- Store file information
 - Continues to store file info in the next available cluster
 - when the file requires space greater than the cluster's size.
 - e.g.
 - Three Files: 1 with (2-3-6-8), 2 with(4-5) 3 with (7)

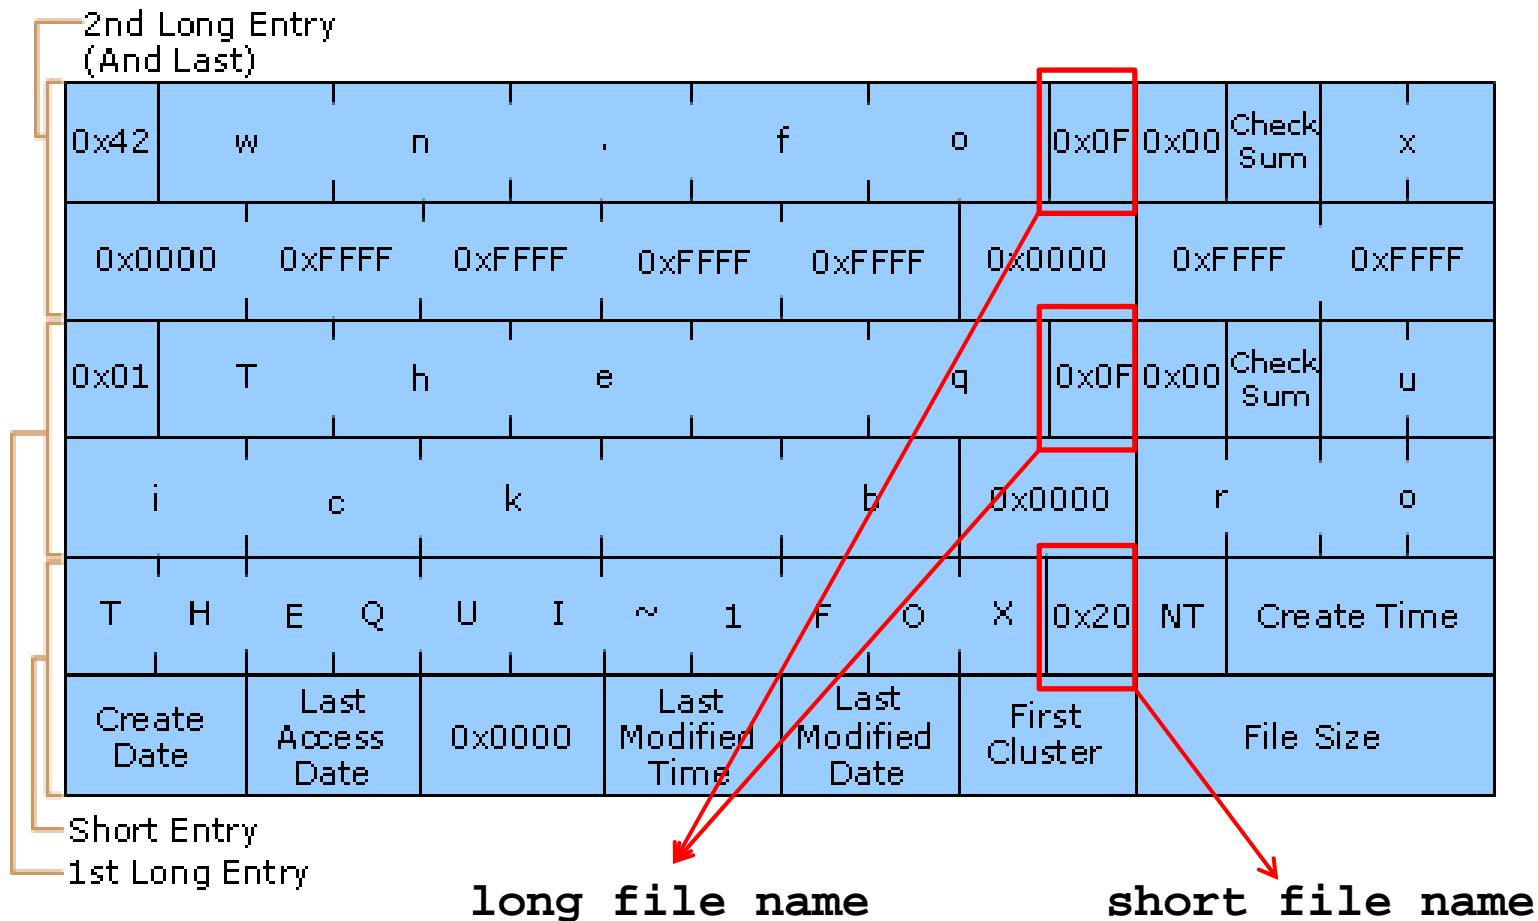


File Naming

- File Names in Windows Server 2003
 - Support long file name, up to 255 characters
 - Generate a 8.3 short file name
- Support Long File Name
 - The main folder entry stores 8.3 short file name
 - The secondary folder entries store long file name
 - Each stores 13 characters in Unicode

File Naming (cont.)

e.g. file with “The quick brown.fox” and “Thequi~1.fox”



FAT: The Good and the Bad

- The Good – FAT supports:
 - Hierarchical tree of directories and files
 - Variable length files
 - Basic file and directory meta-data
- The Bad
 - At most, FAT32 supports 2TB disks
 - Locating free chunks requires scanning the entire FAT
 - Prone to internal and external fragmentation
 - Large blocks → internal fragmentation
 - **Reads require a lot of random seeking**

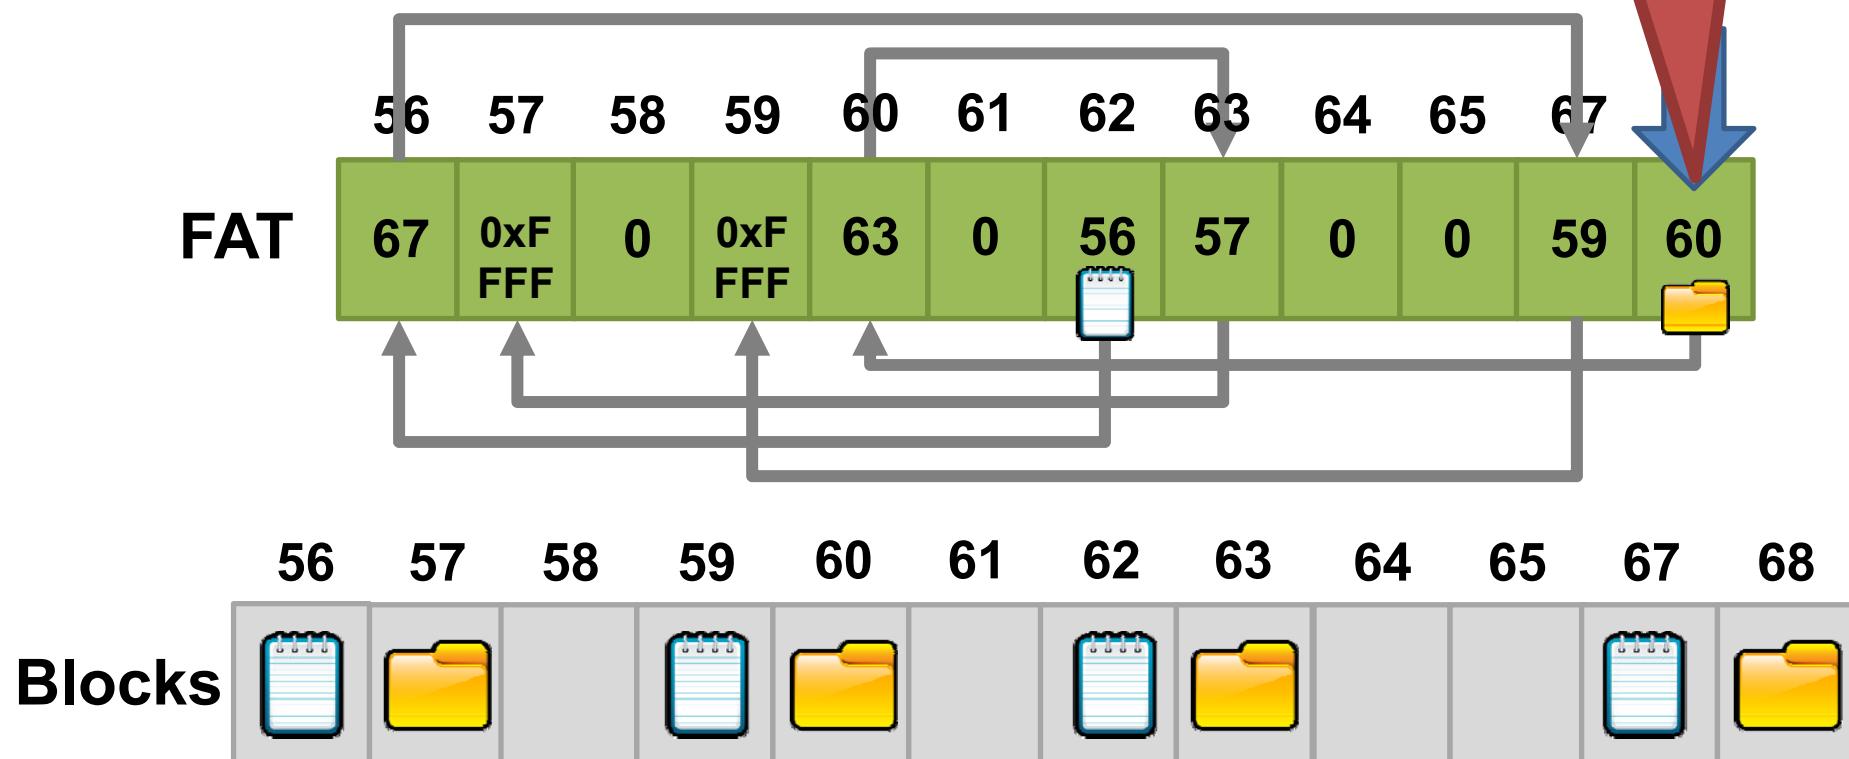
Lots of Seek!

- Consider the following code:

```
int fd = open("my_file.txt", "r");
```

```
int r = read(fd, buffer, 1024 * 4 * 4); // 4 4KB blocks
```

FAT may have very low spatial locality, thus a lot of random seeking

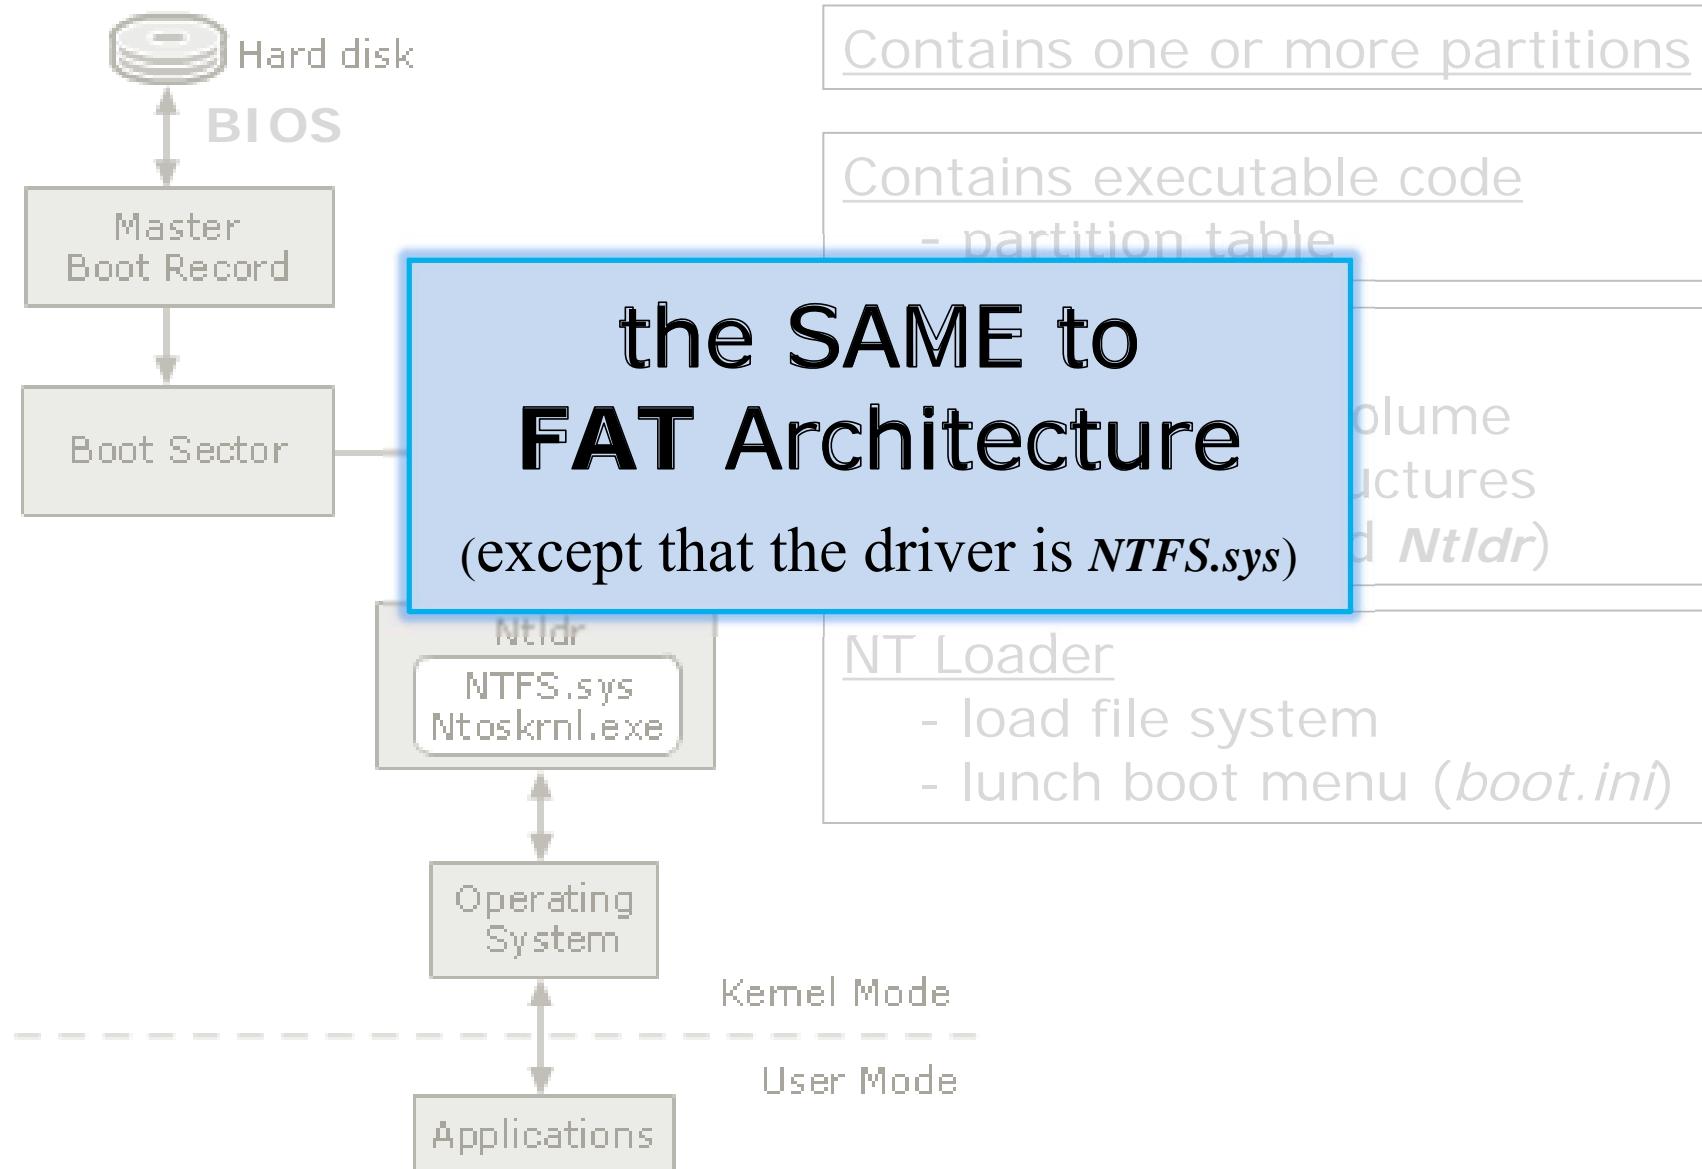


NTFS

NTFS File System

- NTFS Architecture
- NTFS Physical Structure
- NTFS Processes and Interactions

NTFS Architecture



- NTFS File System
 - NTFS Architecture
 - NTFS Physical Structure
 - NTFS Processes and Interactions

NTFS Clusters

- Cluster: smallest allocated disk space to hold file
 - Sequentially logical cluster numbers from the beginning of the partition
 - Clusters start at sector ZERO (*different to FAT*)
 - Floppy disks don't use NTFS
 - Different cluster sizes depending on volume size

Volume Size	NTFS Cluster Size
7MB ~ 512MB	512 bytes
513MB ~ 1GB	1KB
1GB ~ 2GB	2KB
2GB ~ 2TB	4KB

Default NTFS Cluster Sizes

Organization of an NTFS Volume

NTFS Boot Sector	Master File Table	File System Data	Master File Table Copy
------------------	-------------------	------------------	------------------------

NTFS Boot Sector

- Layout of the volume
- File system structure
- Boot code

File System Data

- Data not contained within MFT

Master File Table

- Attributes of files
- Attributes of folders

(Most Important)

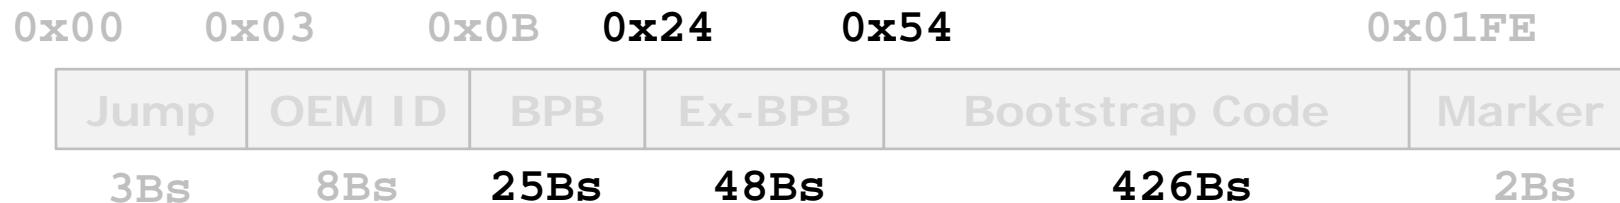
Master File Table Copy

- a copy of MFT

NTFS Boot Sector

- NTFS Boot Sector
 - locate at the first logical sector of each partition
 - Created when you format a volume
 - End with a 2-byte sector marker (always 0x55AA)
 - Components
 - An x86-based boot loader
 - Original Equipment Manufacturer Identification (OEM ID)
 - BIOS Parameter Block (PBP)
 - Extended BPB
 - Executable boot code
- the SAME to
FAT File System**

Component Boot Sector



- Jump: skip the next several non-executable bytes
- OEM ID: a string of 4 bytes
 - the name of the company that formatted the volume
 - (except the size of components)
- BPB & Ex-BPB: information of volume
 - Byte Per Sector, Sectors Per Cluster, Media Descriptor, Total Sectors, Logical Cluster Number of the “MFT”,
- Bootstrap Code: executable code to start the OS
- Marker: end of boot sector

**the SAME to
FAT File System**

Master File Table

- Master File Table (MFT)
 - A relational database
 - Rows: file records
 - Columns: file attributes
 - Contains at least one entry for every file
- Metadata files
 - Describe the MFT
 - The first 16 records of MFT for metadata files
 - Begin with a dollar sign. e.g. \$Mft, \$MftMirr, \$LogFile, ...

Metadata Files in MFT

System File	Name	Rec	Purpose of File
Master File Table	\$Mft	0	Contains file record for each file and folder on a NTFS volume.
Master File Table Mirror	\$MftMirr	1	Duplicate image of the first four records of MFT
Log File	\$LogFile	2	Restores metadata for fast recovery
Volume File	\$Volume	3	Contains information about the volume
Attribute Definitions	\$AttrDef	4	Lists attribute names, numbers, and descriptions
Root File Name Index	.	5	The root file
Cluster Bitmap	\$Bitmap	6	Represents the volume by showing free and unused clusters.

Metadata Files in MFT (cont.)

System File	Name	Rec	Purpose of File
Boot Sector	\$Boot	7	Includes the BPB and additional bootstrap loader code
Bad Cluster file	\$BadClus	8	Contains bad clusters for a volume
Security File	\$Secure	9	Contains unique security descriptors for all files
Upcase Table	\$Upcase	10	Used to match Unicode uppercase characters
NTFS Extension File	\$Extend	11	Optional extensions such as quotas, reparse point data ...
		12-15	Reserved for future use

MFT Zone

- MFT contains a record for each file and folder on the volume
 - File and folder records are 1KB each
- MFT Zone
 - Exclusive use of the MFT
 - Prevent fragmentation
 - Dedicate 12.5% of volume by default
 - On demand allocation for additional MFT zone

NTFS File Record Attributes

- File Attributes
 - View each file and folder as a set of file attributes
 - Resident Attributes vs. Nonresident Attributes
 - Small files and folders are entirely contained within the file's MFT record (*typically, less than 900 bytes*)
 - Types:
 - Standard Info: access mode, timestamp, link count ...
 - Attribute List: locations of all nonresident attributes
 - File Name: long and short file names
 - Data / Index: data of file or index of folder
 - Object ID: volume-unique file identifier
 - ...

NTFS File Record Attributes

- Index of Folder
 - Folder records contain index information (directory entry)
 - Organized B-tree structure for large folders
- Last Access Time
 - The most up-to-date LAT is always stored in memory
 - Written to disk
 - WHERE: *file's attribute AND directory entry for the file*
 - WHEN: *more than an hour OR update other file attributes*

- NTFS File System
 - NTFS Architecture
 - NTFS Physical Structure
 - NTFS Processes and Interactions

Startup Process

- BIOS and CPU initiate the power-on self test
- BIOS find the boot device
- BIOS loads the first physical sector (**MBR**)
- Transfer CPU control to the MBR
- MBR scans all partitions for the active partition
- Load a copy of the **Boot Sector** for active partition
- Transfers CPU execution to Boot Sector

**the SAME to
FAT File System**

File Naming

- File Names in Windows Server 2003
 - Support long file name, up to 255 characters
 - NTFS generate a 8.3 short file name
- HOW to generate short file name
 - Deletes all of unsupported characters (e.g. space)
 - Deletes all extra periods (*abc.def.txt*)
 - Truncates the file name to six characters and appends a tilde(~) and a number
 - Translates all characters in file name and extension to uppercase

“**This is test 1.txt**” -> “**THISIS~1.TXT**”

Compression of Files and Folders

- NTFS Compression
 - Implemented within NTFS
 - Compression only in disk
 - Uncompress before moving data to memory
- Moving and Copying Files or Folders
 - Change compression state
 - Add overhead to the system

Compression of Files and Folders

- Moving Files or Folders within an NTFS Volume
 - keeps its compression state regardless of the compression state of the folder it is moved
- Copying Files or Folders within an NTFS Volume
 - takes on the compression state of the folder it is copied to.
- Between FAT and NTFS Volumes
 - FAT -> NTFS: takes on the compression attribute of the target folder
 - NTFS -> FAT: uncompressed

Hard Links

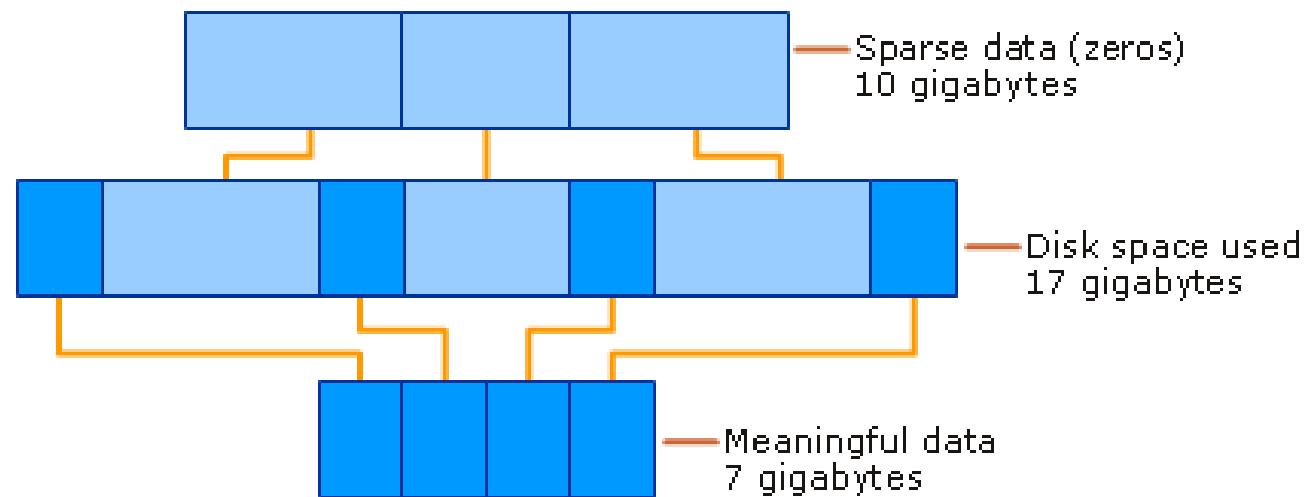
- Hard Links
 - Add a directory entry for the hard link without duplicating the original file
 - Application can modify a file by using any of its hard Links
 - Can't give a file different security descriptors on a per-hard-link basis

Sparse Files

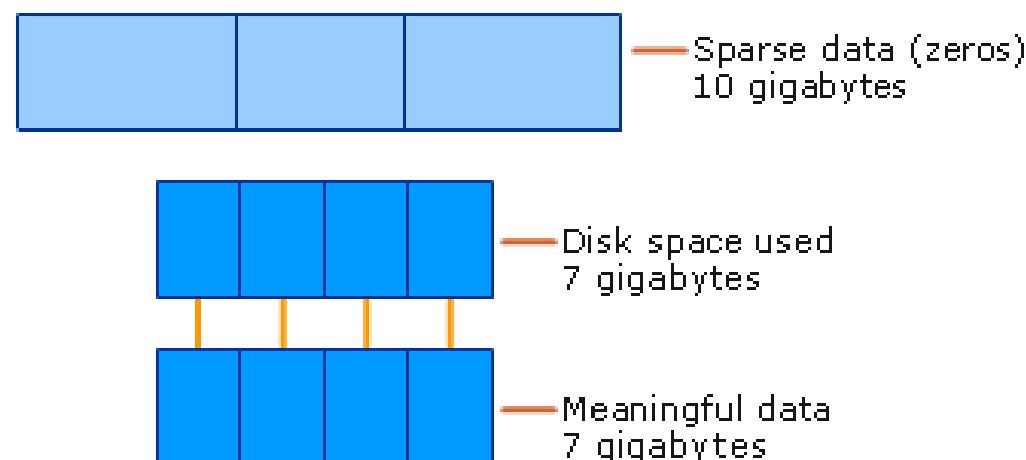
- Sparse Files
 - Focus on the file contain large sections of data composed of zeros
 - Store
 - Allocates disk clusters only for the data explicitly specified by the application
 - Non-specified ranges of the file are represented by non-allocated space on the disk
 - Read
 - From allocated ranges, return the data stored
 - From non-allocated ranges, return ZERO

Sparse Files (cont.)

Without Sparse File Attribute Set



With Sparse File Attribute Set



File System Recoverability

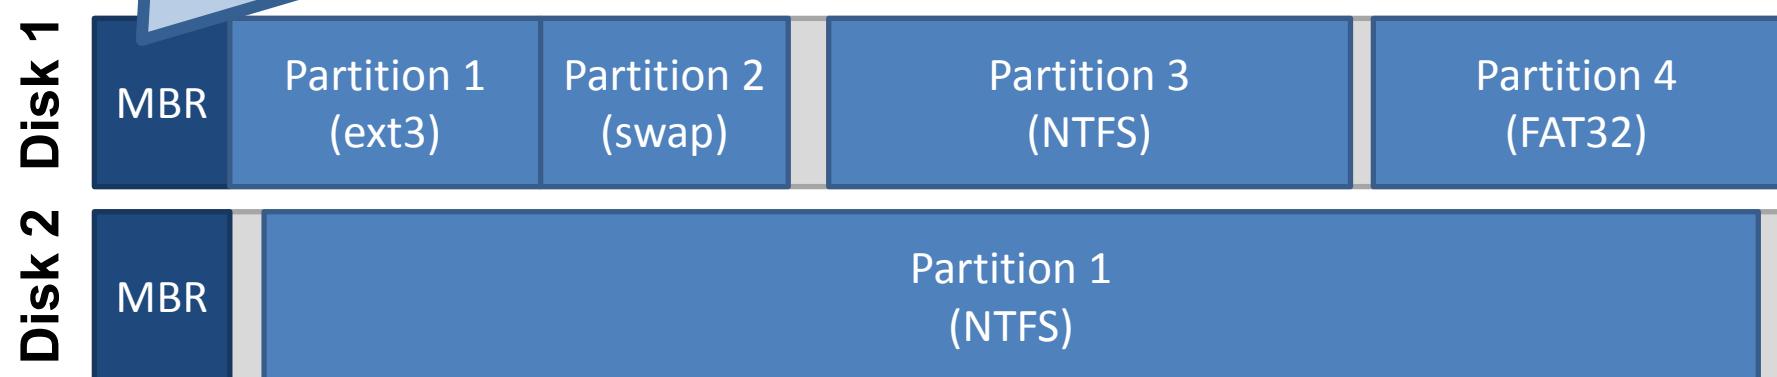
- Recovering NTFS File Structures
 - View each operation as a transaction and manages each one as an integral unit
 - Each transaction
 - Records the metadata operations in a log file cached in memory
 - Records the actual metadata operations in memory
 - Marks the transaction in the cached log file as committed
 - Flushes the [log file](#) to disk
 - Flushes the actual [metadata](#) operations to disk

MBR & MOUNT

The Master Boot Record

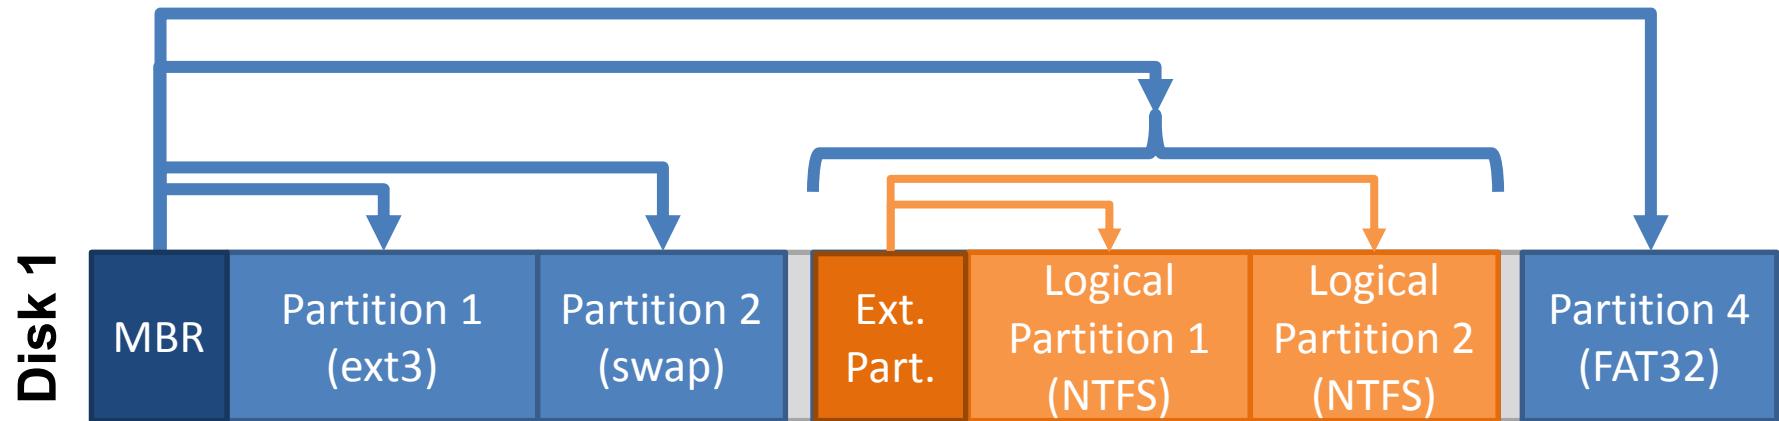
Address		Description	Size (Bytes)
Hex	Dec.		
0x000	0	Bootstrap code area	446
0x1BE	446	Partition Entry #1	16
0x1CE	462	Partition Entry #2	16
0x1DE	478	Partition Entry #3	16
0x1EE	494	Partition Entry #4	16
0x1FE	510	Magic Number	2
		Total:	512

Includes the starting LBA and length of the partition



Extended Partitions

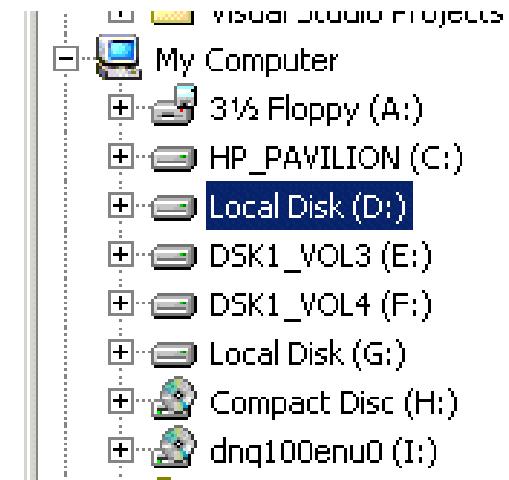
- In some cases, you may want >4 partitions
- Modern OSes support extended partitions



- Extended partitions may use OS-specific partition table formats (meta-data)
 - Thus, other OSes may not be able to read the logical partitions

Types of Root File Systems

- Windows exposes a multi-rooted system
 - Each device and partition is assigned a letter
 - Internally, a single root is maintained
- Linux has a single root
 - One partition is mounted as /
 - All other partitions are mounted somewhere under /
- Typically, the partition containing the kernel is mounted as / or C:



Mounting a File System

1. Read the **super block** for the target file system
 - Contains meta-data about the file system
 - Version, size, locations of key structures on disk, etc.
1. Determine the **mount point**
 - On Windows: pick a drive letter
 - On Linux: mount the new file system under a specific directory

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda5	127G	86G	42G	68%	/media/cbw/Data
/dev/sda4	61G	34G	27G	57%	/media/cbw/Windows
/dev/sdb1	1.9G	352K	1.9G	1%	/media/cbw/NDSS-2013

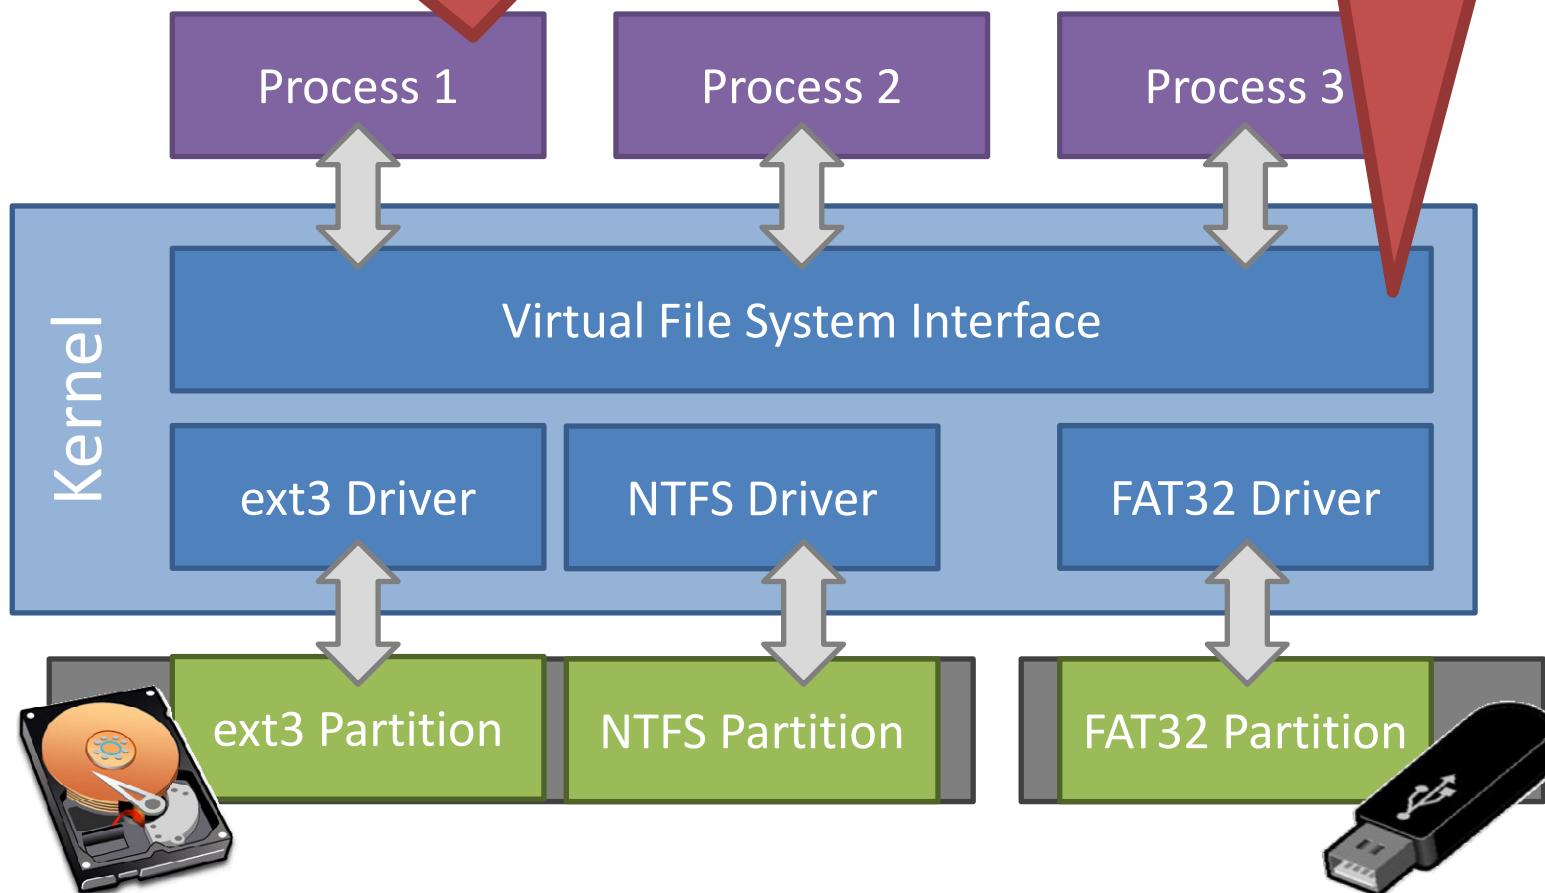
Virtual File System Interface

- Problem: the OS may mount several partitions containing different underlying file systems
 - It would be bad if processes had to use different APIs for different file systems
- Linux uses a Virtual File System interface (VFS)
 - Exposes POSIX APIs to processes
 - Forwards requests to lower-level file system specific drivers
- Windows uses a similar system

VFS Flowchart

Processes (usually) don't need to know about low-level file system details

Relatively simple to add additional file system drivers



Mount isn't Just for Bootup

- When you plug storage devices into your running system, mount is executed in the background
- Example: plugging in a USB stick
- What does it mean to “safely eject” a device?
 - Flush cached writes to that device
 - Cleanly unmount the file system on that device



Homework

- Please learn about **btrfs**, use your own words to describe its system architecture, and states its pros and cons (you can compare it with ext4/fat/ntfs)
- (optional) you can also check ReiserFS and ZFS, they also have very interesting features

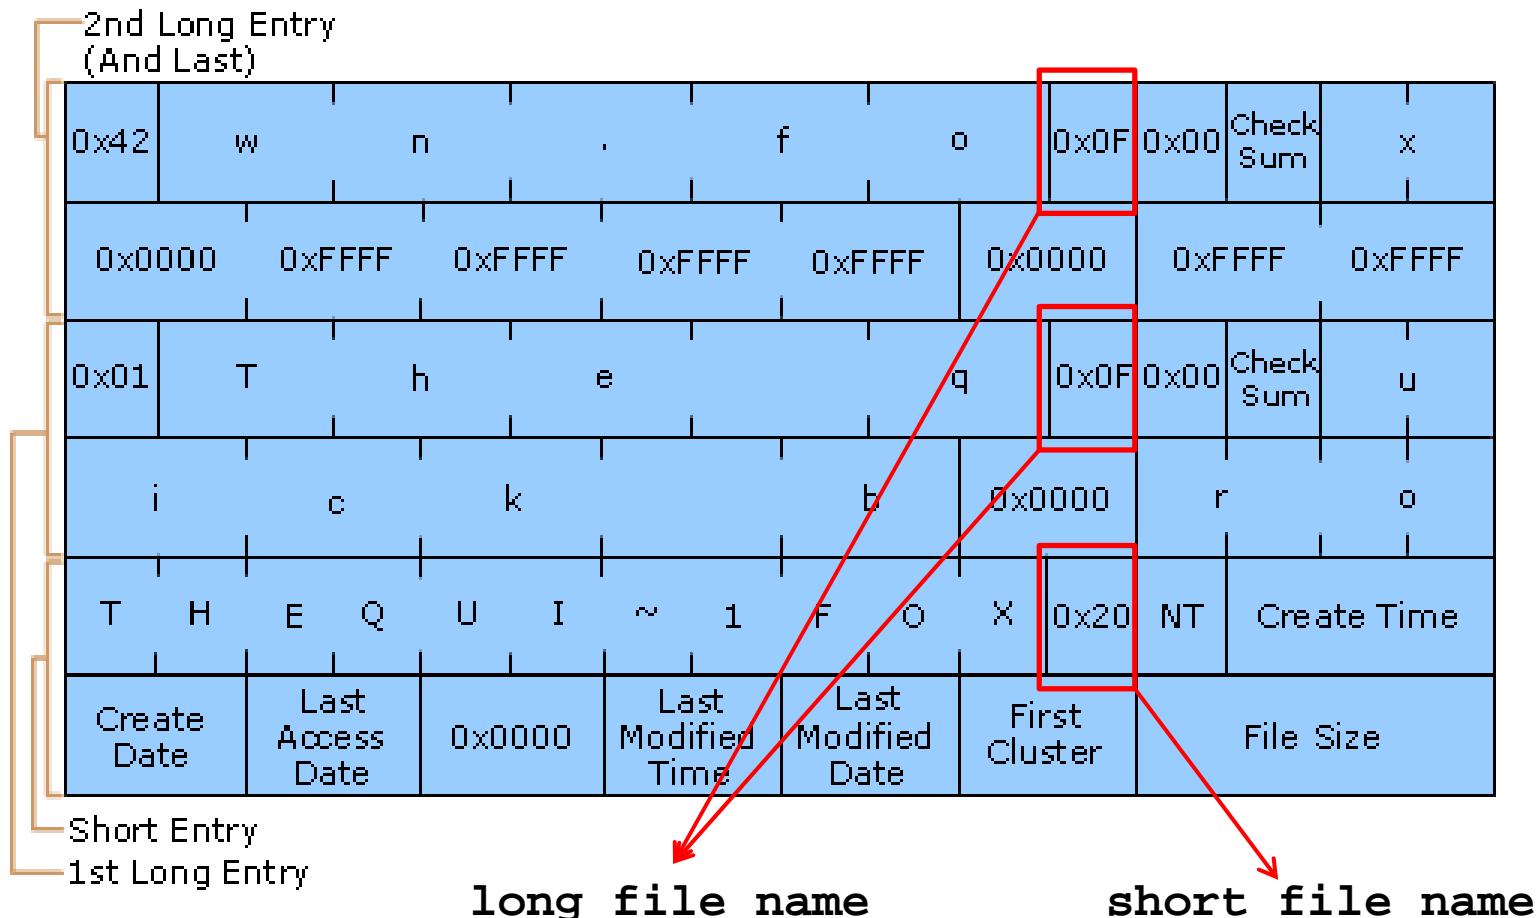
File System Durability & Crash Recovery

Yubin Xia & Rong Chen

Some of the materials are adopted from Frans'
6.828 course

Review: File Naming in FAT32

e.g. file with “The quick brown.fox” and “Thequi~1.fox”



FAT32 Long File Name

- Using 0xF in directory entry
 - Means this entry is part of long file name
 - 0xF is not used in DOS/WIN32 (consider as invalid)
- Using unicode for storage
 - 2 bytes (64 bits)
- The short (8.3) file name is still there
 - The system cannot work without short name

Bind Short Name and Long Name

```
int i,j,chknum=0;  
  
for (i=11; i>0; i--)  
{  
    chksun = ((chksun & 1) ? 0x80 : 0)  
        + (chksun >> 1)  
        + shortname[j++];  
}
```

Short Name Generation

- What if there is already a file with name THEQUI~1FOX?
 - Then use THEQUI~2FOX
 - If still conflict, using THEQUI~3FOX
 - If still conflict, using ...
 - If still conflict, using T~999999FOX
 - If still conflict, error

Question

- What is the process of searching a file by its long name?

File Name Search in ExFAT

- Search name by hash value first
 - Hash the upper case version of the file name
 - Each record in the directory is searched by comparing the hash value
 - When a match is found, the filenames are compared to ensure that the proper file was located in case of collisions
 - Only two characters have to be compared, why?

FS DURABILITY & CRASH CONSISTENCY

File System Durability

Topic: tension between fs perf. and crash recovery

Disk performance is often a #1 bottleneck

"how many seeks will that take?"

Durability != Crash consistency

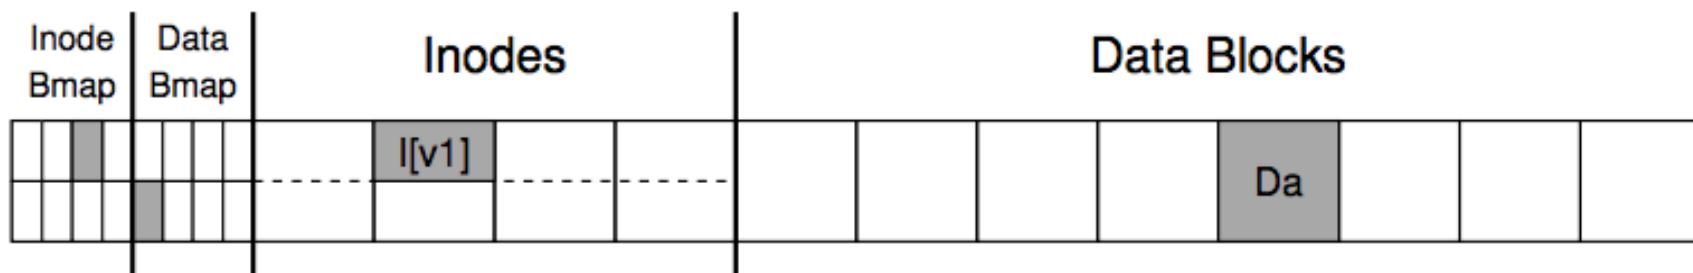
"Here is all of my data. But some of the metadata is wrong."

Crash recovery is much harder than performance

"what if a crash occurred at this point?"

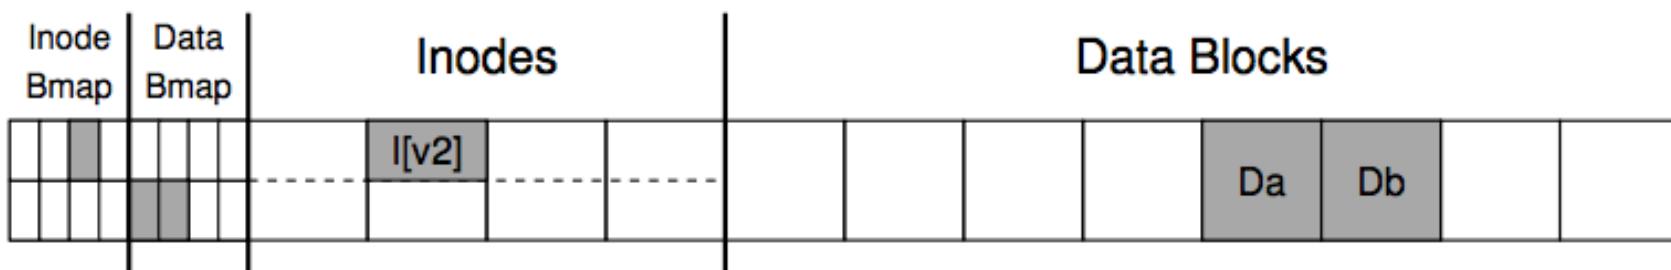
An Example: Append a File

- Inside of I[v1]:
 - owner : yubin
 - permissions : read-only
 - size : 1
 - pointer : 4
 - pointer : null
 - pointer : null
 - pointer : null



An Example: Append a File

- Inside of I[v2]:
 - owner : yubin
 - permissions : read-only
 - size : 2
 - pointer : 4
 - pointer : 5
 - pointer : null
 - pointer : null



Crash Scenarios: 1 Succeeds

- Imagine only a single write succeeds; there are thus three possible outcomes:
 - 1. Just the data block (D_b) is written to disk
 - What will happen?
 - 2. Just the updated inode ($I[v2]$) is written to disk
 - What will happen?
 - 3. Just the updated bitmap ($B[v2]$) is written to disk
 - What will happen?

Crash Scenarios: 2 Succeed

- Two writes succeed and the last one fails:
 - 1. The inode ($I[v2]$) and bitmap ($B[v2]$) are written to disk, but not data (D_b)
 - 2. The inode ($I[v2]$) and the data block (D_b) are written, but not the bitmap ($B[v2]$)
 - 3. The bitmap ($B[v2]$) and data block (D_b) are written, but not the inode ($I[v2]$)

Our Expectation

After rebooting and running recovery code

1. FS internal invariants maintained
 - E.g., no block is both in free list and in a file
2. All but last few operations preserved on disk
 - E.g., data I wrote yesterday are preserved
 - User might have to check last few operations
3. No order anomalies
 - \$ echo 99 > result ; echo done > status

Our Assumptions

Simplifying assumptions:

- Disk is fail-stop disk executes the writes FS sends it, and does nothing else
- Perhaps doesn't perform the very last write
- Thus: no wild writes, no decay of sectors

Why is fs crash recovery hard?

Crash = halt/restart CPU

let disk finish current sector write, assume no h/w damage,
no wild write to disk

Goal: automatic recovery

Can fs always make sense of on-disk **metadata** after restart?

Given that the crash could have occurred at **any point?**

Examples

Crash during mkdir, leave directory without . and ..

Crash during free blocks

Offline and Online Recovery

Offline recovery

file system check utility, such as chkdsk in windows and fsck on linux

E.g., ext3

Online recovery

during operation, check some important inconsistency

E.g., ext4 (also has offline fsck, but much simpler)

Terms for properties of fs ops

What effects will app see after restart + recovery

creat("a"); fd = creat("b"); write(fd,...); crash

Durable: effects of operation are visible

Both a and b are visible

Atomic: all steps of operation visible or none

Either a and b are visible or none is visible

Ordered: exactly a prefix of operations is visible

If b is visible, then a is visible

Recovery approach

1. Synchronous meta-data update + fsck

Used in xv6-rev0

During check, synchronize meta-data, such as file size

2. Soft update (FreeBSD fs modified on FFS)

Soft update, not covered in this course

3. Logging (ext 3/4), xv6-rev6 and following versions

Before doing actual meta-data update, log the event

After crash, recover from log

SYNC METADATA UPDATE + FSCK

Typical set of tradeoffs

FS ensures it can recover its **meta-data** (minimal requirements for a real fs)

- Internal consistency

- No dangling references

- Inode and block free list contain only used (not using) items

- Unique name in one directory, etc.

Weak semantic FS provided limited guarantees

- Atomicity for creat, rename, delete

- Often no durability for anything

- (creat("a"), then crash, no a)

- Often no order guarantees

How do applications handle this weak semantics?

Example

Edit your file, then crash, only half of your file is actually updated?

Fsync and rename (shadow copy)

Fsync force durability, only returned if file is actually written on disk

Rename is an atomic operation, only old name or new name, not half old half new

Mac OS intensively uses rename to ensure atomicity

What Does fsck do?

- 1. Check superblock
 - E.g., making sure the file system size is greater than the number of blocks allocated
 - If error, use an alternate copy of the superblock
- 2. Check free blocks
 - Scans the inodes, indirect blocks, double indirect blocks, etc.
 - Uses this knowledge to produce a correct version of the allocation bitmaps
 - Same for the inode bitmap

What Does fsck do?

- 3. Check inode states
 - Check type: regular file, dir, symbolic link, etc.
 - Clear suspect inodes and clear the inode bitmap
- 4. Check inode links
 - Check link count by scanning the entire fs tree
 - If count mismatches, fix the inode
 - If inode is allocated but no dir contains it, lost+found
- 5. Check duplicates
 - Two inodes refer to the same block
 - If one inode is obviously bad, clear it; otherwise, copy the block and give each a copy

What Does fsck do?

- 6. Check bad blocks
 - E.g., point to some out-of-range address
 - What should fsck do? Just remove the pointer
- 7. Check directories
 - The only file that fsck know more semantic
 - Making sure that “.” and “..” are the first entries
 - Ensure no dir is linked more than once
 - No same filename in one dir

Problem of fsck: Too Slow

How long would fsck take?

an example server: fsck takes 10 minutes per 70GB disk w/ 2 million inodes

clearly reading

still a long time

Consider the case

Scan the disk

Just like find



clearly, not seeking
the size

writes
entire house

Would an xv6 FS be internally consistent after a crash?

Xv6-rev0 strategy: carefully order disk writes to avoid dangling refs

1. initialize a new inode before creating dirent
2. delete dirent before marking inode free
3. mark block in-use before adding it to inode addrs[]
4. remove block from addrs[] before marking free
5. zero block before marking free

Has some visible bugs:

- . and .. during mkdir(), link counts, sizes

Has some invisible loose ends
may lose freed blocks and inodes

Example

file creation: what's the right order of synchronous writes?

1. mark inode as allocated
2. create directory entry

file deletion

1. erase directory entry
2. erase inode addrs[], mark as free
3. mark blocks free

What about app-visible syscall semantics?

- Durable? Yes
 - Use write-through cache, sync I/O, O_SYNC
- Atomic? Often
 - **Mkdir** is an exception
- Ordered? Yes
 - If all writes are sync

Recall: Sync I/O vs. Async I/O

Asynchronous I/O is a **poor** abstraction for:

- Reliability
- Ordering
- Durability
- Ease of programming

Synchronous I/O is superior but **100x slower**

Caller blocked until operation is complete

Issues with Synchronous Write

Main issue

very slow during normal operation

very slow during recovery

Barrier: Flush the Disk

- Disk's write buffer
 - Disk will inform the OS the write is complete when it simply has been placed in the disk's memory cache
 - But the data is not on disk yet! No durability! No order!
- One solution: disable the buffer
- Another solution: using flush operation
 - Force the disk to write data to disk media
 - Aka., disk write barrier
- However, disks may not do as they claim...
 - Some disks just ignore the flush operation to be faster
 - “the fast almost always beats out the slow, even if it is wrong” --- Kahan”

Ordinary perf. of sync meta-data update?

Creating a file and writing a few bytes

Takes 8 writes, probably 80 ms

(ialloc, init inode, write dirent, alloc data block, add to
inode, write data, set length in inode, xxx)

So can create only about a dozen small files per second!

Think about un-tar or rm *

How to get better performance?

Reality

RAM is cheap

disk sequential throughput is high, 50 MB/sec (maybe someday solid state disks will change the landscape)

Why not use a big write-back disk cache?

no sync meta-data update operations

only modify in-memory disk cache (no disk write)

so creat(), unlink(), write() &c return almost immediately bufs written to disk later

if cache is full, write LRU dirty block

write all dirty blocks every 30 seconds, to limit loss if crash

this is how old Linux EXT2 file system worked

Write-back Cache

Would write-back cache improve performance? why, exactly?

after all, you have to write the disk in the end anyway

What can go wrong with write-back cache?

example: `unlink()` followed by `create()` an existing file `x` with some content, all safely on disk

one user runs `unlink(x)`

1. delete `x`'s dir entry **
2. put blocks in free bitmap
3. mark `x`'s inode free; another user then runs `create(y)`
4. allocate a free inode
5. initialize the inode to be in-use and zero-length
6. create `y`'s directory entry **

again, all writes initially just to disk buffer cache

suppose only ** writes forced to disk, then crash

what is the problem?

can `fsck` detect and fix this?

Flash File System and GFS

Rong

Outline

- Flash file system
- Google file system

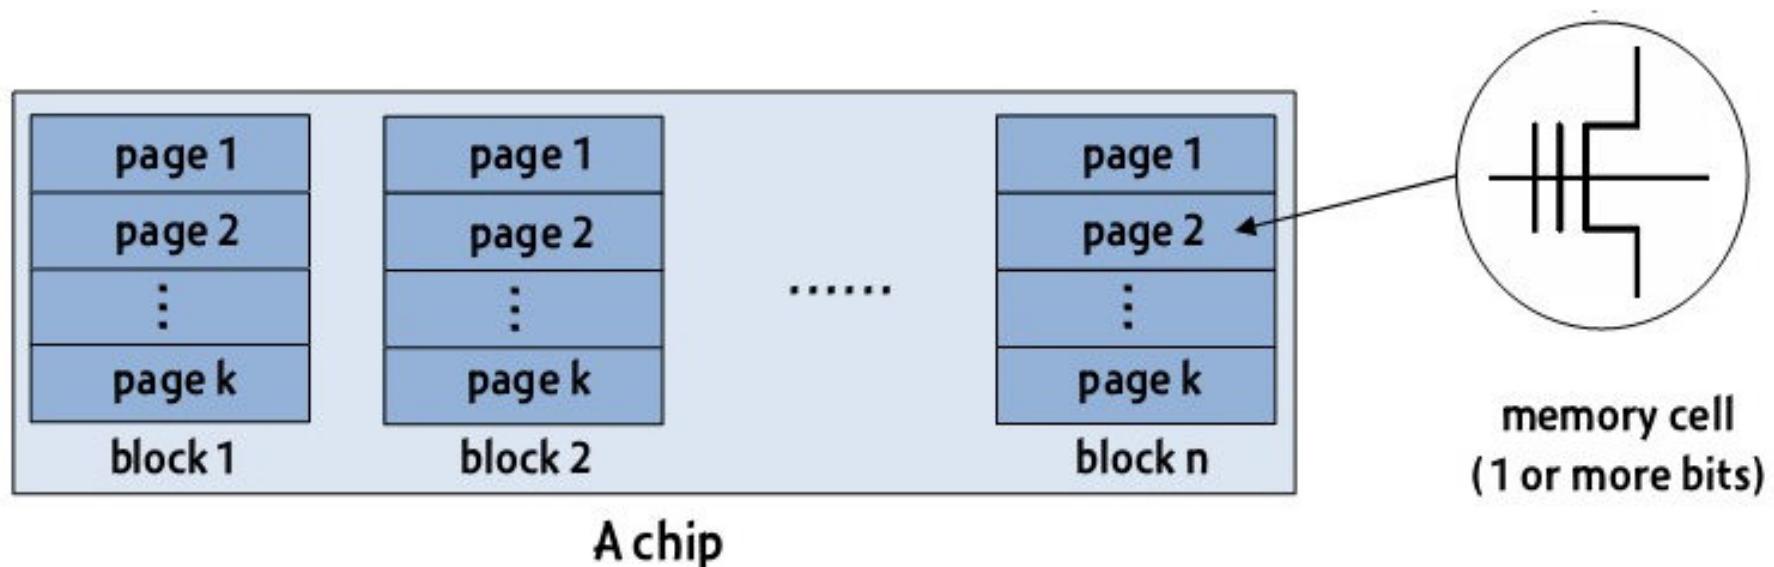
INTRO TO FLASH FILE SYSTEM

What is flash file system?

- Flash file system is designed for storing files on flash memory devices
- Can traditional file system be used on flash file system?
 - Of course, NO.
- So, what is the big difference between flash disk and magetic disk

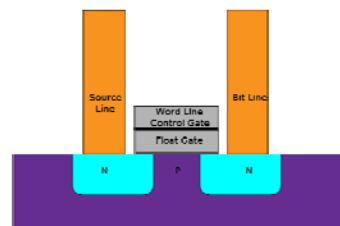
Flash disk organization (1/2)

- Flash disk organization
 - A **chip** (e.g. 1GB) => **blocks** (e.g. 512KB) => **pages** (e.g. 4KB) => **cells**

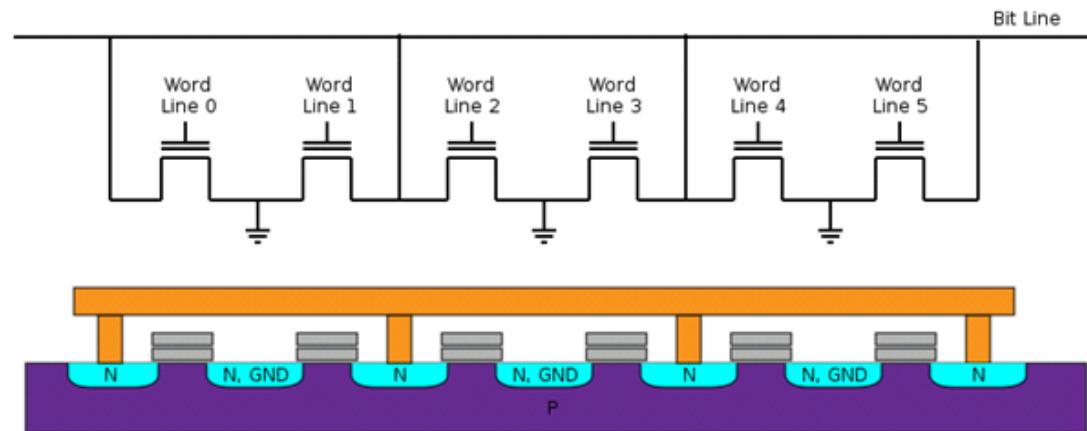


Flash disk organization (2/2)

- Flash cell: a floating gate transistor
 - The number of electrons on the floating gate determines the threshold voltage V
 - The threshold voltage represents a logical bit value (0 or 1)

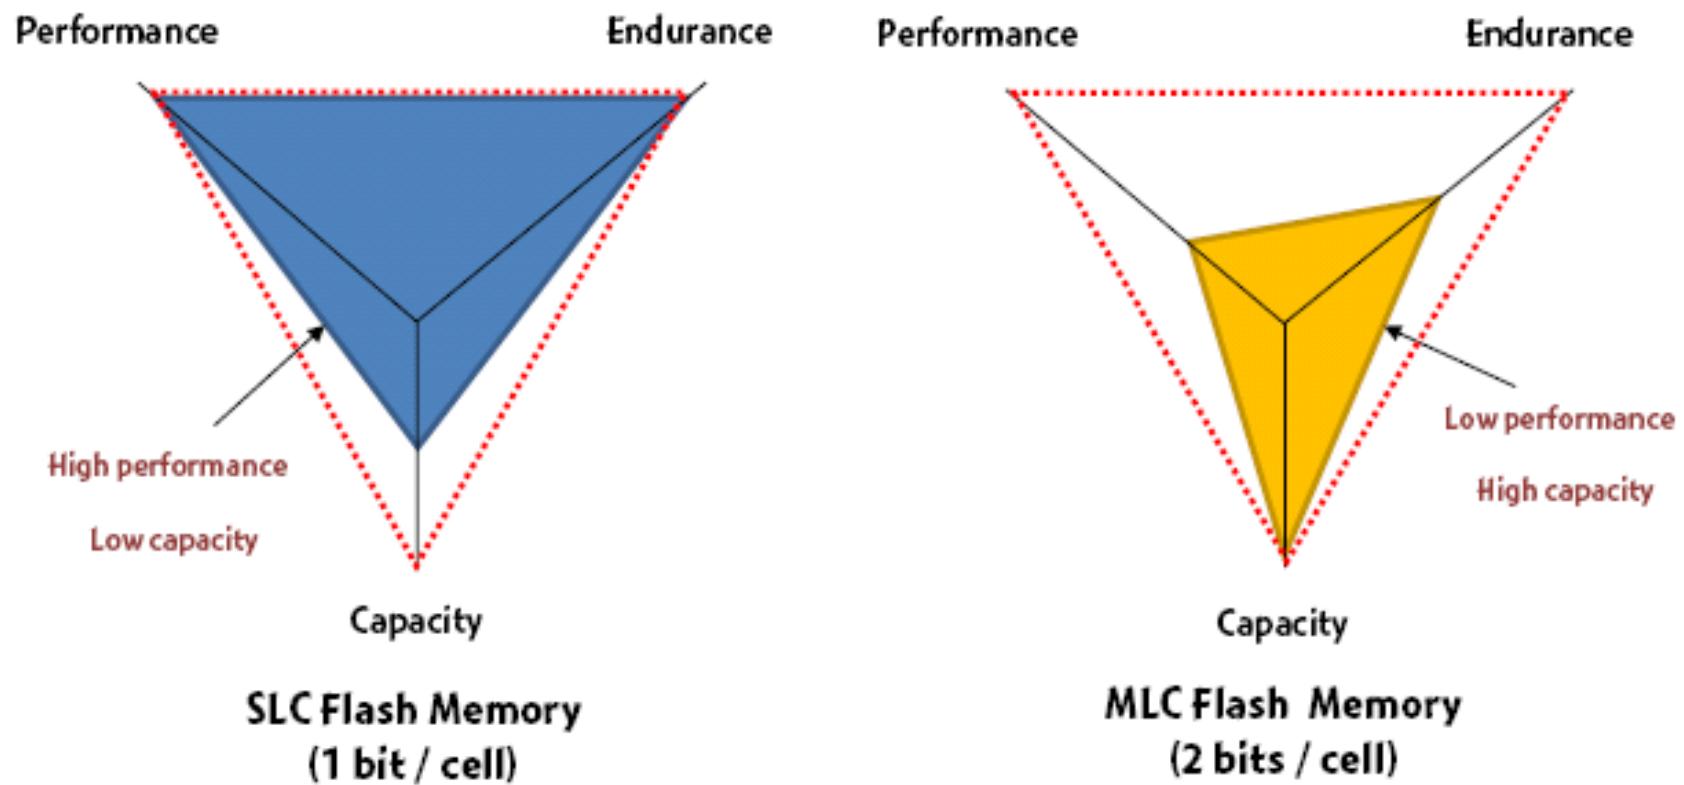


(Single-Level Cell, SLC)



(Multi-Level Cell, MLC)

Comparisons of SLC and MLC Flashes



Flash disk characters (1/2)

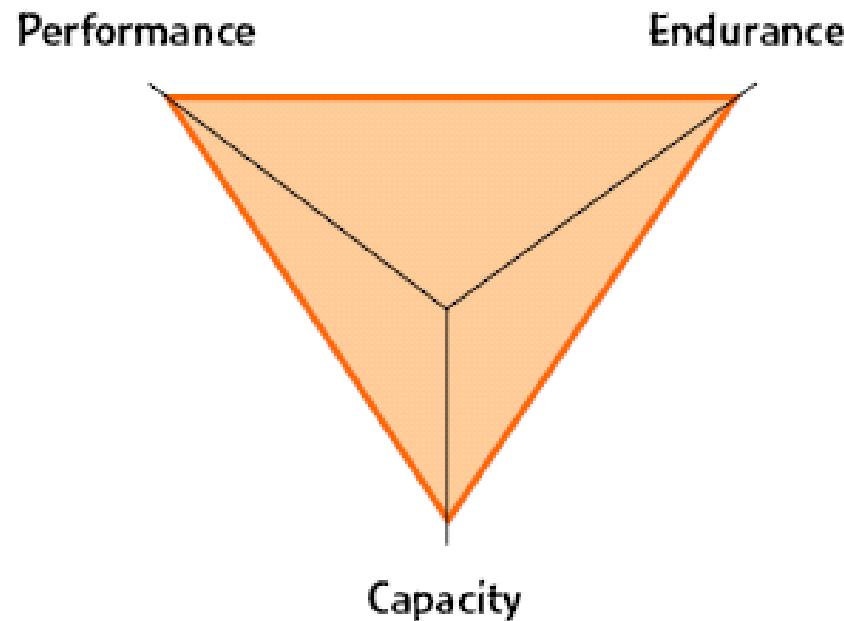
- Asymmetric **read/write** and **erase** operations
 - A “**page**” is a unit of read/write
 - A “**block**” is a unit of erase
- Physical restrictions
 - **Erase-before-write** restriction
 - The number of erase cycles allowed for each block is limited

Flash disk characteristics (2/2)

- Random access
 - Disk file system are optimized to avoid disk seeks whenever possible
 - Flash device impose no seek latency
- Wear leveling
 - Flash device tend to wear out when a single block is repeatedly overwritten
 - Designed to spread out writes evenly
- Heterogeneous cells
 - SLC (single level cell) and MLC (multi level cell)

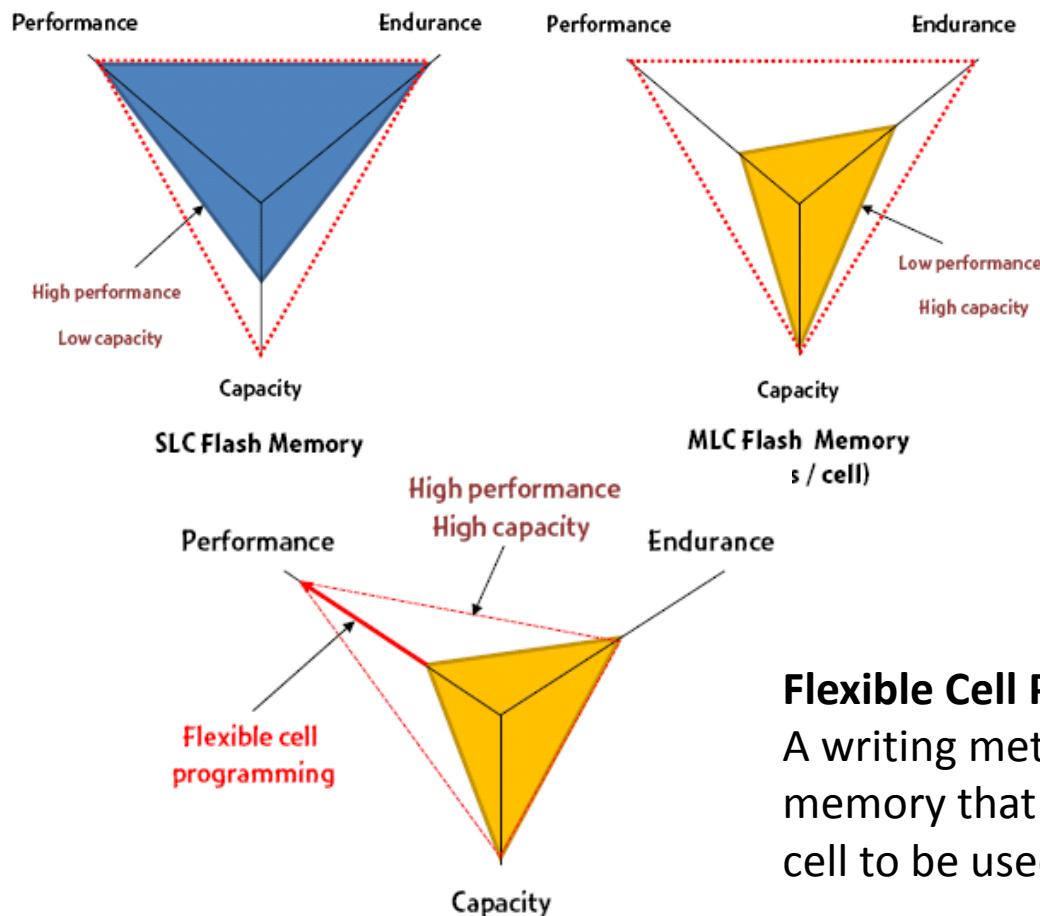
A real flash file system study - flexFS

- Consumers want to have a storage system with high performance, high capacity, and high endurance
- HOW ??



Flexible Cell Programming

- Heterogeneous - Makes it possible to take benefits of two different types of NAND flash memory

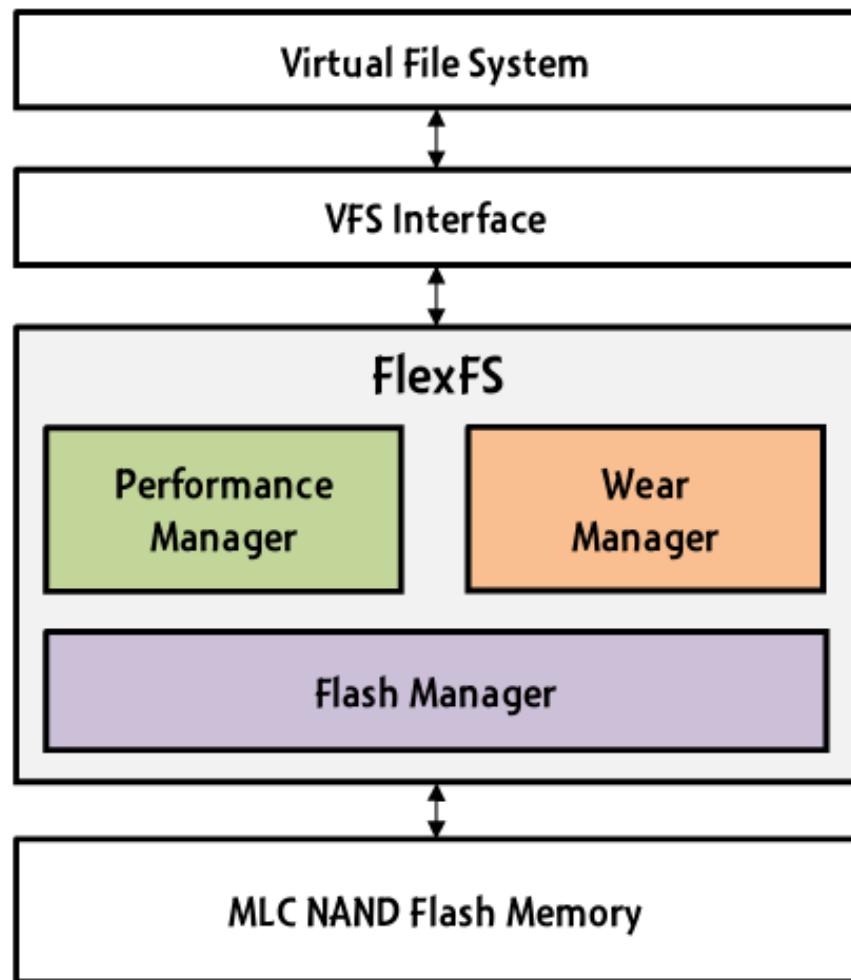


Flexible Cell Programming:
A writing method of MLC flash memory that allows each memory cell to be used as SLC or MLC

Our Approach

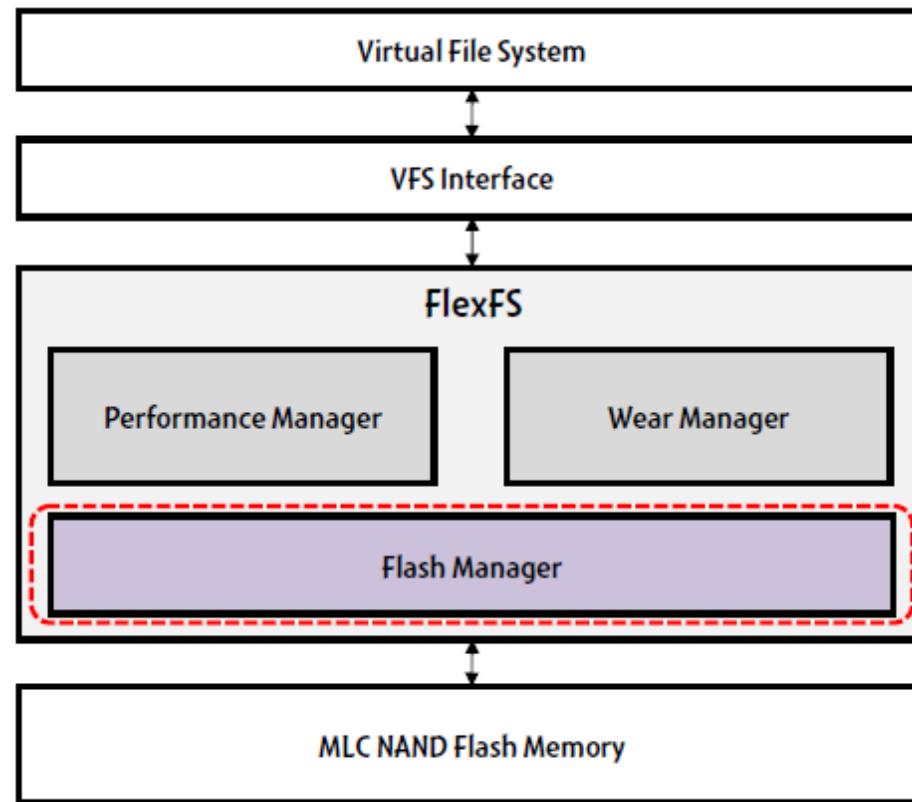
- Proposes a flash file system called FlexFS
 - Exploits **flexible cell programming**
 - Provides the **high performance** of SLC flash memory and the **capacity** of MLC flash memory
 - Provides a mechanism that copes with a poor **wear** characteristic of MLC flash memory
 - Designed for mobile systems, such as mobile phones

Architecture of FlexFS



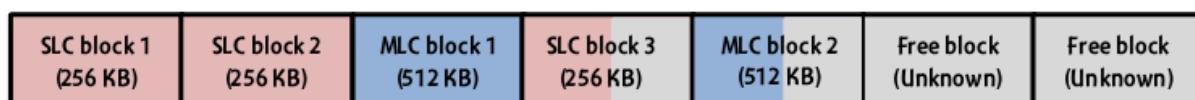
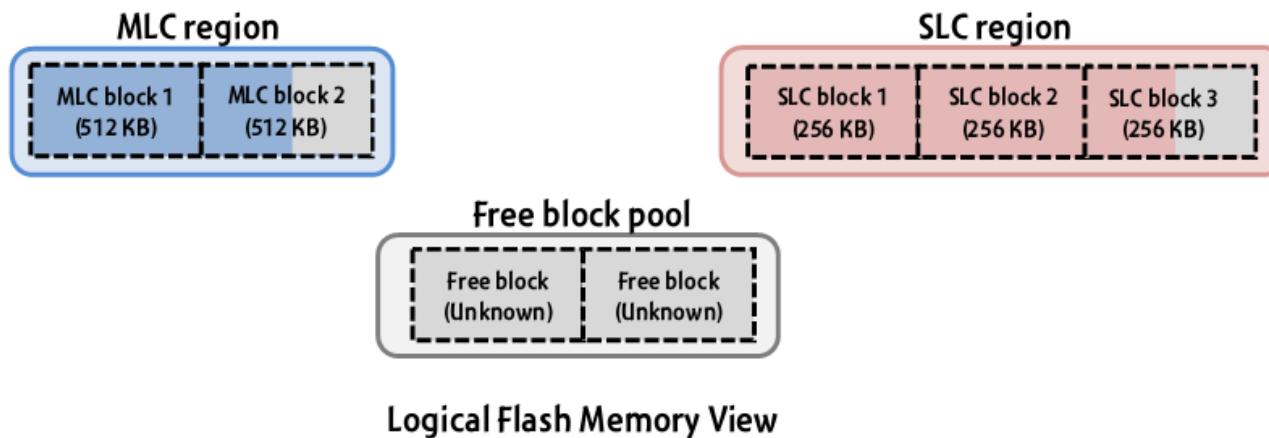
- **Flash Manager**
 - Manages heterogeneous cells
- **Performance manager**
 - Exploits I/O characteristics
 - To achieve the high performance and high capacity
- **Wear manager**
 - Guarantees a reasonable lifetime
 - Distributes erase cycles evenly

Overall Architecture



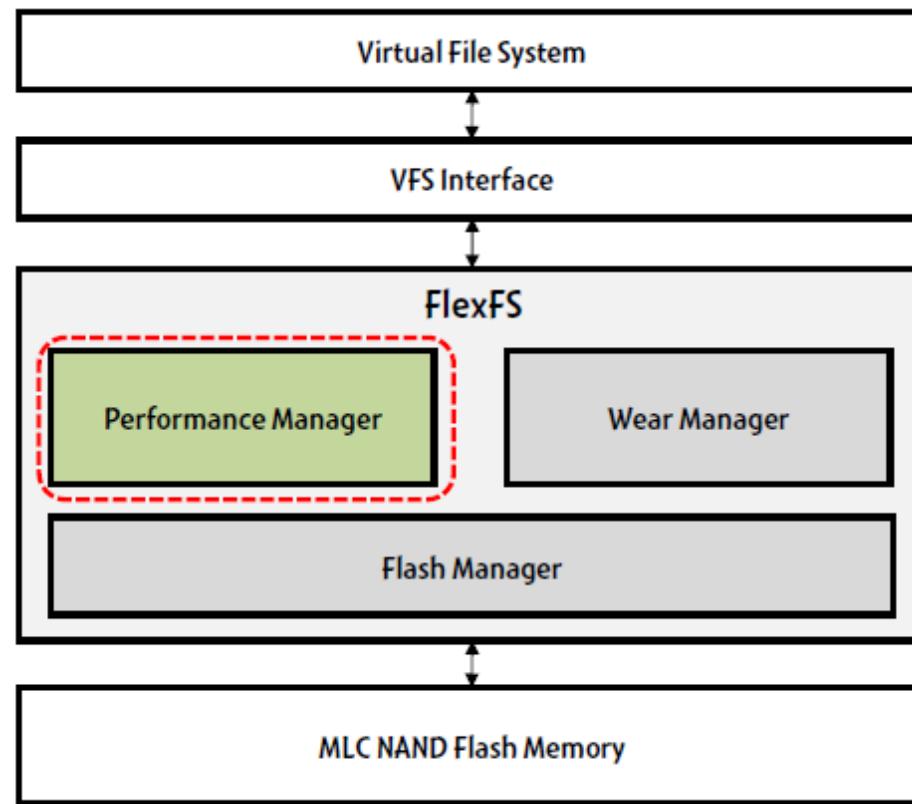
FlexFS - Flash Manager

- Handles Heterogeneous Cells
 - Three types of flash memory block: SLC block, MLC block, and free block
 - Manages them as two regions and one free block pool



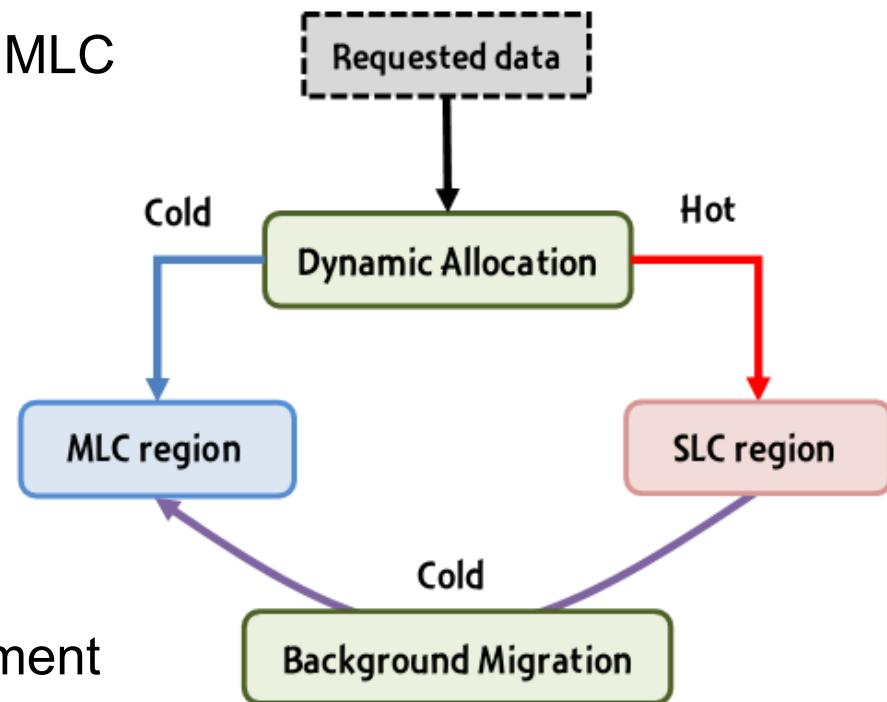
Physical Flash Memory View

Overall Architecture

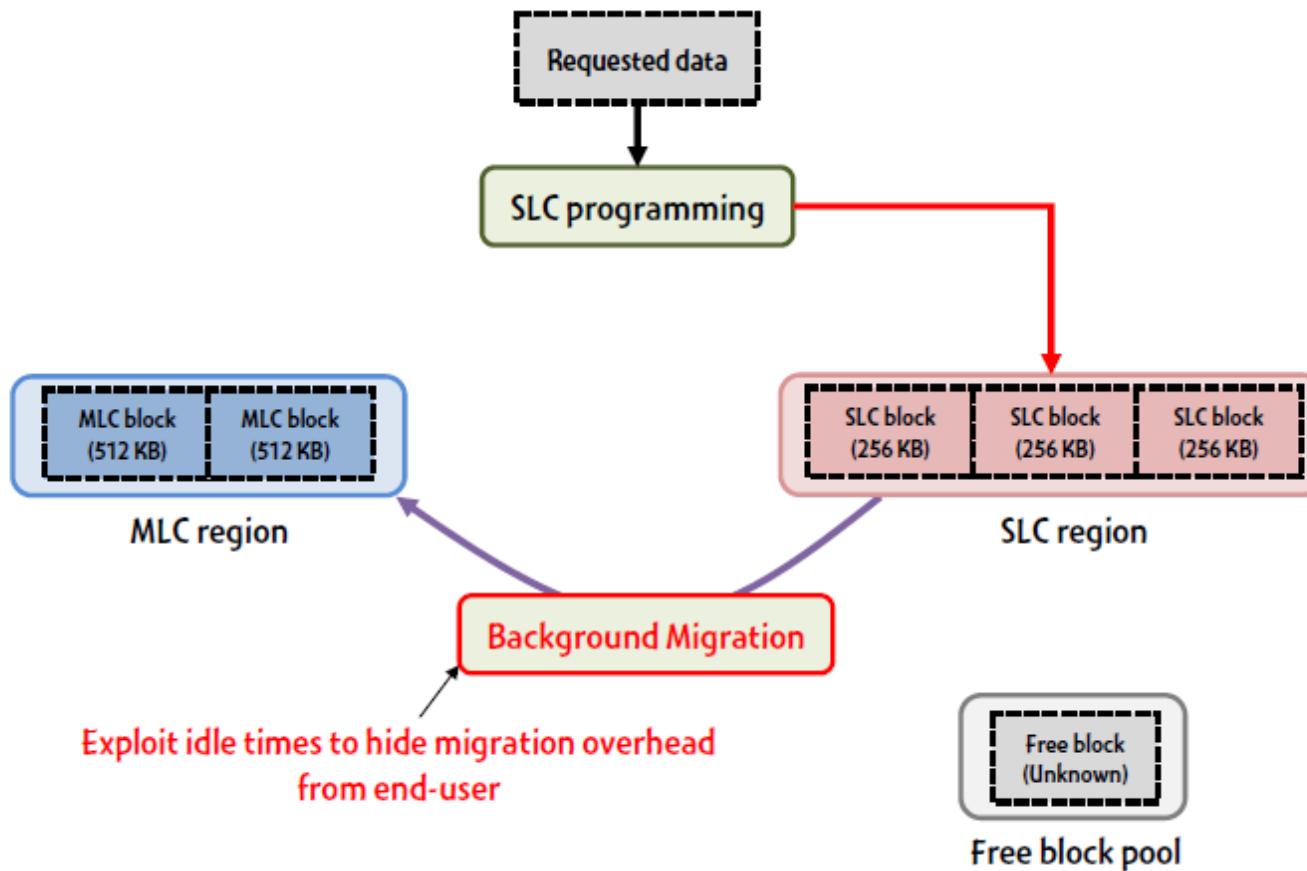


FlexFS – Performance Manager

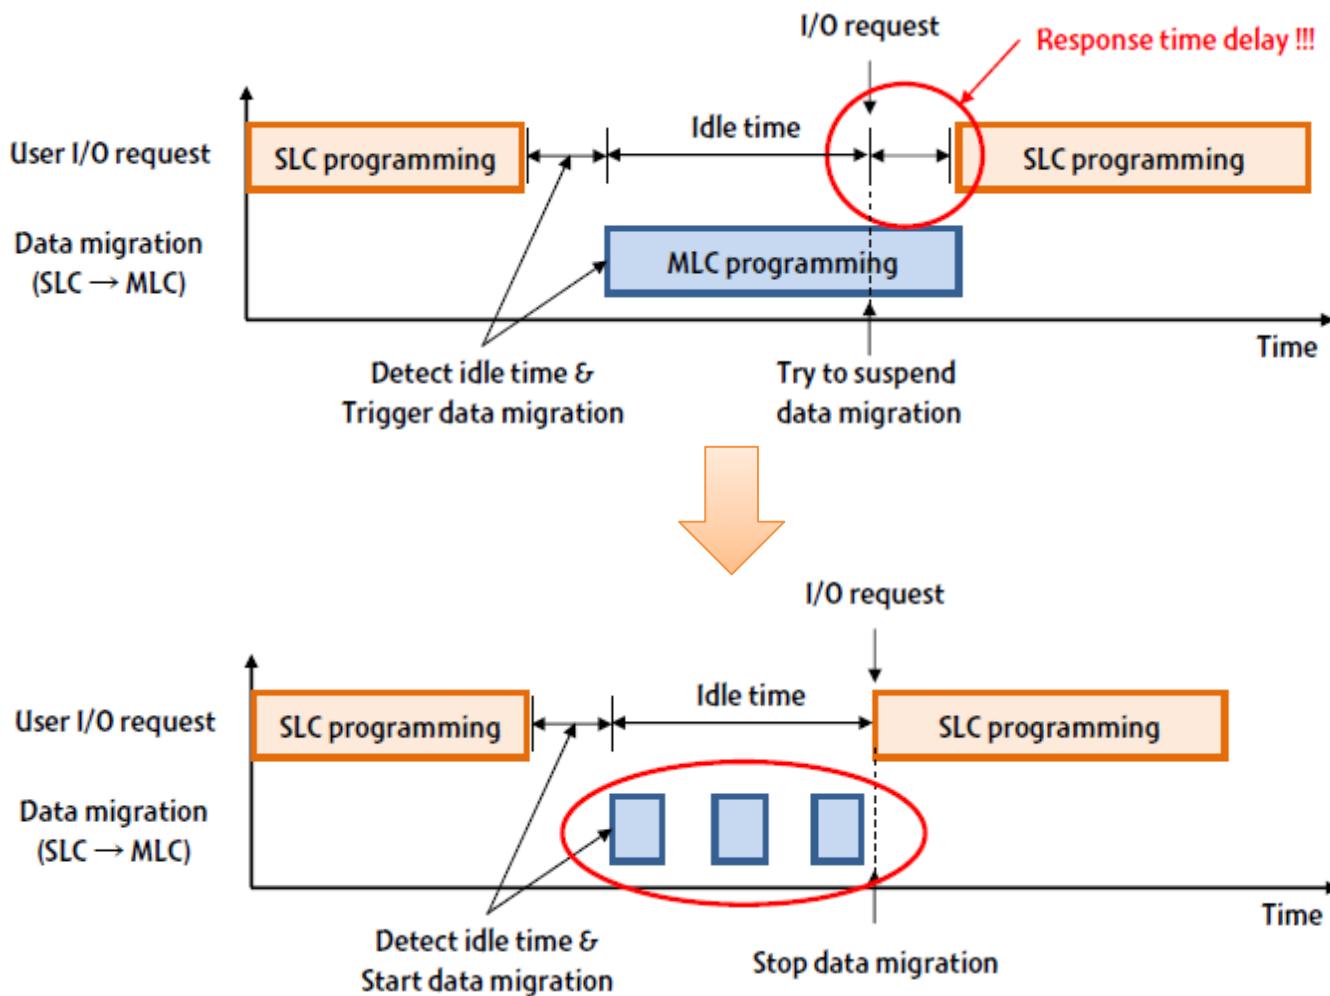
- Manages SLC and MLC regions
 - Provide SLC performance and MLC capacity
 - Exploits I/O characteristics, such as idle time and locality
- Three key techniques
 - Dynamic allocation
 - Background migration
 - Locality – aware data management



Baseline Approach

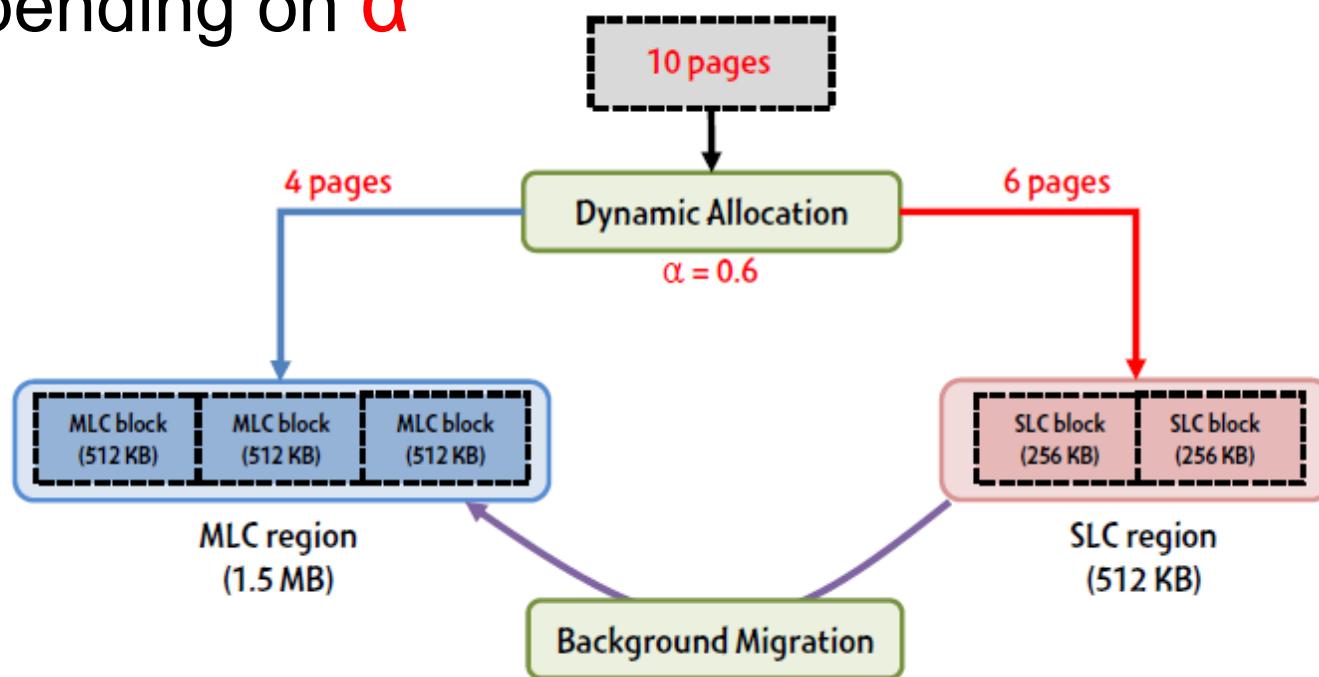


Background Migration



Dynamic Allocation

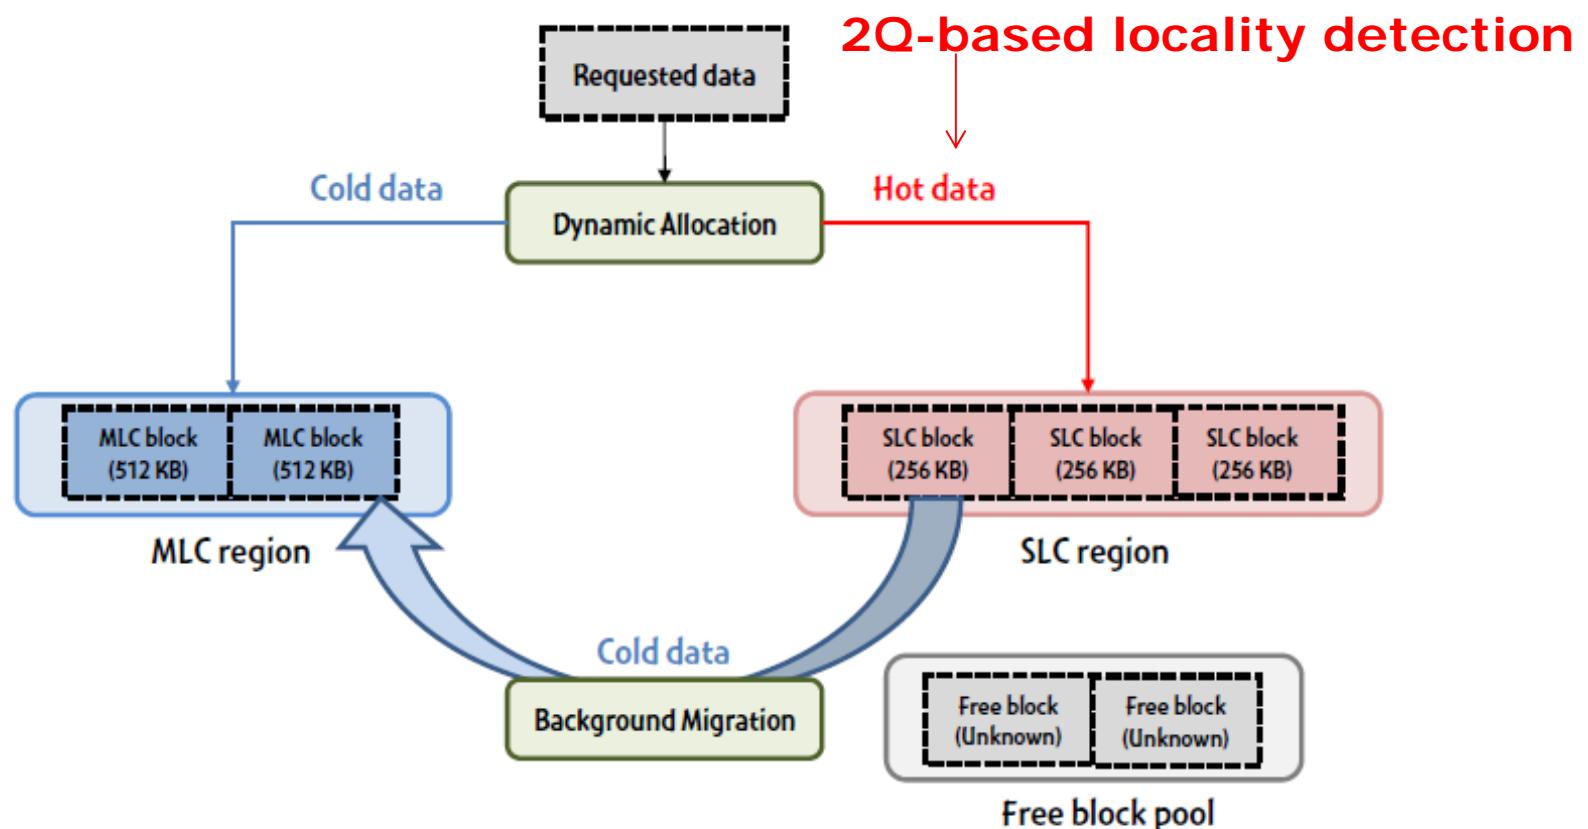
- Distributes the incoming data across two regions depending on α



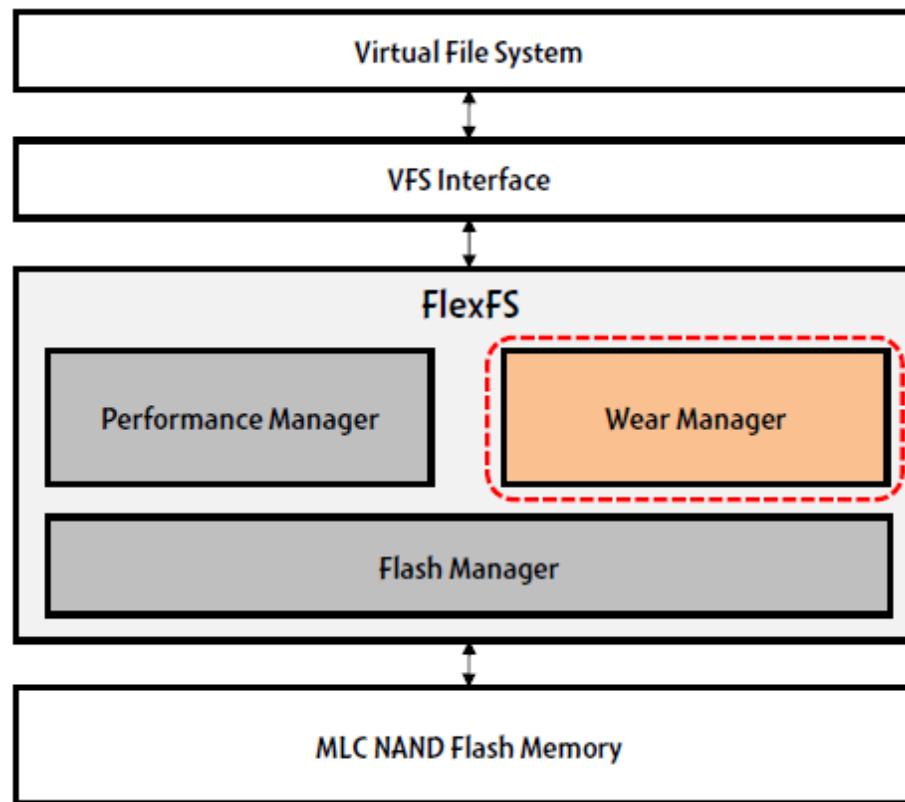
$$\alpha = \frac{T_{\text{predict}}}{N_p \cdot T_{\text{copy}}} \quad (\text{If } T_{\text{predict}} \geq N_p \cdot T_{\text{copy}}, \text{ then } \alpha = 1.0)$$

Locality-aware Data Management

- Data migration for hot data is unnecessary

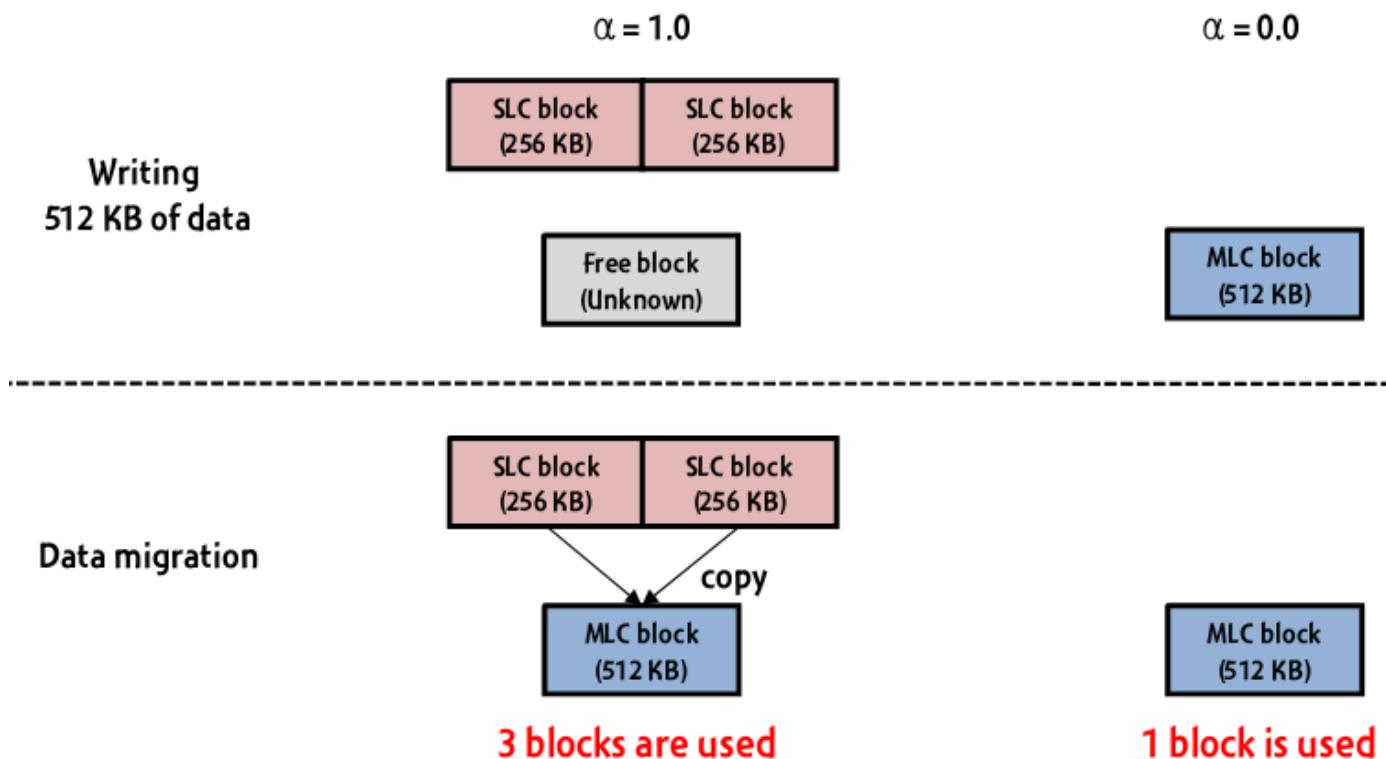


Overall Architecture

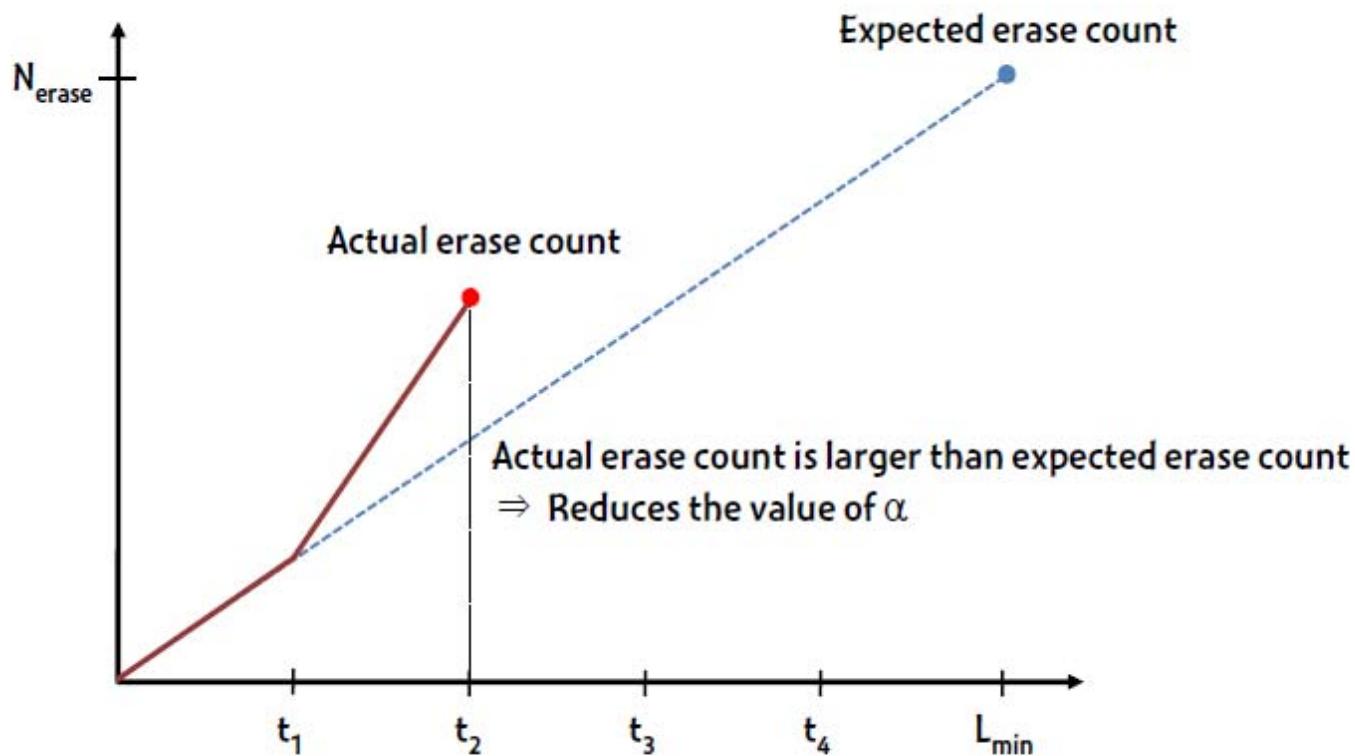


FlexFS– wearing rate control

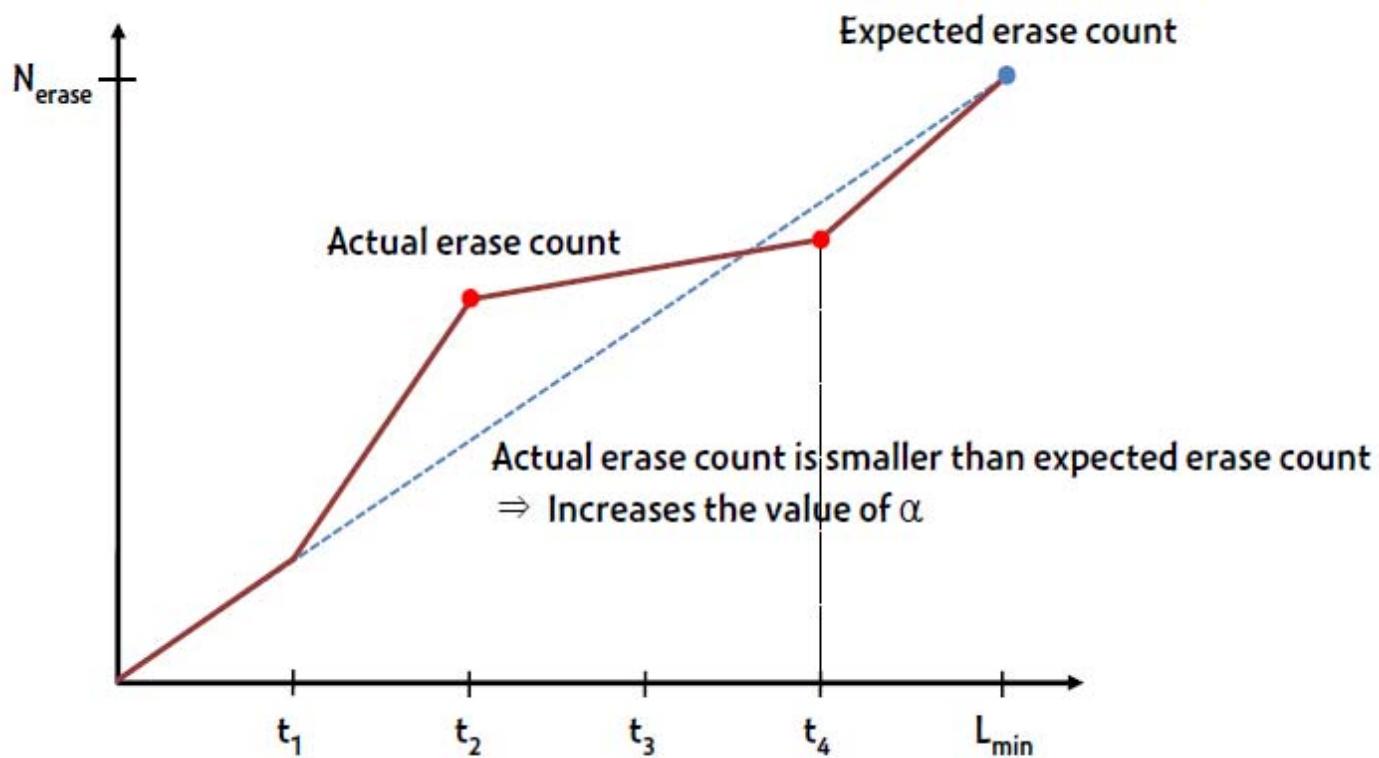
- How FlexFS control the wearing rate
 - The wearing rate is directly proportional to the value of α



Wearing Rate Control

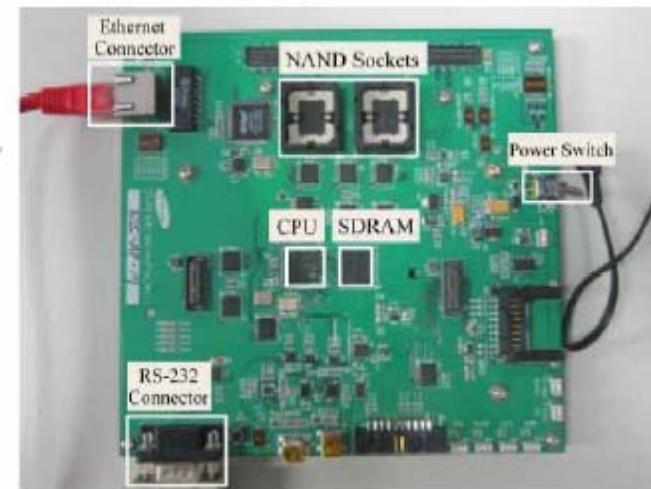
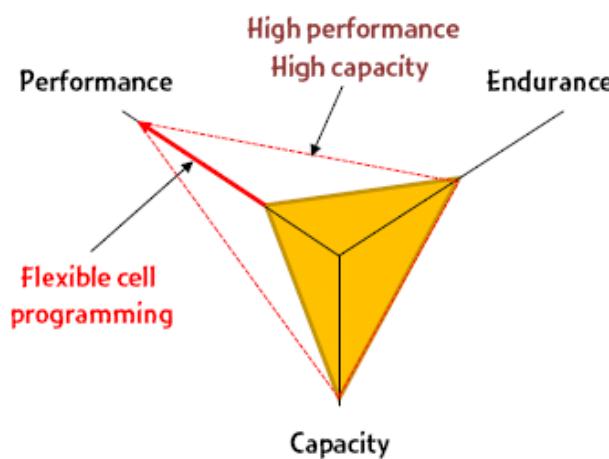


Wearing Rate Control



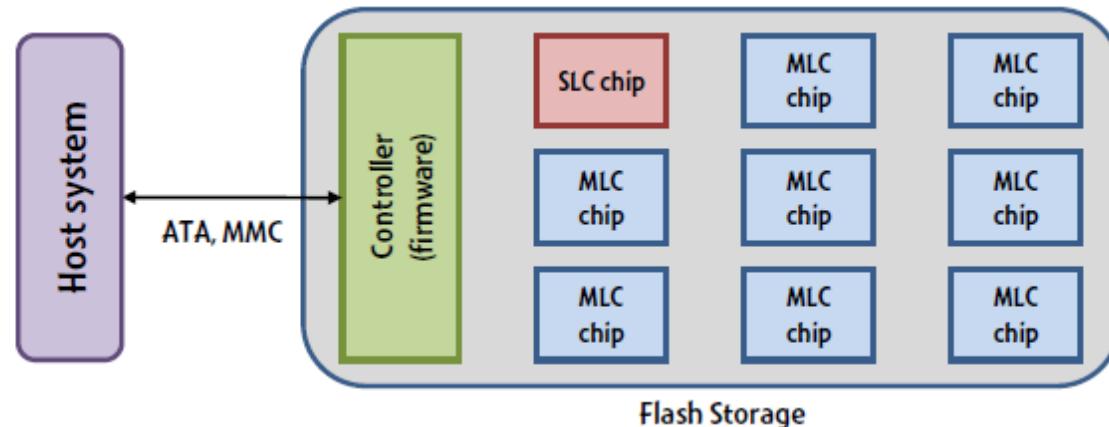
Conclusion

- Propose a new file system for MLC NAND flash memory
 - Exploits the flexible cell programming to achieve the SLC **performance** and MLC **capacity** while ensuring a reasonable **lifetime**



Other Solution

- SLC/MLC hybrid storage
[Chang et al (2008), Park et al (2008), Im et al (2009)]
 - Composed of a single SLC chip and many MLC chips
 - Uses the SLC chip as a write buffer for MLC chips
 - Redirects frequently accessed small data into the SLC chip
 - Redirects bulk data into the MLC chips



INTRO TO GFS

Introduction to google file system

- Shares same goals as previous distributed file systems, but with some differences.
 - Component **failures** are considered the norm rather than the exception.
 - Files are **huge** by traditional standards.
 - Most files are mutated by **appending** new data rather than overwriting existing data.
 - **Co-designing** the file system and applications increases flexibility in development.

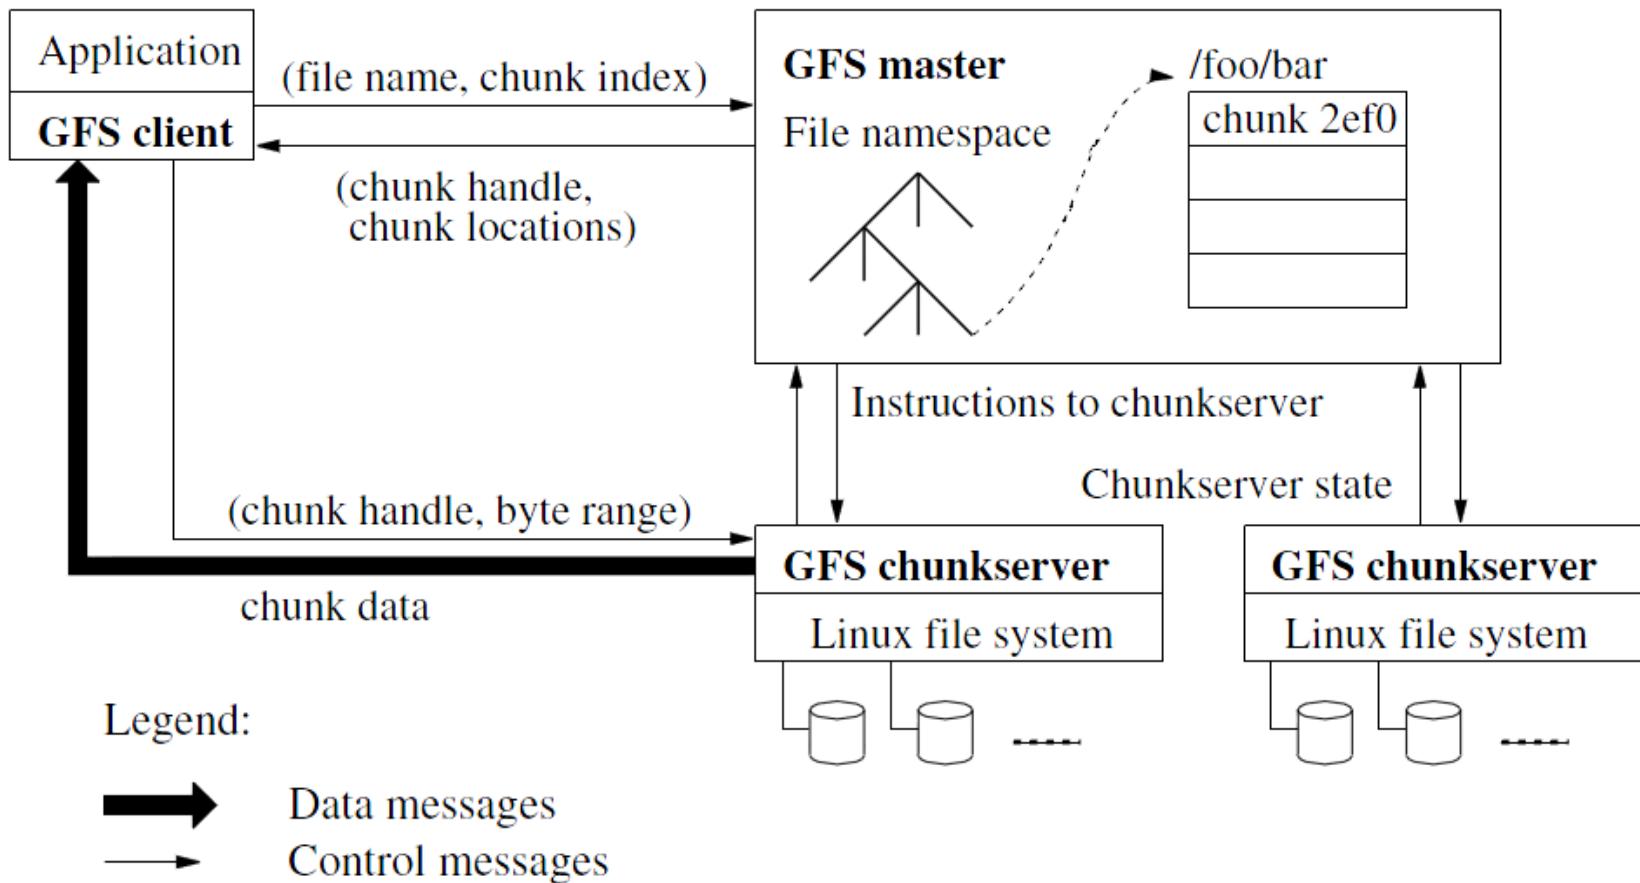
Design Assumptions

- System is built from **commodity hardware** which fails as the norm.
 - The system must be able to detect and recover from such occurrences.
- The system must be optimized to deal with **large** files. (Multi-GB)
- Two types of **reading**
 - Large streaming reads
 - Small random reads

Design Assumptions

- Large **sequential** writes
 - Appending data to files which are seldom altered again.
- Atomicity in **concurrent** writes from multiple clients
- High sustained **bandwidth** is more valuable than low **latency**

Architecture



Architecture Design

- **GFS Cluster** (accessed by clients)
 - Single Master + Multiple Chunkservers
- **Chunkservers**
 - Files of fixed sized chunks
 - Each chunk has a globally unique 64 bit chunk handle.
- **Master**
 - Maintains file system metadata
 - Namespace
 - Access Control Information
 - Mapping from files to chunks
 - Current locations of chunks

Single Master

- Clients do not read or write through the master
- The master relays relevant chunkserver location information to the client
- The client temporarily caches the chunkserver data and directly accesses the chunkserver
- Shadow Masters (fault tolerance)

Metadata

- Three major types:
 - Chunk namespaces
 - Mapping from files to chunks
 - Location of chunk replicas
- Chunk location is not stored persistently and is instead read on each startup
- Metadata is stored in **memory**
- Master monitors chunk location through **heartbeat** messages with each chunk.

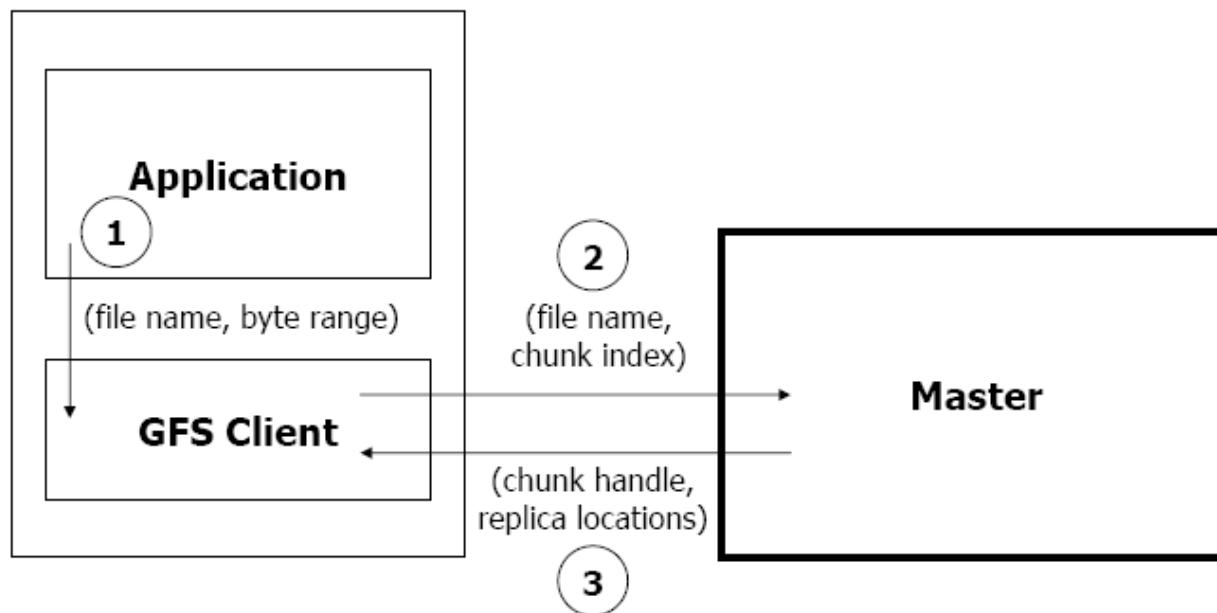
Chunkserver

- Contains chunks (blocks) of a fixed size
 - 64 MB
 - Fewer chunk location, fewer metadata
- Chunks are replicated regularly
 - Default: **3-way mirror** (across machines and racks)
- Talks with master through ***heartbeat*** messages

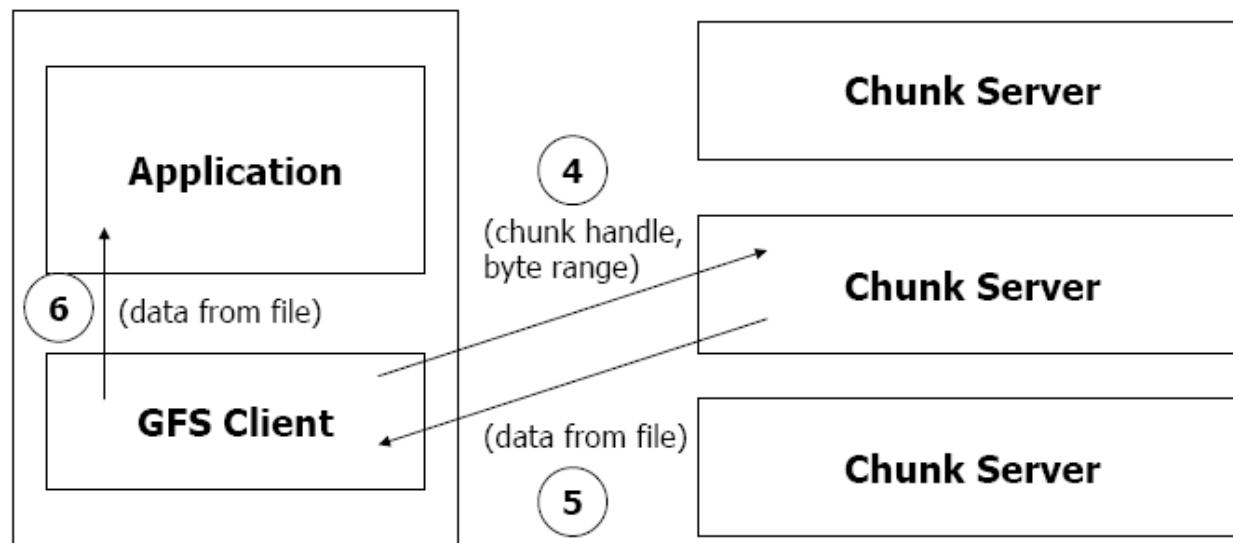
System Interactions

- Mutation
 - Write
 - Append
 - Acts on all of the chunk's replicas
- Lease
 - Initial timeout of 60 seconds, which can be renewed or revoked.
- Data Flow
 - Decoupled **control** flow from **data** flow
 - Data pushed linearly to avoid bottlenecks and high latency links

Read Algorithm



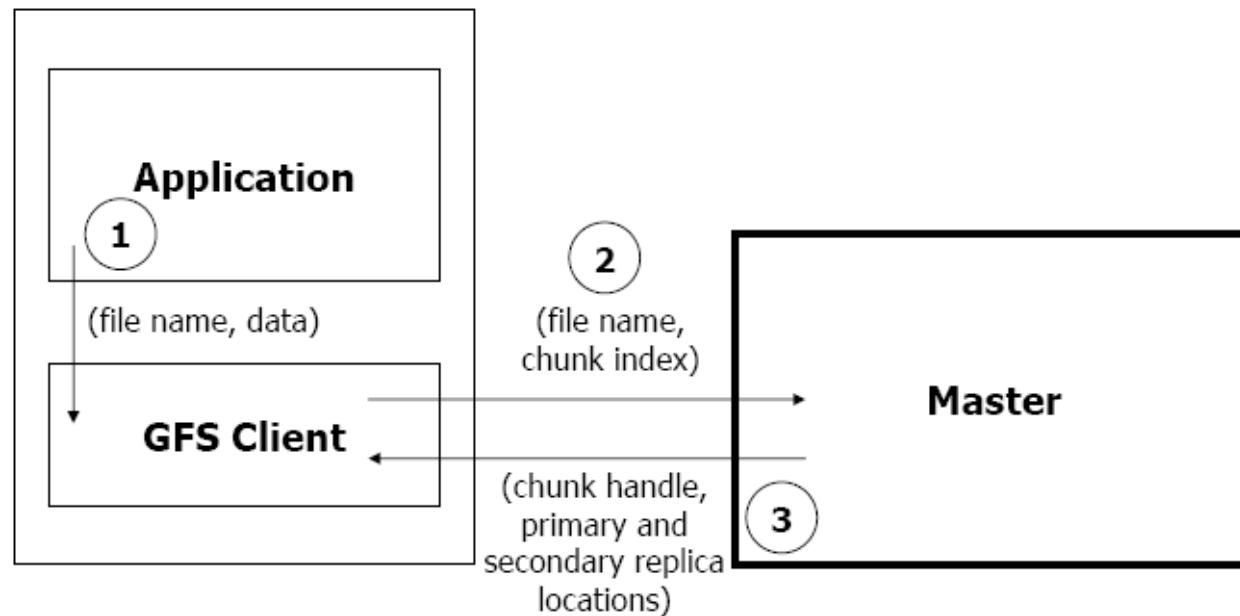
Read Algorithm



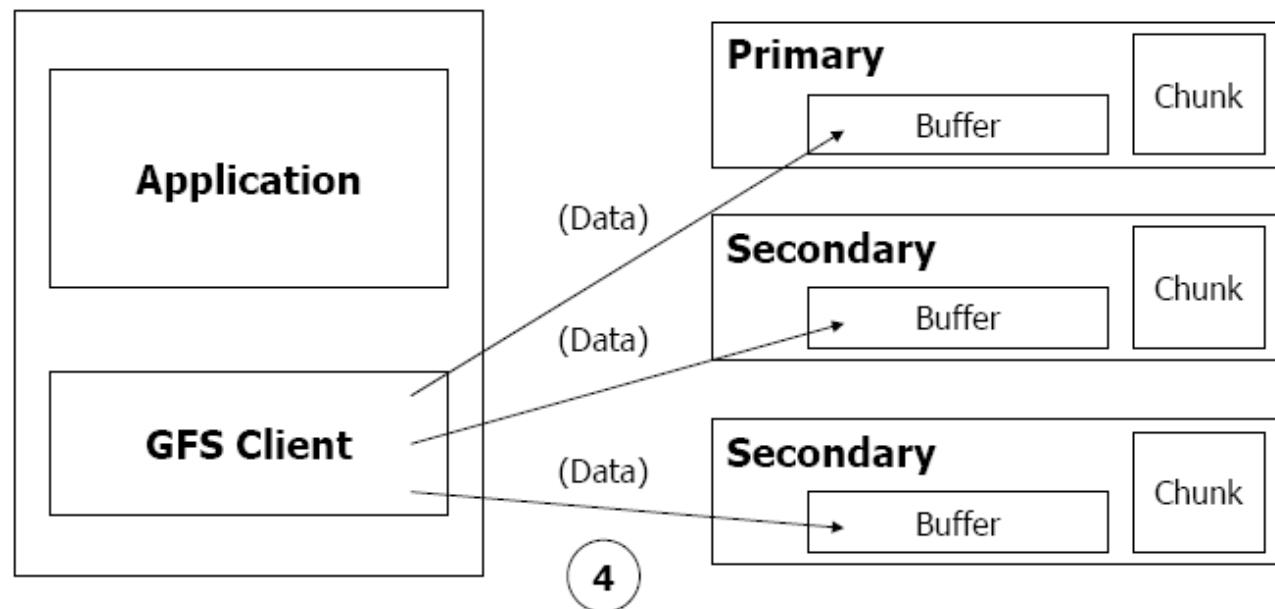
Read Algorithm

1. Application originates the read request.
2. GFS client translates the request from (filename, byte range) -> (filename, chunk index), and sends it to master.
3. Master responds with chunk handle and replica locations (i.e. chunk servers where the replicas are stored).
4. Client picks a location and sends the (chunk handle, byte range) request to that location.
5. Chunk server sends requested data to the client.
6. Client forwards the data to the application.

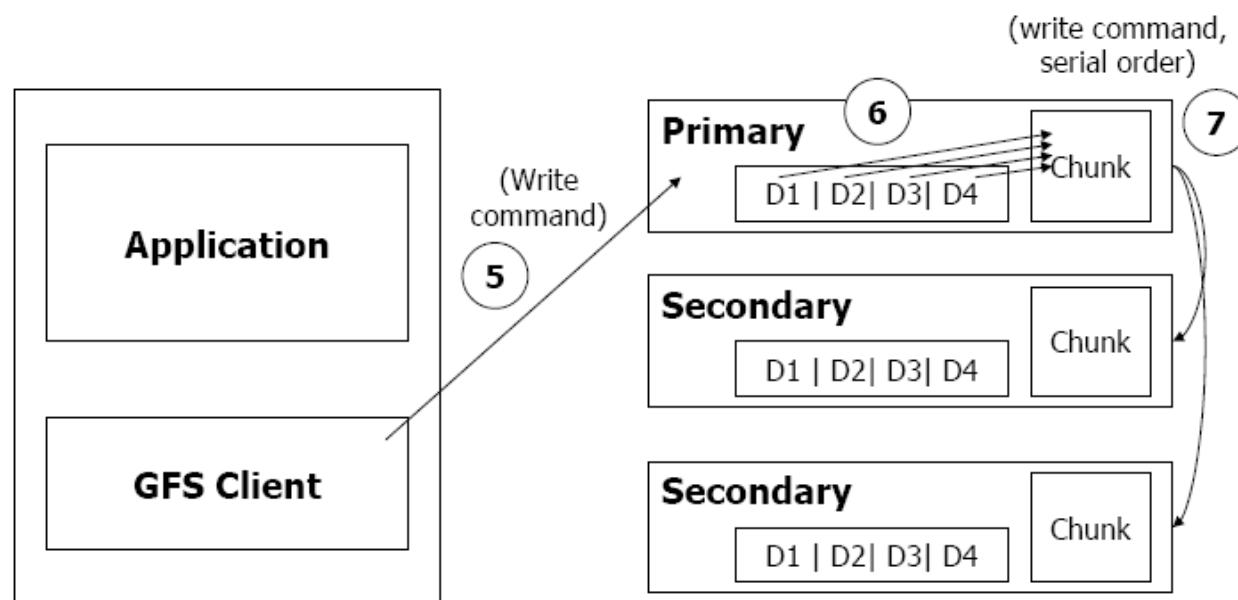
Write Algorithm



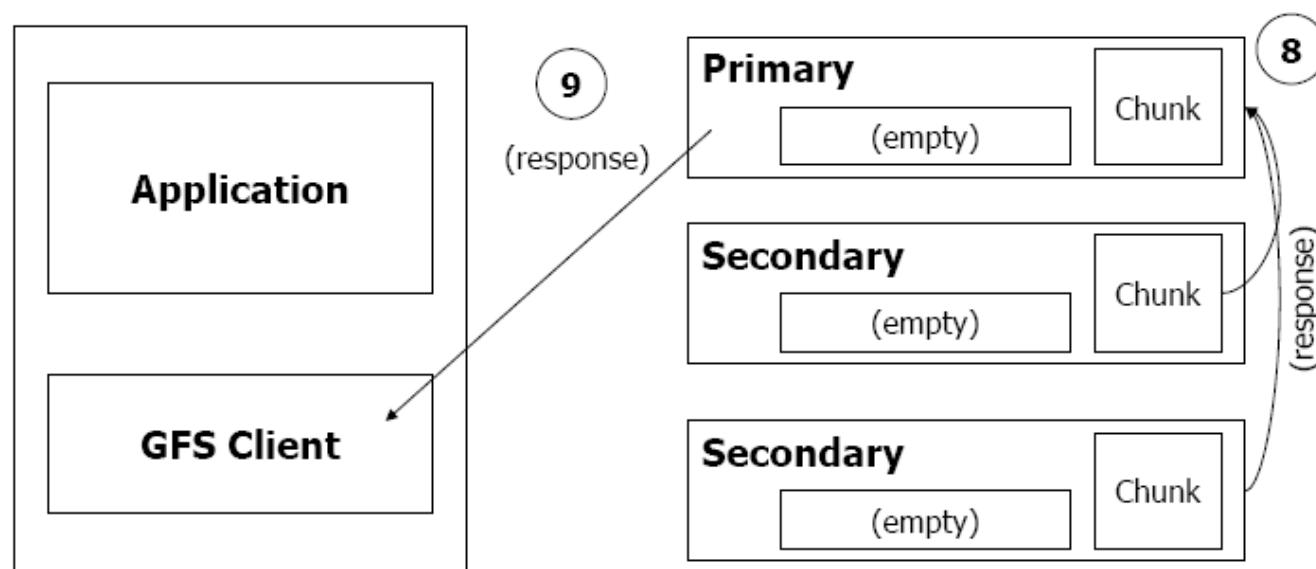
Write Algorithm



Write Algorithm



Write Algorithm



Write Algorithm

1. Application originates write request.
2. GFS client translates request from (filename, data) -> (filename, chunk index), and sends it to master.
3. Master responds with chunk handle and (primary + secondary) replica locations.
4. Client pushes write data to all locations. Data is stored in chunkservers' internal buffers.
5. Client sends write command to primary.

Write Algorithm

6. Primary determines serial order for data instances stored in its buffer and writes the instances in that order to the chunk.
7. Primary sends serial order to the secondaries and tells them to perform the write.
8. Secondaries respond to the primary.
9. Primary responds back to client.

Note: If write fails at one of chunk servers, client is informed and retries the write.

Master Operation

1. Replica Placement
2. Creation, Re-replication, Rebalancing
3. Garbage Collection
4. Stale Replica Detection

Fault Tolerance and Diagnosis

- Fast recovery
 - Master and chunkserver are designed to restore their states and start in seconds
- Chunk replication : 3-way mirrors
 - across multiple machines, across multiple racks.

Fault Tolerance and Diagnosis

- Master Mechanisms:
 - Log of all changes made to metadata.
 - Periodic checkpoints of the log.
 - Log and checkpoints replicated on multiple machines.
 - Master state is replicated on multiple machines.
 - “Shadow” masters for reading data if “real” master is down.

Fault Tolerance and Diagnosis

- Data integrity
 - A chunk is divided into 64-KB blocks
 - Each with its 32 bit checksum
 - Verified at read and write times
 - Also background scans for rarely used data

Summary of GFS

- Runs on commodity hardware and is scalable
- Performs well for the specified tasks and assumptions previously mentioned
- Innovation
 - File system API tailored to stylized workload
 - Single-master design to simplify coordination
 - Metadata fit in memory

Summary of GFS

- Metadata fit in memory
- Flat namespace
- Dedicated Care for Component Failure
 - hard disk failure, data corruption, network disconnection, etc.
- High-throughput
 - Minimized the master involvement
 - Chunk servers themselves send and receive the client data
 - The master leases authority to mutate chunks

Thanks

- Next class
 - Journaling

Journaling: xv6 and ext3

Yubin Xia & Rong Chen

Ack: slides adopted from Frans' 6.828 course

Outline

- Recap
- Journaling
- Xv6 journaling
- Ext3 journaling

Recap: File system durability

Topic: tension between fs perf. and crash recovery

Disk performance is often a #1 bottleneck
"how many seeks will that take?"

But many obvious fixes make crash / power failure hard to recover from
"what if a crash occurred at this point?"

Crash recovery is much harder than performance

Recap: Recovery approach

Synchronous meta-data update + fsck

- Used in xv6-rev0

- During check, synchronize meta-data, such as file size

Logging (ext 3/4), xv6-rev6 and following versions

- Before doing actual meta-data update, log the event

- After crash, recover from log

Soft update (FreeBSD fs modified on FFS)

- Soft update, not covered in this course

Recap: Terms for properties of fs ops

What effects will app see after restart + recovery

creat("a"); fd = creat("b"); write(fd,...); crash

Durable: effects of operation are visible

Both a and b are visible

Atomic: all steps of operation visible or none

Either a and b are visible or none is visible

Ordered: exactly a prefix of operations is visible

If b is visible, then a is visible

Recap: synchronous metadata update?

Xv6-rev0 strategy: carefully order disk writes to avoid dangling refs

1. initialize a new inode before creating dirent
2. delete dirent before marking inode free
3. mark block in-use before adding it to inode addrs[]
4. remove block from addrs[] before marking free
5. zero block before marking free

Has some visible bugs:

- . and .. during mkdir(), link counts, sizes

Has some invisible loose ends
may lose freed blocks and inodes

Logging (Journaling)

Goal:

- atomic system calls with respect to crashes
- fast recovery (no hour-long fsck)
- speed of write-back cache for normal operations

Approach

- will introduce logging in two steps
 - first xv6's log, which only provides safety
 - then Linux EXT3, which is also fast

Basic idea behind logging

To ensure atomicity:

- all of a system call's writes, or none

- let's call an atomic operation a "transaction"

- record all writes the sys call *will* do in the log

- then record "done"

- then do the writes

- on crash+recovery:

- if "done" in log, replay all writes in log

- if no "done", ignore log

- this is a WRITE-AHEAD LOG

Xv6's simple logging

FS has a log on disk syscall:

```
begin_trans()
```

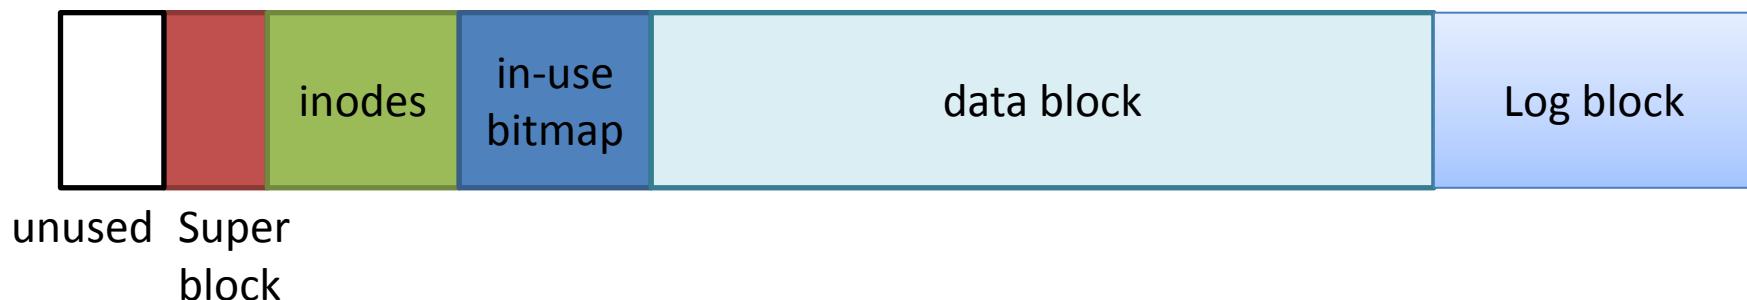
```
bp = bread()
```

```
bp->data[] = ...
```

```
log_write(bp)
```

```
more writes ...
```

```
commit_trans()
```



Recap: balloc

```
// Allocate a zeroed disk block.
static uint
balloc(uint dev)
{
    int b, bi, m;
    struct buf *bp;
    struct superblock sb;

    bp = 0;
    readsb(dev, &sb);
    for(b = 0; b < sb.size; b += BPB){
        bp = bread(dev, BBLLOCK(b, sb.ninodes));
        for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
            m = 1 << (bi % 8);
            if((bp->data[bi/8] & m) == 0){ // free?
                bp->data[bi/8] |= m;           Update bitmap
                log_write(bp);
                brelse(bp);
                bzero(dev, b + bi);
                return b + bi;
            }
        }
        brelse(bp);
    }
    panic("balloc: out of blocks");
}
```

Update bitmap

```
void
log_write(struct buf *b)
{
    int i;

    if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
        panic("too big a transaction");
    if (!log.busy)
        panic("write outside of trans");

    for (i = 0; i < log.lh.n; i++) {
        if (log.lh.sector[i] == b->sector) // log absorbtion?
            break;
    }
    log.lh.sector[i] = b->sector;
    struct buf *lbuf = bread(b->dev, log.start+i+1);
    memmove(lbuf->data, b->data, BSIZE);
    bwrite(lbuf);
    brelse(lbuf);
    if (i == log.lh.n)
        log.lh.n++;
    b->flags |= B_DIRTY; // XXX prevent eviction
}
```

Transaction Semantics in xv6

`begin_trans:`

- need to indicate which group of writes must be atomic!

- lock -- xv6 allows only one transaction at a time

`log_write:`

- record sector #

- append buffer content to log

- leave modified block in buffer cache (but do not write)

`commit_trans():`

- record "done" and sector #s in log

- do the writes

- erase "done" from log

`recovery:`

- if log says "done": copy blocks from log to real locations on disk

```
void
begin_trans(void)
{
    acquire(&log.lock);
    while (log.busy) {
        sleep(&log, &log.lock);
    }
    log.busy = 1;
    release(&log.lock);
}

void
commit_trans(void)
{
    if (log.lh.n > 0) {
        write_head(); // Write header to disk -- the real commit
        install_trans(); // Now install writes to home locations
        log.lh.n = 0;■
        write_head(); // Erase the transaction from the log
    }
    acquire(&log.lock);
    log.busy = 0;
    wakeup(&log);
    release(&log.lock);
}
```

```
static void
recover_from_log(void)
{
    read_head();■
    install_trans(); // if committed, copy from log to disk
    log.lh.n = 0;
    write_head(); // clear the log
}
```

```
// Copy committed blocks from log to their home location
static void■
install_trans(void)
{
    int tail;

    for (tail = 0; tail < log.lh.n; tail++) {
        struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block
        struct buf *dbuf = bread(log.dev, log.lh.sector[tail]); // read dst
        memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst
        bwrite(dbuf); // write dst to disk
        brelse(lbuf);■
        brelse(dbuf);
    }
}
```

Sys_unlink

```
struct inode *dp =
    nameiparent(path, name)
```

```
begin_trans();

ilock(dp);

// Cannot unlink .. or ...
if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
    goto bad;

if((ip = dirlookup(dp, name, &off)) == 0)
    goto bad;
ilock(ip);

if(ip->nlink < 1)
    panic("unlink: nlink < 1");
if(ip->type == T_DIR && !isdirempty(ip)){
    iunlockput(ip);
    goto bad;
}

memset(&de, 0, sizeof(de));
if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
    panic("unlink: writei");
if(ip->type == T_DIR){
    dp->nlink--;
    iupdate(dp);
}
iunlockput(dp);

ip->nlink--;
iupdate(ip);
iunlockput(ip);

commit_trans();
```

Limitation with xv6 logging

Only one transaction at a time

two system calls might be modifying different parts of the FS
synchronous write to on-disk log

each write takes one disk rotation time

commit takes another

a file create/delete involves around 10 writes

thus 100 ms per create/delete -- very slow!

tiny update -> whole block write

creating a file only dirties a few dozen bytes

but produces many kilobytes of log writes

synchronous writes to home locations after commit

i.e. write-through, not write-back

makes poor use of in-memory disk cache

Questions

How can we get both performance and safety?

we'd like system calls to proceed at in-memory speeds
using write-back disk cache
i.e. have typical system call complete w/o actual disk
writes

Linux's ext3's Journaling

case study of the details required to add logging to a file system

Stephen Tweedie 2000 talk transcript "EXT3, Journaling Filesystem"

ext3 adds a log to ext2, a previous xv6-like log-less file system

has many modes, start with "journaled data"
log contains both metadata and file content blocks

Ext3 Structures

in-memory write-back block cache

in-memory list of blocks to be logged, per-transaction

on-disk FS

on-disk circular log file

What's in the ext3 log?

superblock: starting offset and starting seq #

descriptor blocks: magic, seq, block #s

data blocks (as described by descriptor)

commit blocks: magic, seq

How does ext3 get good perf.?

batches many syscalls per commit

defers copying cache block to log until it commits
log to disk

hopes multiple syscalls modified same block
thus many syscalls but only one copy of block in log
"write absorption"

Syscall process

h = start()

get(h, block #)

 warn logging system we'll modify cached block

 added to list of blocks to be logged

 prevent writing block to disk until after xaction commits

modify the blocks in the cache

stop(h)

guarantee: all or none

stop() does *not* cause a commit

Ext3 transaction

[circle set of cache blocks in this transaction]

while "open", adds new syscall handles, and remembers their block #s

only one open transaction at a time

ext3 commits current transaction every few seconds (or fsync())

Committing a transaction to disk

open a new transaction, for subsequent syscalls

mark transaction as done

wait for in-progress syscalls to stop()

(maybe it starts writing blocks, then waits, then writes again if needed)

write descriptor to log on disk w/ list of block #s

write each block from cache to log on disk

wait for all log writes to finish

append the commit record

now cached blocks allowed to go to homes on disk (but not forced)

Is log correct if concurrent syscalls?

e.g. create of "a" and "b" in same directory

inode lock prevents race when updating directory

other stuff can be truly concurrent (touches
different blocks in cache)

transaction combines updates of both system calls

What if syscall B reads uncommitted result of syscall A?

A: echo hi > x

B: ls > y

Could B commit before A, so that crash would reveal anomaly?

case 1: both in same xaction -- ok, both or neither

case 2: A in T1, B in T2 -- ok, A must commit first

case 3: B in T1, A in T2

could B see A's modification?

ext3 must wait for all ops in prev xaction to finish

before letting any in next start

so that ops in old xaction don't read modifications of next xaction

Performance

create 100 small files in a directory

would take xv6 over 10 seconds (many disk writes per syscall)

repeated modifications to same direntry, inode, bitmap
blocks in cache

write absorption...

then one commit of a few metadata blocks plus 100 file
blocks

how long to do a commit?

seq write of 100×4096 at 50 MB/sec: 10 ms

wait for disk to say “writes are on disk”

then write the commit record

that wastes one rotation, another 10 ms

modern disk interfaces can avoid wasted rotation

What if a crash?

crash may interrupt writing last xaction to log on disk

so disk may have a bunch of full xactions, then maybe one partial

may have written some block cache to disk (home)
but only for fully committed xactions, not partial last one

How does recovery work?

1. find the start and end of the log

log "superblock" at start of log file

log superblock has start offset and seq# of first transaction

scan until bad record or not the expected seq #

go back to last commit record

crash during commit -> last transaction ignored during recovery

2. replay all blocks through last completed xaction, in log order

How does recovery work?

what if block after last valid log block looks like a log descriptor?

perhaps left over from previous use of log? (**seq...**)

perhaps some file data happens to look like a descriptor?
(magic #...)

when can ext3 free a transaction's log space?

after cached blocks have been written to FS on disk
free == advance log superblock's start pointer/seq

what if block in T1 has been dirtied in cache by T2?

- can't write that block to FS on disk
- note ext3 only does copy-on-write while T1 is committing
 - after T1 commit, T2 dirties only block copy in cache
- so can't free T1 until T2 commits, so block is in log
 - T2's logged block contains T1's changes

Question

what if not enough free space in log for a syscall?

suppose we start adding syscall's blocks to T2
half way through, realize T2 won't fit on disk

we cannot commit T2, since syscall not done

can we free T1 to free up log space?

maybe not, due to previous issue, T2 maybe dirtied a
block in T1

deadlock!

Solution: Reservations

syscall pre-declares how many block of log space it might need

block the syscall from starting until enough free space

may need to commit open transaction, then free older transaction

OK since reservations mean all started sys calls can complete + commit

Durability of ext3

ext3 not as immediately durable as xv6

creat() returns -> maybe data is not on disk! crash will undo it.

need fsync(fd) to force commit of current transaction, and wait

would ext3 have good performance if commit after every sys call?

would log many more blocks, no absorption

10 ms per syscall, rather than 0 ms

("Rethink the Sync" [OSDI'06] addresses this problem)

Issues with ext3

disks usually have write caches and re-order writes,
for performance

- sometimes hard to turn off (the disk lies)

- often with re-ordering enabled for speed, out of ignorance

bad news if disk writes commit block before
preceding stuff

- then recovery replays "descriptors" with random block #s!
and writes them with random content!

how to solve this problem?

- using checksum

Ordered Mode vs Journaled Mode

journaling file content is slow, every data block written twice

perhaps not needed to keep FS internally consistent

can we just lazily write file content blocks?

no:

if metadata updated first, crash may leave file pointing to blocks with someone else's data

ext3 ordered mode:

write content block to disk before committing inode w/ new block #

thus won't see stale data if there's a crash

most people use ext3 ordered mode

Xv6 vs. ext3

does ext3 fix the xv6 log performance problems?

only one transaction at a time -- yes

synchronous write to on-disk log -- yes, but 5-second window

tiny update -> whole block write -- yes (indirectly)

synchronous writes to home locations after commit – yes

ext3 very successful

but: no checksum -- ext4

but: not efficient for applications that use fsync()

Questions

- In `commit_trans`, there are three steps
 - `Write_head`: log “done” to disk
 - `Install_trans`: checkpointing metadata to home location
- Can these two operations get reordered or mixed together?
- How can we make them reordered?

```
void
commit_trans(void)
{
    if (log.lh.n > 0) {
        write_head();      // Write header to disk -- the real commit
        install_trans();  // Now install writes to home locations
        log.lh.n = 0;■
        write_head();      // Erase the transaction from the log
    }
}
```

T2 starts while T1 is committing to log on disk

- what if syscall in T2 wants to write block in prev xaction?
- can't be allowed to write buffer that T1 is writing to disk
 - then new syscall's write would be part of T1
 - crash after T1 commit, before T2, would expose update
- T2 gets a separate copy of the block to modify
 - T1 holds onto old copy to write to log
- are there now *two* versions of the block in the buffer cache?
 - no, only the new one is in the buffer cache, the old one isn't
- does old copy need to be written to FS on disk?
 - no: T2 will write it

Correctness w/ ordered mode

- A. rmdir, re-use block for file, ordered write of file, crash before rmdir or write committed
 - now scribbled over the directory block

fix: defer free of block until freeing operation forced to log on disk

Correctness w/ ordered mode

B. rmdir, commit, re-use block in file, ordered file
write, commit,

crash, replay rmdir

file is left w/ directory content e.g. . and ..

fix: revoke records, prevent log replay of a given
block

Final Tidbit

- open a file, then unlink it
 - unlink commits
 - file is open, so unlink removes dir ent but doesn't free blocks
 - crash
 - nothing interesting in log to replay
 - inode and blocks not on free list, also not reachable by any name
- will never be freed! oops
- solution: add inode to linked list starting from FS superblock commit that along with remove of dir ent
 - recovery looks at that list, completes deletions

Scalable Locking

Rong Chen

Institute of Parallel and Distributed Systems

Shanghai Jiao Tong University

http://ipads.se.sjtu.edu.cn/haibo_chen

Some slides adjusted from. Elsa L Gunter (UIUC), Jonathan Walpole (PSU)
Paul McKenney (IBM) Tom Hart (University of Toronto), Frans Kaashoek (MIT)

Outline

Scalability Tutorial

Non-scalable locks

Scalable locks

SCALABILITY TUTORIALS

What is scalability?

Application does N times as much work on N cores
as it could on 1 core

Scalability may be limited by Amdahl's Law:

Locks, shared data structures, ... Shared hardware
(DRAM, NIC, ...)

$$S_{\text{latency}}(s) = \frac{1}{1 - p + \frac{p}{s}}$$

Locking

Why do kernels normally use locks?

Locks support a concurrent programming style based on mutual exclusion

- Acquire lock on entry to critical sections

- Release lock on exit

- Block or spin if lock is held

- Only one thread at a time executes the critical section

Locks prevent concurrent access and enable sequential reasoning about critical section code

Sample: SEND & RECEIVE

Bounded Buffer

Shared between modules
Producers and consumers

Race Condition

If multiple writers
Using lock to ensure single writer principle

Locking

ACQUIRE & RELEASE
Need atomicity
RSM: read-set-memory

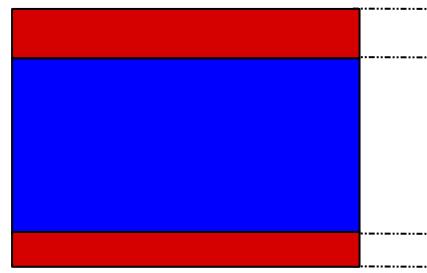
Amdahl's Law

$$\frac{1}{(1 - P) + \frac{P}{S}}$$

Two independent parts **A** **B**



Critical-section efficiency



$$\text{Critical-section efficiency} = \frac{T_c}{T_c + T_a + T_r}$$

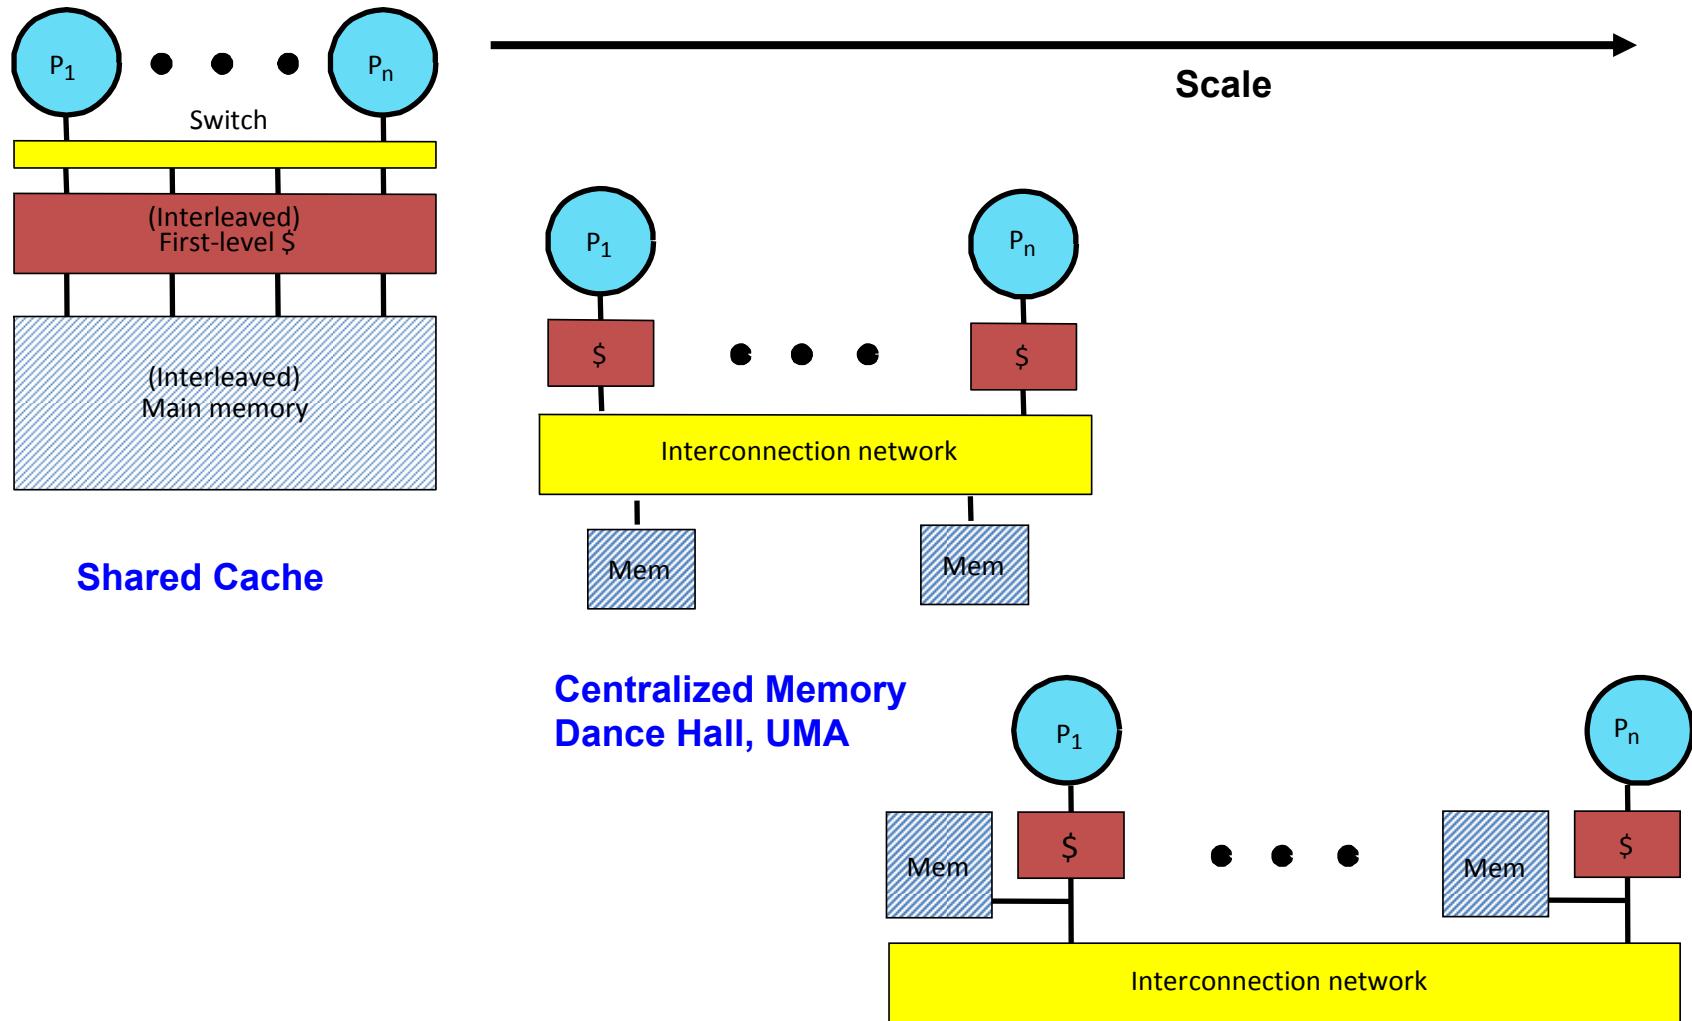
Ignoring lock contention and cache conflicts in the critical section

Some numbers you may want to know

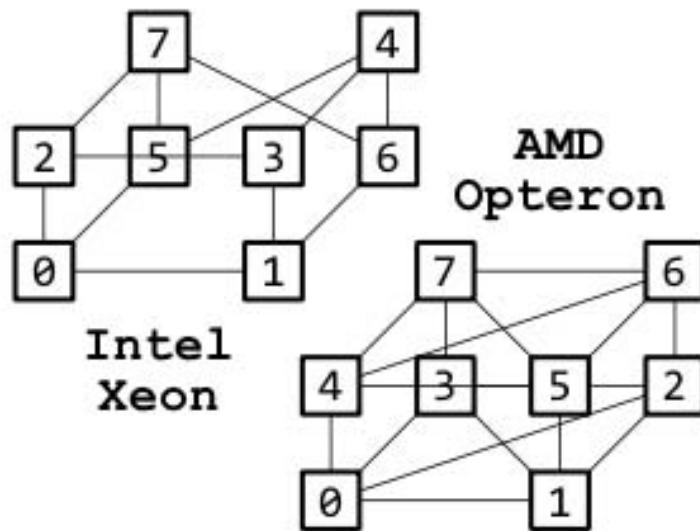
	Opteron	Xeon	Niagara	Tilera
L1	3	5	3	2
L2	15	11		11
LLC	40	44	24	45

Source: *Everything you always wanted to know about synchronization, but were afraid to ask.*

Shared Memory Multiprocessors



More numbers you may want to know



(a) System topology

Inst.	0-hop	1-hop	2-hop
80-core Intel Xeon machine			
Load	117	271	372
Store	108	304	409
64-core AMD Opteron machine			
Load	228	419	498
Store	256	463	544

(b) Latencies (cycles) on the distance

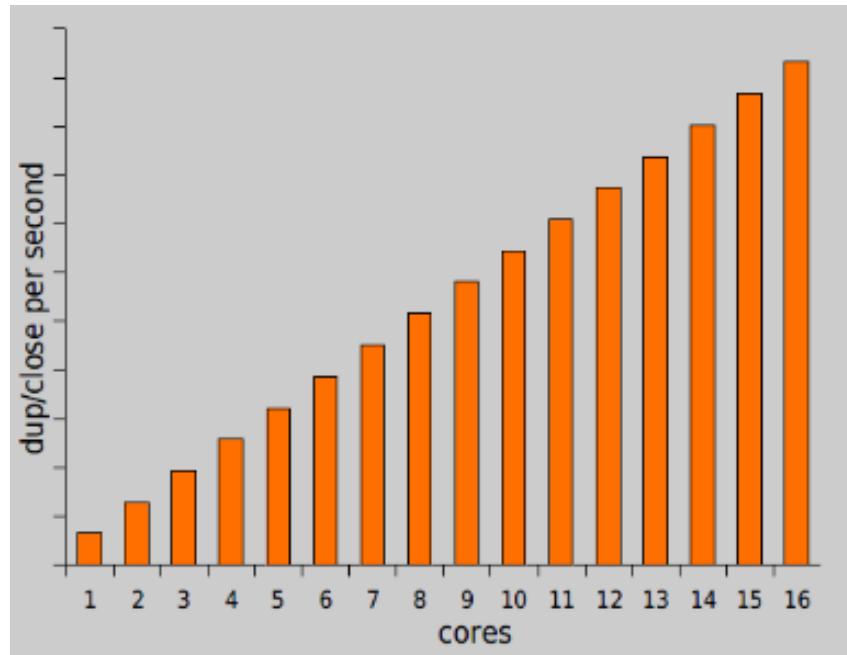
```

fd_alloc(void) {
    lock(fd_table);
    fd = get_free_fd();
    set_fd_used(fd);
    fix_smallest_fd()
    unlock(fd_table);
}

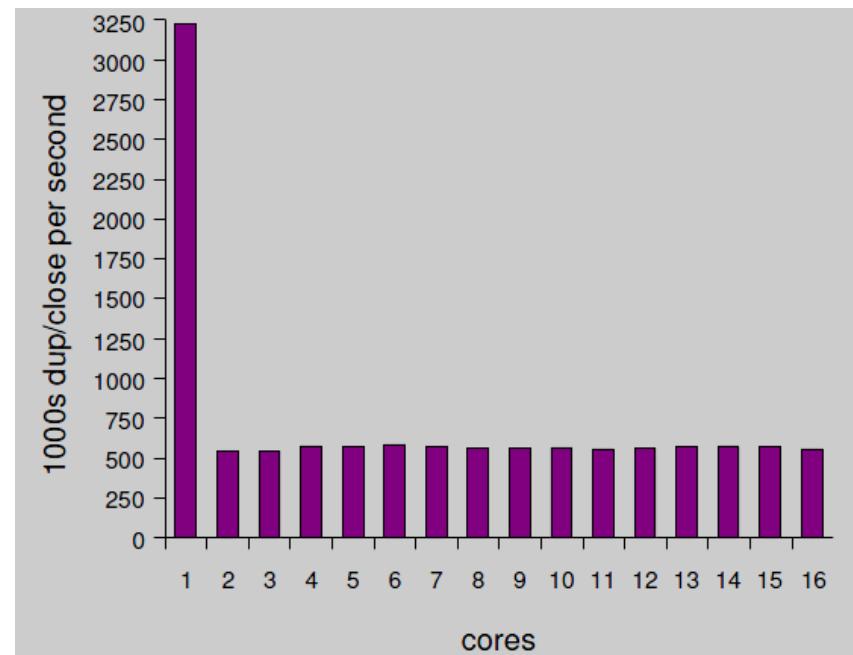
```

Motivating example: file descriptors

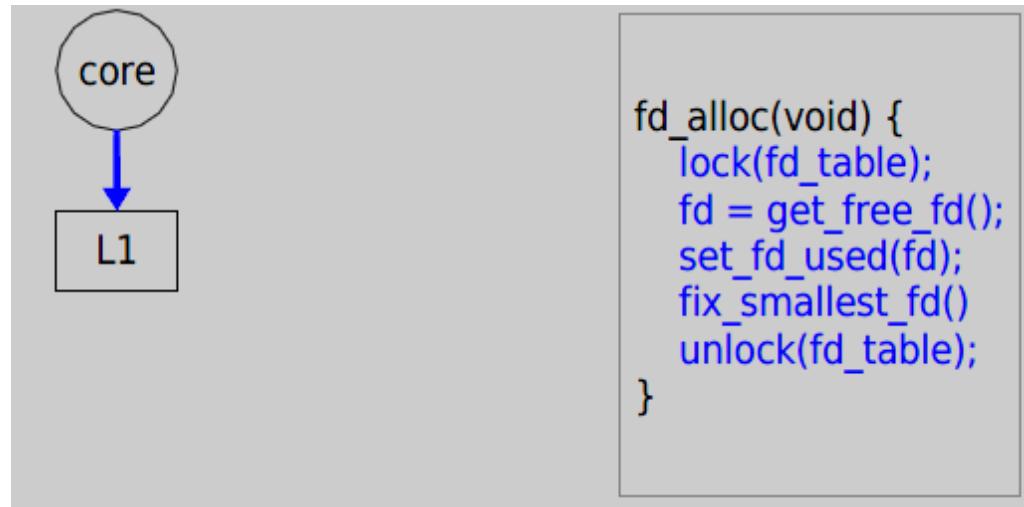
Ideal FD performance graph



Actual FD performance

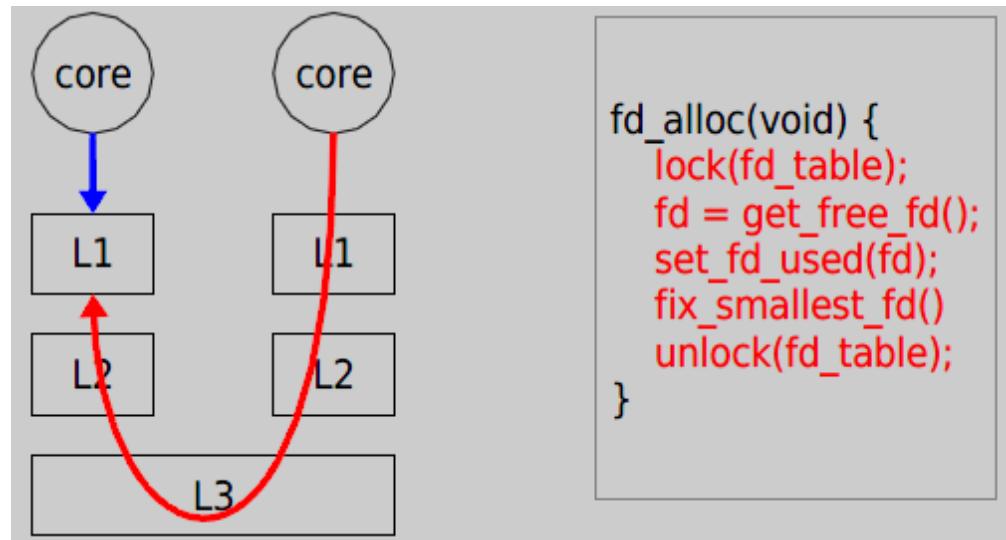


Why throughput drops?



Load fd_table data from L1 in 3 cycles.

Why throughput drops?



Now it takes 121 cycles!

**NON-SCALABLE LOCKING
ARE DANGEROUS**

Cause: Non-scalable locks

Non-scalable locks

Such as spin locks

Poor performance when highly contended

Many systems are using non-scalable locks

But they are dangerous

Why they are dangerous

Lead to performance collapse when adding a few more cores

Even tiny critical section will also lead to this performance collapse

```

// Acquire the lock.
// Loops (spins) until the lock is acquired.
// Holding a lock for a long time may cause
// other CPUs to waste time spinning to acquire it.
void
acquire(struct spinlock *lk)
{
    pushcli(); // disable interrupts to avoid deadlock.
    if(holding(lk))
        panic("acquire");

    // The xchg is atomic.
    // It also serializes, so that reads after acquire are not
    // reordered before it.
    while(xchg(&lk->locked, 1) != 0)
        ;

    // Record info about lock acquisition for debugging.
    lk->cpu = cpu;
    getcallerpcs(&lk, lk->pcs);           // Release the lock.
}

void
release(struct spinlock *lk)
{
    if(!holding(lk))
        panic("release");

    lk->pcs[0] = 0;
    lk->cpu = 0;

    // The xchg serializes, so that reads before release are
    // not reordered after it. The 1996 PentiumPro manual (Volume 3,
    // 7.2) says reads can be carried out speculatively and in
    // any order, which implies we need to serialize here.
    // But the 2007 Intel 64 Architecture Memory Ordering White
    // Paper says that Intel 64 and IA-32 will not move a load
    // after a store. So lock->locked = 0 would work here.
    // The xchg being asm volatile ensures gcc emits it after
    // the above assignments (and after the critical section).
    xchg(&lk->locked, 0);

    popcli();
}

```

Xv6 locking

Case study: ticket spinlock

Normal spinlock has extremely noticeable unfairness

- Even on a 8 core CPU

Ticket spinlock guarantees lock are granted to acquirers in order

- Used in Linux kernel

But it's non scalable lock

Pseudo code for ticket lock

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}

void spin_lock(spinlock_t *l) {
    int t = atomic_xadd(&l->next_ticket);
    while (t != l->current_ticket)
        ; // spin
}

void spin_unlock(spinlock_t *l) {
    l->current_ticket++;
}
```

Pseudo code for ticket lock

```
struct spinlock_t {  
    int current_ticket;           ← Currently serving which one  
    int next_ticket;             ← Ticket for the next comer  
}  
  
void spin_lock(spinlock_t *l) {  
    int t = atomic_xadd(&l->next_ticket);           ← Add and get  
    while (t != l->current_ticket)  
        ; // spin  
}  
  
void spin_unlock(spinlock_t *l) {  
    l->current_ticket++;  
}
```

Questions

How do current_ticket and next_ticket work together to guarantee mutual exclusion?

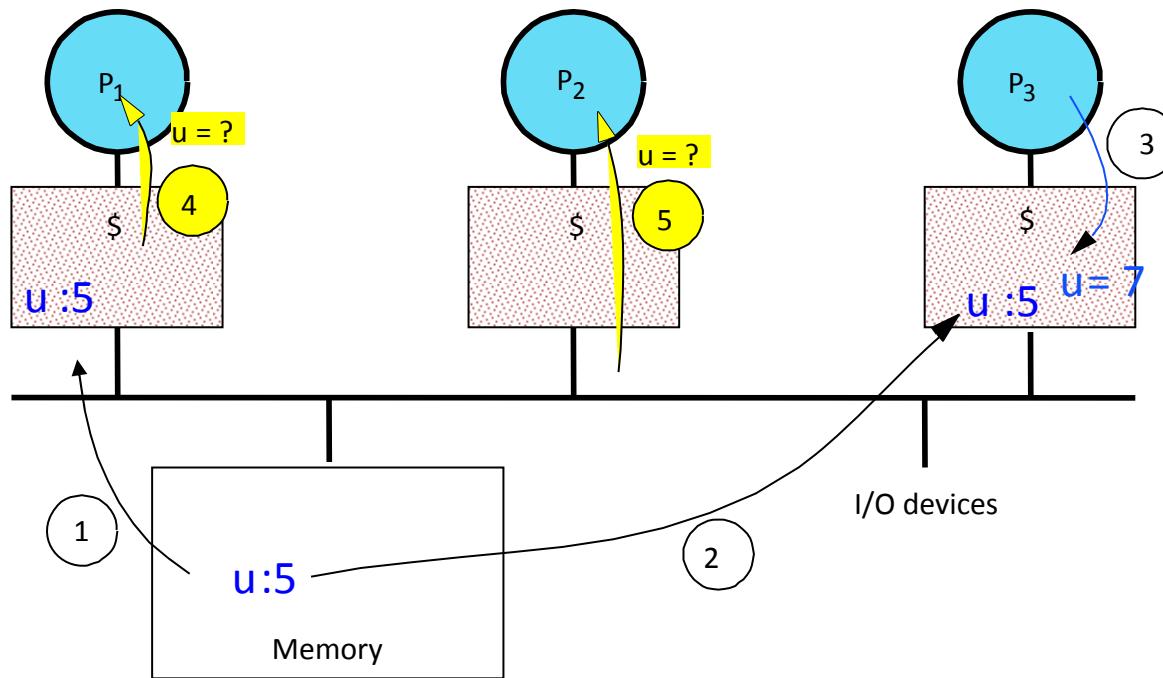
Background on cache coherence

On multi-core processors, each core has it's own private cache

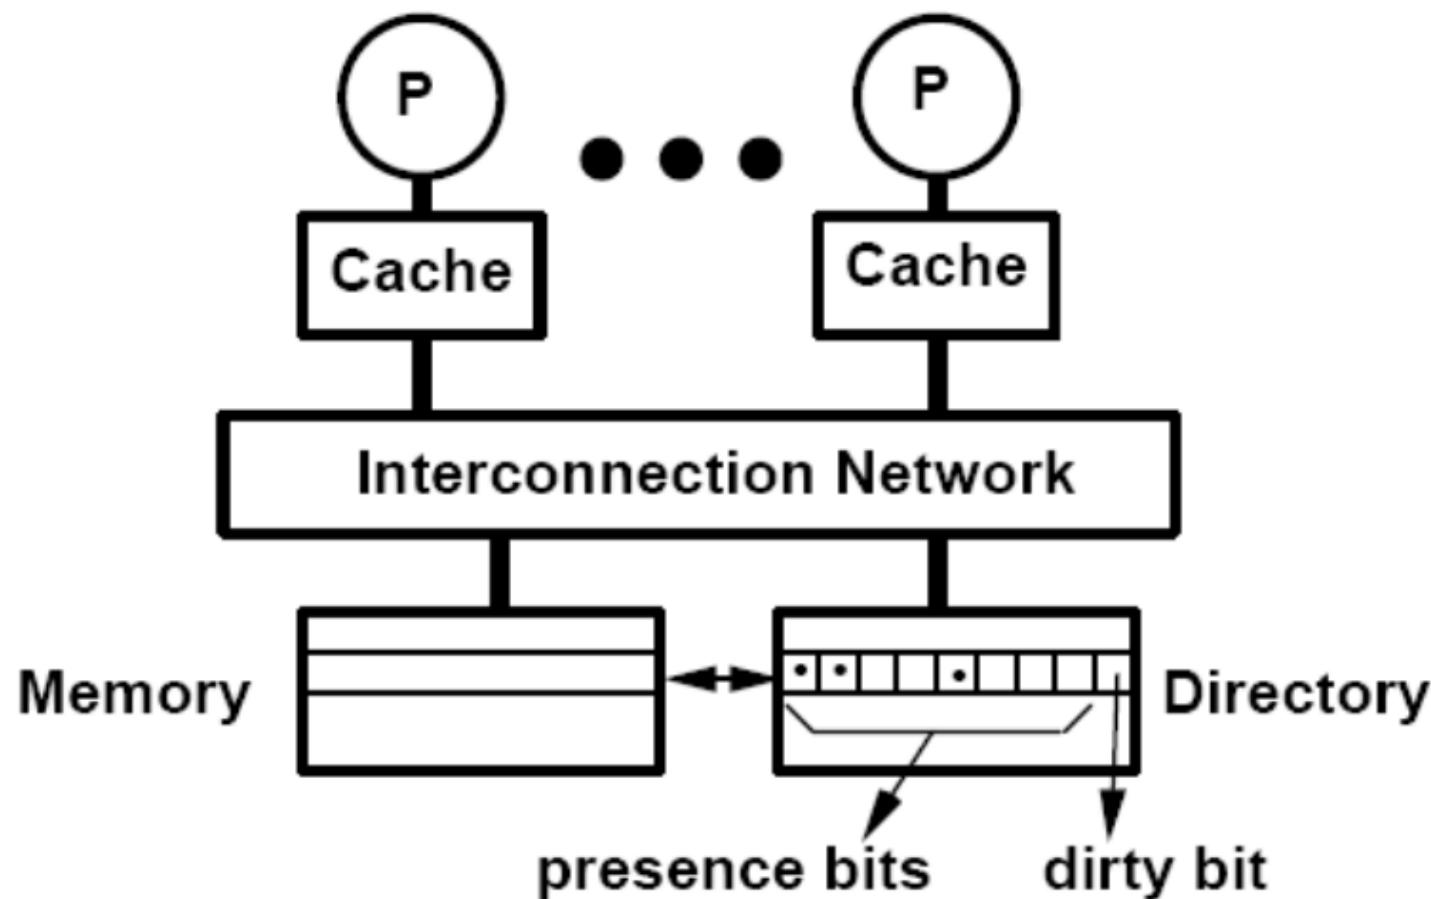
Need to keep the cache content in sync when some CPU modifies some memory

Accomplished by cache coherence protocol

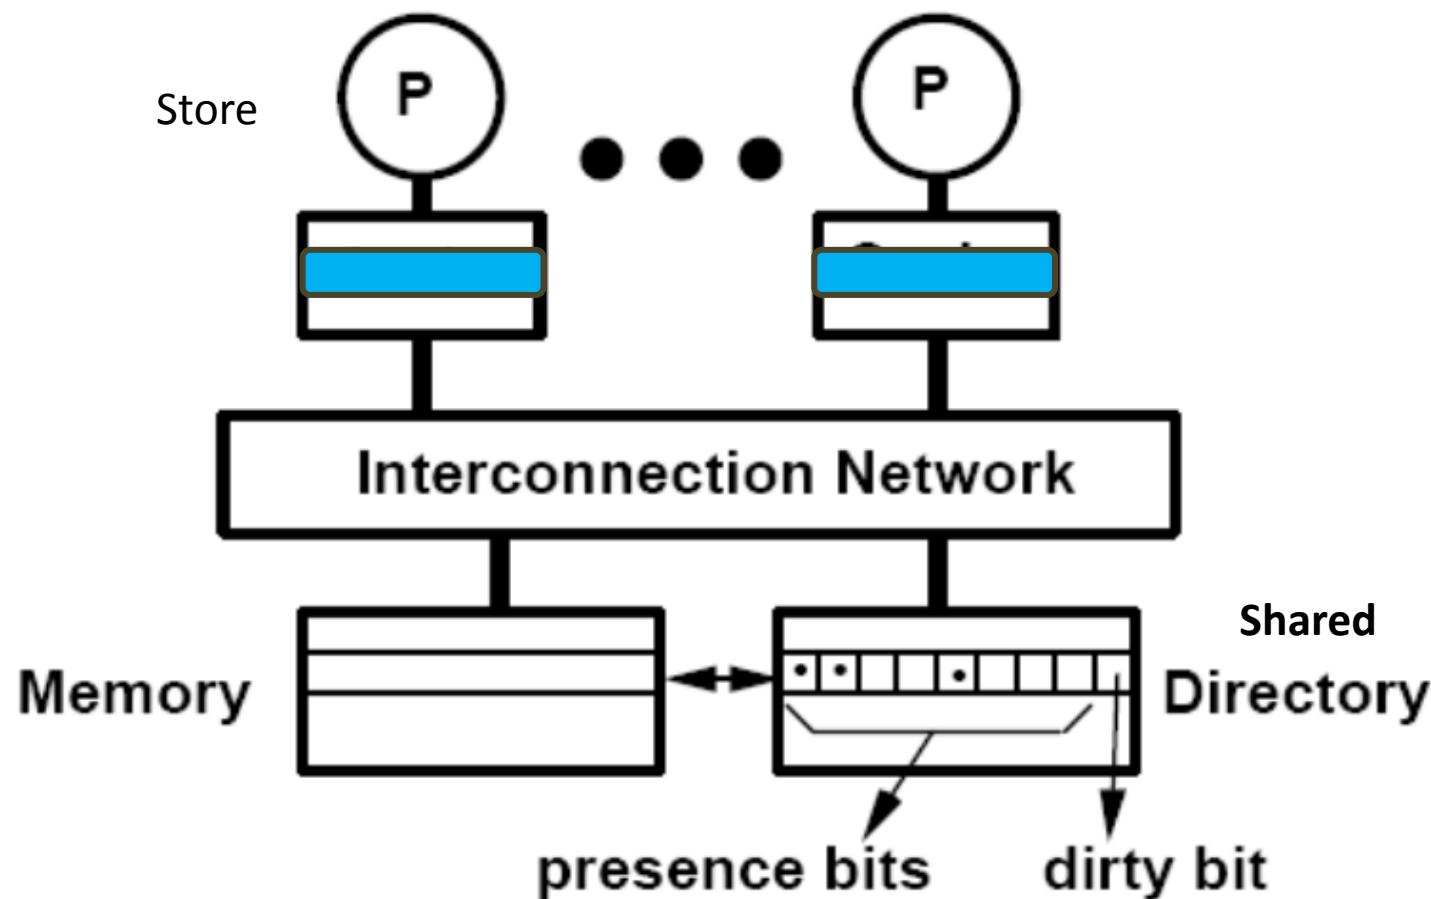
Example Cache Coherence Problem



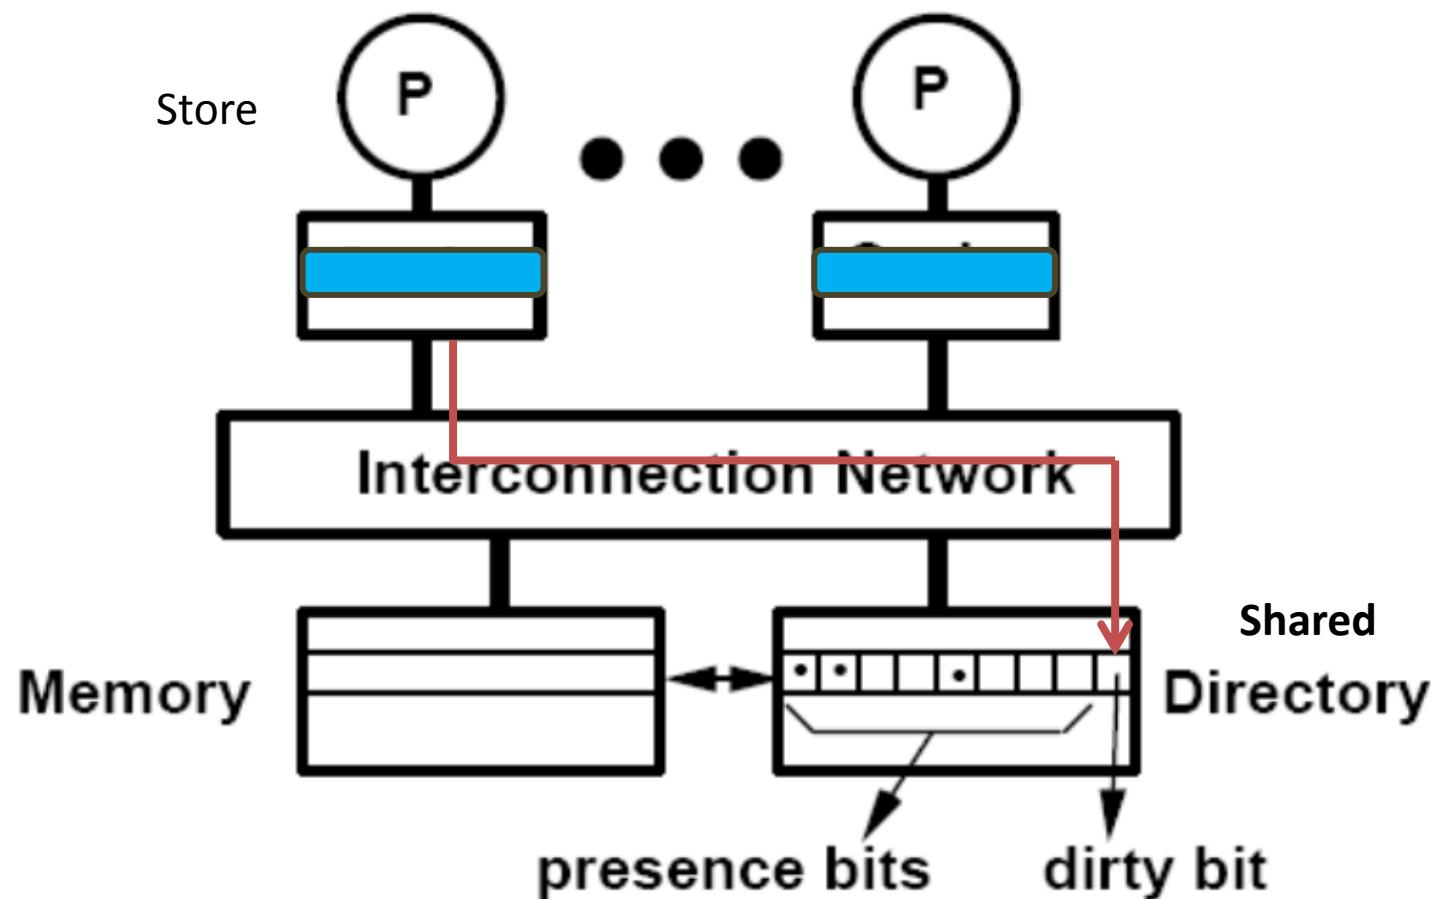
Directory-based cache coherence



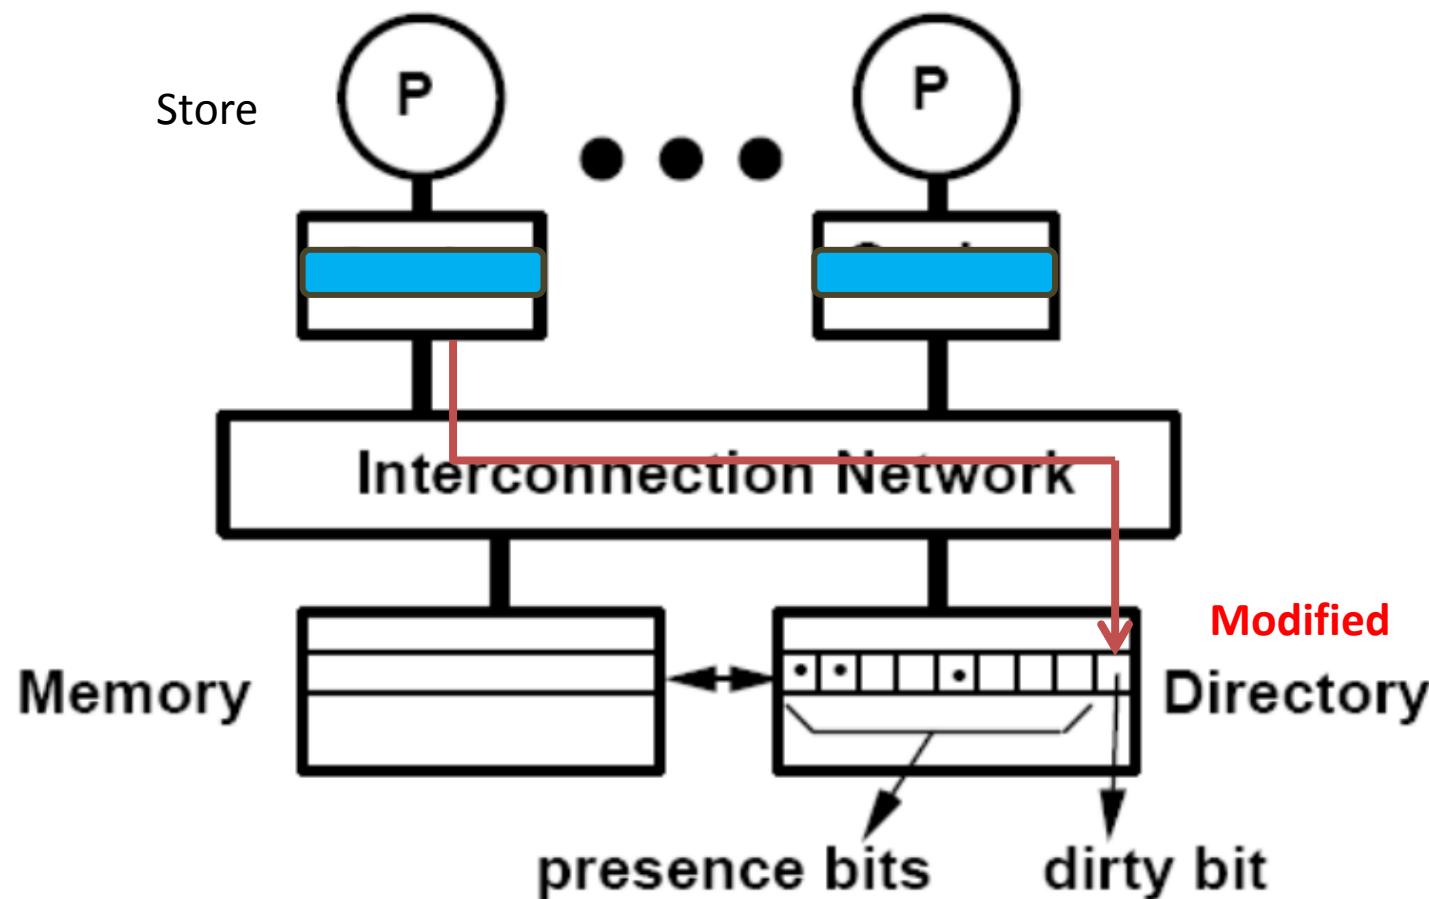
Directory-based cache coherence



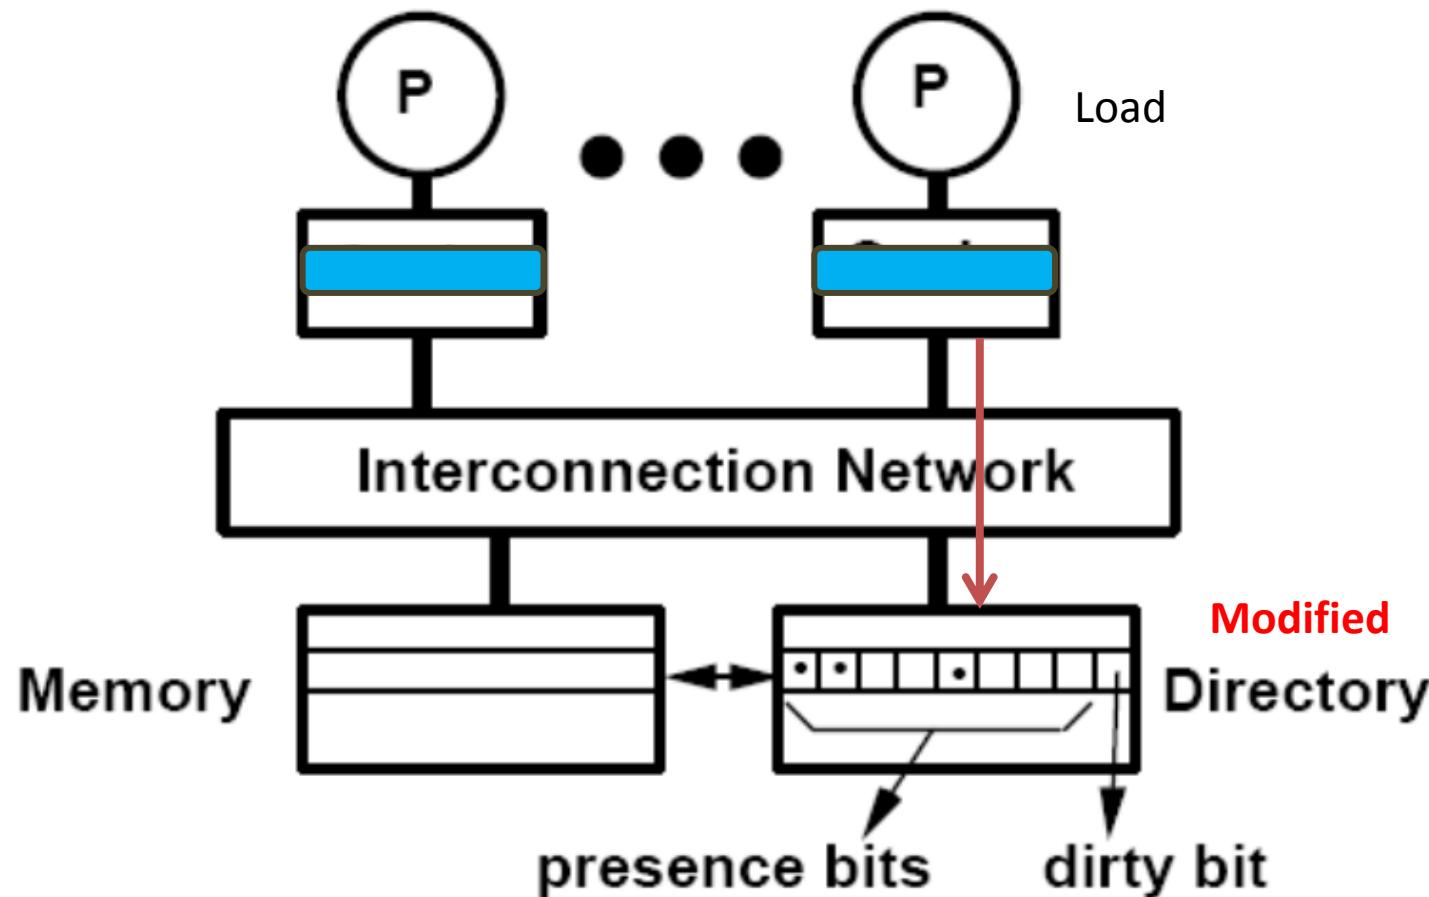
Directory-based cache coherence



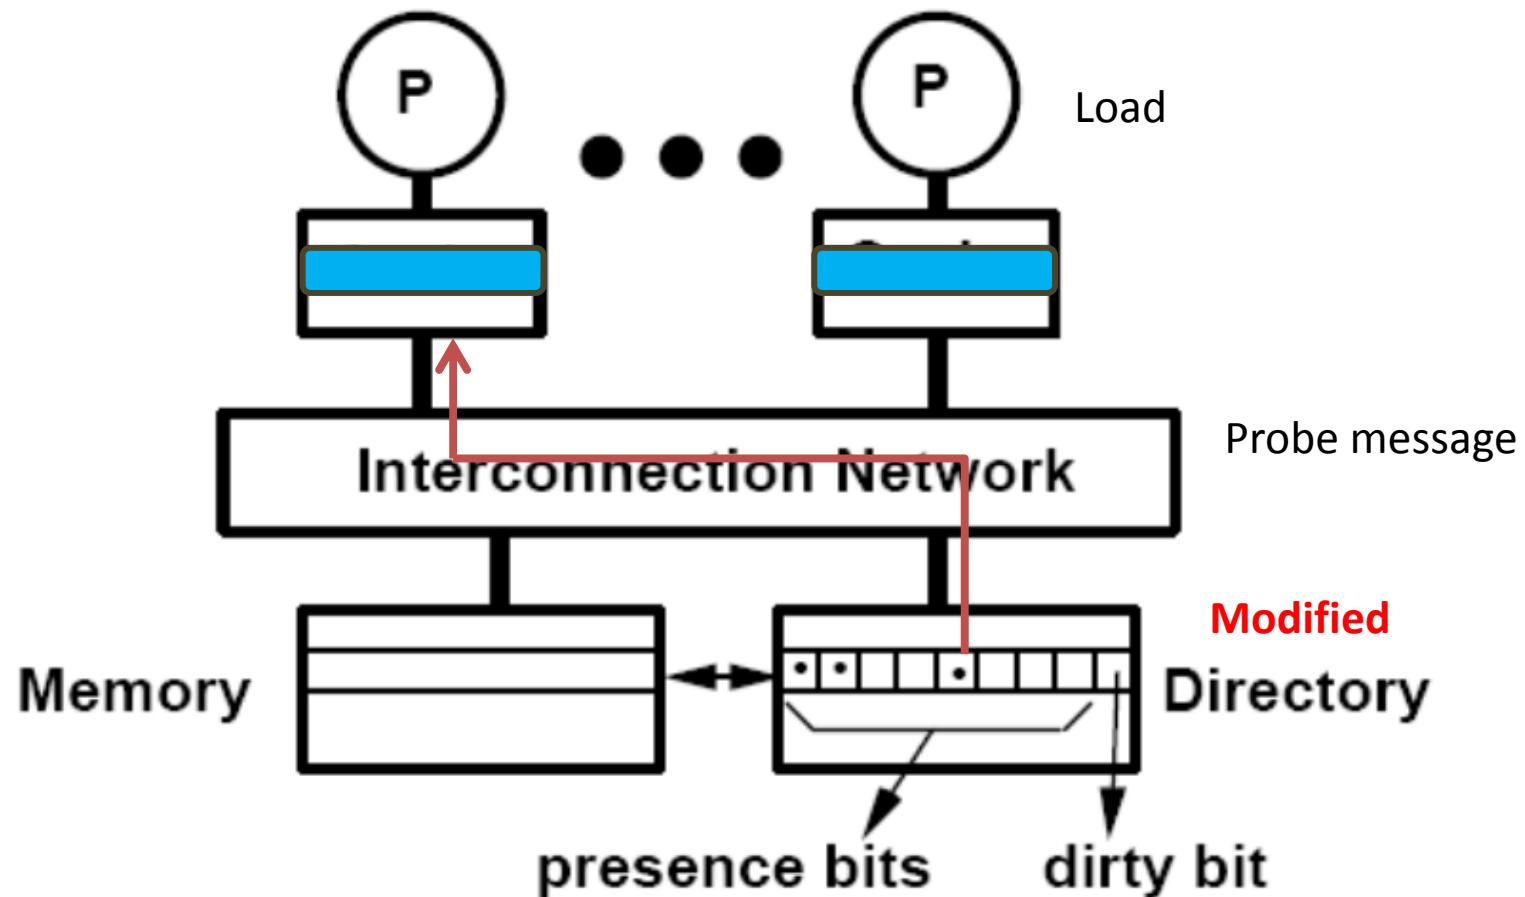
Directory-based cache coherence



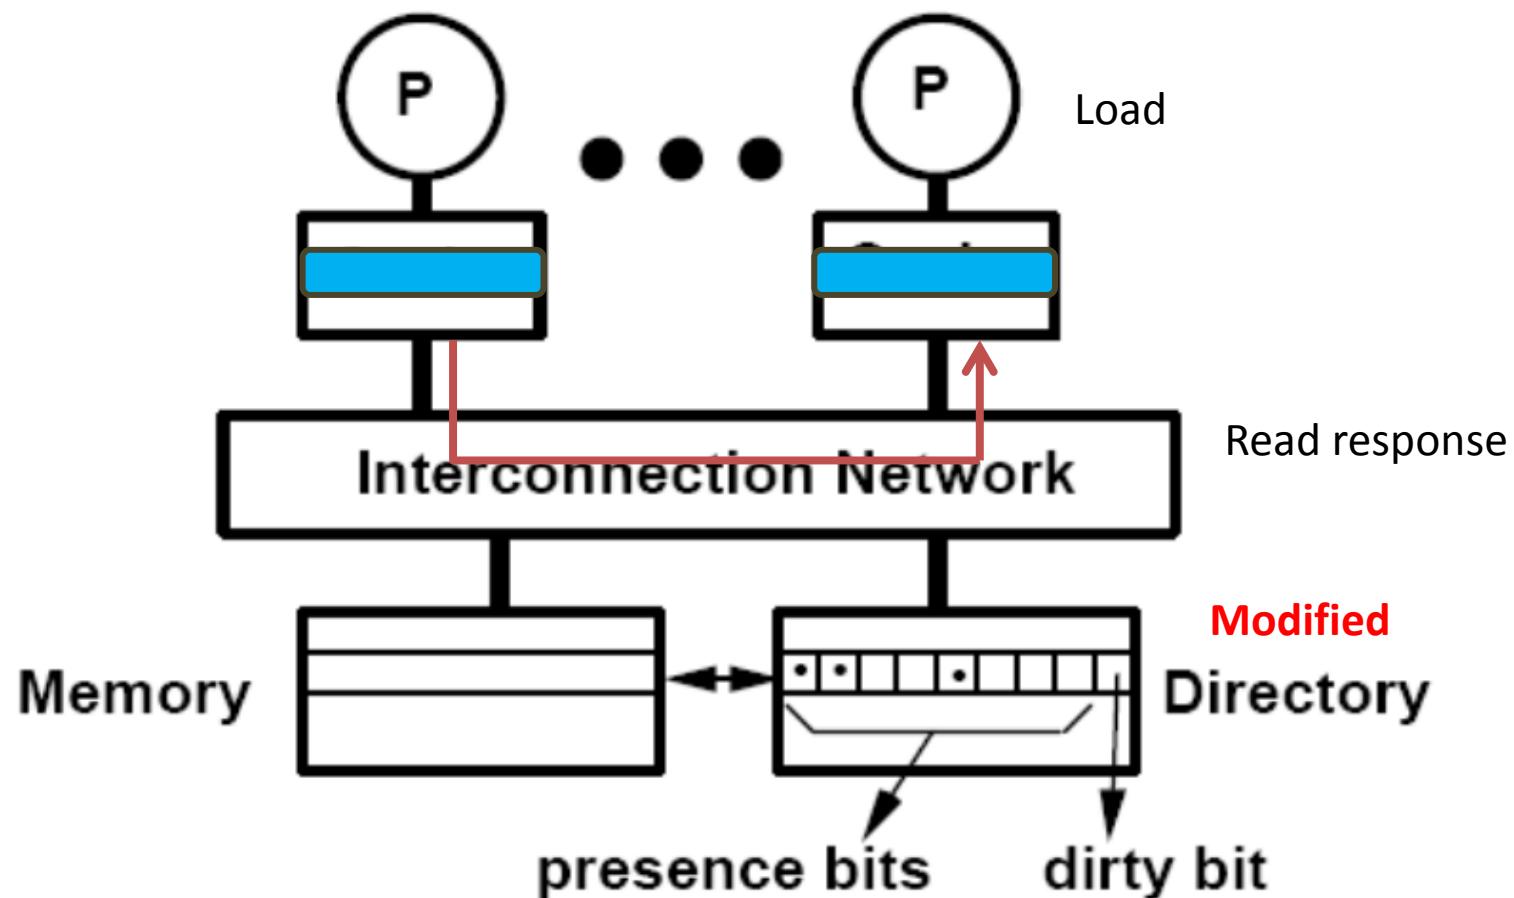
Directory-based cache coherence



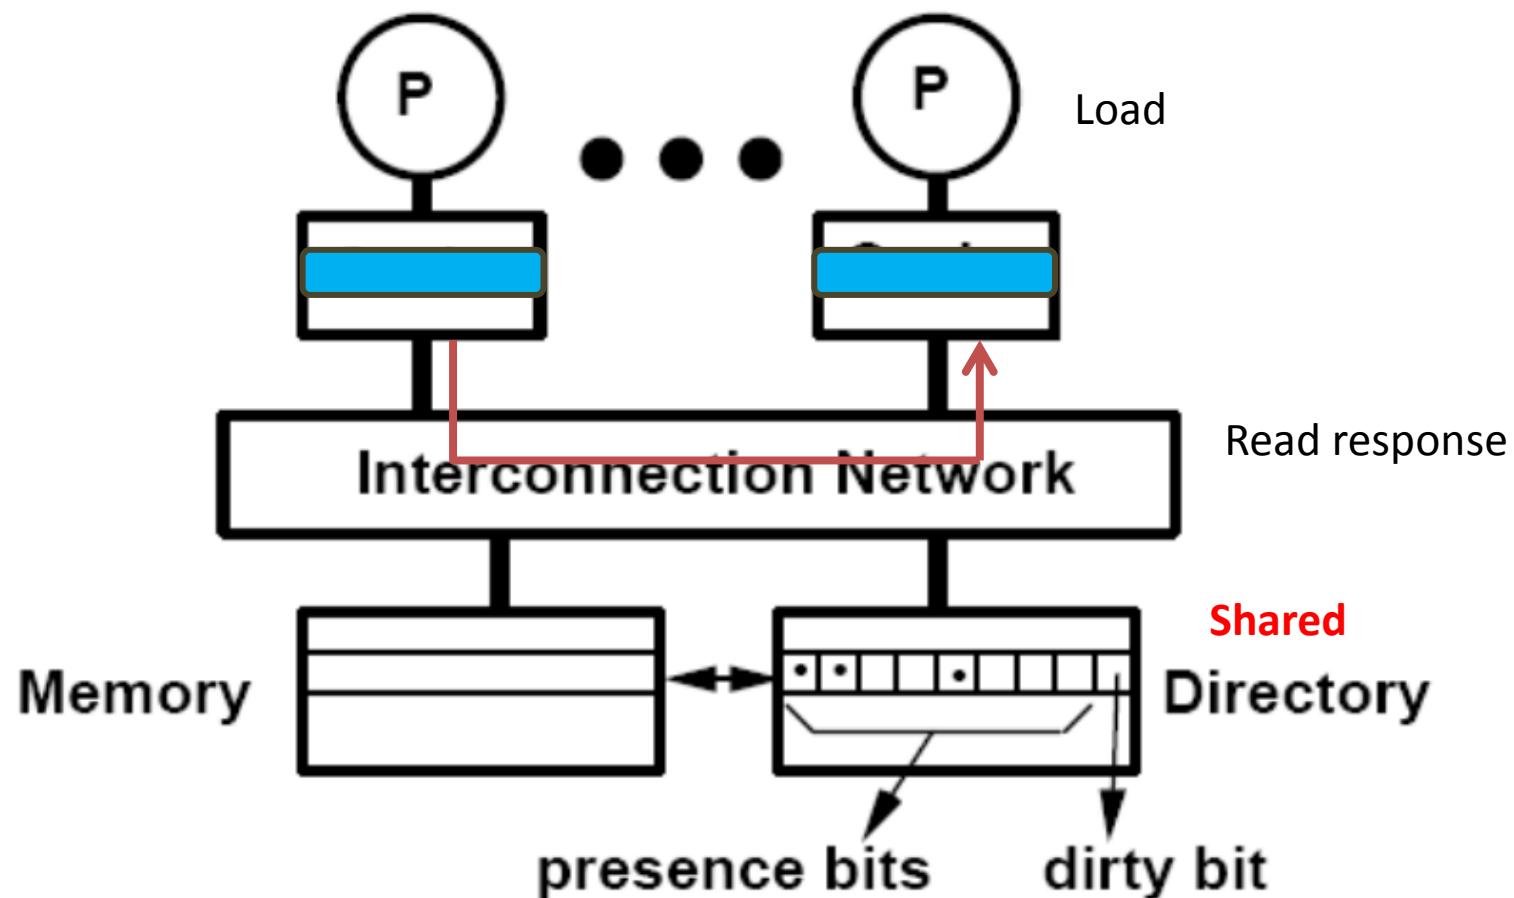
Directory-based cache coherence



Directory-based cache coherence



Directory-based cache coherence



A few notes cache coherence

There may be more than a single directory

Especially for NUMA systems

Interconnect structure affects cache coherence performance

Directory is just one possible implementation

Snooping is another commonly used approach

Intuition of the collapse

Key point: read with modified cache line have to get data back from the owner

Coherence message are processed **sequentially**

Lock holder modifies cache holding the lock

Waiter is trying to read the lock

They get value of the lock from the lock holder

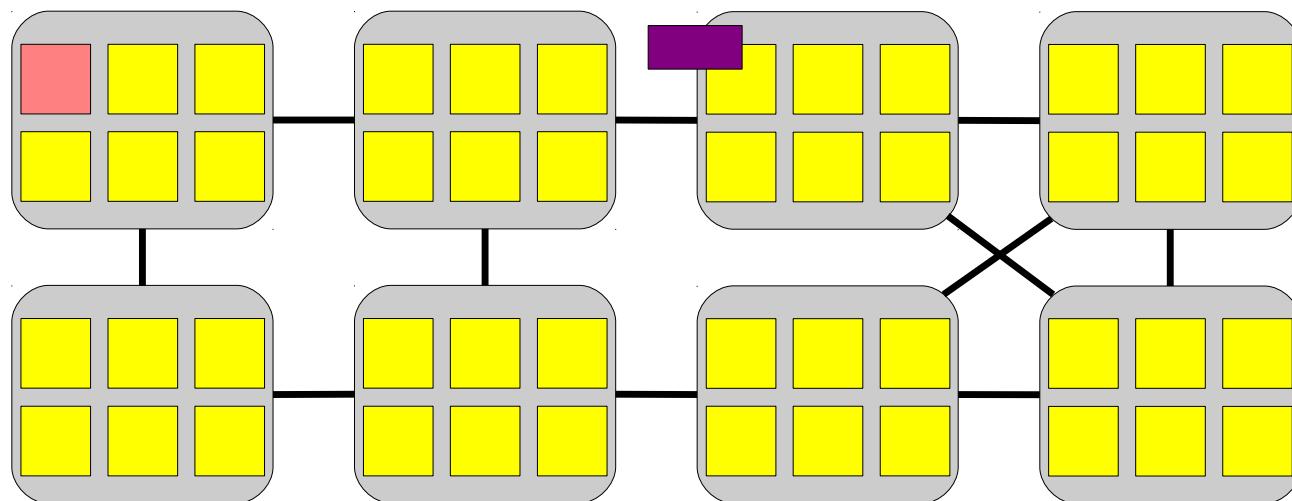
More readers means the next lock holder need to wait more time to get the lock

Allocate a ticket
read current ticket and spin

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```



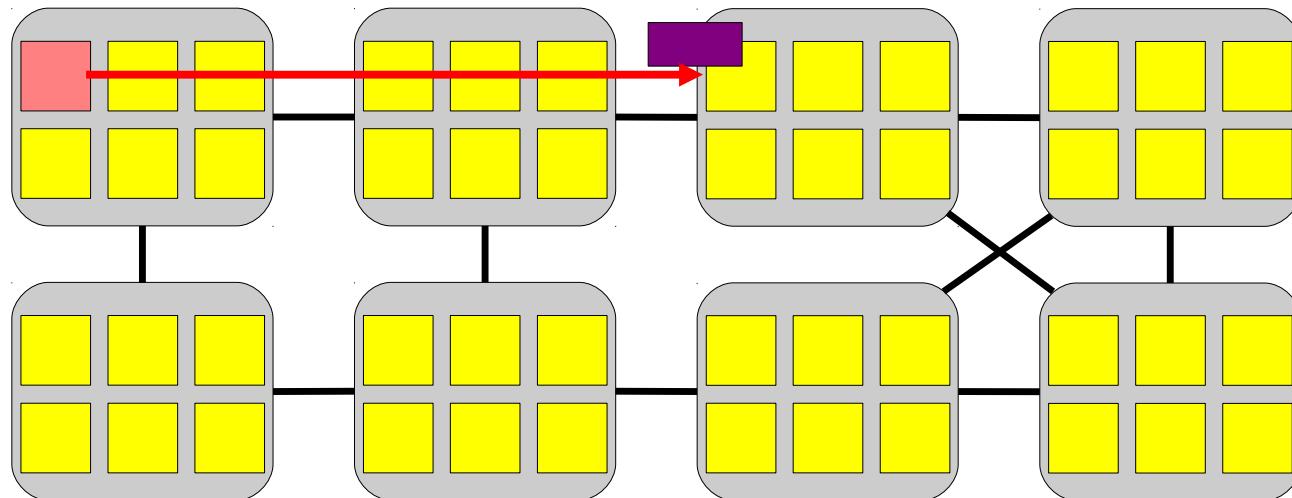
Allocate a ticket
read current ticket and spin

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

cache coherence message



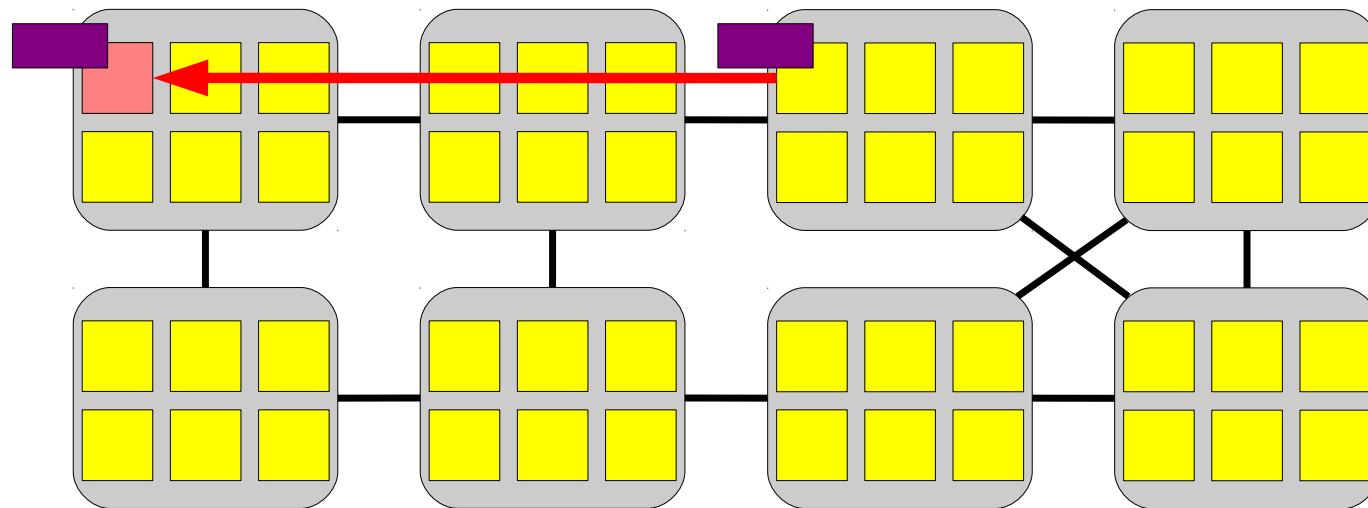
Allocate a ticket
read current ticket and spin

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

cache coherence message



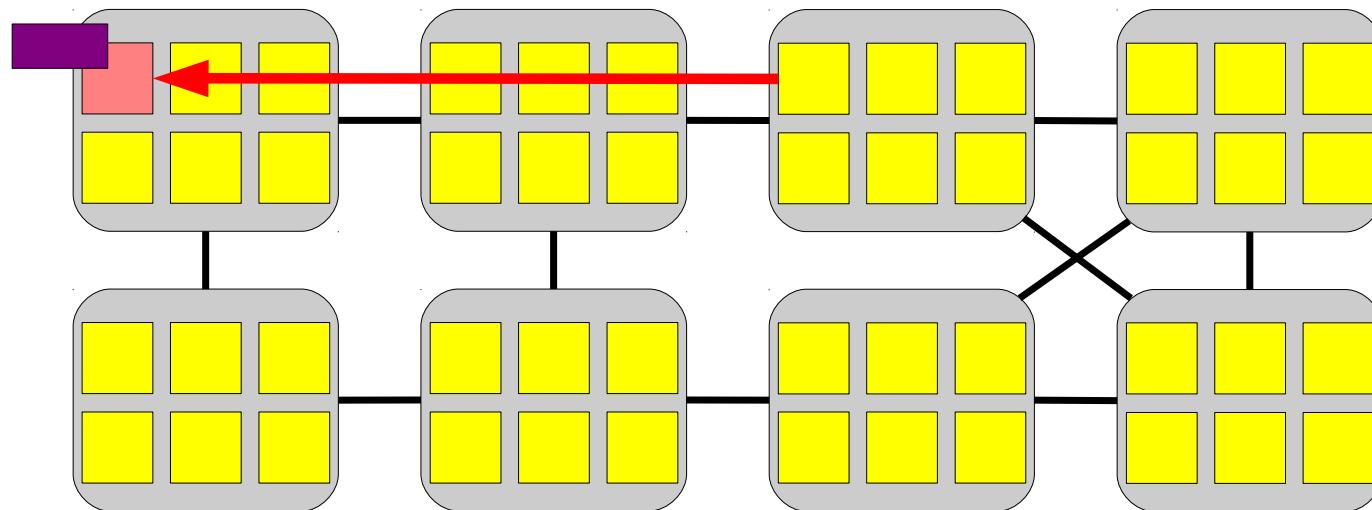
Allocate a ticket
read current ticket and spin

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

120 ~ 420 cycles

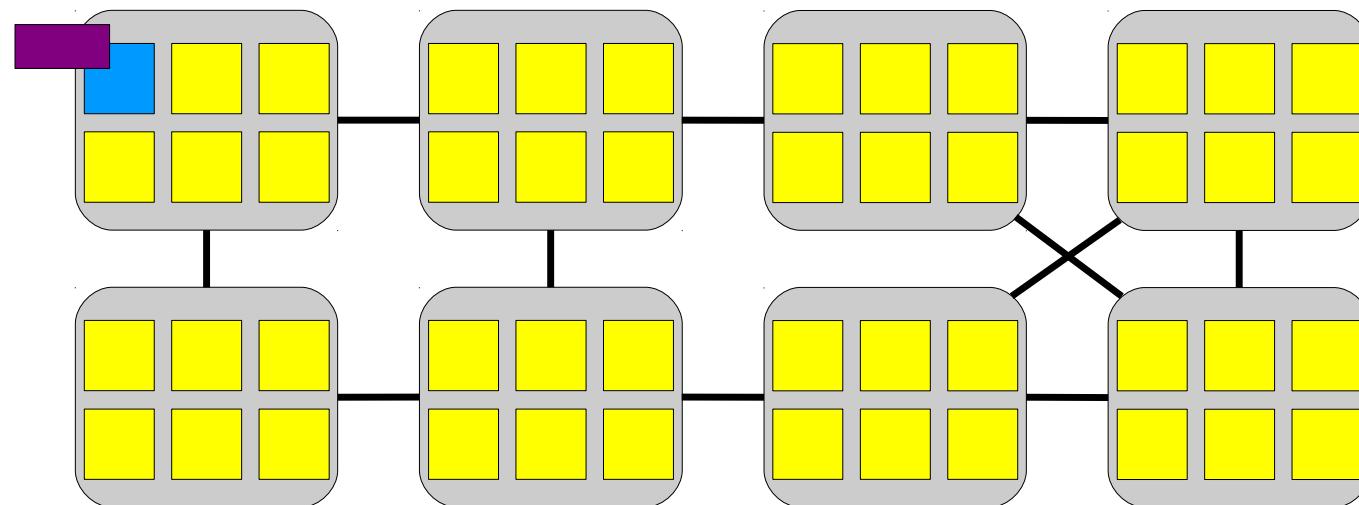


update the ticket value

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

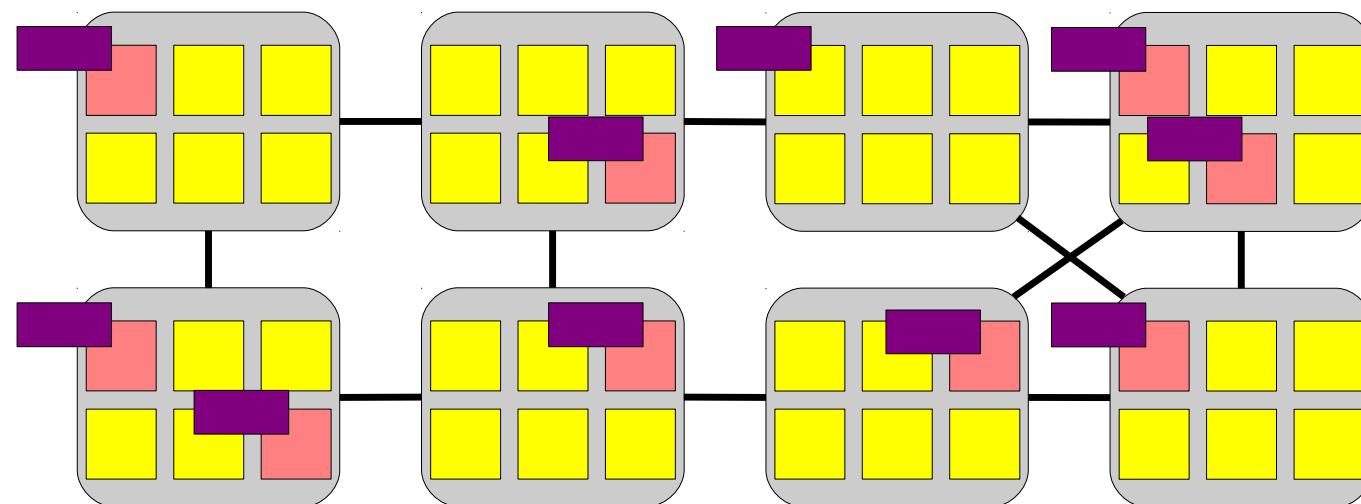
```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```



```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

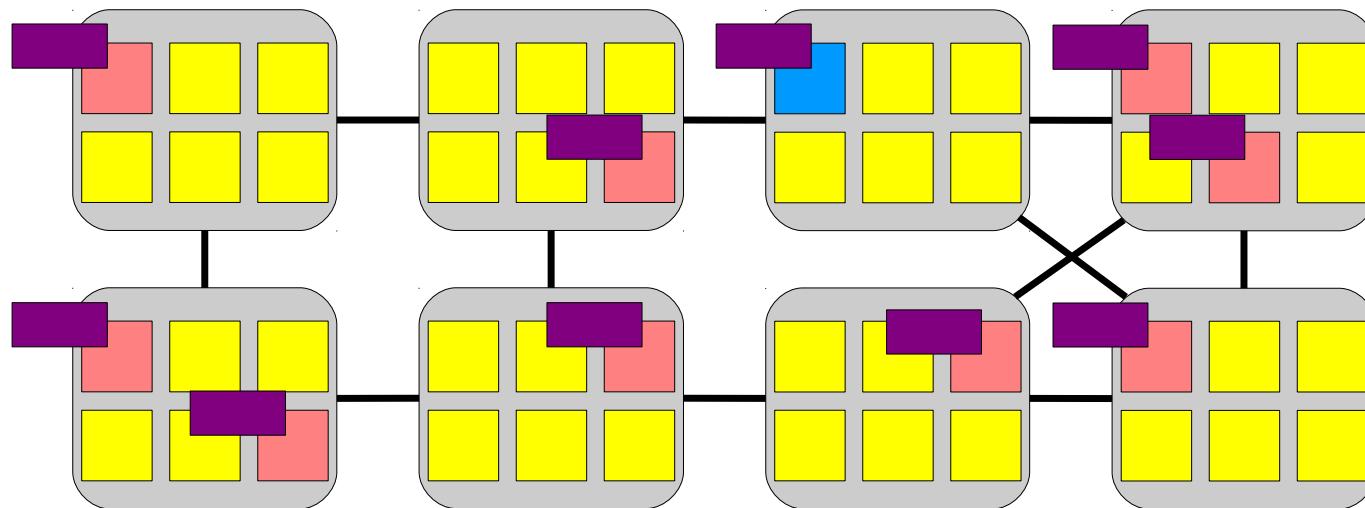


Lock shared by all core

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```



Lock holder update ticket

```

void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}

```

```

void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

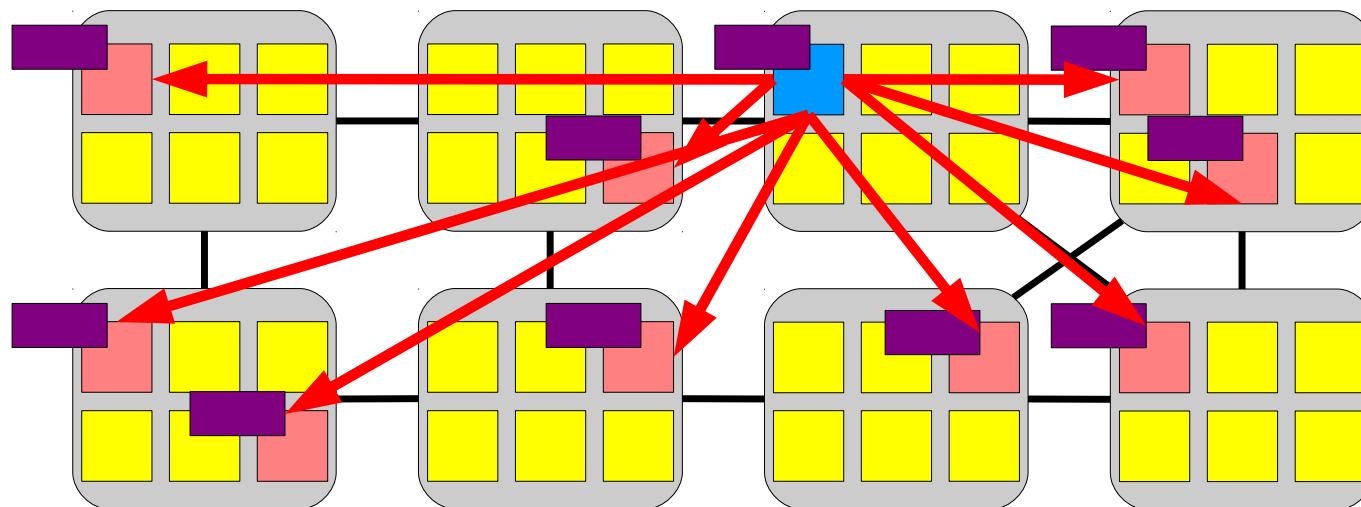

```

```

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}

```

Invalidate message



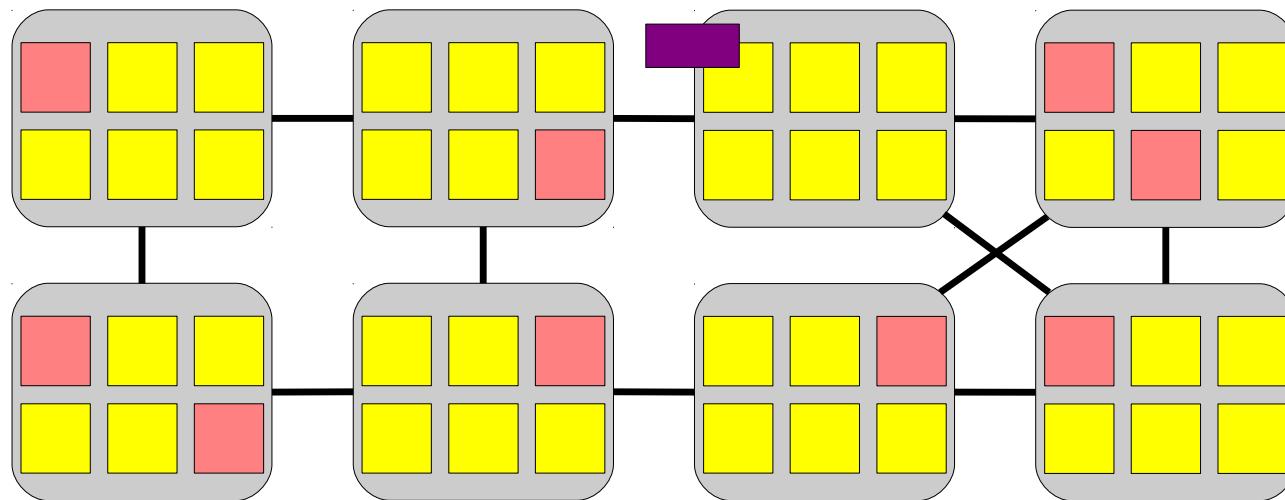
Lock holder update ticket

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

Only lock holder has lock in cache

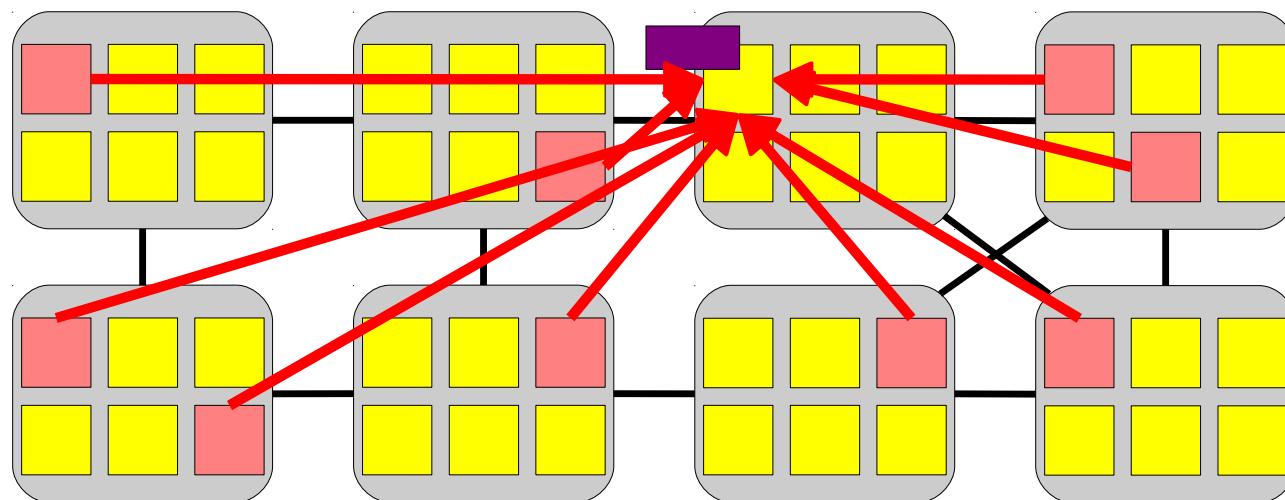


Lock holder update ticket

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```



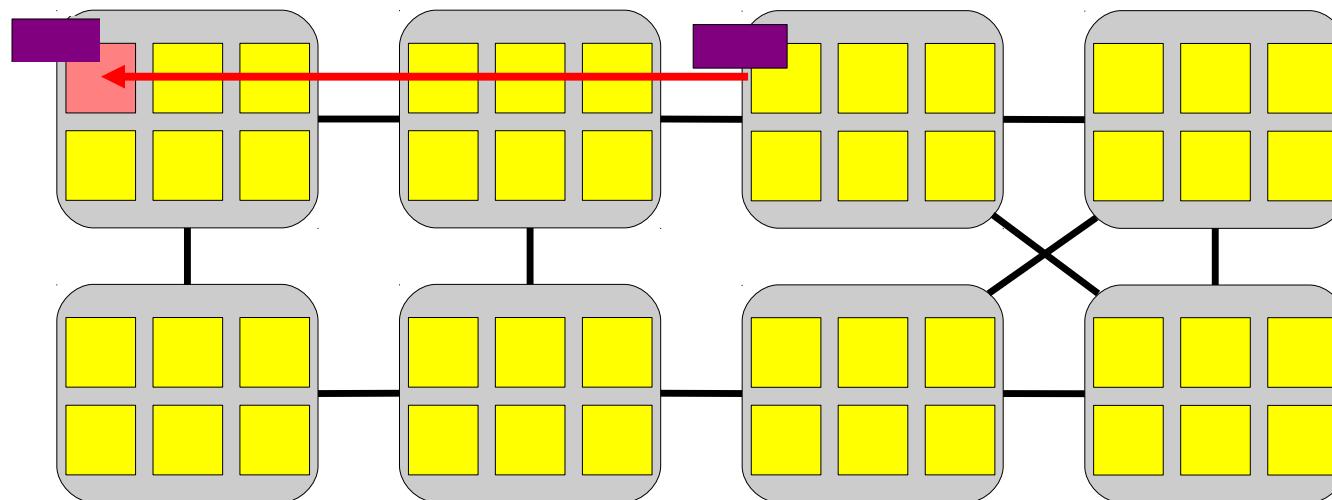
All waiters read the lock

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

500 ~ 4000 cycles!



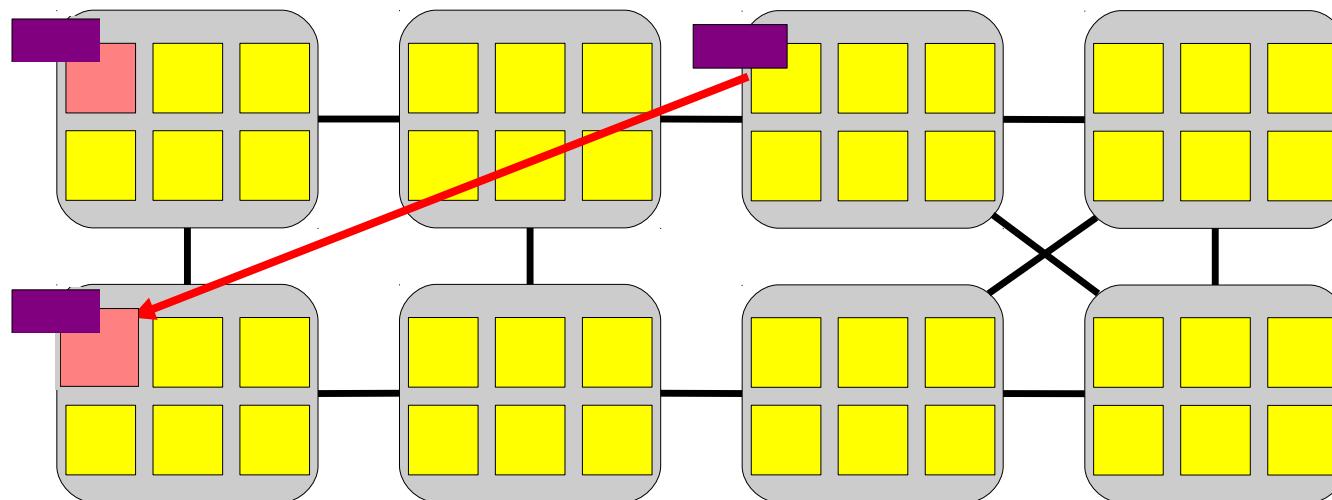
All waiters read the lock

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

Reply read request one by one



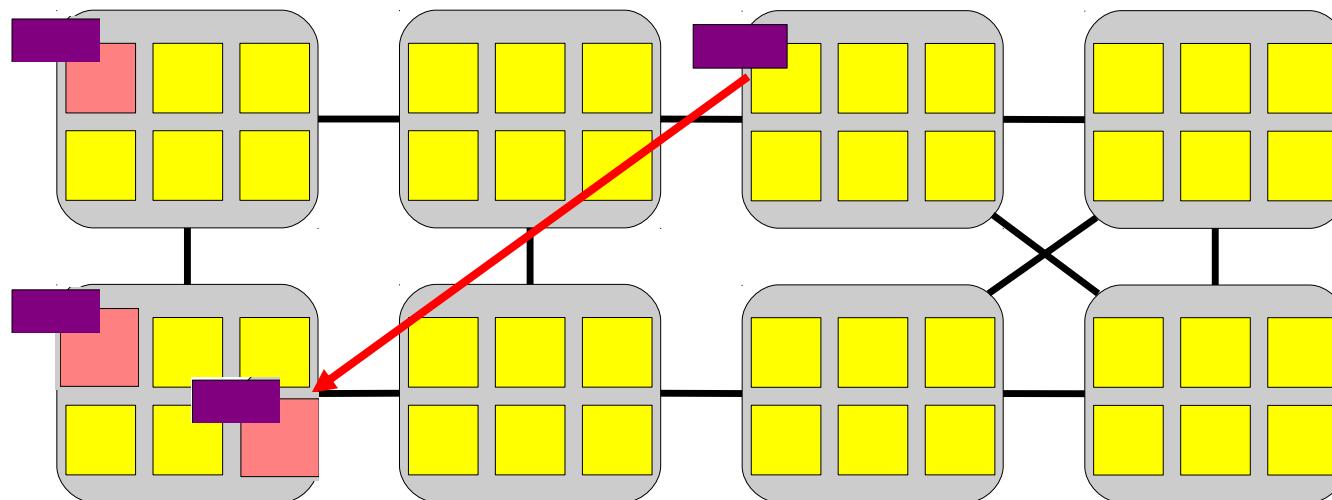
All waiters read the lock

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

Reply read request one by one



All waiters read the lock

```

void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}

```

Previous lock holder notifies next
lock holder after sending out
 $N/2$ replies

```

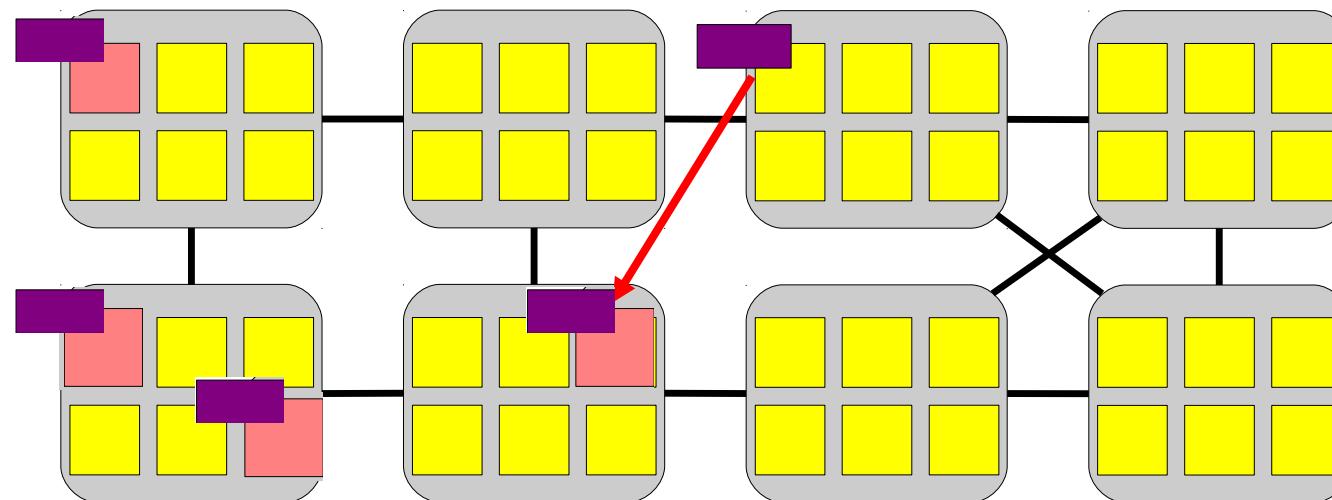
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

```

```

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}

```



All waiters read the lock

```
// Acquire the lock.  
// Loops (spins) until the lock is acquired.  
// Holding a lock for a long time may cause  
// other CPUs to waste time spinning to acquire it.  
void  
acquire(struct spinlock *lk)  
{  
    pushcli(); // disable interrupts to avoid deadlock.  
    if(holding(lk))  
        panic("acquire");  
  
    // The xchg is atomic.  
    // It also serializes, so that reads after acquire are not  
    // reordered before it.  
    while(xchg(&lk->locked, 1) != 0)  
    ;  
  
    // Record info about lock acquisition for debugging.  
    lk->cpu = cpu;  
    getcallerpcs(&lk, lk->pcs); // Release the lock.  
}  
  
void  
release(struct spinlock *lk)  
{  
    if(!holding(lk))  
        panic("release");  
  
    lk->pcs[0] = 0;  
    lk->cpu = 0;  
  
    // The xchg serializes, so that reads before release are  
    // not reordered after it. The 1996 PentiumPro manual (Volume 3,  
    // 7.2) says reads can be carried out speculatively and in  
    // any order, which implies we need to serialize here.  
    // But the 2007 Intel 64 Architecture Memory Ordering White  
    // Paper says that Intel 64 and IA-32 will not move a load  
    // after a store. So lock->locked = 0 would work here.  
    // The xchg being asm volatile ensures gcc emits it after  
    // the above assignments (and after the critical section).  
    xchg(&lk->locked, 0);  
  
    popcli();  
}
```

Is xv6 locking scalable?

SCALABLE LOCKING

Making ticket spinlock scalable

Common way: use proportional back-off

```
void spin_lock(spinlock_t *l) {
    int t = atomic_xadd(&l->next_ticket);
    while (t != lock->current_ticket)
        // wait more time with each failure
        ;
}
```

Why this would work?

Example Fix

```
arch/x86/kernel/smp.c | 23 ++++++-----+
 1 file changed, 20 insertions(+), 3 deletions(-)
diff --git a/arch/x86/kernel/smp.c b/arch/x86/kernel/smp.c
index 20da354..aa743e9 100644
--- a/arch/x86/kernel/smp.c
+++ b/arch/x86/kernel/smp.c
@@ -117,11 +117,28 @@ static bool smp_no_nmi_ipi = false;
 */
void ticket_spin_lock_wait(arch_spinlock_t *lock, struct __raw_tickets inc)
{
+    __ticket_t head = inc.head, ticket = inc.tail;
+    __ticket_t waiters_ahead;
+    unsigned loops;
+
+    for (;;) {
-        cpu_relax();
-        inc.head = ACCESS_ONCE(lock->tickets.head);
+        waiters_ahead = ticket - head - 1;
+        /*
+         * We are next after the current lock holder. Check often
+         * to avoid wasting time when the lock is released.
+         */
+        if (!waiters_ahead) {
+            do {
+                cpu_relax();
+            } while (ACCESS_ONCE(lock->tickets.head) != ticket);
+            break;
+        }
+        loops = 50 * waiters_ahead;
+        while (loops--)
+            cpu_relax();

-        if (inc.head == inc.tail)
+        head = ACCESS_ONCE(lock->tickets.head);
+        if (head == ticket)
+            break;
    }
}
```

Linus' Response

<http://linux-kernel.2935.n7.nabble.com/PATCH-v5-0-5-x86-smp-make-ticket-spinlock-proportional-backoff-w-auto-tuning-td596698i20.html>

So I claim:

- it's *really* hard to trigger in real loads on common hardware.
- if it does trigger in any half-way reasonably common setup (hardware/software), we most likely should work really hard at fixing the underlying problem, not the symptoms.
- we absolutely should *not* pessimize the common case for this

Thanks

- Next class
 - MCS Lock
 - Read-writer lock, Lock-free synchronization

Synchronization Constructs

Rong Chen
Institute of Parallel and Distributed Systems
Shanghai Jiao Tong University
http://ipads.se.sjtu.edu.cn/haibo_chen

Some slides adjusted from. Elsa L Gunter (UIUC), Jonathan Walpole (PSU)
Paul McKenney (IBM) Tom Hart (University of Toronto)

Outline

Recap

MCS Lock

Read-write Lock

Lock-free coordination

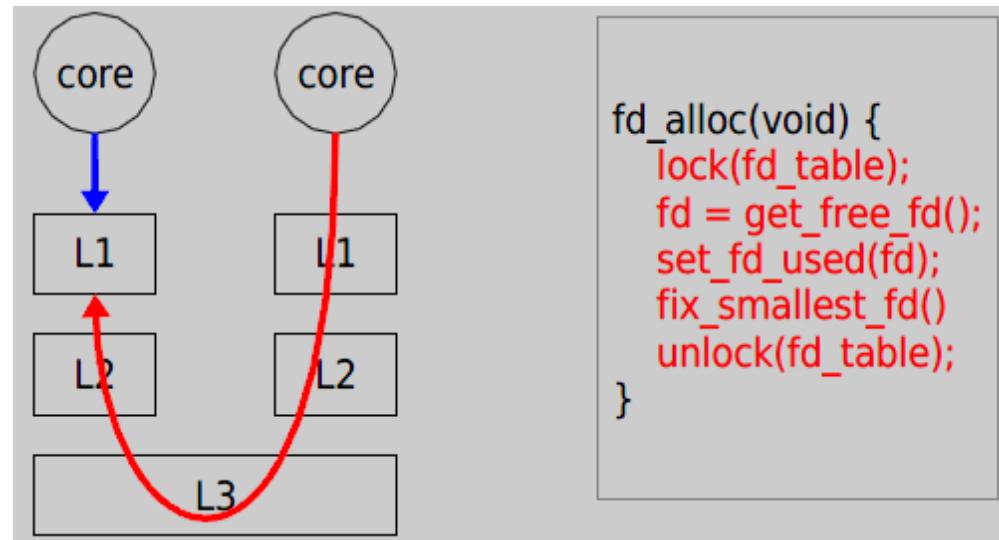
Recap: What is scalability?

Application does N times as much work on N cores
as it could on 1 core

Scalability may be limited by Amdahl's Law:

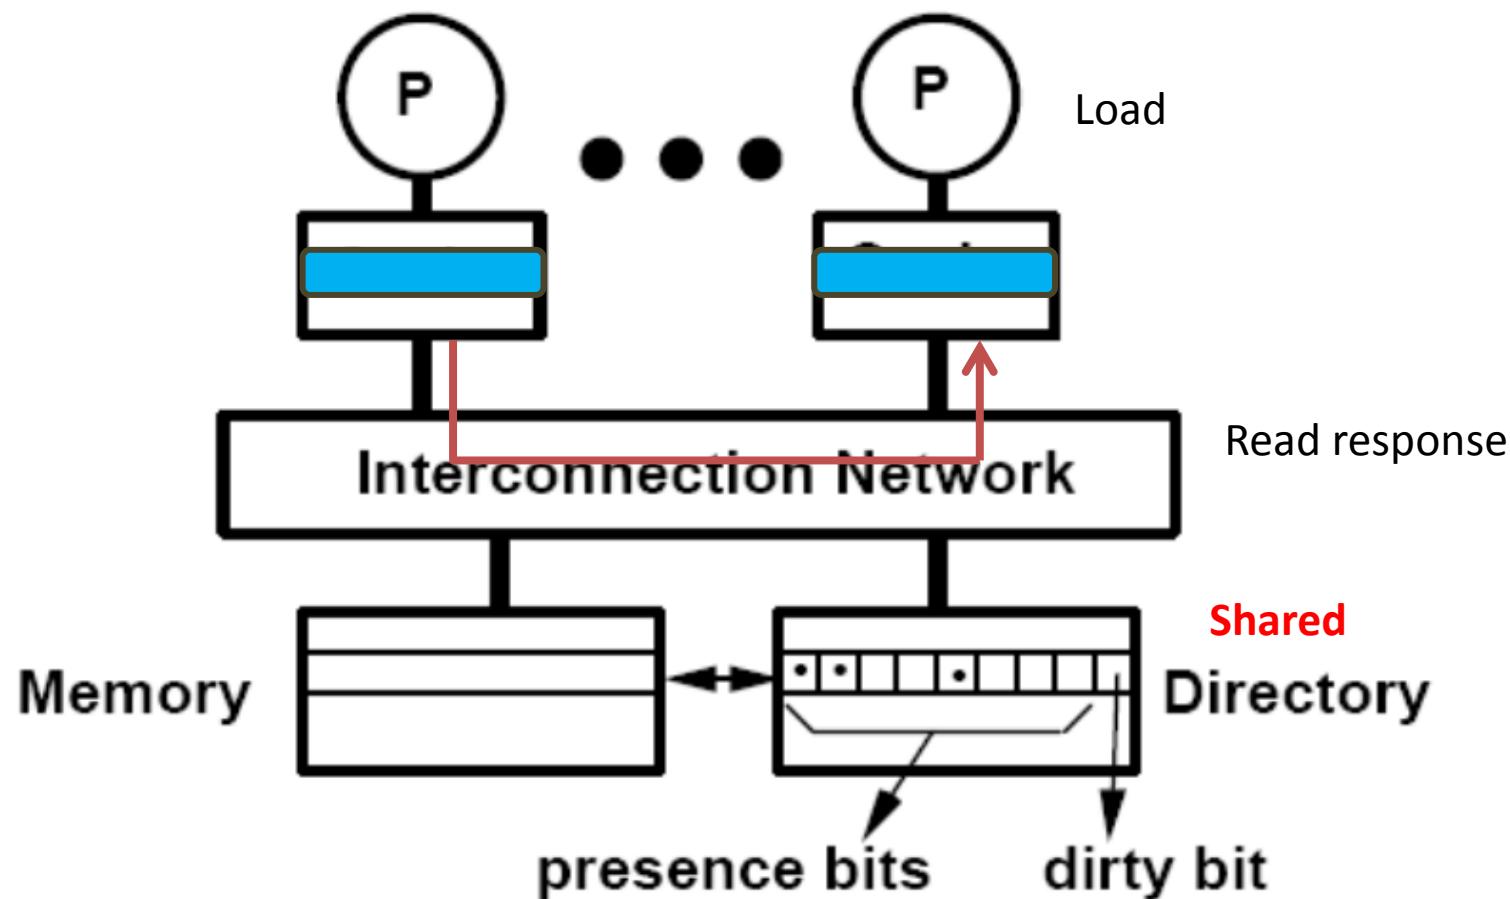
Locks, shared data structures, ... Shared hardware
(DRAM, NIC, ...)

Recap: Why throughput drops?



Now it takes 121 cycles!

Recap: Directory-based cache coherence

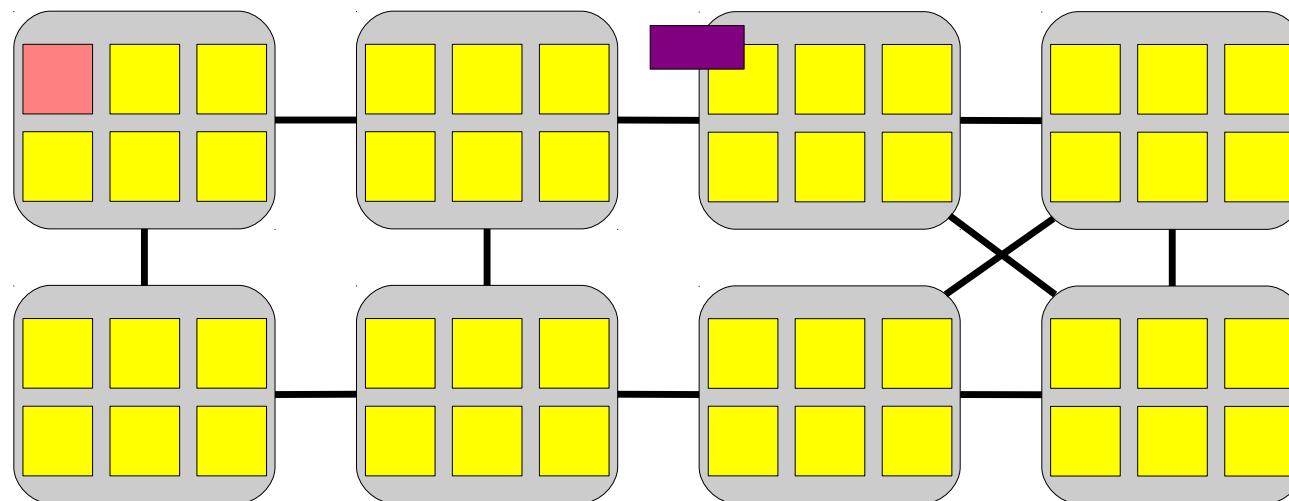


Recap: Allocate a ticket read current ticket and spin

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}
```

```
struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```



Synchronization Constructs

MCS lock

Read-write Lock

Lock-free

Using scalable locks

Many existing scalable locks

Main idea is to avoid contending on a single cache line

Example

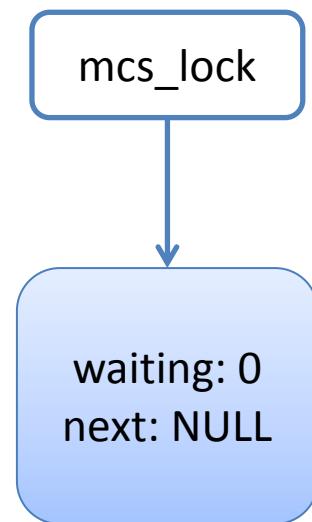
MCS (John M. Mellor-Crummey and Michael L. Scott)
K42

General idea of MCS lock

mcs_lock

NULL

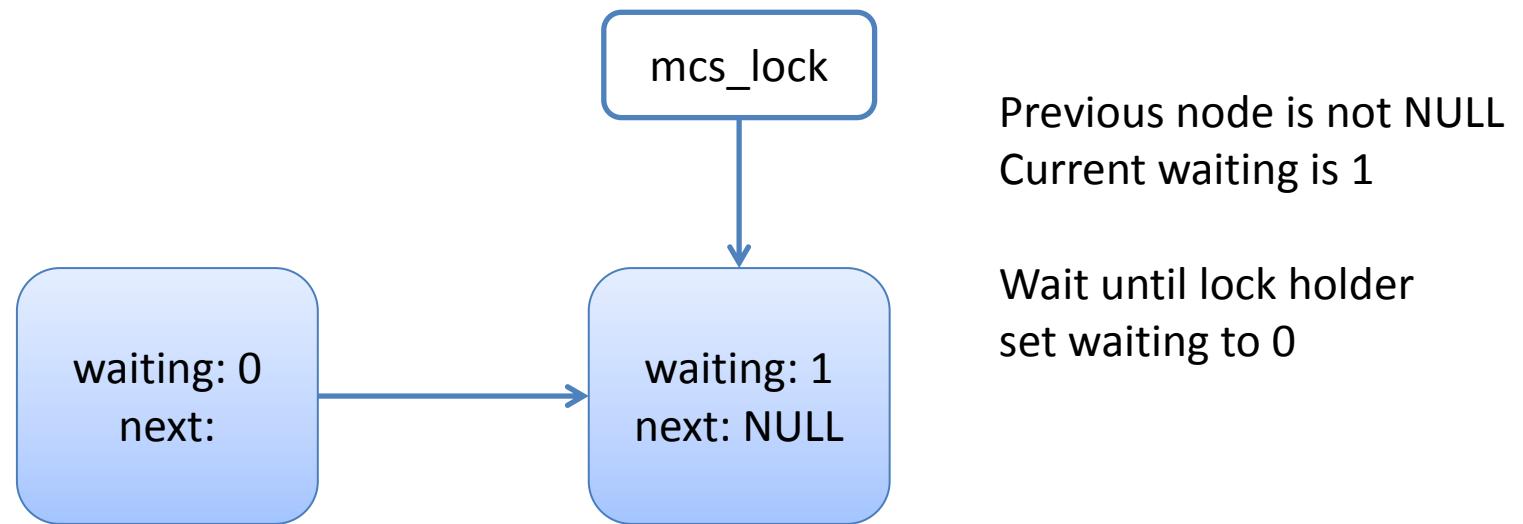
General idea of MCS lock



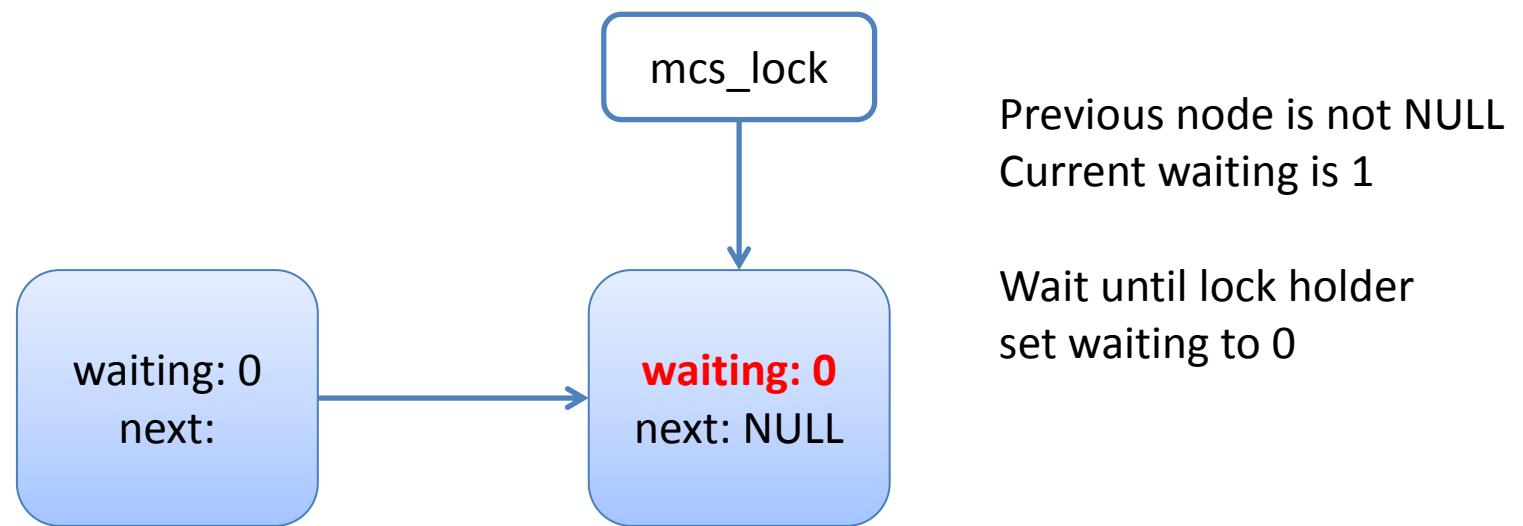
Use compare and swap to change
mcs_lock point to self node

Check previous node
NULL in this case, no need to wait

General idea of MCS lock



General idea of MCS lock



General idea of MCS lock



Previous node is not NULL
Current waiting is 1

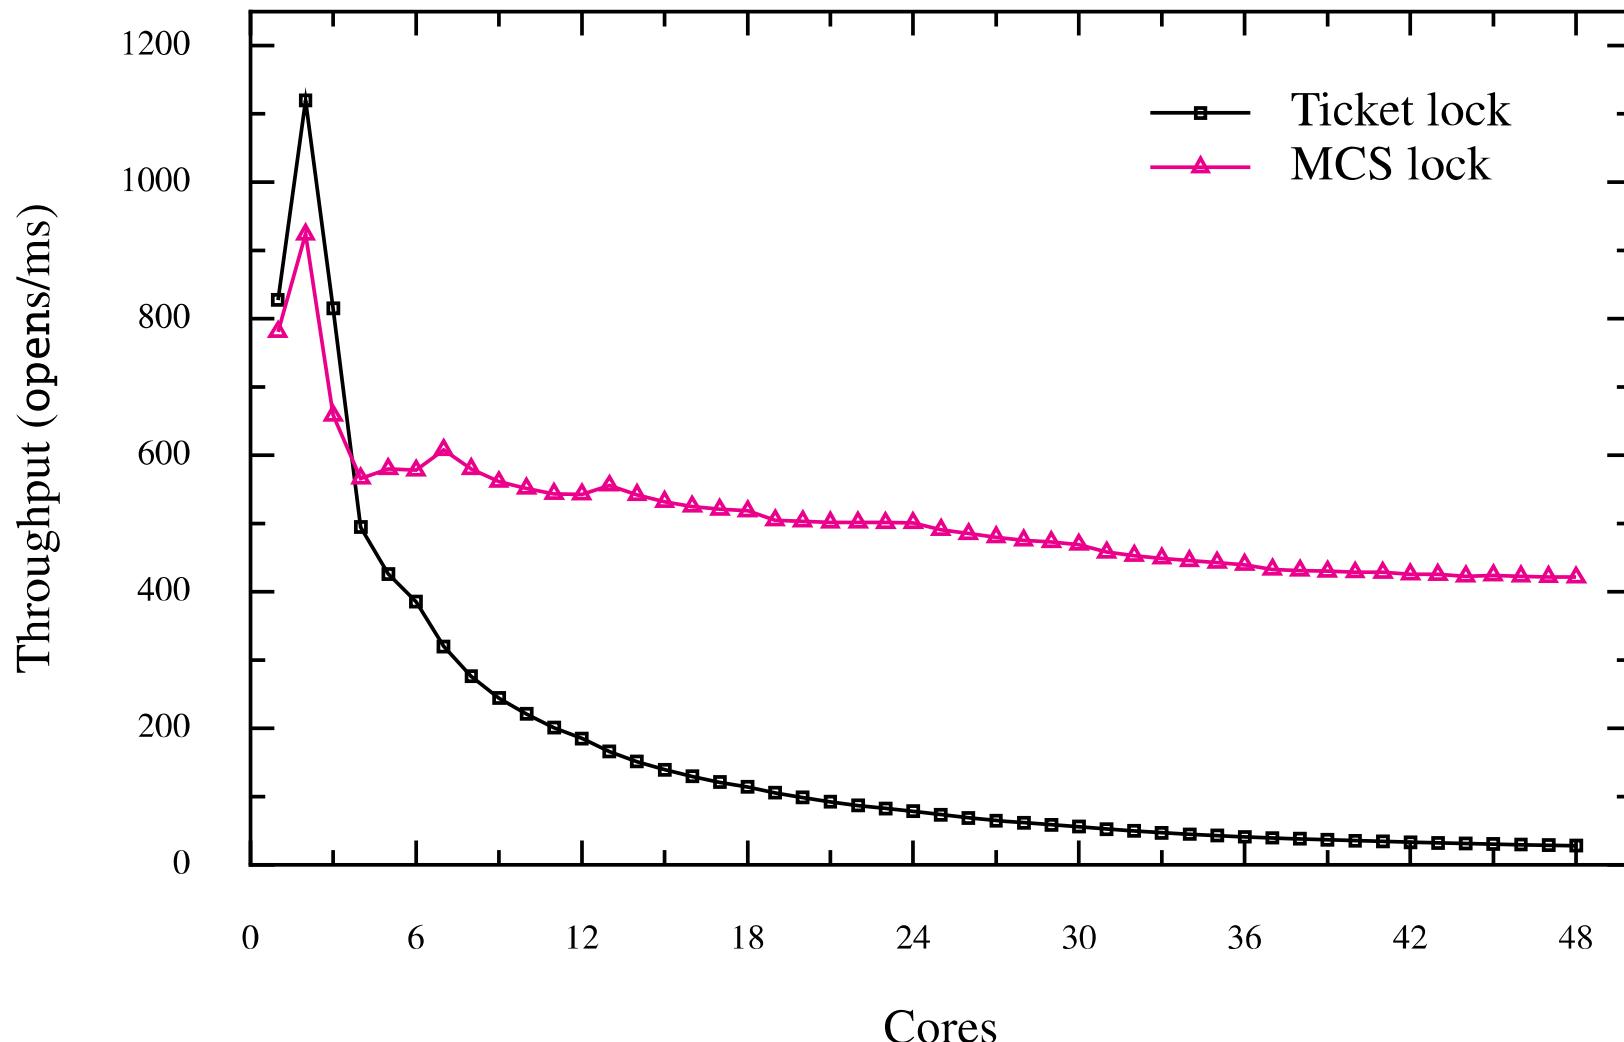
Wait until lock holder
set waiting to 0

Code

```
mcs_node{  
    mcs_node next;  
    int is_locked;  
}  
  
mcs_lock{  
    mcs_node queue;  
}  
  
function lock(mcs_lock lock, mcs_node my_node) {  
    my_node.next = NULL;  
    mcs_node predecessor =  
        fetch_and_store(lock.queue, my_node);  
    if (predecessor != NULL) {  
        my_node.is_locked = true;  
        predecessor.next = my_node;  
        while (my_node.is_locked) {};  
    }  
}
```

```
function unlock(mcs_lock lock, mcs_node my_node) {
    if(my_node.next == NULL) {
        if(compare_and_swap(lock.queue, my_node, NULL) {
            return;
        } else{
            while(my_node.next == NULL) { } ;
        }
    }
    my_node.next.is_locked = false;
}
```

Using scalable locks: FOPS



READER-WRITER LOCK

Motivating Example: Calendar

Goal: online calendar for a class

Lots of people may read it at the same time

Only one person updates it (Prof, TAs)

Shared data

`map<date, listOfEvents> EventMap`

`listOfEvents GetEvents(date)`

`AddEvent(data, newEvent)`

Basic Code – Single Threaded

```
GetEvents(date) {  
    List events = EventMap.find(date).copy();  
    return events;  
}
```

```
AddEvent(data, newEvent) {  
    EventMap.find(date) += newEvent;  
}
```

Inefficient Multi-threaded code

```
GetEvents(date) {  
    lock(mapLock);  
    List events = EventMap.find(date).copy();  
    unlock(mapLock);  
    return events;  
}
```

```
AddEvent(data, newEvent) {  
    lock(mapLock);  
    EventMap.find(date) += newEvent;  
    unlock(mapLock);  
}
```

How to do with reader – writer locks?

```
GetEvents(date) {  
    List events = EventMap.find(date).copy();  
    return events;  
}
```

```
AddEvent(data, newEvent) {  
    EventMap.find(date) += newEvent;  
}
```

Reader – Writer Locks

Problem definition:

Shared data that will be read and written by multiple threads

Allow multiple readers to access shared data when no threads are writing data

A thread can write shared data only when no other thread is reading or writing the shared data

Interface

readerStart()

readerFinish()

writerStart()

writerFinish()

Many threads can be in between a `readStart` and `readerFinish`

Only one thread can be between `writerStart` and `writerFinish`

How to do with reader – writer locks?

```
GetEvents(date) {  
    readerStart(mapRWLock);  
    List events = EventMap.find(date).copy();  
    readerFinish(mapRWLock);  
    return events;  
}  
  
AddEvent(data, newEvent) {  
    writerStart(mapRWLock);  
    EventMap.find(date) += newEvent;  
    writerFinish(mapRWLock);  
}
```

Additional Layer of Synchronization

Concurrent programs

Even higher-level synchronization

High-level synchronization provided by software

Low-level atomic operations provided by hardware

Reader – Writer Locks using Monitors

Note: Implement Reader/Writer Locks as an abstractions, *not* as an integrated part of code

Central Questions:

Shared Data?

Ordering Constraints?

How many Condition Variables?

Reader – Writer Locks using Monitors

Note: Implement Reader/Writer Locks as an abstractions, *not* as an integrated part of code

Central Questions:

Shared Data? NumReaders, NumWriters

Ordering Constraints?

How many Condition Variables?

Reader – Writer Locks using Monitors

Note: Implement Reader/Writer Locks as an abstractions, *not* as an integrated part of code

Central Questions:

Shared Data? NumReaders, NumWriters

Ordering Constraints?

readerStart must wait if there are writers

writerStart must wait if there are readers or writes

How many Condition Variables?

Reader – Writer Locks using Monitors

Note: Implement Reader/Writer Locks as an abstractions, *not* as an integrated part of code

Central Questions:

Shared Data? NumReaders, NumWriters

Ordering Constraints?

readerStart must wait if there are writers

writerStart must wait if there are readers or writes

How many Condition Variables?

One: condRW (no readers or writers)

Basic Implementation

```
readerStart() {
```

```
readerFinish() {
```

```
}
```

```
}
```

Basic Implementation

```
readerStart() {  
    lock(lockRW);  
  
    while(numWriters > 0){  
        wait(lockRW, condRW);  
    };  
  
    numReaders++;  
  
    unlock(lockRW);  
}  
  
readerFinish() {  
    lock(lockRW);  
  
    numReaders--;  
  
    broadcast(lockRW, condWR);  
  
    unlock(lockRW);  
}
```

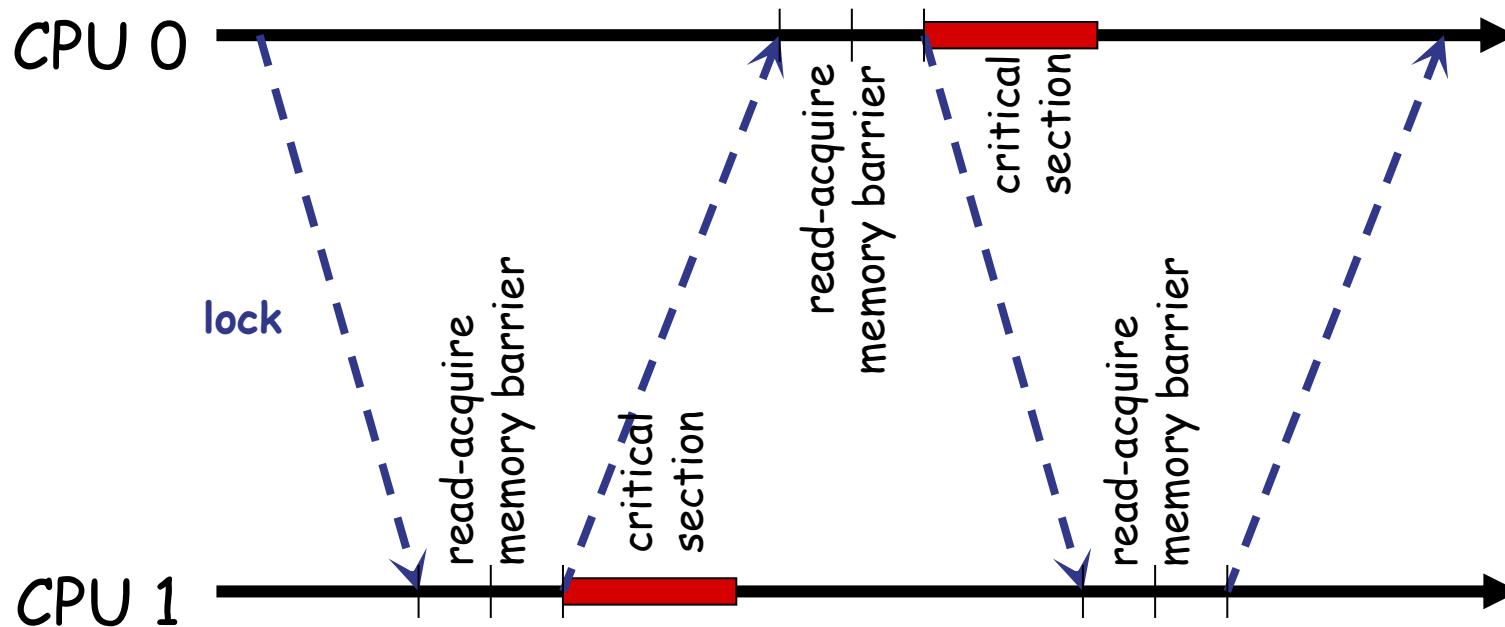
Basic Implementation

```
writerStart() {  
    lock(lockRW);  
  
    while(numReaders > 0||  
          numWriters >0){  
        wait(lockRW, condRW);  
    };  
  
    numWriters++;  
  
    unlock(lockRW);  
}  
  
writerFinish() {  
    lock(lockRW);  
  
    numWriters--;  
  
    broadcast(lockRW, condWR);  
  
    unlock(lockRW);  
}
```

Better Implementation

```
readerStart() {  
    lock(lockRW);  
  
    while(numWriters > 0){  
        wait(lockRW, condRW);  
    };  
  
    numReaders++;  
  
    unlock(lockRW);  
}  
  
readerFinish() {  
    lock(lockRW);  
  
    numReaders--.  
  
    if(numReaders == 0){  
        signal(lockRW, condWR);  
    };  
    unlock(lockRW);  
}
```

Scalability of reader/writer locking



Reader/writer locking does not scale due to critical section efficiency!

Reader/writer spin locks (rwlock) in Linux

```
rwlock_t mr_rwlock = RW_LOCK_UNLOCKED;  
  
read_lock(&mr_rwlock);  
/* critical section (read only) ... */  
read_unlock(&mr_rwlock);  
  
write_lock(&mr_rwlock);  
/* critical section (read and write) ... */  
write_unlock(&mr_rwlock);
```

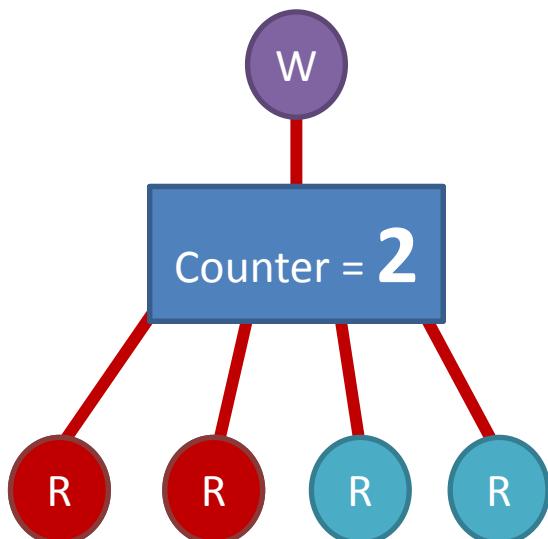
Scalability of rwlock?

Reading

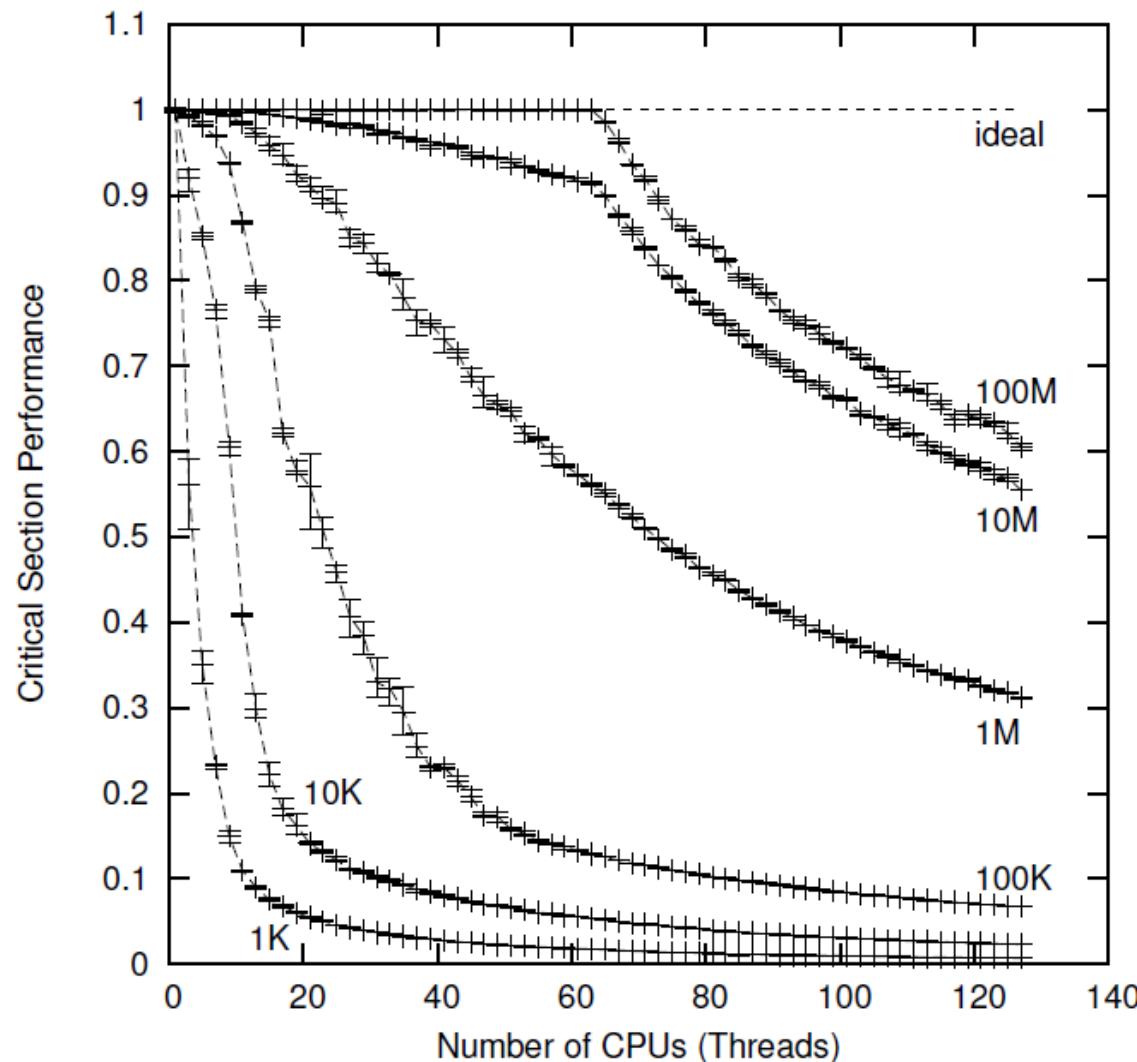
Counter -> #readers

Writing

Counter = -1



Traditional RW lock performance



Problem

Reader Lock Acquisition

add_and_featch(&counter)

Need to wait for message (current #reader)

Need to send message (next reader)

Acquisition serialized by messages!!

Occupies a great deal of message bandwidth!!

Solutions

Idea 1

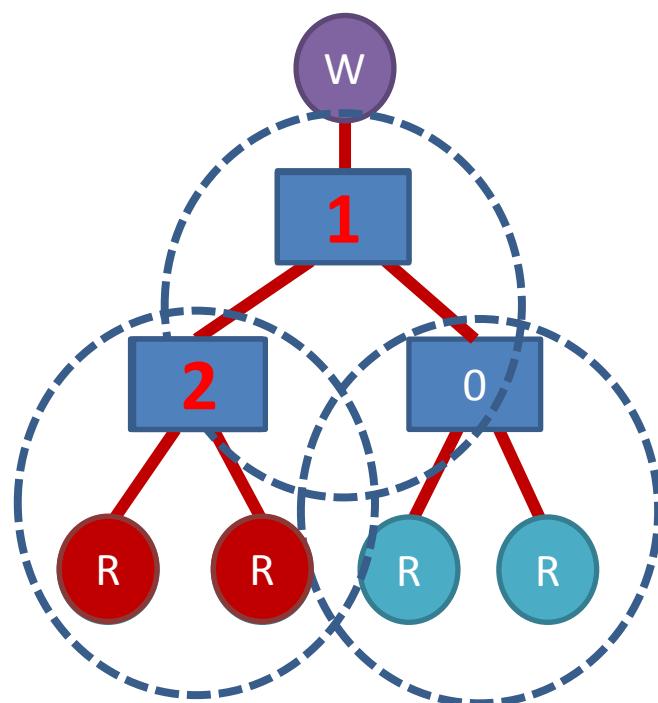
Reduce time to wait for message (GOLL)

Idea 2

Reduce #message (BR lock)

GOLL

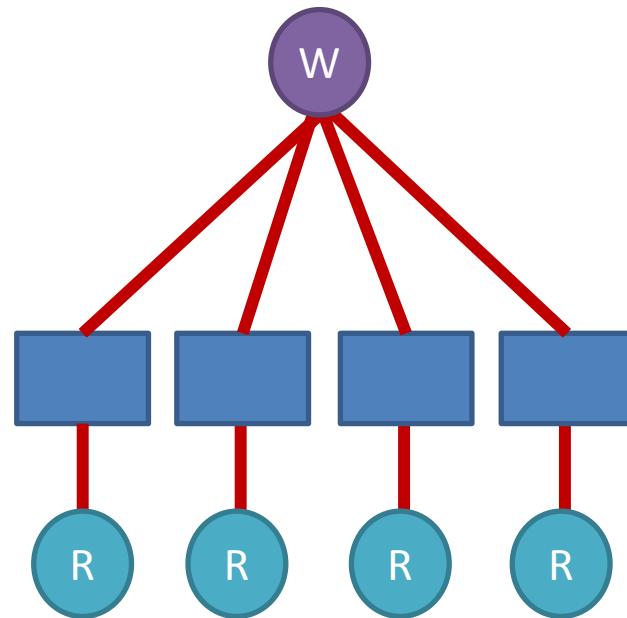
Reduce time to wait for message



Waiting domain is split into pieces
Reduce #remote messages

BR lock

Reduce #message



No reader messages if there's no writer

Thanks

Next class

Concurrency bugs, race detection

Bug Survey

Rong Chen
Institute of Parallel and Distributed Systems
Shanghai Jiao Tong University
http://ipads.se.sjtu.edu.cn/rong_chen

Some slides adjusted from Frans Kaashoek (MIT), Shan Lu (U. Wisc) and Ding Yuan (U. Toronto)

Outline

Recap

Understanding Concurrency bugs

OS Bugs

Data race detection

Recap: How to do with reader – write locks?

```
GetEvents(date) {  
    readerStart(mapRWLock);  
    List events = EventMap.find(date).copy();  
    readerFinish(mapRWLock);  
    return events;  
}  
  
AddEvent(data, newEvent) {  
    writerStart(mapRWLock);  
    EventMap.find(date) += newEvent;  
    writerFinish(mapRWLock);  
}
```

Recap: Basic Implementation

```
readerStart() {  
    lock(lockRW);  
  
    while(numWriters > 0){  
        wait(lockRW,condRW);  
    };  
  
    numReaders++;  
  
    unlock(lockRW);  
}
```

```
readerFinish() {  
    lock(lockRW);  
  
    numReaders--;  
  
    broadcast(lockRW,condRW);  
  
    unlock(lockRW);  
}
```

Recap: Basic Implementation

```
writerStart() {  
    lock(lockRW);
```

```
    while(numReaders > 0 ||  
          numWriters > 0){  
        wait(lockRW,condRW);  
    };
```

```
    numWriters++;  
  
    unlock(lockRW);  
}
```

```
writerFinish() {  
    lock(lockRW);
```

```
    numWriters--;  
  
    broadcast(lockRW,condRW);
```

```
    unlock(lockRW);  
}
```

Recap: Better Implementation

```
readerStart() {  
    lock(lockRW);  
  
    while(numWriters > 0){  
        wait(lockRW,condRW);  
    }  
  
    numReaders++;  
  
    unlock(lockRW);  
}
```

```
readerFinish() {  
    lock(lockRW);  
  
    numReaders--.  
  
    if(numReaders == 0){  
        signal(lockRW,condWR);  
    };  
    unlock(lockRW);  
}
```

CONCURRENCY BUG CHARACTERISTICS

Summary

105 **real-world** concurrency bugs from 4 large open source programs

Study from 4 dimensions

- Bug patterns
- Manifestation condition
- Diagnosing strategy
- Fixing methods

Implications for:

- Bug detection
- Software testing
- PL design

Outline

Methodology

Findings and implications

- Bug pattern

- Bug manifestation

- Bug fixing

Conclusions

Application sources

	MySQL	Apache	Mozilla	OpenOffice
Software Type	Server	Server	Client	GUI
Language	C++/C	Mainly C	C++	C++
LOC (M line)	2	0.3	4	6
Bug DB history	6 years	7 years	10 years	8 years

**Different types of
real world applications**

Bug sources

	MySQL	Apache	Mozilla	OpenOffice	Total
Non-deadlock	14	13	41	6	74
Deadlock	9	4	16	2	31

Limitations

- No scientific computing applications
- No JAVA programs
- Never-enough bug samples

Non-Deadlock Bug Pattern

Classified based on root causes

Categories:

Atomicity violation

Desired atomicity of certain code region is violated

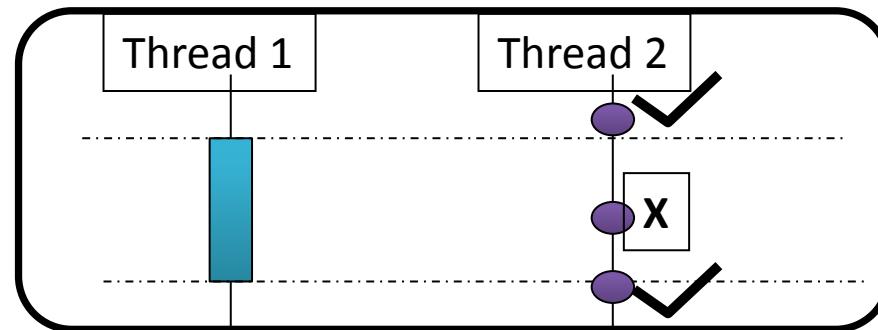
Order violation

The desired order between two (sets of) accesses is flipped

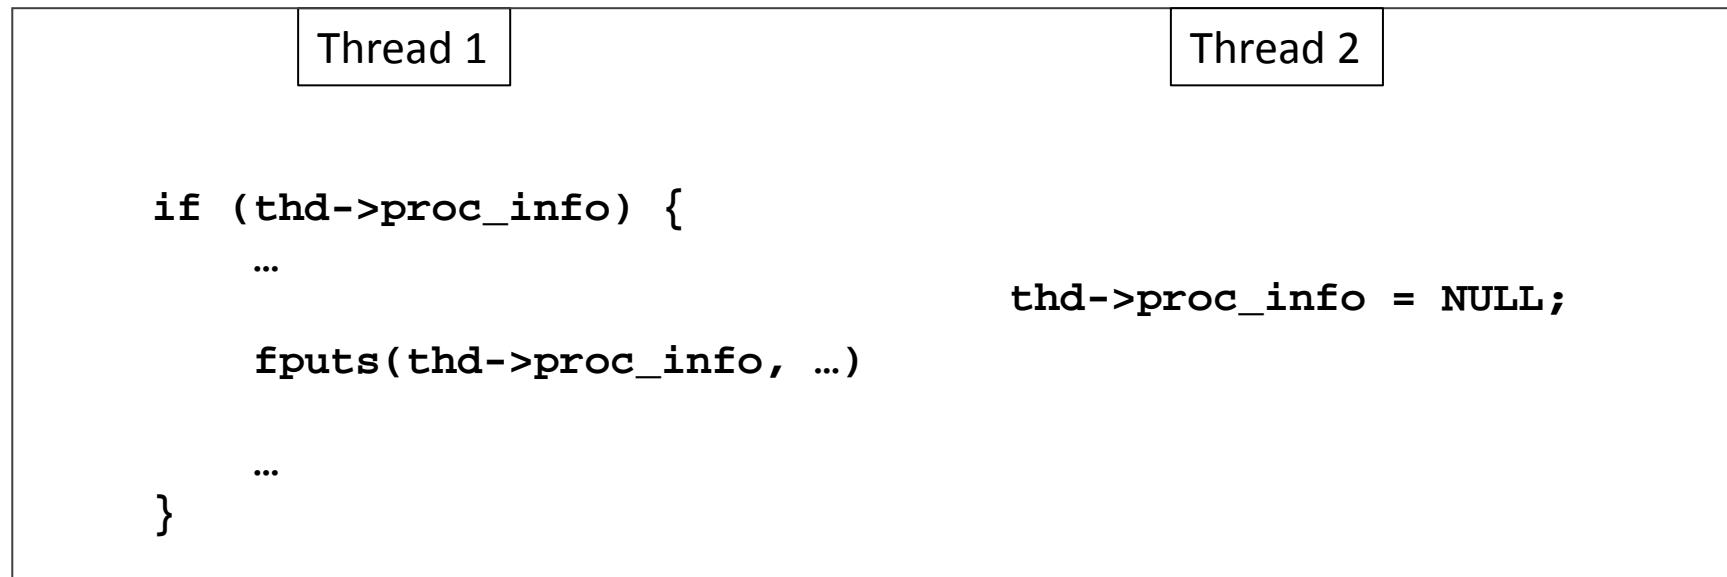
Others

Atomicity violation

The desired atomicity of certain code region is violated

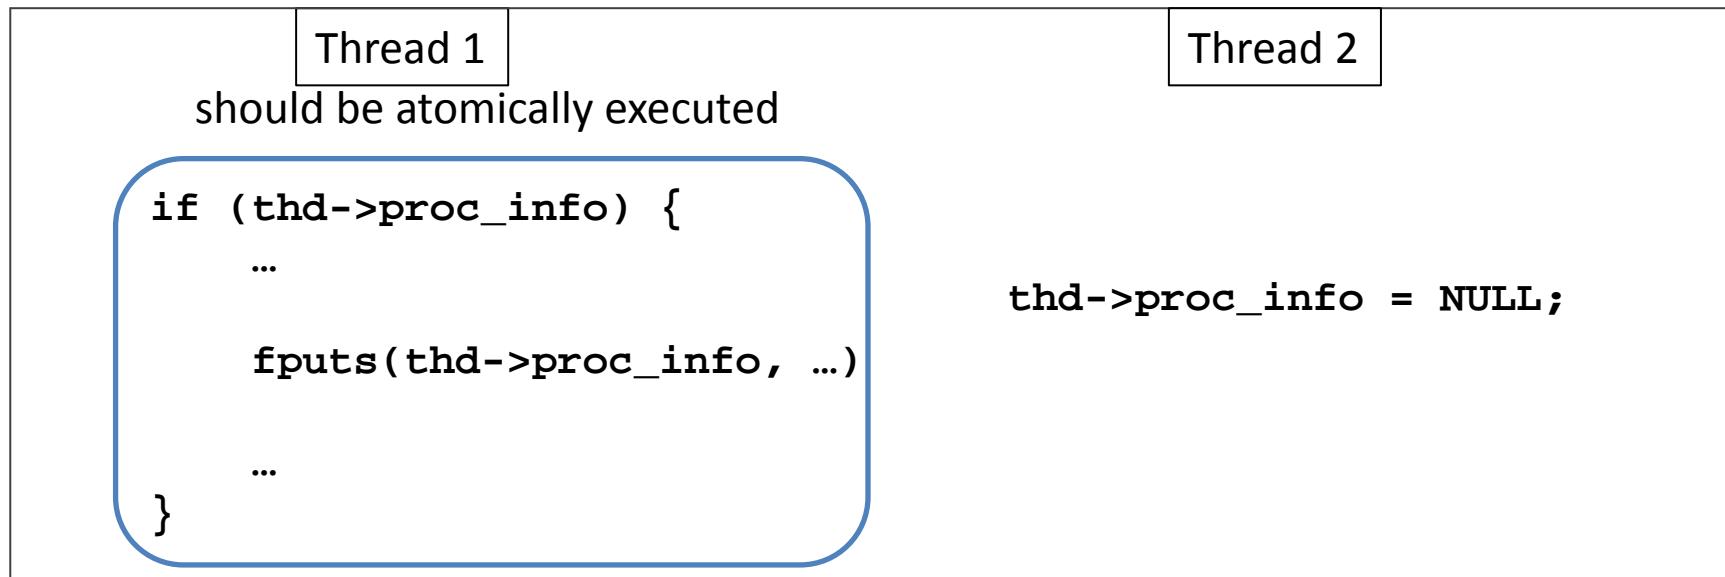


Example of atomicity violation



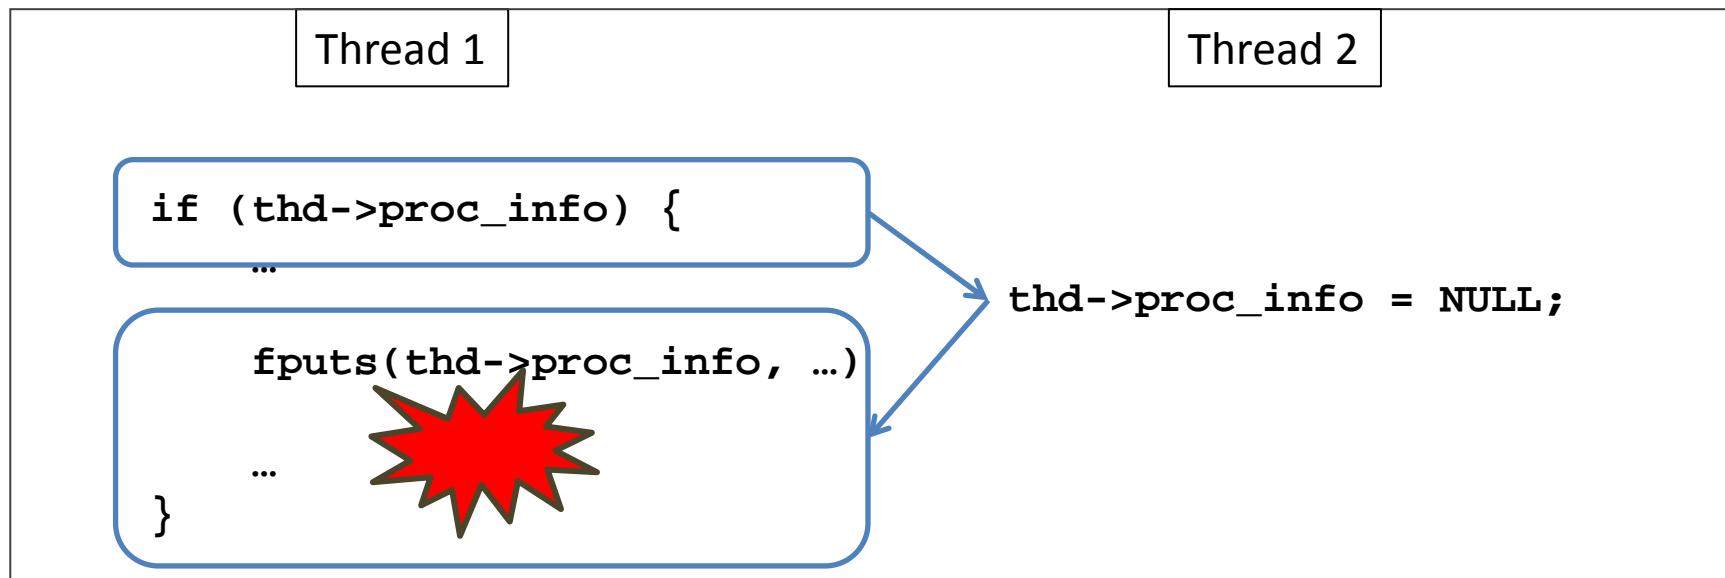
MySQL ha_innodb.hpp

Example of atomicity violation



MySQL ha_innodb.hpp

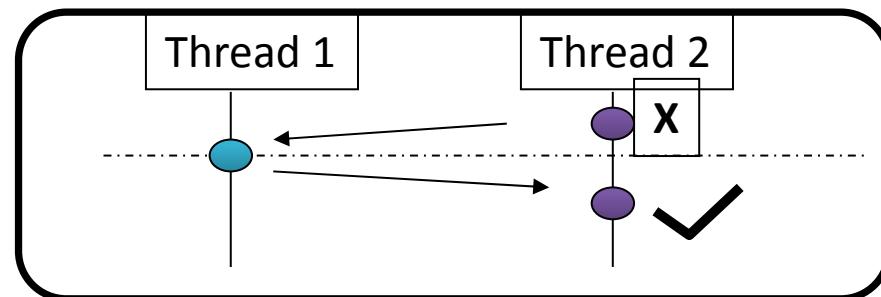
Example of atomicity violation



MySQL ha_innodb.hpp

Order violation

- The desired order between two (sets of) accesses is flipped



Example of order violation

Thread 1

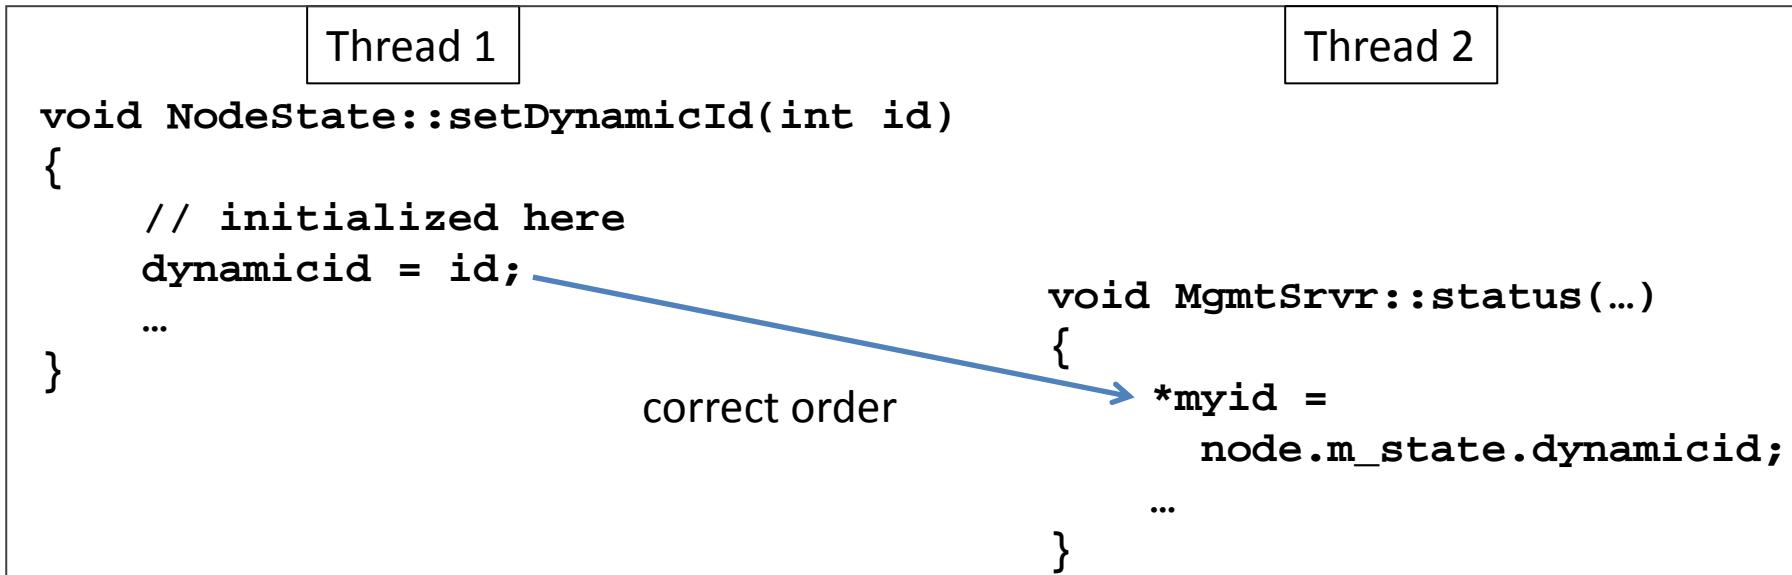
```
void NodeState::setDynamicId(int id)
{
    // initialized here
    dynamicid = id;
    ...
}
```

Thread 2

```
void MgmtSrvr::status(...)
{
    *myid =
        node.m_state.dynamicid;
    ...
}
```

MySQL NodeState.hpp

Example of order violation



MySQL NodeState.hpp

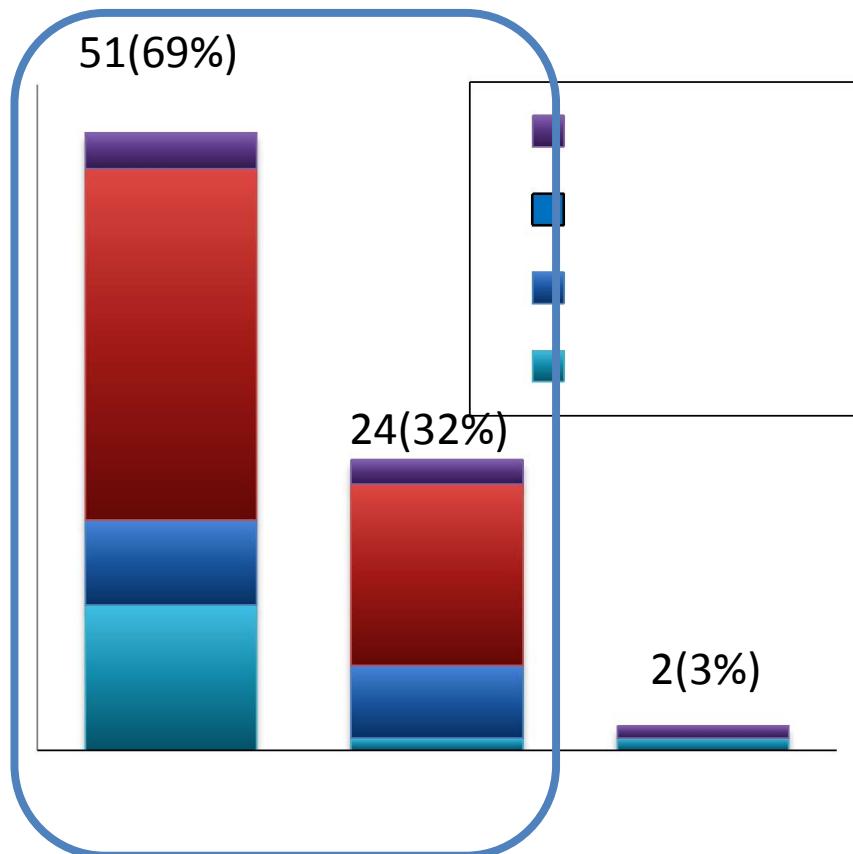
Example of order violation

Thread 1	Thread 2
<pre>void NodeState::setDynamicId(int id) { // initialized here dynamicid = id; ← ... }</pre>	<pre>void MgmtSrvr::status(...) { *myid = node.m_state.dynamicid; ... }</pre>

A blue arrow points from the assignment statement in Thread 1 to the brace of the status() function in Thread 2, with the text "buggy order" written below the arrow.

MySQL NodeState.hpp

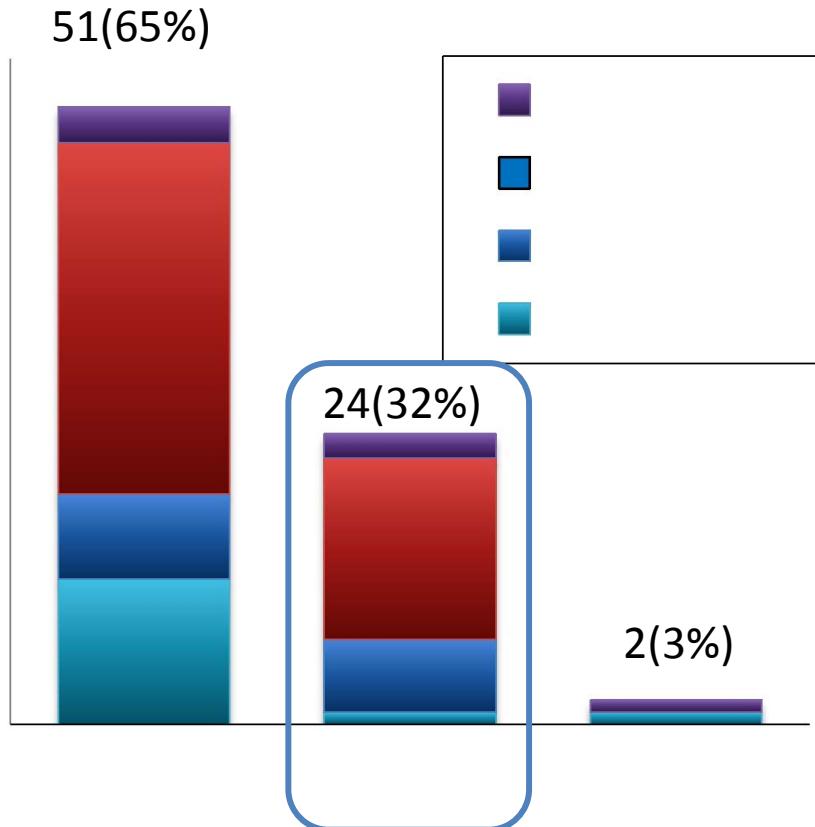
Non-deadlock bug pattern



Implications

We should focus on atomicity violation and order violation

Non-deadlock bug pattern



Implications

We should focus on atomicity violation and order violation

Bug detection tools for order violation bugs are desired

How to trigger a bug

Bug manifestation condition

A specific execution order among a smallest set of memory accesses

The bug is guaranteed to manifest, as long as the condition is satisfied

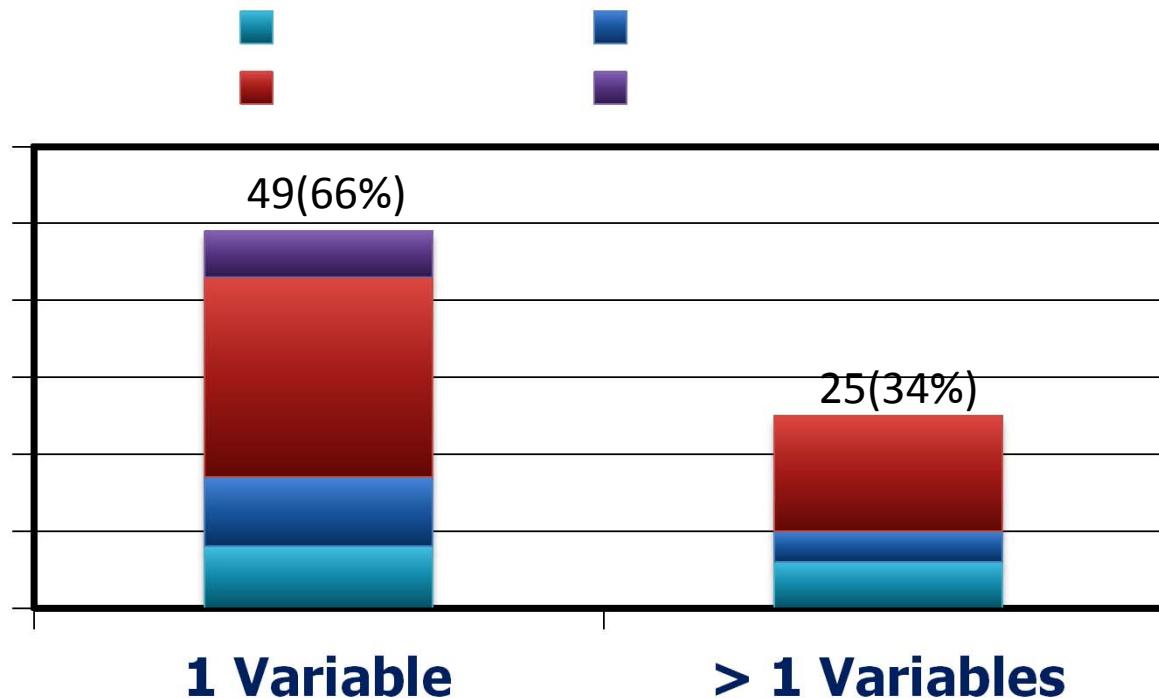
How many threads are involved?

How many variables are involved?

How many accesses are involved?

Single Variable vs. Multiple Variable

Findings



Single Variable vs. Multiple Variable

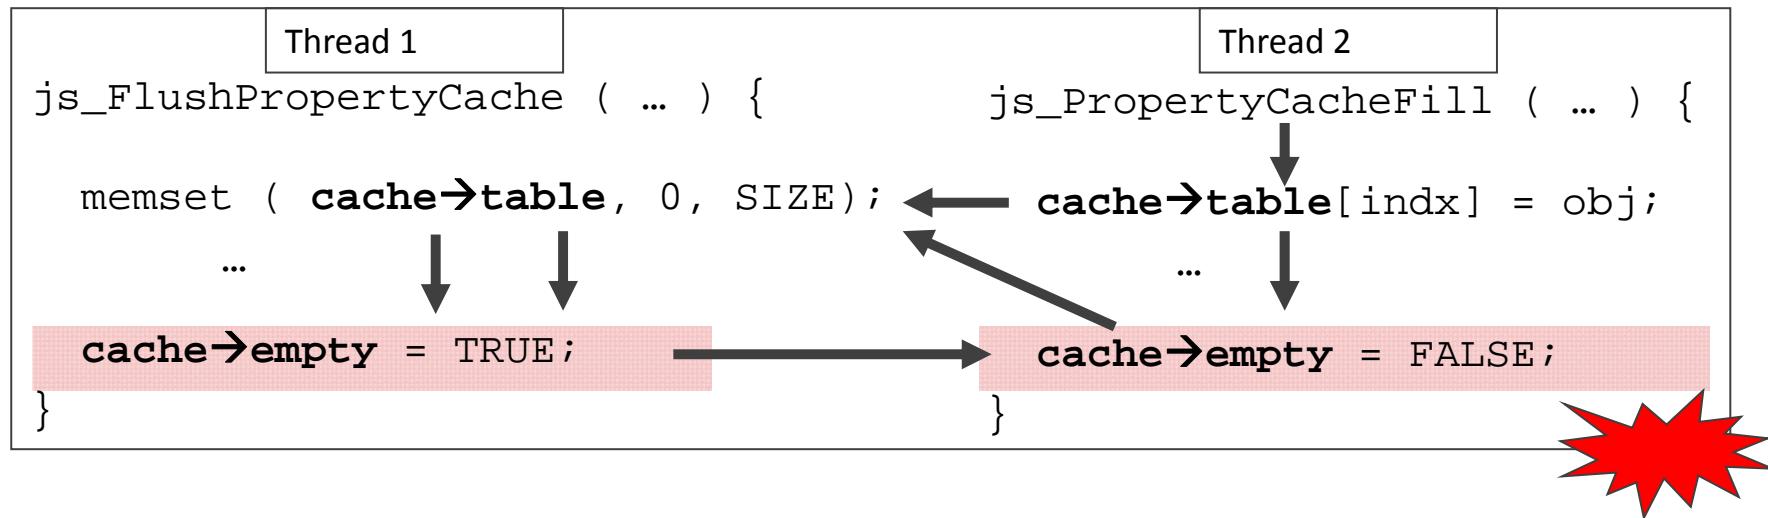
Single variables are more common

The widely-used simplification is reasonable

Multi-variable concurrency bugs are non-negligible

Techniques to detect multi-variable concurrency bugs
are needed

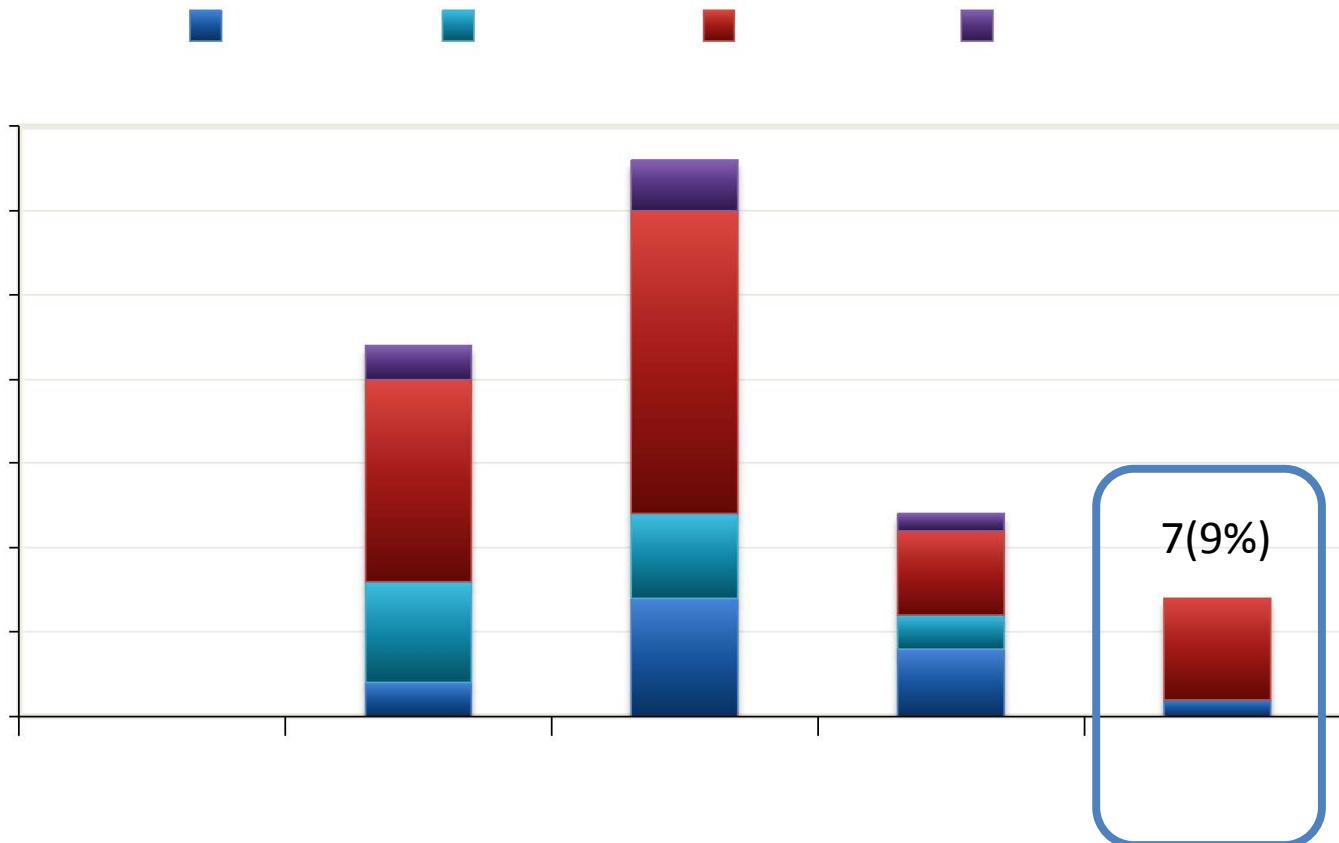
Multi-Variable Concurrency Bug Example



Control the order among accesses to any one variable
can **not** guarantee the bug manifestation

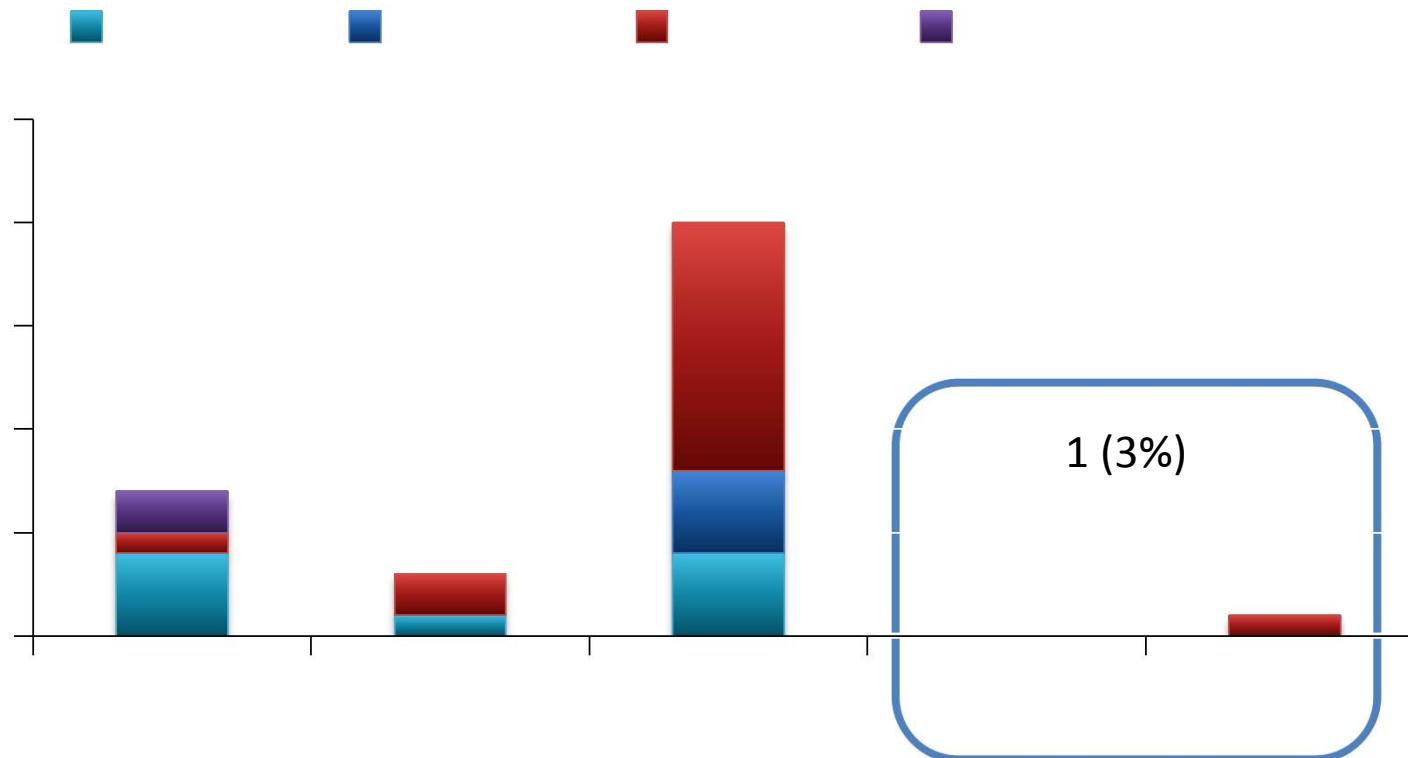
Non-deadlock bugs

Number of Accesses



Deadlock bugs

Number of Accesses



Implications

Only a few percentage bugs need more than 4 access to trigger

Concurrent program testing can focus on small groups of accesses

The testing target shrinks from exponential to polynomial

Number Threads Involved

101 out of 105 (96%) bugs involve at most two threads

Most bugs can be reliably disclosed if we check all possible interleaving between each pair of threads

Few bugs cannot

Example: Intensive resource competition among many threads causes unexpected delay

How Were Non-Deadlock Bugs Fixed?

Adding/changing locks 20 (27%)

Condition check 19 (26%)

Data-structure change 19 (26%)

Code switch 10 (13%)

Other 6 (8%)

Implications

No silver bullet for fixing concurrency bugs.

Lock usage information is not enough to fix bugs.

How Were Deadlock Bugs Fixed?

Give up resource acquisition	19 (61%)
Change resource acquisition order	7 (23%)
Split the resource to smaller ones	1 (3%)
Others	4 (13%)

Might introduce
non-dead locks

We need to pay attention to the correctness of
“fixed” deadlock bugs

Other findings

Impact of concurrency bugs

- ~ 70% leads to program crash or hang

Reproducing bugs are critical to diagnosis

Programmers lack diagnosis tools

- Most are diagnosed via code review

- Reproduce bugs are extremely hard and directly determines the diagnosing time

60% 1st-time patches contain concurrency bugs
(old or new)

Summary

Bug detection needs to look at order-violation bugs and multi-variable concurrency bugs

Testing can target at more realistic interleaving coverage goals

Fixing concurrency bugs is not trivial and not easy to get right

Support from automated tools is needed

BUGS IN EXCEPTION HANDLERS

multiple events are required

User: “Sudden outage on the entire HBase cluster.”

Event 1: Load balance: transfer Region R from slave A to B



Slave B opens R

/ Master: delete the
* ZooKeeper znode after
* the region is opened */*

try {

deleteZNode();

} catch (KeeperException e) {
cluster.abort("...");

}

Event 2: Slave B dies



R is assigned to slave C

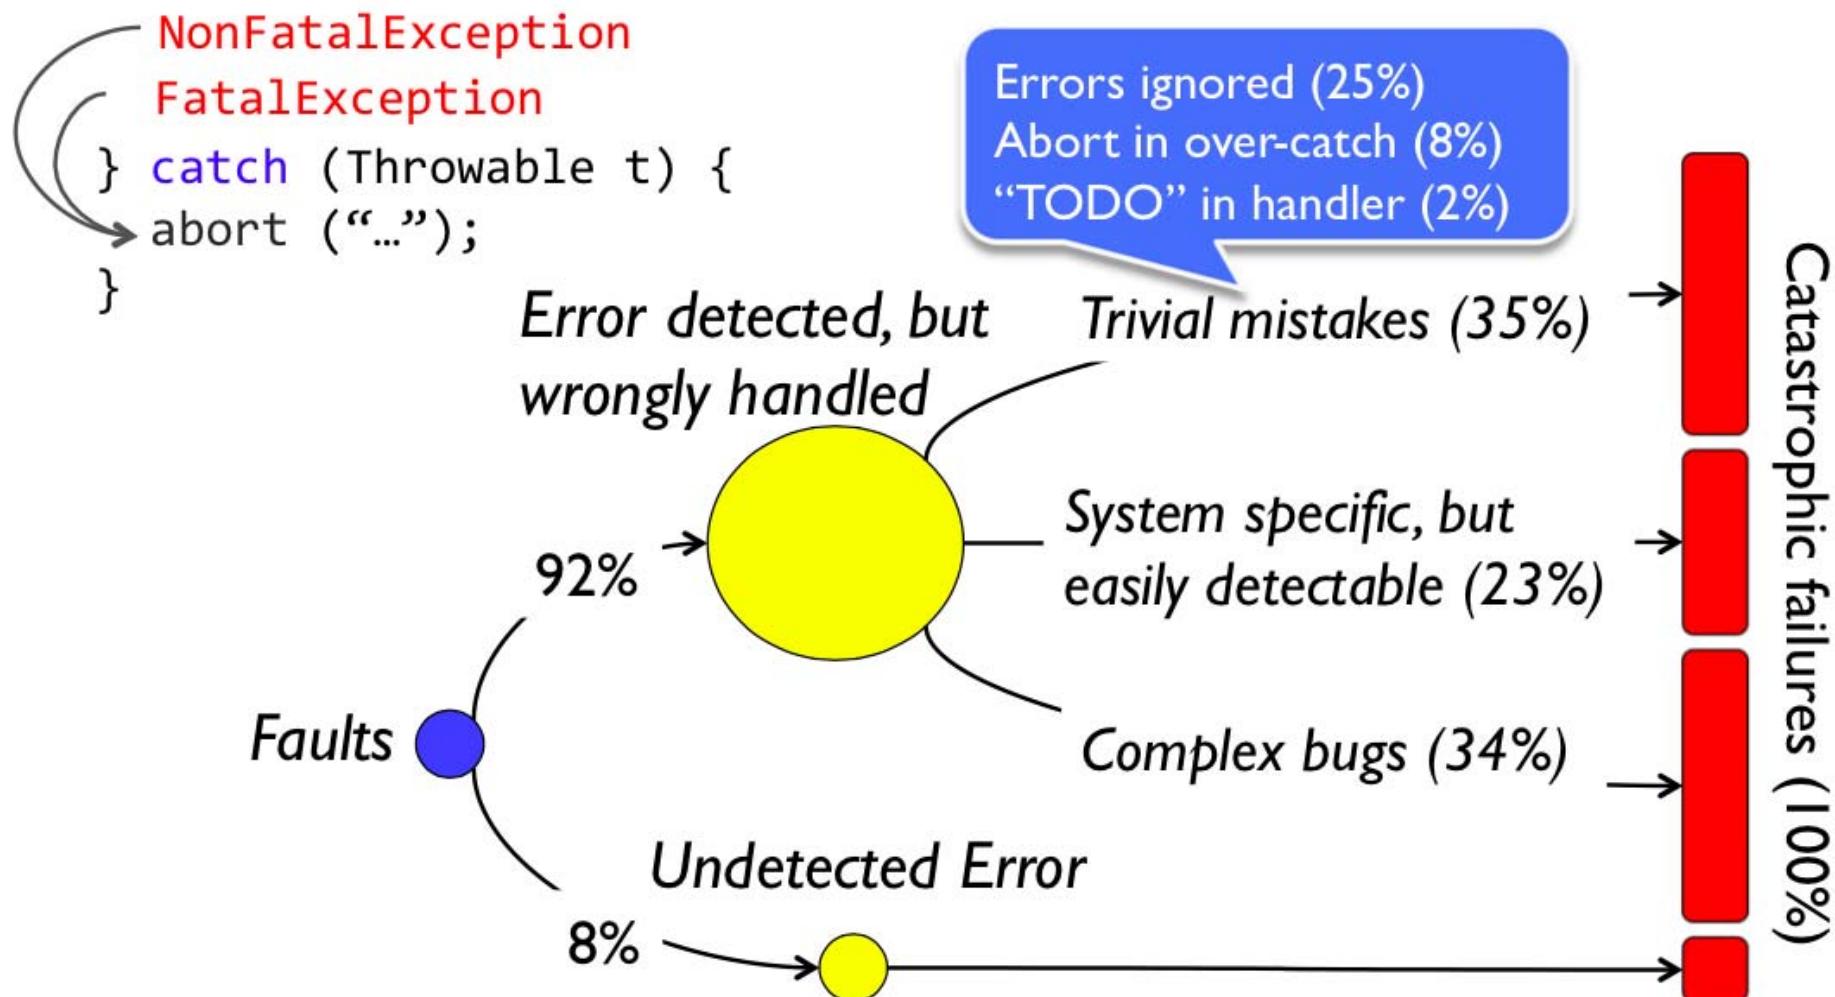


Slave C opens R

Not handled
properly

Breakdown of catastrophic failures

92% of catastrophic failures are the result of incorrect error handling



A failure caused by trivial mistake

User:

“MapReduce jobs hang when a rare Resource Manager restart occurs. *I have to ssh to every one of our 4000 nodes in a cluster and kill all jobs.*”

```
catch (RebootException) {  
    // TODO  
    LOG("Error event from RM: shutting down...");  
    + eventHandler.handle(exception_response);  
}
```

Why do developers ignore error handling?

- ▶ Developers think the errors *will never happen*
 - ▶ Code evolution may enable the errors
 - ▶ The judgment can be wrong

```
    } catch (IOException e) {  
        // will never happen  
    }
```
- ▶ Error handling is difficult
 - ▶ Errors can be returned by 3rd party libraries

```
    } catch (NoTransitionException e) {  
        /* Why this can happen? Ask God not me. */  
    }
```
- ▶ Feature development is prioritized

Other Findings

- ▶ Failures require no more than 3 nodes to manifest
 - * *almost all (98%)*
- ▶ Failures can be reproduced offline by unit tests
 - ▶ The triggering events are recorded in system log
 - * *Logs are noisy: the median of the number of log messages printed by each failure is 824.*
- ▶ Non-deterministic failures can still be deterministically reproduced
 - * *at least one part of the timing dependency can be controlled by testers*

Unexpected fun: comments in error handlers

```
/* If this happens, hell will unleash on earth. */  
  
/* FIXME: this is a buggy logic, check with alex. */  
  
/* TODO: this whole thing is extremely brittle. */  
  
/* TODO: are we sure this is OK? */  
  
/* I really thing we should do a better handling of these  
 * exceptions. I really do. */  
  
/* I hate there was no piece of comment for code  
 * handling race condition.  
 * God knew what race condition the code dealt with! */
```

OS BUGS

Bugs Cost??

Patriot missile defense system

28 **dead** soldiers, 98 wounded

Therac-25 medical device

Several people dead, others wounded

General Electric XA/21

50 million people left without water, electricity.

What Bugs Means to You?



This Lecture

An example analysis of typical OS bugs

Describe several rules on OS bug finding

Based on Invariants

Invariants in JOS

How to find bugs?

What is your belief set?

MUST set

MAY set

What is the implied sets?

Inconsistency means possible bugs!!

Trivial consistency: NULL pointers

$*p$ implies MUST belief:

p is not null

A check ($p == \text{NULL}$) implies two MUST beliefs:

POST: p is null on true path, not null on false path

PRE: p was unknown before check

```
/* 2.4.1: drivers/isdn/svmb1/capidrv.c */
if (!card)
    printk(KERN_ERR, "capidrv-%d: ...", card->contrnr...)
```

```
/* drivers/net/wan/sdla_chdlc.c:3948 */
if (!card){
    lock_adapter_irq(&card->wandev.lock, &smp_flags);
    card->tty=NULL;
```

Null pointer fun

Use-then-check

```
/* 2.4.7: drivers/char/mxser.c */
struct mxser_struct *info = tty->driver_data;
unsigned flags;
if(!tty || !info->xmit_buf)
    return 0;
```

Contradiction/redundant checks

```
/* 2.4.7/drivers/video/tdfxfb.c */
fb_info.regbase_virt = ioremap_nocache(...);
if(!fb_info.regbase_virt)
    return -ENXIO;
fb_info.bufbase_virt = ioremap_nocache(...);

if(!fb_info.regbase_virt) {
    iounmap(fb_info.regbase_virt);
```

Internal Consistency: finding security holes

Applications are bad:

Rule: “do not dereference user pointer `<p>`”

One violation = security hole

Big Problem: which are the user pointers???

Sol’n: forall pointers, cross-check two OS beliefs

`“*p” implies safe kernel pointer`

`“copyin(p)/copyout(p)” implies dangerous user pointer`

Error: pointer `p` has both beliefs.

Statistical: Deriving deallocation routines

Use-after free errors are horrible.

Problem: lots of undocumented sub-system free functions

Soln: derive behaviorally: pointer “p” not used after call
“foo(p)” implies **MAY** belief that “foo” is a free function

Conceptually: Assume all functions free all arguments
(in reality: filter functions that have suggestive names)

A bad free error

```
/* drivers/block/cciss.c:cciss_ioctl */  
if (ioccommand.Direction == XFER_WRITE) {  
    if (copy_to_user(...)) {  
        cmd_free(NULL, c);  
        if (buff != NULL) kfree(buff);  
        return( -EFAULT );  
    }  
}  
if (ioccommand.Direction == XFER_READ) {  
    if (copy_to_user(...)) {  
        cmd_free(NULL, c);  
        kfree(buff);  
    }  
}  
cmd_free(NULL, c);  
if (buff != NULL) kfree(buff);
```

“A must be followed by B”

“a(); ... b();” implies MAY belief that a() follows b()

You might believe a-b paired, or might be a coincidence

Checking derived lock functions

```
/* fs/proc/inode.c:41:de_put: */
lock_kernel();
Simplest: if (!de->count) {
    printk("de_put: entry already free!\n");
    return;
}
unlock_kernel();
```

Evilest:

```
/* 2.4.1: drivers/sound/trident.c:trident_release:
lock_kernel();
card = state->card;
dmabuf = &state->dmabuf;
VALIDATE_STATE(state);

#define VALIDATE_MAGIC(FOO,MAG)
({
    if (!(FOO) || (FOO)->magic != MAG) {
        printk(invalid magic, __FUNCTION__);
        return -ENXIO;
    }
})
#define VALIDATE_STATE(a) VALIDATE_MAGIC(a,TRIDENT STATE MAGIC)
```

General Rules

Do **not** call blocking functions with interrupts disabled or spin lock held

check for **NULL** results

Do **not** allocate large stack variables

Do not **re-use** already-allocated memory

Check **user** pointers before using them in kernel mode

Release acquired locks

Unstated Rules in JOS

Interrupts are **disabled** in kernel mode

Only **env 1** has access to disk

All registers are saved & restored on context switch

Application code is **never** executed with CPL 0

Don't **allocate** an already-allocated physical page

Propagate error messages to user applications
(e.g., out of resources)

Unstated Rules in JOS

A spawned program should have open **only** file descriptors 0, 1, and 2.

User pointers should be run through TRUP before used by the kernel

Next Class

Data Race

Deadlock

Data Race & DeadLock

Rong Chen

Outline

Happens-before Race Detection

Lockset-Based Race Detection

Deadlocks

Detecting deadlocks using Lockset Analysis

Data Race Detection

What is Race?



Data Race

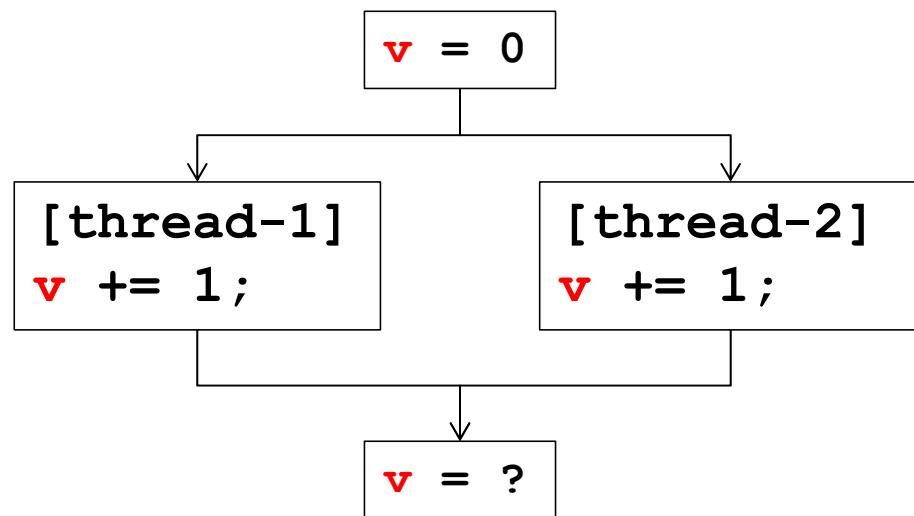
An undesirable situation that occurs when a device or system attempts to perform **two or more** operations at the **same time**, but the operations must be done in the **proper sequence** in order to be done correctly

- multithread
- distributed Programs

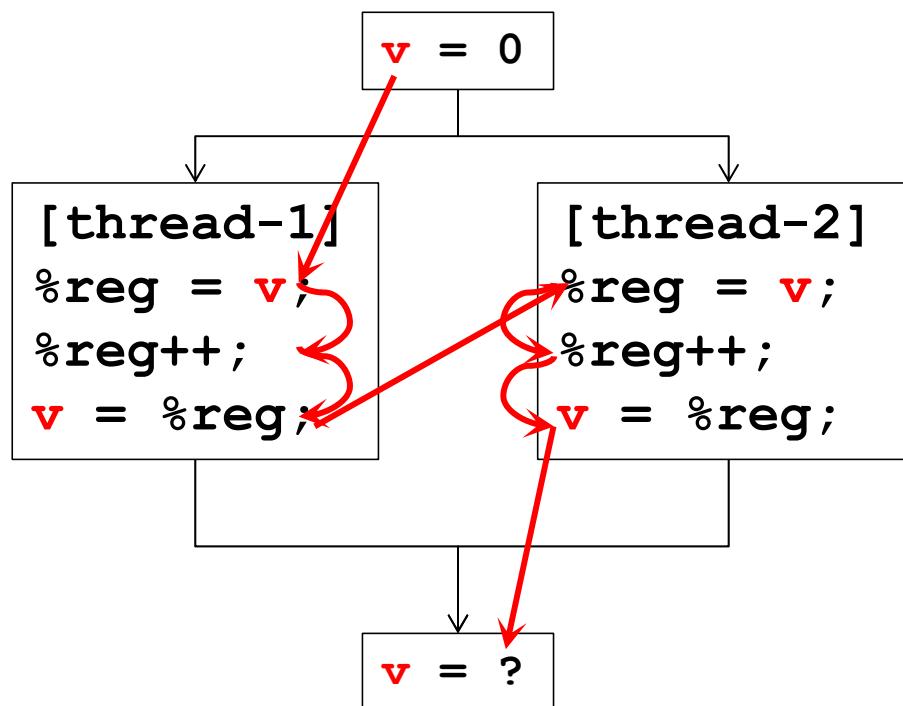
Key reason

- separate processes or threads of execution depends on **same shared state**

Data Race



Data Race



v = 0 (memory)

T1 reads 0 to %reg

T1 increments %reg to 1

T1 stores %reg to v (v = 1)

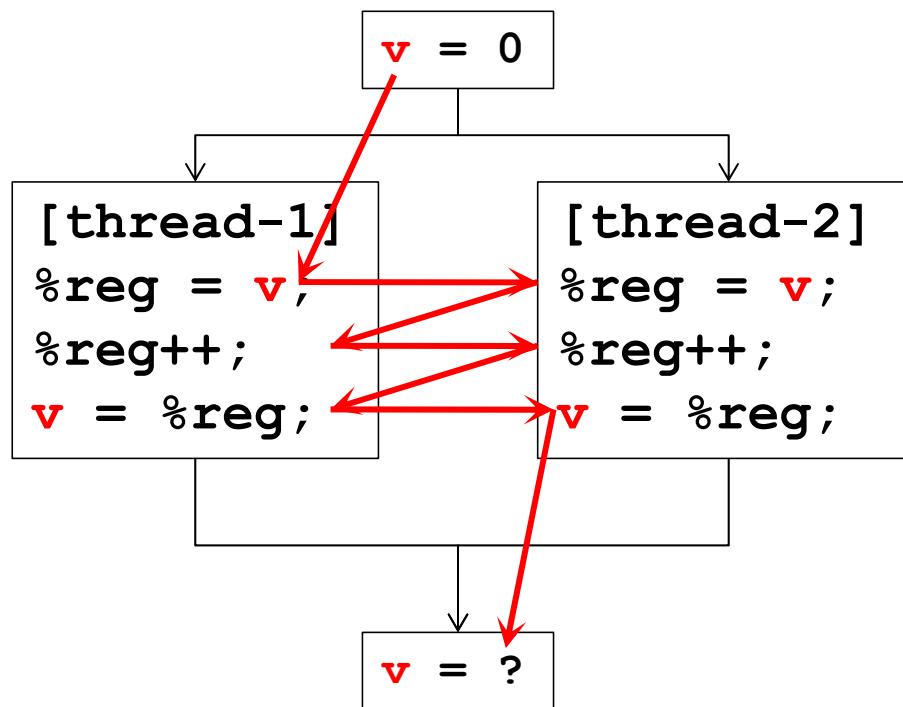
T2 reads 1 to %reg

T2 increments %reg to 2

T2 stores %reg to v (v = 2)

v = 2; (memory)

Data Race



v = 0 (memory)

T1 reads 0 to %reg

T2 reads 0 to %reg

T1 increments %reg to 1

T2 increments %reg to 1

T1 stores %reg to v (v = 1)

T2 stores %reg to v (v = 1)

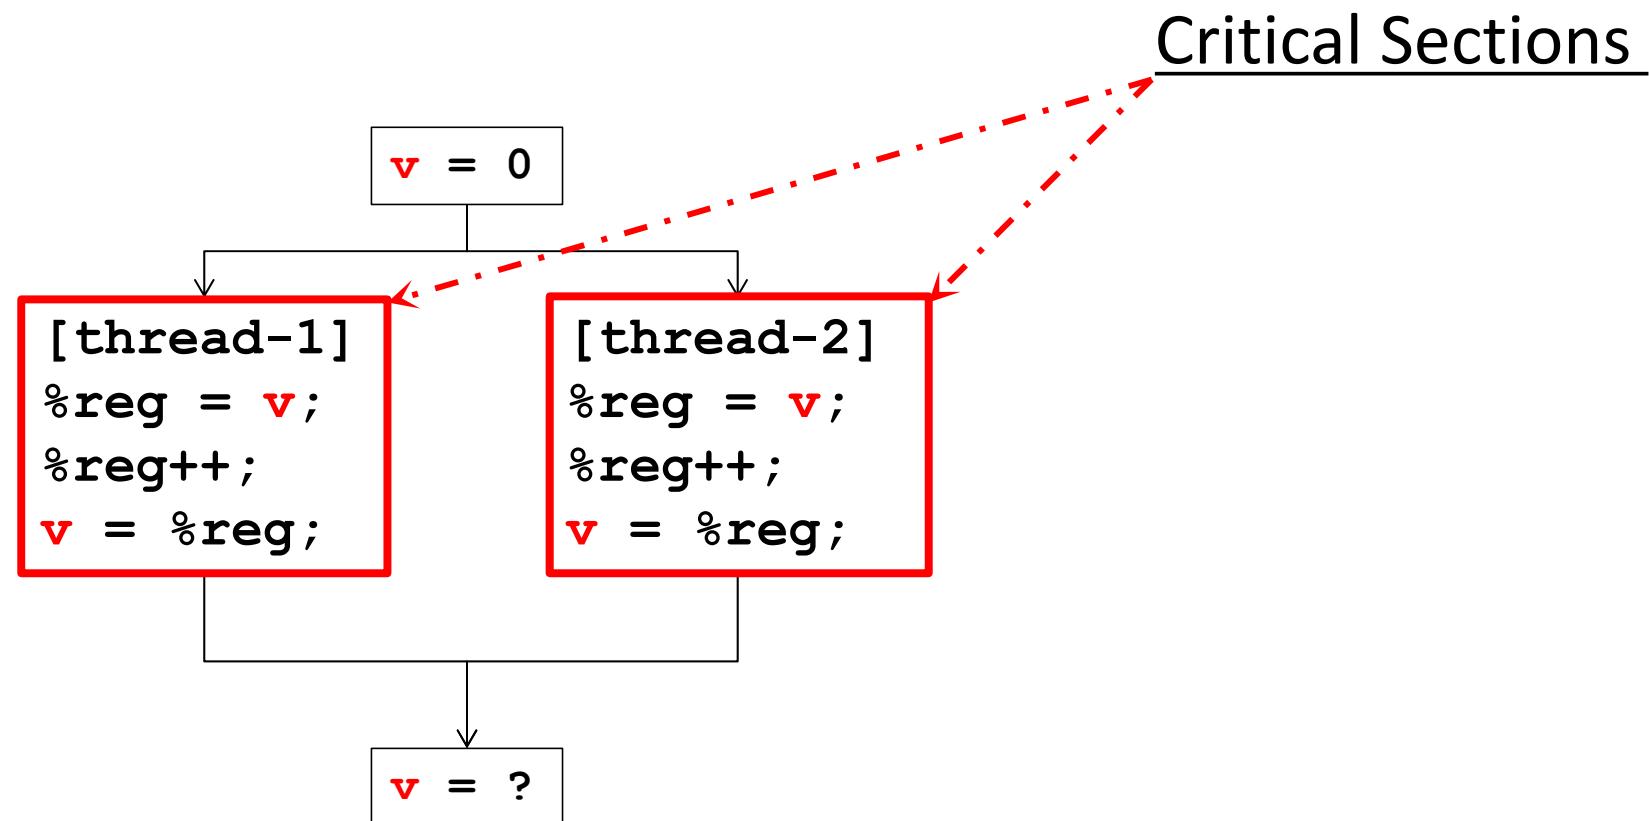
v = 1; (memory)

Data Race

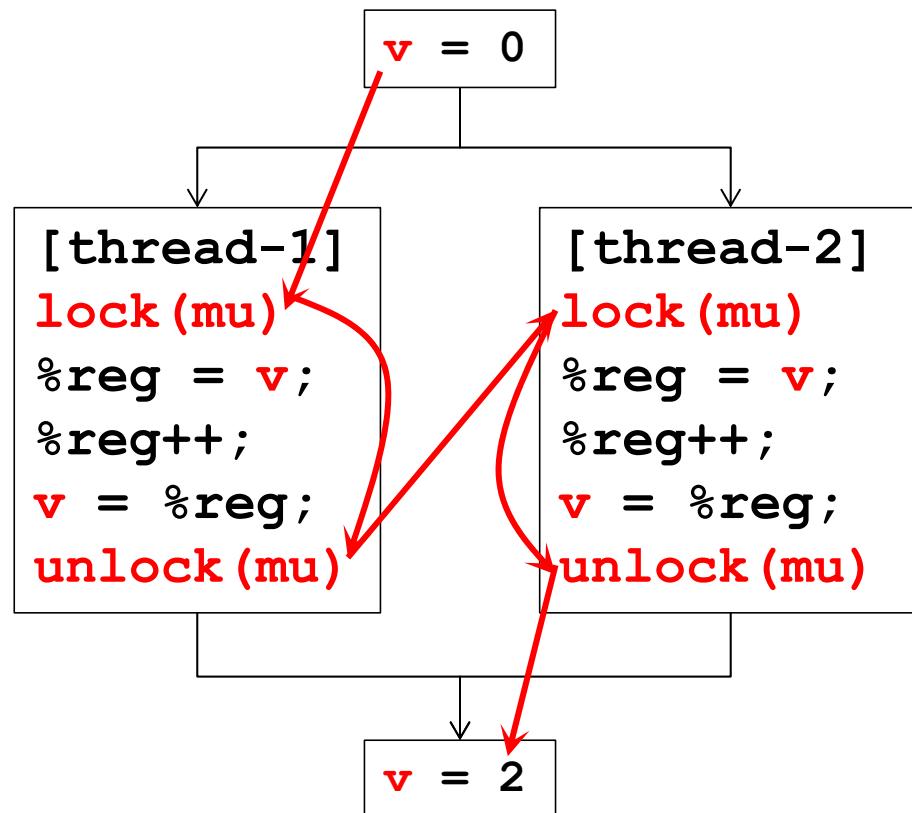
The definition of data race:

1. two concurrent threads access a **shared variable**
2. at least one access is a **write**
3. the threads use no explicit mechanism to prevent the accesses from being **simultaneous**.

Remove Race Condition



Remove Race Condition



Mutually-exclusive

- lock & unlock
- cli & sti

Data Race Bug

Thread-1

```
OpenInputStream()
{
    PostEvent();
    ...
    m_inputStream = ...
    ...
}
```

*file:
nsSocketTransport.cc*

Thread-2

```
ProcessCurrentURL()
{
    WaitEvent();
    ...
    if (m_inputStream) {
        AsyncRead(m_inputStream);
    }
    ...
}
```

file: nsImapProtocol.cpp

*Data race bug in Mozilla

Data Race Detectors

Two major categories:

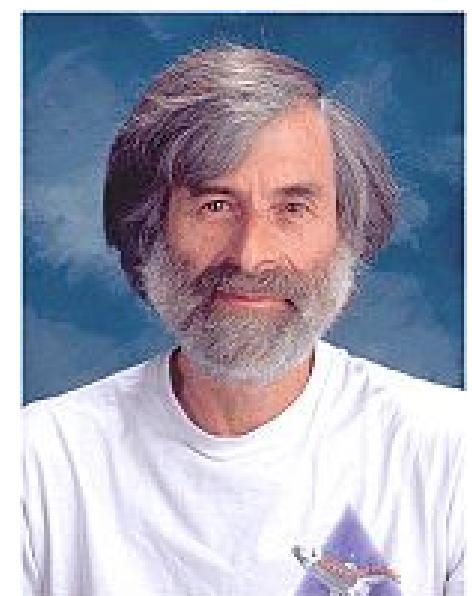
- Happens-before based
- Lockset based

HAPPENS-BEFORE BASED

Happens-before Relation

The **happens-before** relation is a means of **ordering** events based on the causal relationship of pairs of events in a concurrent system

- denoted: →
- Formulated by Leslie Lamport
- strict **partial order** on events
- without using physical clocks



*Times, Clocks, and Ordering of Events in a Distributed System

Rules in HB Relation

For the events A and B,

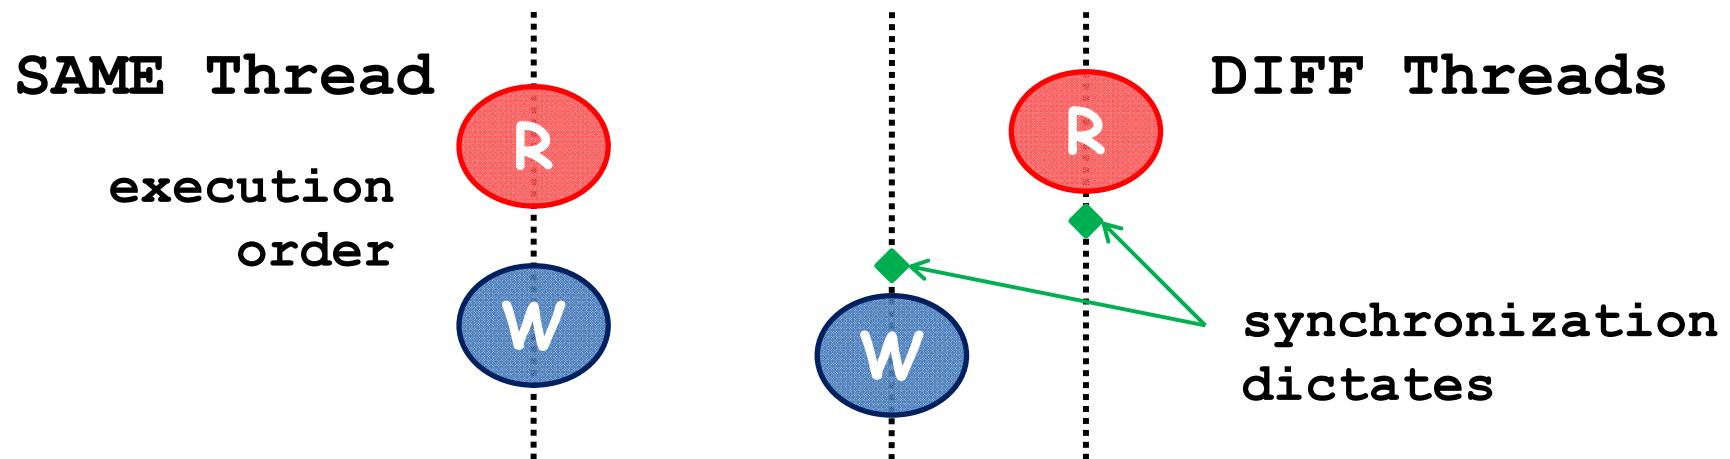
HB1: On the **same** sequential thread,

$A \rightarrow B$ if A executes before B.

HB2: On the **different** threads,

$A \rightarrow B$ if there is a synchronization that dictates A precedes B.

HB3: If $A \rightarrow B$ and $B \rightarrow C$ then $A \rightarrow C$

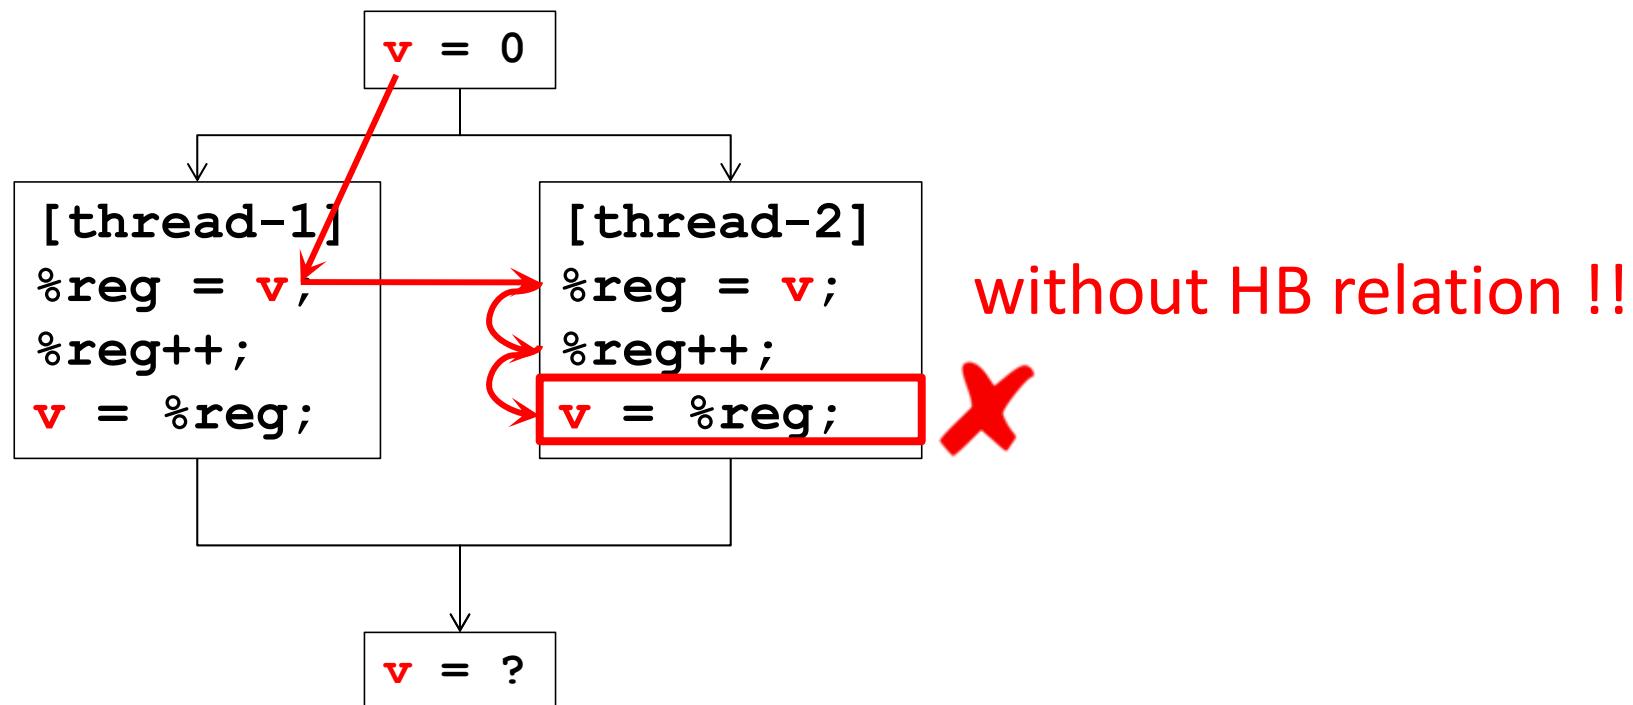


HB based Detectors

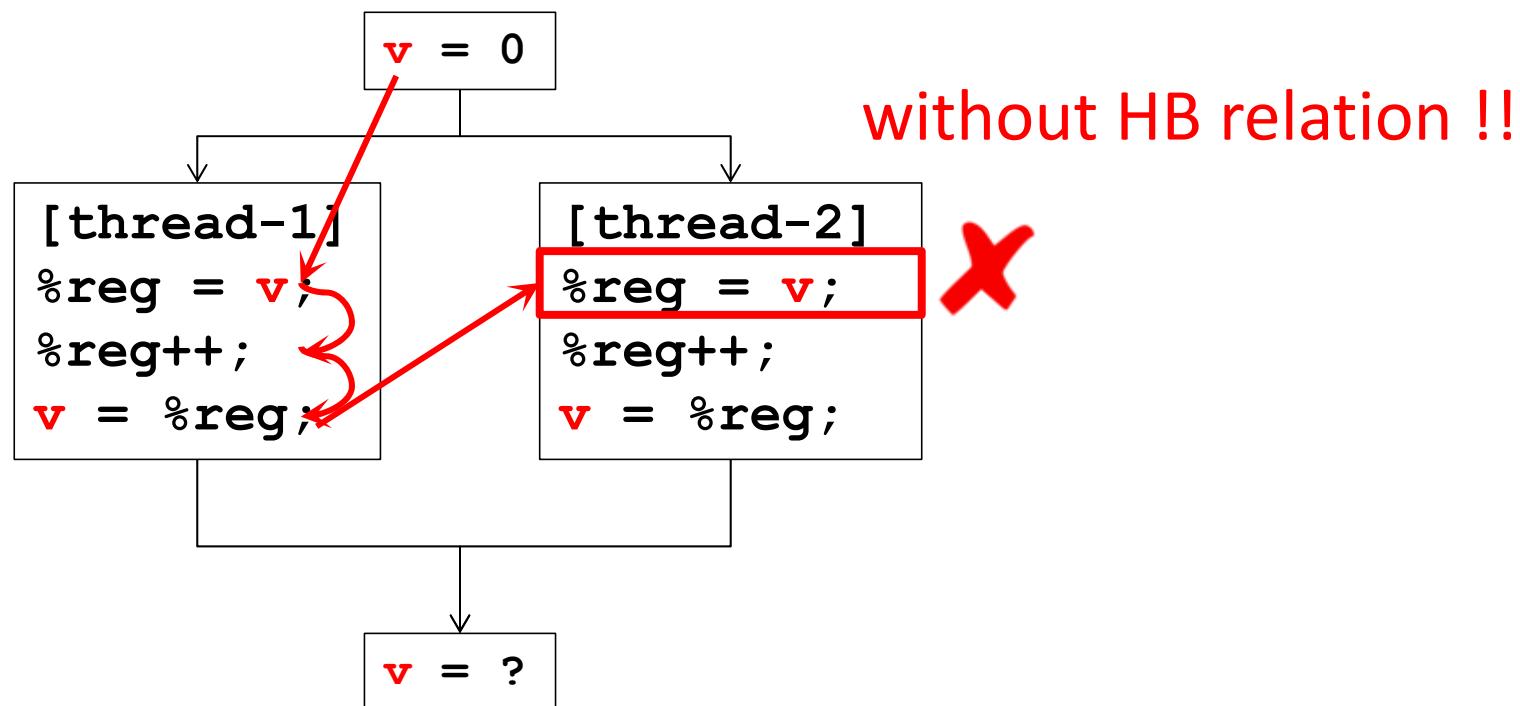
The definition of data race:

1. a pair of accesses to the **same** memory location
2. at least one access is a **write**
3. neither one **happens-before** the other

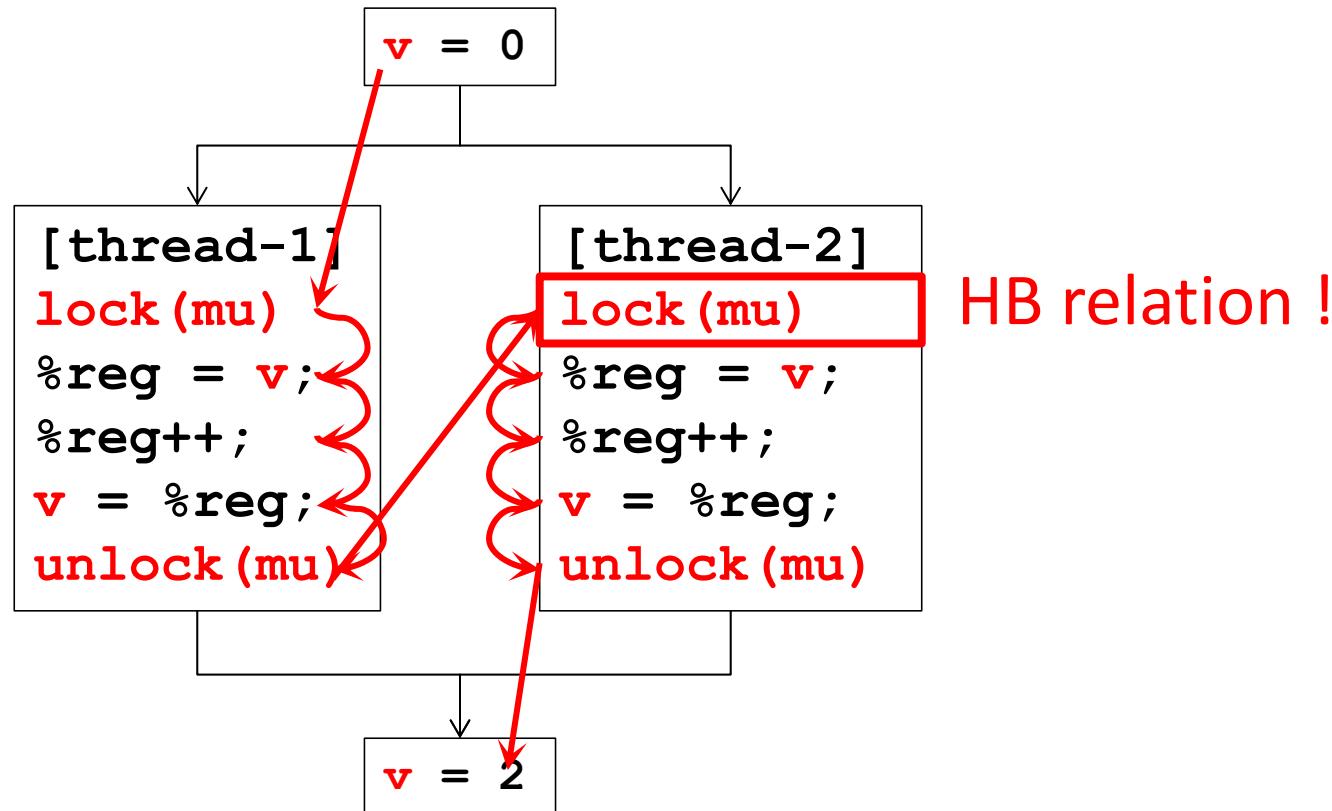
Example



Example



Example



Pros and Cons

Pros:

Detect **true** data race

Cons:

Difficult to implement efficiently

- Each thread, shared-memory location, and concurrent access

Depend on the interleaving produced by
the scheduler

- Miss data race

LOCKSET BASED

Lock

The definition of lock:

1. a synchronization object used for **mutual exclusion**
2. a lock is either **available** or **owned** by a thread.
3. the operations on a lock **m** are **lock(m)** and **unlock(m)**

Lockset based Detectors

The definition of data race:

1. a pair of accesses to the **same** memory location
2. at least one access is a **write**
3. **No** lock protects all accesses to the same data

Lockset

Locking Discipline

- A programming policy that ensures the absence of data races
 - e.g. “**every variable shared between threads is protected by a mutual exclusion lock**”

Principle

- Check all shared memory accesses follow a consistent lock discipline
 - monitors all reads and writes, and infer the protection relation from the execution history

Algorithm of Lockset

The summary:

1. Let `locks_held(t)` be the set of locks held by thread `t`
2. For each `v`, initialize `C(v)` to the set of all locks
3. On each access to `v` by thread `t`,
 set `C(v) := C(v) ∩ locks_held(t)`
 if `C(v) = {}`, then issue a warning

Example

<u>Programs</u>	<u>locks held(t)</u>	<u>C(v)</u>
	{ }	{ mu1 , mu2 }
lock(mu1);	{ mu1 }	
lock(mu2);	{ mu1 , mu2 }	
v = v + 1;		{ mu1 , mu2 }
unlock(mu2);		
	{ mu1 }	
v = v + 2;		{ mu1 }
unlock(mu1);		
	{ }	
lock(mu2);	{ mu2 }	
v = v + 1;		{ } Warning!!
unlock(mu2);		

Three Challenges

#1 Initialization

- Shared variables are frequently initialized without holding a lock

#2 Read-only Shared Variable

- write once and read all the time without lock

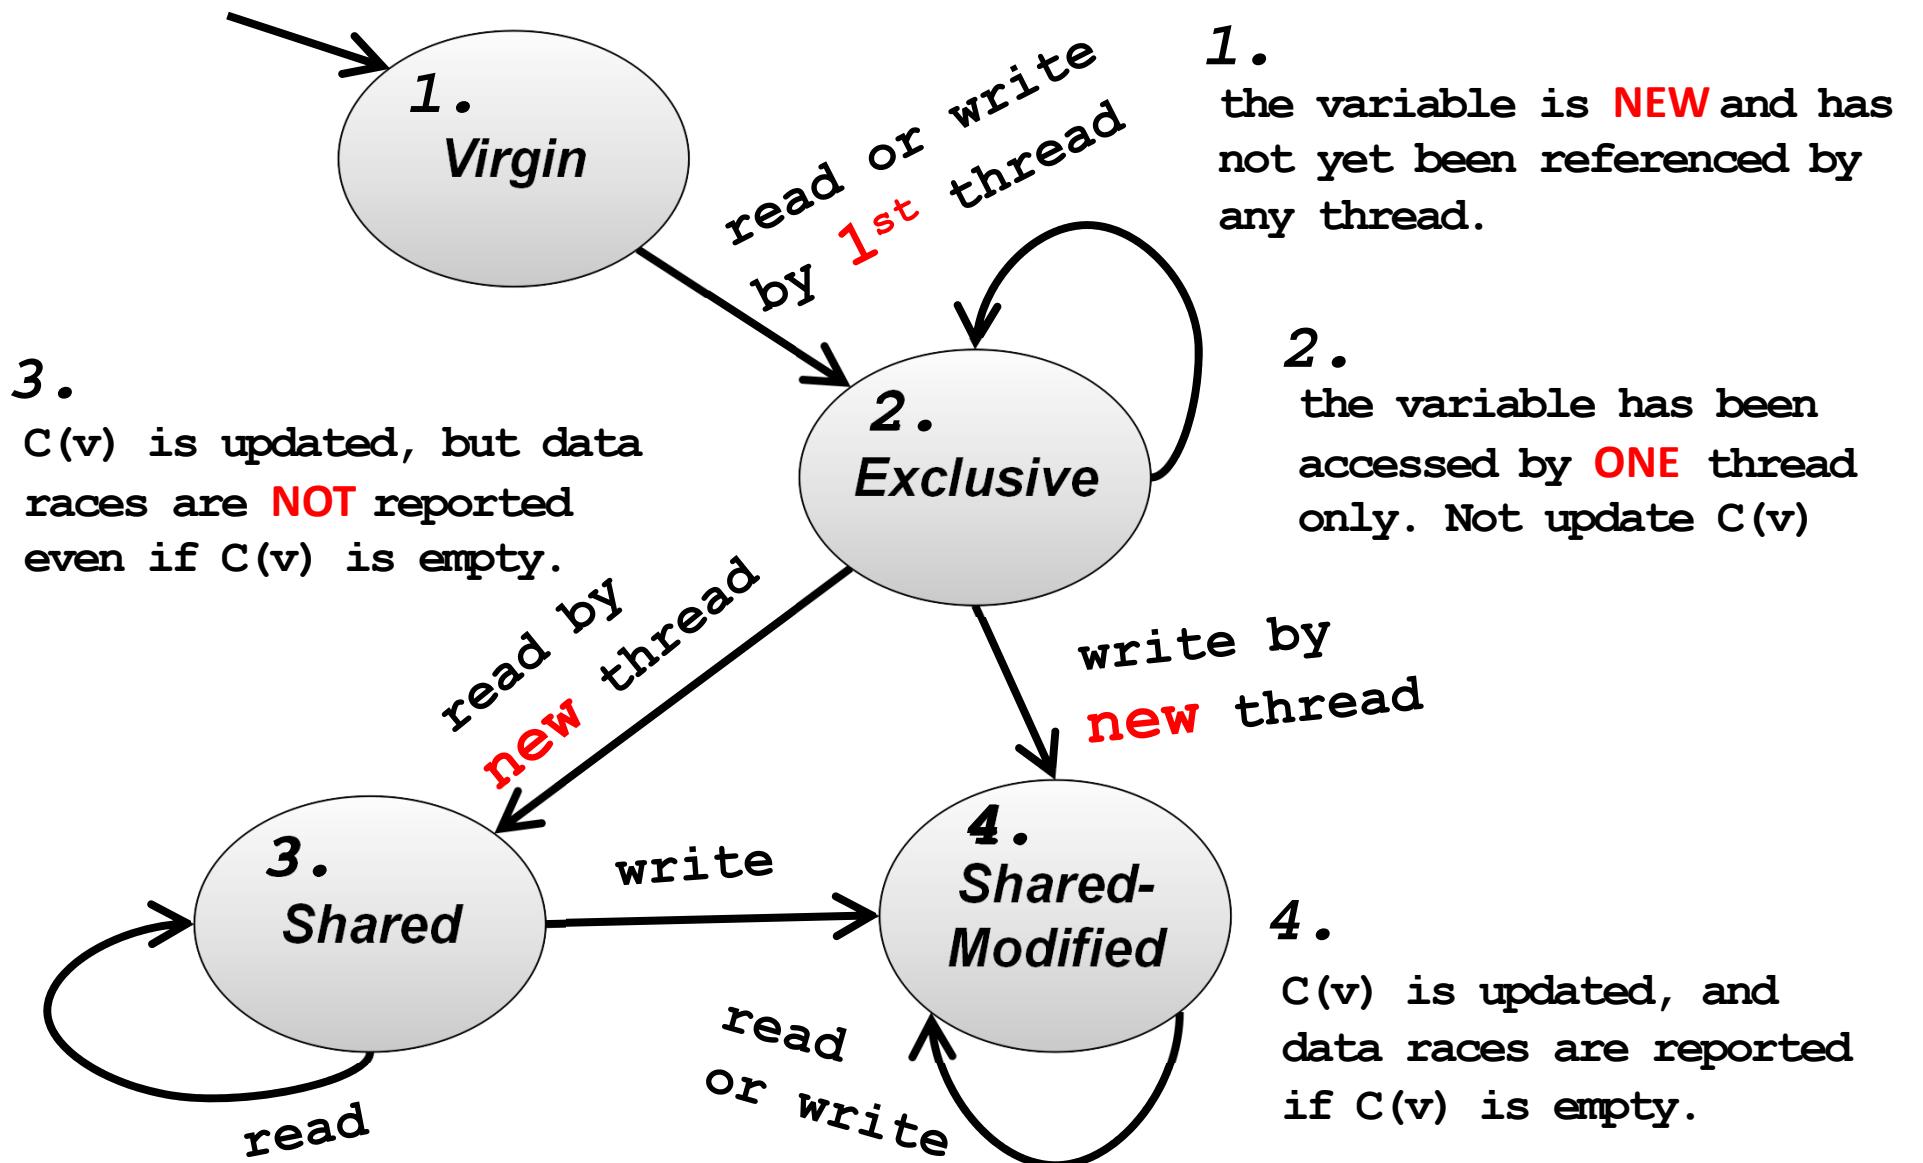
#3 Read-Write Lock

Solution to #1 and #2

Delay the refinement until the shared variable has been initialized

- No easy way to know when initialization is done
- Heuristic way: the first time to be accessed by the 2nd thread

Solution to #1 and #2



Solution to #3

Modify the algorithm:

1. Let `locks_held(t)` and `write_locks_held(t)` be the set of locks held in any mode and `write` mode by thread `t`
2. For each `v`, initialize `Cr(v)` and `Cw(v)` to the set of all locks
3. On each `read` to `v` by thread `t`,
set `Cr(v) := Cr(v) ∩ locks_held(t)`
if `Cr(v) = {}`, then issue a warning
4. On each `write` to `v` by thread `t`,
set `Cw(v) := Cw(v) ∩ write_locks_held(t)`
if `Cw(v) = {}`, then issue a warning

Pros and Cons

Pros:

More efficient way to detect data race

Predict data race that have not manifest

Cons:

Exists report false positive

- Memory reuse

Limits the synchronization method to lock

Deadlock

What is a deadlock



"Thanks for coming in. We'll get back to you as soon as we lower our expectations."

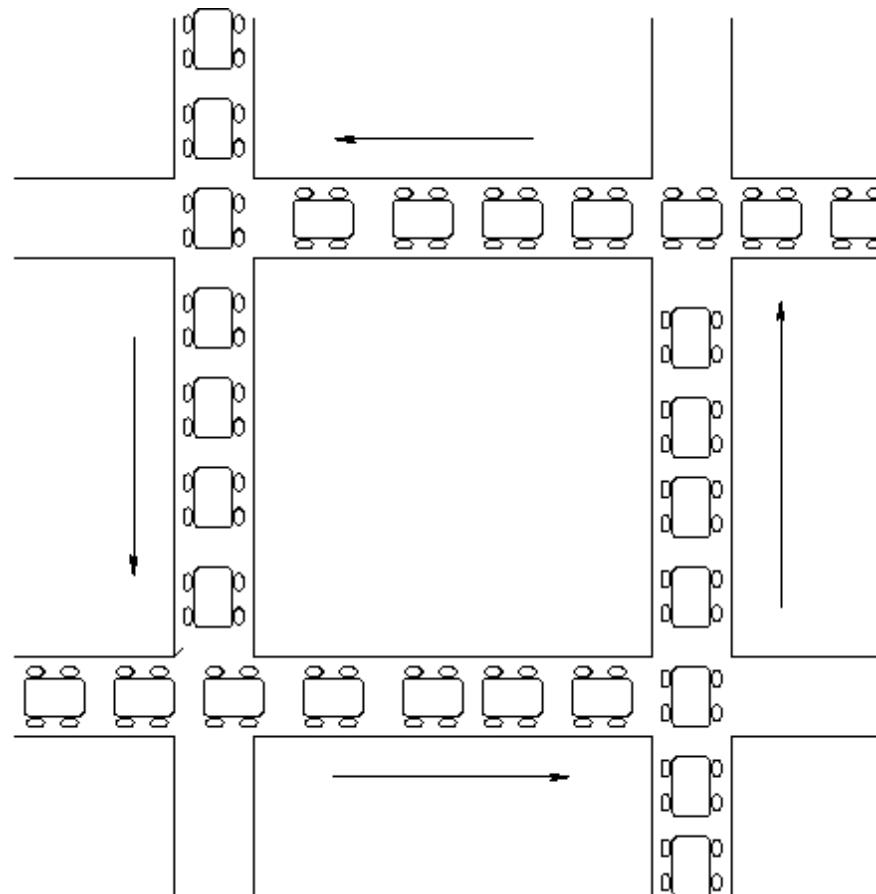
This guy has no experience!!!

How can I get experiences without a job!!!



Deadlock Definition

A set of processes is **deadlocked** when
every process in the set is waiting for an event that can
only be generated by some process in the set

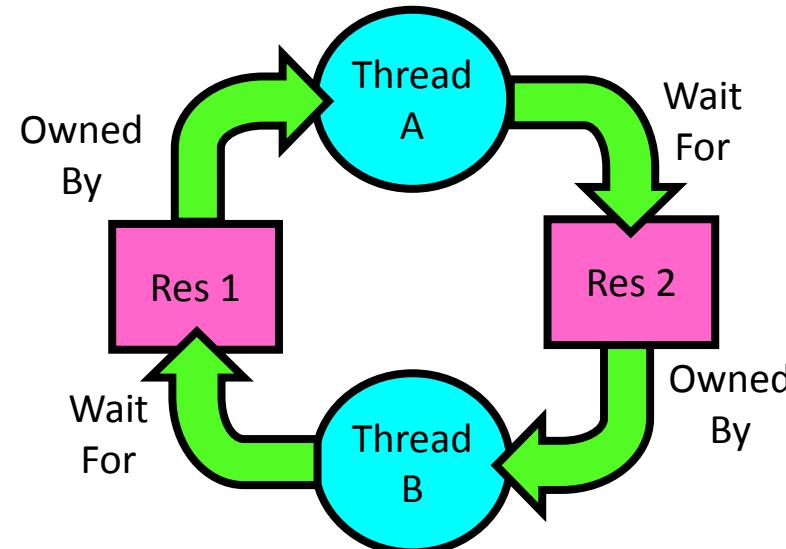


Deadlock

Circular waiting for resources

Thread A owns Res 1 and is waiting for Res 2

Thread B owns Res 2 and is waiting for Res 1



Four requirements for Deadlock

Mutual exclusion

Only one thread at a time can use a resource.

Hold and wait

Thread holding at least one resource is waiting to acquire additional resources held by other threads

No preemption

Resources are released only voluntarily by the thread holding the resource, after thread is finished with it

Circular wait

There exists a set $\{T_1, \dots, T_n\}$ of waiting threads

T_1 is waiting for a resource that is held by T_2

T_2 is waiting for a resource that is held by T_3

...

T_n is waiting for a resource that is held by T_1

A Graph Theoretic Model of Deadlock

Basic components of any resource allocation problem

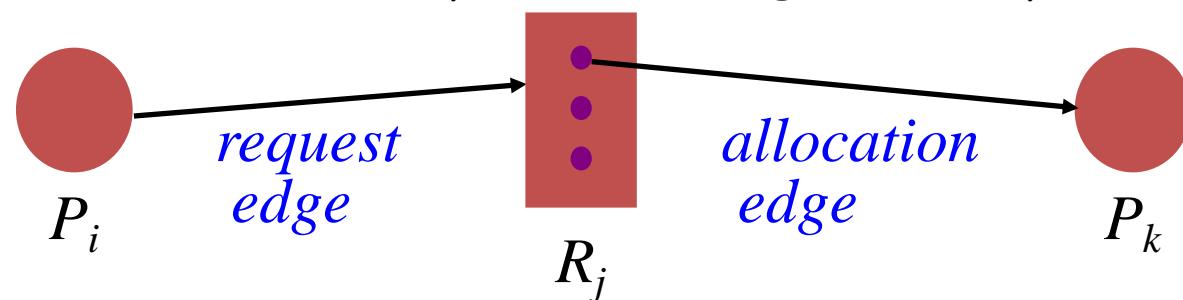
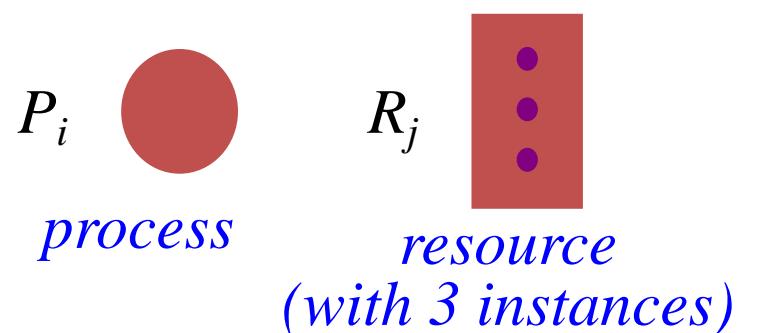
Processes and resources

Model the state of a computer system as a directed graph (called resource allocation graph, or RAG)

$$G = (V, E)$$

V = the set of vertices =
 $\{P_1, \dots, P_n\} \cup \{R_1, \dots, R_m\}$

E = the set of edges =
 $\{\text{edges from a resource to a process}\} \cup \{\text{edges from a process to a resource}\}$

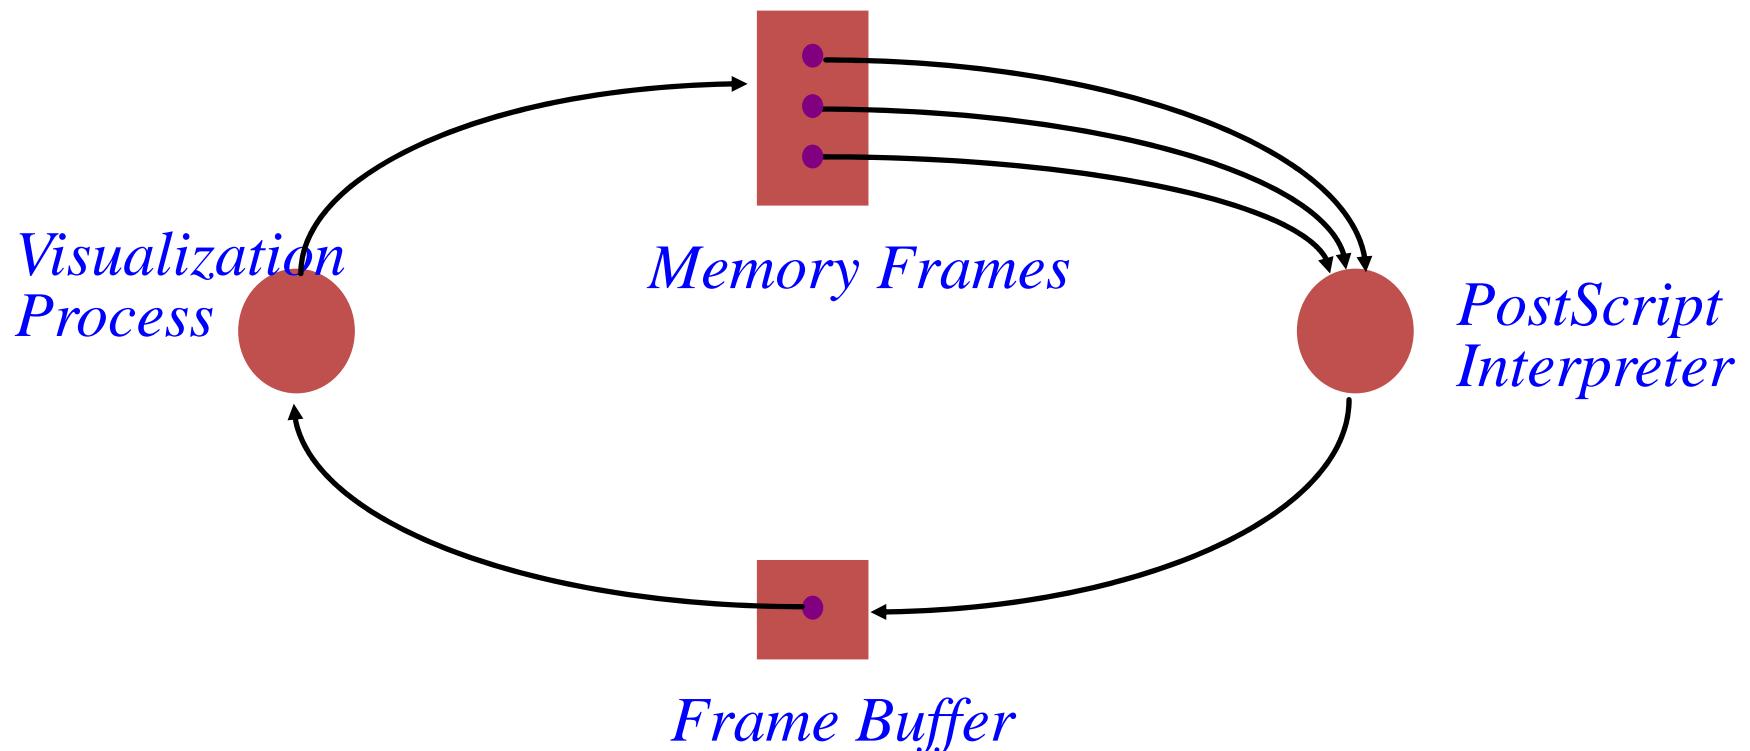


Resource Allocation Graphs

Example

A PostScript interpreter that is waiting for the **frame buffer lock** and a visualization process that is waiting for memory

$$V = \{PS\ interpret, visualization\} \cup \{memory\ frames, frame\ buffer\ lock\}$$

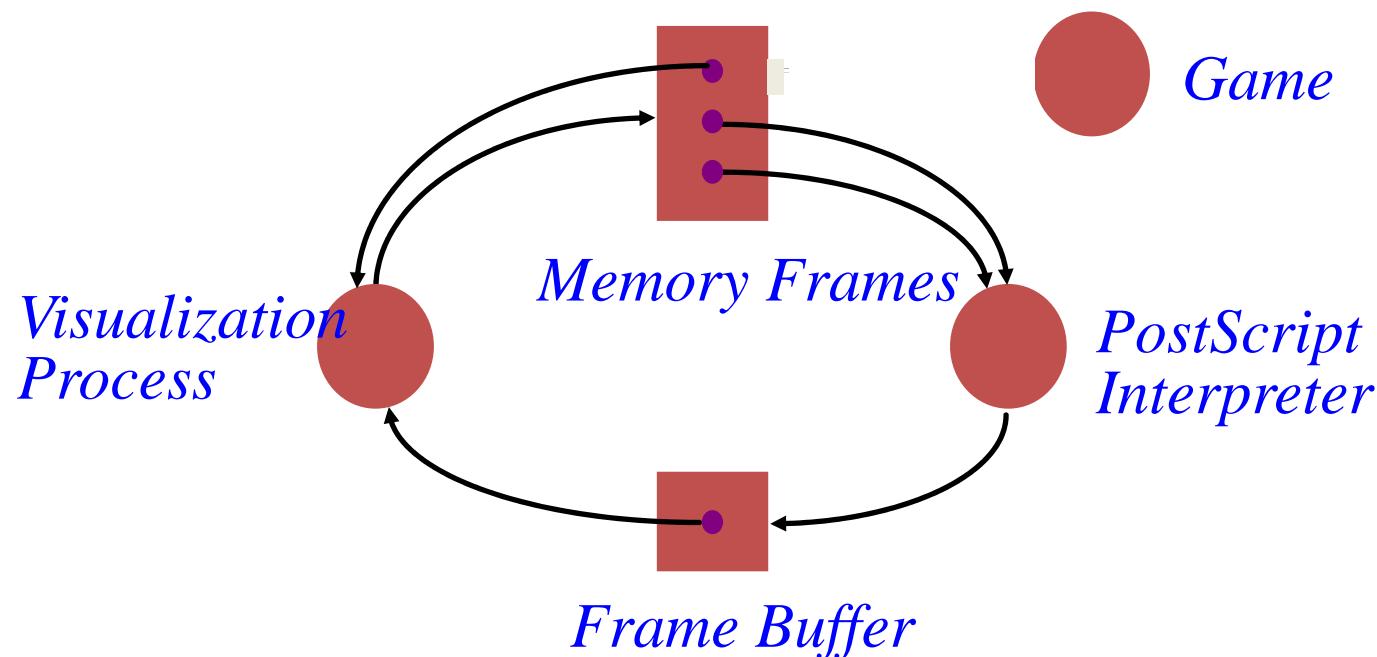


Resource Allocation Graphs & Deadlock

Theorem: If a resource allocation graph does not contain a **cycle** then no processes are deadlocked

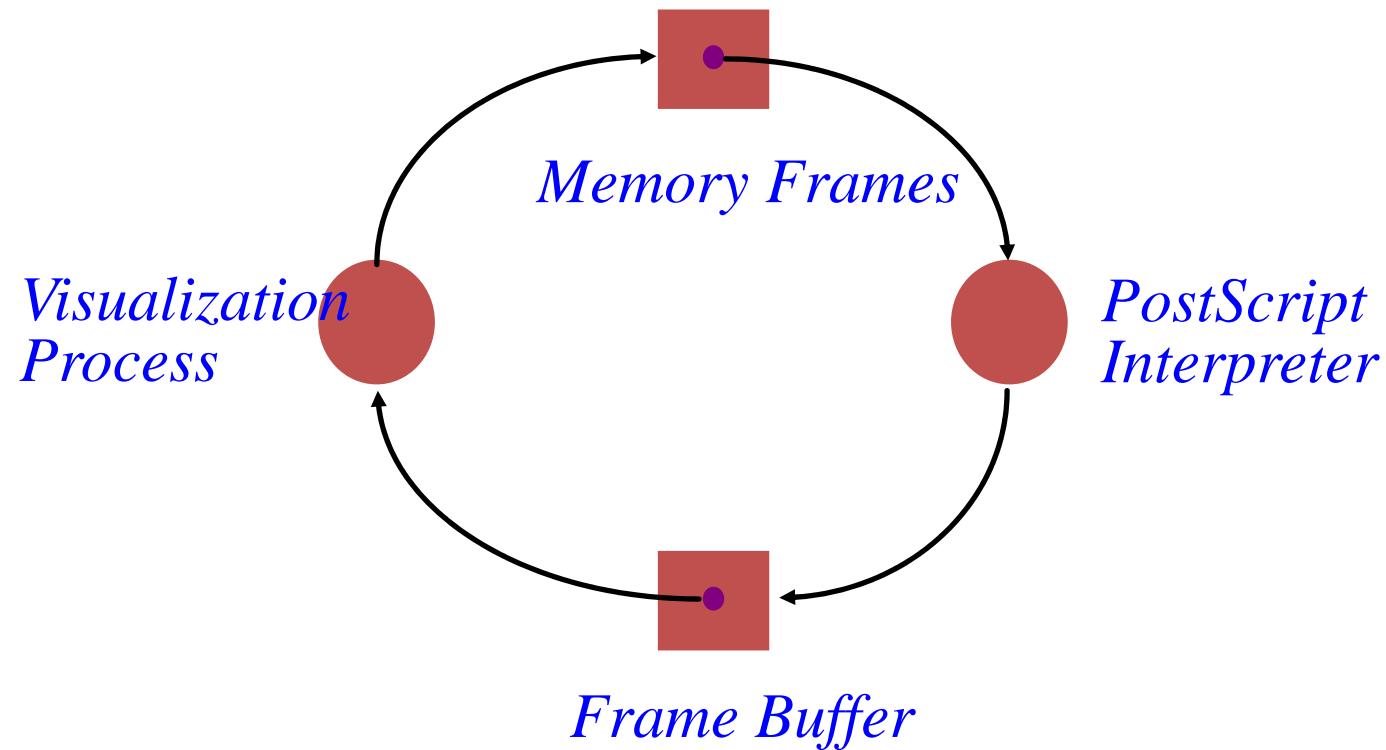
A cycle in a RAG is a necessary condition for deadlock

Is the existence of a cycle a sufficient condition?



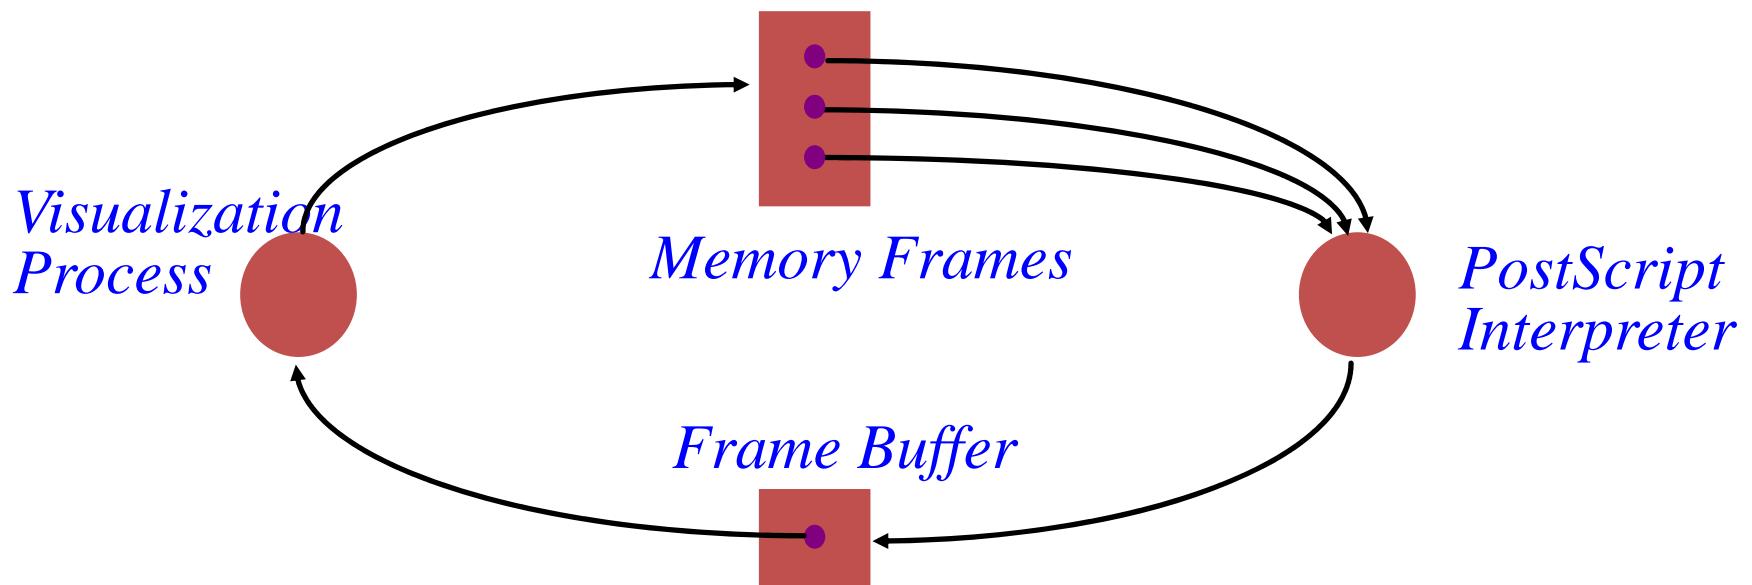
Single Resource RAG & Deadlocks

Theorem: If there is only a single unit of all resources then a set of processes are deadlocked iff there is a cycle in the resource allocation graph



An operational definition of deadlock

- A set of processes are deadlocked **iff** the following conditions hold simultaneously
 1. Mutual exclusion is required for resource usage
 2. A process is in a “hold-and-wait” state
 3. Preemption of resource usage is not allowed
 4. Circular waiting exists (a cycle exists in the RAG)



Dealing With Deadlock

Adopt some resource allocation protocol that ensures deadlock can never occur

Deadlock prevention/avoidance

Guarantee that deadlock will never occur

Generally breaks one of the following conditions:

 Mutex, Hold-and-wait, No preemption, Circular wait

Deadlock detection and recovery

Admit the possibility of deadlock occurring and periodically check for it

On detecting deadlock, abort

 Breaks the no-preemption condition

Deadlock Avoidance by Resource Ordering

Eliminate circular waiting by ordering all locks (or semaphores, or resources)

All code grabs locks in a predefined order

Problems?

Maintaining global order is difficult, especially in a large project

Global order can force a client to grab a lock earlier than it would like, tying up a resource for too long

Deadlock Detection & Recovery

Detection: periodic check RAG for cycles

How often should the OS check for deadlock?

- After every resource request?

- Only when we suspect deadlock has occurred?

Recovery: break the deadlock

Abort all deadlocked processes & reclaim their resources?

Abort one process at a time until all cycles in the RAG are eliminated?

- Start from low priority process?

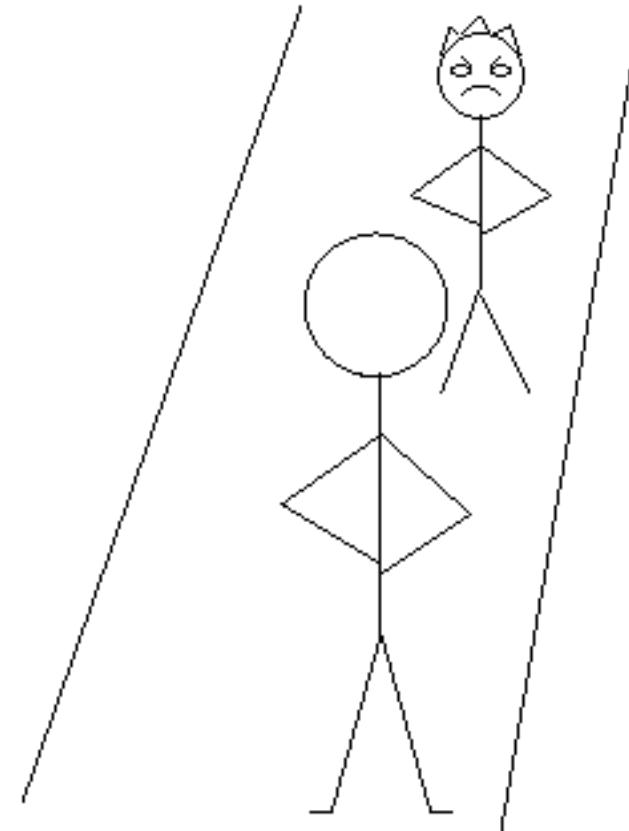
- Start from processes with most allocation of resources?

LiveLock

Similar to a deadlock

The states of the processes constantly change with regard to one another

none progressing



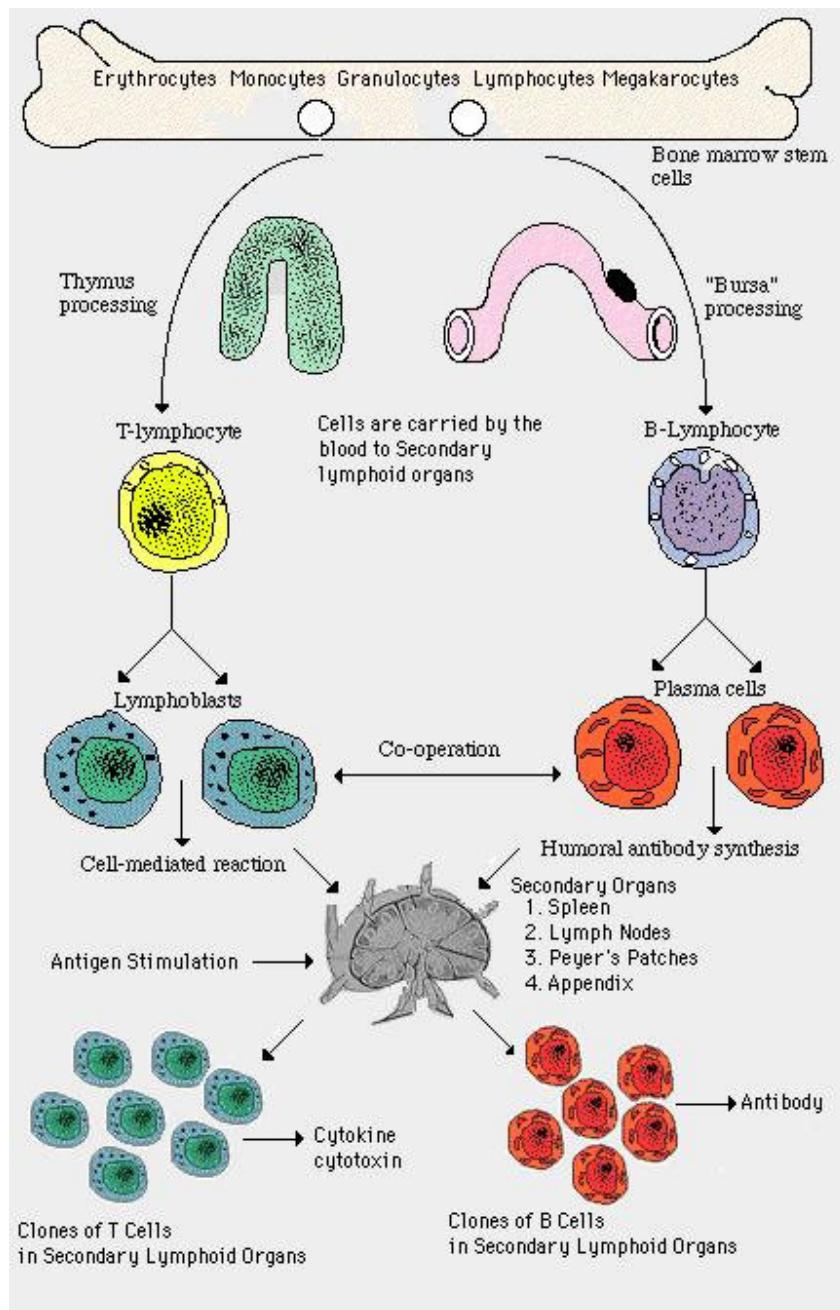
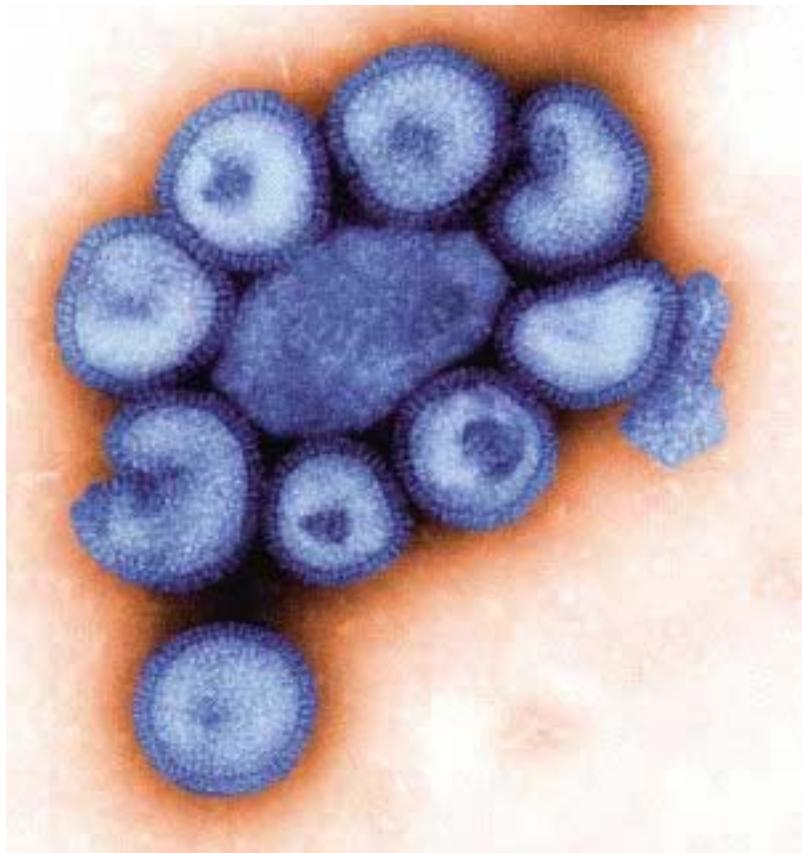
Deadlock Pitfalls in Reality

Even **deadlock-free** code would deadlock once deployed

Due to dependencies on deadlock-prone third party libraries or runtimes

Examples: web browsers plug-ins, Java beans.

Upgrade of such libraries/runtimes can introduce new deadlocks during execution.

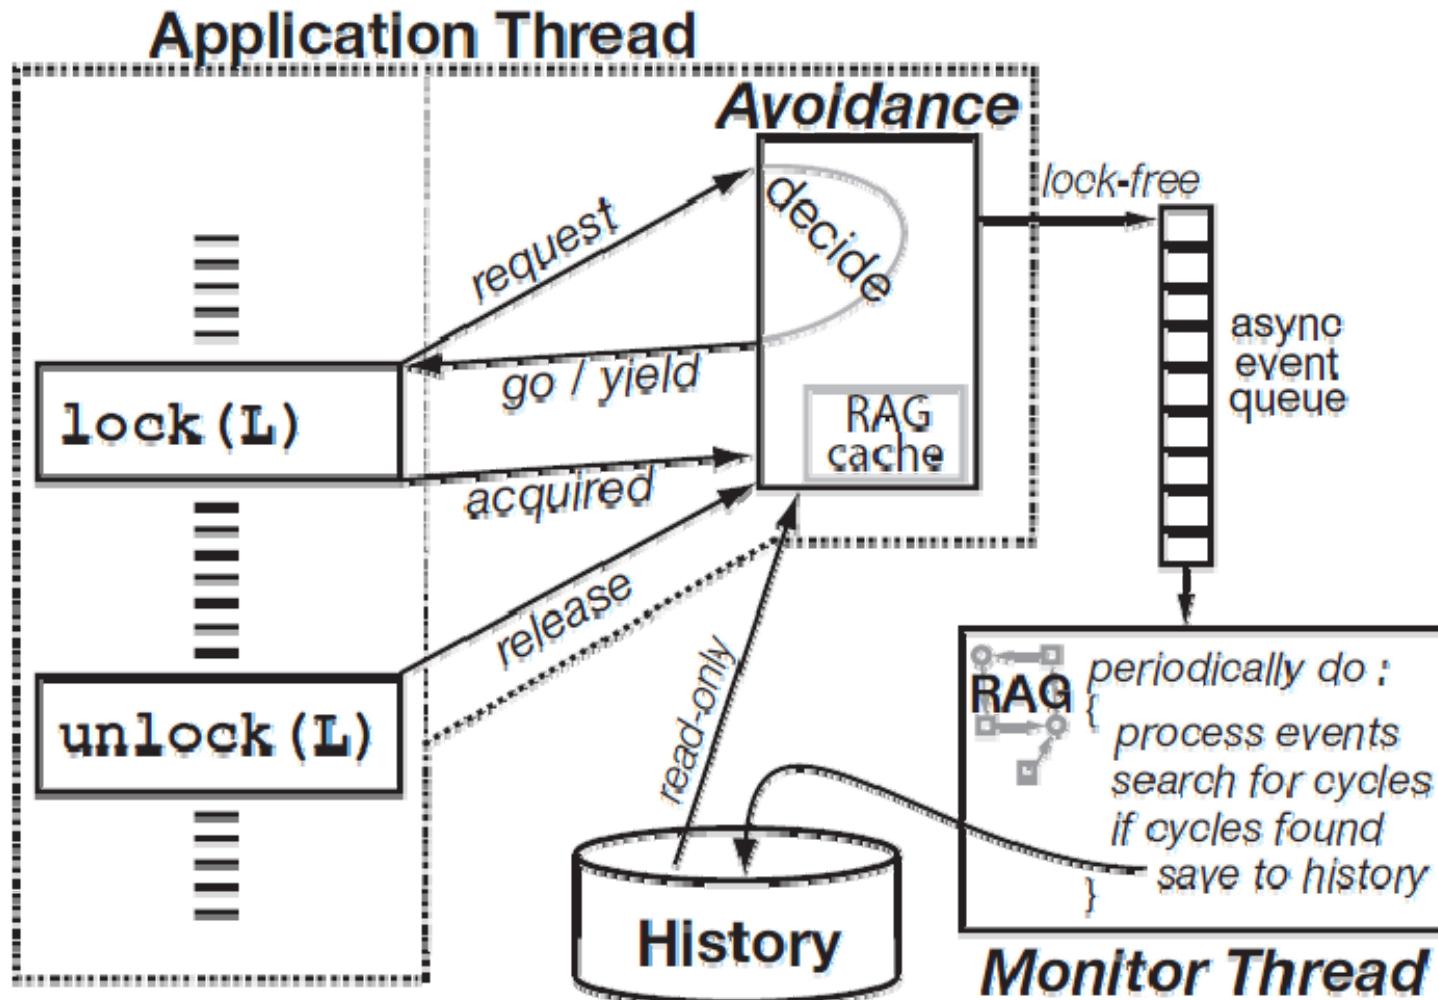


Dimmunix - Teaser

So, it is often **hard** to handle a deadlock. But what if...

- *Deadlock immunity:*
 - once afflicted by a given deadlock, develop resistance against future concurrences of that and similar deadlocks
- *Dimmunix* - a tool for developing deadlock immunity.
- The first time a deadlock pattern manifests
 - Dimmunix automatically captures its *signature* and subsequently avoids entering the same pattern.

Dimmunix Architecture

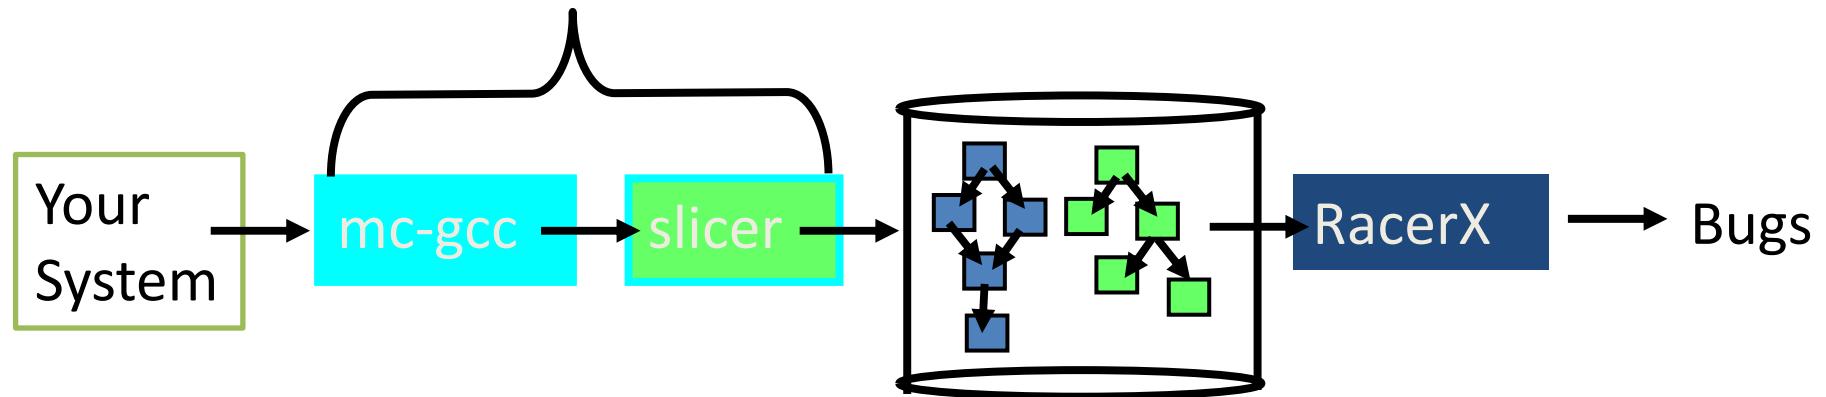


RaceX: Static LockSet Analysis for DeadLock

Slides from Dawson Engler@Stanford

The RacerX experience

- How to use:
 - List locking functions & entry points. Small:
 - Linux: 18 + 31, FreeBSD: 30 + 36, System X: 50 + 52
 - Emit trees from source code (2x cost of compile)



- Run RacerX over emitted trees
 - Links all trees into global control flow graph (CFG)
 - Checks for deadlocks & races
 - ~2-20 minutes for Linux.
- Post-process to rank errors (most of IQ spent here)
- Inspect

Lockset analysis

- Lockset: set of locks currently held [Eraser]
 - For each root, do a flow-sensitive, inter-procedural DFS traversal computing lockset at each statement
 - initial \rightarrow lockset = {}
 - lock(l) \rightarrow lockset = lockset \cup { l }
 - unlock(l) \rightarrow lockset = lockset - { l }
 - Speed: If stmt **s** was visited before with lockset **ls**, stop.
- Inter-procedural:
 - Routine can exit with multiple locksets: resume DFS w/ each after callsite.
 - Record <in-ls, {out-ls}> in fn summary. If ls in summary, grab cached out-ls' s and skip fn body.

Lockset

```
connect() {  
    → lock(a);      { a }  
    → open_conn();  { a }  
    send();  
}
```

{ a }

summary:
{ a } → ?

```
open_conn() {  
    → if (x)          { a }  
    → lock(b);        { a, b }  
    → else             { a }  
    → lock(c);        { a, c }  
}
```

→

{ a, b } { a, c }

Lockset

```
connect() {  
    lock(a);      { a }  
    open_conn();  { a }  
    send();       { a, b }, {a, c }  
}
```

summary:
 $\{ a \} \rightarrow \{ a, b \}, \{a, c \}$

```
open_conn() {  
    if (x)          { a }  
    lock(b);       { a, b }  
    else           { a }  
    lock(c);       { a, c }  
}
```

{ a, b }, {a, c }

Big picture: Deadlock detection

- Pass 1: constraint extraction
 - emit 1-level locking dependencies during lockset analysis



- Pass 2: constraint solving
 - Compute transitive closure & flag cycles.
 - “**a → b → a**” : T1 acquires a, T2 acquires b, boom.
- Ranking:
 - Global locks over local
 - Depth of callchain & number of conditionals (less better)
 - Number of threads involved (fewer MUCH better)

Simplest deadlock example

```
// 2.5.62/drivers/char/rtc.c
int rtc_register(rtc_task_t *task) {
    spin_lock_irq(&rtc_lock);
    //...
    spin_lock(&rtc_task_lock);
    if (rtc_callback) {
        spin_unlock(&rtc_task_lock);
        spin_unlock_irq(&rtc_lock);
```

```
//2.5.62/drivers/char/rtc.c
rtc_unregister(rtc_task_t *task)
{
    spin_lock_irq(&rtc_task_lock);
    //...
    spin_lock(&rtc_lock);
```

- Constraint extraction emits “`rtc_lock`→`rtc_task_lock`” and “`rtc_task_lock`→`rtc_lock`”
- Constraint solving flags cycle: T1 acquires `rtc_lock`, T2 acquires `rtc_task_lock`. Boom.
- Ranked high: only two threads, global locks, local error.

False positive trouble

- Most FPs from bogus locks in lockset
 - Typically caused by mishandled data dependencies
- Oversimplified typical example
 - Naïve analysis will think four paths rather than two, including false one that holds lock **a** at line 5.

```
1: if(x)          {}
2:   lock(a);    {a}
3: if(x)          {a}
4:   unlock(a);
5: lock(b);      {a}  "a→b"
```

- Inter-procedural analysis makes this much worse.
- Could add path-sensitivity, but undecidable in general

Unlockset analysis

- Observations:
 - In practice, all false positives due to the A in “A→B”, most because A goes “too far”
 - We had unconsciously adopted pattern of inspecting errors where there was an explicit unlock of “A” after “A→B” since that strongly suggested “A” was held.

```
// 2.5.62/drivers/char/rtc.c
rtc_register(rtc_task_t *task) {
    spin_lock_irq(&rtc_lock);
    //...
    spin_lock(&rtc_task_lock); → rtc_lock → rtc_task_lock
    if (rtc_callback) {
        spin_unlock(&rtc_task_lock);
        spin_unlock_irq(&rtc_lock);
```

Unlockset analysis

- At statement S remove any lock L from lockset if there exists no successor statement S' reachable from S that contains an unlock of L.

1: if(x)	{}
2: lock(a);	{a}
3: if(x)	{a}
4: unlock(a);	
5: lock(b);	{a} → {}

- Key: lockset holds exactly those locks the analysis can handle. Scales with analysis sophistication.

Unlockset implementation sketch

- Essentially compute reaching definitions
 - Run lockset analysis in reverse from leaves to roots
 - Unlockset holds all locks that will be released
 - initial \rightarrow unlockset = {}
 - lock(l) \rightarrow unlockset = unlockset - {l}
 - unlock(l) \rightarrow unlockset = unlockset \cup {l}
 - s.unlockset = s.unlockset \cup unlockset
 - During lockset analysis:
 $lockset = \text{intersect}(s.unlockset, lockset);$
- Main complication: function calls.
 - Different locks released after different callsites. Don't want to mix these up (context sensitivity)

Deadlock results

System	Confirmed	Unconfirmed	False
System X	2	3	7
Linux 2.5.62	4	8	6
FreeBSD	2	3	6
Total	8	14	19

- A bit surprised at the low bug counts
 - Main reason seems to be not that many locks held simultaneously
 - < 1000 unique constraints, only so many chances for error.

The most surprising error

```
// Entered holding scsiLock
int FindHandle(int handleID) {
    prevIRQL = SP_LockIRQ(&handleArrayLock, ...);
    Validate(handle);

    ...
int Validate(handle) {
    ASSERT(SP_IsLocked(&scsiLock));
    while (adapter->openInProgress) {
        CpuSched_Wait(&adapter->openInProgress,
                      CPUSCHED_WAIT_SCSI, &scsiLock);
        SP_Lock(&scsiLock);
```

- T1 enters FindHandle with scsiLock, calls Validate, calls CpuSched_wait (rel scsiLock, sleep w/ handleArrayLock)
- T2 acquires scsiLock and calls FindHandle. Boom.

Lock Analysis in Real World

- Lockdep analysis
 - <https://www.kernel.org/doc/Documentation/lockdep-design.txt>

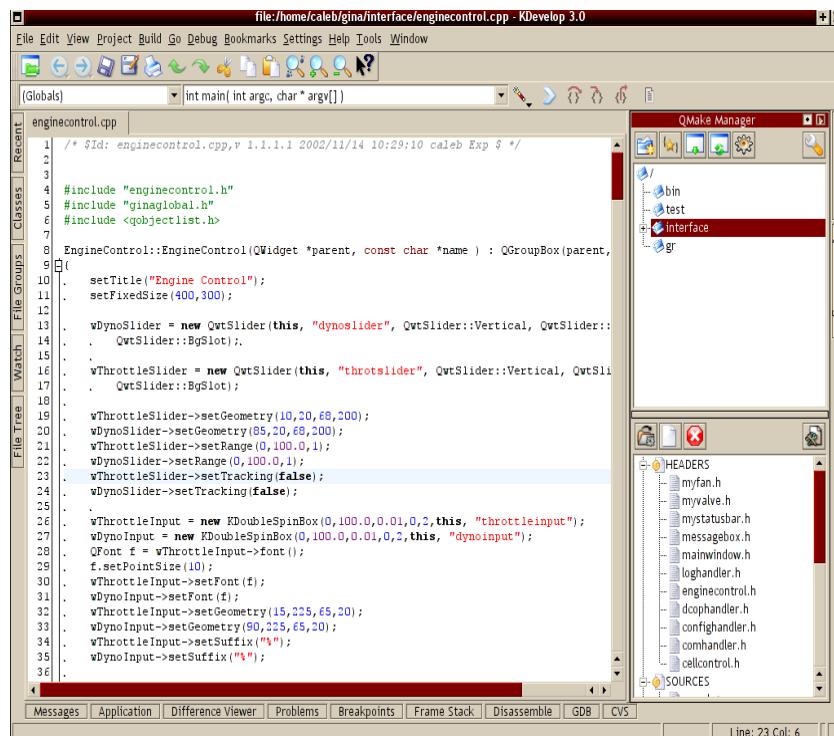
```
Chain exists of:  
  &of->mutex --> sr_mutex --> &mm->mmap_sem  
  
Possible unsafe locking scenario:  
  
          CPU0                         CPU1  
          ----                         ----  
          lock(&mm->mmap_sem);           lock(sr_mutex);  
                                      lock(&mm->mmap_sem);  
          lock(&of->mutex);  
*** DEADLOCK ***
```

```
===== [ INFO: possible circular locking dependency detected ]  
3.12.0+ #2 Not tainted  
-----  
trinity-child0/9004 is trying to acquire lock:  
  (&of->mutex){+.+.+}, at: [<ffffffffff8123c0cf>] sysfs_bin_mmap+0x4f/0x120  
  ready holding lock:  
    (&mm->mmap_sem){++++++}, at: [<ffffffffff8116b5ff>] vm_mmap_pgoff+0x6f/0xc0  
  which lock already depends on the new lock.
```

**Valgrind: a useful tool for
program debugging/analysis**

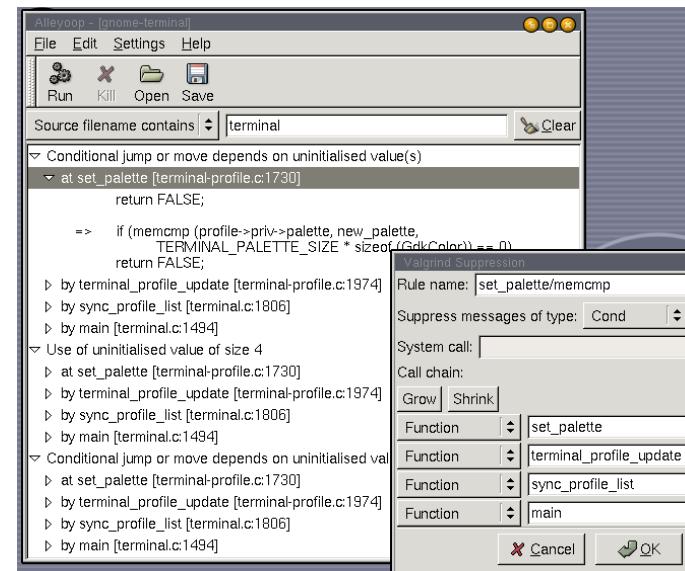
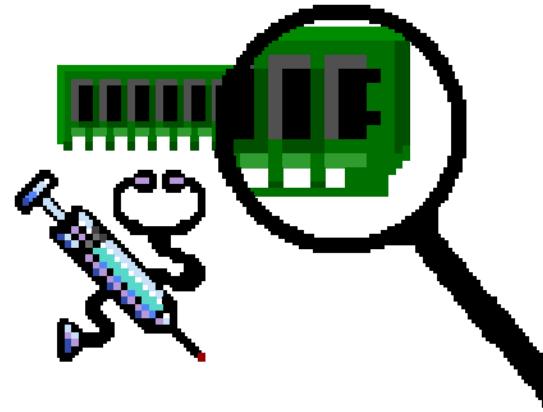
What is Valgrind

- Linux, x86, program analyzer
- Advantages
 - powerful
 - no recompiles, just binaries
 - any language (!)
 - extensible, configurable
 - free, GPL'd goodness
- Speed slowdown

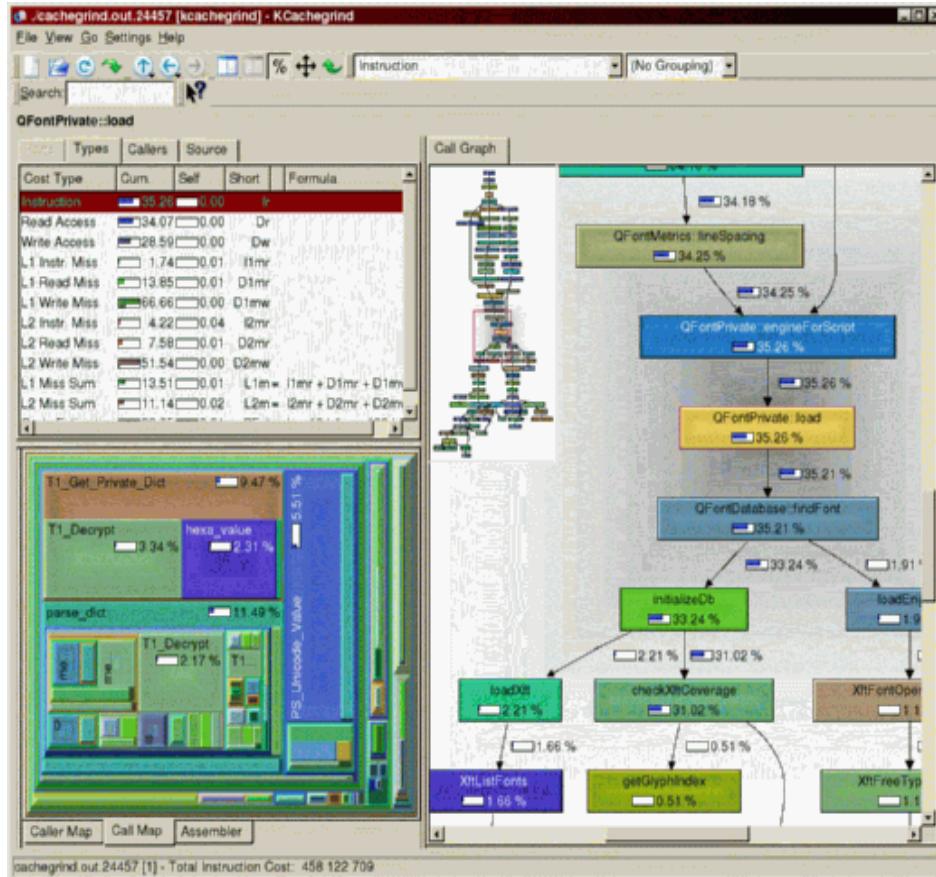


What are the existing skins

- MemCheck
 - Heavy duty memory checker
- AddrCheck
 - Faster, cheaper
- Cachegrind
 - Caching info
- Helgrind
 - Thread semantics
- Others



Other features



- Callgraphs
- Multiple other patches
 - Open file descriptors
- Suppressions
 - Selective understanding
- Ignore System libs
 - Yes, bugs
- /dev/random

How to Use it

- Apt-get install valgrind

You invoke Valgrind like this:

```
valgrind [valgrind-options] your-prog [your-prog-options]
```

The most important option is `--tool` which dictates which Valgrind tool to run.

```
valgrind --tool=memcheck ls -l
```

Example code

```
int main ( void )
{
    int r;
    pthread_mutex_t mx1, mx2;
    r = pthread_mutex_init( &mx1, NULL ); assert(r==0);
    r = pthread_mutex_init( &mx2, NULL ); assert(r==0);

    r = pthread_mutex_lock( &mx1 ); assert(r==0);
    r = pthread_mutex_lock( &mx2 ); assert(r==0);

    r = pthread_mutex_unlock( &mx1 ); assert(r==0);
    r = pthread_mutex_unlock( &mx2 ); assert(r==0);

    r = pthread_mutex_lock( &mx2 ); assert(r==0); /* error */
    r = pthread_mutex_lock( &mx1 ); assert(r==0);

    r = pthread_mutex_unlock( &mx1 ); assert(r==0);
    r = pthread_mutex_unlock( &mx2 ); assert(r==0);

    r = pthread_mutex_destroy( &mx1 );
    r = pthread_mutex_destroy( &mx2 );

    return 0;
}
```

How Valgrind Detect it

```
Thread #1: lock order "0x7FEFFFAB0 before 0x7FEFFFA80" violated
  at 0x4C23C91: pthread_mutex_lock (hg_intercepts.c:388)
  by 0x40081F: main (tc13_laogl.c:24)
```

Required order was established by acquisition of lock at 0x7FEFFFAB0

```
  at 0x4C23C91: pthread_mutex_lock (hg_intercepts.c:388)
  by 0x400748: main (tc13_laogl.c:17)
```

followed by a later acquisition of lock at 0x7FEFFFA80

```
  at 0x4C23C91: pthread_mutex_lock (hg_intercepts.c:388)
  by 0x400773: main (tc13_laogl.c:18)
```

**MANY OTHER USEFUL
FEATURES**

Example Uses of MemCheck

```
#include <stdlib.h>

void f(void)
{
    int* x = malloc(10 * sizeof(int));
    x[10] = 0;
}

int main(void)
{
    f();
    return 0;
}
```

```
==19182== Invalid write of size 4
==19182==     at 0x804838F: f (example.c:6)
==19182==     by 0x80483AB: main (example.c:11)
==19182== Address 0x1BA45050 is 0 bytes after a block of size 40 alloc'd
==19182==     at 0x1B8FF5CD: malloc (vg_replace_malloc.c:130)
==19182==     by 0x8048385: f (example.c:5)
==19182==     by 0x80483AB: main (example.c:11)
```

```
==19182== 40 bytes in 1 blocks are definitely lost in loss record 1 of 1
==19182==     at 0x1B8FF5CD: malloc (vg_replace_malloc.c:130)
==19182==     by 0x8048385: f (a.c:5)
==19182==     by 0x80483AB: main (a.c:11)
```

Cachegrind: a cache profiler

Cachegrind

- Cache behavior is crucial
 - L1 misses: ~10 cycles
 - L2 misses: ~200 cycles
- But difficult to predict
- Cachegrind gives three outputs:
 - Total hit/miss counts and ratios (I1, D1, L2)
 - Per-function hit/miss counts (sorted from most to least)
 - Per-line hit/miss counts (source code annotations)
- Source code annotations are the most useful
 - Most fine-grained data
 - Data that programmers can act on to speed up their programs

Sample output

```
-----  

Ir I1mr I2mr      Dr D1mr D2mr      Dw      D1mw      D2mw  

-----  

14,789,396 547 544 6,329,792 751 689 2,111,757 1,113,292 1,094,855 PROGRAM TOTALS  

-----  

Ir I1mr I2mr      Dr D1mr D2mr      Dw      D1mw      D2mw file:function  

-----  

14,688,273 1 1 6,294,531 0 0 2,098,178 1,113,088 1,094,656 example.c:main  

-----  

-- Auto-annotated source: example.c  

-----  

Ir I1mr I2mr      Dr D1mr D2mr      Dw      D1mw      D2mw  

-----  

. . . . . . . . int main(void)  

10 0 0 0 0 0 1 0 0 {  

. . . . . . . . int i, j, a[1024][1024];  

. . . . . . . .  

4,100 1 1 2,049 0 0 1 0 0 for (i = 0; i < 1024; i++) {  

4,198,400 0 0 2,098,176 0 0 1,024 0 0 for (j = 0; j < 1024; j++) {  

5,242,880 0 0 2,097,152 0 0 1,048,576 65,536 56,320 a[i][j] = 0; // fast  

5,242,880 0 0 2,097,152 0 0 1,048,576 1,047,552 1,038,336 a[j][i] = 0; // slow  

. . . . . . . . }  

. . . . . . . . }  

1 0 0 0 0 0 0 0 0 return 0;  

2 0 0 2 0 0 0 0 0 }
```

How Cachegrind works

- Each instruction is instrumented
 - Call to a C cache simulation function
 - Different functions for loads, stores, modifies
 - Some combining of C calls for efficiency
- Each source code line gets a *cost centre*
 - Holds counters: accesses, hits and misses
 - Uses debug info to map each instruction to a cost centre
- Online simulation (i.e. no trace gathering)
- Cost centres dumped to file at end
 - Simple but compact text format
 - Post-processing script produces previous slide's output

Cache simulation

- Approximates an AMD Athlon hierarchy
 - I1, D1, inclusive L2
 - Write-allocate
 - LRU replacement
- Each cache is command-line configurable:
 - Cache size
 - Line size
 - Associativity
- On x86/AMD64 can use CPUID to auto-detect these parameters
- Simulation can be replaced easily

Inaccuracies

- Imperfect address trace
 - No kernel code
 - Other processes ignored (arguably good)
 - Conversion to Valgrind's IR changes a very small number of loads/stores
- Incorrect addresses
 - Virtual addresses
 - Memory layout and thread scheduling is different under Cachegrind compared to native
- Prefetches and cache-bypassing are ignored
 - Difficult to handle well without detailed microarchitectural simulation
- Still useful for general insights

How is it used?

- Characterization:
 - Program A vs. program B
 - Cache hierarchy A vs. cache hierarchy B
- Optimisation:
 - Identifies cache-unfriendly code
 - Fixing such code requires non-trivial insight
 - But easier (i.e. not impossible!) than fixing without this data
- Evaluation of optimizations:
 - Program A vs. optimized program A

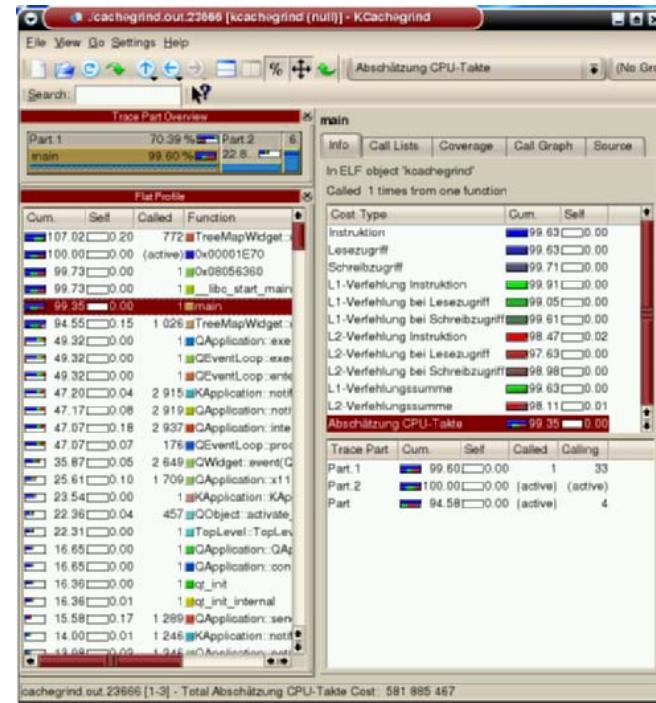
Cachegrind summary

- Cachegrind is a cache simulator
- Gives total, per-function and per-line hit/miss counts
- Simulation is imperfect, but still useful
- Used for characterization, optimisation and evaluation

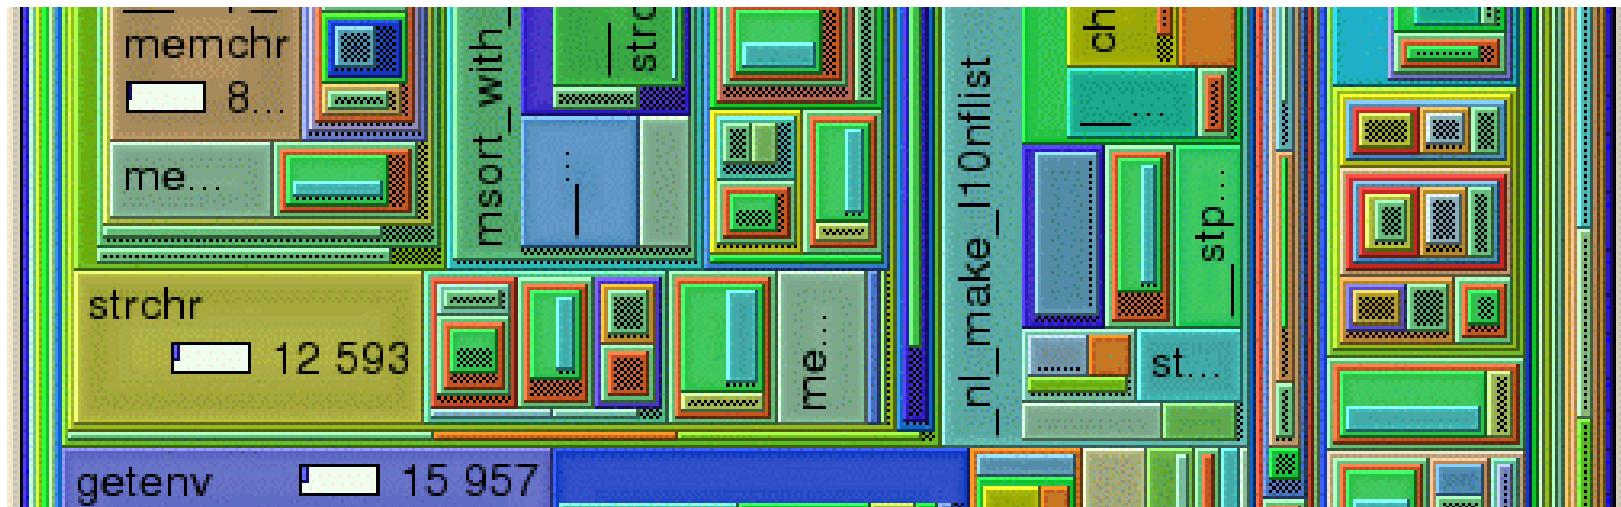
Callgrind: a call graph profiler

Callgrind

- Extension of Cachegrind
- By Josef Weidendorfer
- Also provides:
 - Call graph information
 - Graphical results viewer (KCachegrind)
 - Allows interactive browsing of results
 - Accepts Cachegrind results also
 - Greater selectivity of what code is profiled

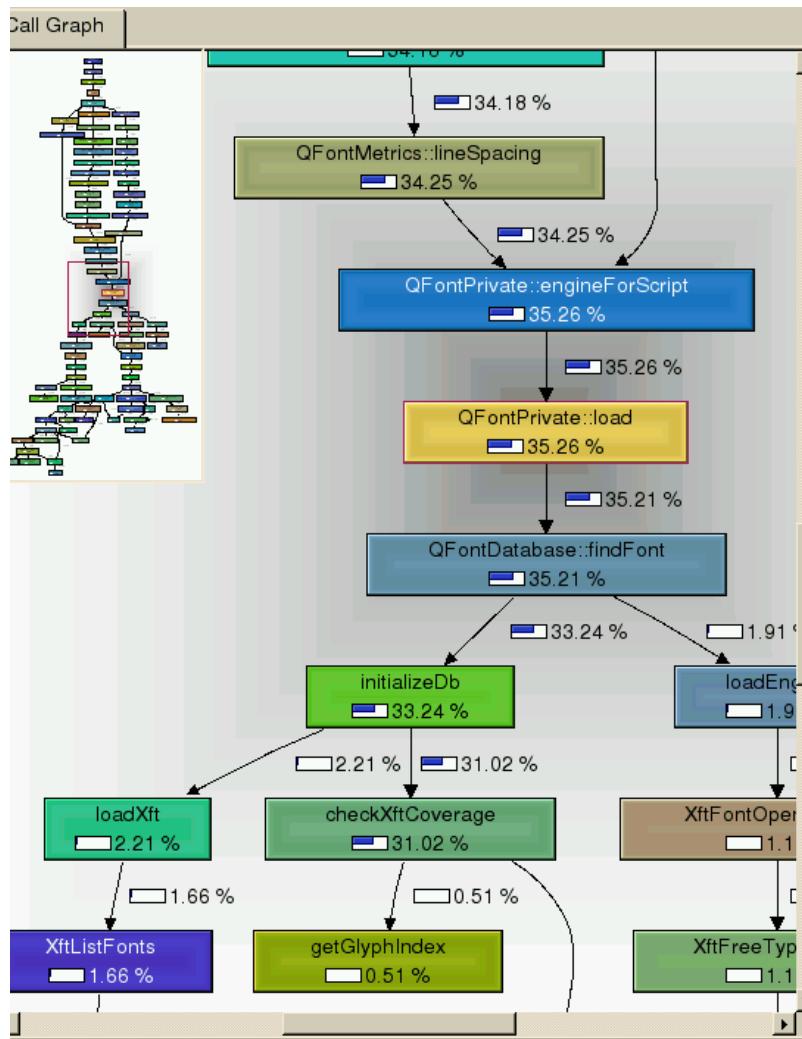


KCacheGrind's tree-map view



- Box sizes represent relative counts
- Nesting of boxes represents call chains
- Interactive: can drill down through boxes

KCacheGrind's call graph view



- Shows whole call graph
- Boxes show count proportions
- Interactive

Selective profiling

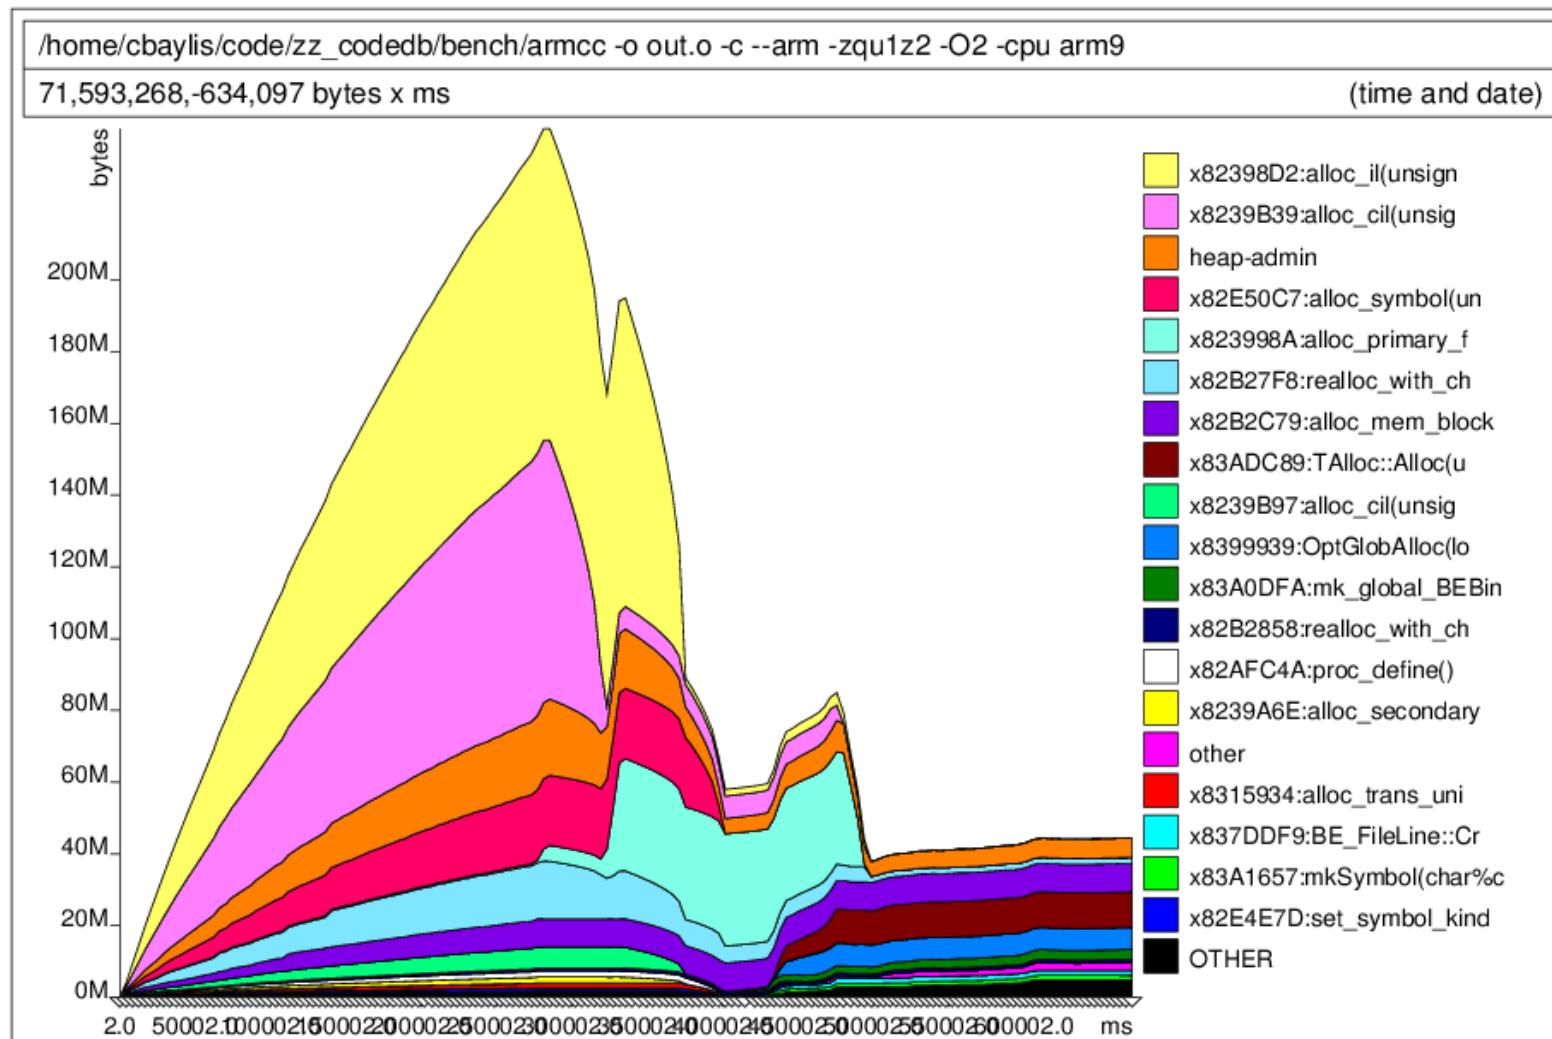
- Can dump counts at particular times
 - At termination (same as Cachegrind)
 - Periodically (every N code blocks)
 - At entry/exit of named functions
 - At particular program points (using client requests)
 - At any time (by invoking a separate script)
- Counters are zeroed after each dump
- Can choose which events to count
 - Instructions
 - Memory events (for cache simulation)
 - Function entries/exits

An interesting difficulty

- Callgrind maintains a call stack
 - For tracking function entries/exits
- Several difficulties:
 - `setjmp/longjmp`
 - Tail recursion
 - Dynamic linking
 - Calls through jump tables
 - Jump table patched on first call after loading
 - Stack switching
- Missed entries/exits can throw everything out

Massif: a heap profiler

Massif heap graph



Massif

- Measures heap and stack
 - Each heap allocation site is a band
 - Stack is a band
- Also produces HTML output
 - Represents the call graph underlying allocations
 - Users can drill down through calling chains from allocation sites
- Simple interaction with Valgrind's core
 - Only uses function wrapping
 - No instrumentation of code blocks
 - Complexity in the tool, not at the core/tool boundary

Thanks

Next class

Virtualization

System Virtualization

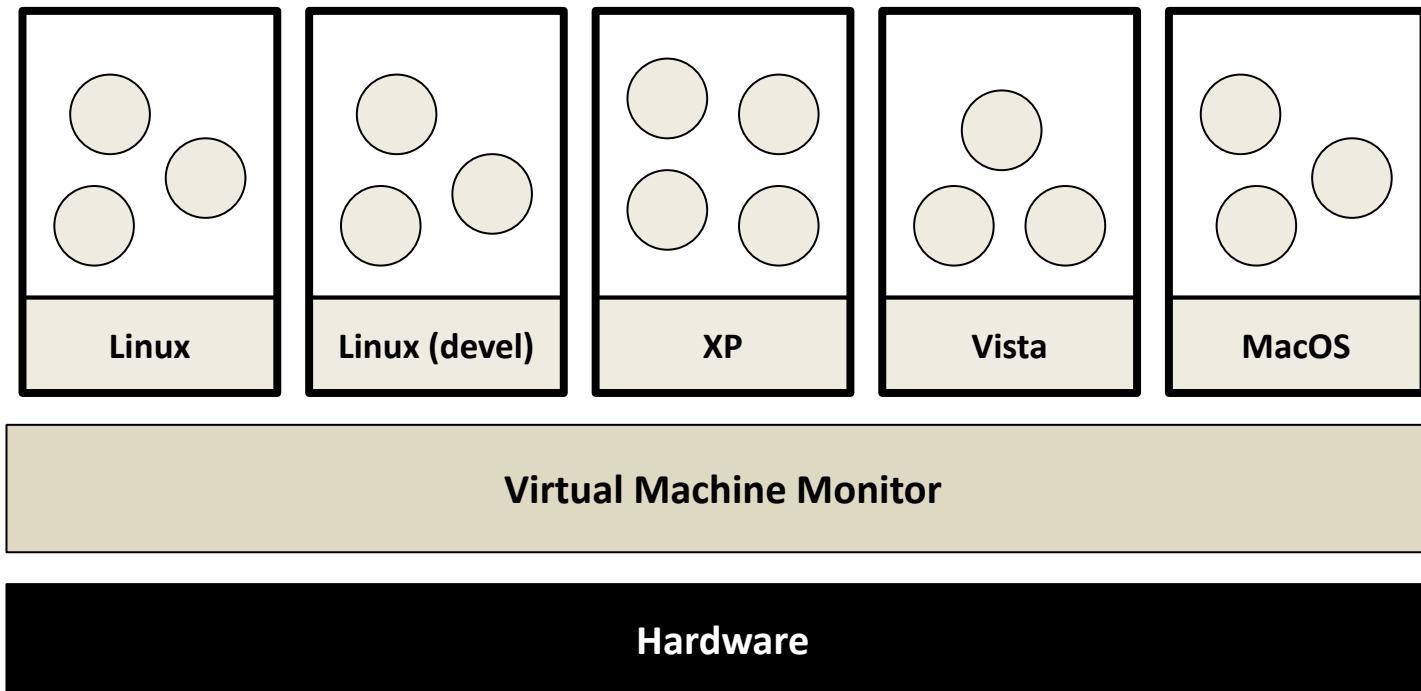
Yubin Xia
Software School
Shanghai Jiao Tong University

Some Slides adapted from
VMWare's academic course plan

Outline

- Virtualizing CPU
- Virtualizing Memory
- Virtualizing I/O

What is Virtualization



Virtualization Properties

- Isolation
- Encapsulation
- Interposition

Isolation

- Fault Isolation
 - Fundamental property of virtualization
- Software Isolation
 - Software versioning
 - DLL Hell
- Performance Isolation
 - Accomplished through scheduling and resource allocation

Encapsulation

- All VM state can be captured into a file
 - Operate on VM by operating on file
 - mv, cp, rm
- Complexity
 - Proportional to virtual HW model
 - Independent of guest software configuration

Interposition

- All guest actions go through monitor
- Monitor can inspect, modify, deny operations
- Ex
 - Compression
 - Encryption
 - Profiling
 - Translation

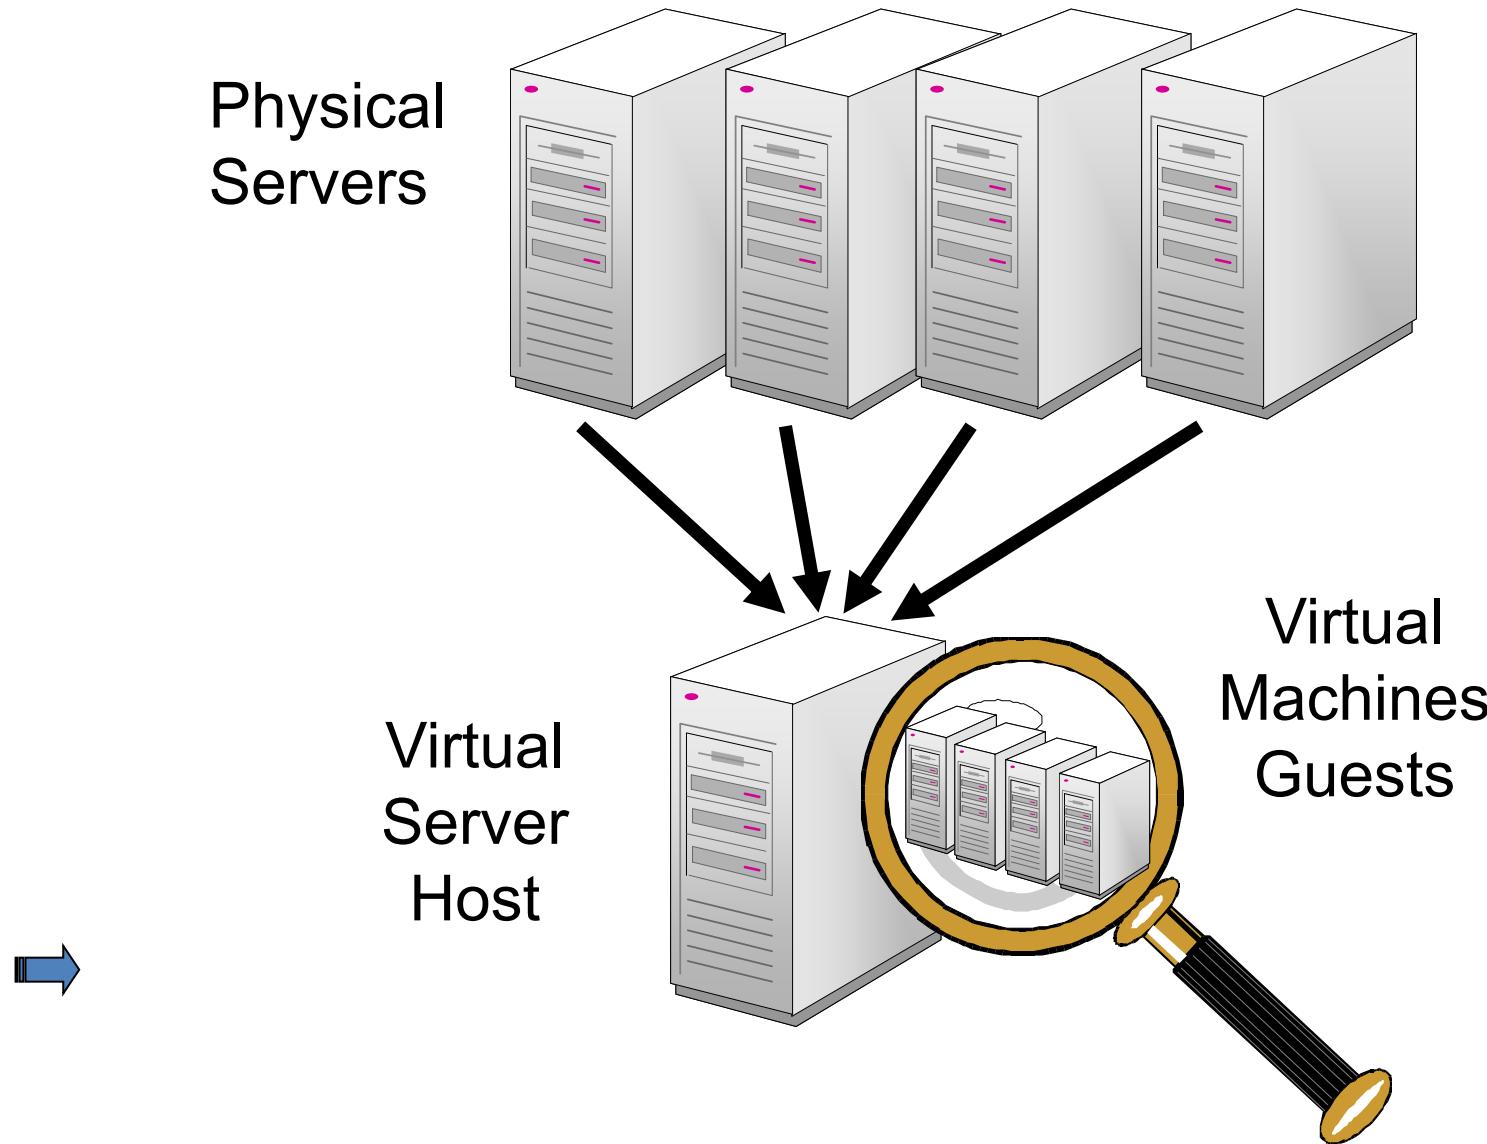
Why Not the OS?

- It about interfaces
 - VMMs operate at the hardware interface
 - Hardware interface are typically **smaller**, better defined than software interfaces
- Microkernel for commodity Operating Systems
- Disadvantages of being in the monitor
 - Low visibility into what the guest is doing

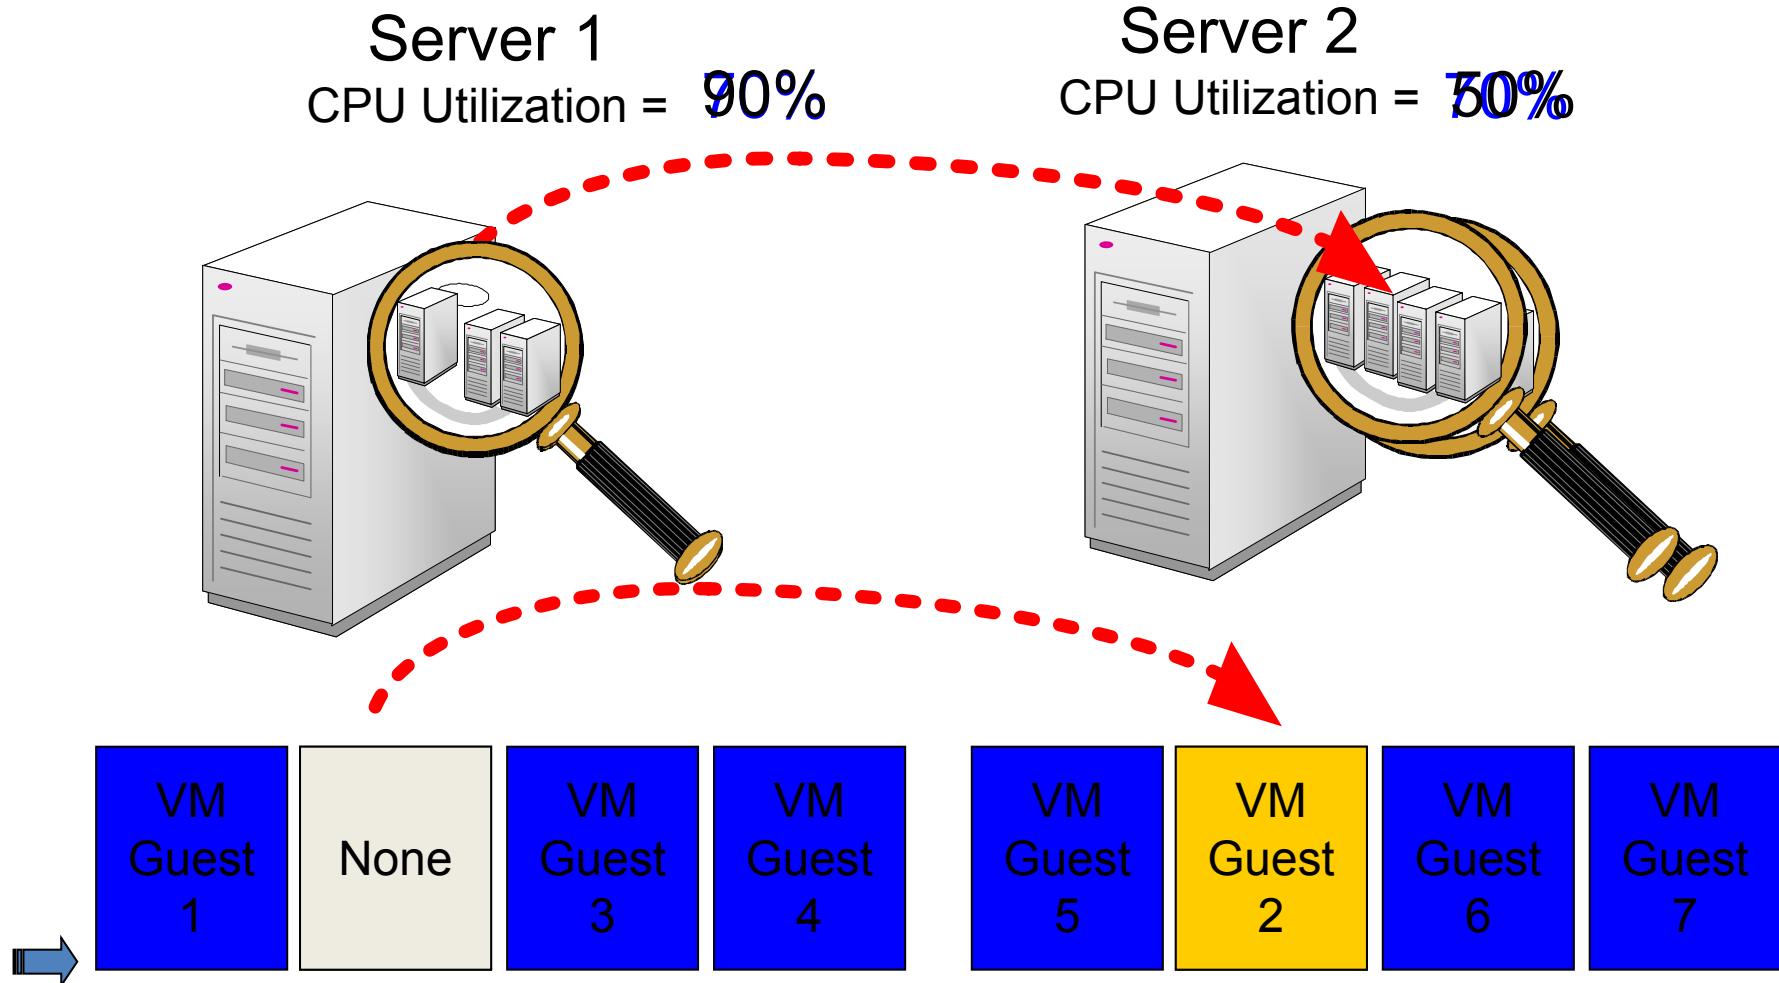
Virtualization benefits

- Increased resource utilization:
 - Server consolidation
- Mobility
- Enhanced Security
- Trusted Computing
- Test and Deployment
- ...

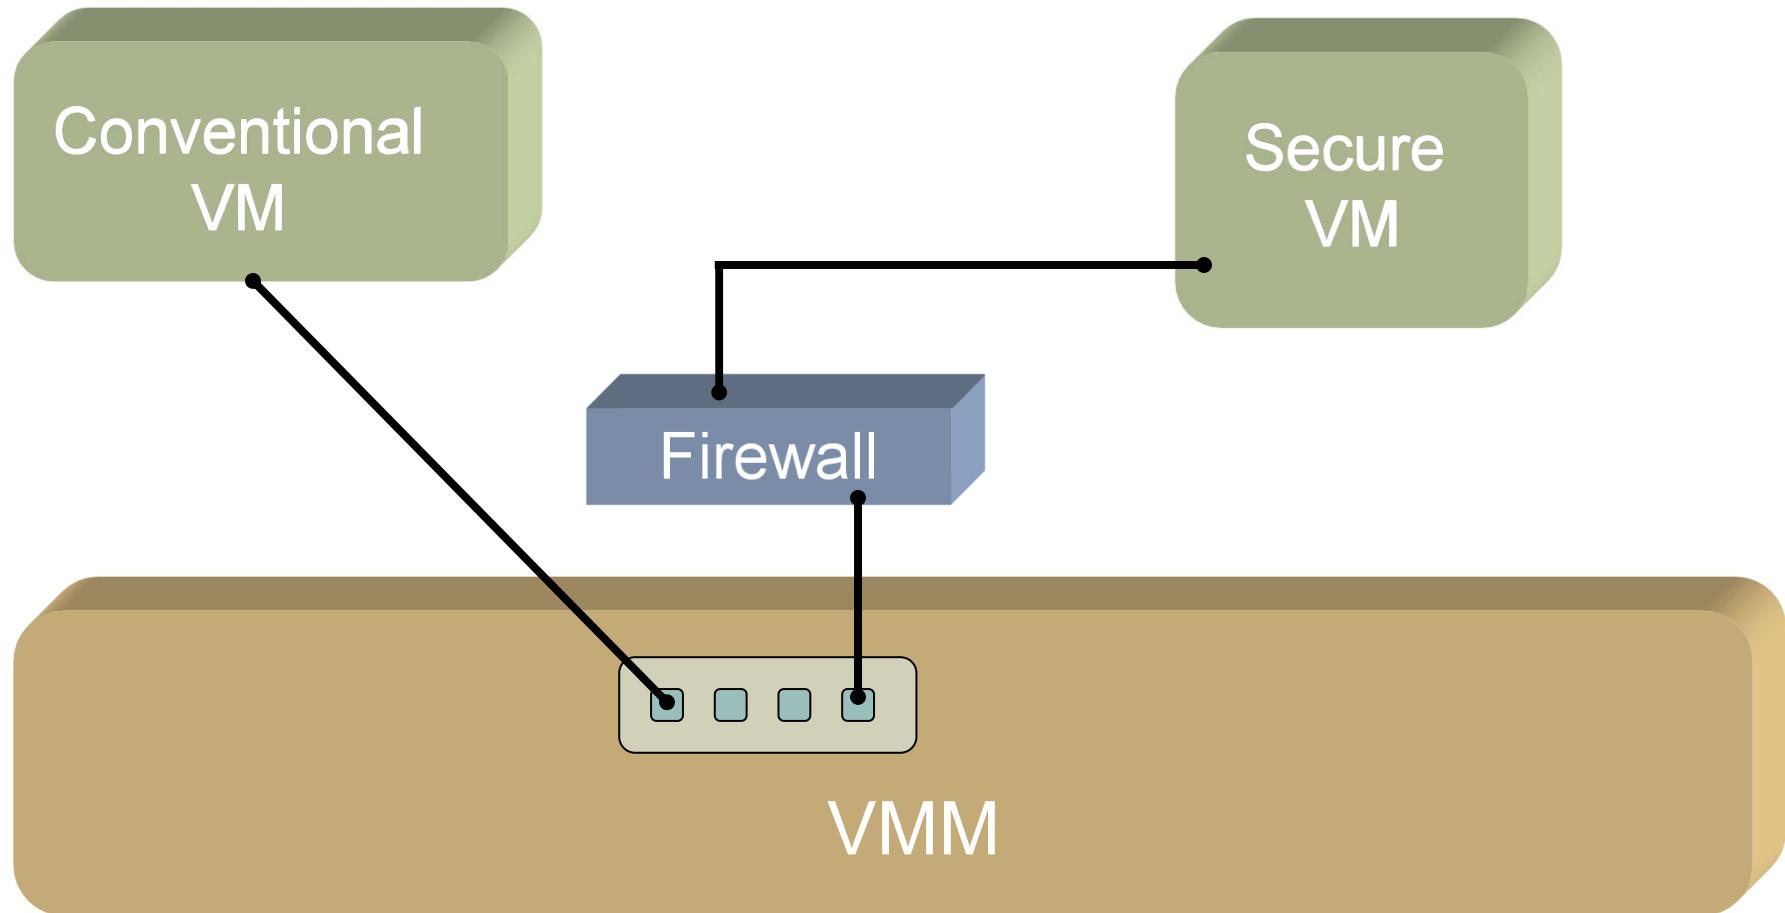
Virtualization: server consolidation



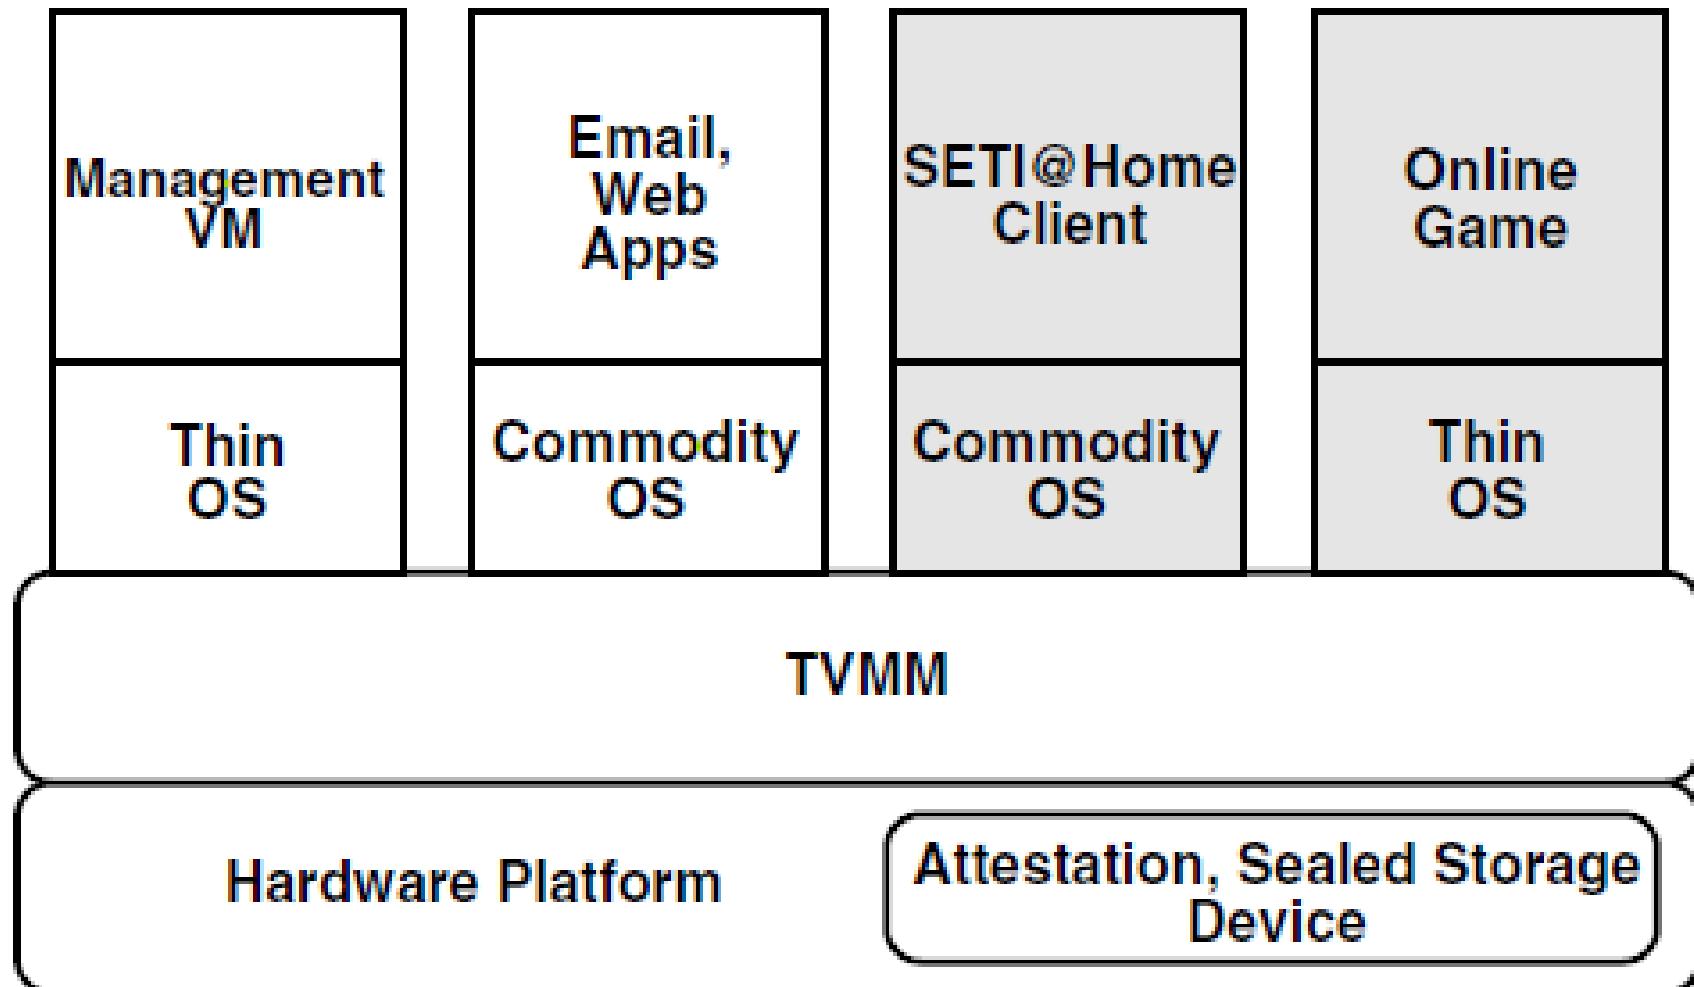
Mobility: load balance



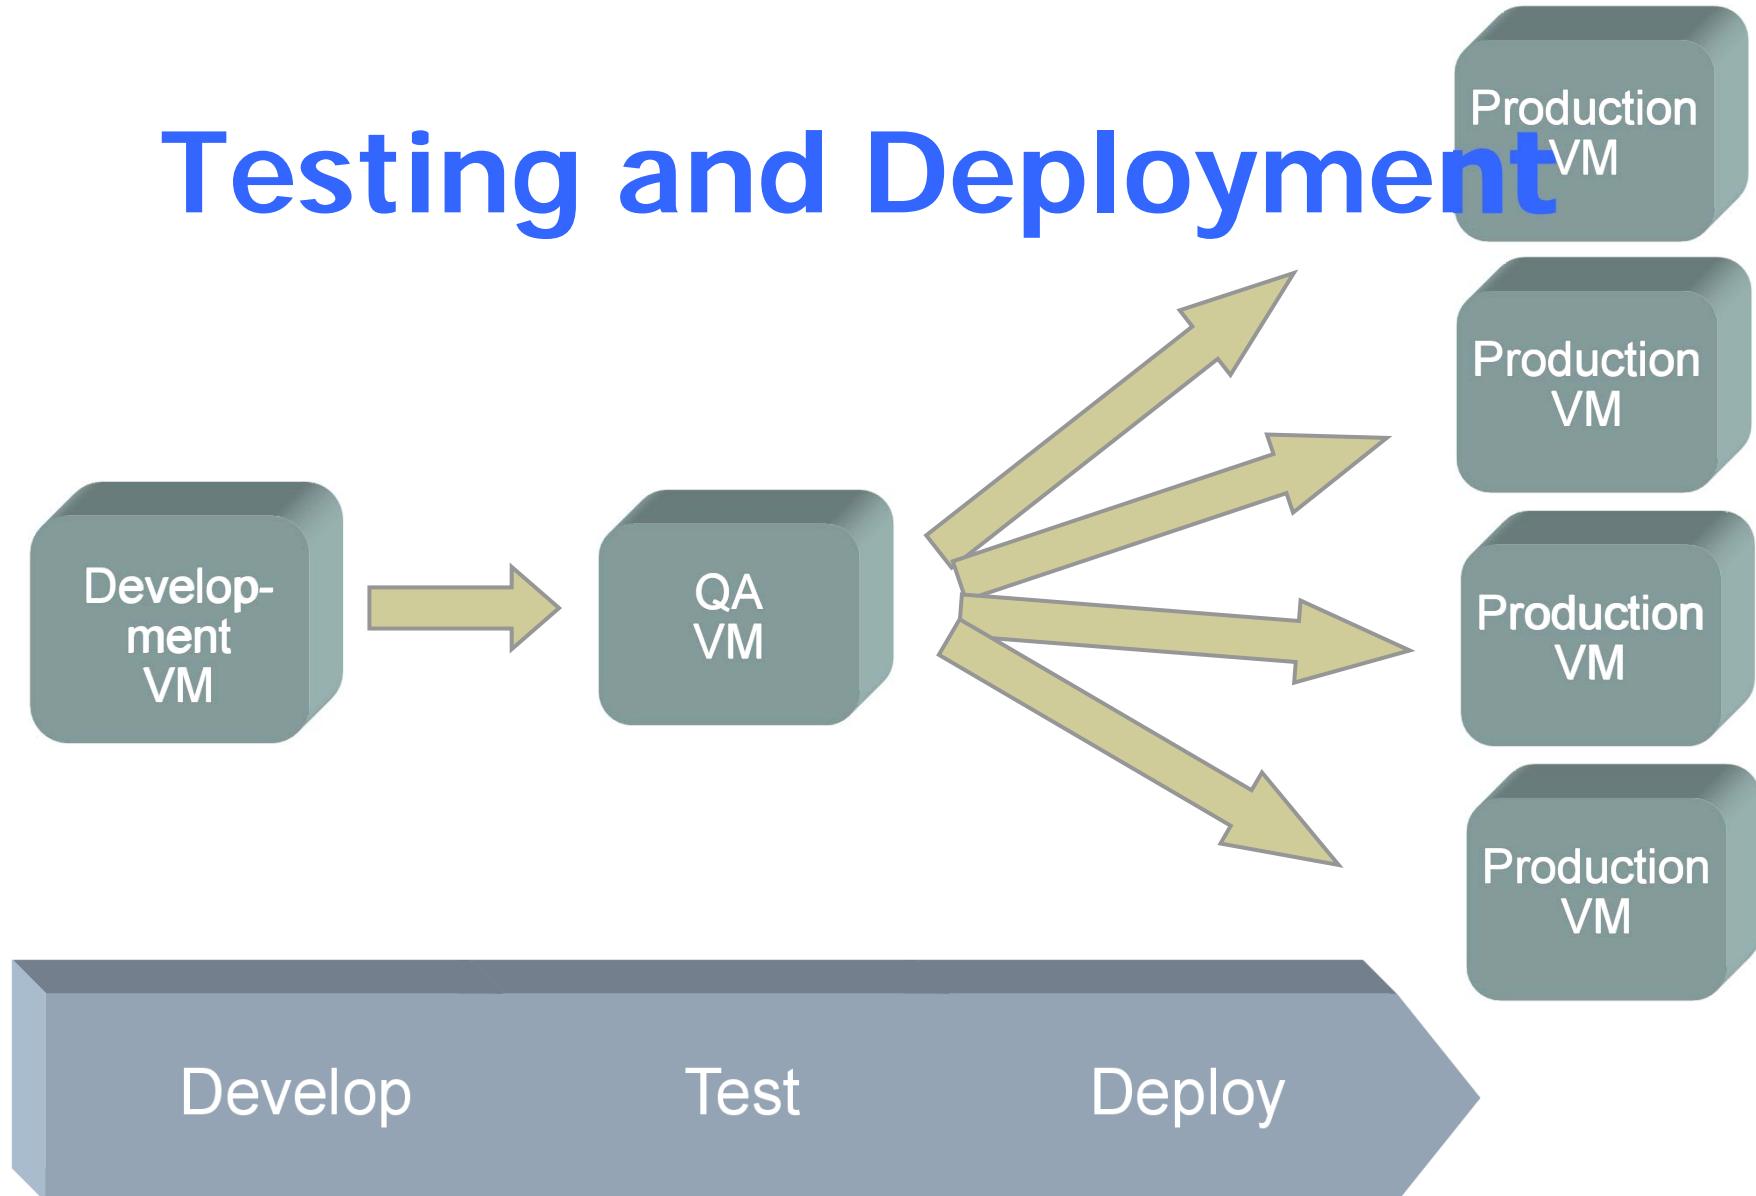
Enhanced Security



Trusted Computing



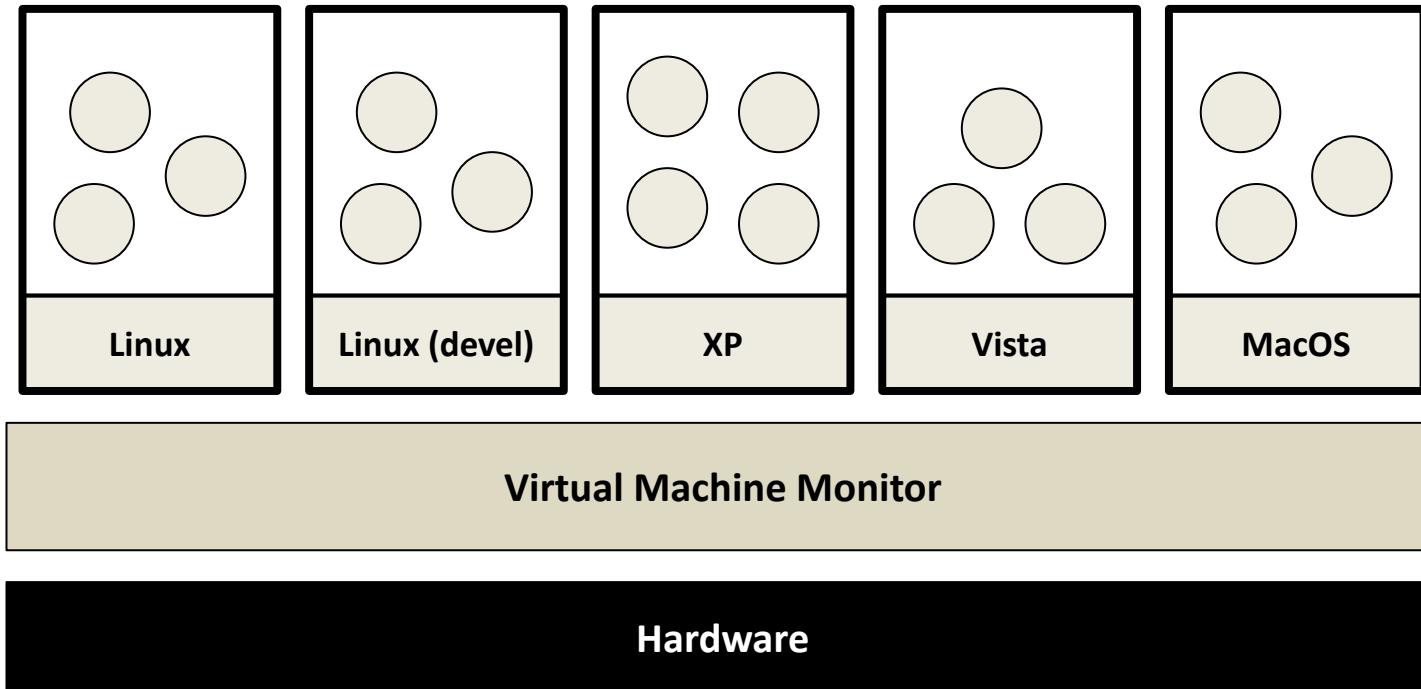
Testing and Deployment



System Virtual Machine Monitor Architectures

- Traditional
- Hosted
 - VMware Workstation
- Hybrid
 - VMware ESX
 - Xen
- Hypervisor

Traditional

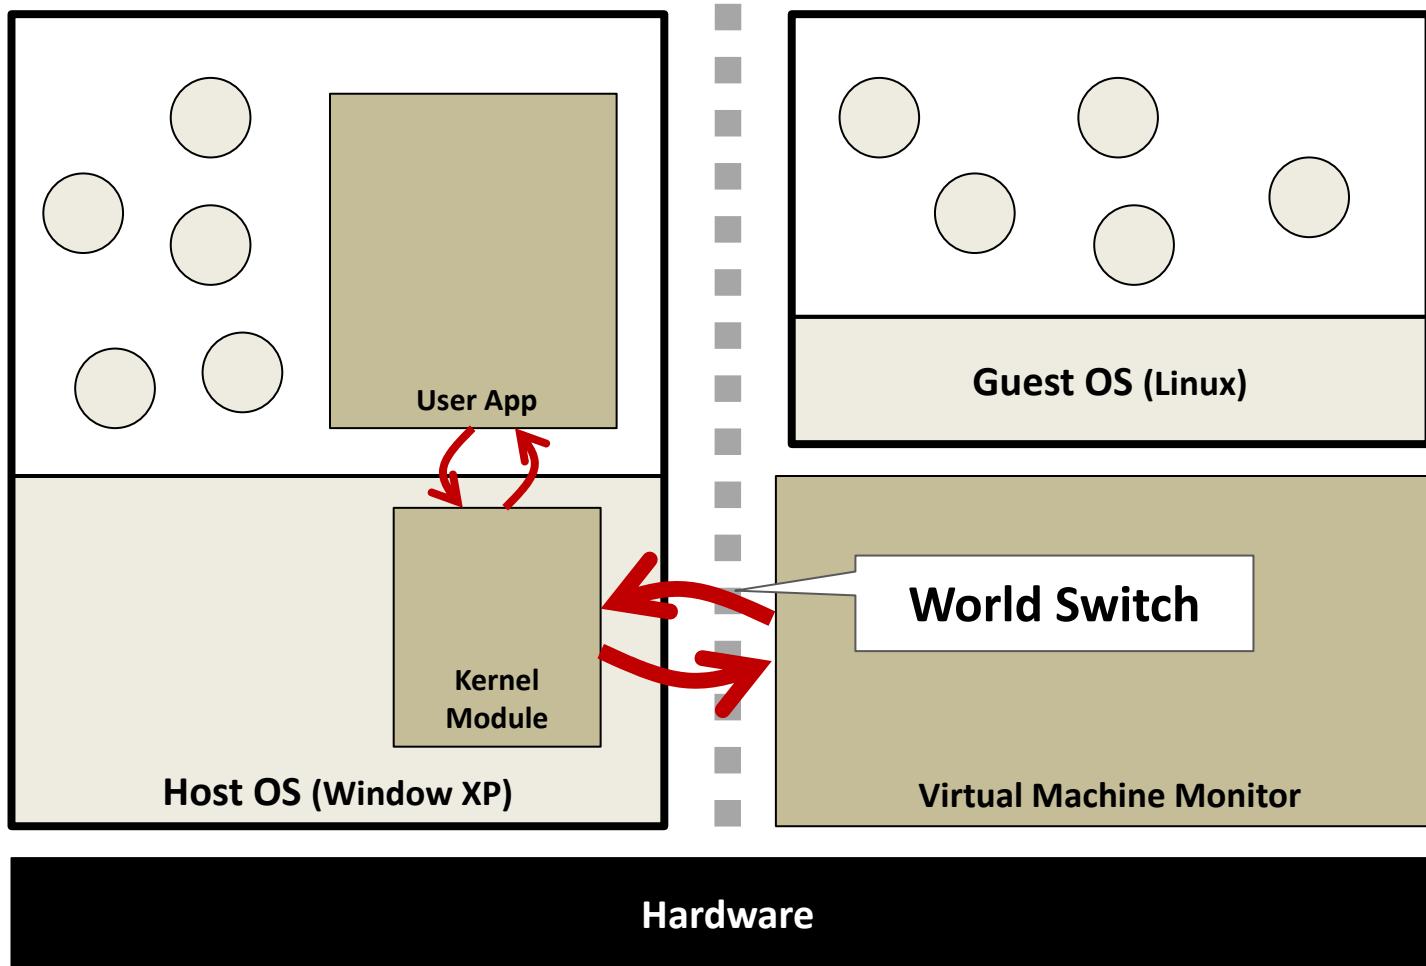


- Examples: IBM VM/370, Stanford DISCO

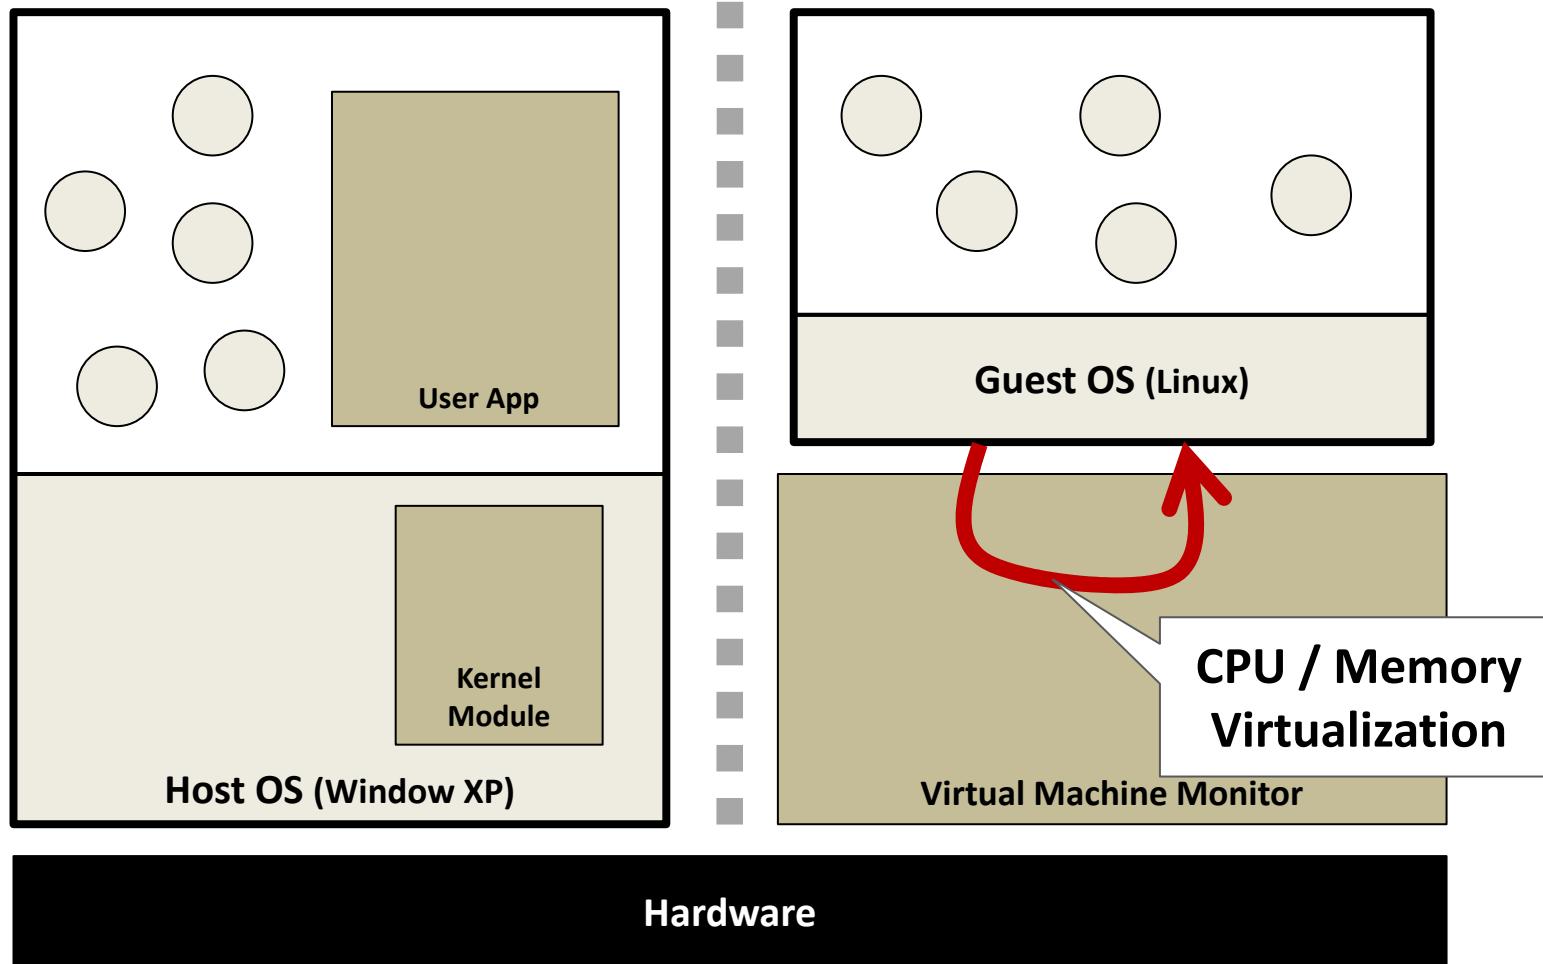
Hosted Virtual Machines

- Goal:
 - Run Virtual Machines as an application on an existing Operating System
- Why
 - Application continuity
 - Reuse existing device drivers
 - Leverage OS support
 - File system
 - CPU Scheduler
 - VM management platform

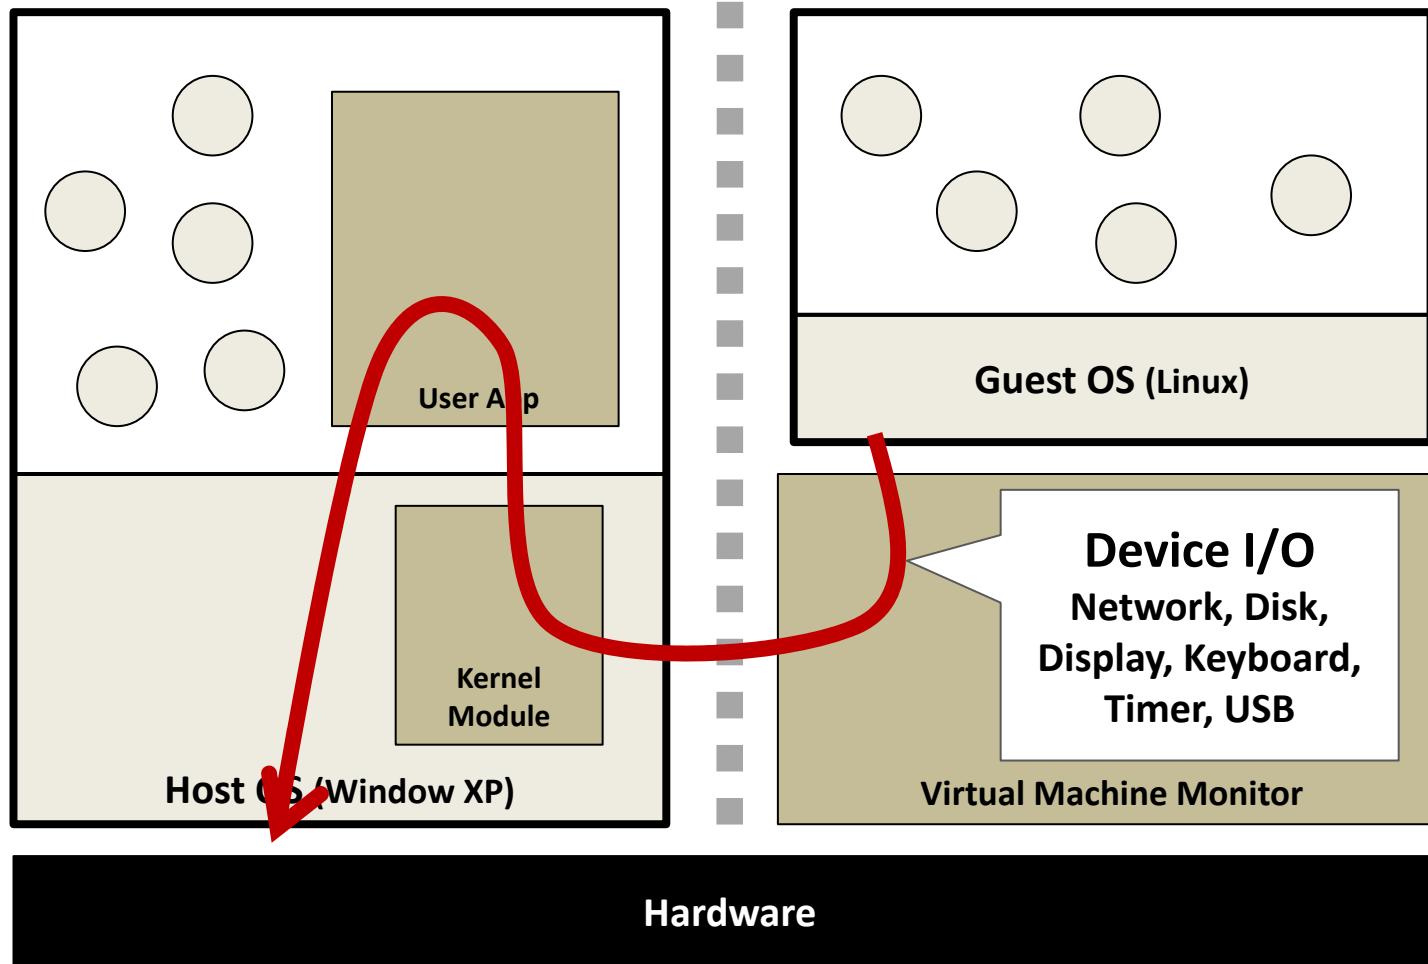
Hosted Monitor Architecture



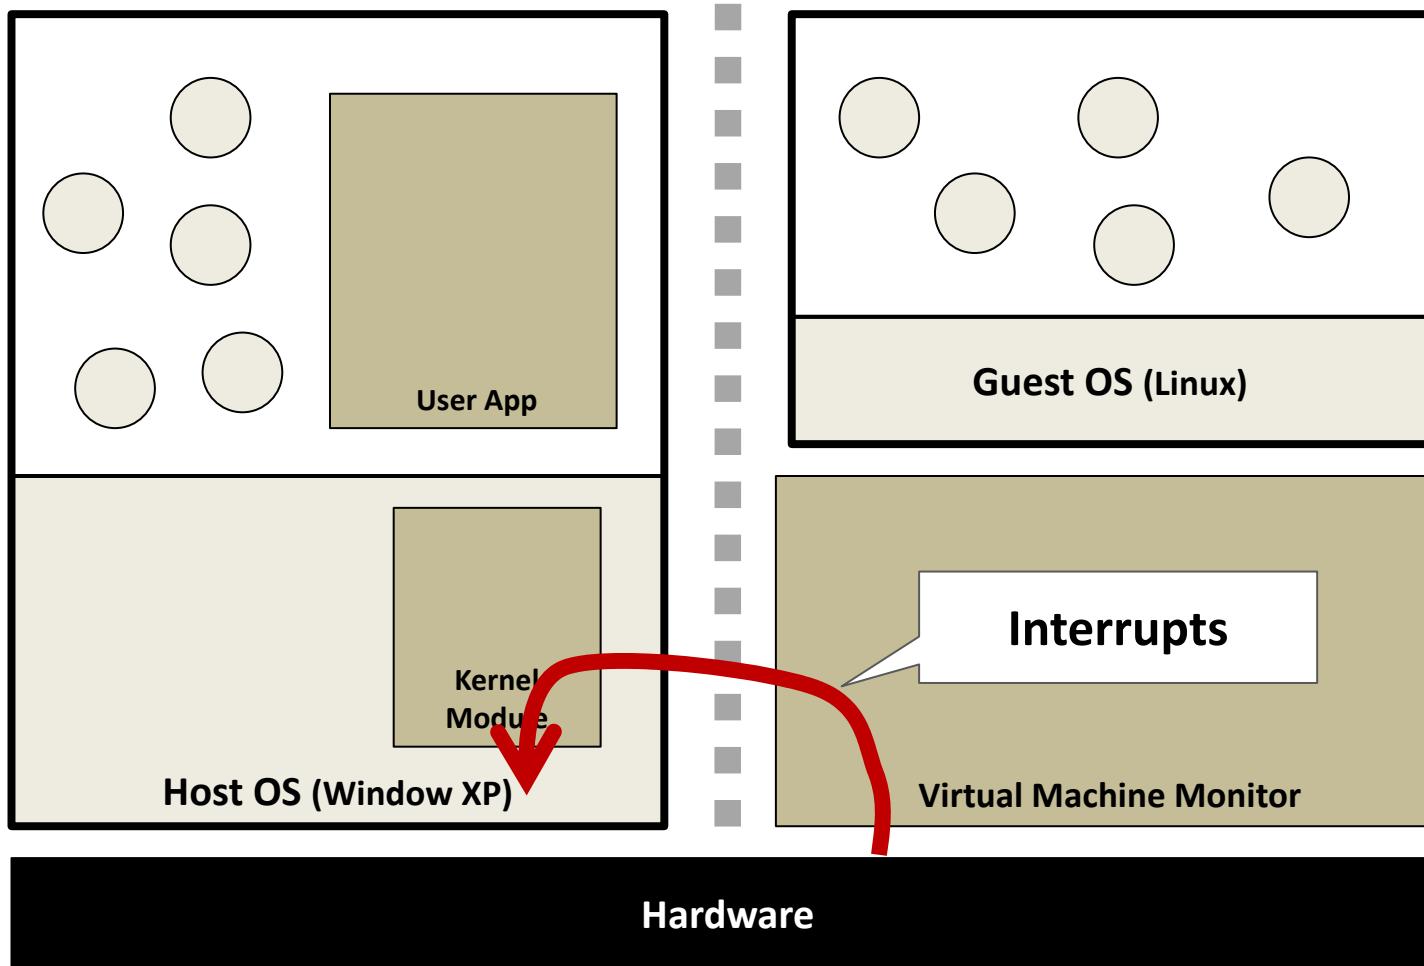
Hosted Monitor Architecture



Hosted Monitor Architecture



Hosted Monitor Architecture



Hosted Architecture Tradeoffs

- Positives
 - Installs like an application
 - No disk partitioning needed
 - Virtual disk is a file on host file system
 - No host reboot needed
 - Runs like an application
 - Uses host schedulers
- Negatives
 - I/O path is slow
 - Requires world switch
 - Relies on host scheduling
 - May not be suitable for intensive VM workloads

VMware ESX 2.0

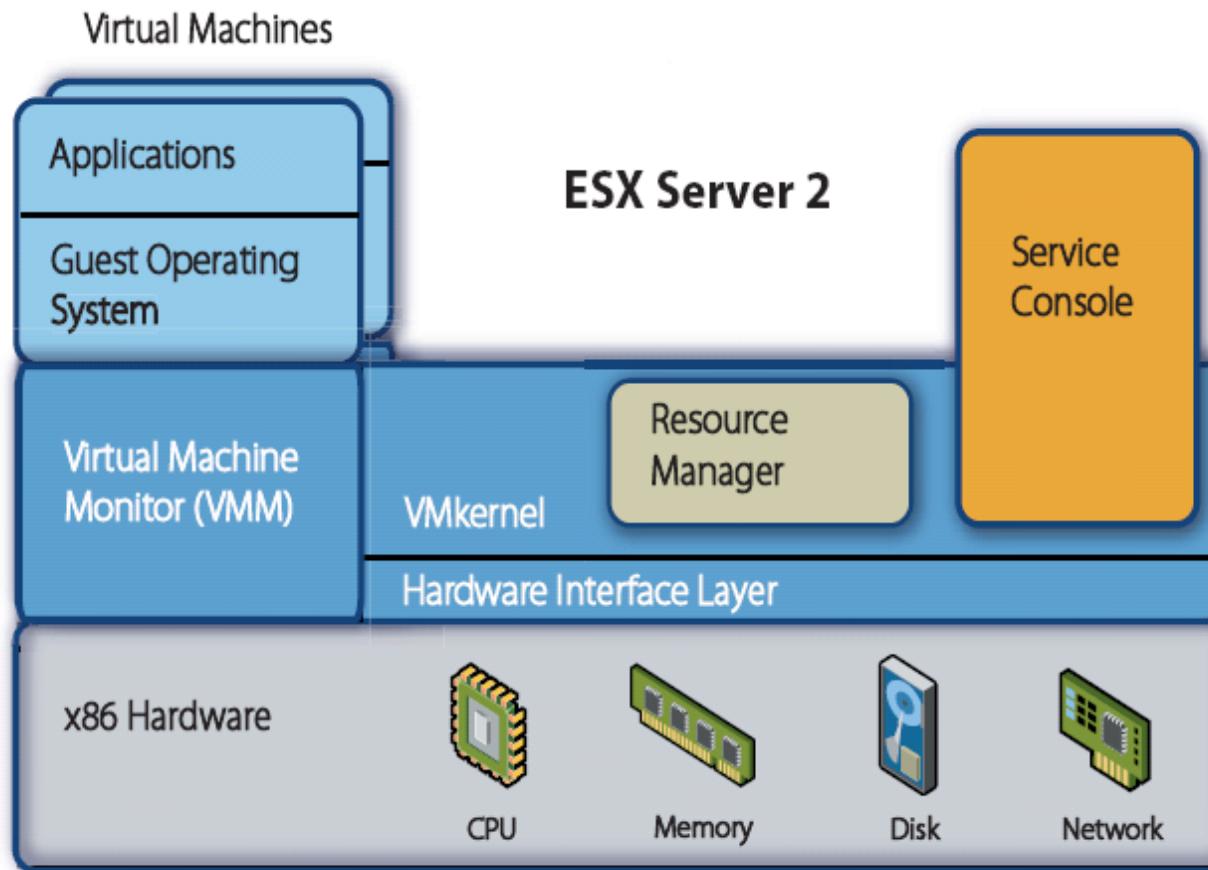
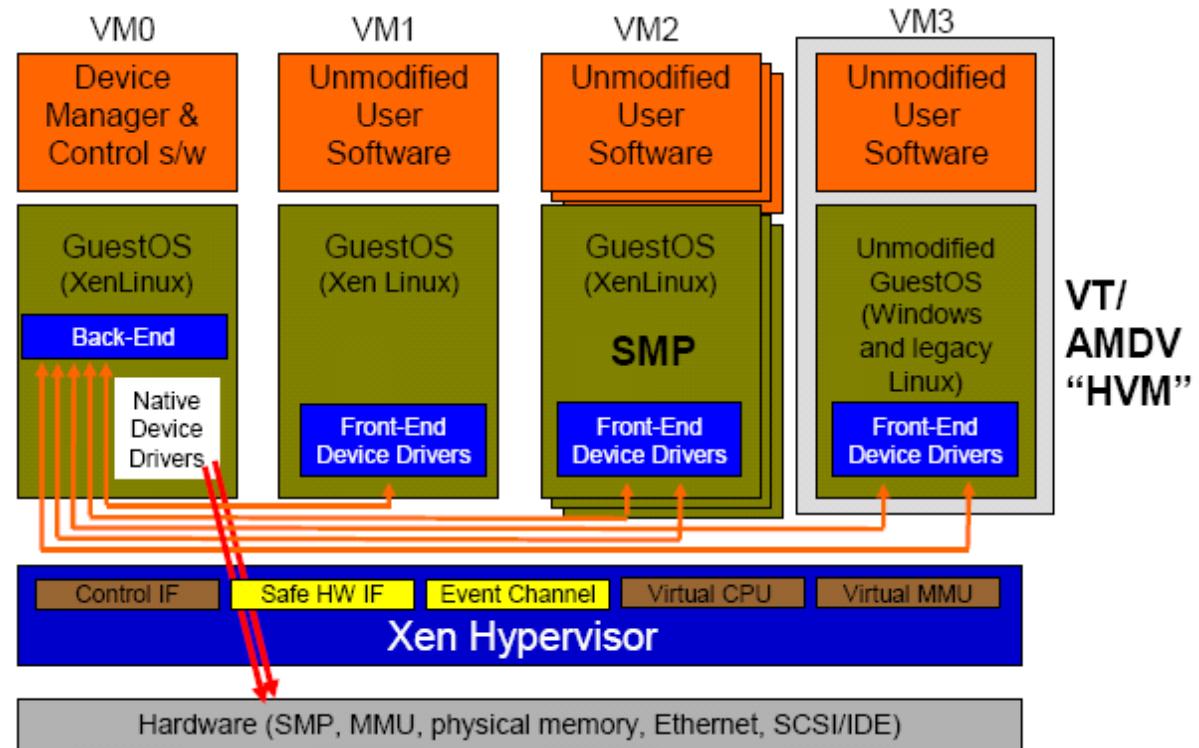


Figure 1: ESX Server architecture

Source: http://www.vmware.com/pdf/esx2_performance_implications.pdf

Hybrid Ex 2 - Xen 3.0

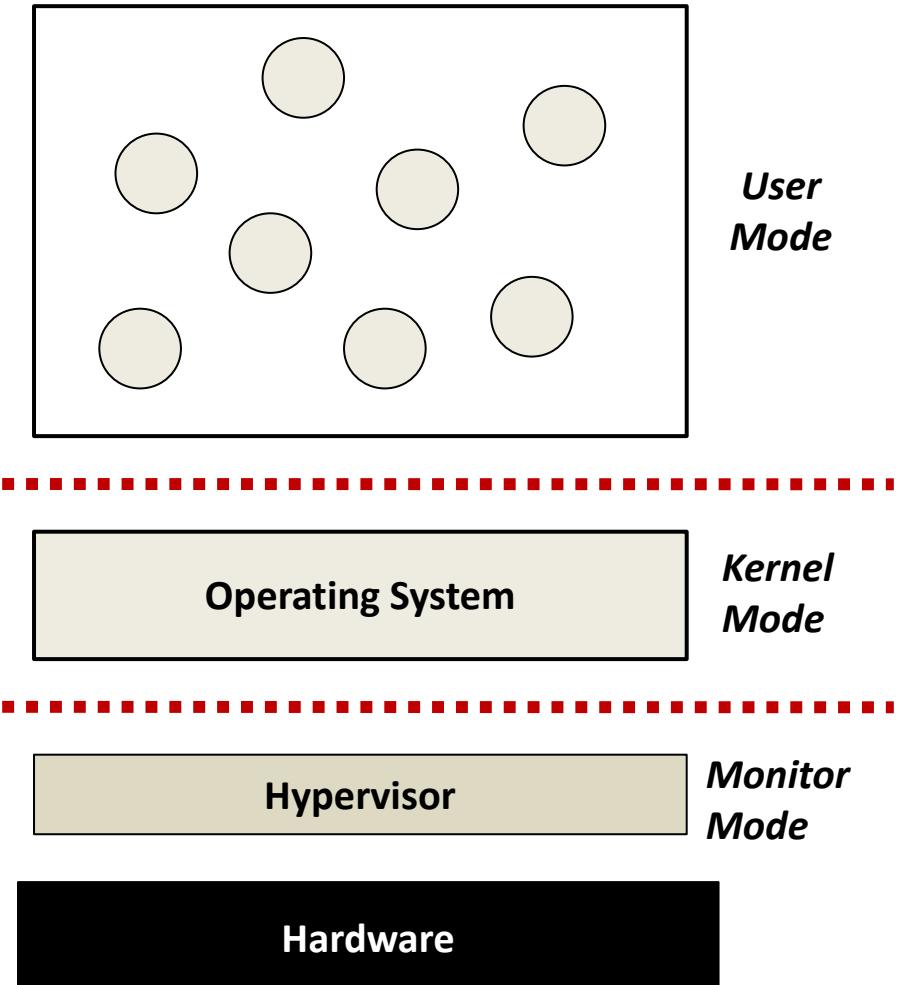
- Para
–virtualization
 - Linux Guest
- Hardware-
supported
virtualization
 - Unmodified
Windows
- Isolated Device
Drivers



*Source: Ottawa Linux Symposium 2006 presentation.
<http://www.cl.cam.ac.uk/netos/papers/>*

Hypervisor

- Hardware-supported single-use monitor
- Characteristics
 - Small size
 - Runs in a special hardware mode
 - Guest OS runs in normal privileged level
- Uses
 - Security
 - System management
 - Fault tolerance

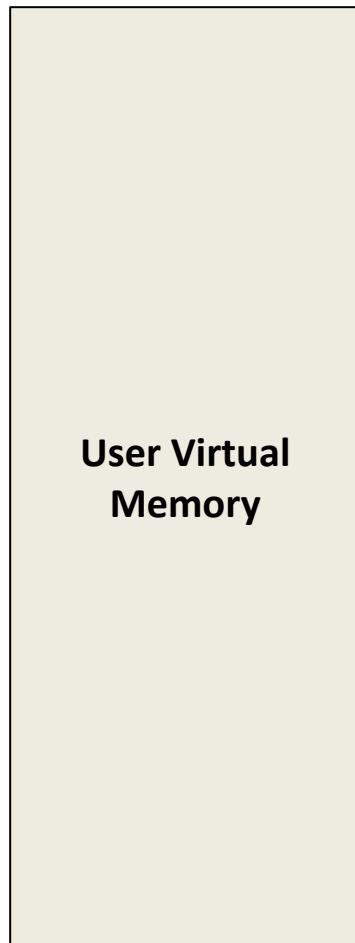


CPU VIRTUALIZATION

Recap: CPU Organization

- Instruction Set Architecture (ISA)
 - Defines:
 - the state visible to the programmer
 - registers and memory
 - the instruction that operate on the state
- ISA typically divided into 2 parts
 - User ISA
 - Primarily for computation
 - System ISA
 - Primarily for system resource management

Recap: User ISA - State



Special-Purpose Registers

Program Counter

Condition Codes

General-Purpose Registers

Reg 0

Reg 1

Reg n-1

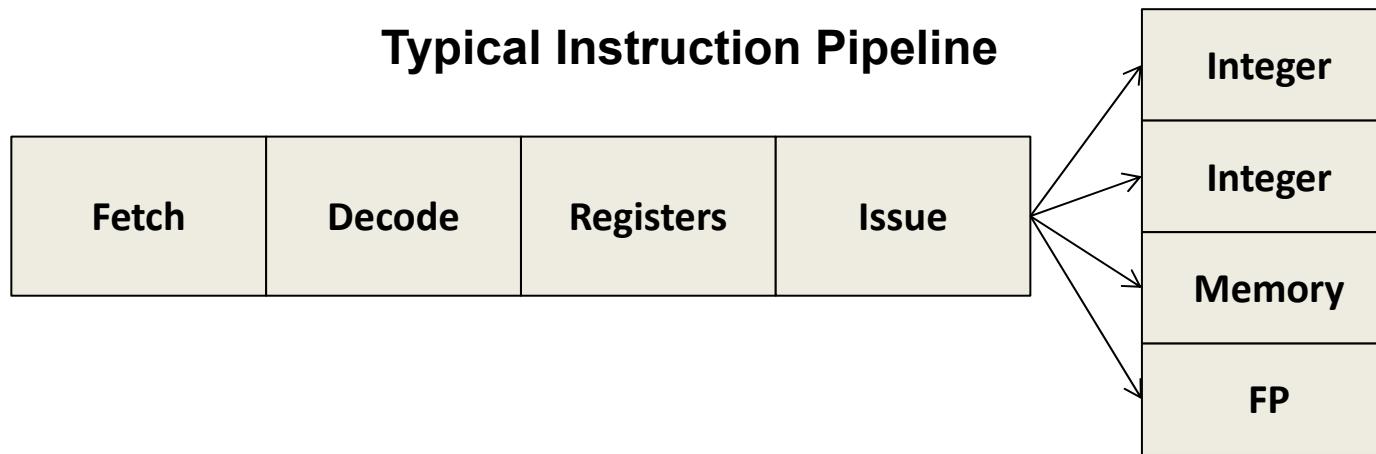
Floating Point Registers

FP 0

FP 1

FP n-1

Recap: User ISA – Instructions

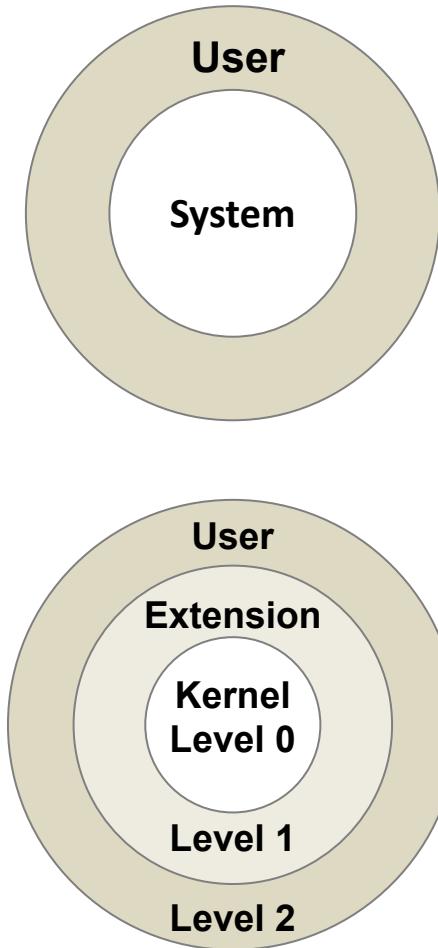


Integer	Memory	Control Flow	Floating Point
Add Sub And Compare ...	Load byte Load Word Store Multiple Push ...	Jump Jump equal Call Return ...	Add single Mult. double Sqrt double ...

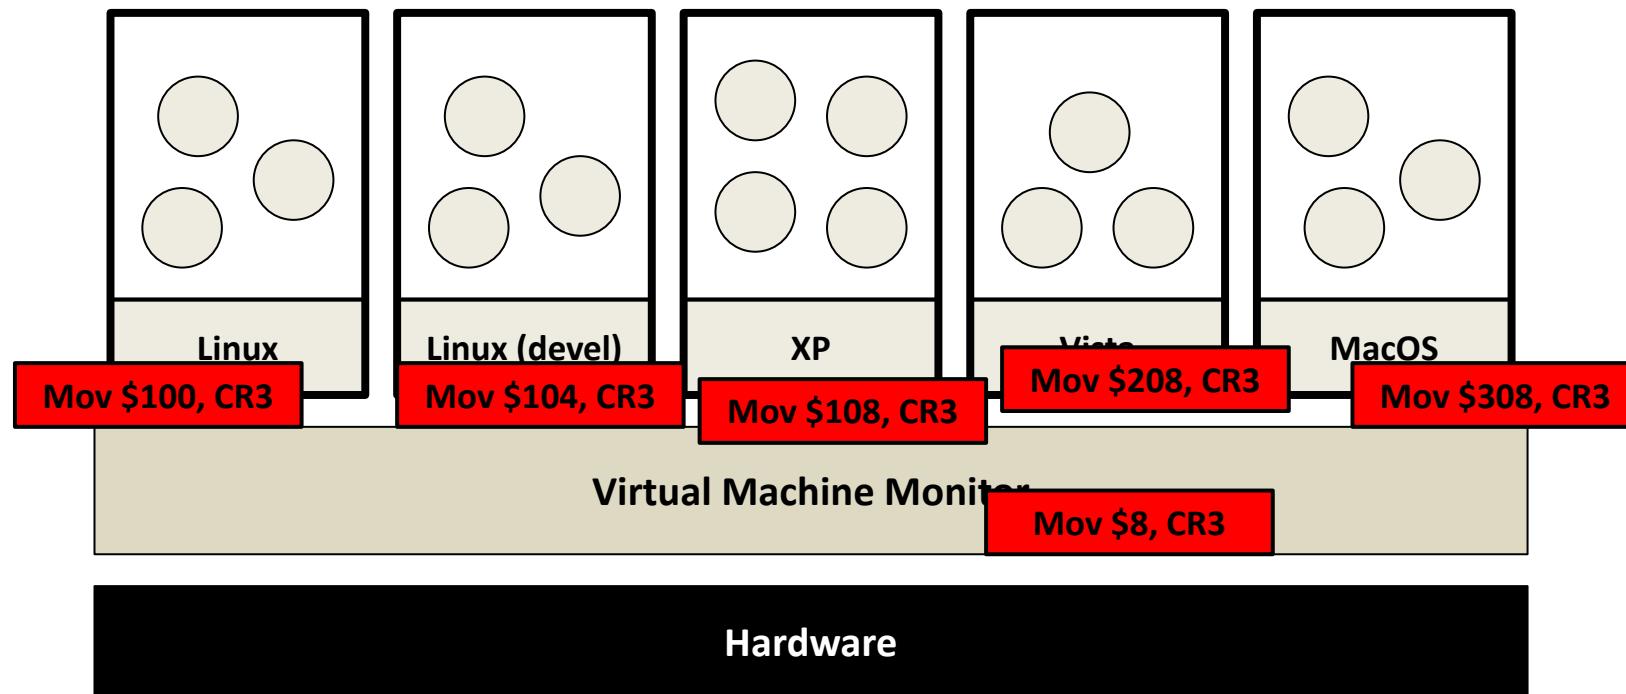
Instruction Groupings

Recap: System ISA

- Privilege Levels
- Control Registers
- Traps and Interrupts
 - Hardcoded Vectors
 - Dispatch Table
- System Clock
- MMU
 - Page Tables
 - TLB
- I/O Device Access



Challenges to Virtualized System ISA



Virtualizing the System ISA

- Hardware needed by monitor
 - Ex: monitor must control real hardware interrupts
- Access to hardware would allow VM to compromise isolation boundaries
 - Ex: access to MMU would allow VM to write any page
- So...
 - All access to the virtual System ISA by the guest must be **emulated** by the monitor in software.
 - System state kept in **memory**.
 - System instructions are implemented as **functions** in the monitor.

Example: CPUSState

```
static struct {
    uint32 GPR[16];
    uint32 LR;
    uint32 PC;
    int IE;
    int IRQ;
} CPUSState;
```

```
void CPU_CLI(void)
{
    CPUSState.IE = 0;
}
```

- Goal for CPU virtualization techniques
 - Process normal instructions as fast as possible
 - Forward privileged instructions to emulation routines

Virtualization VS. Multiplexing

- Why not run virtual machines just as user apps?
 - Virtual machine includes guest kernel and guest apps
 - What is the problem?

Definition

- **Sensitive Instructions**
 - Instructions that inspect or change system states of a processor
 - E.g., mov to CR registers
- **Privileged Instructions**
 - Instructions that executed at high privileged level
- What if an instruction is sensitive but not privileged?

Formal Requirement of Virtualization

- Popek & Goldberg, 1974 “Formal Requirements for Virtualizable Third Generation Architectures”
- Provide by “virtual machine monitor” with three essential characteristics:
 - Essentially identical execution environment (as real machine)
 - Minor performance penalty for programs in VM
 - VMM has complete control over system resources

x86 virtualization challenge

- The IA-32 instruction set contains 17 **sensitive, unprivileged** instructions:
 - Sensitive register instructions: read or change sensitive registers and/or memory locations such as a clock register or interrupt registers:
 - SGDT, SIDT, SLDT, SMSW, PUSHF, POPF
 - Protection system instructions: reference the storage protection system, memory or address relocation system:
 - LAR, LSL, VERR, VERW, POP, PUSH, CALL, JMP, INT n, RET, STR, MOV

Solutions

- Instruction Interpretation
- Trap-and-emulate
- Binary Translation
- Para-virtualization

Instruction Interpretation

- Emulate Fetch/Decode/Execute pipeline in software
- Positives
 - Easy to implement
 - Minimal complexity
- Negatives
 - Slow!

Example: Virtualizing the Interrupt Flag w/ Instruction Interpreter

```
void CPU_Run(void)
{
    while (1) {
        inst = Fetch(CPUState.PC);

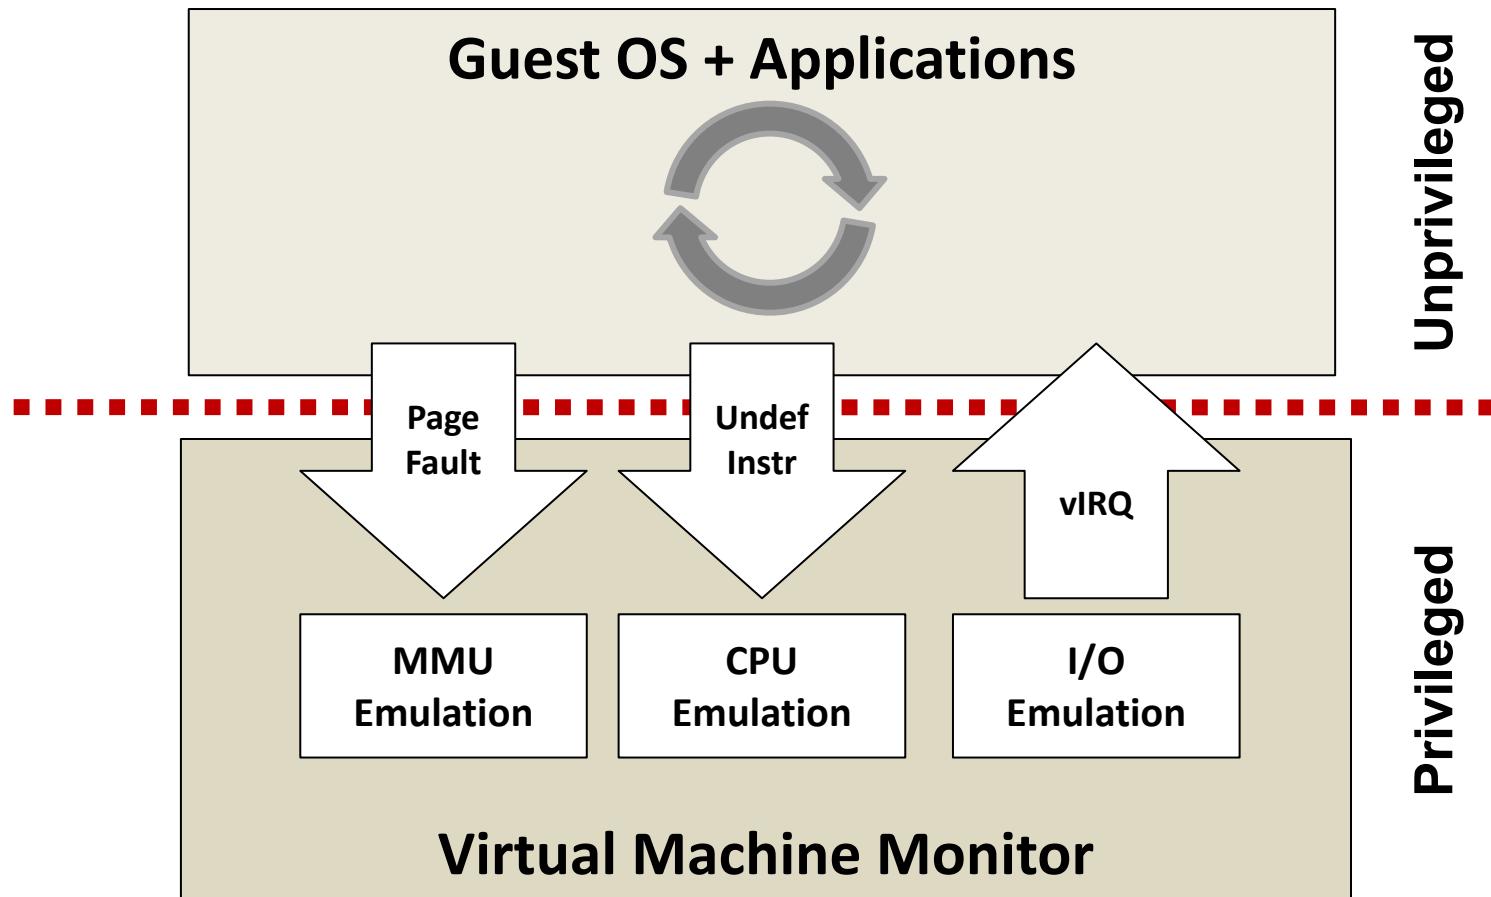
        CPUState.PC += 4;

        switch (inst) {
        case ADD:
            CPUState.GPR[rd]
                = GPR[rn] + GPR[rm];
            break;
        ...
        case CLI:
            CPU_CLI();
            break;
        case STI:
            CPU_STI();
            break;
        }

        if (CPUState.IRQ
            && CPUState.IE) {
            CPUState.IE = 0;
            CPU_Vector(EXC_INT);
        }
    }
}

void CPU_CLI(void)
{
    CPUState.IE = 0;
}
```

Trap and Emulate



Issues with Trap and Emulate

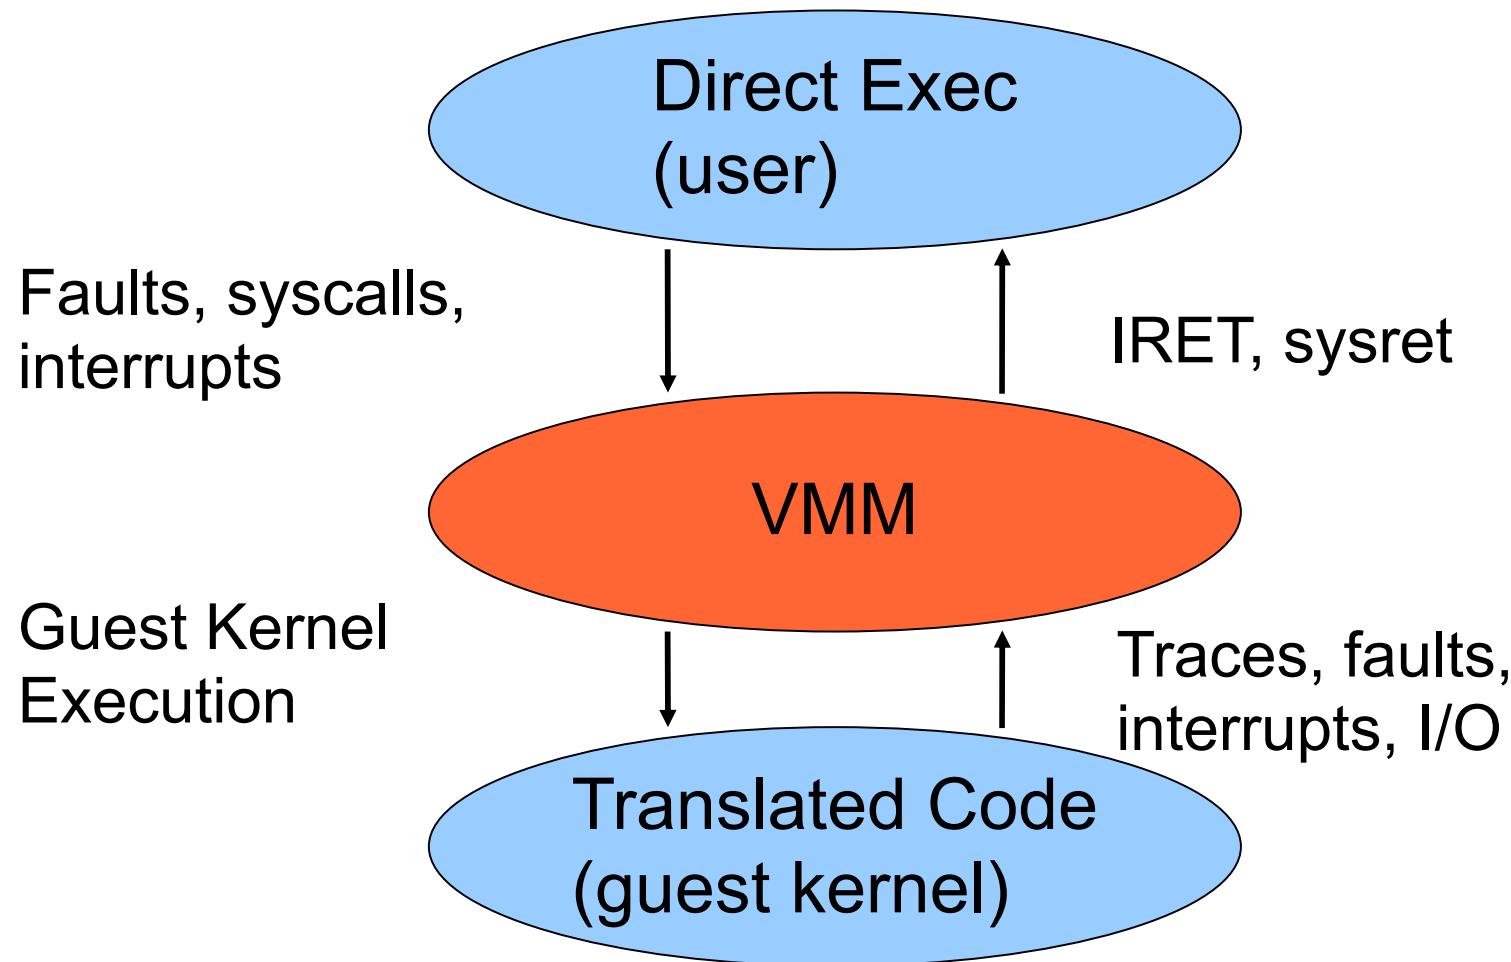
- Not all architectures support it
- Trap costs may be high
- Monitor uses a privilege level
 - Need to virtualize the protection levels

“Strictly Virtualizable”

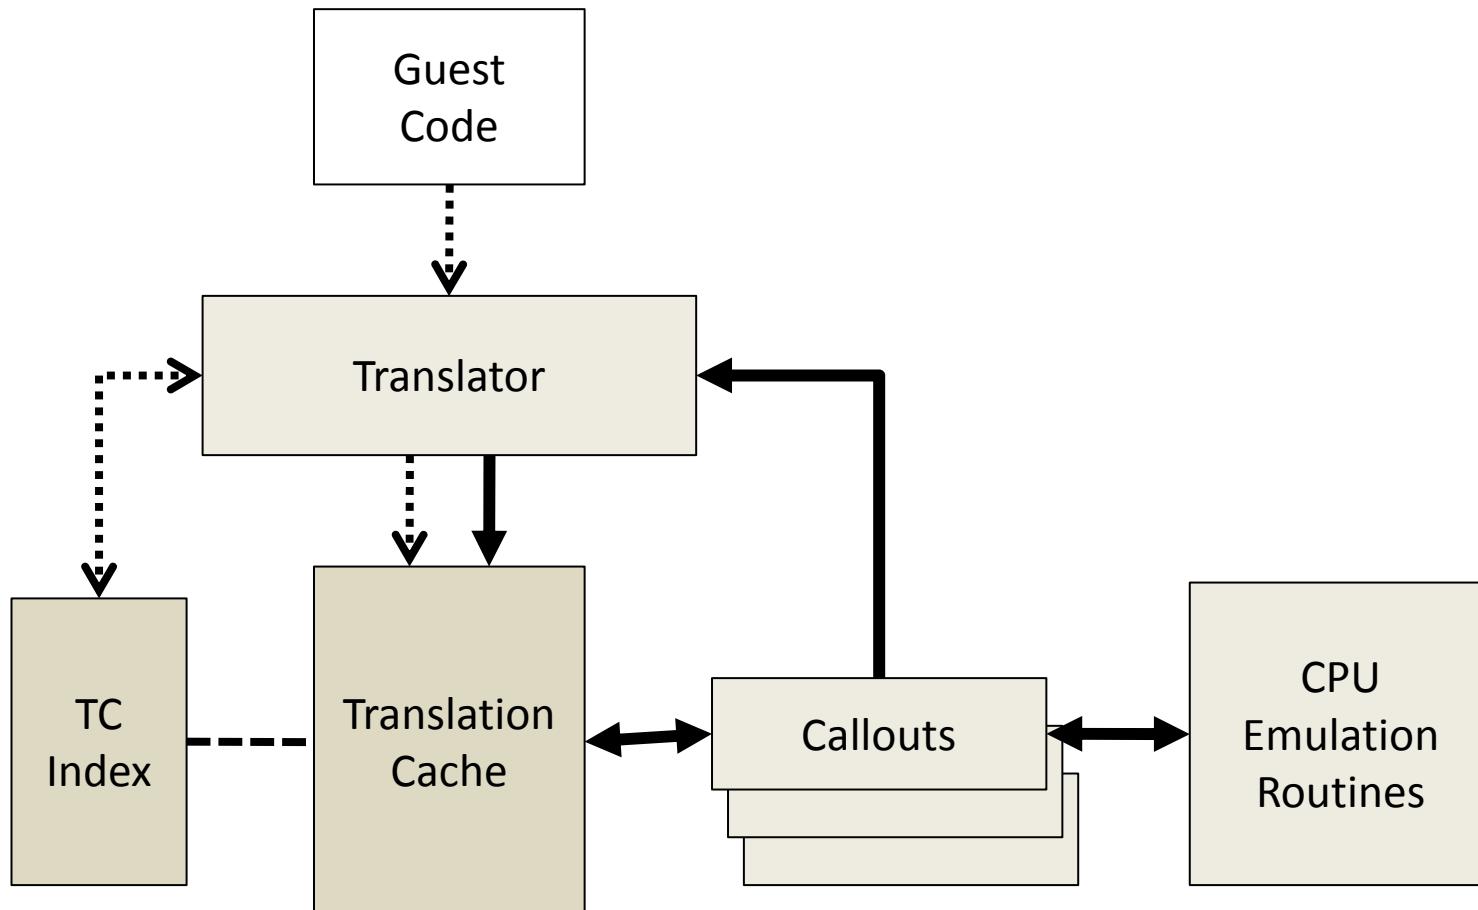
A processor or mode of a processor is strictly virtualizable if, when executed in a lesser privileged mode:

- all instructions that access privileged state trap
- all instructions either trap or execute identically
- ...

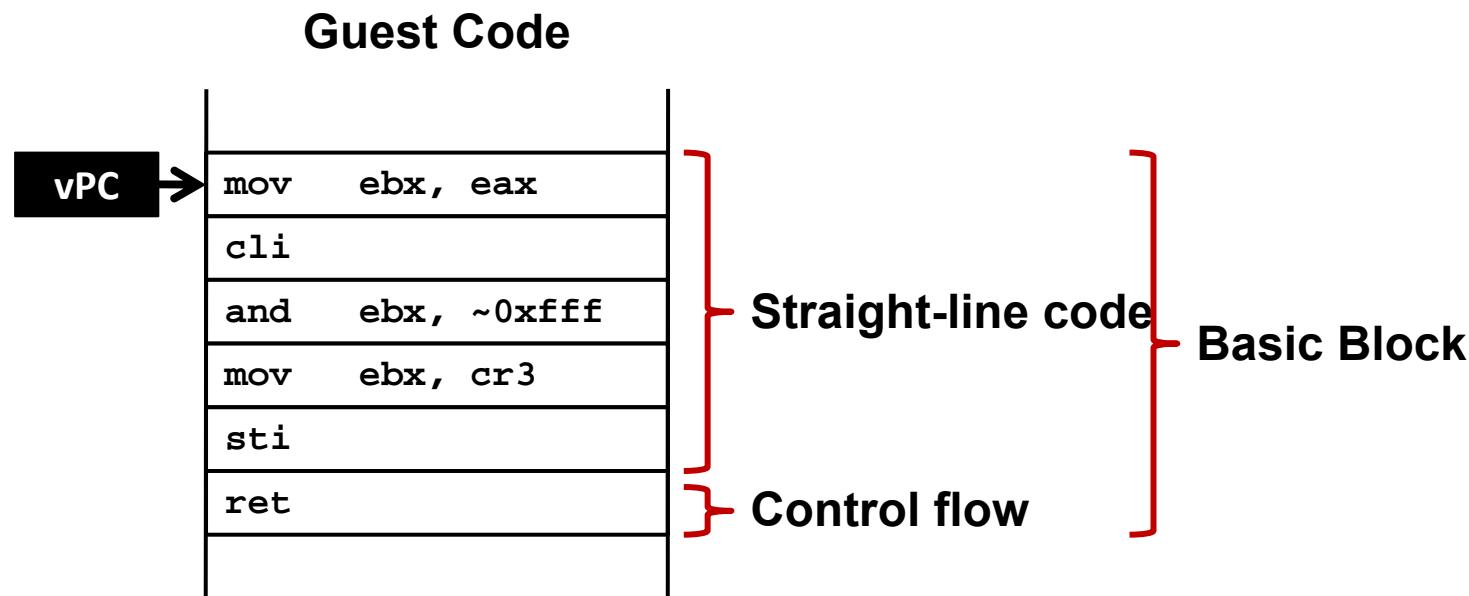
Software VMM



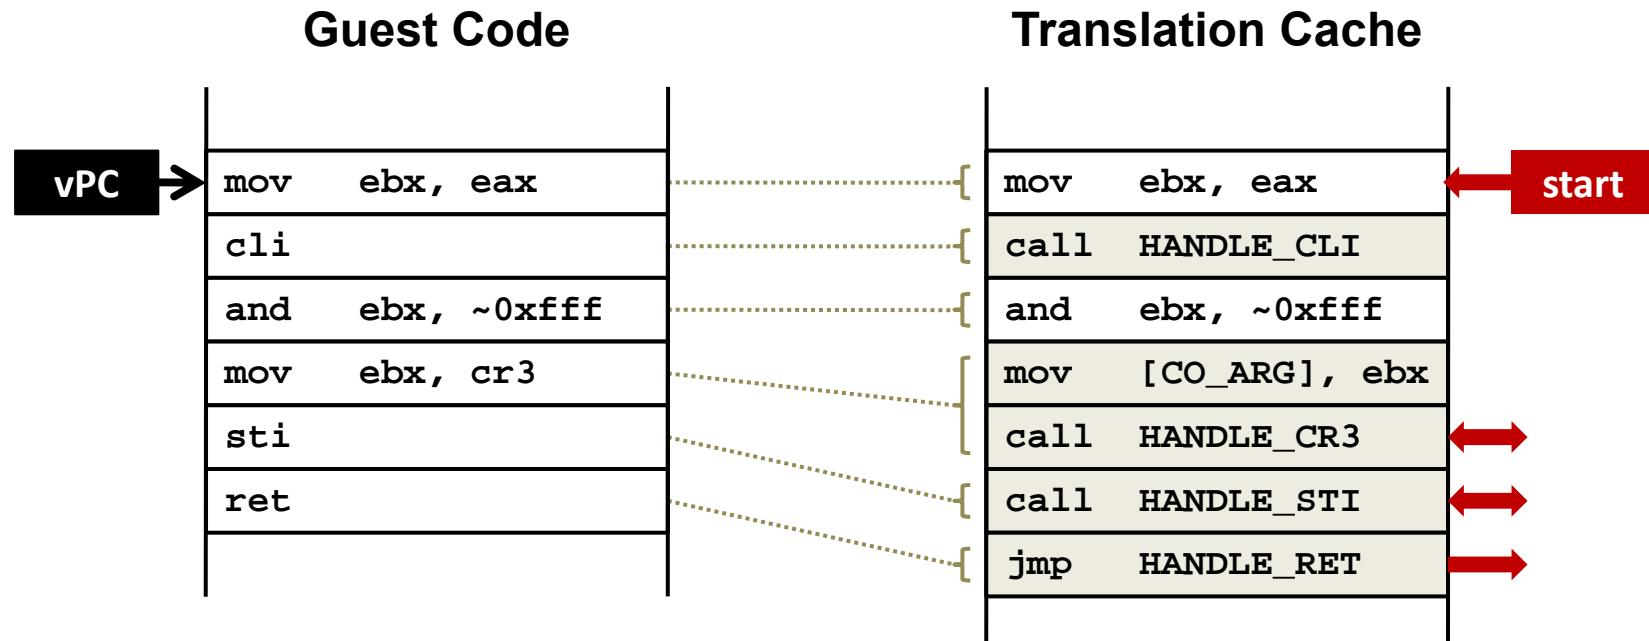
Binary Translator



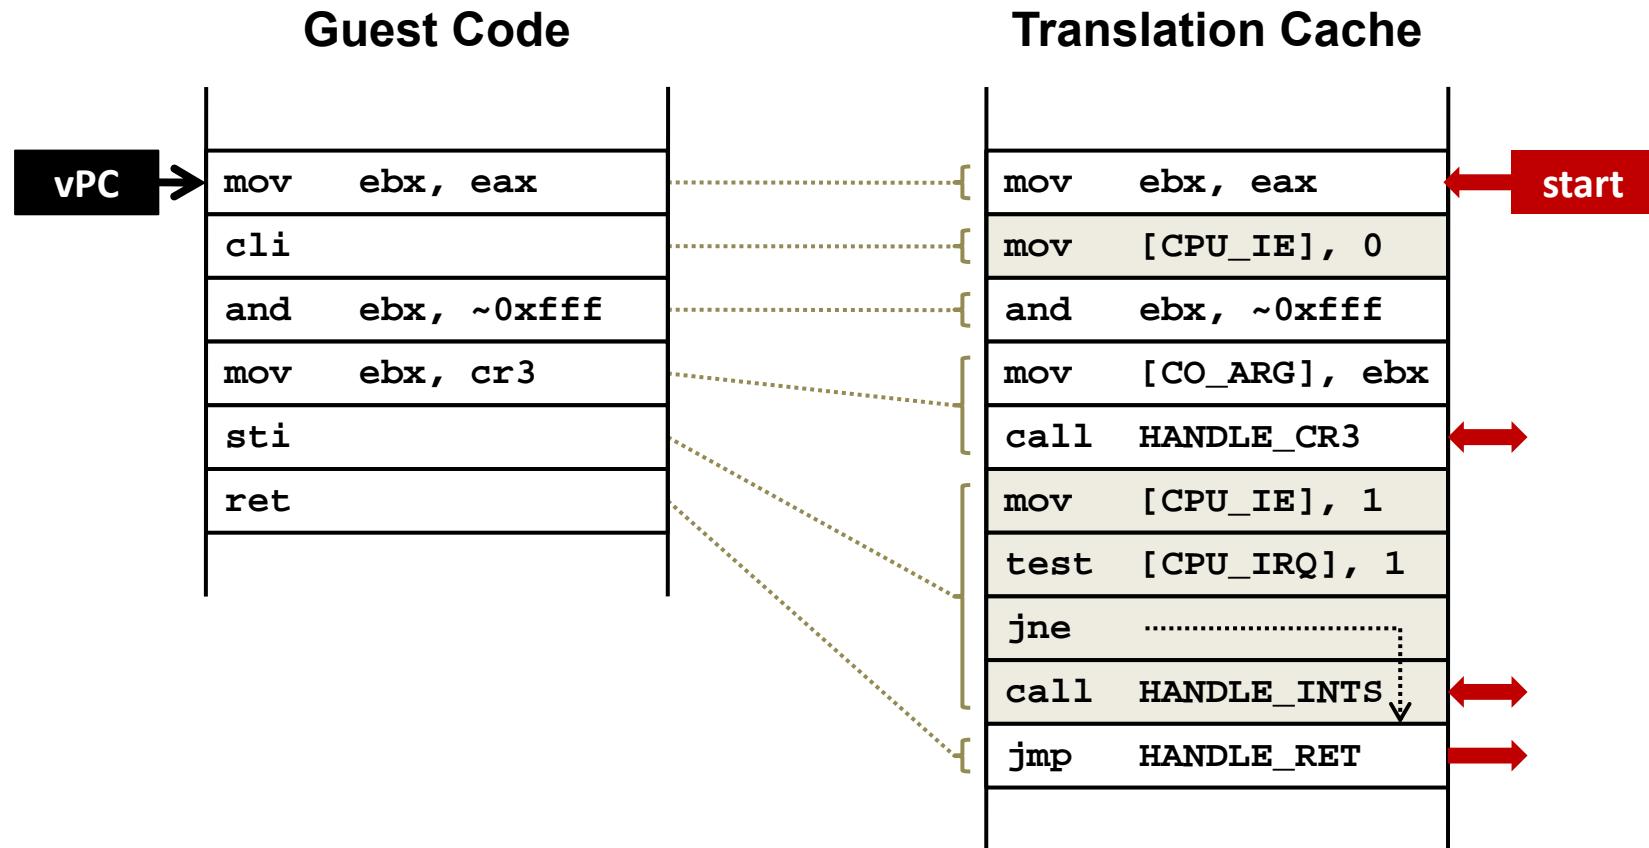
Basic Blocks



Binary Translation



Binary Translation



Basic Binary Translator

```
void BT_Run(void)
{
    CPUState.PC = _start;
    BT_Continue();
}

void BT_Continue(void)
{
    void *tcpc;

    tcpc = BTFindBB(CPUState.PC);

    if (!tcpc) {
        tcpc = BTTranslate(CPUState.PC);
    }

    RestoreRegsAndJump(tcpc);
}

void *BTTranslate(uint32 pc)
{
    void *start = TCTop;
    uint32 TCPC = pc;
```

Basic Binary Translator – Part 2

```
void BT_CalloutSTI(BTSavedRegs regs)
{
    CPUState.PC = BTFindPC(regs.tcpc);
    CPUState.GPR[] = regs.GPR[];

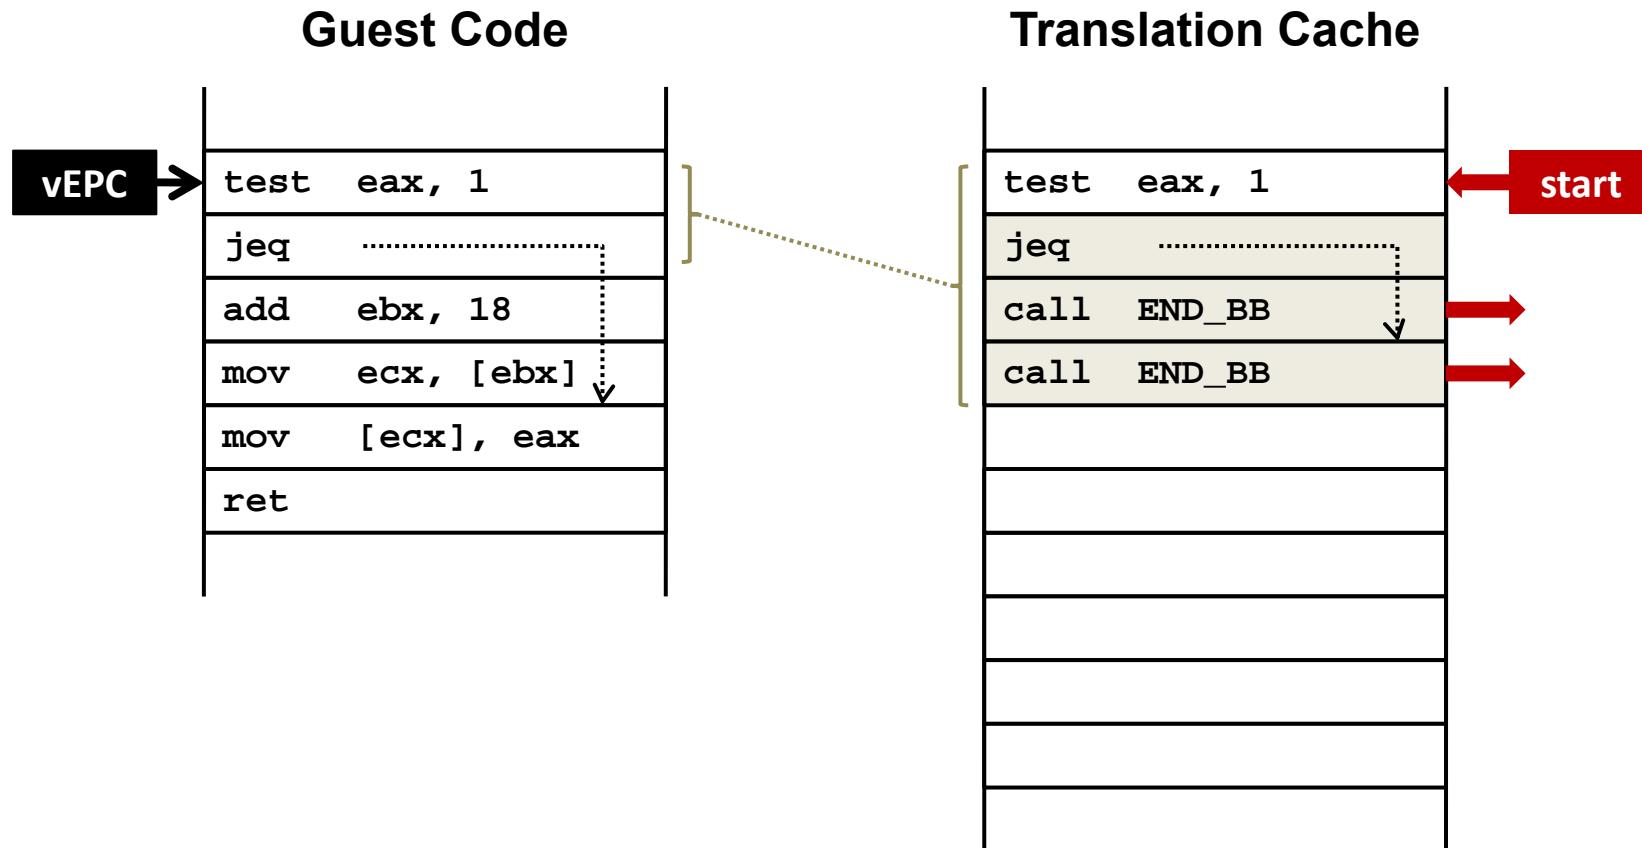
    CPU_STI();

    CPUState.PC += 4;

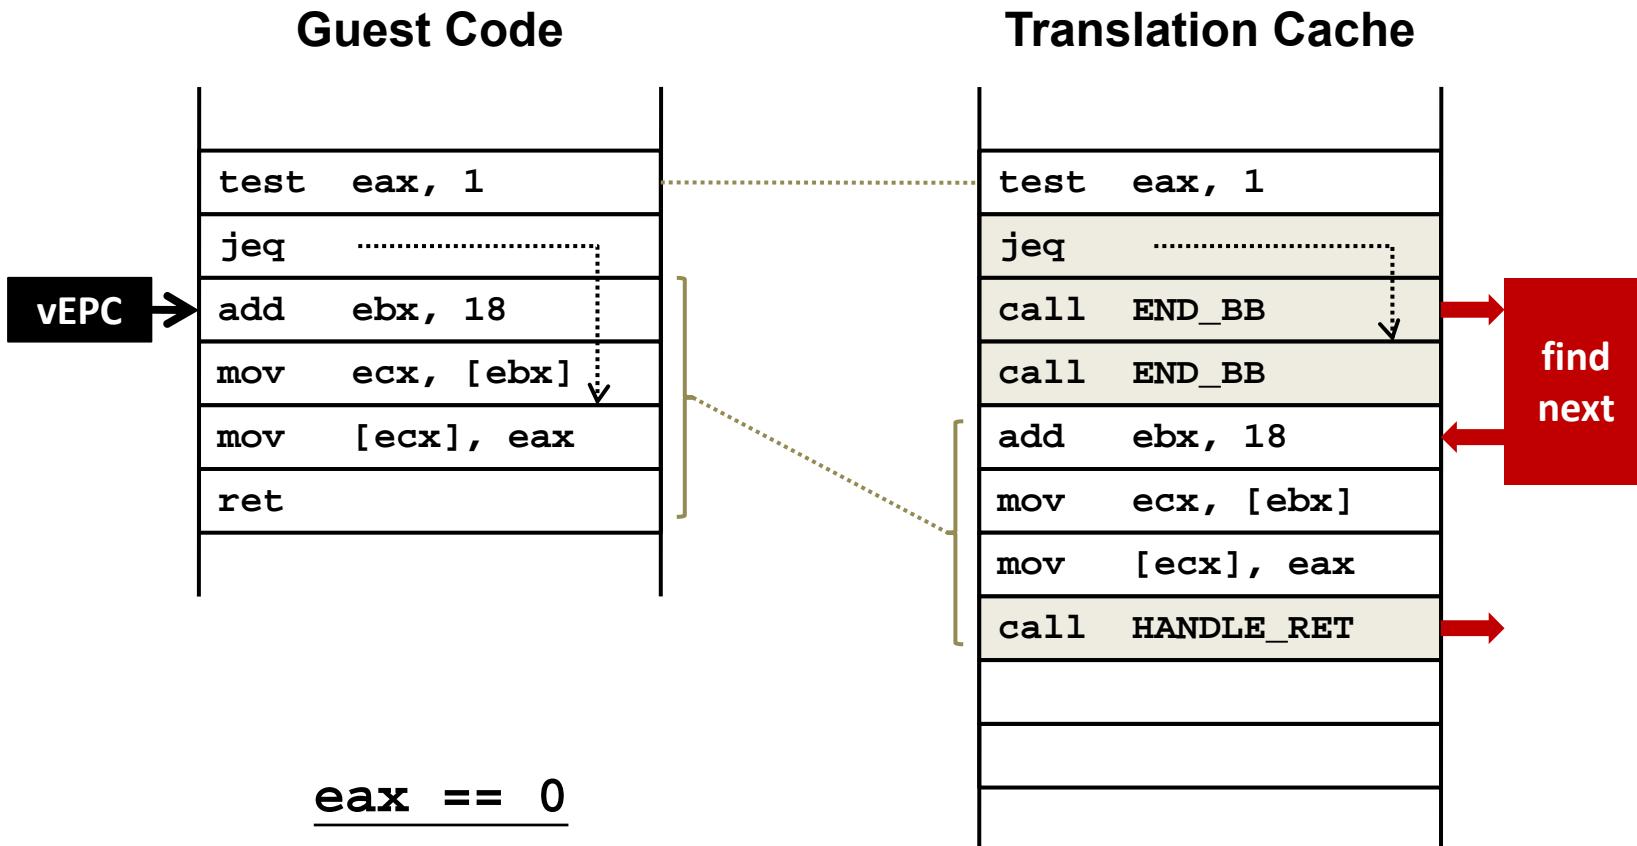
    if (CPUState.IRQ
        && CPUState.IE) {
        CPUVector();
        BT_Continue();
        /* NOT_REACHED */
    }

    return;
}
```

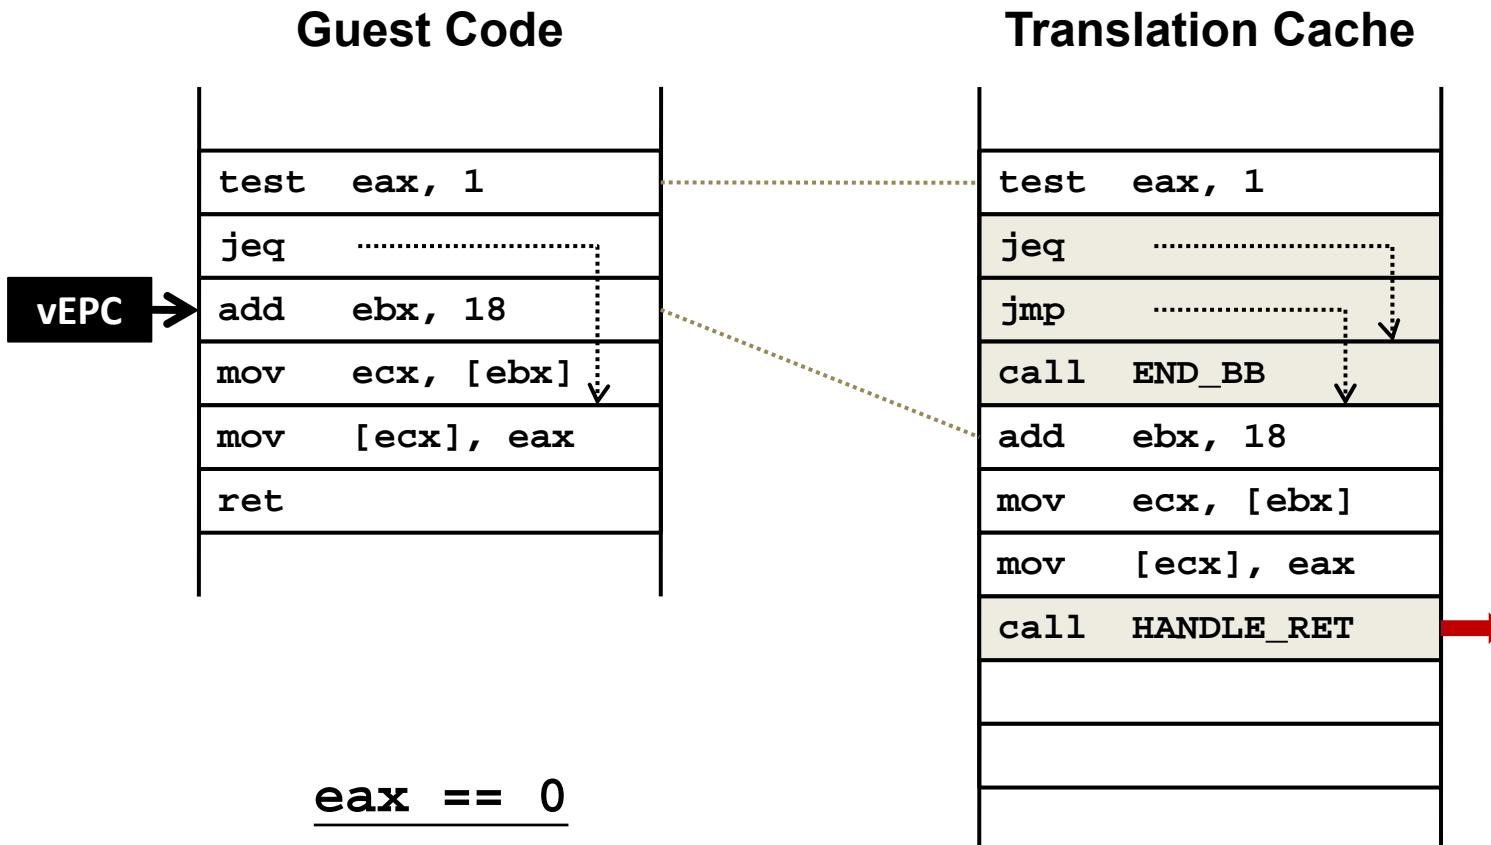
Controlling Control Flow



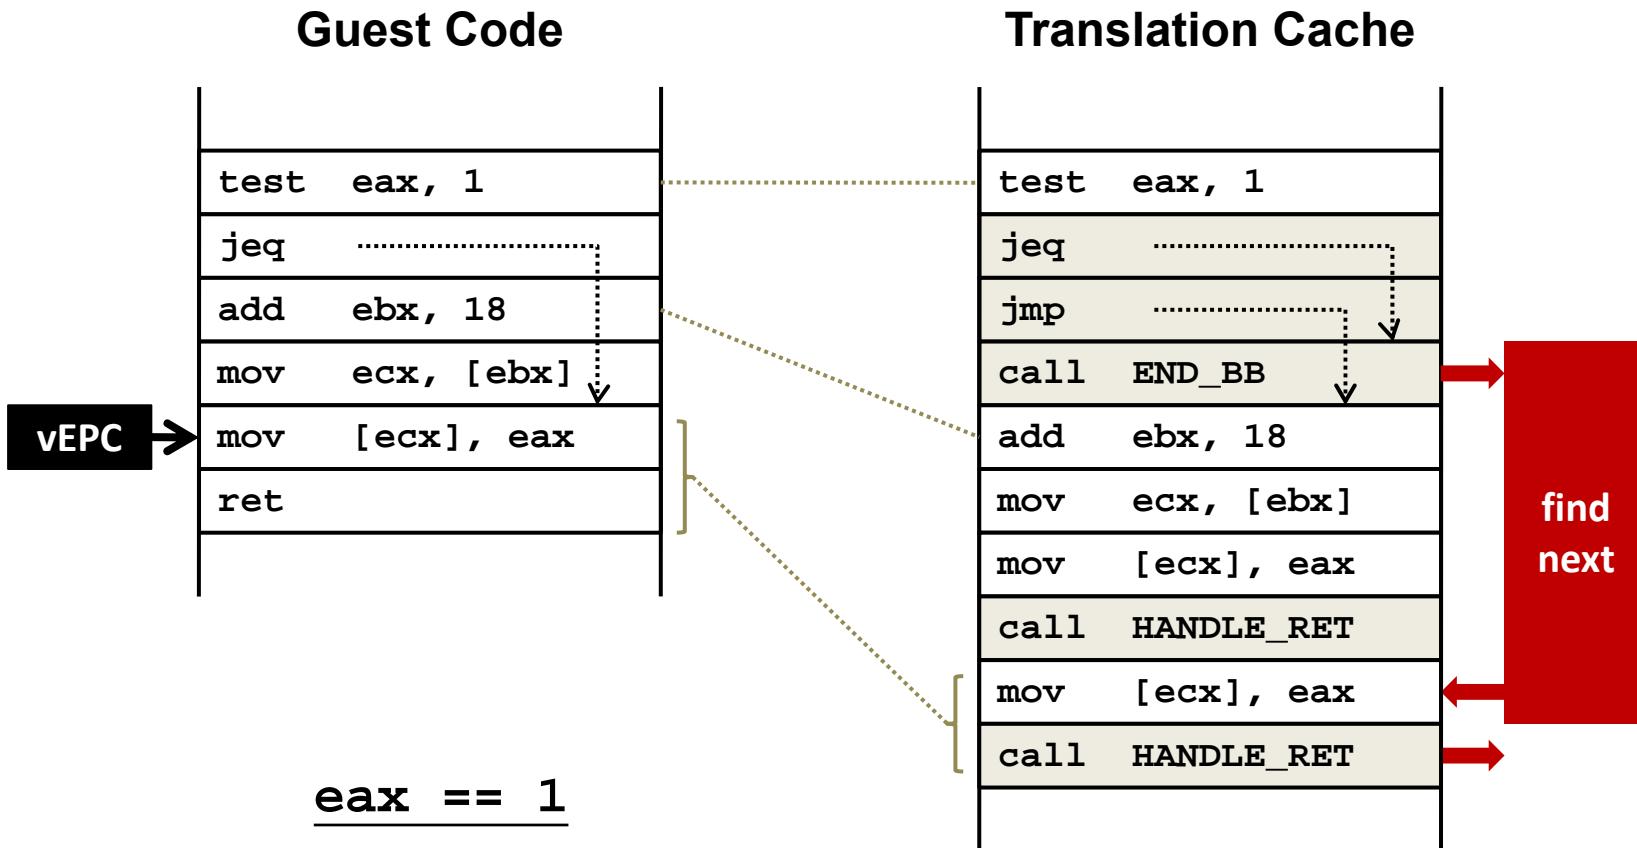
Controlling Control Flow



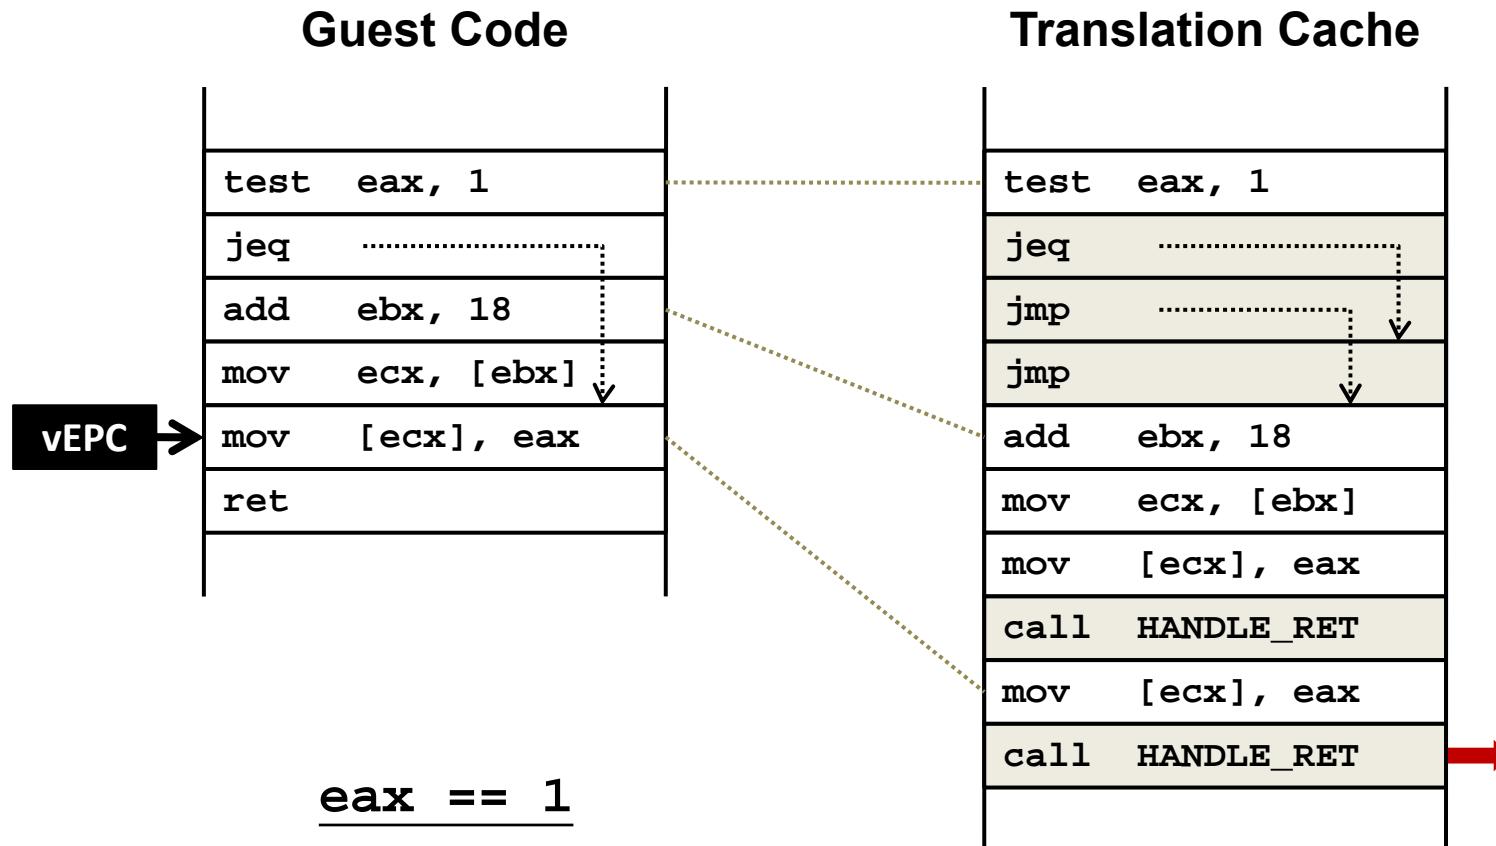
Controlling Control Flow



Controlling Control Flow



Controlling Control Flow



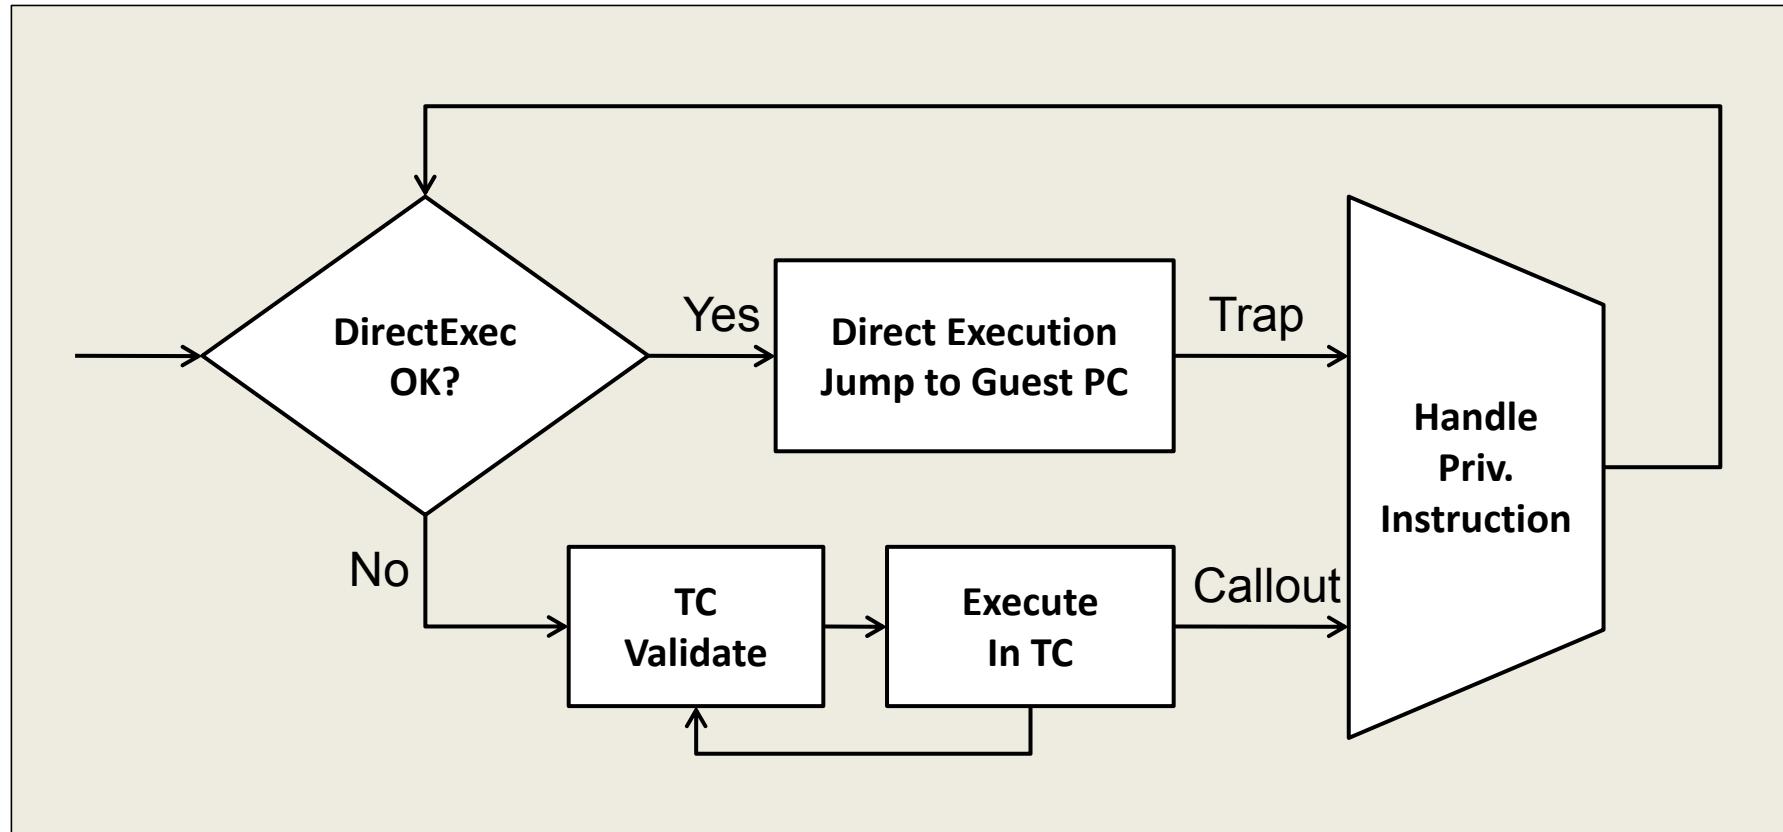
Issues with Binary Translation

- Translation cache index data structure
- PC Synchronization on interrupts
- Self-modifying code
 - Notified on writes to translated guest code

Other Uses for Binary Translation

- Cross ISA translators
 - Digital FX!32
- Optimizing translators
 - H.P. Dynamo
- High level language byte code translators
 - Java
 - .NET/CLI

Hybrid Approach



- Binary Translation for the Kernel
- Direct Execution (Trap-and-emulate) for the User
- U.S. Patent 6,397,242

Para-virtualization (e.g., Xen)

- Modify operating systems to let OS to cooperate with VMM
 - Let OS runs in Ring1 and user apps run in Ring3
 - Change sensitive instructions to explicit calls to the VMM
 - No need to trap and emulate

MEMORY VIRTUALIZATION

Virtualizing Memory

- VMM constructs a page table that maps guest address to host physical address
 - E.g., if guest VM has 1GB of memory, it can access memory address 0~1GB
 - Each guest VM has its own mapping for memory address 0, etc.
 - Different host physical address used to store data for those memory locations

Virtualizing the CR3/Page Tables

- Terminology: 3 types of address now
 - GVA->GPA->HPA (Guest virtual. Guest physical. Host physical)
 - Guest VM's page table contains guest physical address
 - Hardware page table must point to host physical address
- Setting CR3 to point to guest page table would not work
 - Processes in guest VM might access host physical address 0~1GB
 - But those HPA might not belong to that guest VM

One Solution: Shadow Pages

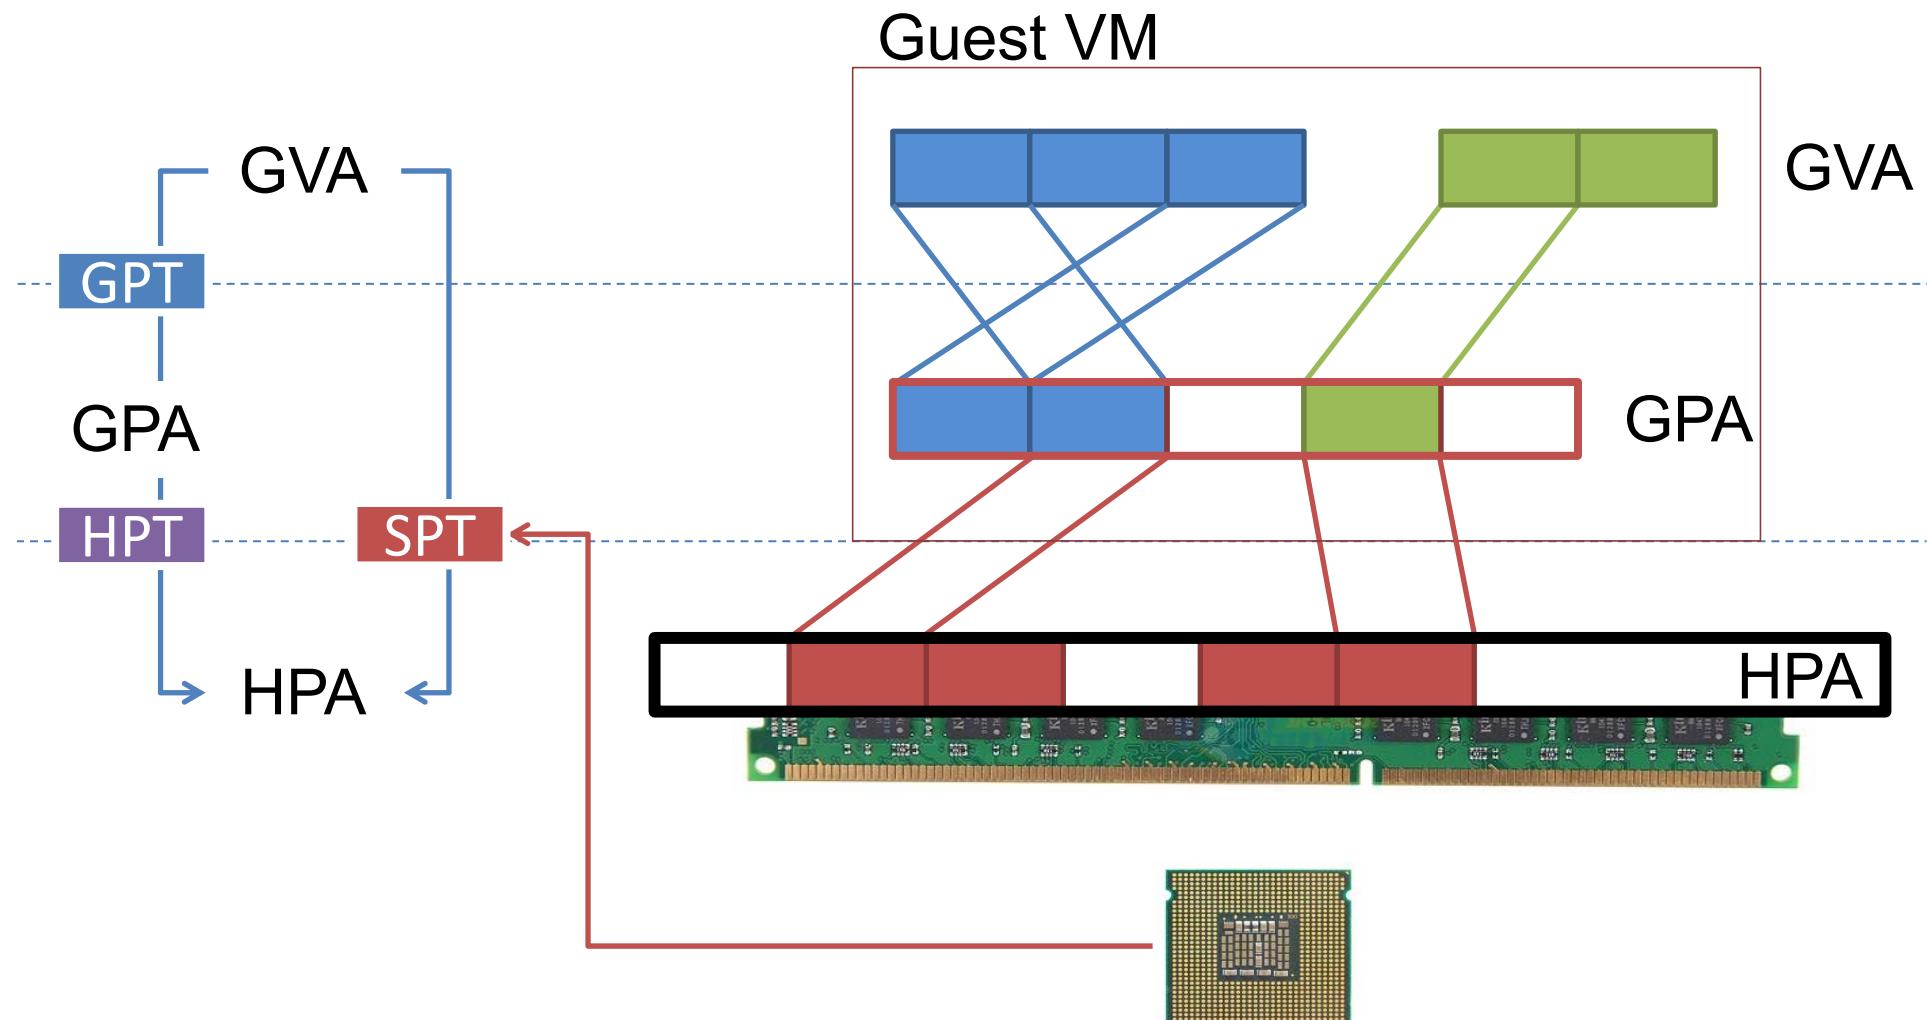
1. VMM intercepts guest OS setting the virtual CR3
2. VMM iterates over the guest page table, constructs a corresponding shadow PT
3. In shadow PT, every guest physical address is translated into host physical address
4. Finally, VMM loads the host physical address of the shadow PT

(Shadow PT is per process)

Shadow Page Table



```
set_ptp(guest_pt):
    for gva in 0 .. 220:
        if guest_pt[gva] & PTE_P:
            gpa = guest_pt[gva] >> 12
            hpa = host_pt[gpa] >> 12
            shadow_pt[gva] = (hpa << 12) | PTE_P
        else:
            shadow_pt[gva] = 0
    PTP = shadow_pt
```



What if Guest OS Modifies Its Page Table?

- Real hardware would start using the new page table's mappings
 - Virtual machine monitor has a separate shadow page table
- Goal:
 - VMM needs to intercept when guest OS modifies page table, update shadow page table accordingly
- Technique:
 - Use the read/write bit in the PTE to mark those pages read-only
 - If guest OS tries to modify them, hardware triggers page fault
 - Page fault handled by VMM: update shadow page table & restart guest

Protect Kernel-only Pages



```
set_ptp(guest_pt, kmode):
    for gva in 0 .. 220:
        if guest_pt[gva] & PTE_P and
            (kmode or guest_pt[gva] & PTE_U):
            gpa = guest_pt[gva] >> 12
            hpa = host_pt[gpa] >> 12
            shadow_pt[gva] = (hpa << 12) | PTE_P | PTE_U
        else:
            shadow_pt[gva] = 0
    PTP = shadow_pt
```

Protect Kernel-only Pages

- How do we selectively allow / deny access to kernel-only pages in guest PT?
 - Hardware doesn't know about our virtual U/K bit
- Idea:
 - Generate two shadow page tables, one for U, one for K
 - When guest OS switches to U mode, VMM must invoke `set_ptp(current, 0)`

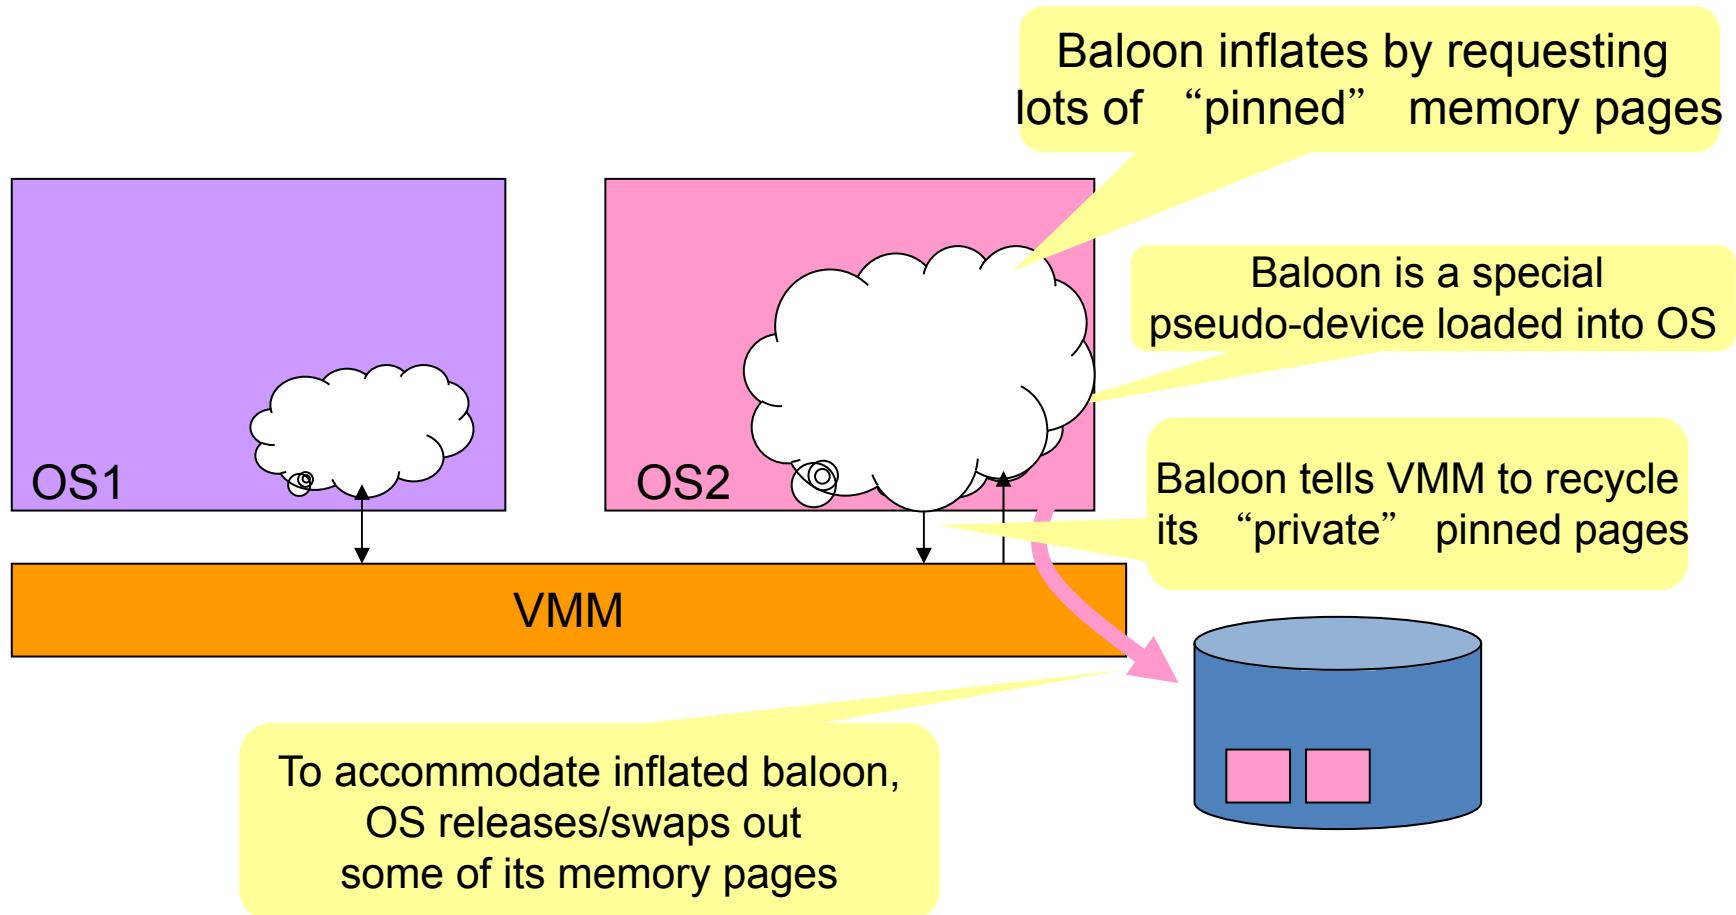
Managing Memory in VMM

- Configure VMs to use more memory than actually available
- What happens when running out of memory?
- Strawman: use LRU paging at VMM
 - OS already uses LRU → doubling paging
 - OS will recycle whatever page VMM just paged out
 - Better to do random eviction

Vmware ESX: Reclaiming Pages

- Idea: trick OS to return memory to VMM
- OS is better at deciding what to swap
 - Normally OS uses all available memory
 - E.g. buffer cache contains old pages, OS won't discard if it does not need memory
- ESX trick: **balloon** driver

Memory Ballooning



ESX: Sharing Pages across VMs

- Many VMs run same OS and programs
 - Many Linux boxes with Apache server
- Idea: use 1 machine page for identical physical pages
- Periodically scan to find identical machine pages
 - Do copy-on-write to eliminate redundancy
- Optimization: use a hash table keyed by hash(content)
 - Allows quick lookup based on page content

System Virtualization

Yubin Xia & Rong Chen
Software School
Shanghai Jiao Tong University

Some Slides adapted from
VMWare's academic course plan

Review: CPU Virtualization

- Instruction Interpretation
- Trap-and-emulate
- Binary Translation
- Para-virtualization

MEMORY VIRTUALIZATION

Review: Virtualizing Memory

- VMM constructs a page table that maps guest address to host physical address
 - E.g., if guest VM has 1GB of memory, it can access memory address 0~1GB
 - Each guest VM has its own mapping for memory address 0, etc.
 - Different host physical address used to store data for those memory locations

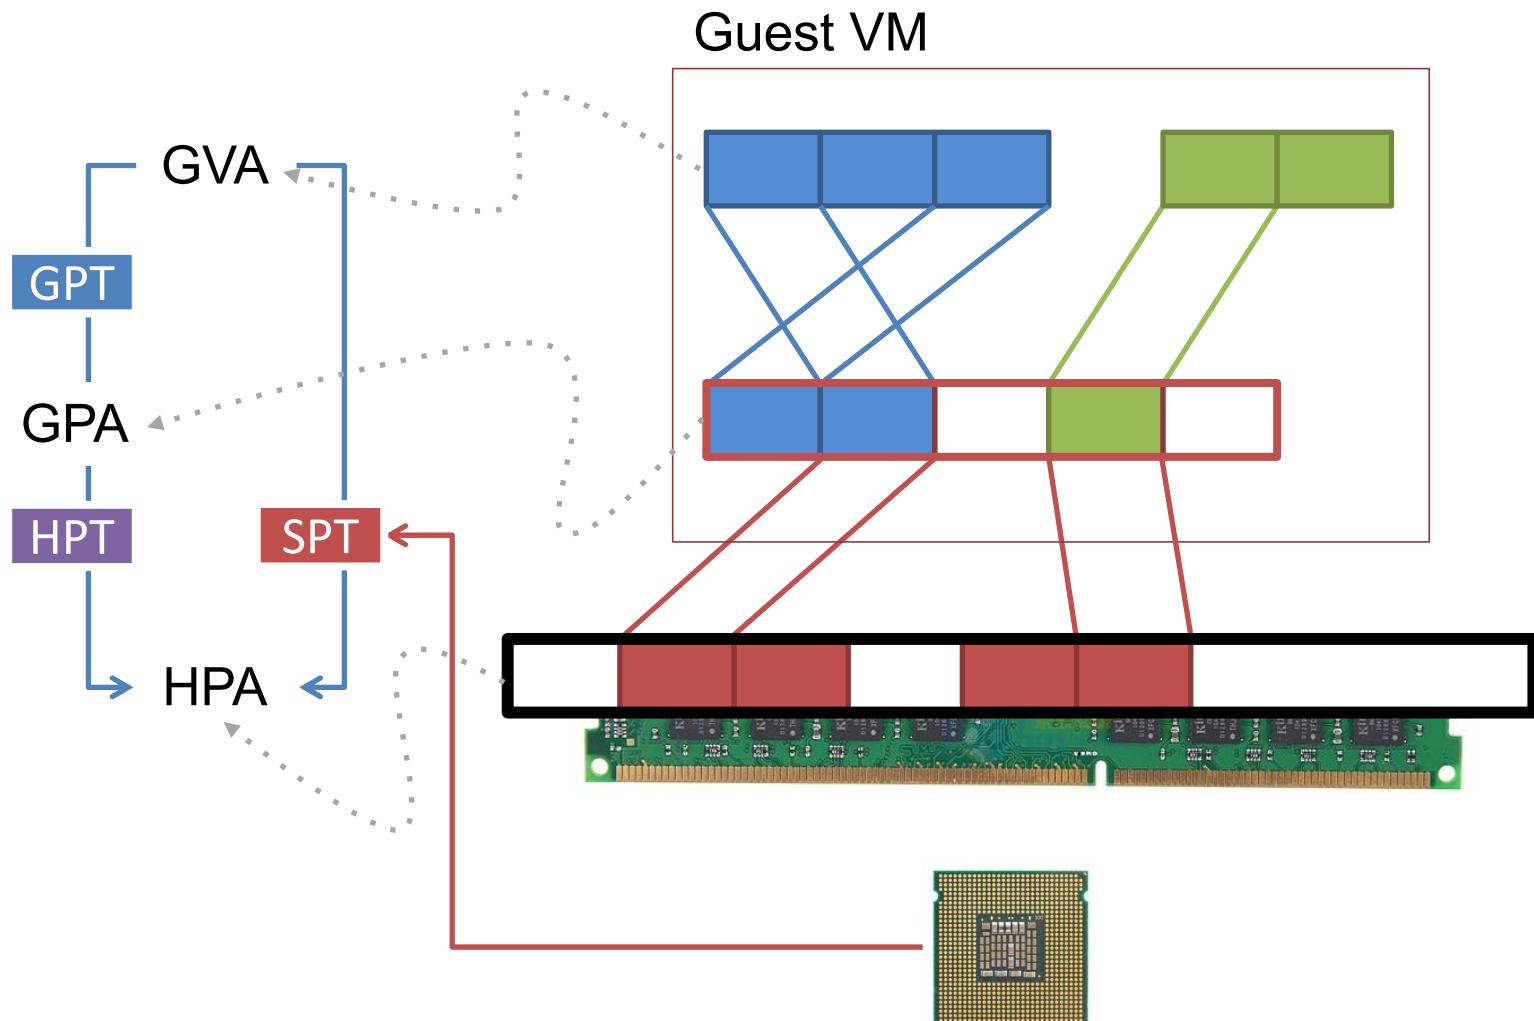
Review: Shadow Pages

- Shadow Paging
 - VMM intercepts guest OS setting the virtual PTP register
 - VMM iterates over the guest page table, constructs a corresponding shadow PT
 - In shadow PT, every guest physical address is translated into host physical address
 - Finally, VMM loads the host physical address of the shadow PT
 - Shadow PT is per process

Shadow Page Table

31	12	11	9	8	7	6	5	4	3	2	1	0
Physical-Page Base Address	AVL	G T	P A	D A	A C D	P W T	P U S	P R /	U /	R W	P	

```
set_ptp(guest_pt):
    for gva in 0 .. 220:
        if guest_pt[gva] & PTE_P:
            gpa = guest_pt[gva] >> 12
            hpa = host_pt[gpa] >> 12
            shadow_pt[gva] = (hpa << 12) | PTE_P
        else:
            shadow_pt[gva] = 0
    PTP = shadow_pt
```



Protect Kernel-only Pages



```
set_ptp(guest_pt, kmode):
    for gva in 0 .. 220:
        if guest_pt[gva] & PTE_P and
            (kmode or guest_pt[gva] & PTE_U):
            gpa = guest_pt[gva] >> 12
            hpa = host_pt[gpa] >> 12
            shadow_pt[gva] = (hpa << 12) | PTE_P | PTE_U
        else:
            shadow_pt[gva] = 0
    PTP = shadow_pt
```

Protect Kernel-only Pages

- How do we selectively allow / deny access to kernel-only pages in guest PT?
 - Hardware doesn't know about our virtual U/K bit
- Idea:
 - Generate two shadow page tables, one for U, one for K
 - When guest OS switches to U mode, VMM must invoke `set_ptp(current, 0)`

Review: Hardware Support for Memory Virtualization

- Hardware Support for Memory Virtualization
 - Intel's EPT (Extended Page Table)
 - AMD's NPT (Nested Page Table)
- Another Table
 - EPT for translation from GPA to HPA
 - EPT is controlled by the hypervisor
 - EPT is per-VM

Hardware-assisted Virtualization

VT-x Modes

VMX root operation:

Full privileged, intended for Virtual Machine Monitor

VMX non-root operation:

Not fully privileged, intended for guest software

Both forms of operation support all four privilege levels from 0 to 3

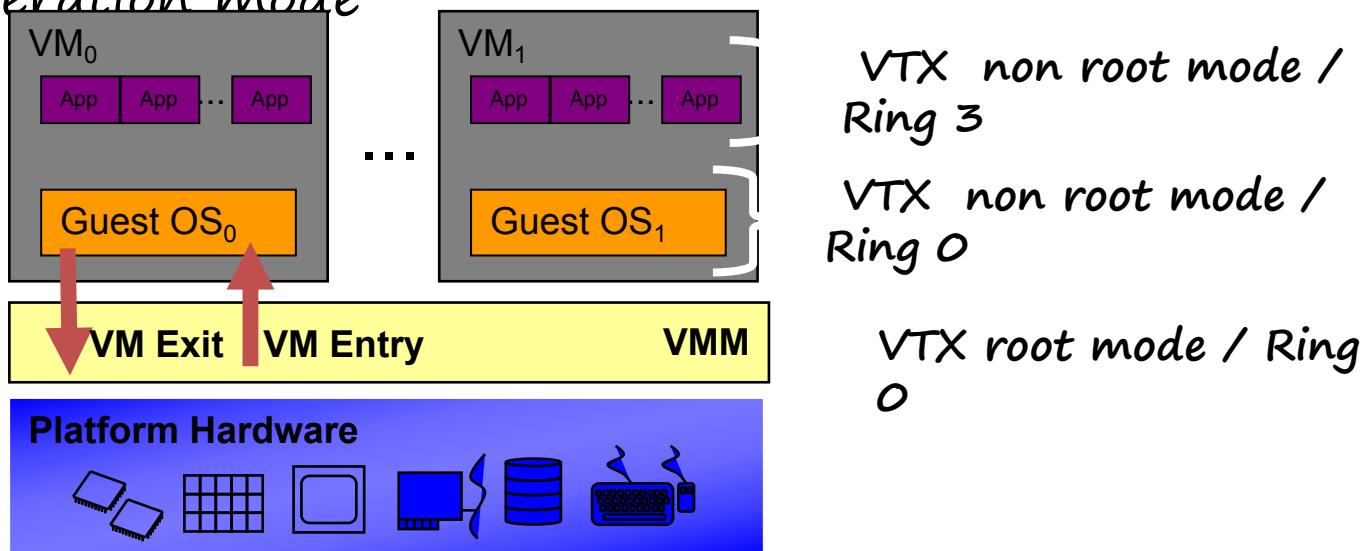
VT-x transition mechanisms

VM exit

From VMX non-root operation mode to VMX root operation mode

VM entry

from VMX root operation mode to VMX non-root operation mode



Virtual Machine Control Structure (VMCS)

Data structure that manages VM entries and VM exits.

VMCS is logically divided into:

- Guest-state area

- Host-state area.

- VM-execution control fields

- VM-exit control fields

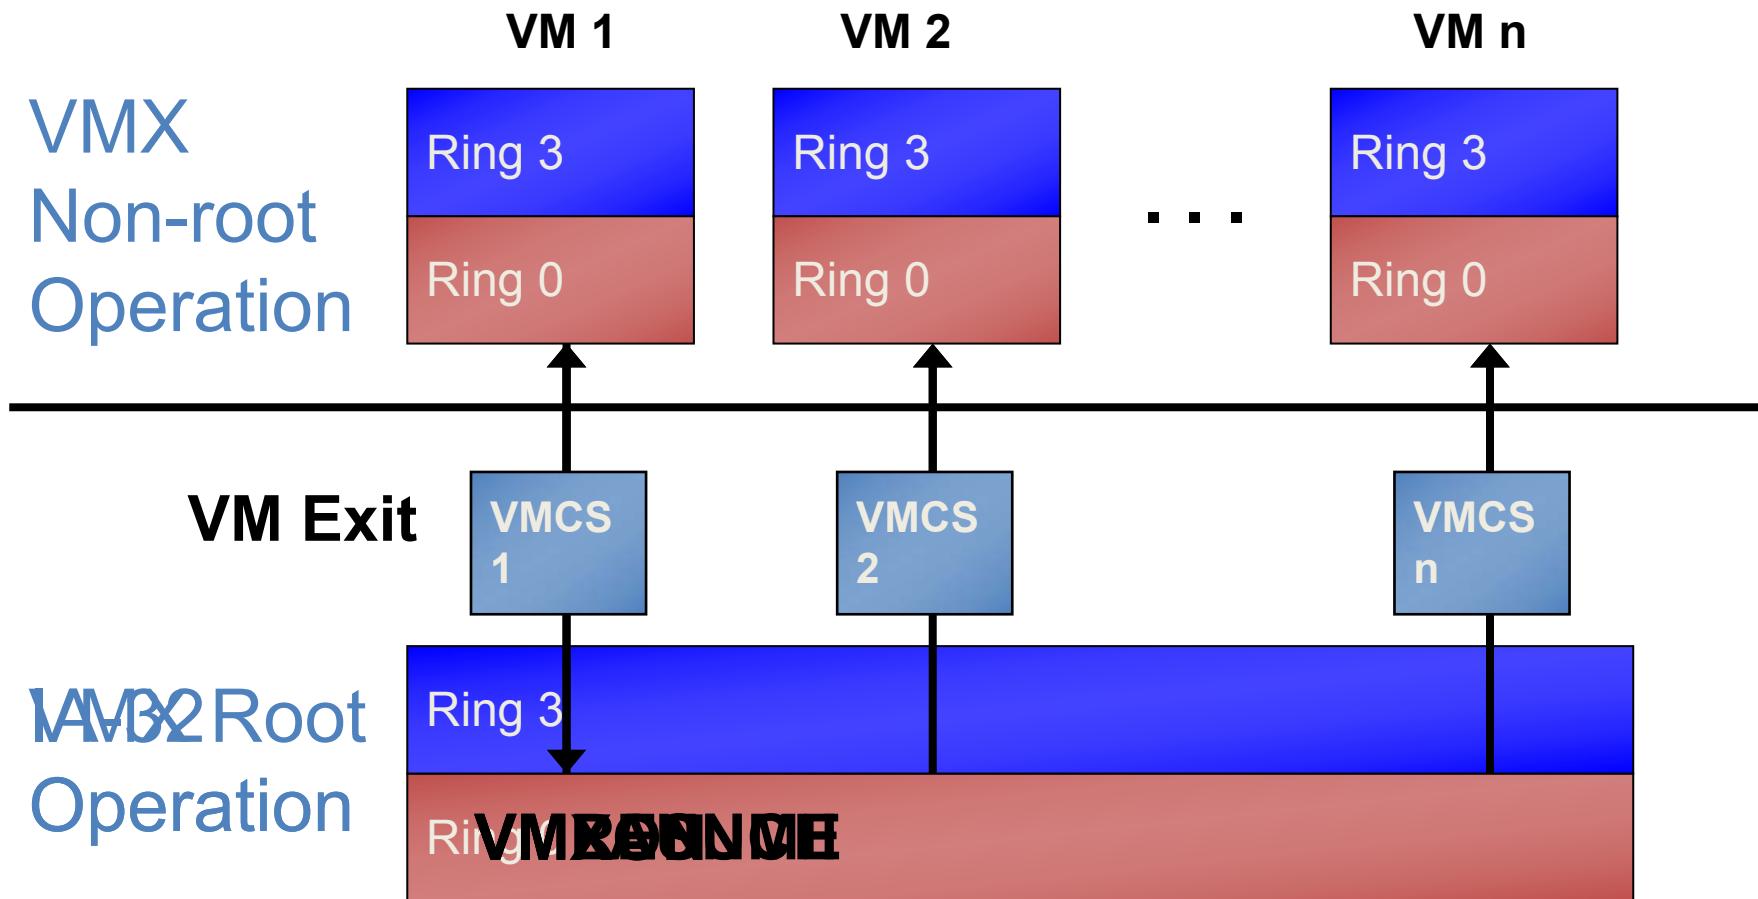
- VM-entry control fields

- VM-exit information fields

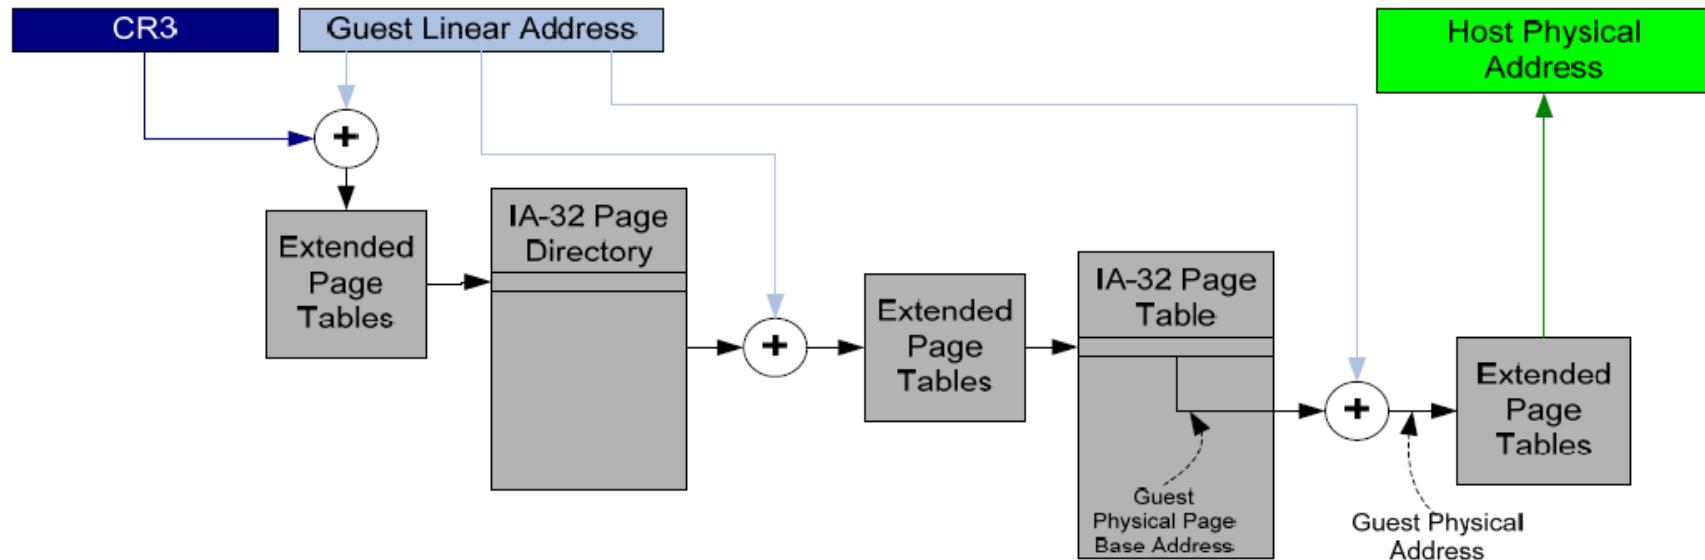
VM entries load processor state from the guest-state area.

VM exits save processor state to the guest-state area and the exit reason, and then load processor state from the host-state area.

VT-x Operations



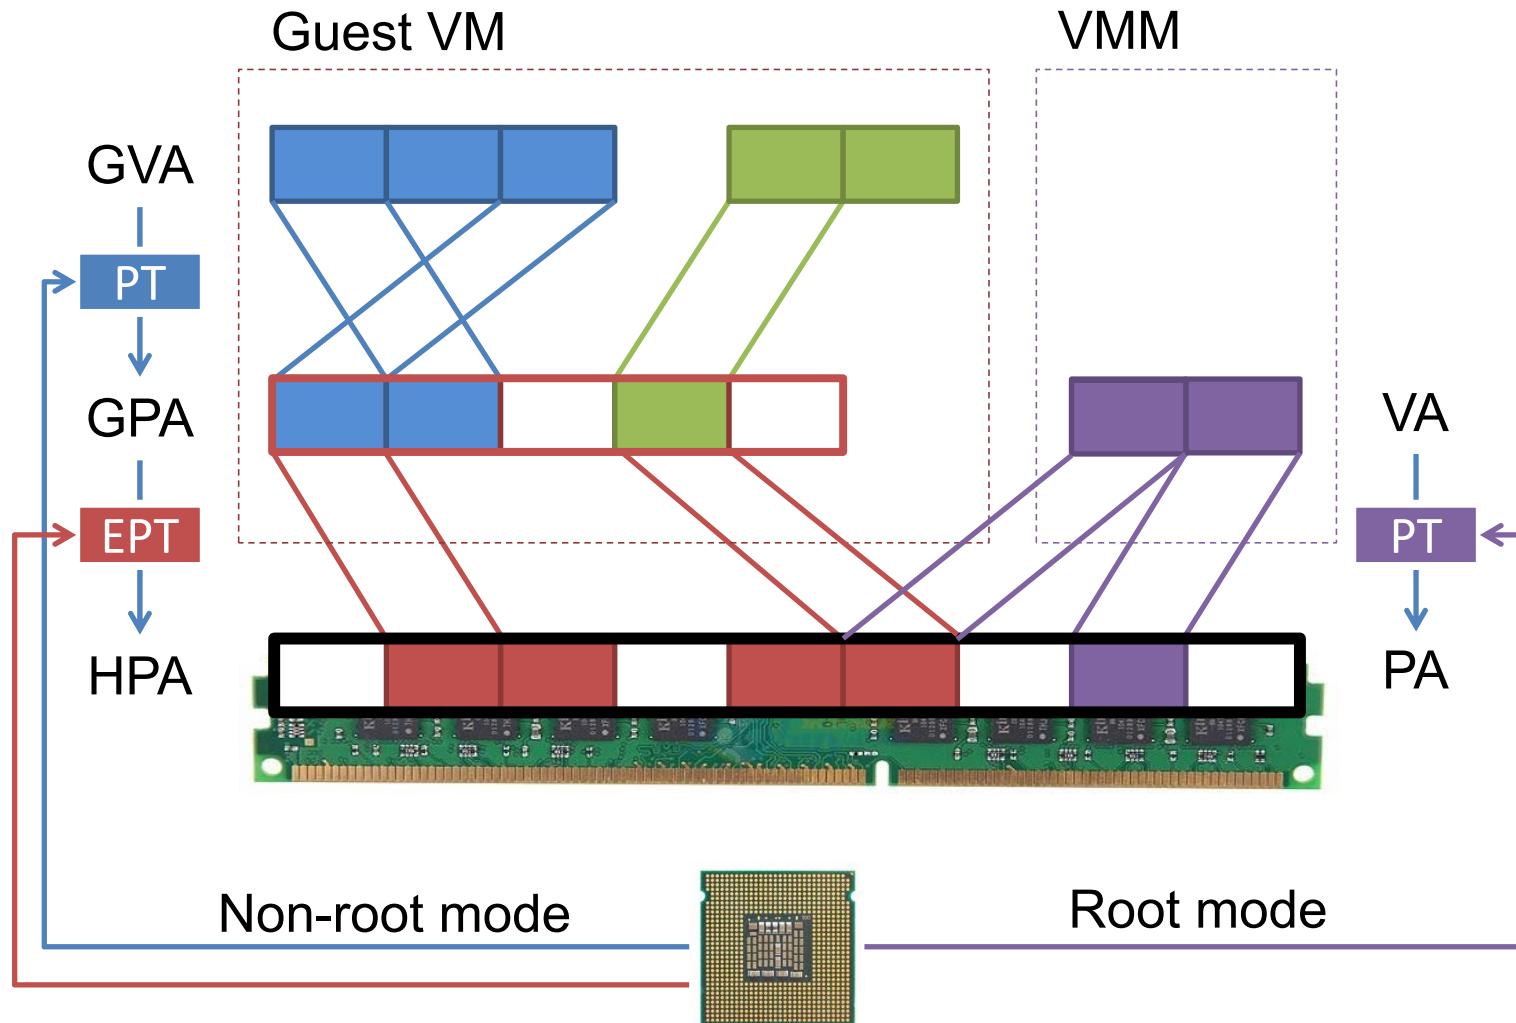
VT-x extension: Extended Page Table (EPT)



All guest-physical addresses go through extended page tables
Includes address in CR3, address in PDE, address in PTE, etc.

Reduces the frequency of VM exits to VMM.

The net effect of both implementations (EPT or NPT) is to allow the guest OS to own and manage its own page table, and not force the host to get involved.



VT-x New instructions

VMXON and VMXOFF

To enter and exit VMX-root mode.

VMLAUNCH: Used on initial transition from VMM to Guest

Enters VMX non-root operation mode

VMRESUME: Used on subsequent entries

Enters VMX non-root operation mode

Loads Guest state and Exit criteria from VMCS

VMEXIT

Used on transition from Guest to VMM

Enters VMX root operation mode

Saves Guest state in VMCS

Loads VMM state from VMCS

VMPTRST and VMPTRLD

To Read and Write the VMCS pointer.

VMREAD, VMWRITE, VMCLEAR

Read from, Write to and clear a VMCS.

MANAGE VM MEMORY

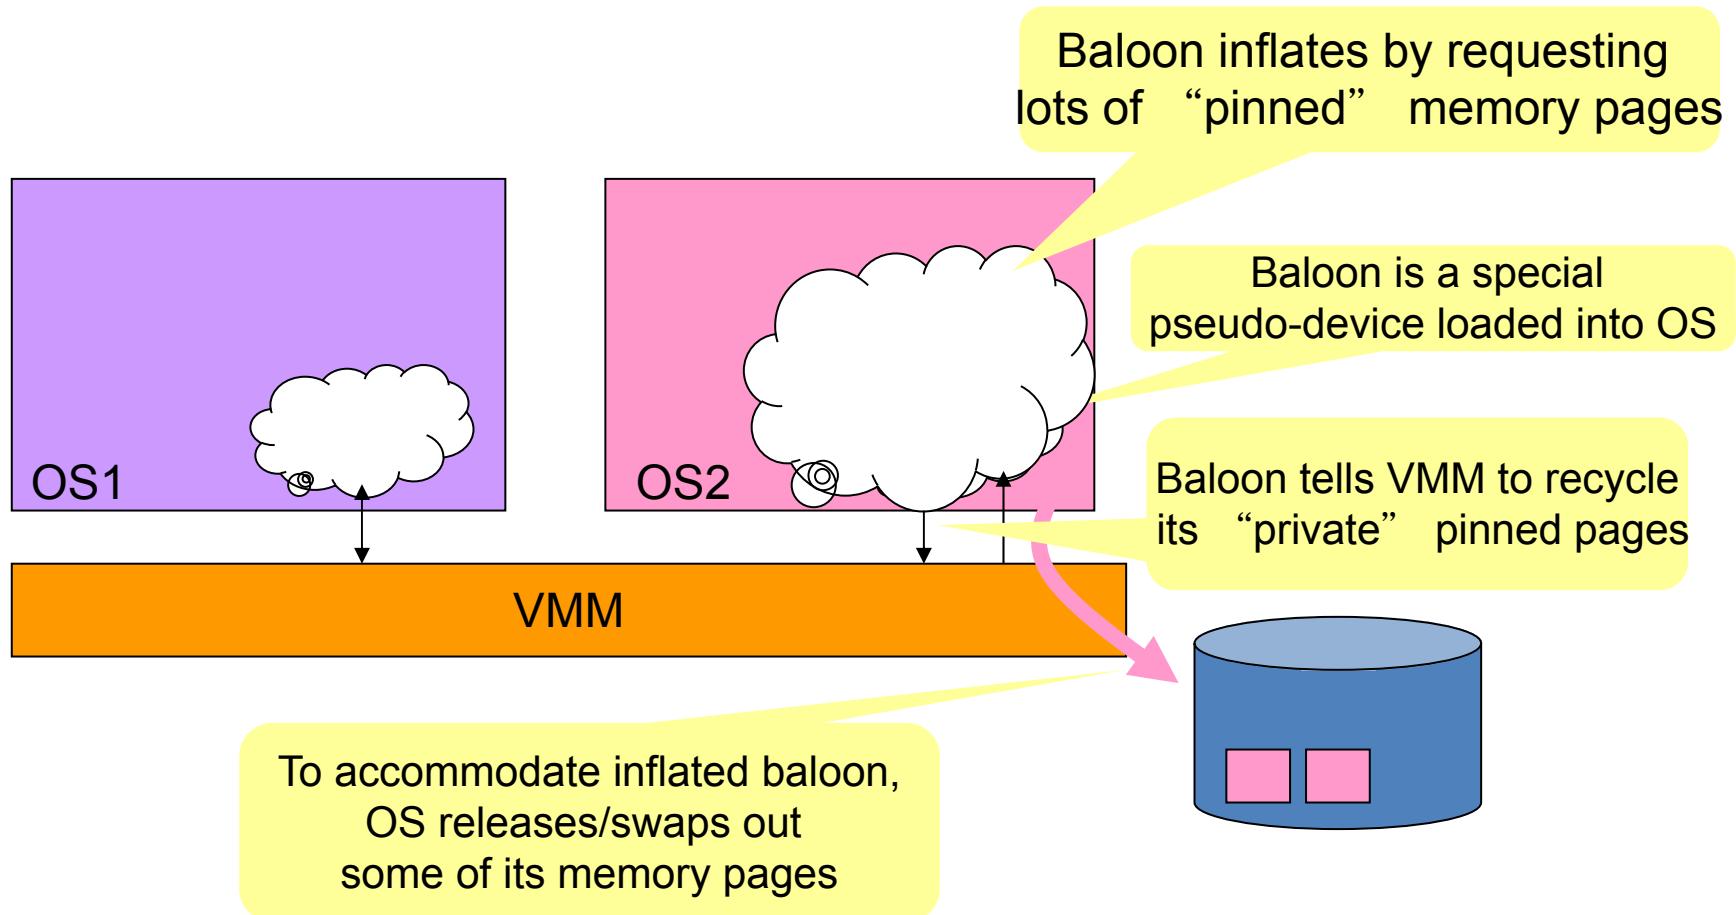
Managing Memory in VMM

- Configure VMs to use more memory than actually available
- What happens when running out of memory?
- Strawman: use LRU paging at VMM
 - OS already uses LRU → doubling paging
 - OS will recycle whatever page VMM just paged out
 - Better to do random eviction

Vmware ESX: Reclaiming Pages

- Idea: trick OS to return memory to VMM
- OS is better at deciding what to swap
 - Normally OS uses all available memory
 - E.g. buffer cache contains old pages, OS won't discard if it does not need memory
- ESX trick: **balloon** driver

Memory Ballooning



ESX: Sharing Pages across VMs

- Many VMs run same OS and programs
 - Many Linux boxes with Apache server
- Idea: use 1 machine page for identical physical pages
- Periodically scan to find identical machine pages
 - Do copy-on-write to eliminate redundancy
- Optimization: use a hash table keyed by hash(content)
 - Allows quick lookup based on page content

I/O VIRTUALIZATION

Outline

Types of Device Virtualization

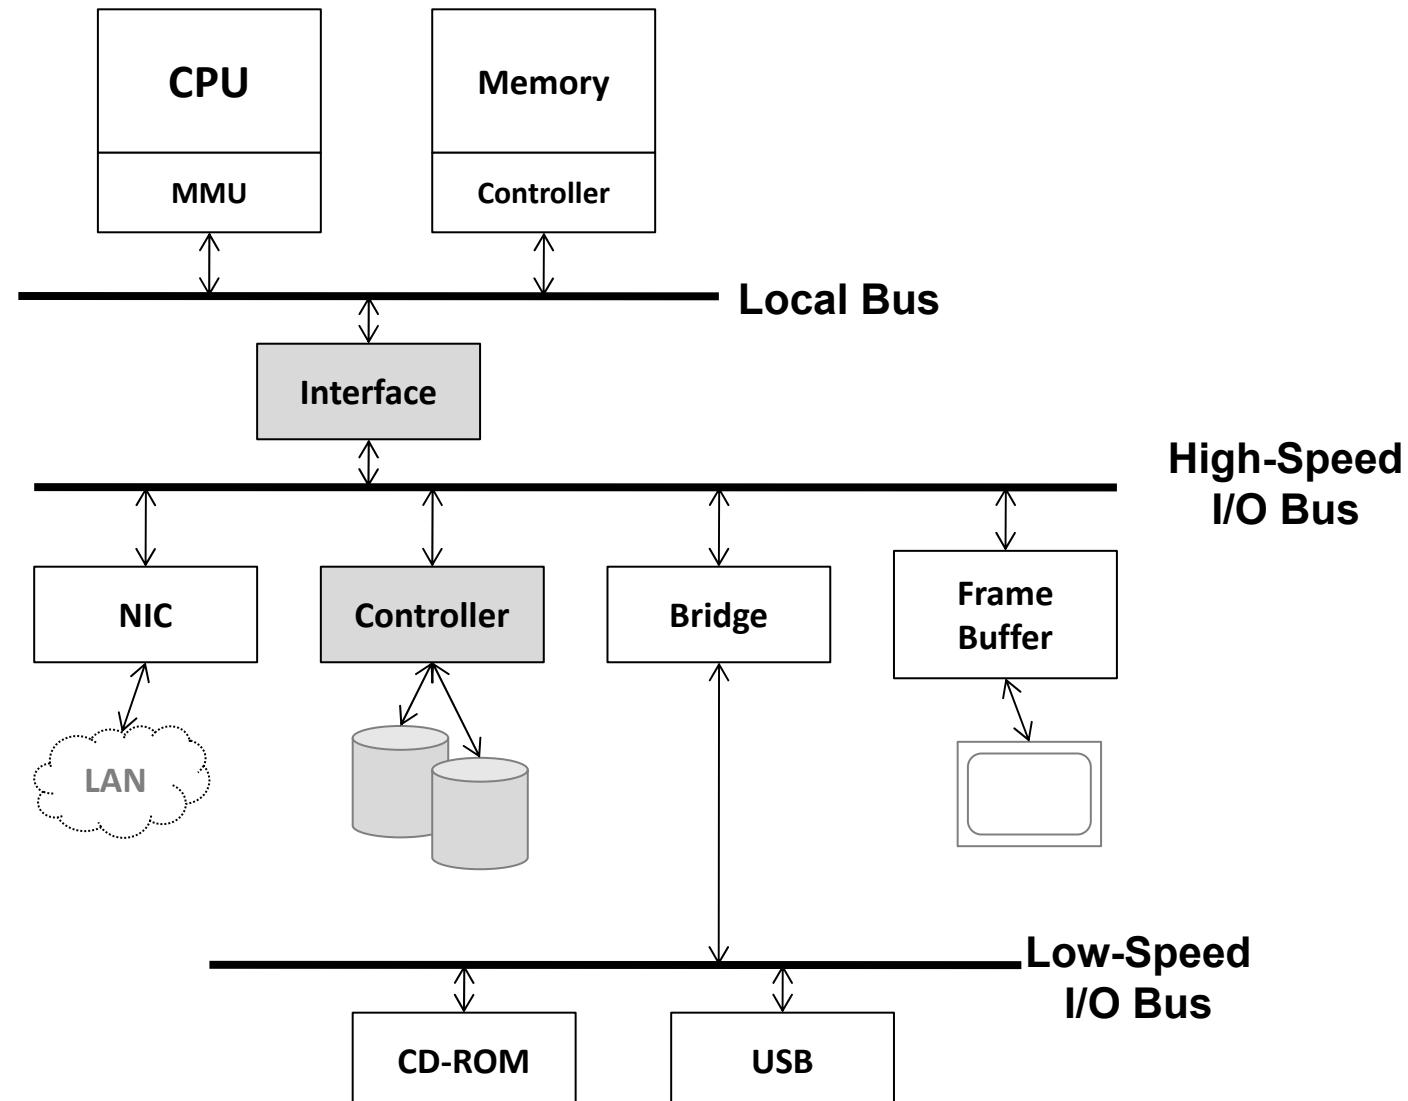
Direct Access

Emulated

Para-virtualized

Hardware assisted I/O virtualization

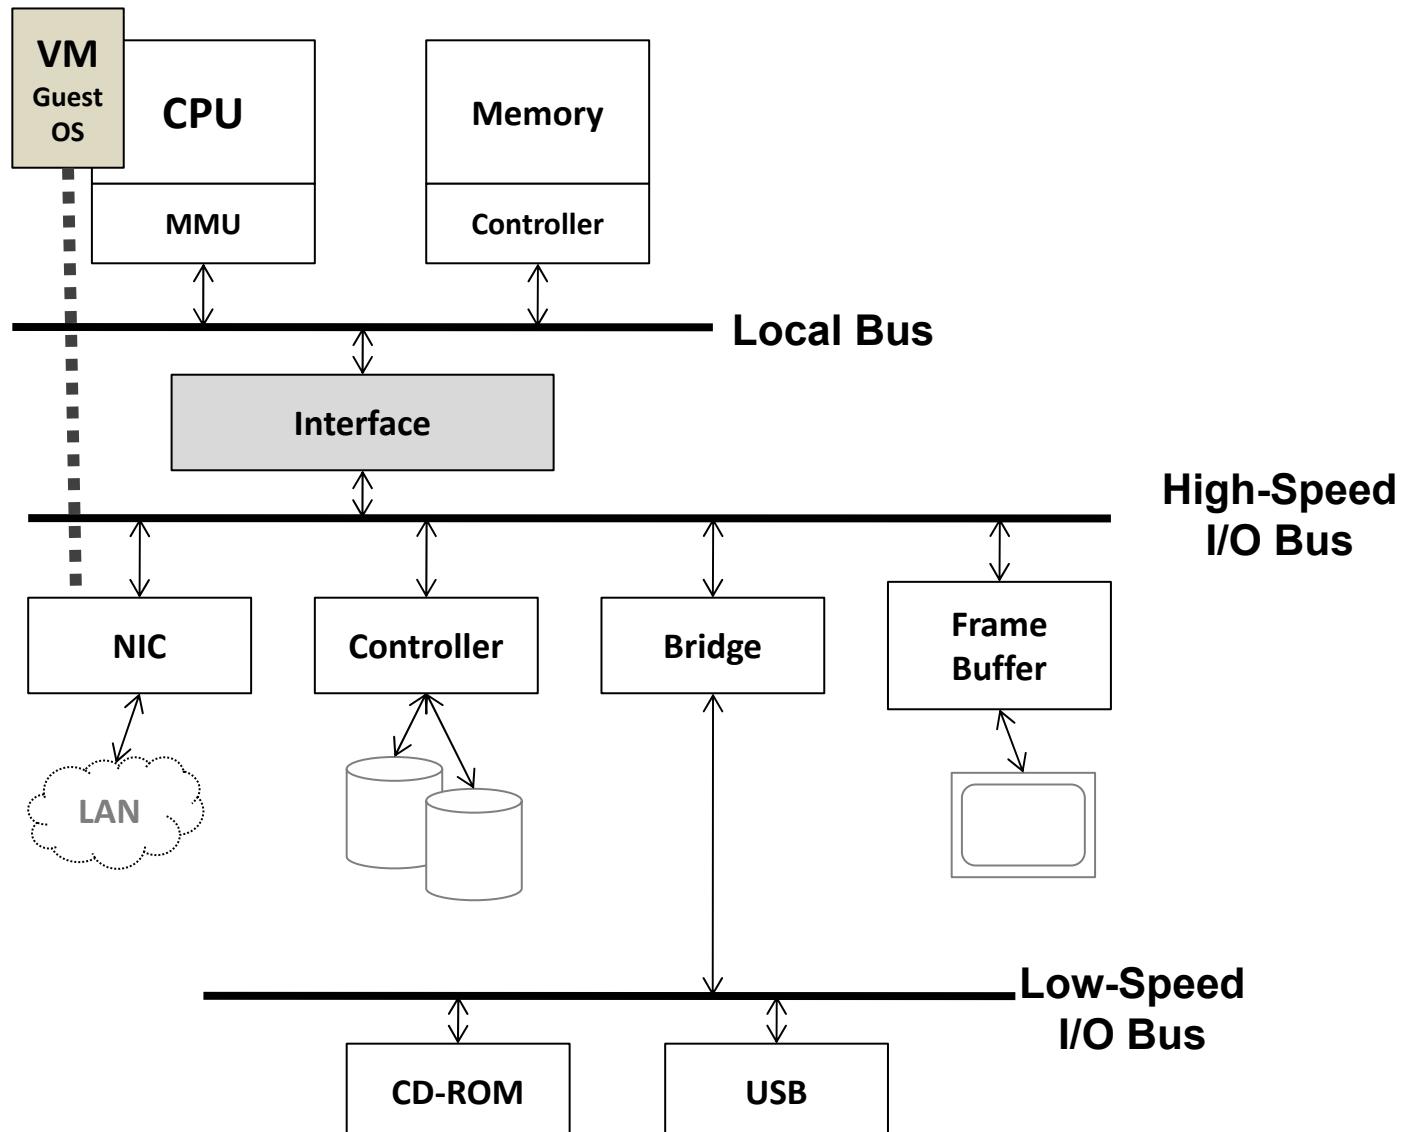
Computer System Organization



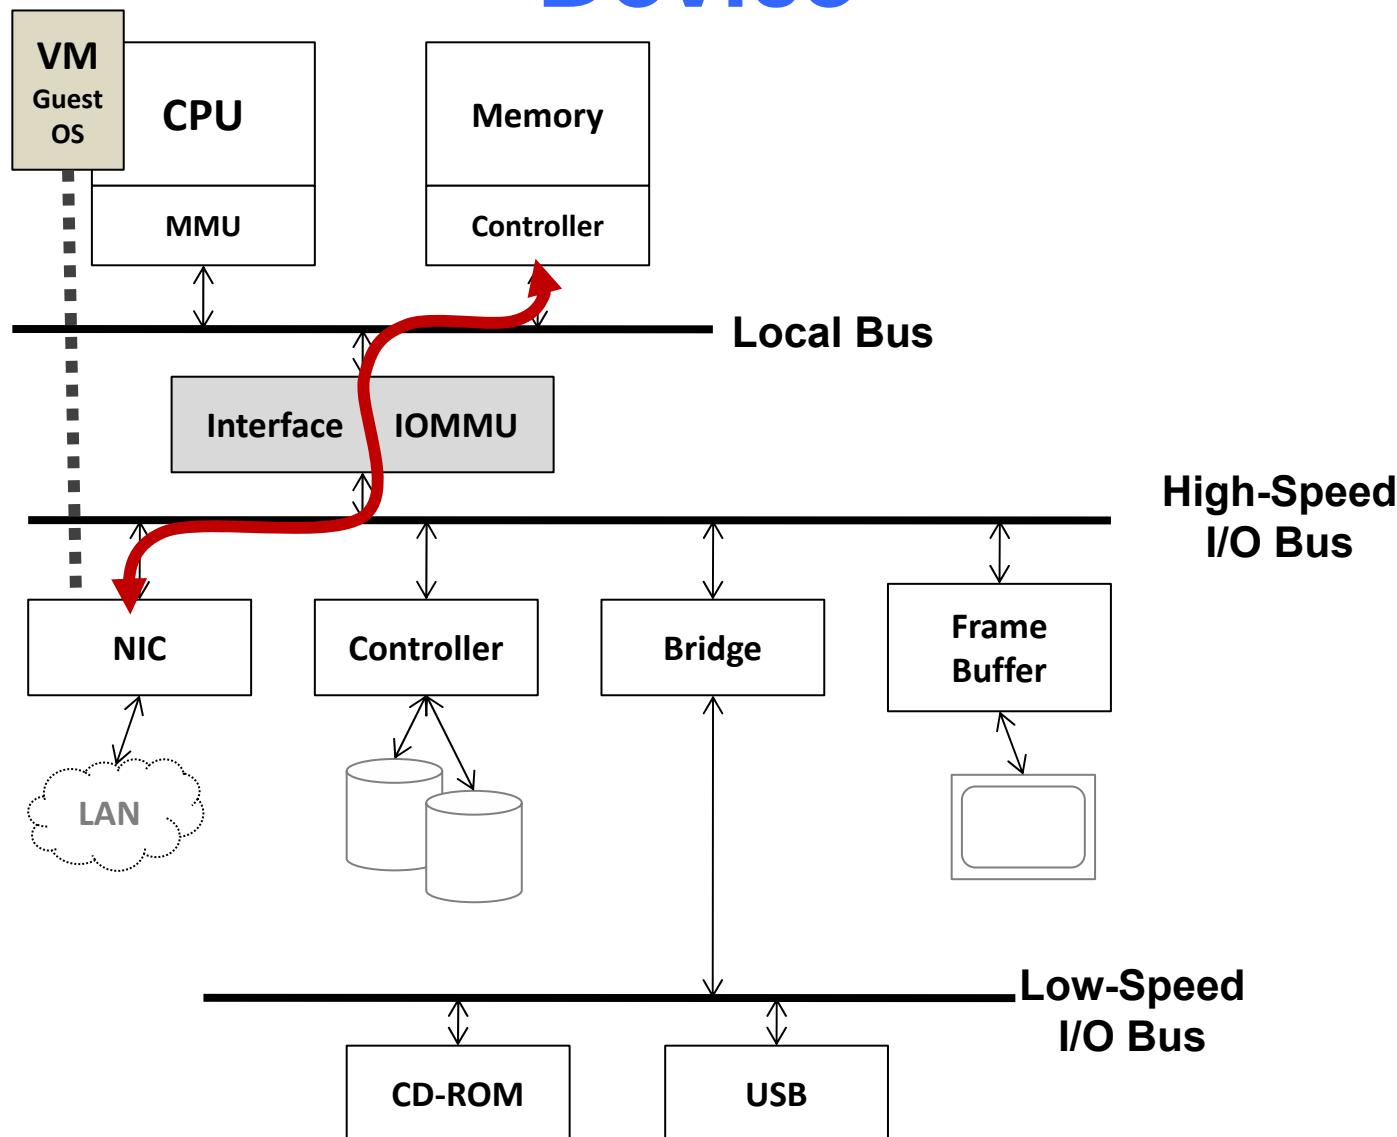
Device Virtualization

- Goals
 - Isolation
 - Multiplexing
 - Speed
 - Mobility
 - Interposition
- Device Virtualization Strategies
 - Direct Access
 - Emulation
 - Para-virtualization

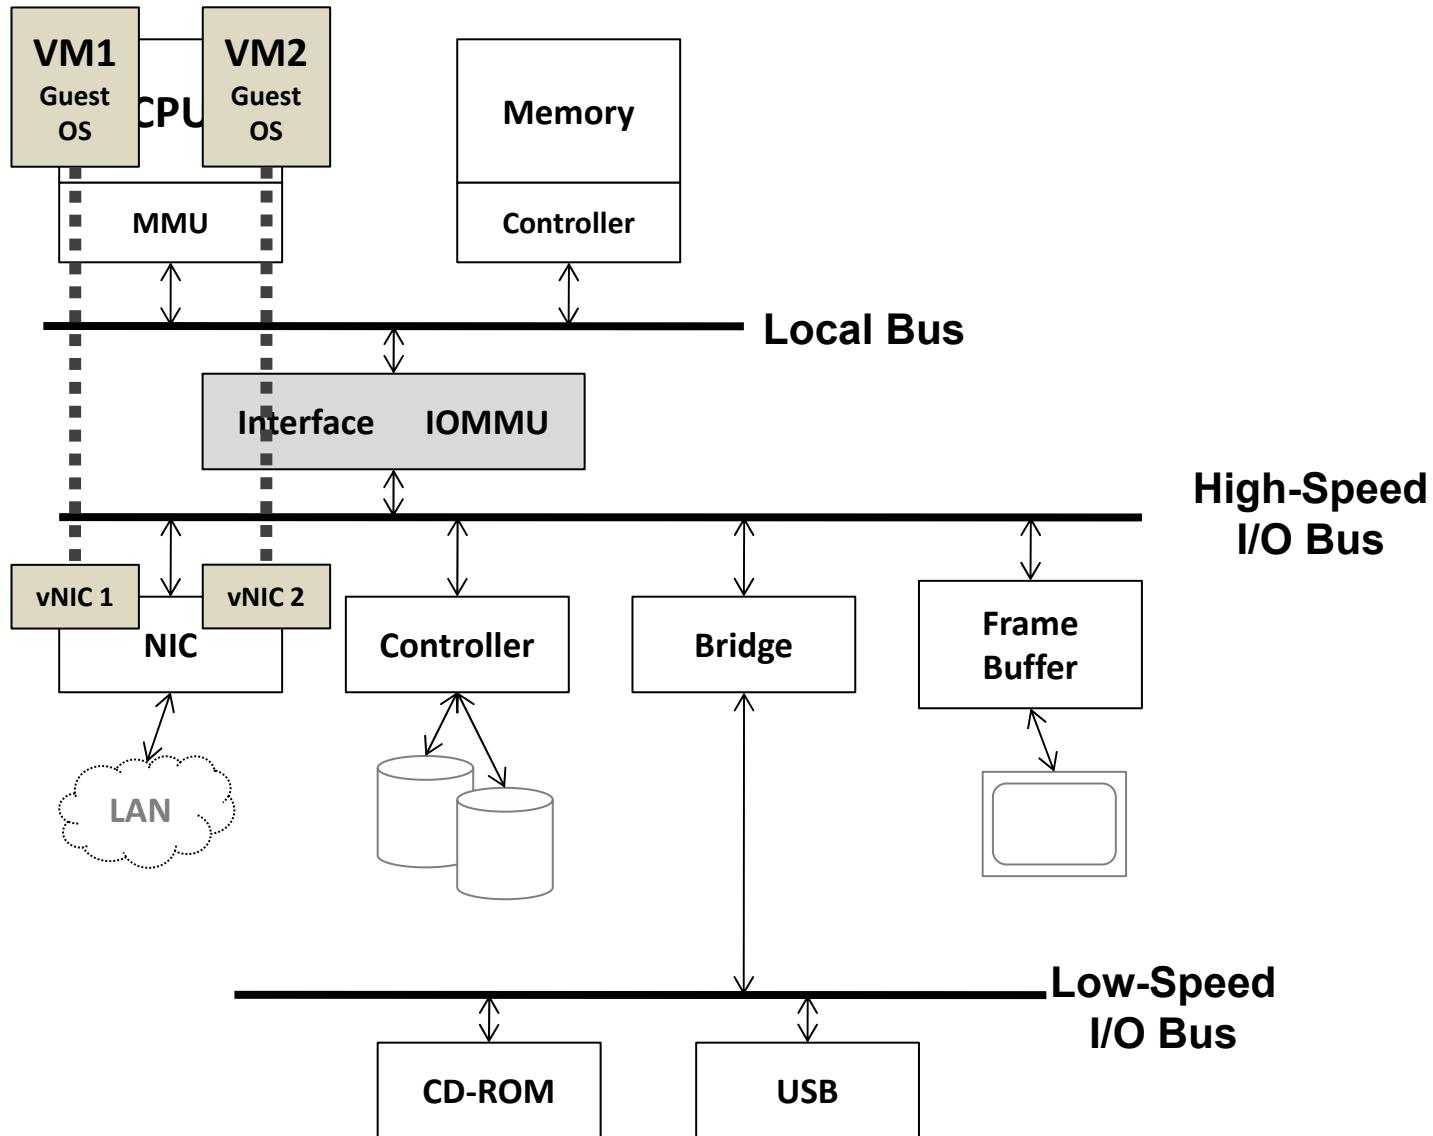
Direct Access Device



Memory Isolation w/ Direct Access Device



Virtualization Enabled Device



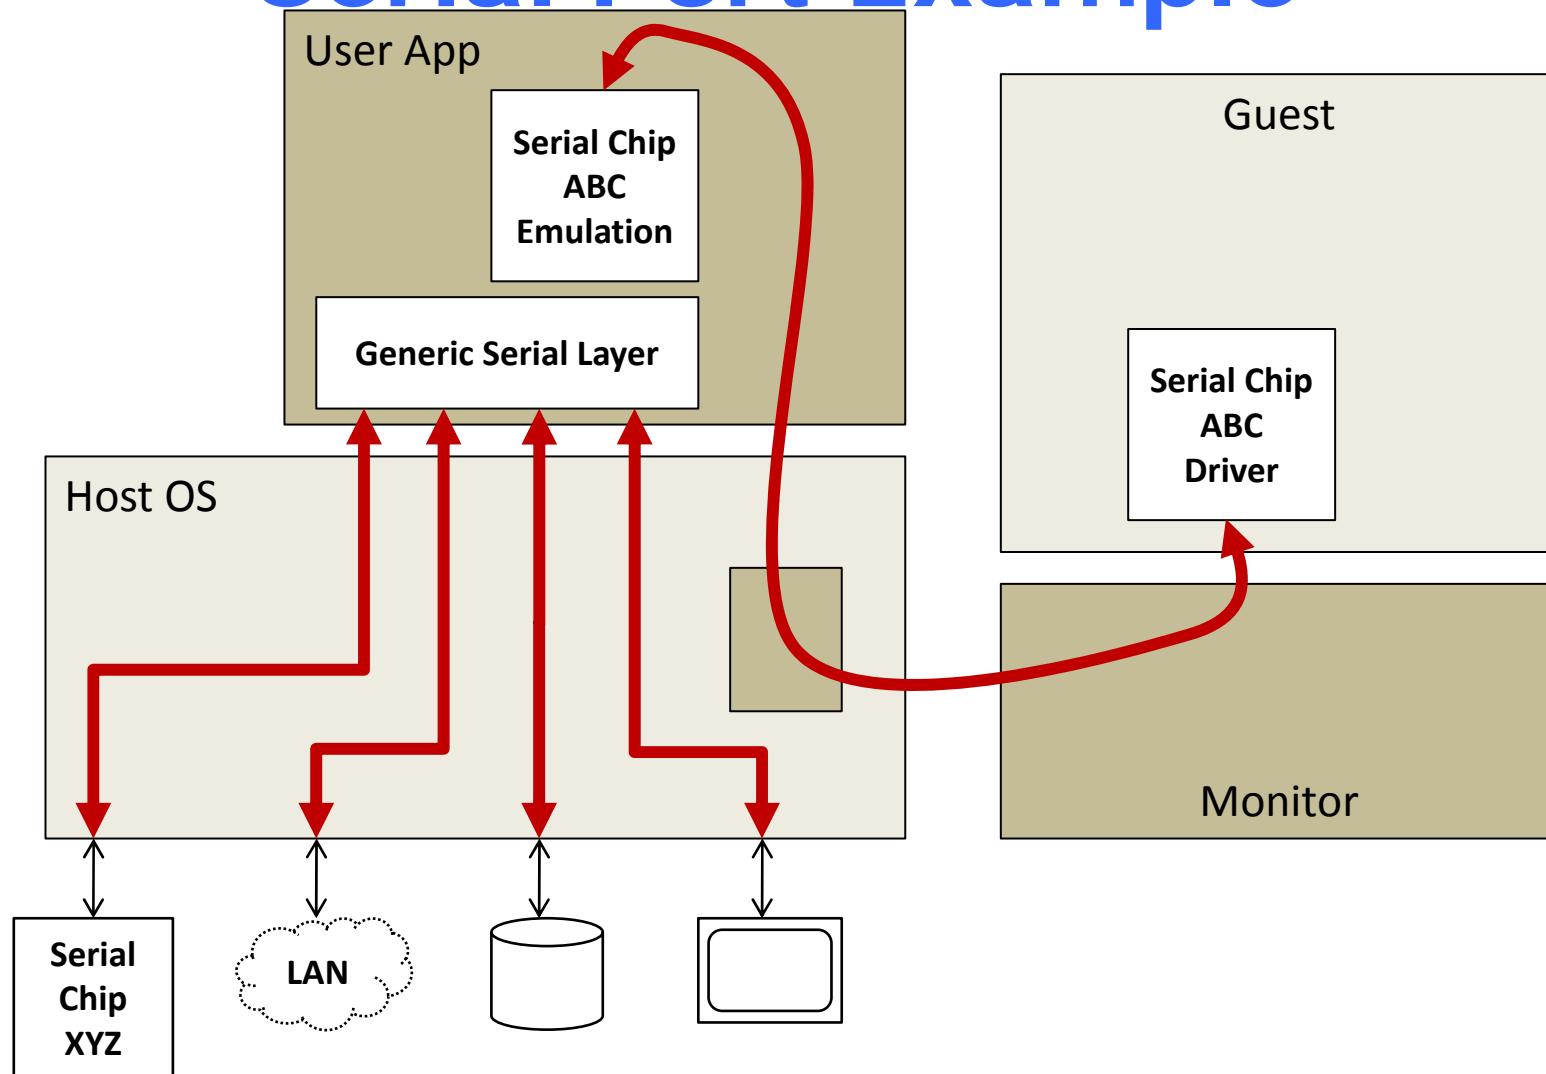
Direct Access Device Virtualization

- Allow Guest OS direct access to underlying device
- Positives
 - Fast
 - Simplify monitor
 - Limited device drivers needed
- Negatives
 - Need hardware support for safety (IOMMU)
 - Need hardware support for multiplexing
 - Hardware interface visible to guest
 - Limits mobility of VM
 - Interposition hard by definition

Emulated Devices

- Emulate a device in class
 - Emulated registers
 - Memory mapped I/O or programmed I/O
- Convert
 - Intermediate representation
- Back-ends per real device

Serial Port Example



Emulated Devices

Positives

Platform stability

Allows interposition

No special hardware support needed

Isolation, multiplexing implemented by monitor

Negatives

Can be slow

Drivers needed in monitor or host

Para-Virtualized Devices

Guest passes requests to Monitor at a higher abstraction level

- Monitor calls made to initiate requests

- Buffers shared between guest / monitor

Positives

- Simplify monitor

- Fast

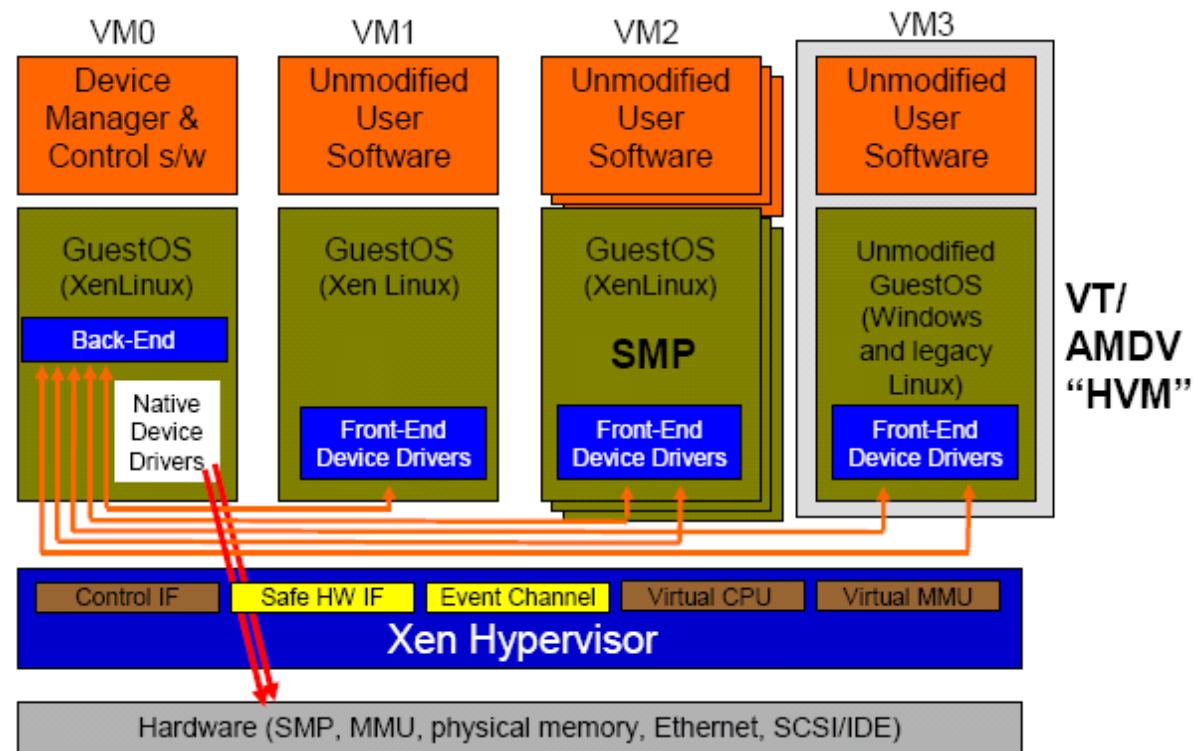
Negatives

- Monitor needs to supply guest-specific drivers

- Bootstrapping issues

Para-virtualized Devices: Front-end/Back-end

- Para
–virtualization
 - Linux Guest
- Hardware-
supported
virtualization
 - Unmodified
Windows
- Isolated Device
Drivers



*Source: Ottawa Linux Symposium 2006 presentation.
<http://www.cl.cam.ac.uk/netos/papers/>*

VirtIO: Unified Para-virtualized I/O

Motivation: Linux has support for at least 8 virtualization platforms, each with its own para-virtualized I/O design/interfaces

VirtIO: to provide a unified I/O mode for para-virtualized device

Adopted by KVM and Iguest

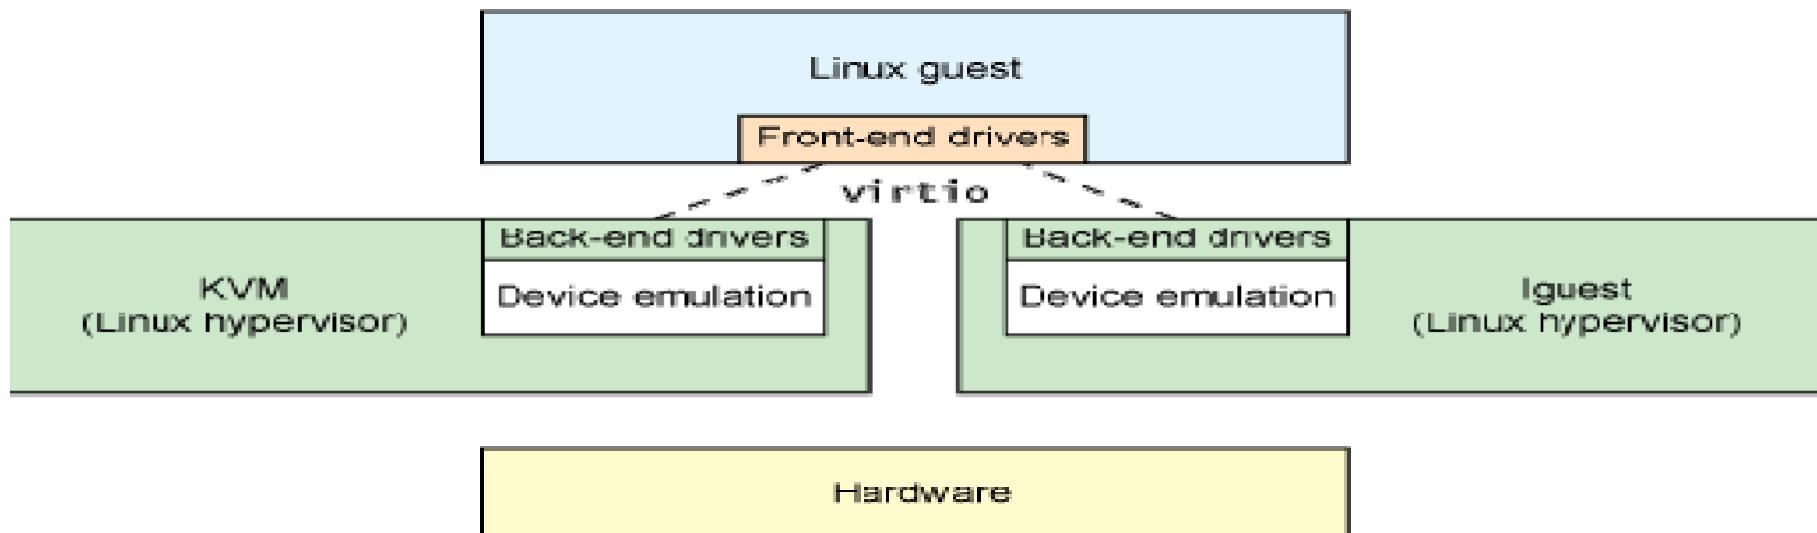
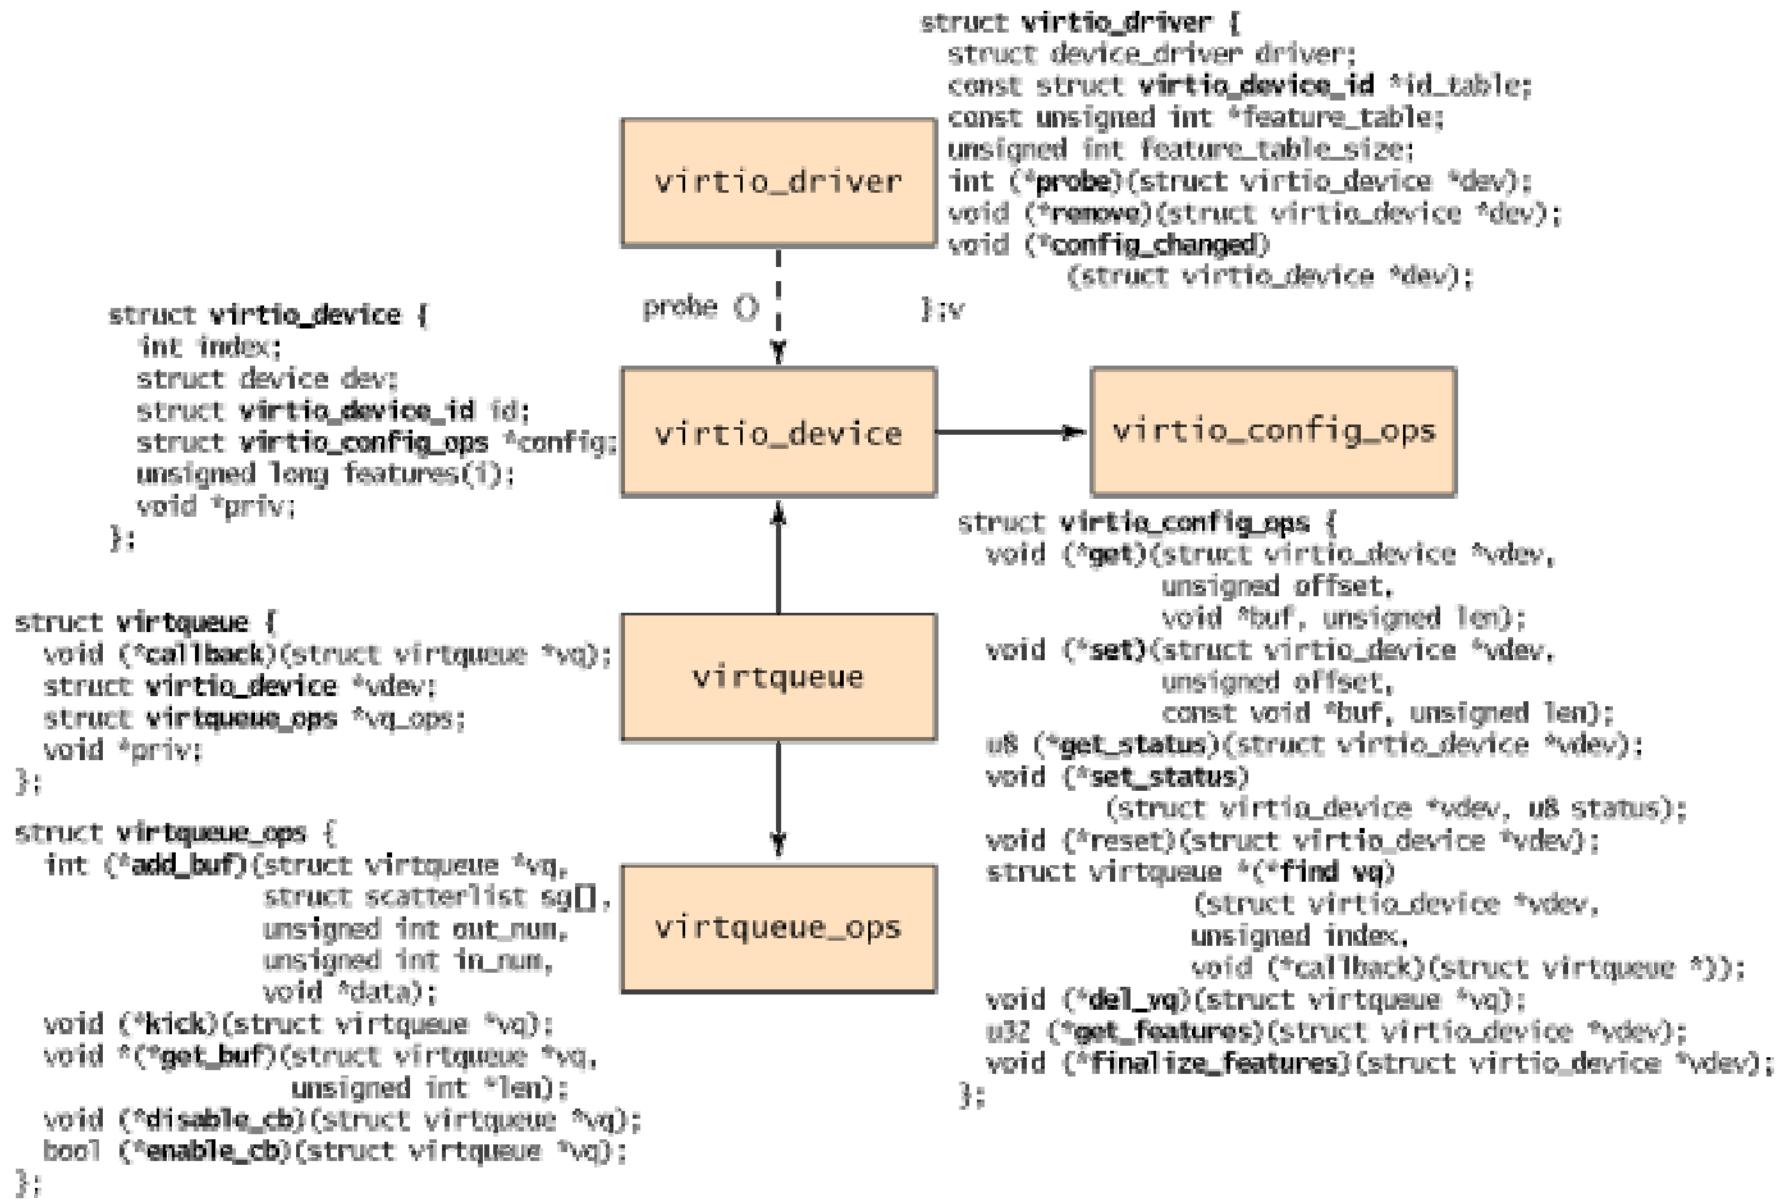


Figure from: <http://www.ibm.com/developerworks/cn/linux/l-virtio/figure2.gif>

VirtIO



Hardware-assisted I/O Virtualization

Motivation

- Emulated devices
 - Slow performance
 - High implementation complexity
- Para-virtualized I/O virtualization
 - Needs modification to guest OS, e.g., Linux, Windows
 - May still have suboptimal performance
- Hardware-assisted I/O virtualization
 - Provide efficient hardware support for I/O virtualization
 - E.g., Intel VT-d

Issues to Address

I/O address translation

How to translate I/O address to host physical address

Interrupt mapping

How to route an interrupt correctly to a guest VM

Device multiplexing

How to multiplex a single hardware device among multiple VMs

Mostly importantly

Provide strong isolation, while reduce hypervisor involvement

VT-d: Intel® Virtualization Technology for Directed I/O

Provides the capability to ensure improved isolation of I/O resources for greater reliability, security, and availability.

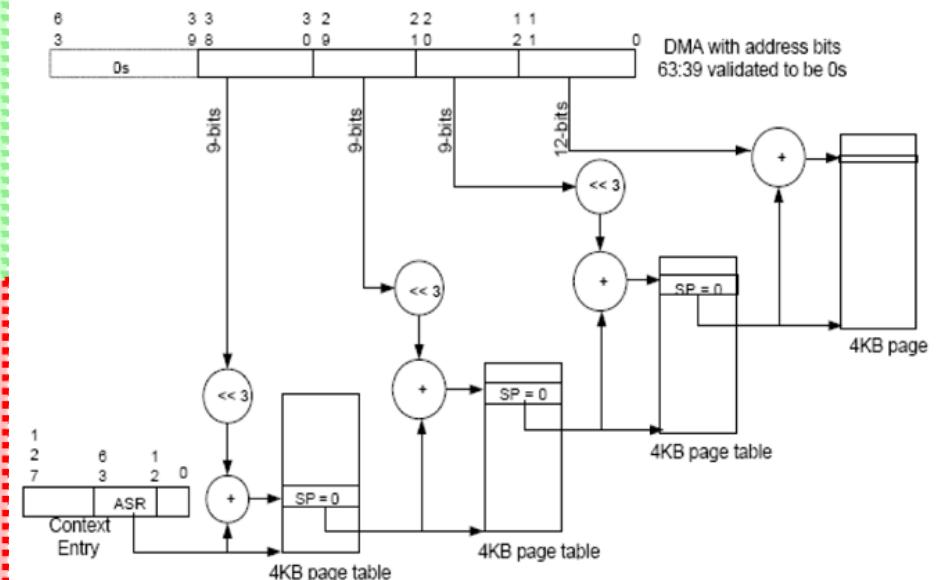
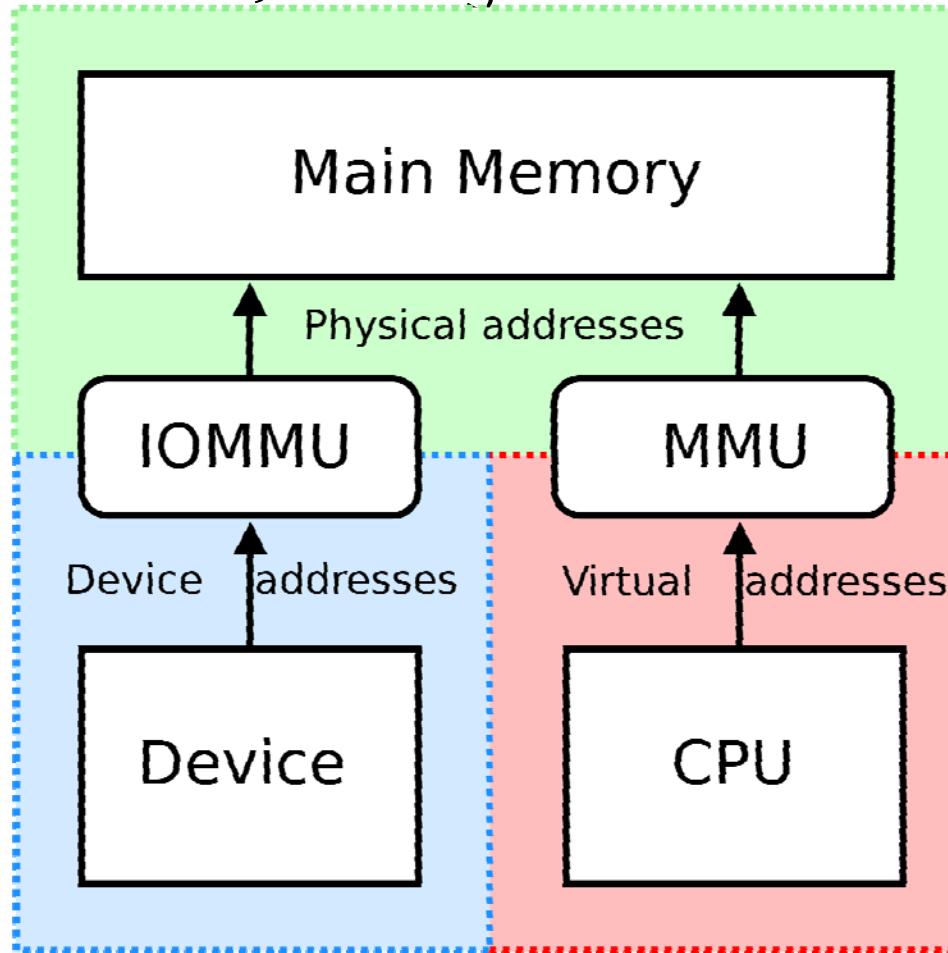
Supports the remapping of I/O DMA transfers and device-generated interrupts.

Provides flexibility to support multiple usage models that may run un-modified, special-purpose, or "virtualization aware" guest OSs.

Address Translation Services

VT-d architecture defines a multi-level page-table structure for DMA address translation

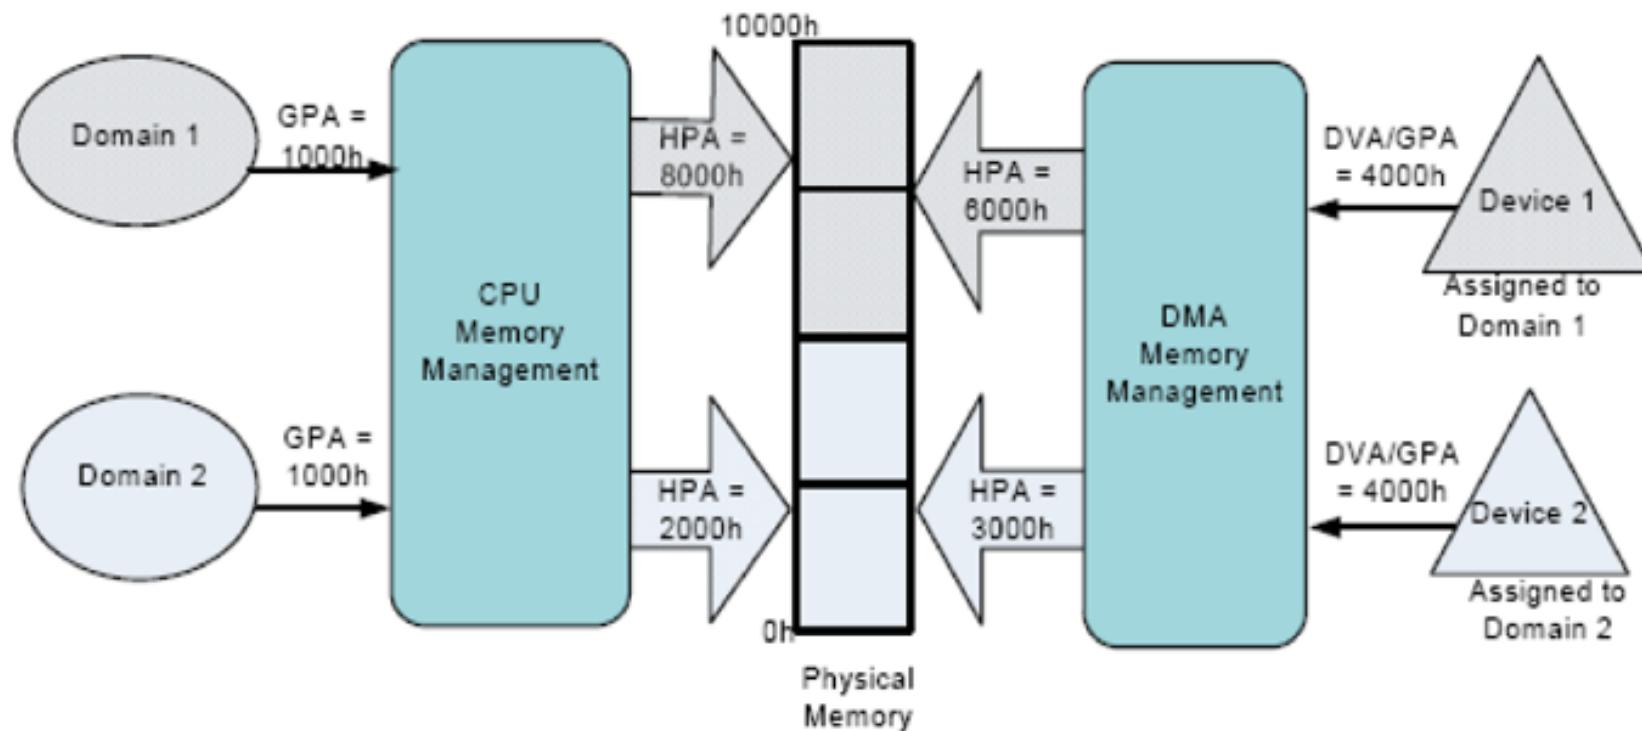
The multi-level page tables are similar to IA-32 processor page-tables, enabling software to manage memory at 4 KB or larger



DMA Remapping

DMA-remapping translates the address of the incoming DMA request to the correct physical memory address and perform checks for permissions to access that physical address

DMA-remapping hardware logic in the chipset sits between the DMA capable peripheral I/O devices and the computer's physical memory



VI - a feature: interrupt Remapping

The interrupt requests generated by I/O devices must be controlled by the VMM

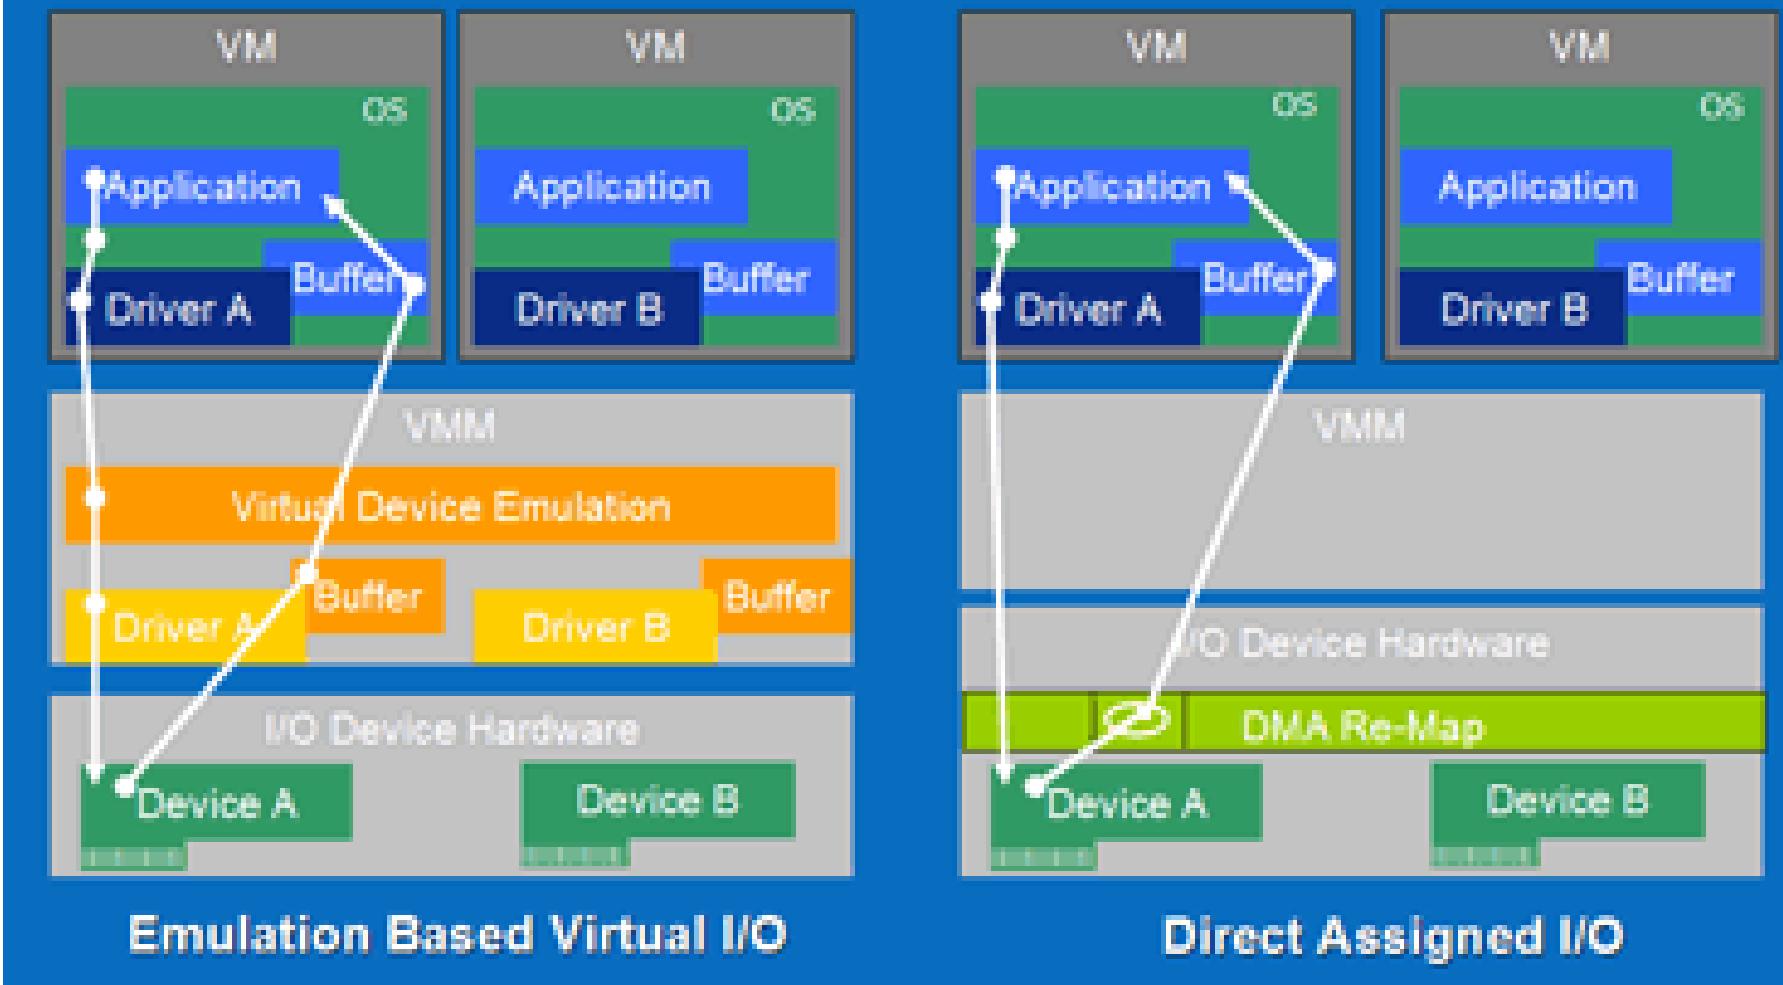
When the interrupt occurs, the VMM must present the interrupt to the guest. This is not accomplished through hardware.

The VT-d interrupt-remapping architecture addresses this problem by redefining the interrupt-message format.

Interrupt requests specify a requester-ID and interrupt-ID, and remap hardware transforming these requests to a physical

Interrupt Remapping

Intel® Virtualization Technology (Intel® VT) for
Directed I/O (Intel® VT-d) – Direct assignment



Device multiplexing: Single Root I/O Virtualization

- Single Root I/O Virtualization (SR-IOV) is a Peripheral Component Interconnect Special Interest Group (PCI-SIG) specification.
- SR-IOV provides a standard mechanism for devices to advertise their ability to be simultaneously shared among multiple virtual machines.
- SR-IOV allows for the partitioning of a PCI function into many virtual interfaces for the purpose of sharing the resources of a PCI Express* (PCIe) device in a virtual environment.
- With SR-IOV:
 - SI's will get direct access to PCIe device functions
 - No more need for hypervisor (VI) to manage all system resources
- PCIe devices will have multiple virtual functions (VF's)
 - utilizable by multiple SI's
 - a single SI may also use multiple virtual functions
- Security of I/O Streams ensured by
 - Independency of control structures between VF's within one PCIe device
 - I/O address translation controls

Single Root I/O Virtualization

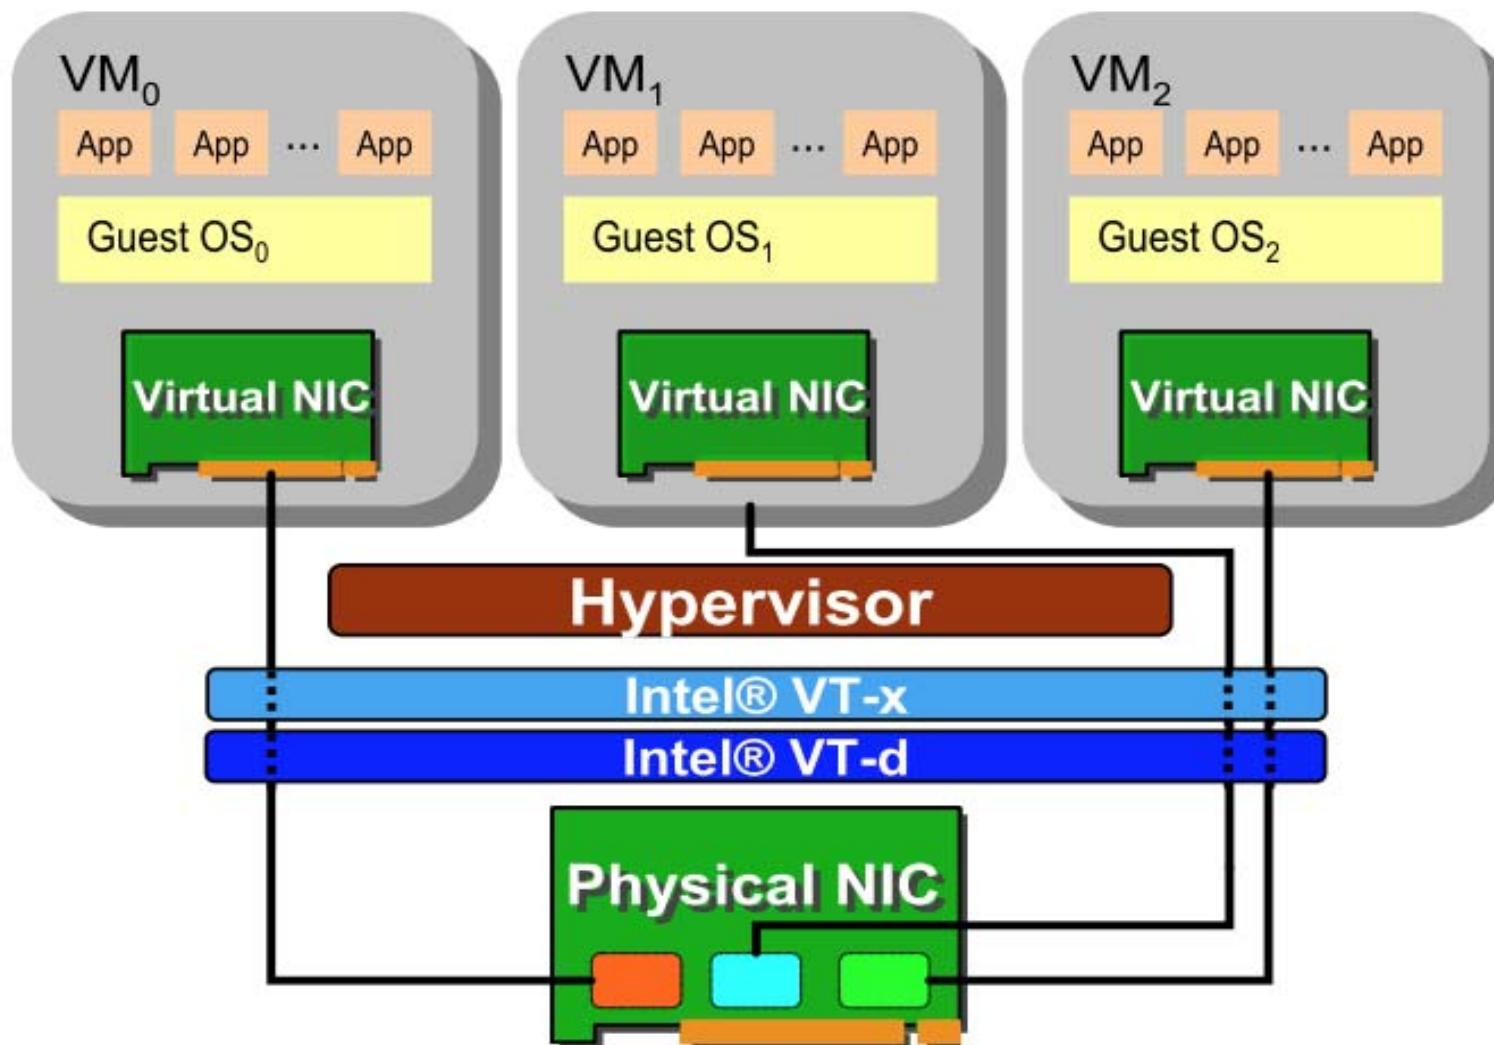


Figure 1. Natively Shared

Intel Virtualization Technology Evolution

